



<https://hao-ai-lab.github.io/dsc204a-f25/>

DSC 204A: Scalable Data Systems

Fall 2025

Staff

Instructor: Hao Zhang

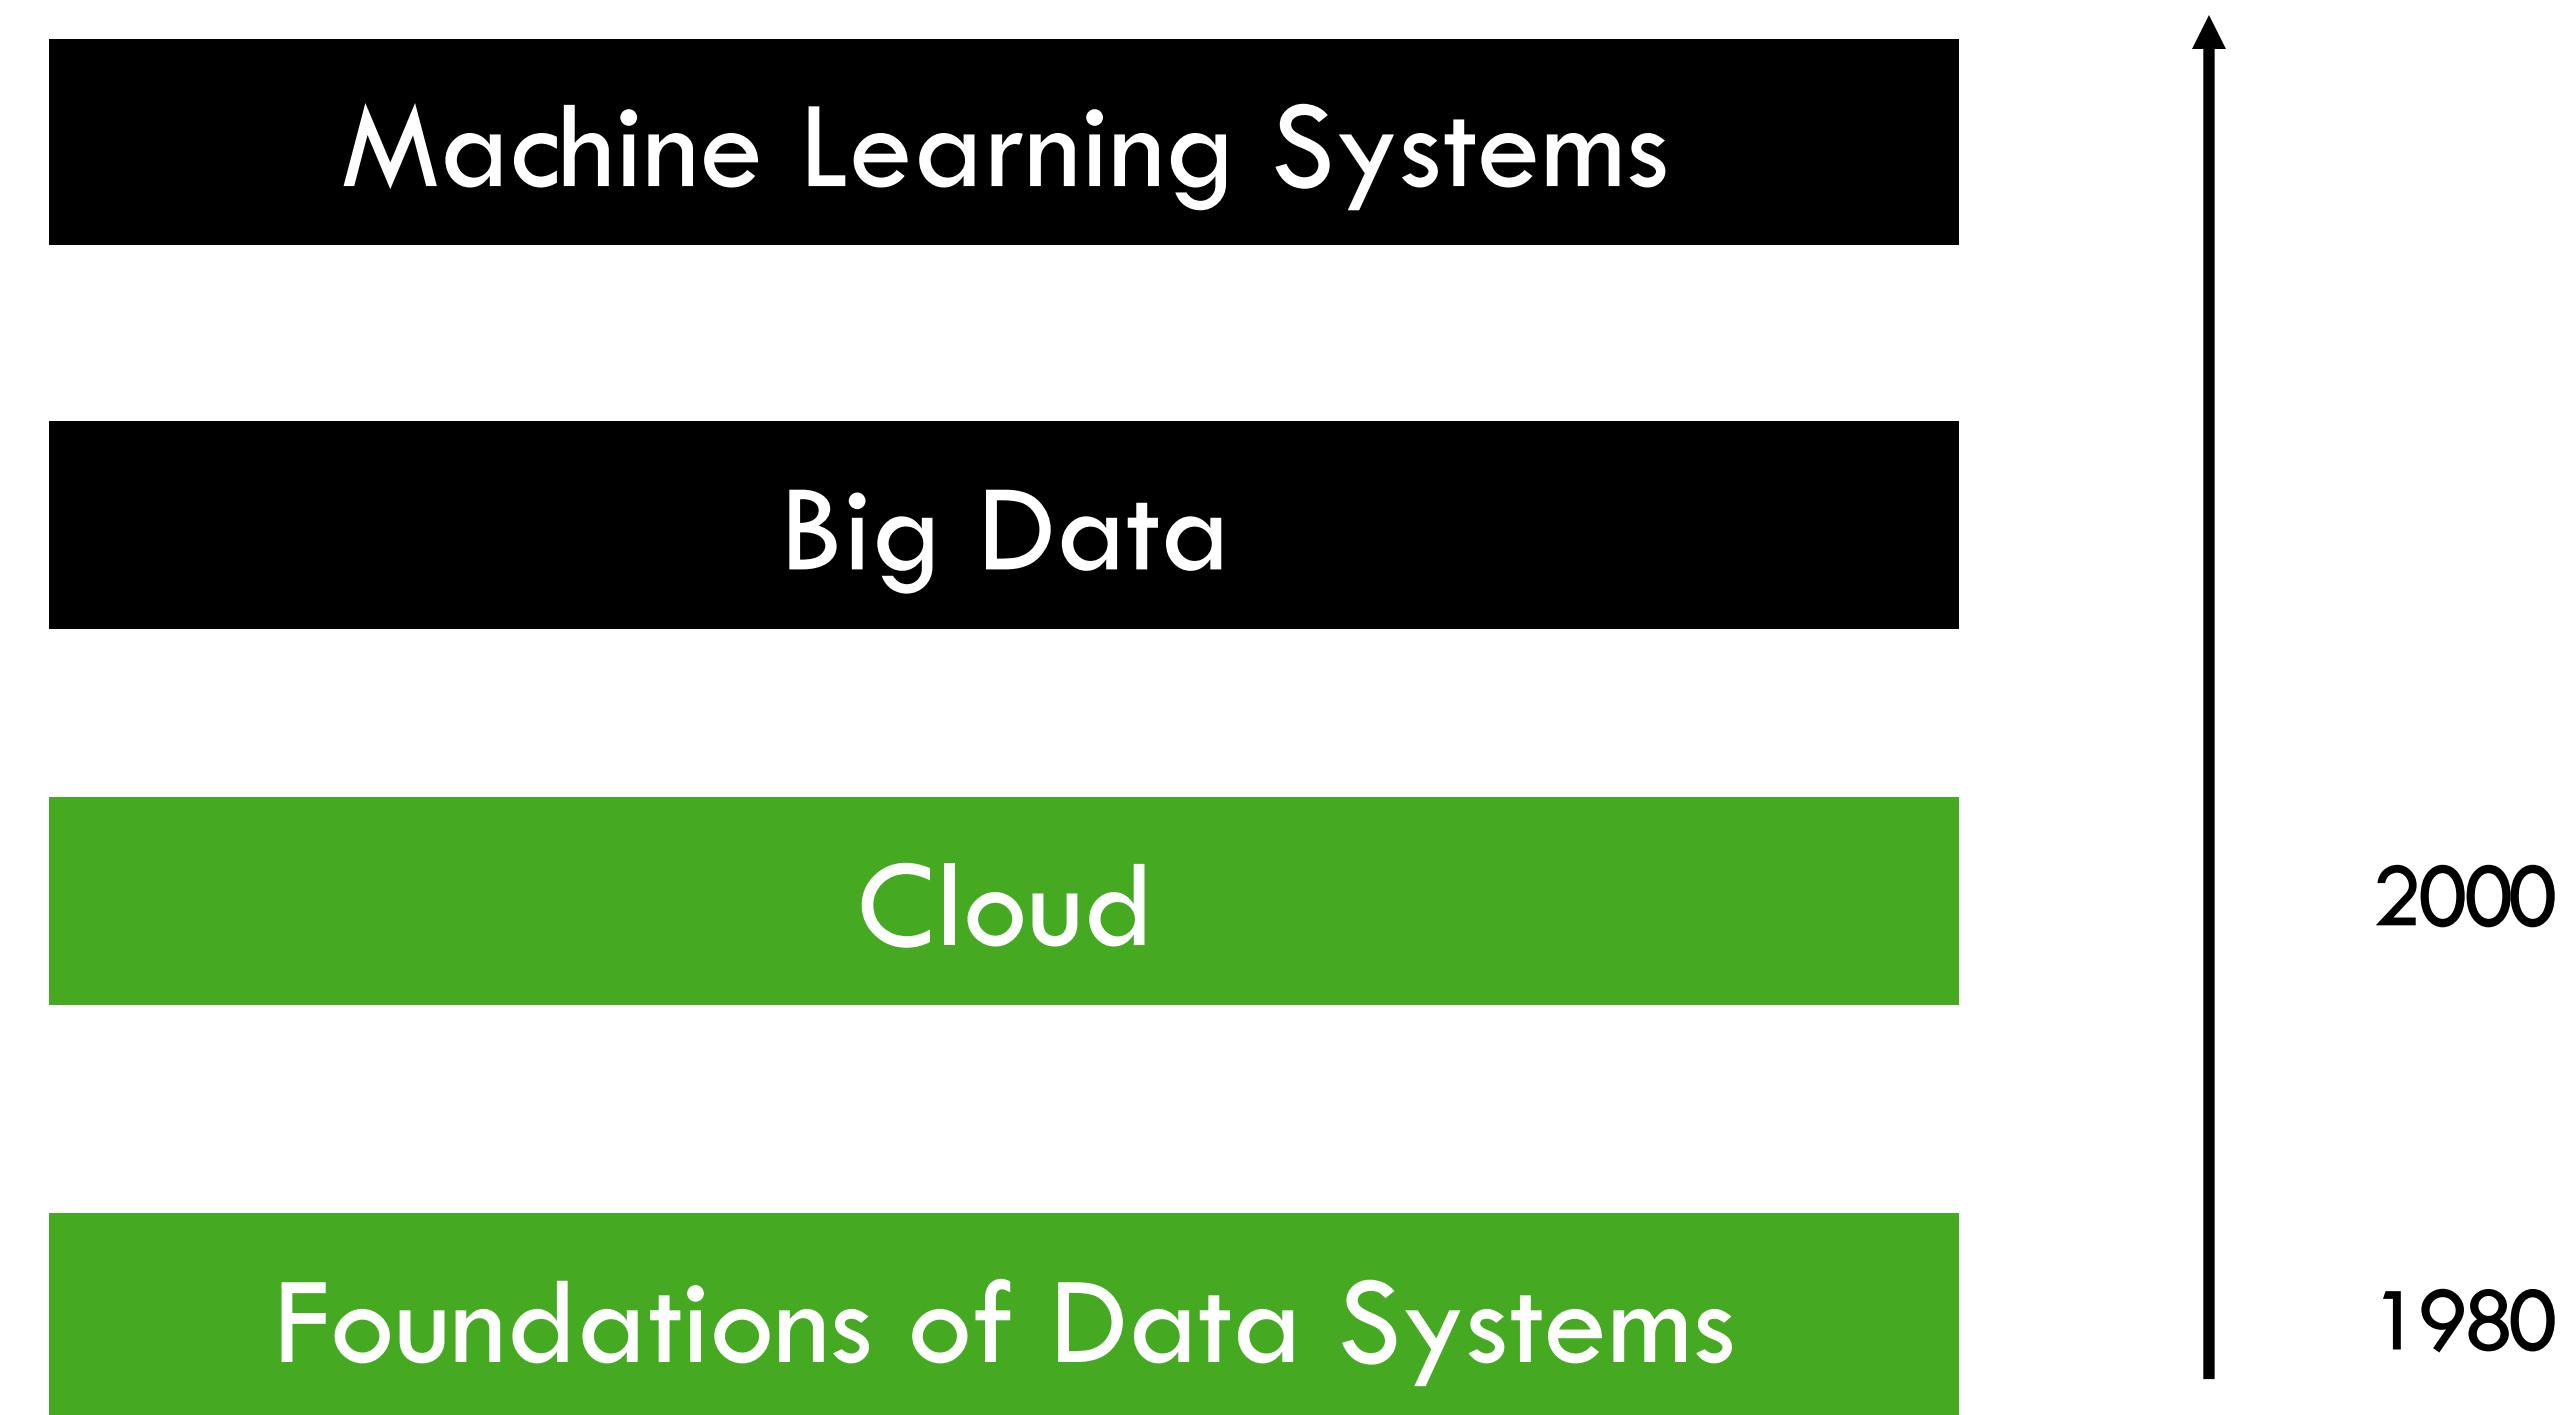
TAs: Mingjia Huo, Yuxuan Zhang

 [@haozhangml](https://twitter.com/haozhangml)

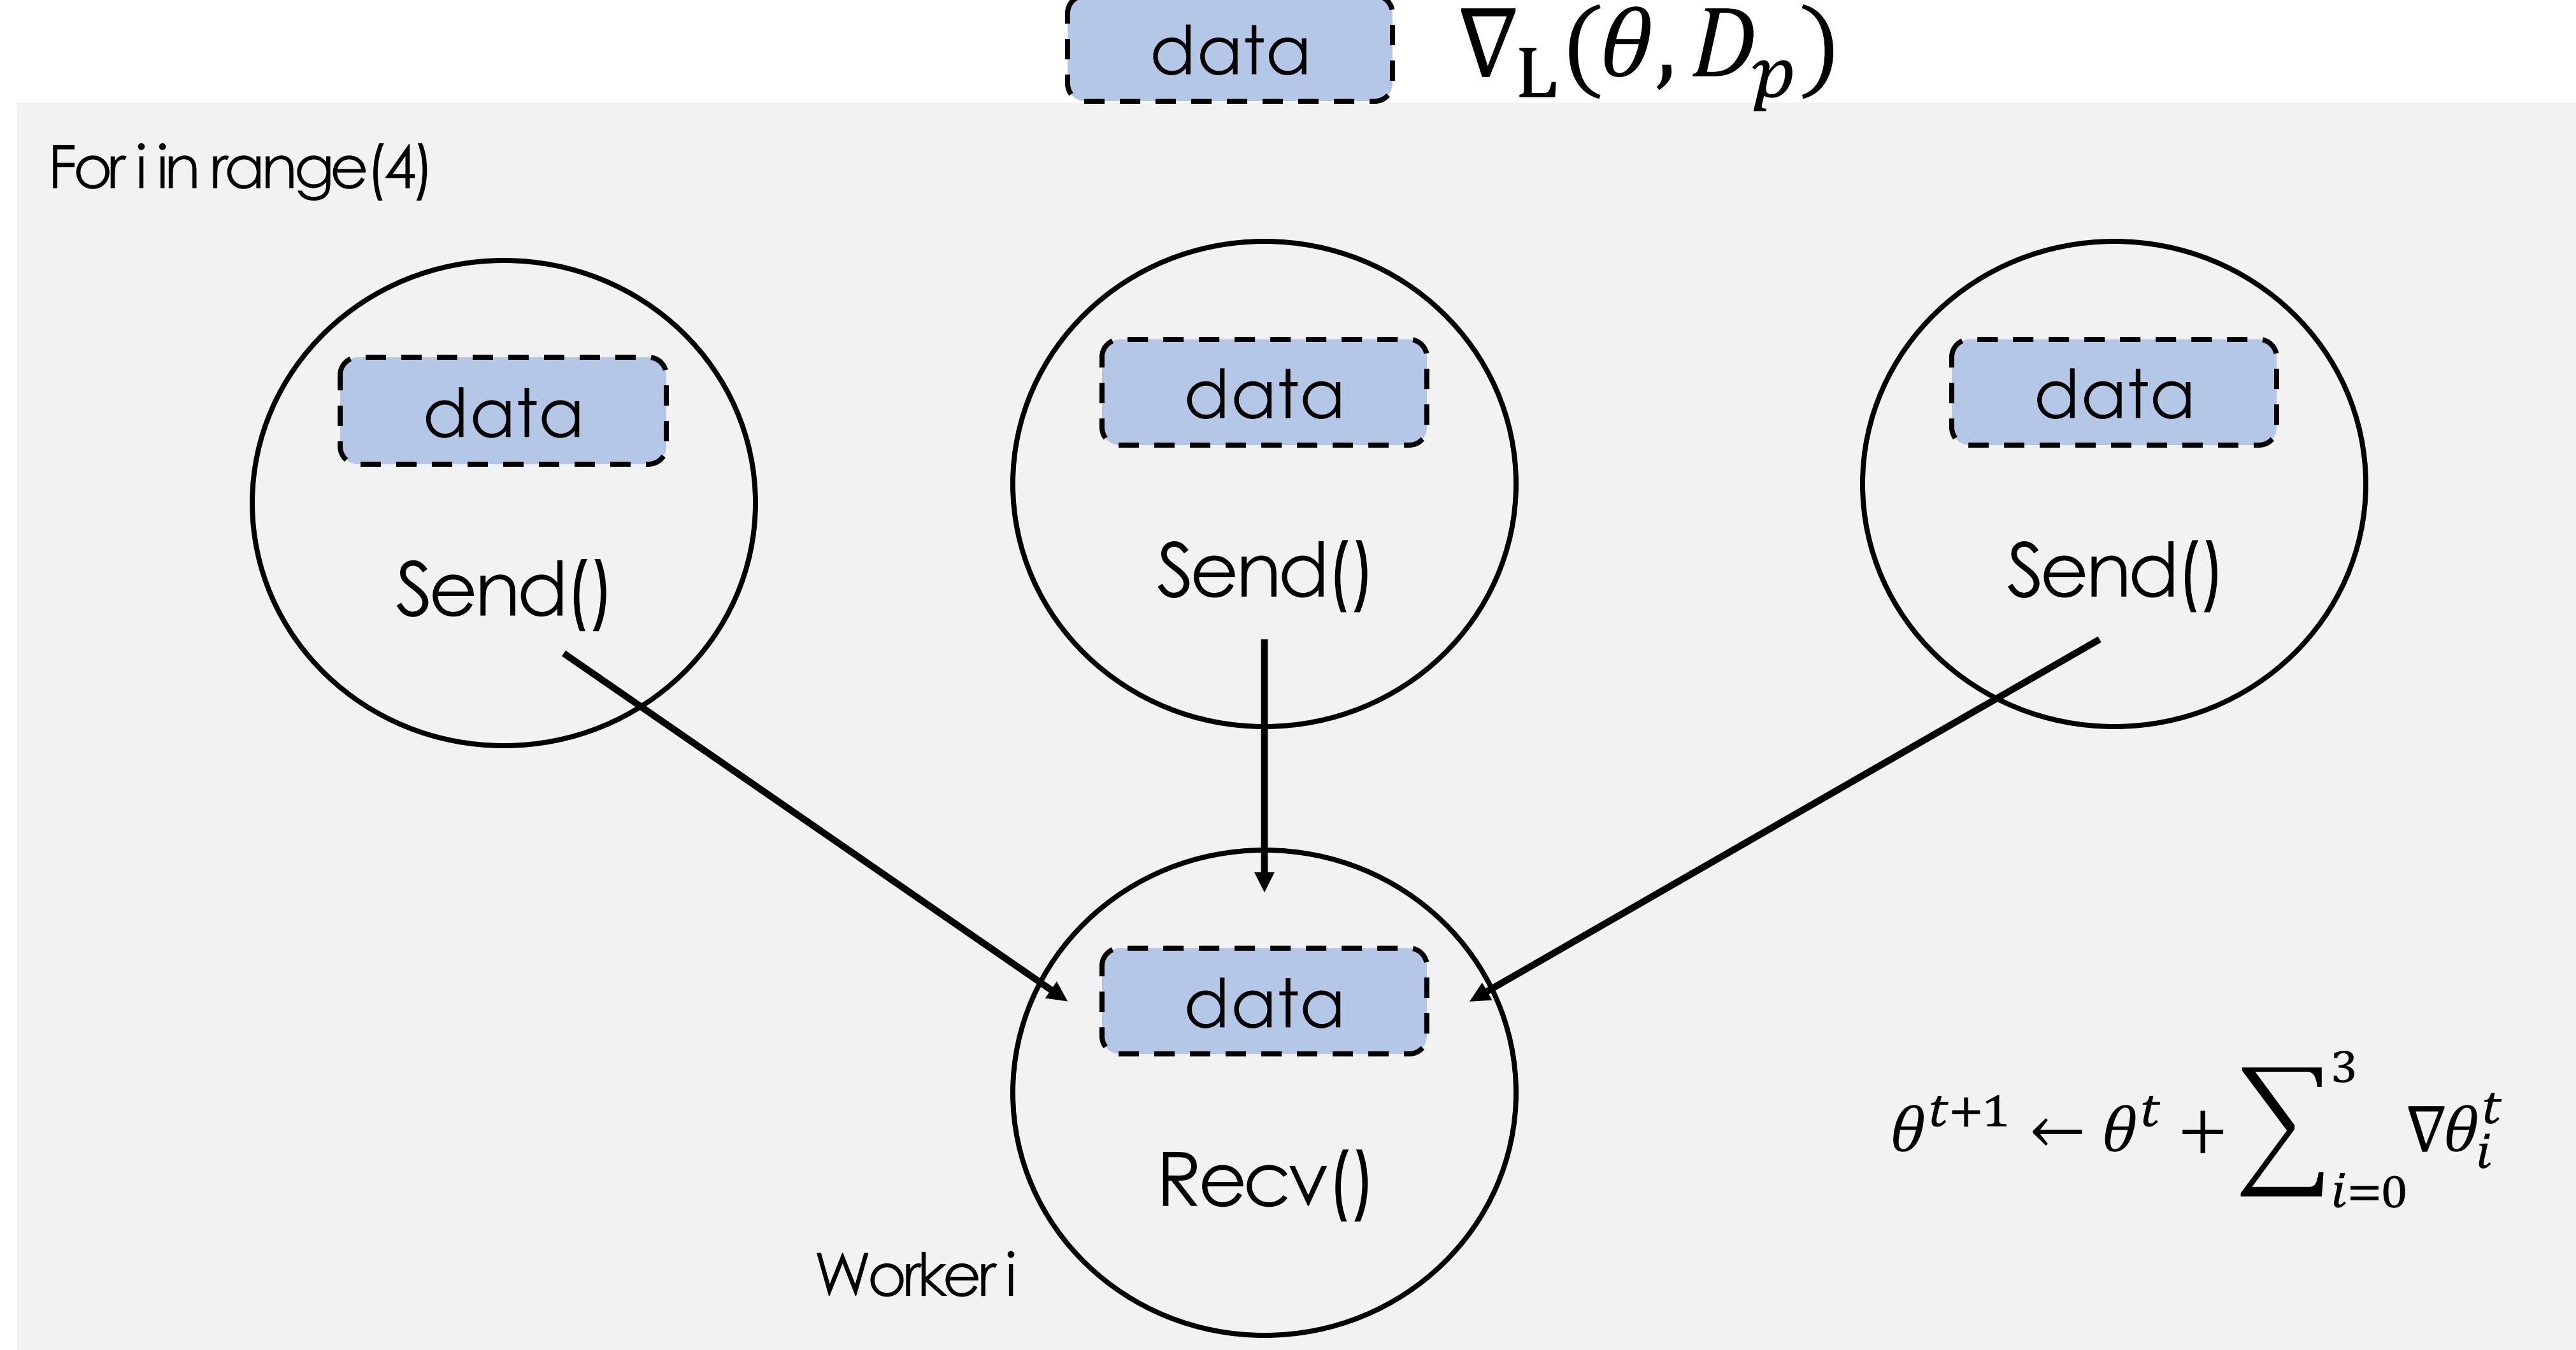
 [@haoailab](https://twitter.com/haoailab)

 haozhang@ucsd.edu

Where We Are



Problem: We need All-Reduce



Program This? Will be in PA2!

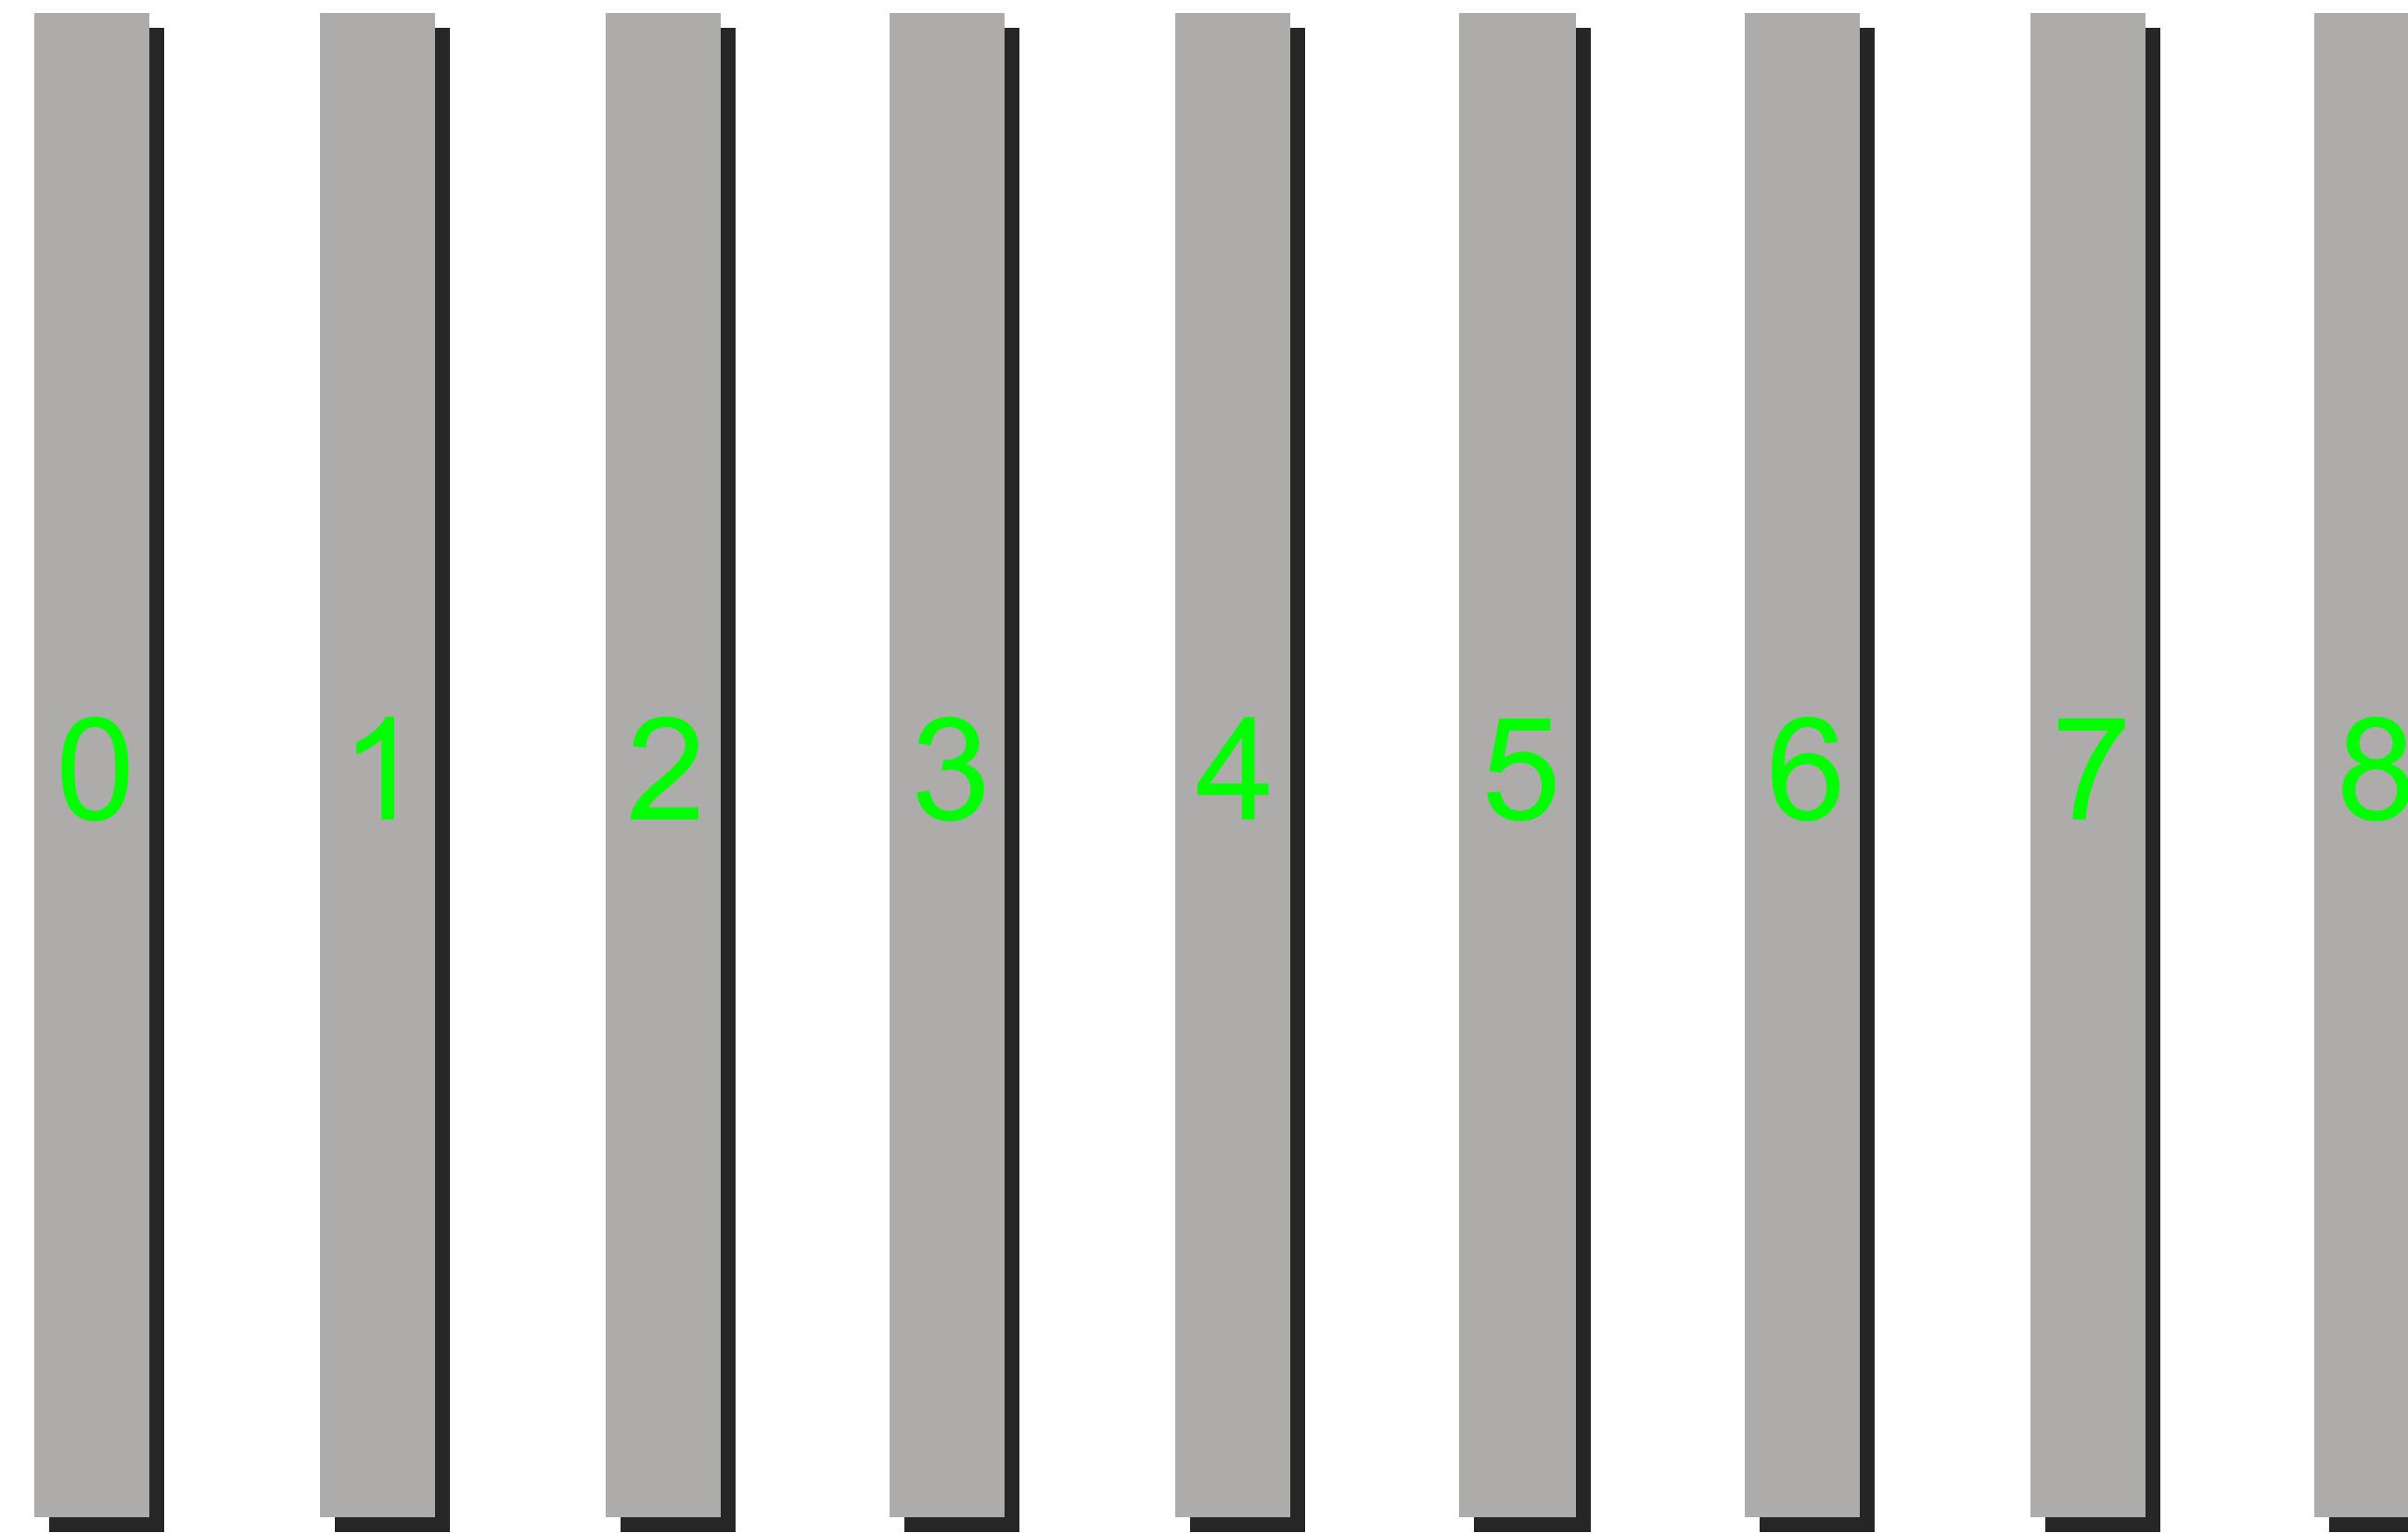
Performance

- Message size over networks:
 - Sum: $3N$
 - Send Sum back: $3N$
 - $= 6N$
- Can we do better?
 - Hint: we cannot do better than $3N$

Why Collective Communication?

- Programming Convenience
 - Use a set of well-defined communication primitives to express complex communication patterns
- Unification and Performance
 - Since they are well defined and well structured, we can optimize them to the extreme
- ML Systems ❤️ Collective communication

Make it Formal

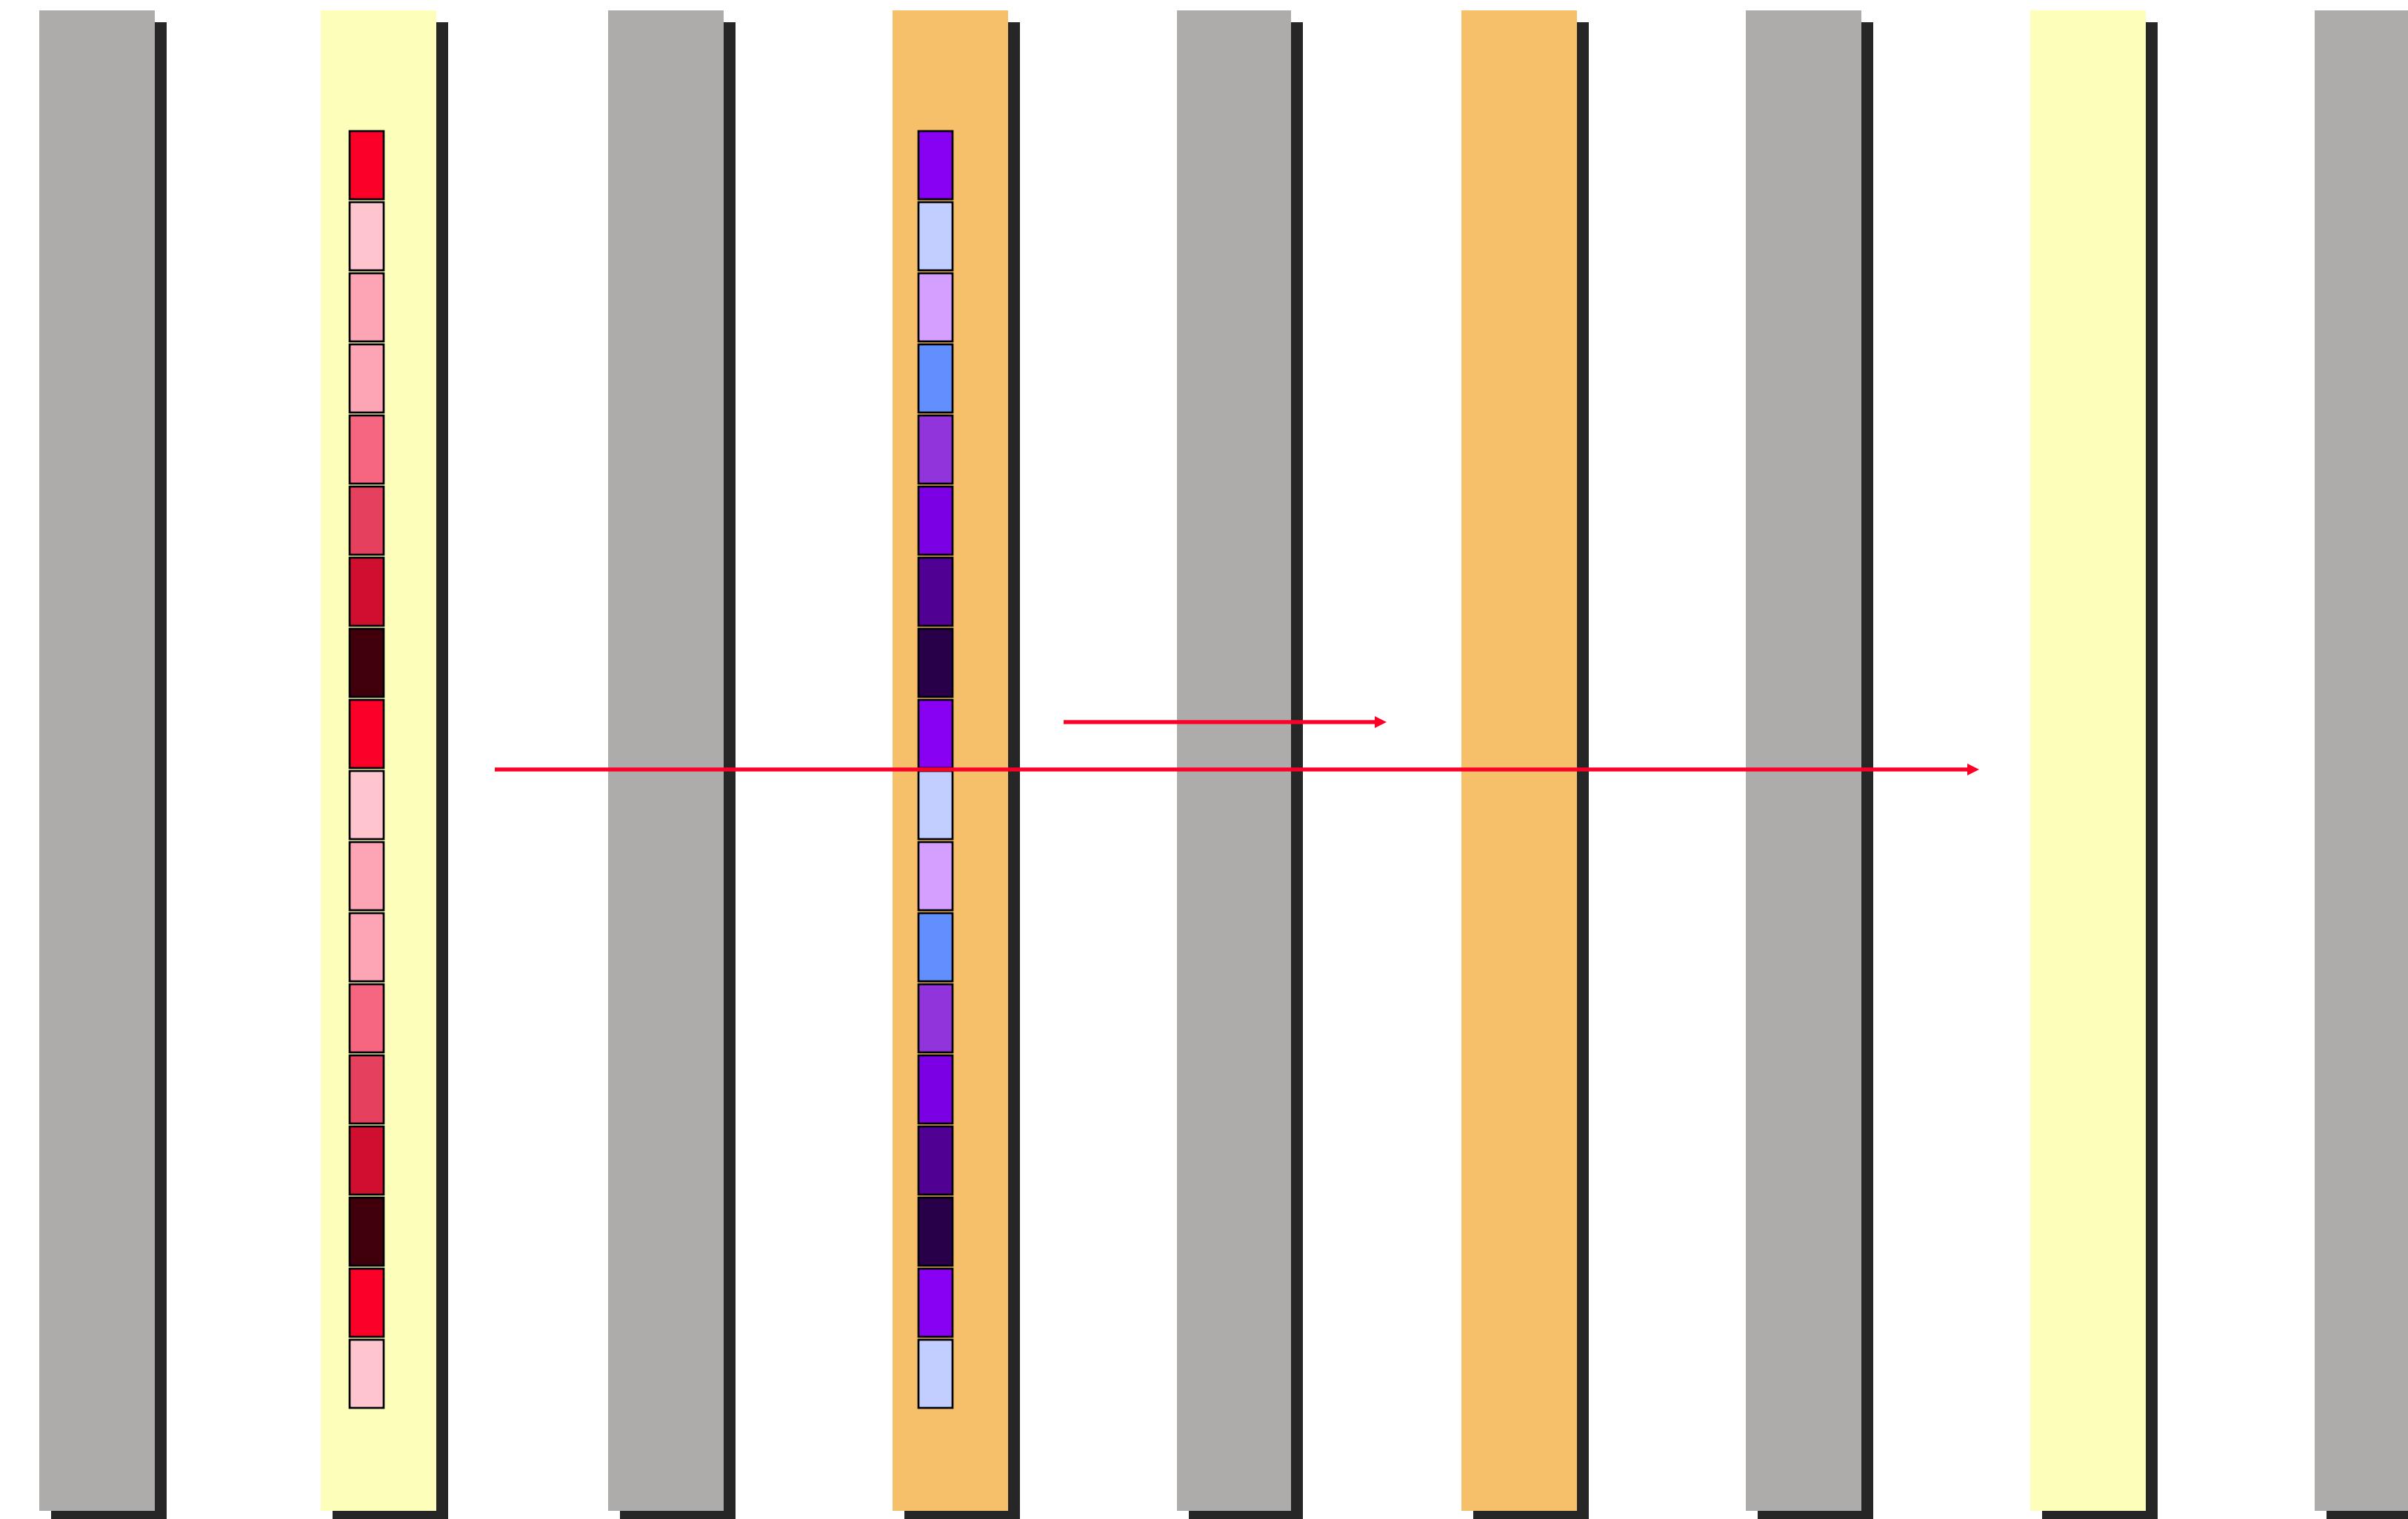


- A 1D **Mesh** of workers (or devices, or nodes)

Model of Parallel Computation

- a node can send directly to any other node (maybe not true)
- a node can simultaneously receive and send
- cost of communication
 - sending a message of length n between any two nodes

$$\alpha + n \beta$$

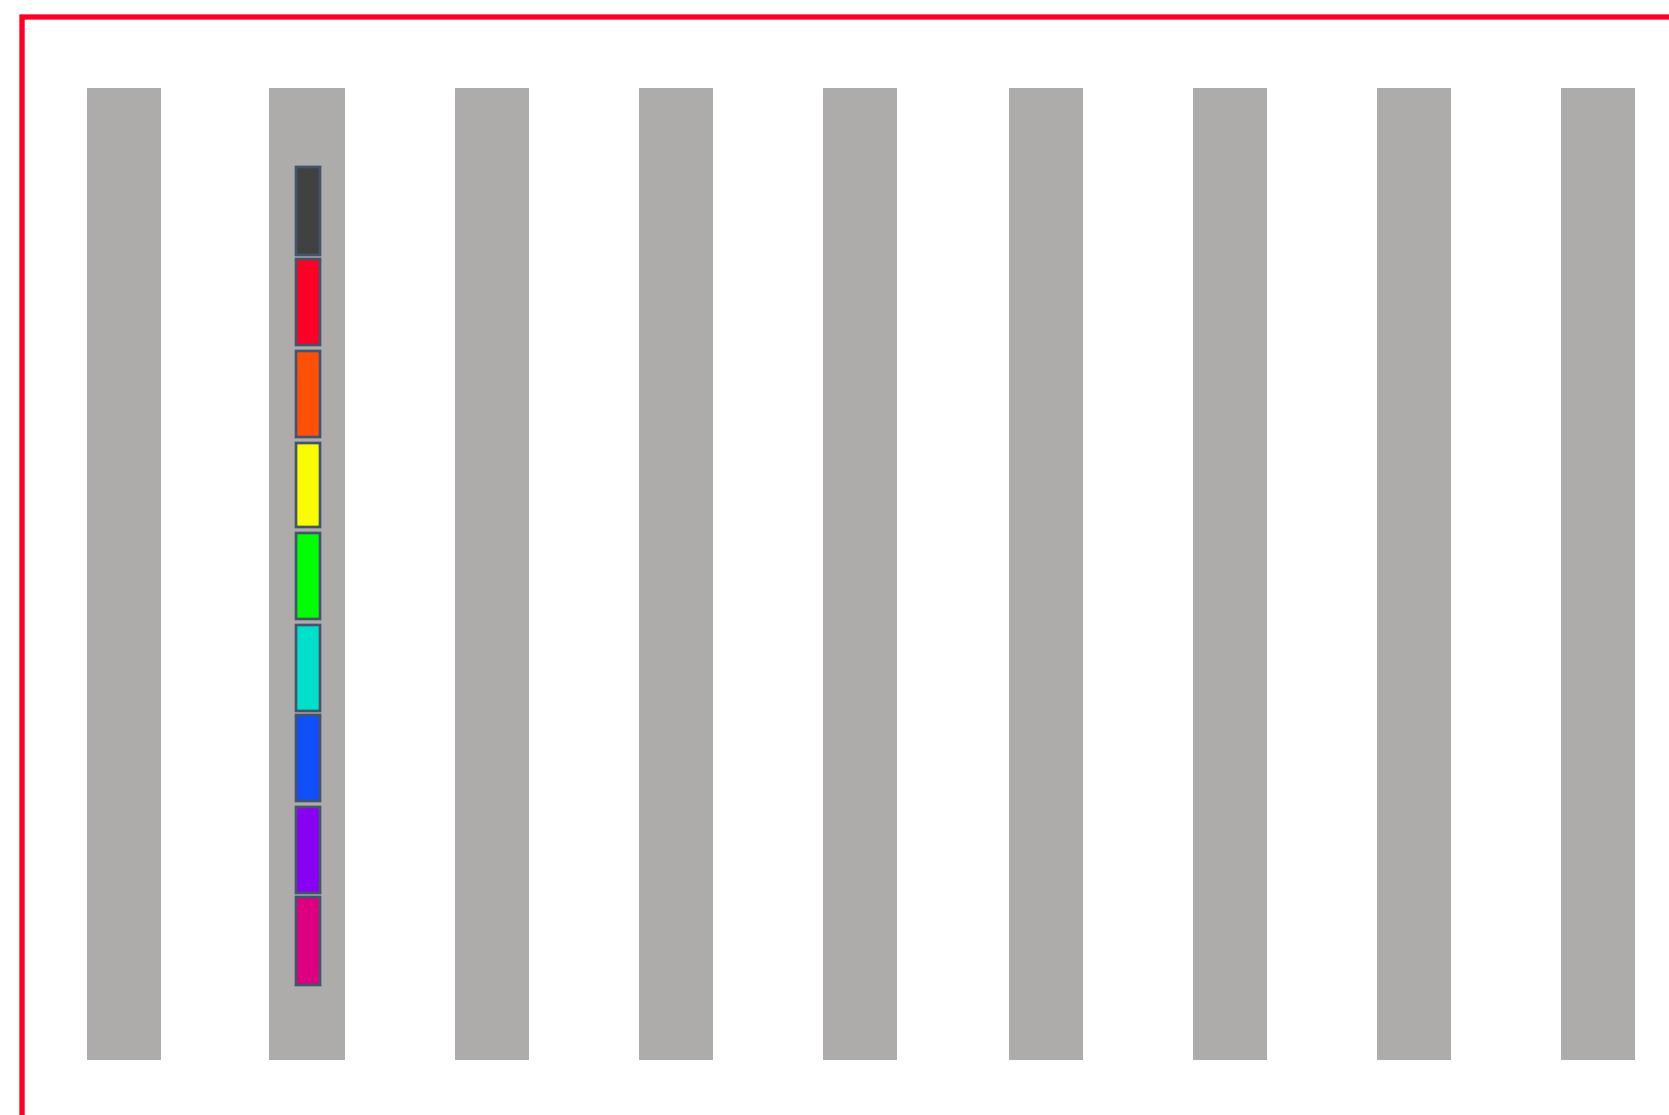


Collective Communications

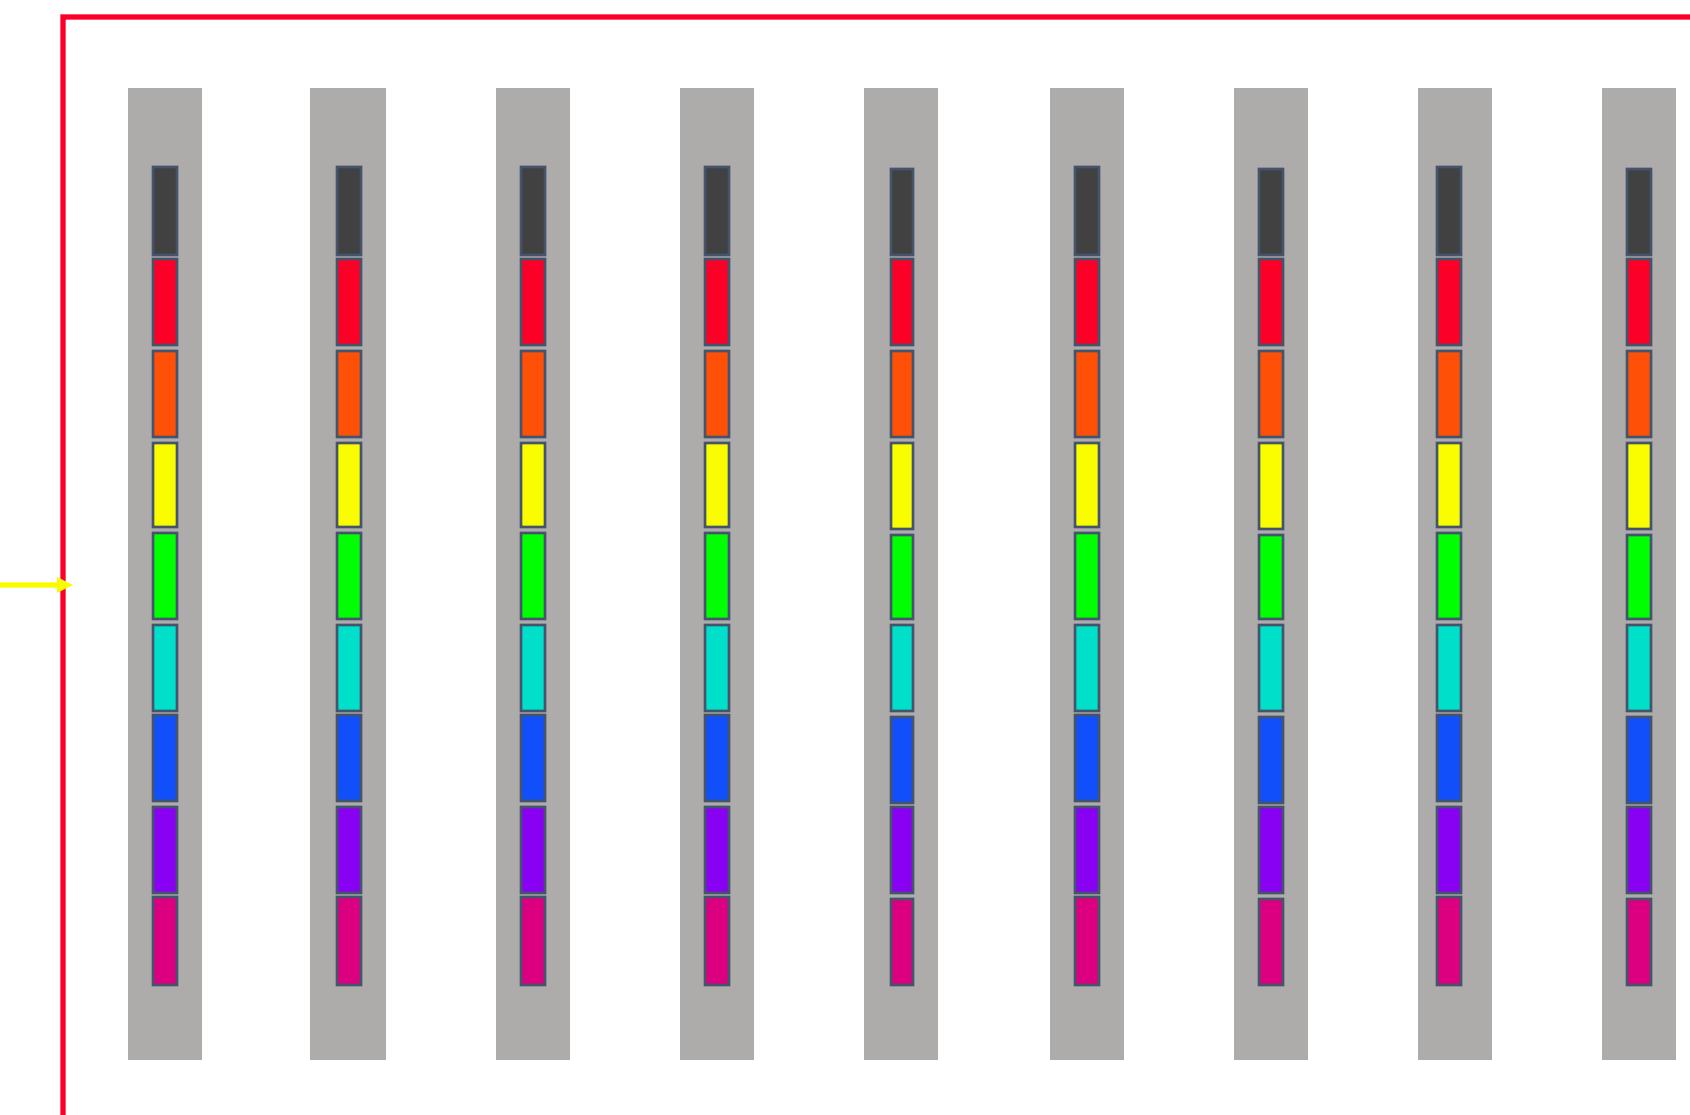
- Broadcast
- Reduce(-to-one)
- Scatter
- Gather
- Allgather
- Reduce-scatter
- Allreduce
- All-2-All

Broadcast

Before

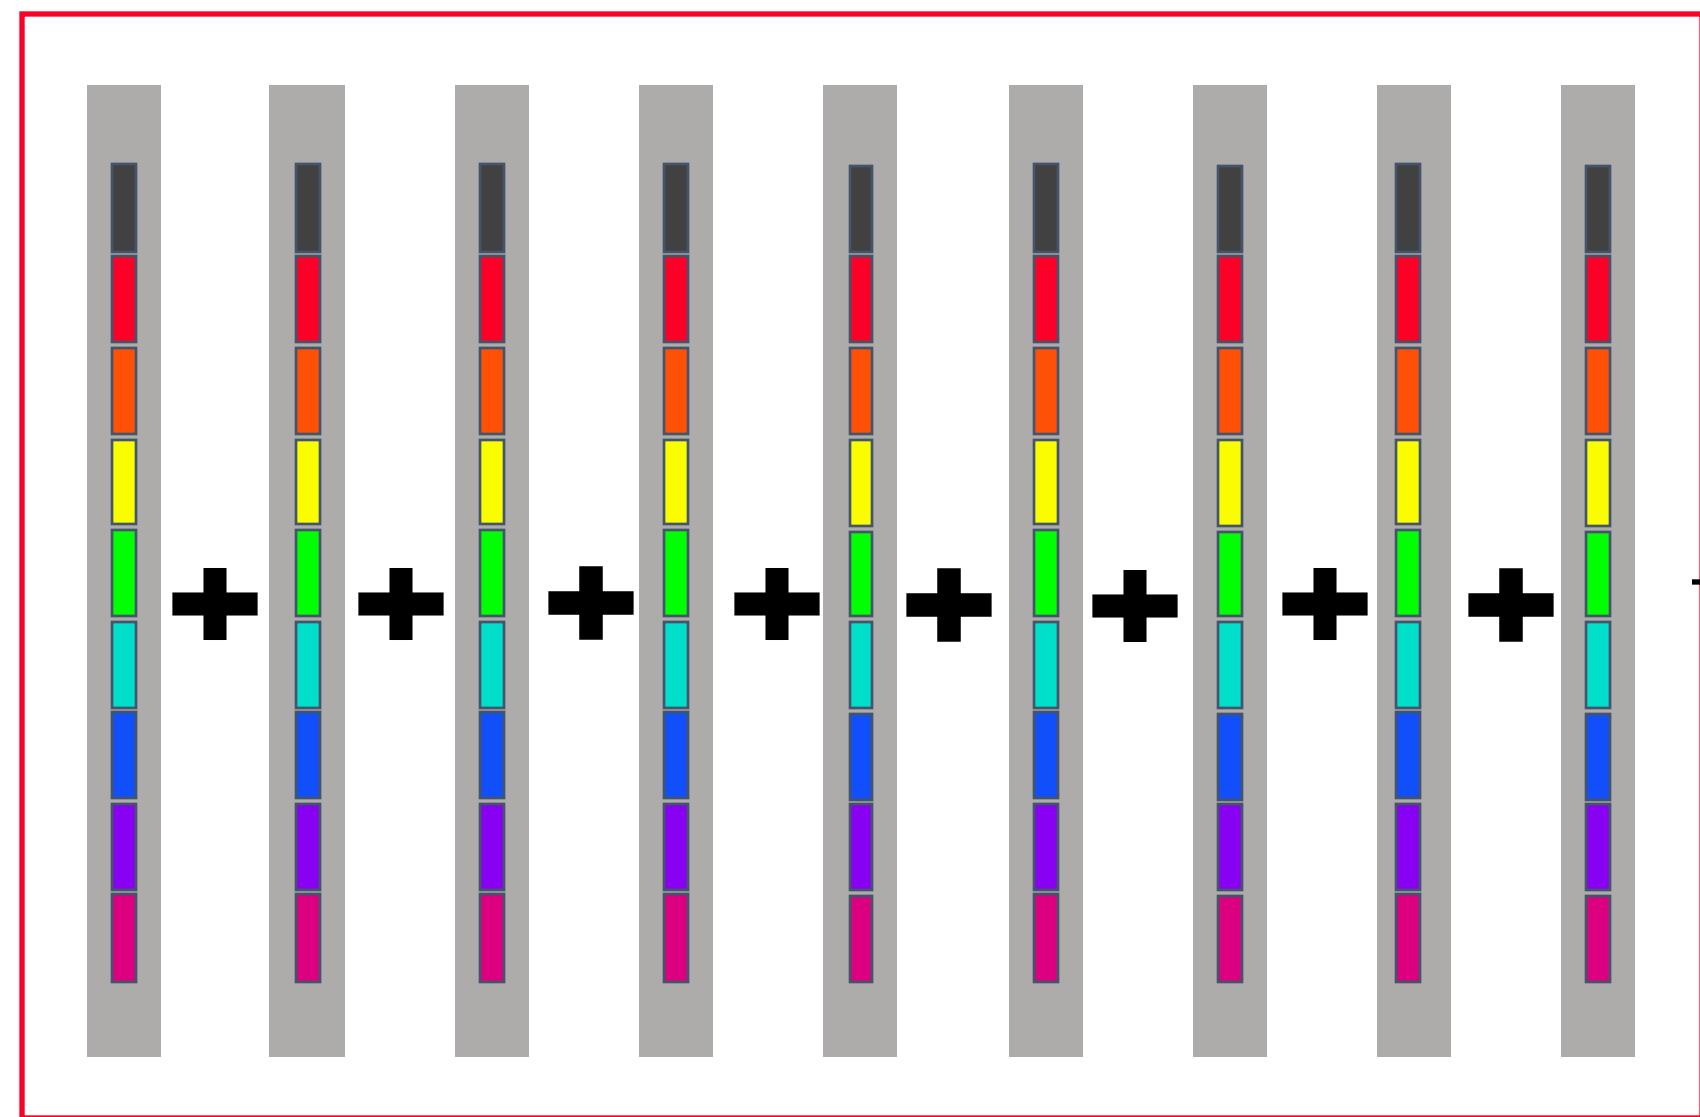


After

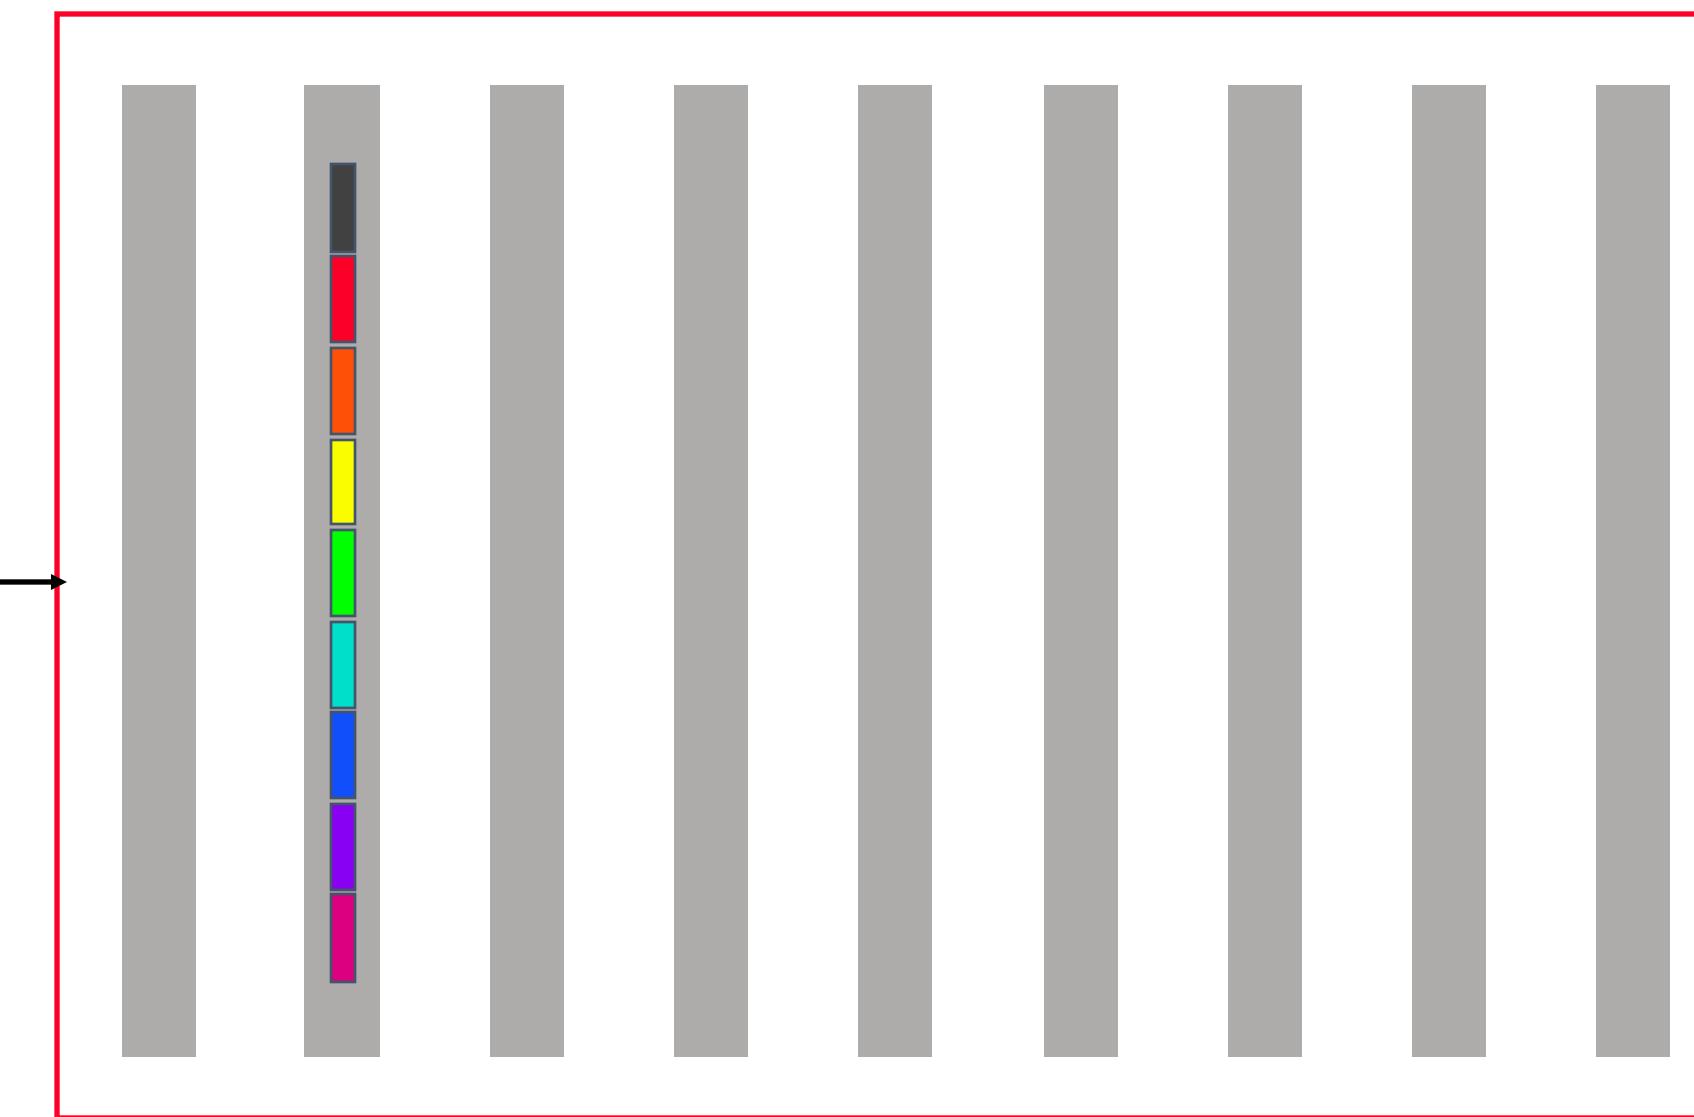


Reduce(-to-one)

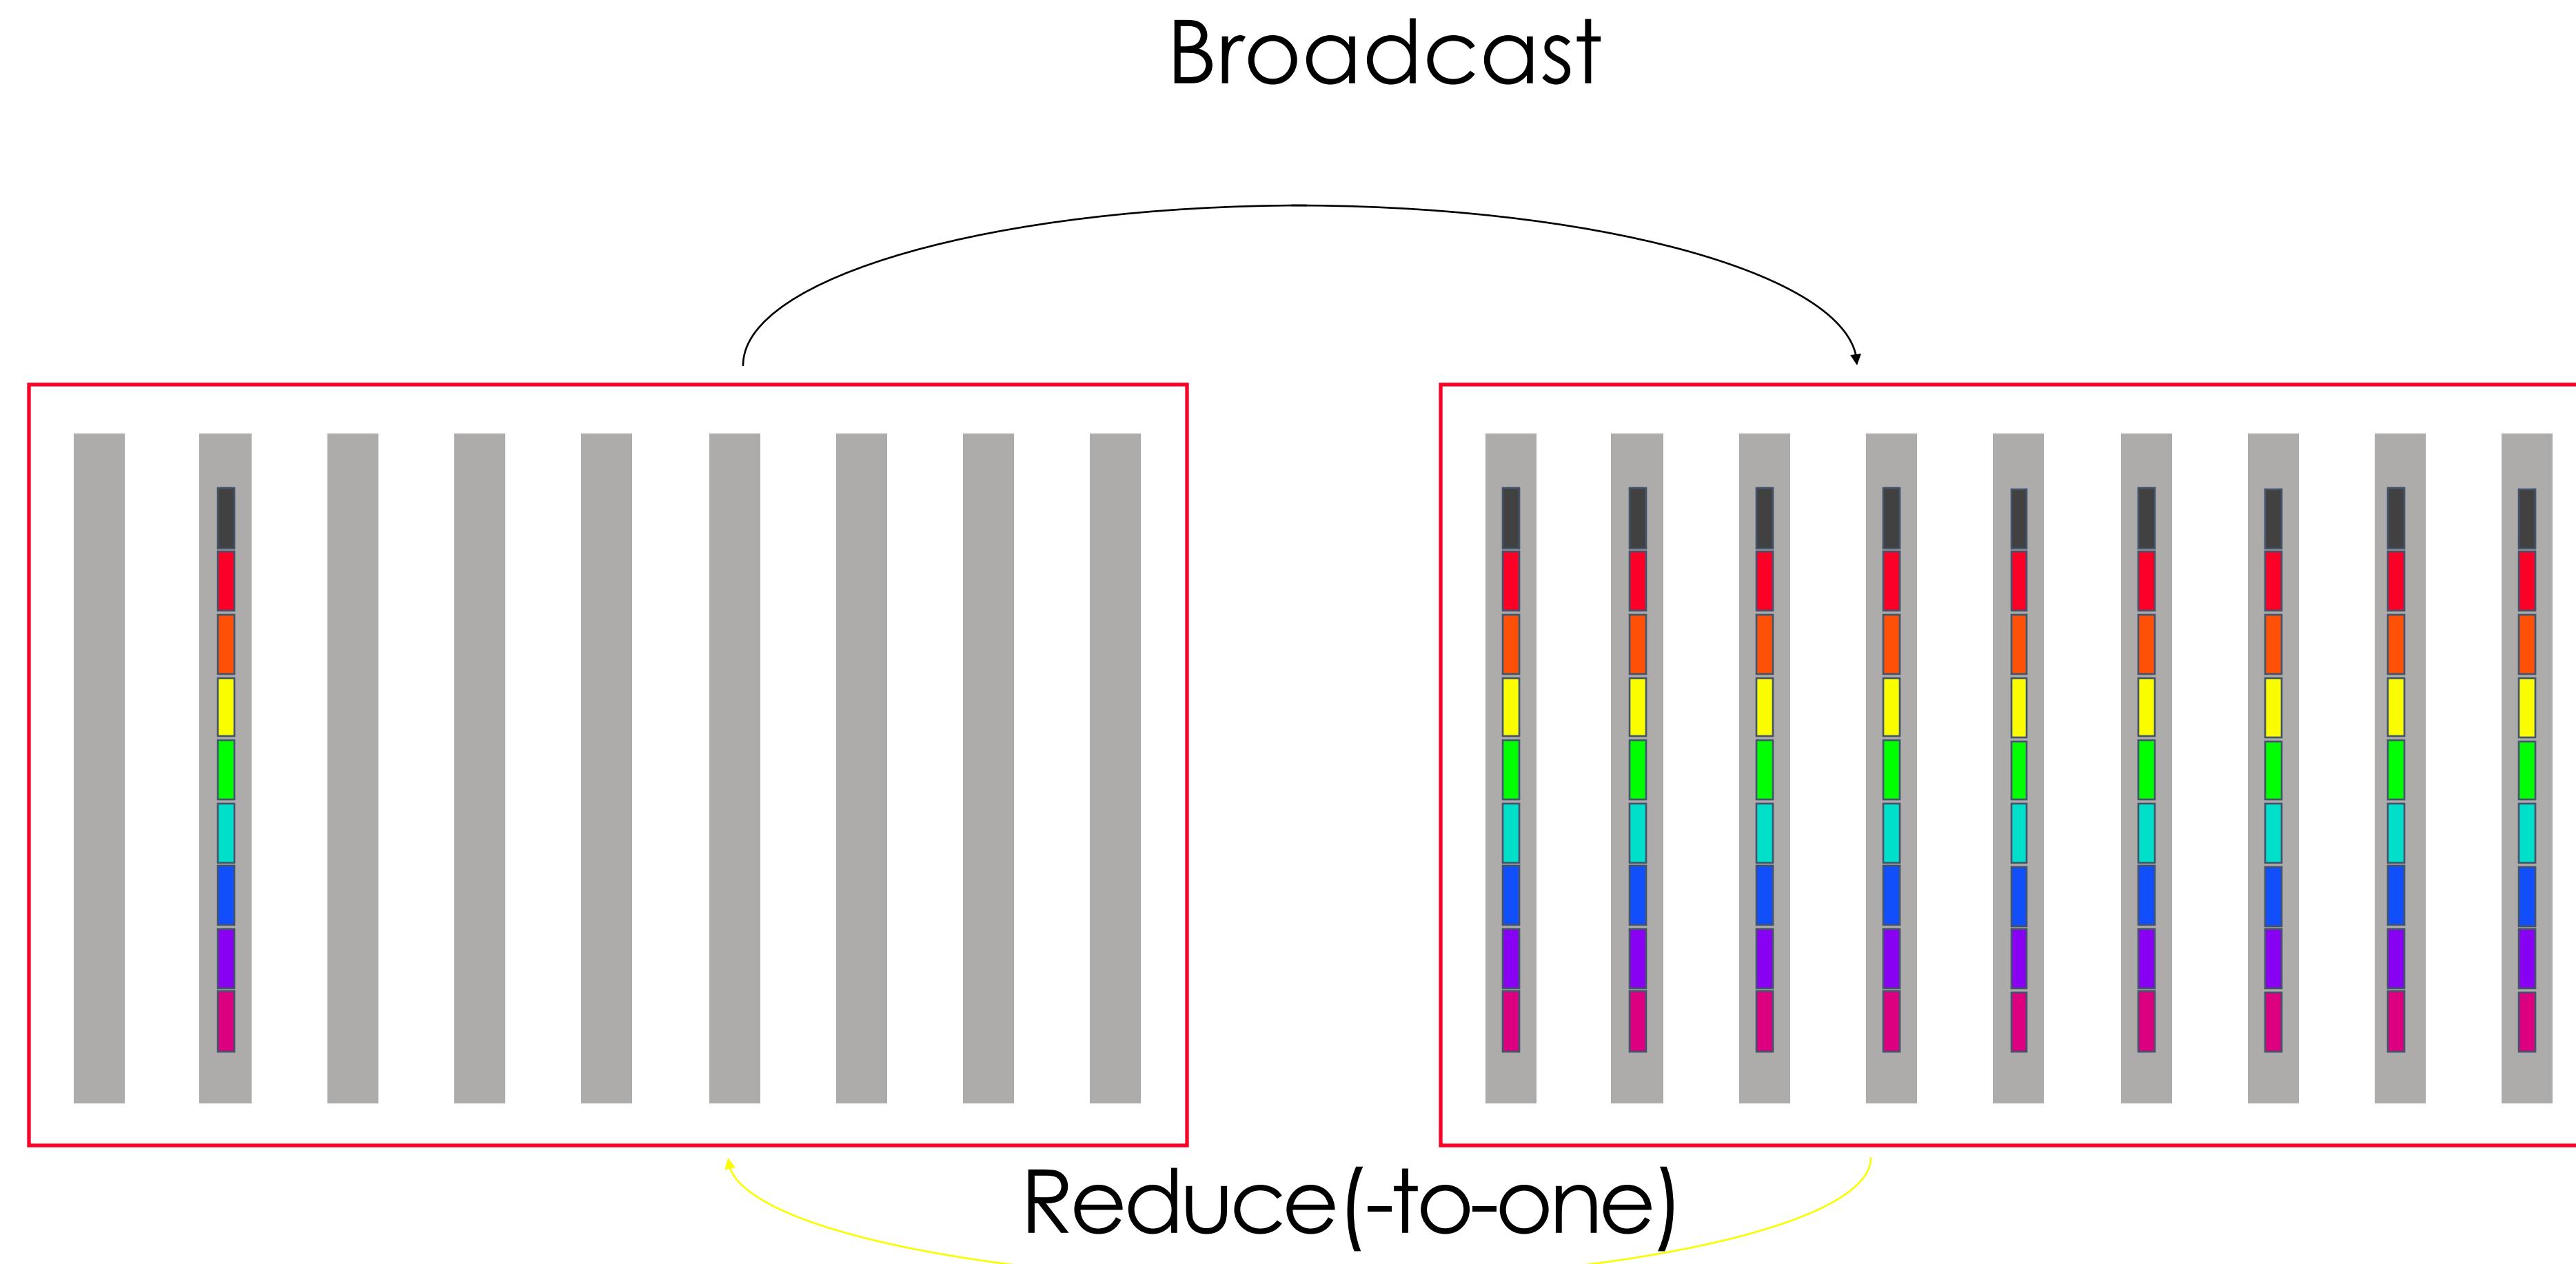
Before



After

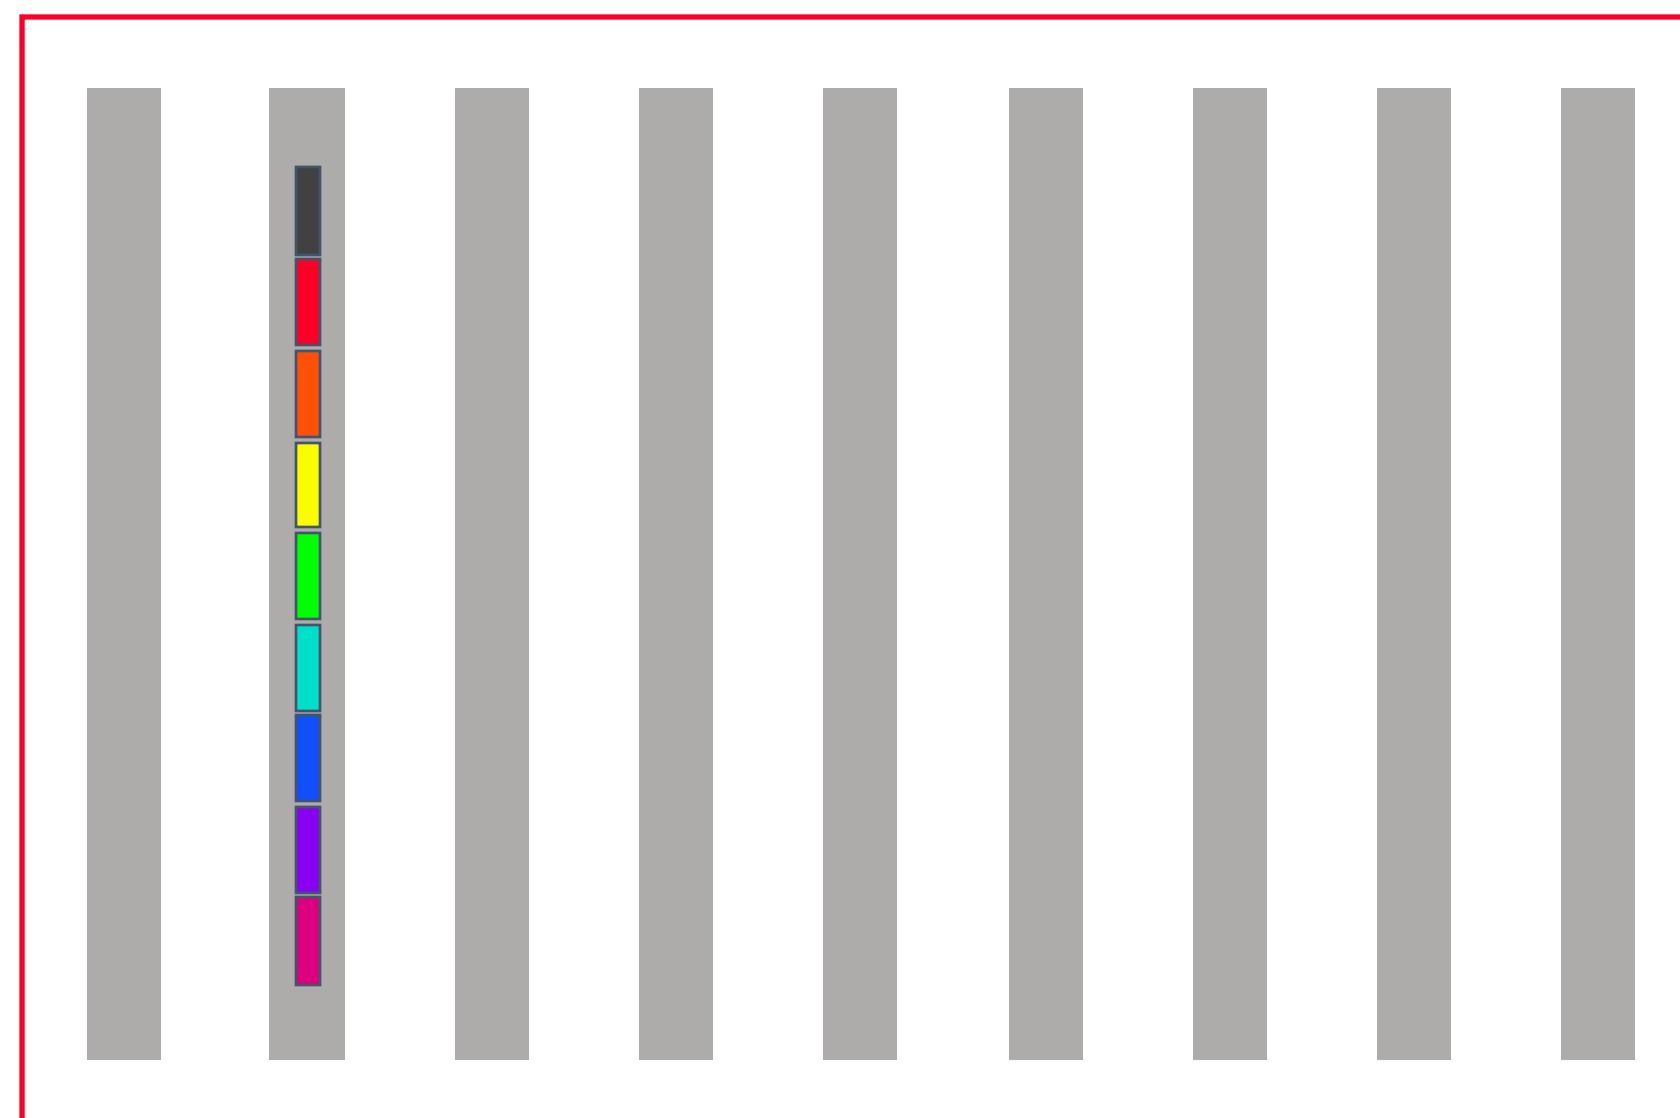


Broadcast/Reduce(-to-one)

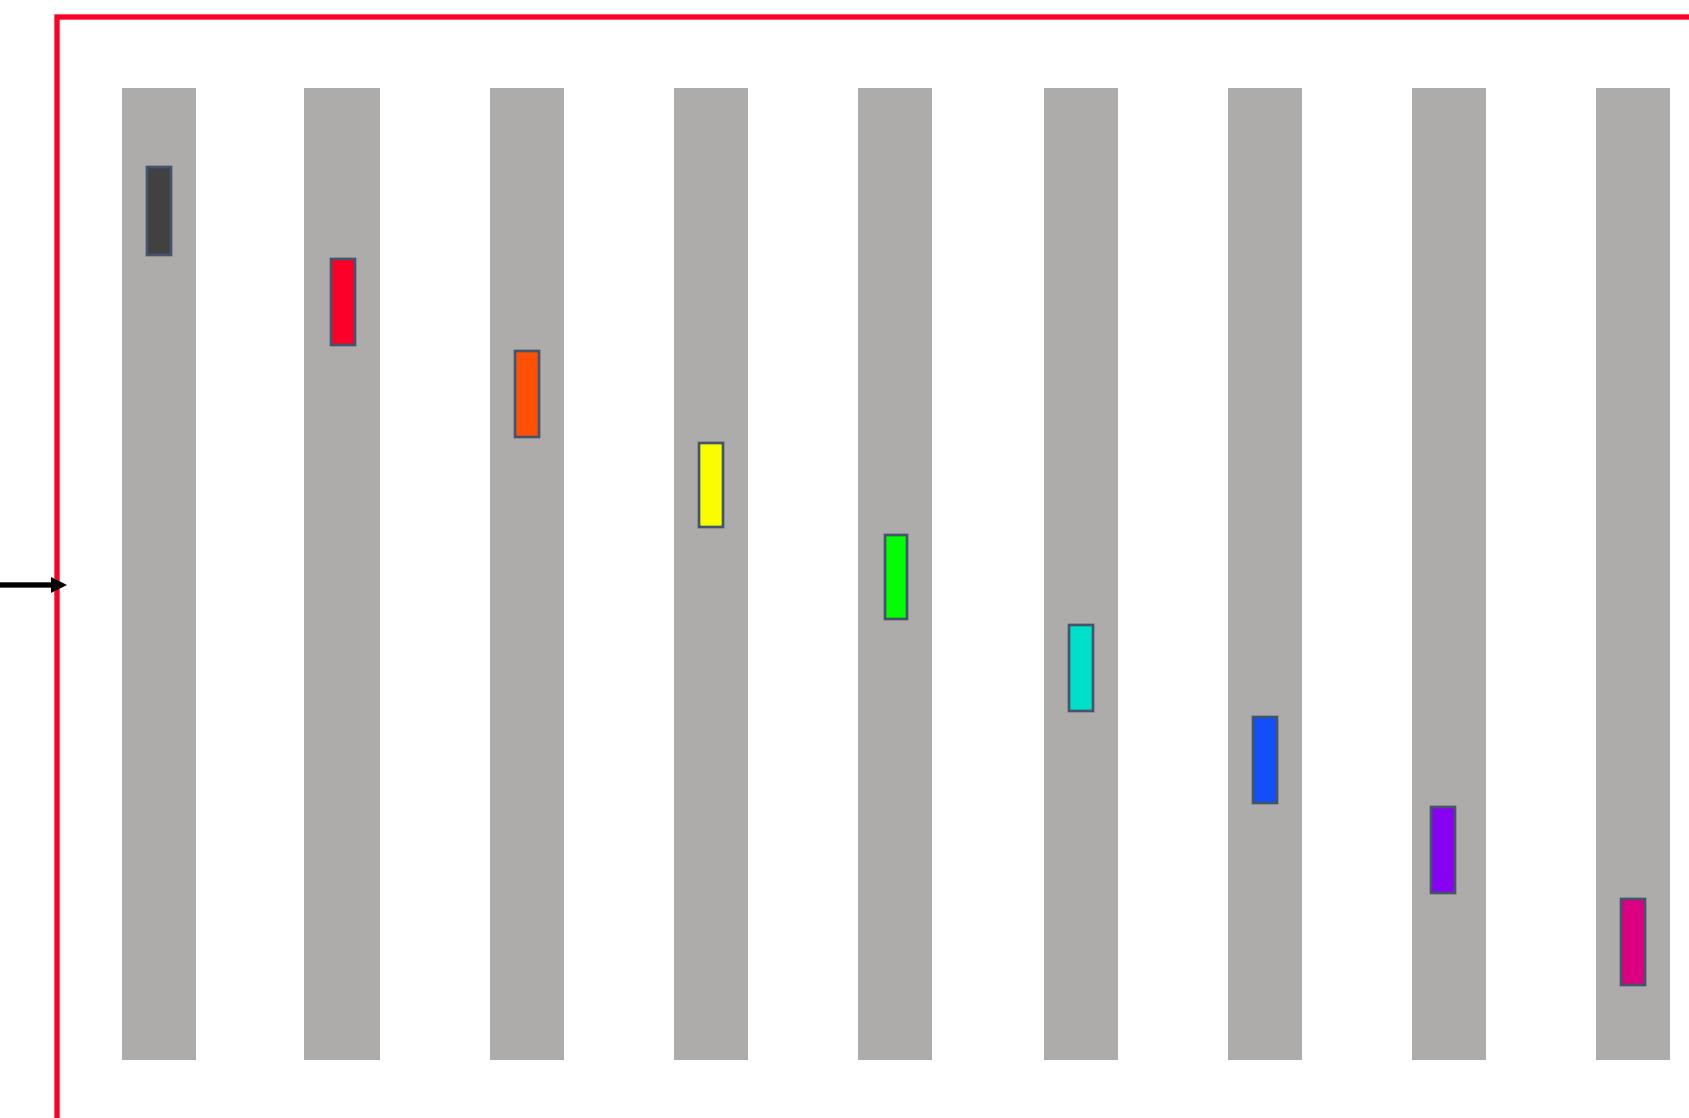


Scatter

Before

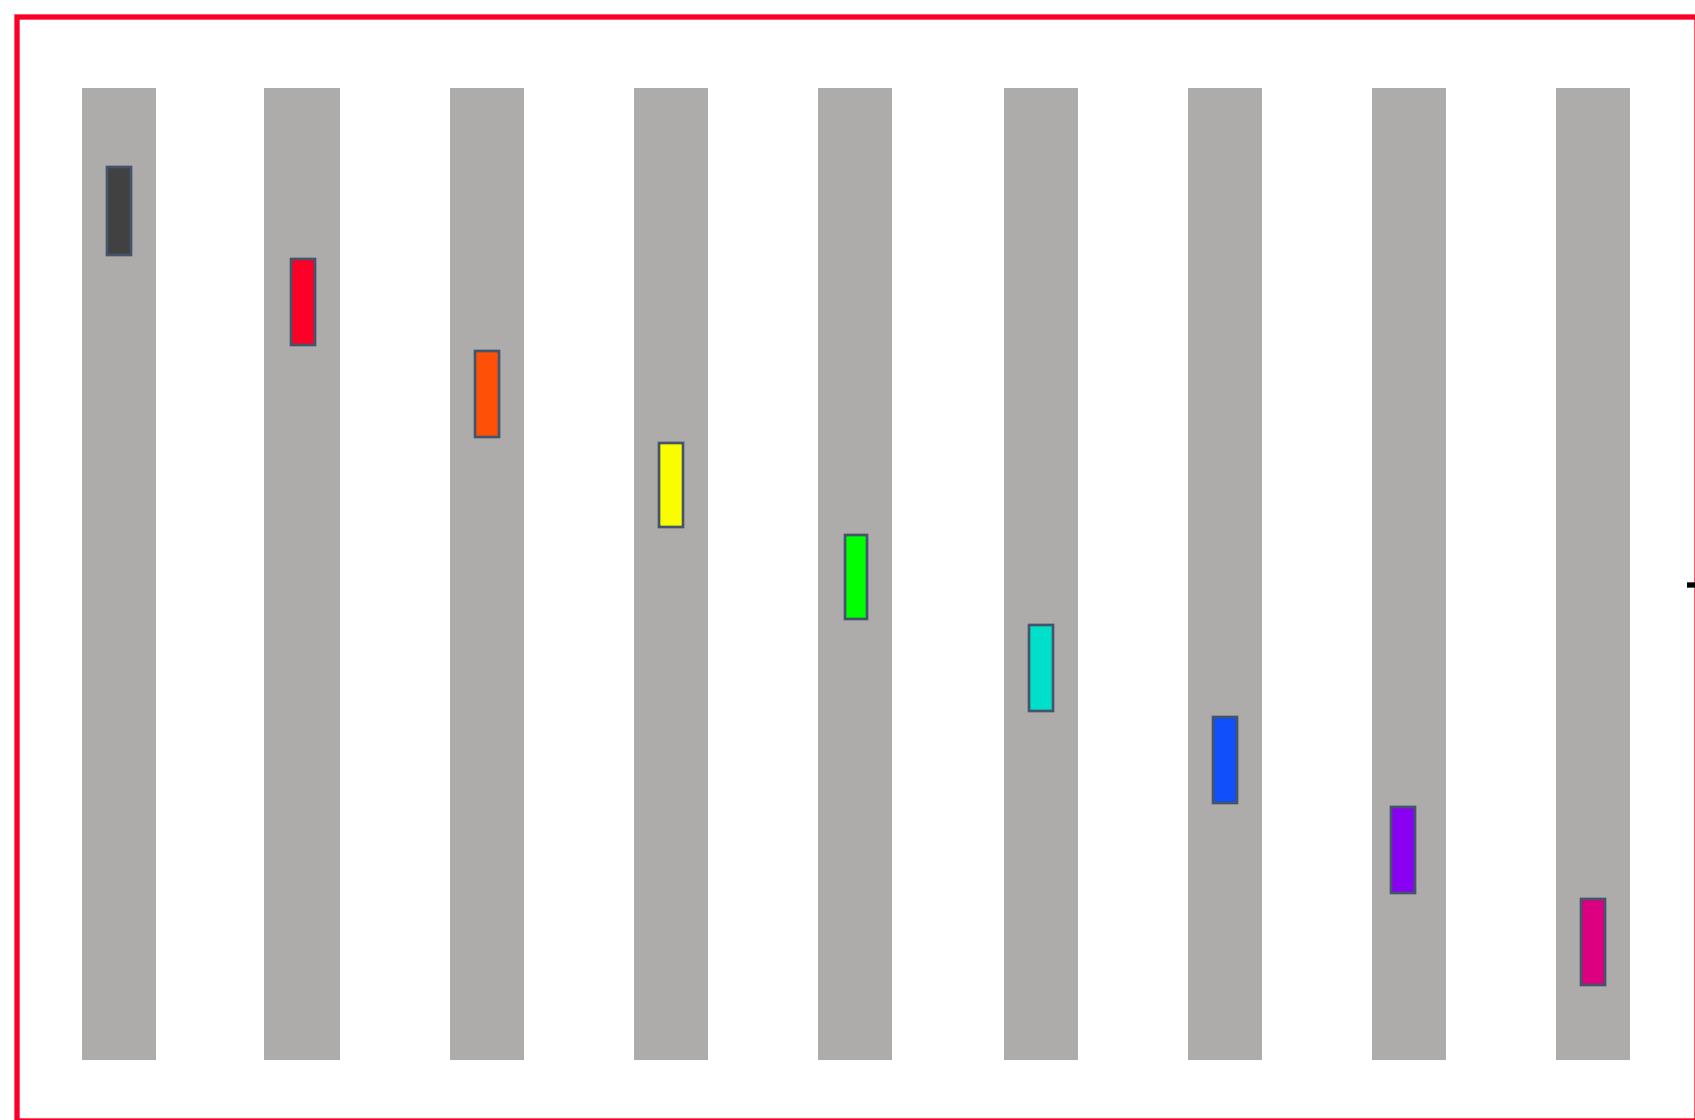


After

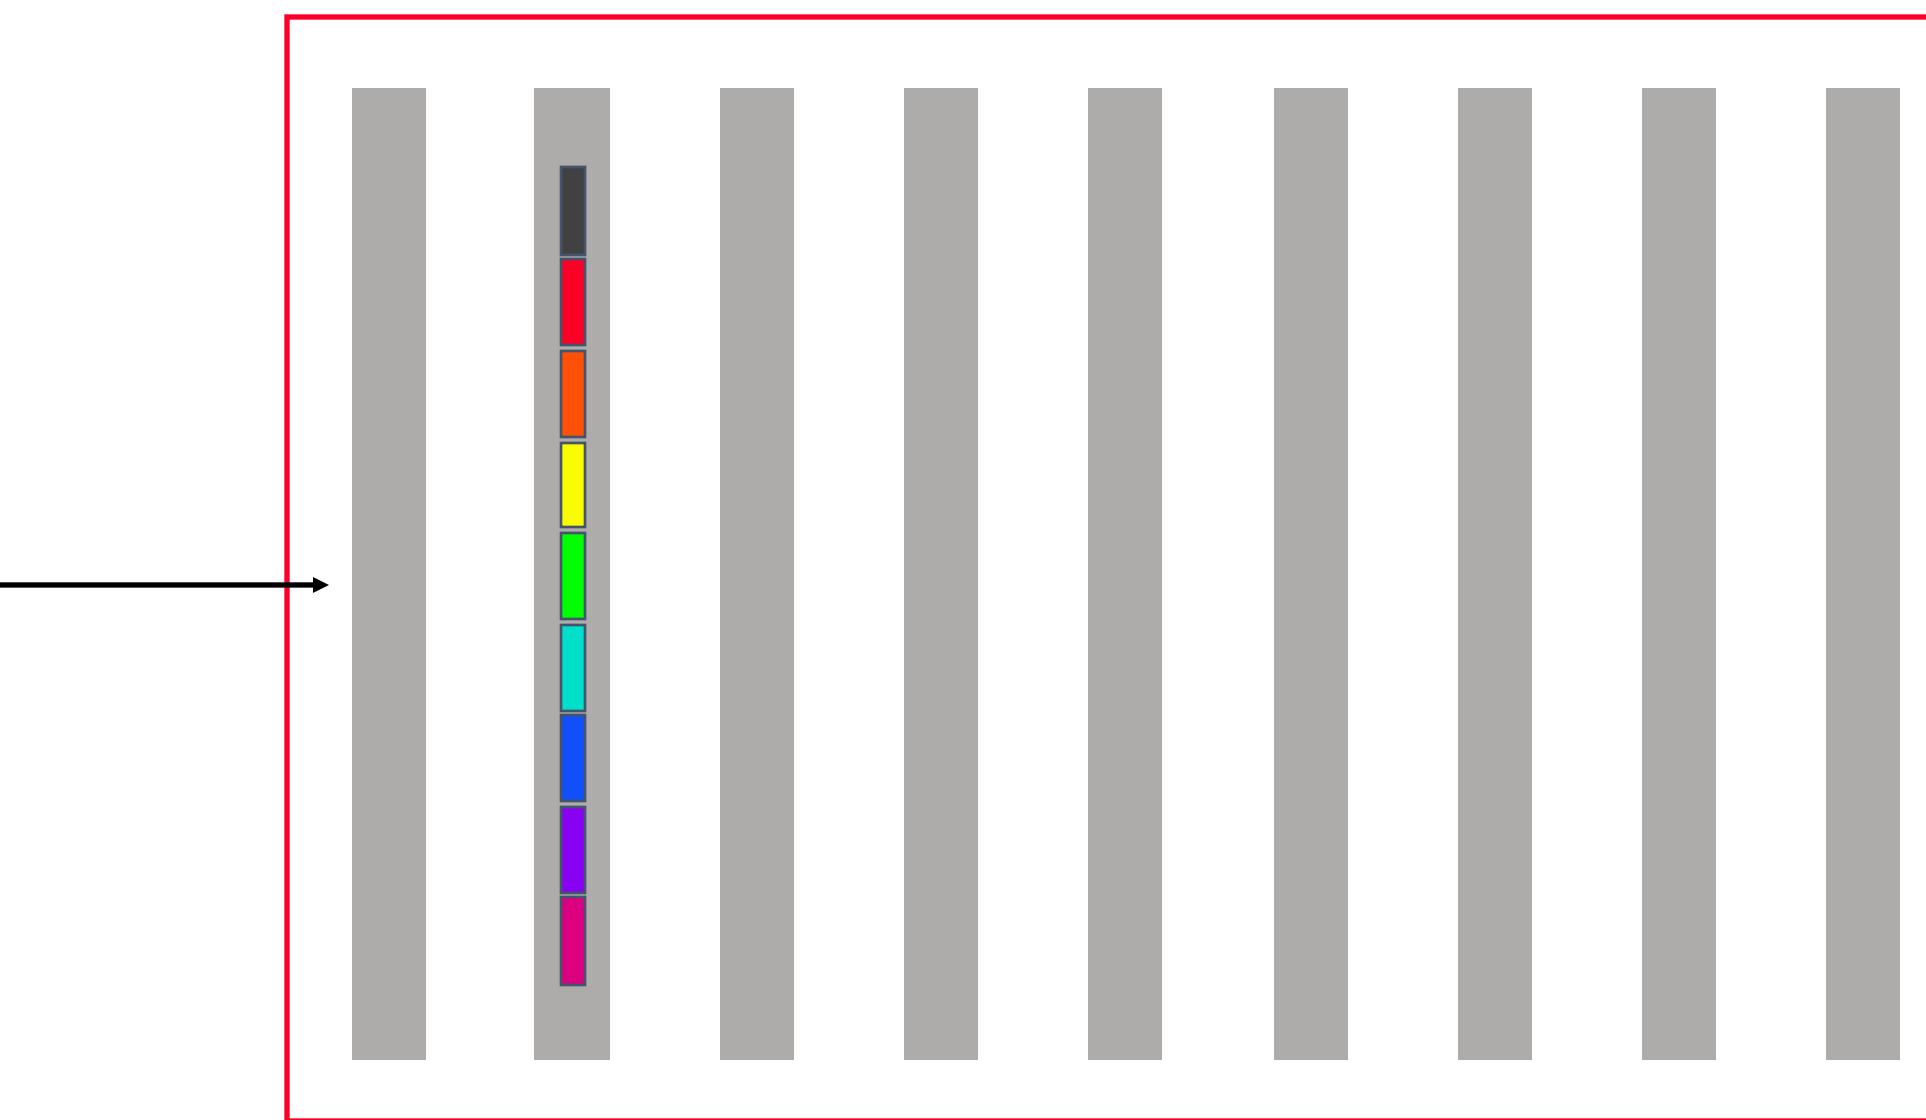


Gather

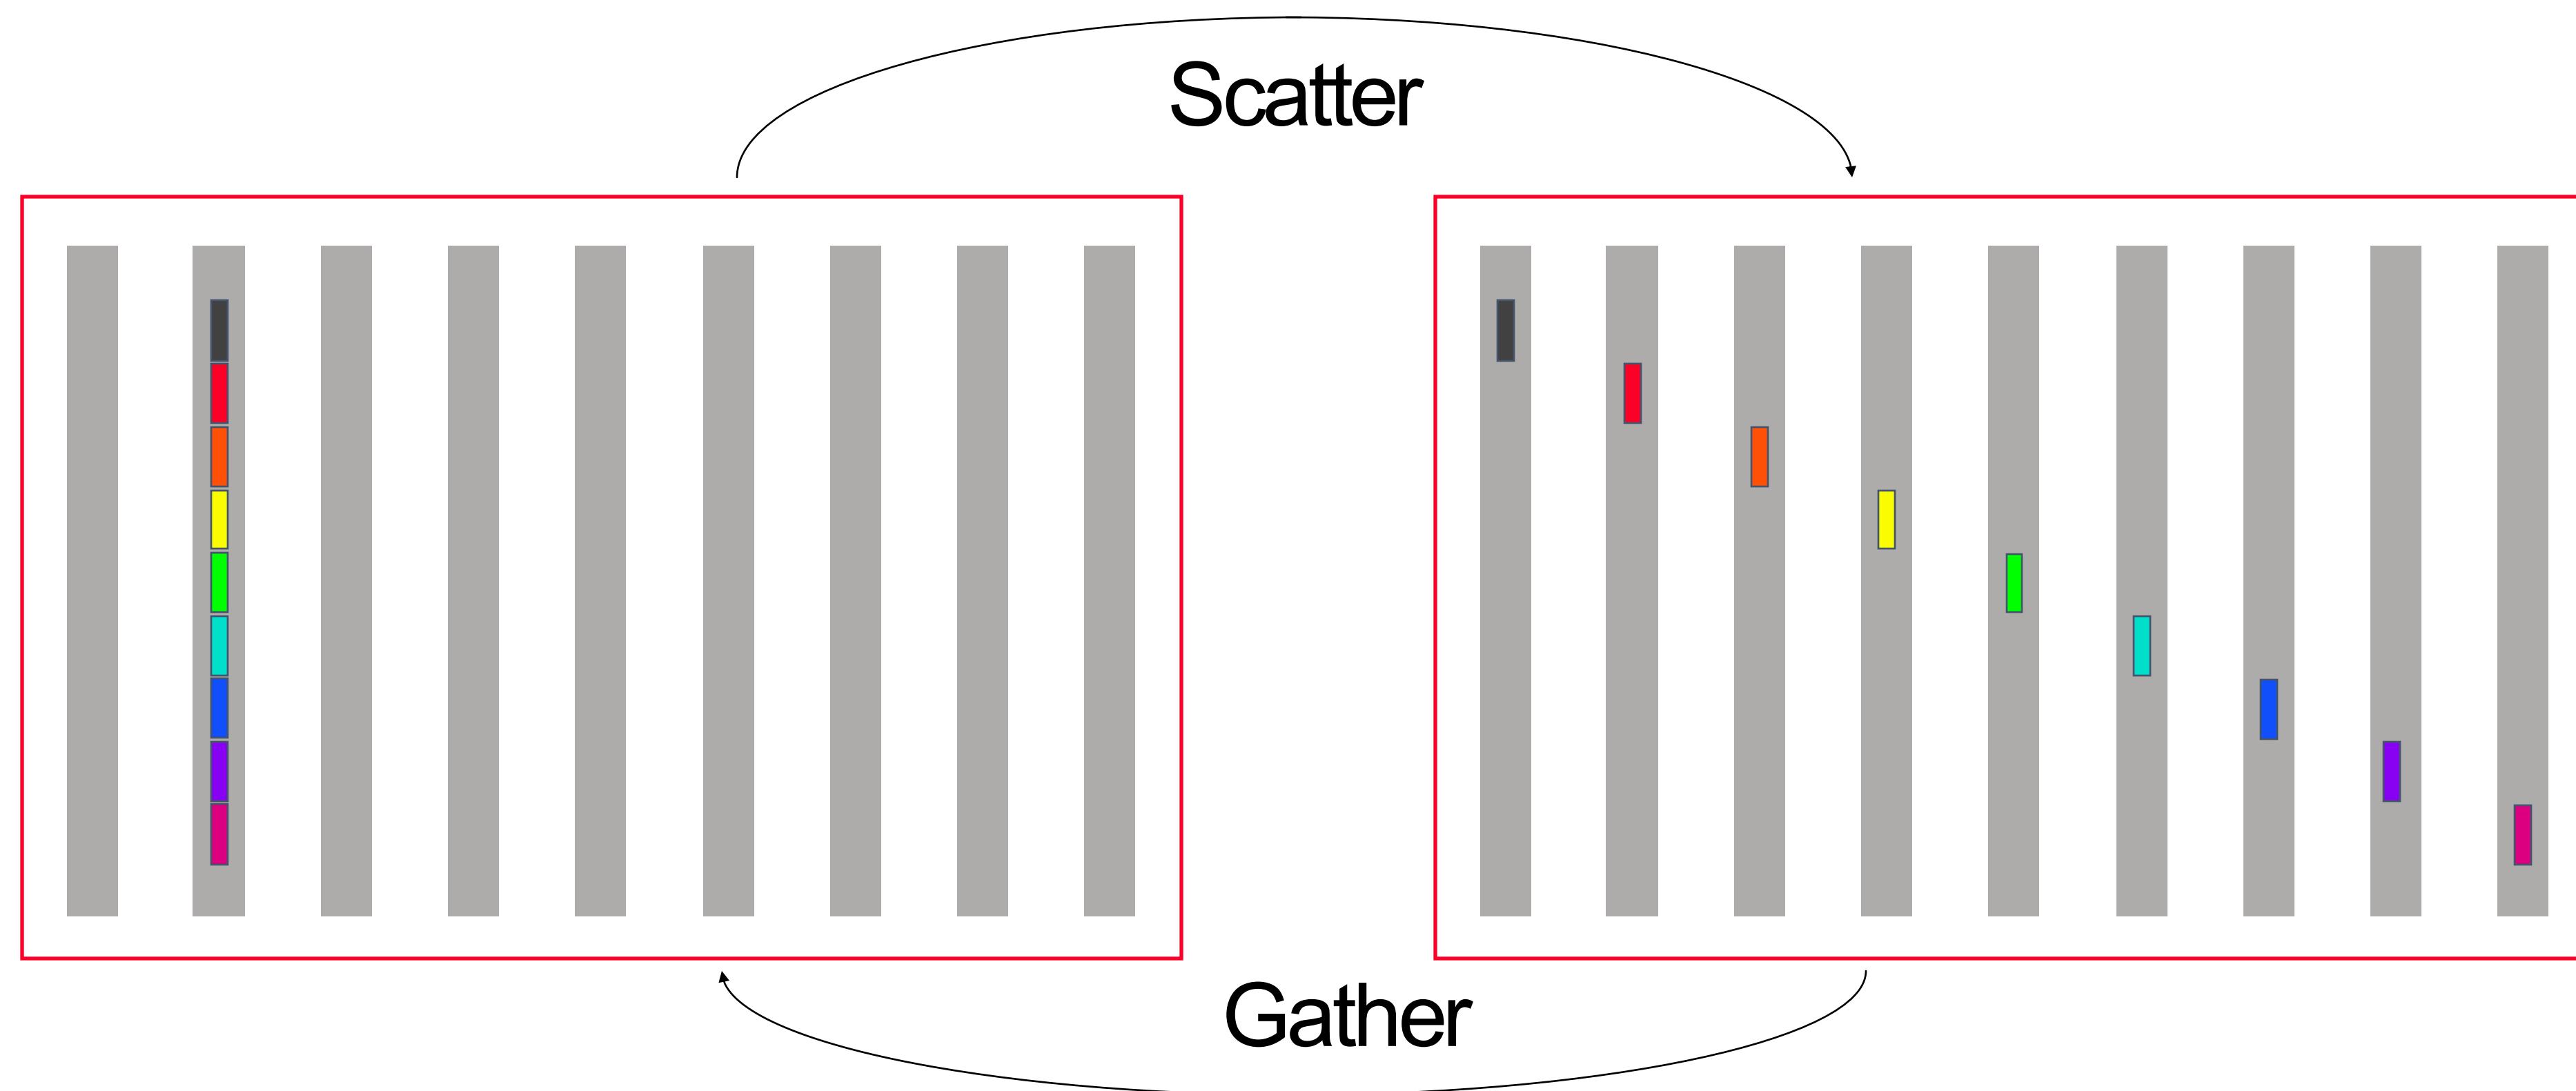
Before



After

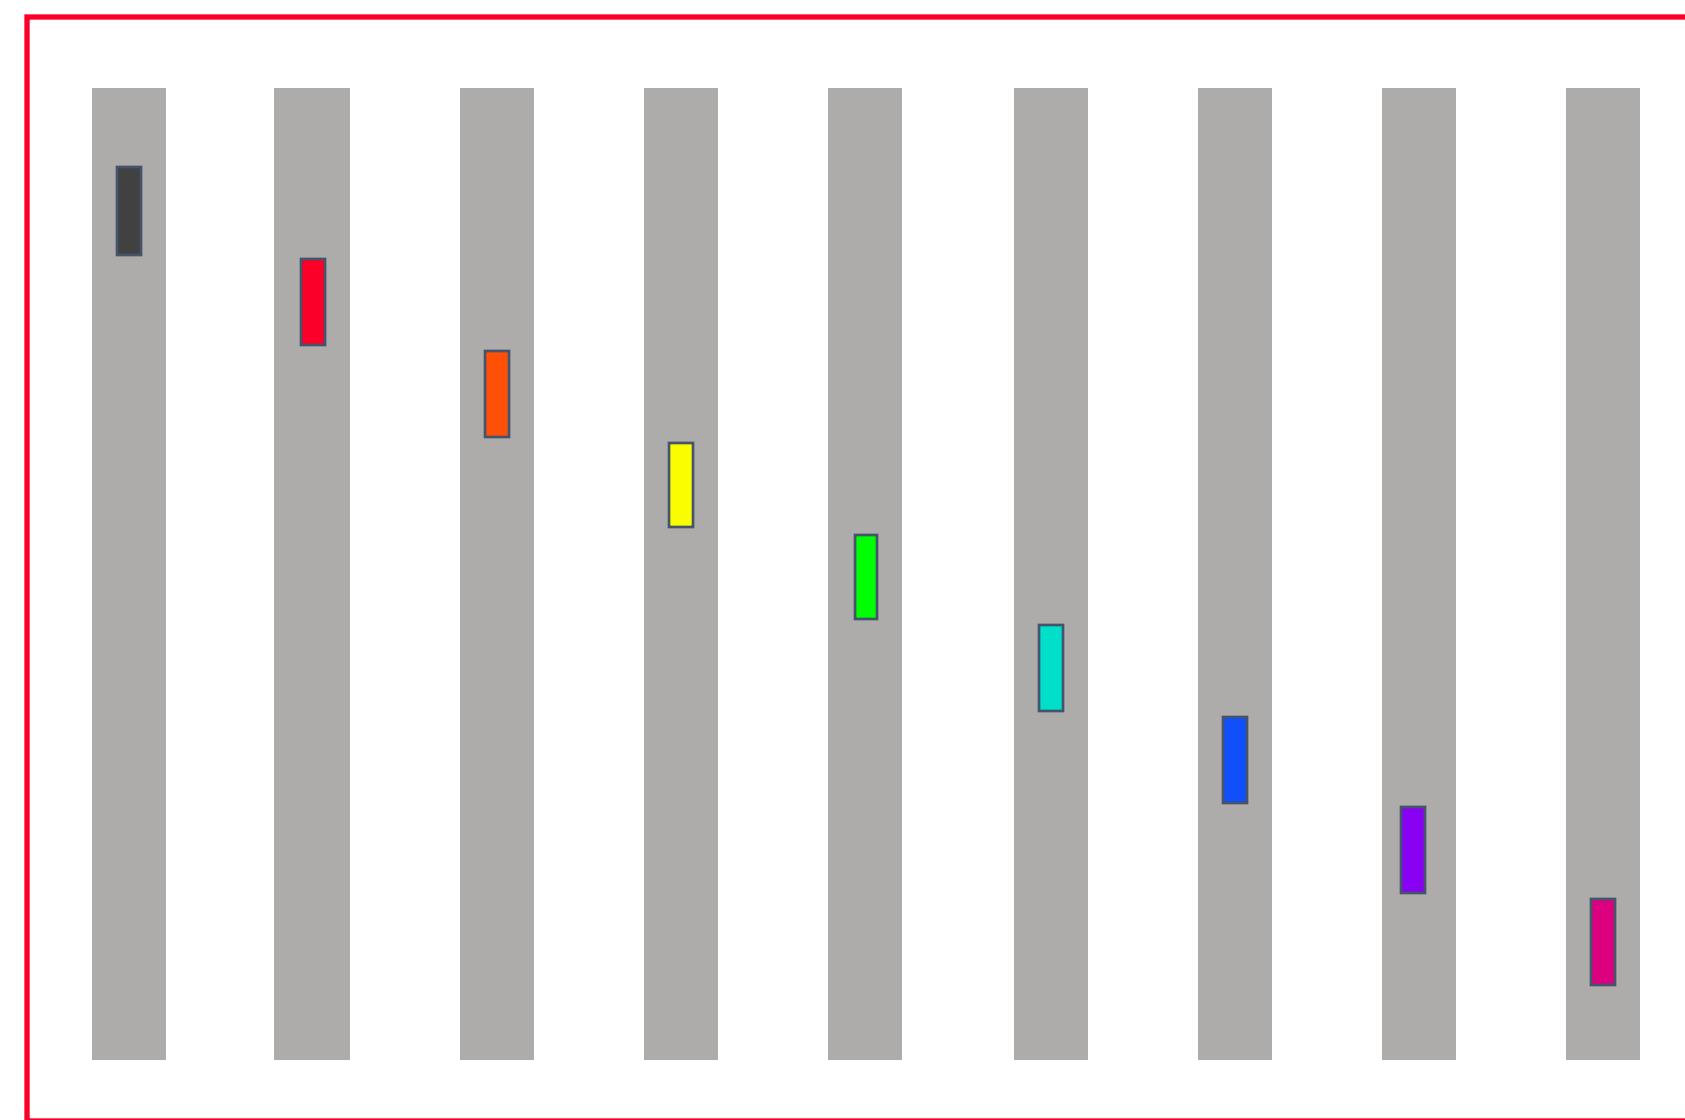


Scatter/Gather

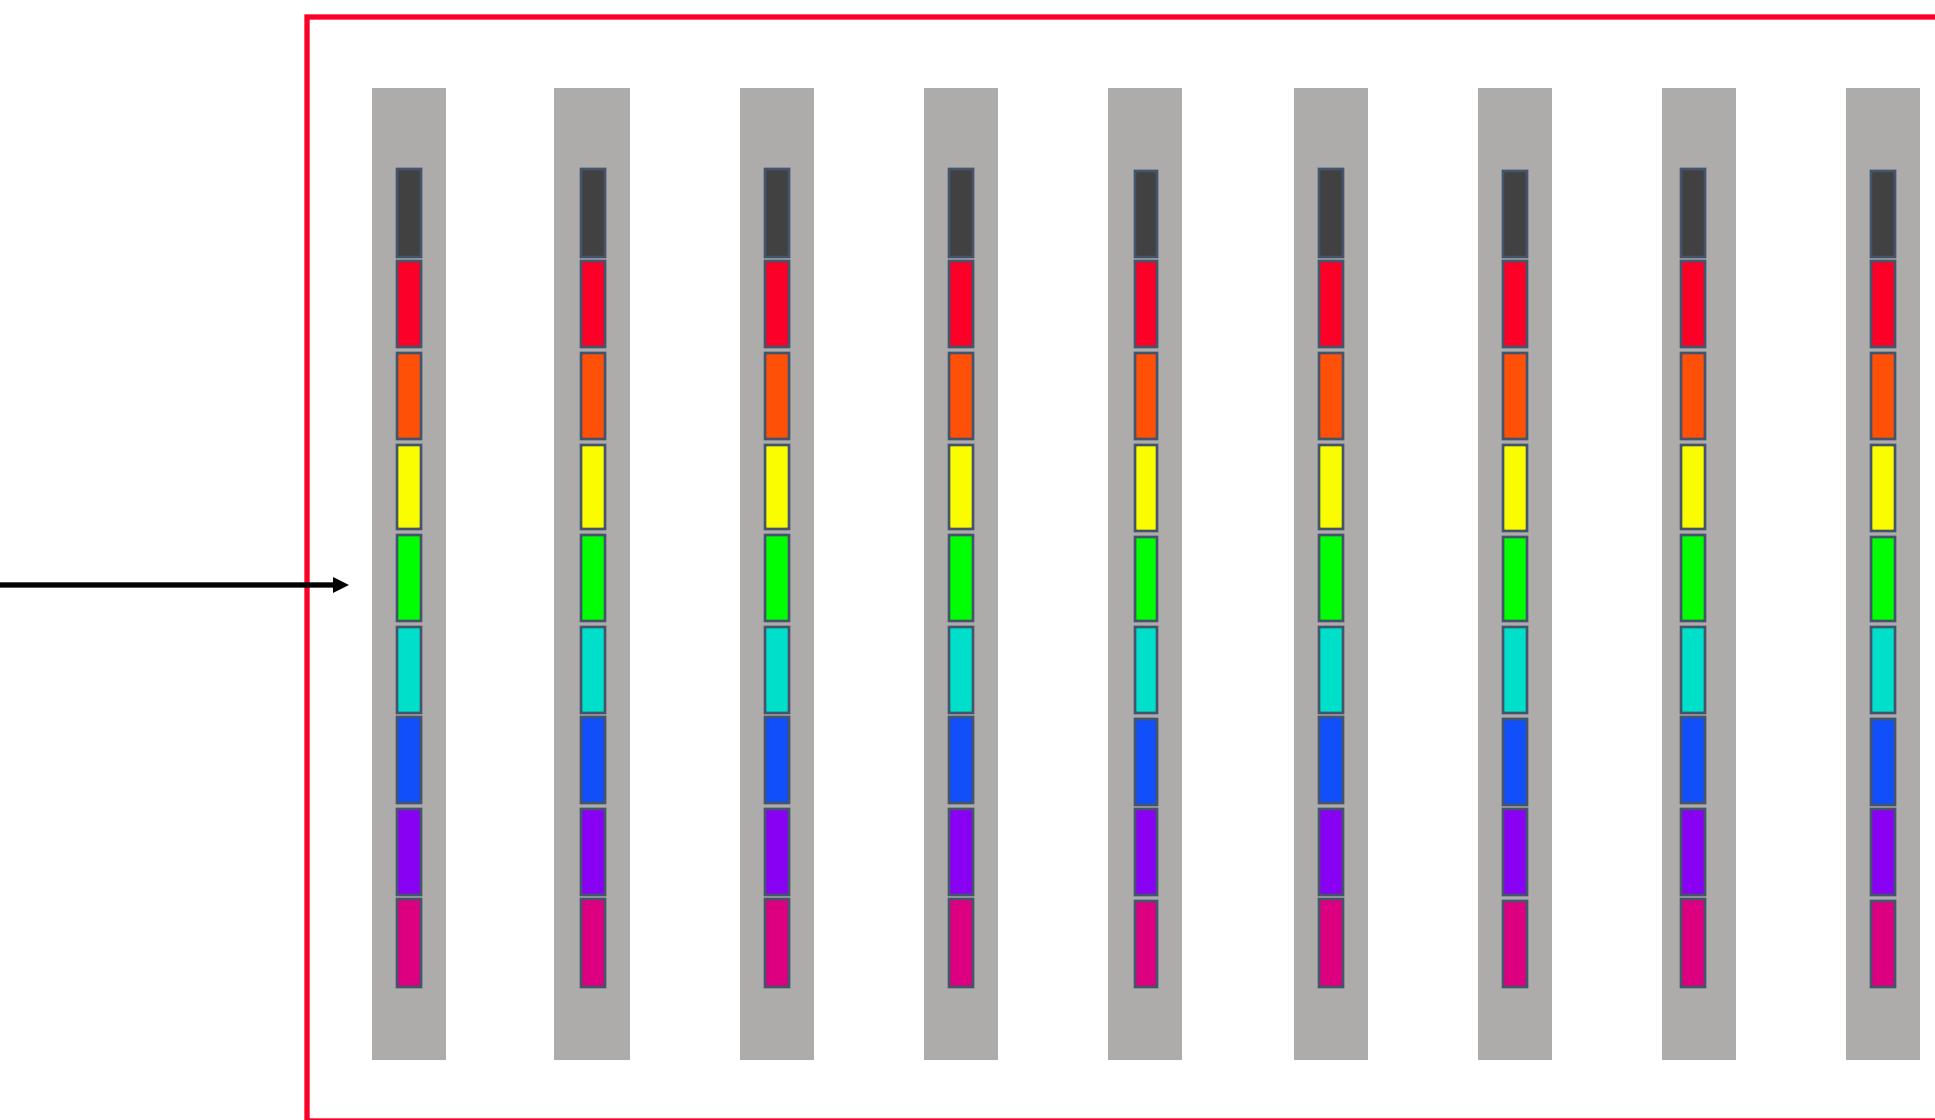


Allgather

Before

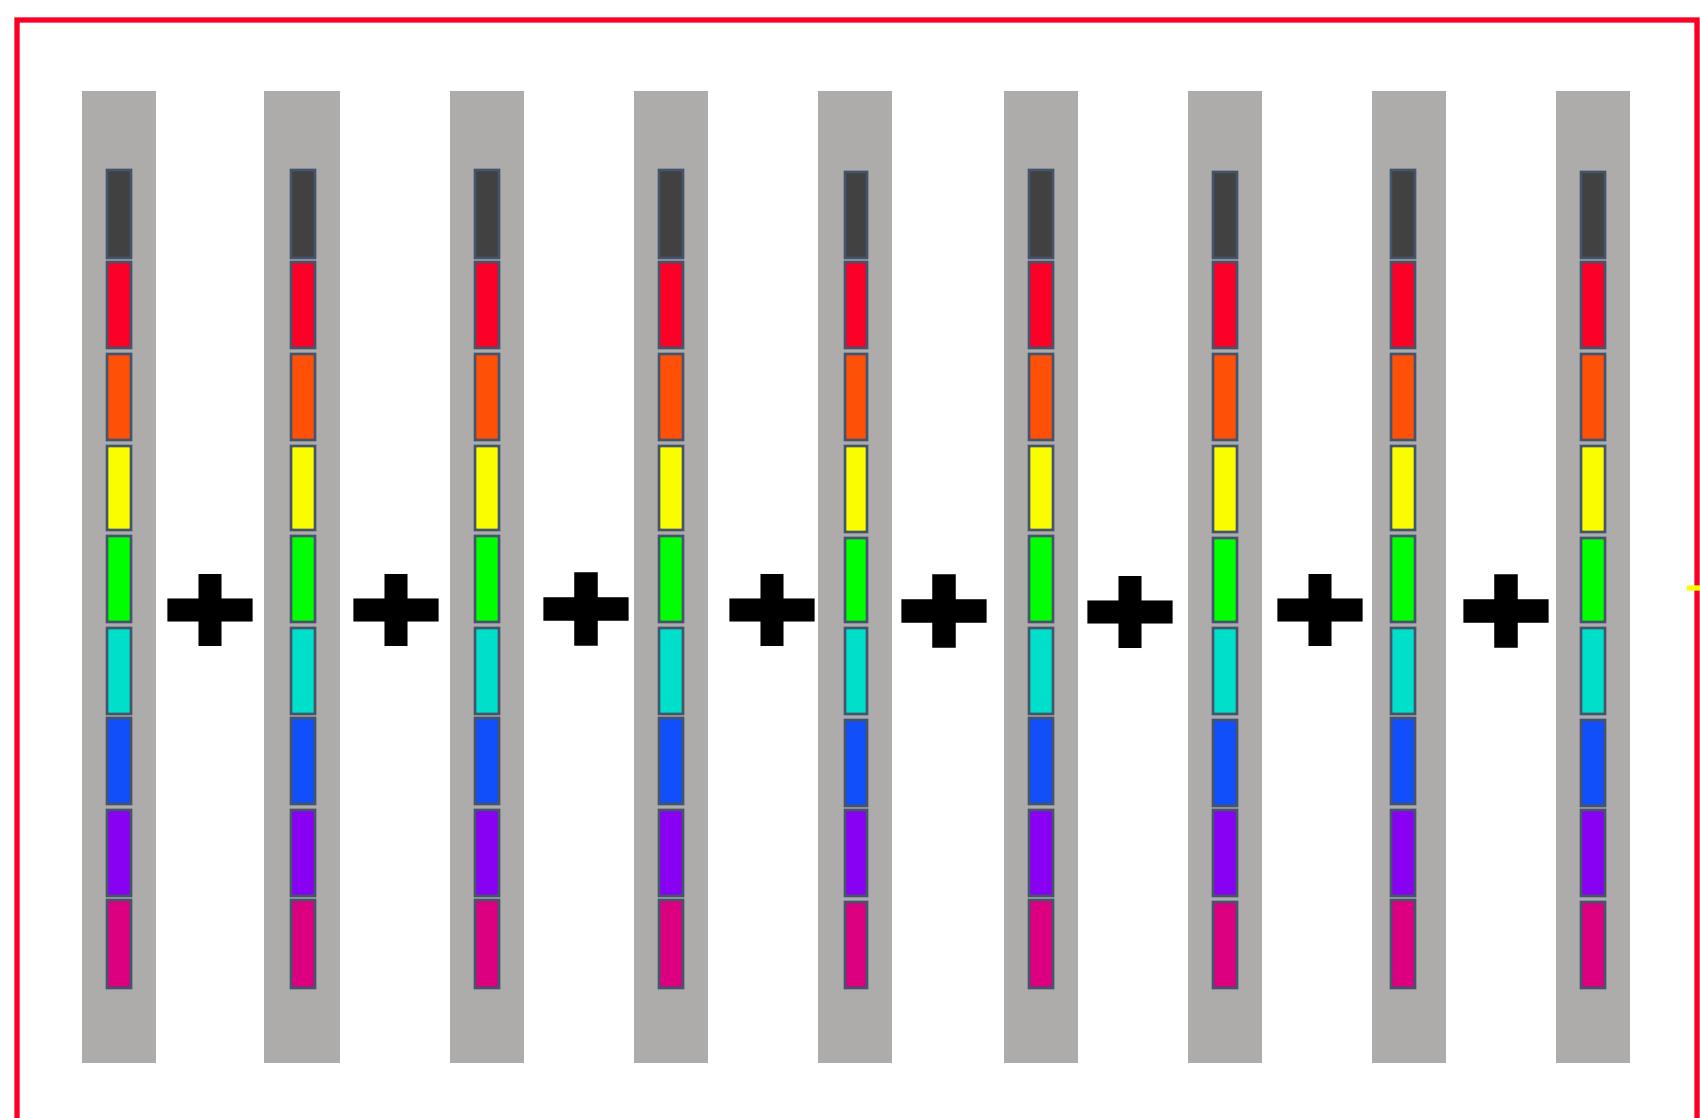


After

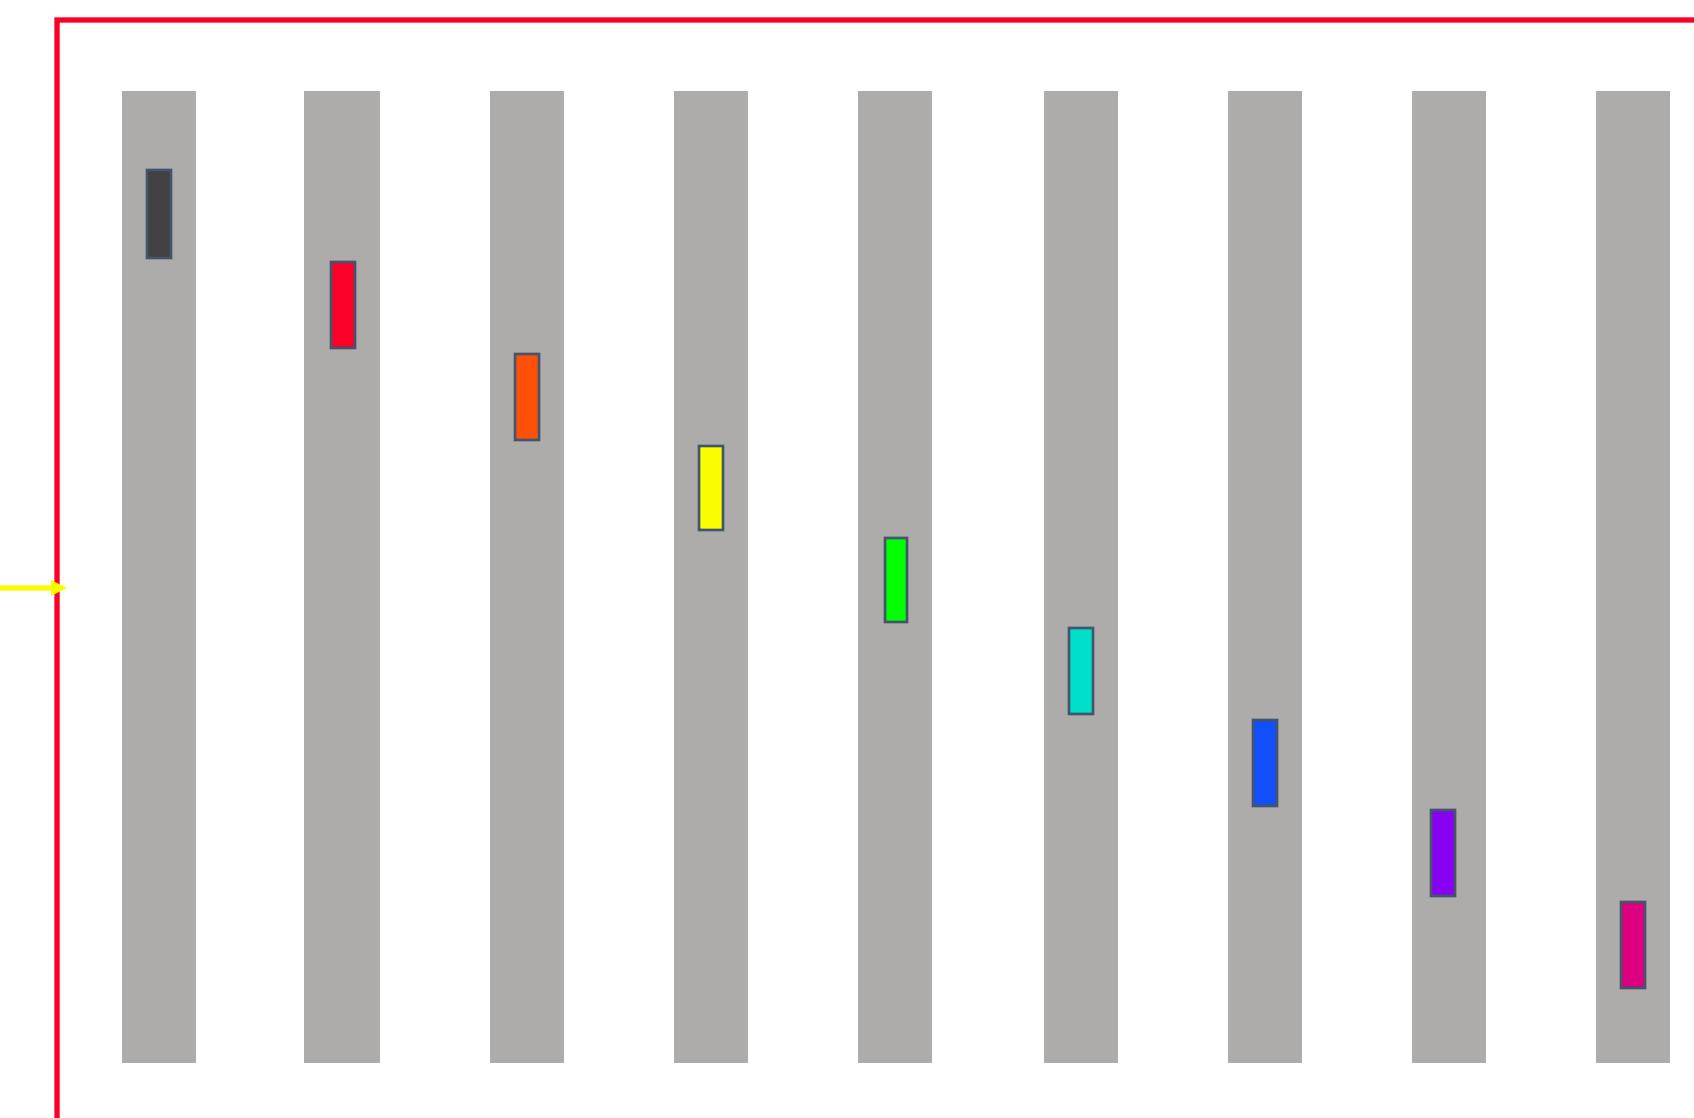


Reduce-scatter

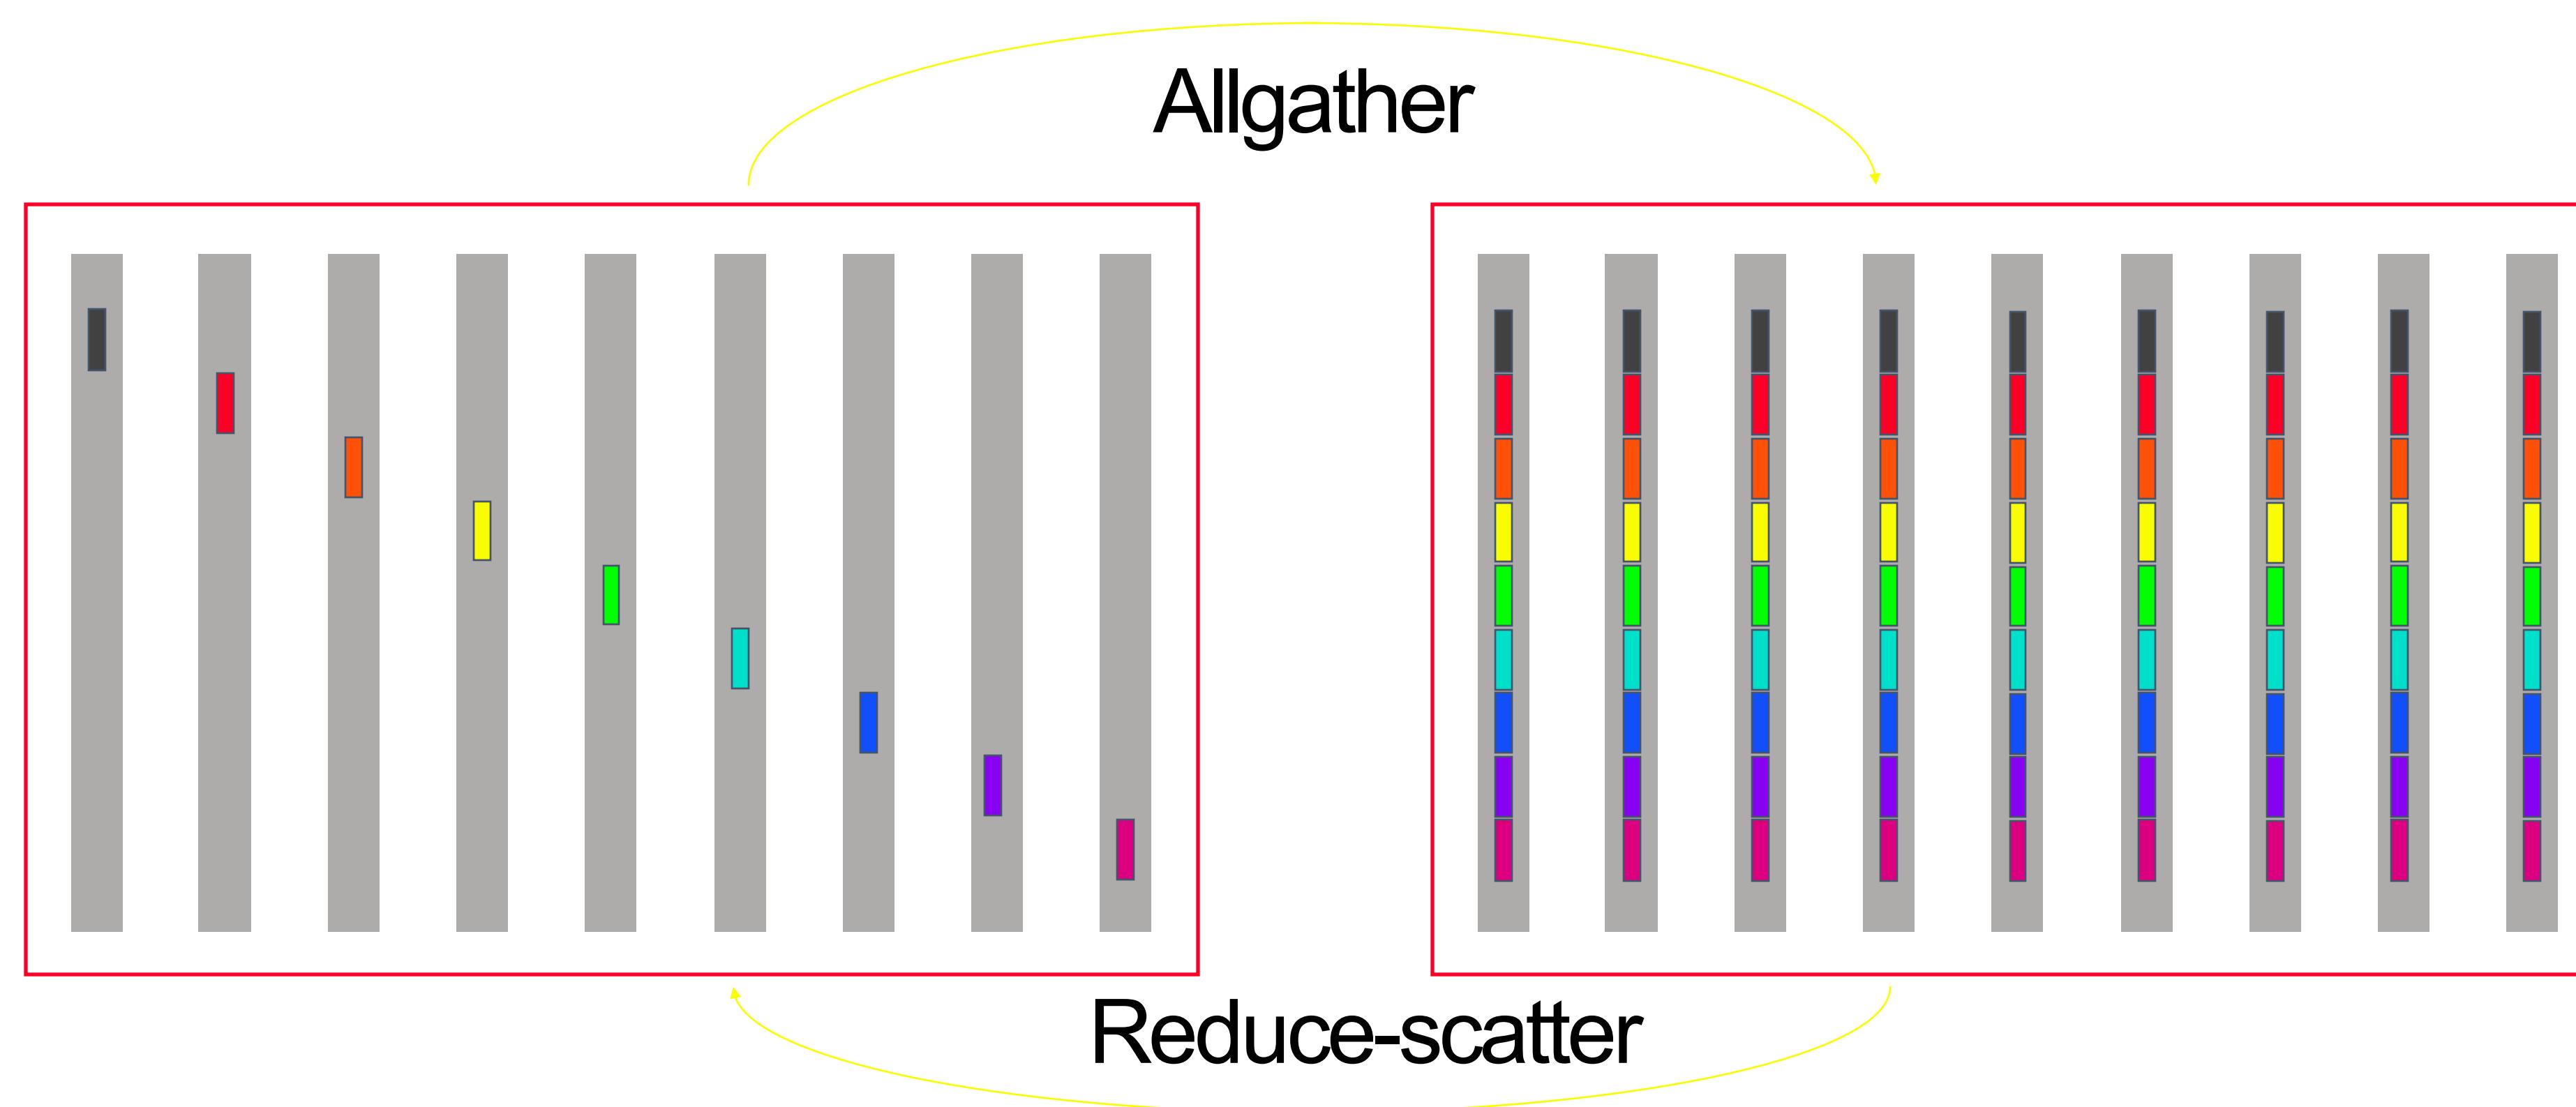
Before



After

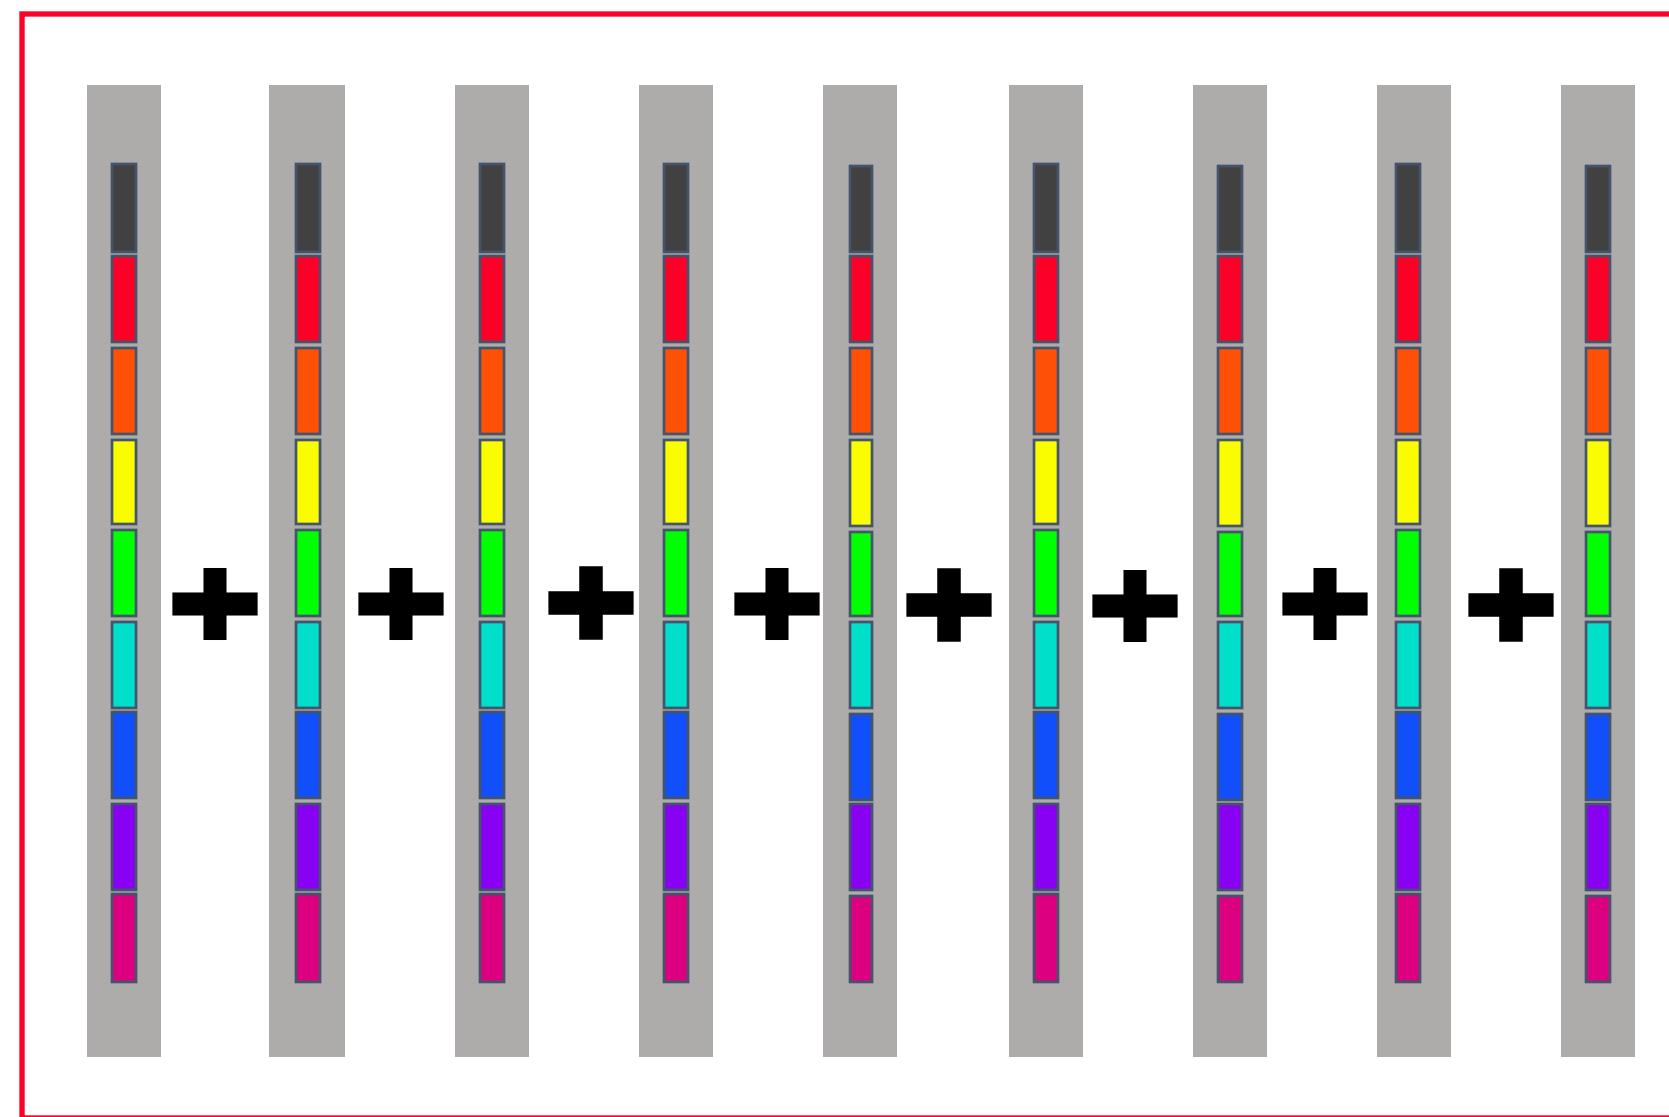


Allgather/Reduce-scatter

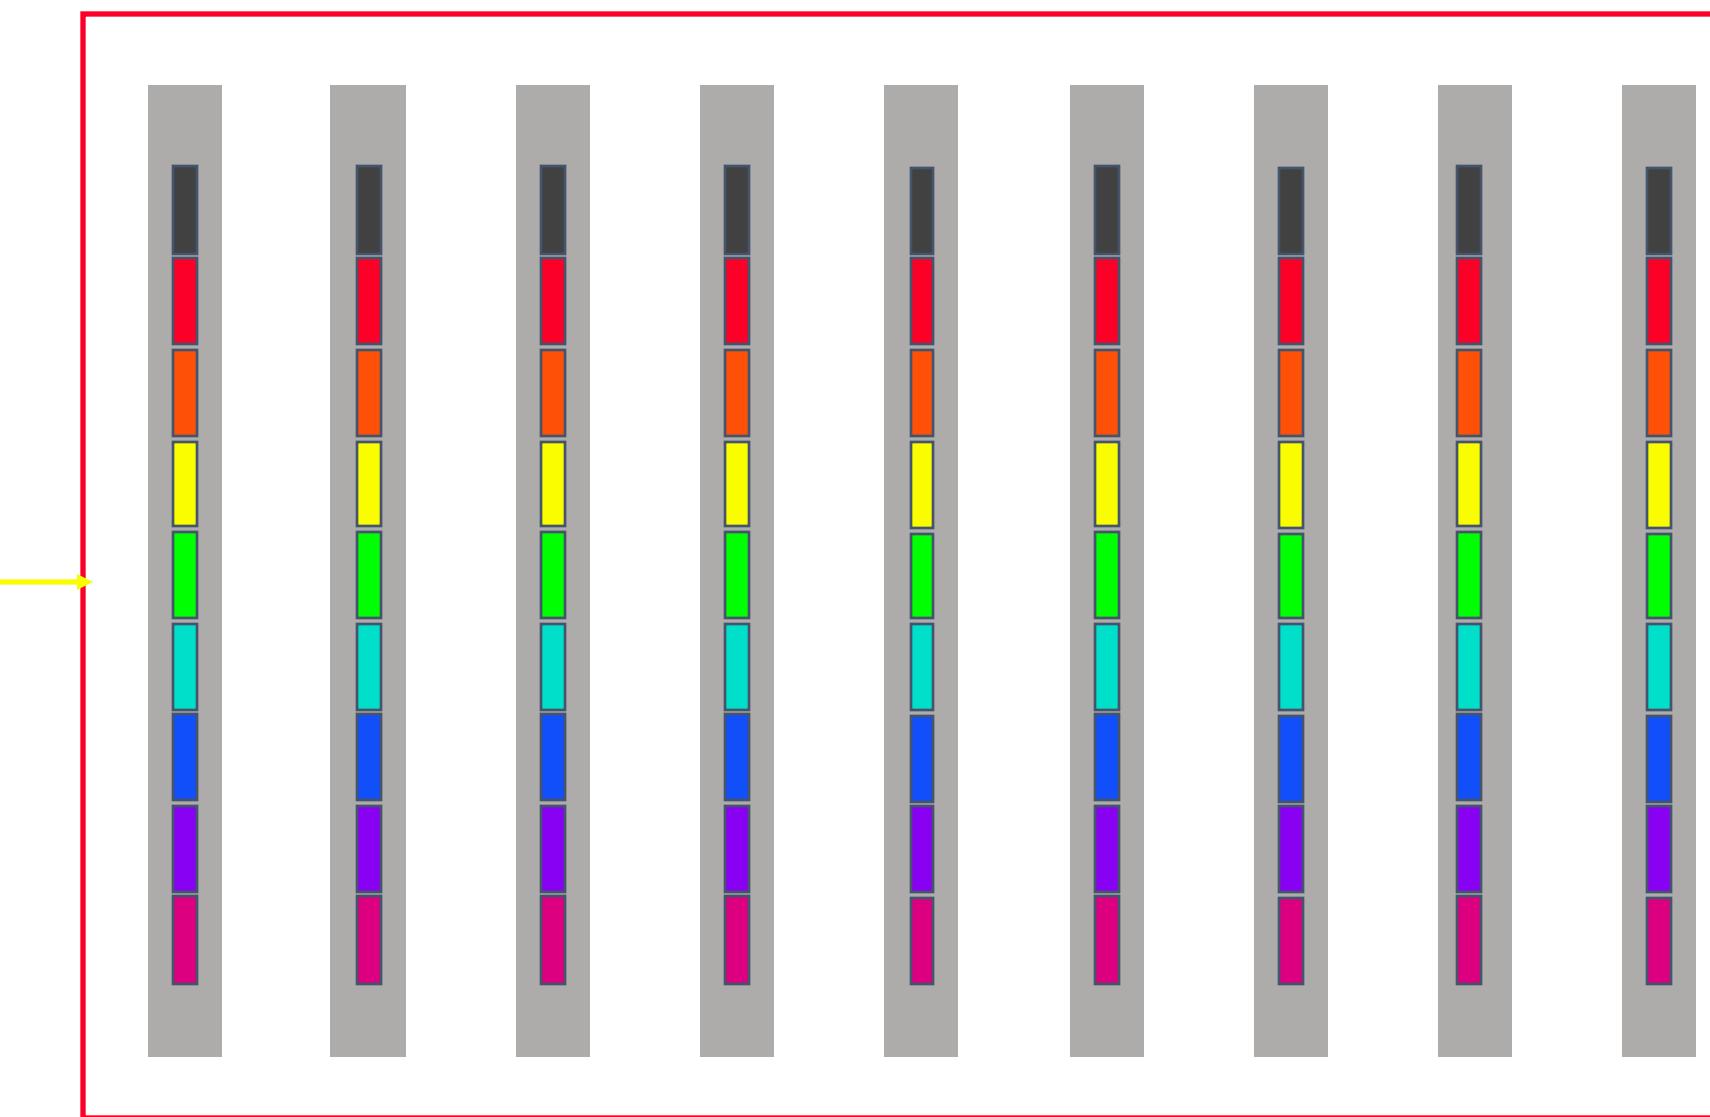


Allreduce

Before

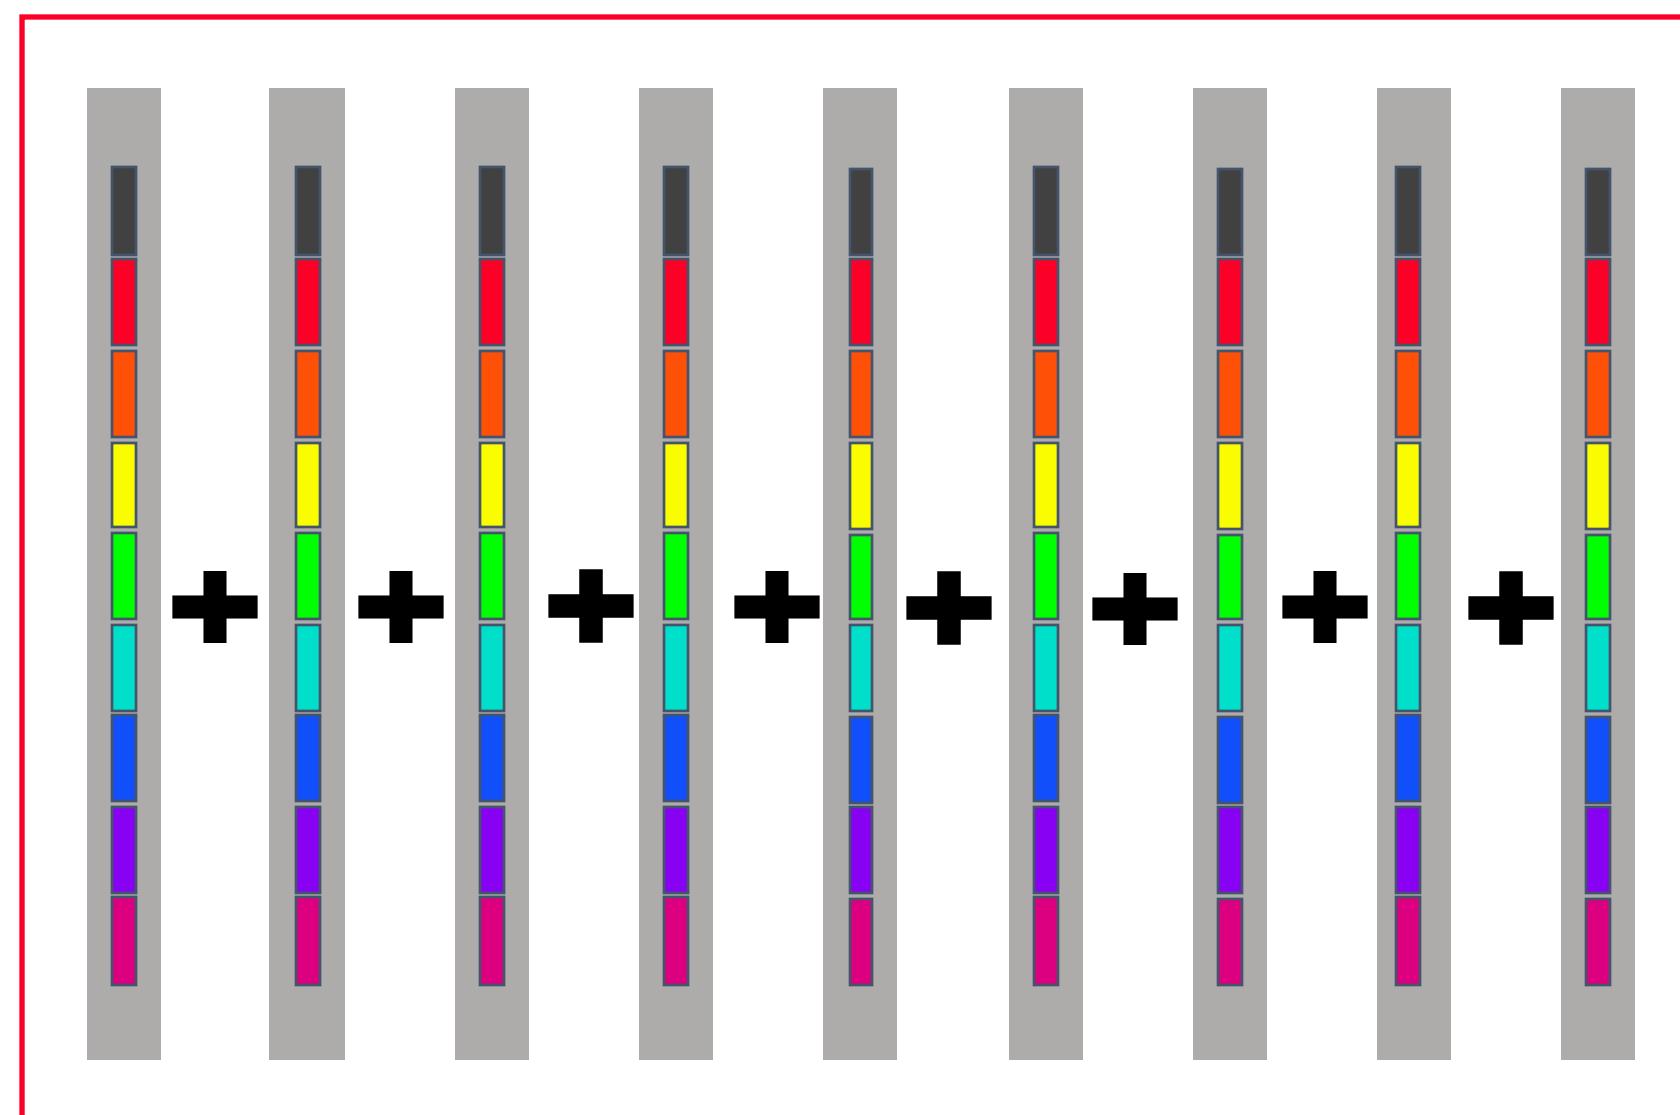


After

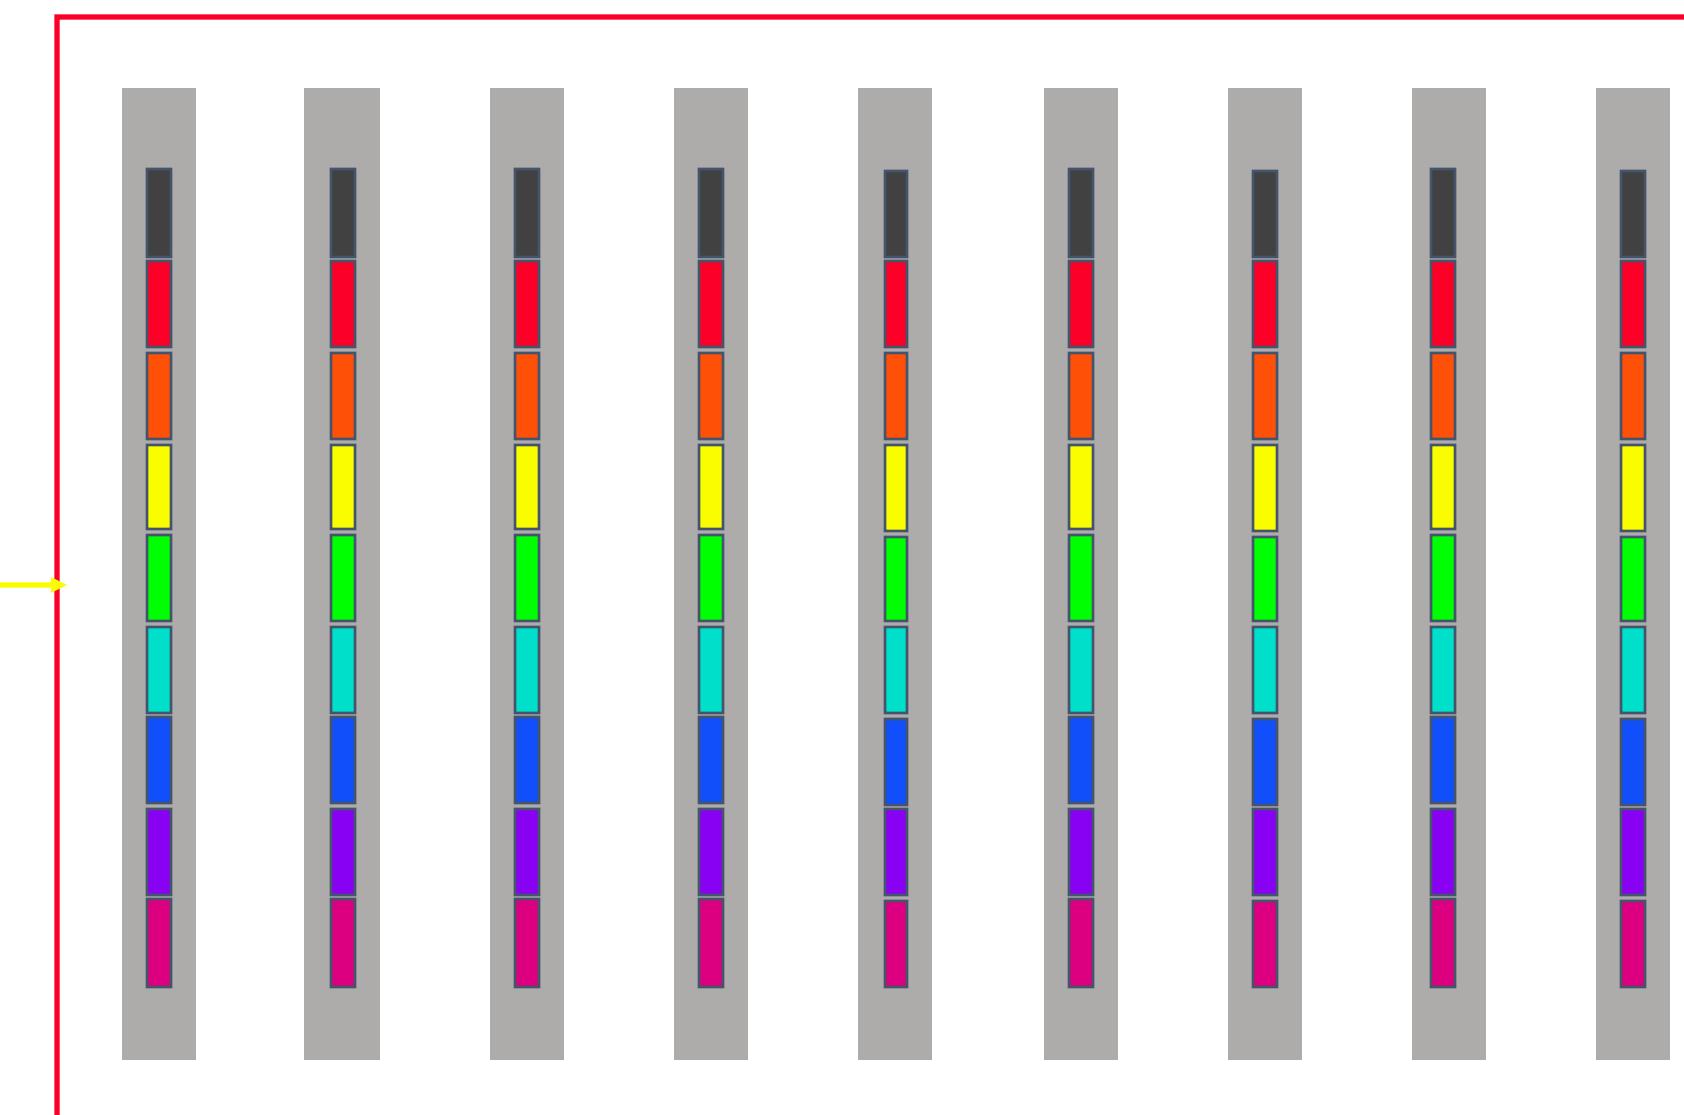


All2All

Before



After



Some Facts

- Collective is much more expensive than P2P
 - Collective can be assembled using many P2P
 - Collective is cheaper than realizing collective using P2P (we'll see)
- Collective is highly optimized in the past 20 years
 - Look out for "X" CCL libraries
 - NCCL, MCCL, OneCCL, UCCL
 - Collective is not fault-tolerant
 - A major sources of faults in ML systems

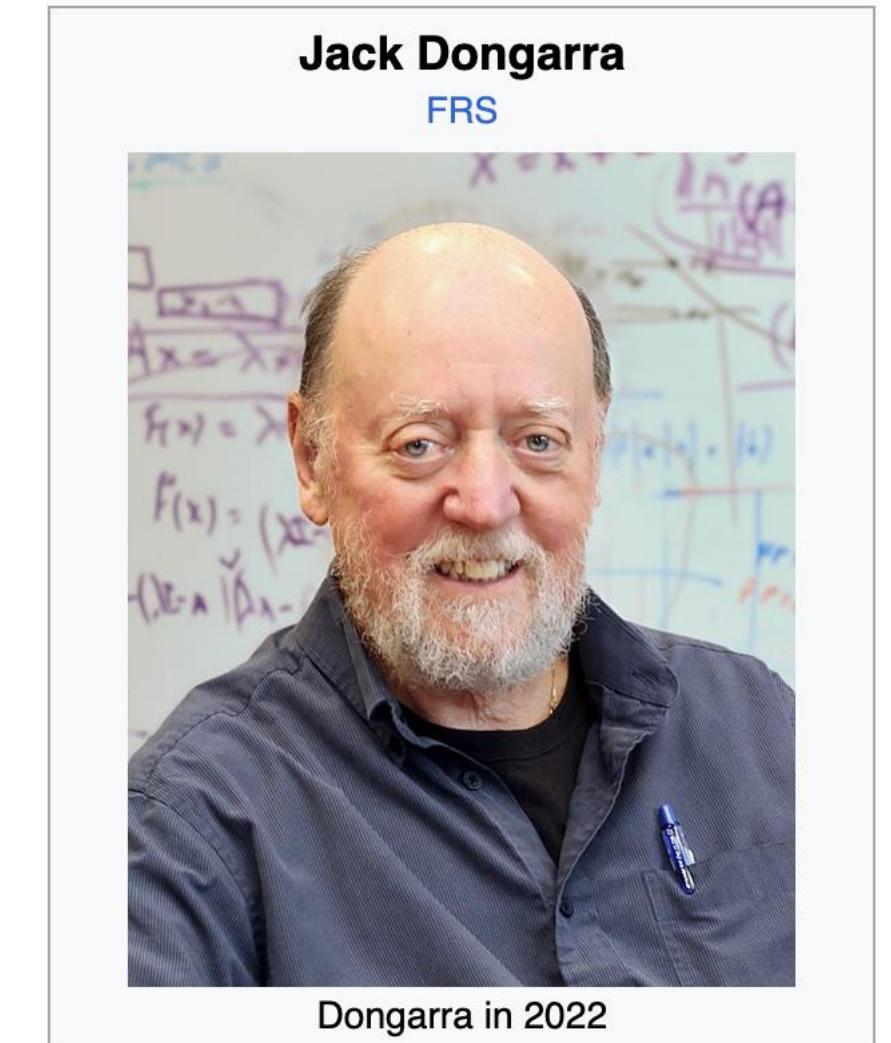
Communication Model: $\alpha\beta$ model

Communication Model: $\alpha + n\beta, \beta = \frac{1}{B}$

- Small Message size ($n \rightarrow 0$): α dominates, emphasize latency
- Large Message Size ($n \rightarrow +\infty$): $n\beta$ dominate, emphasize bandwidth utilization

Two Family of Mainstream Algorithms/Implementations

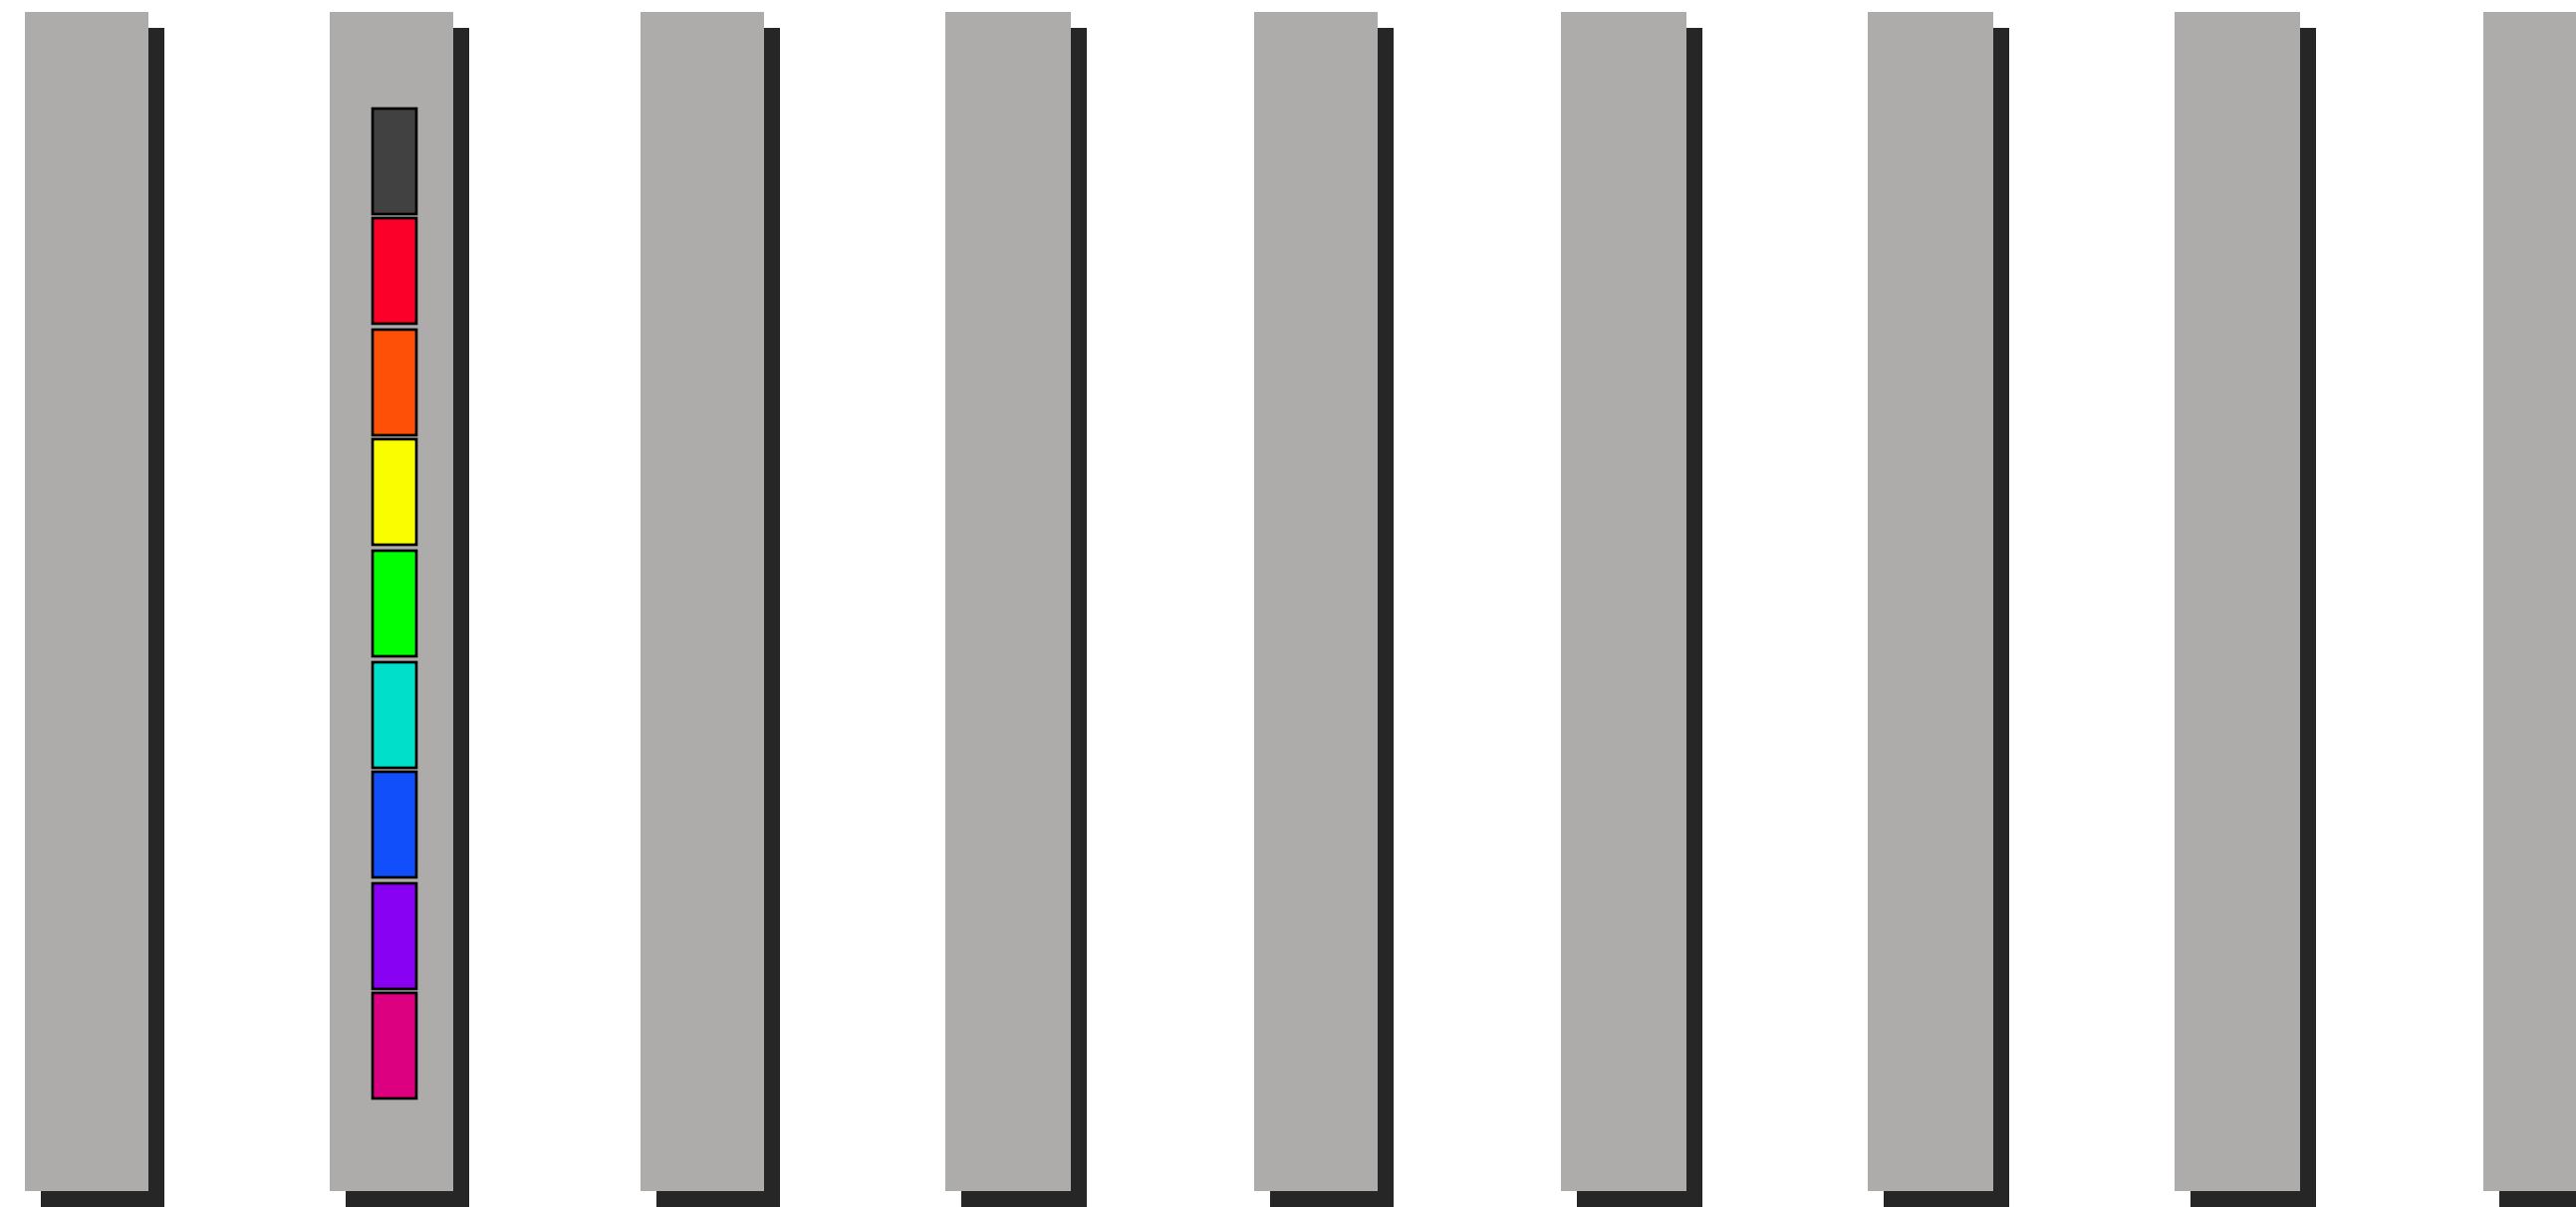
- Small message: Minimum Spanning Tree algorithm
 - Emphasize **low latency**
- Large Message: Ring algorithm
 - Emphasize **bandwidth utilization**
- There are 50+ different algorithms developed in the past 50 years by a community called “High-performance computing”
 - 2021 Turing award



General principles: Low Latency

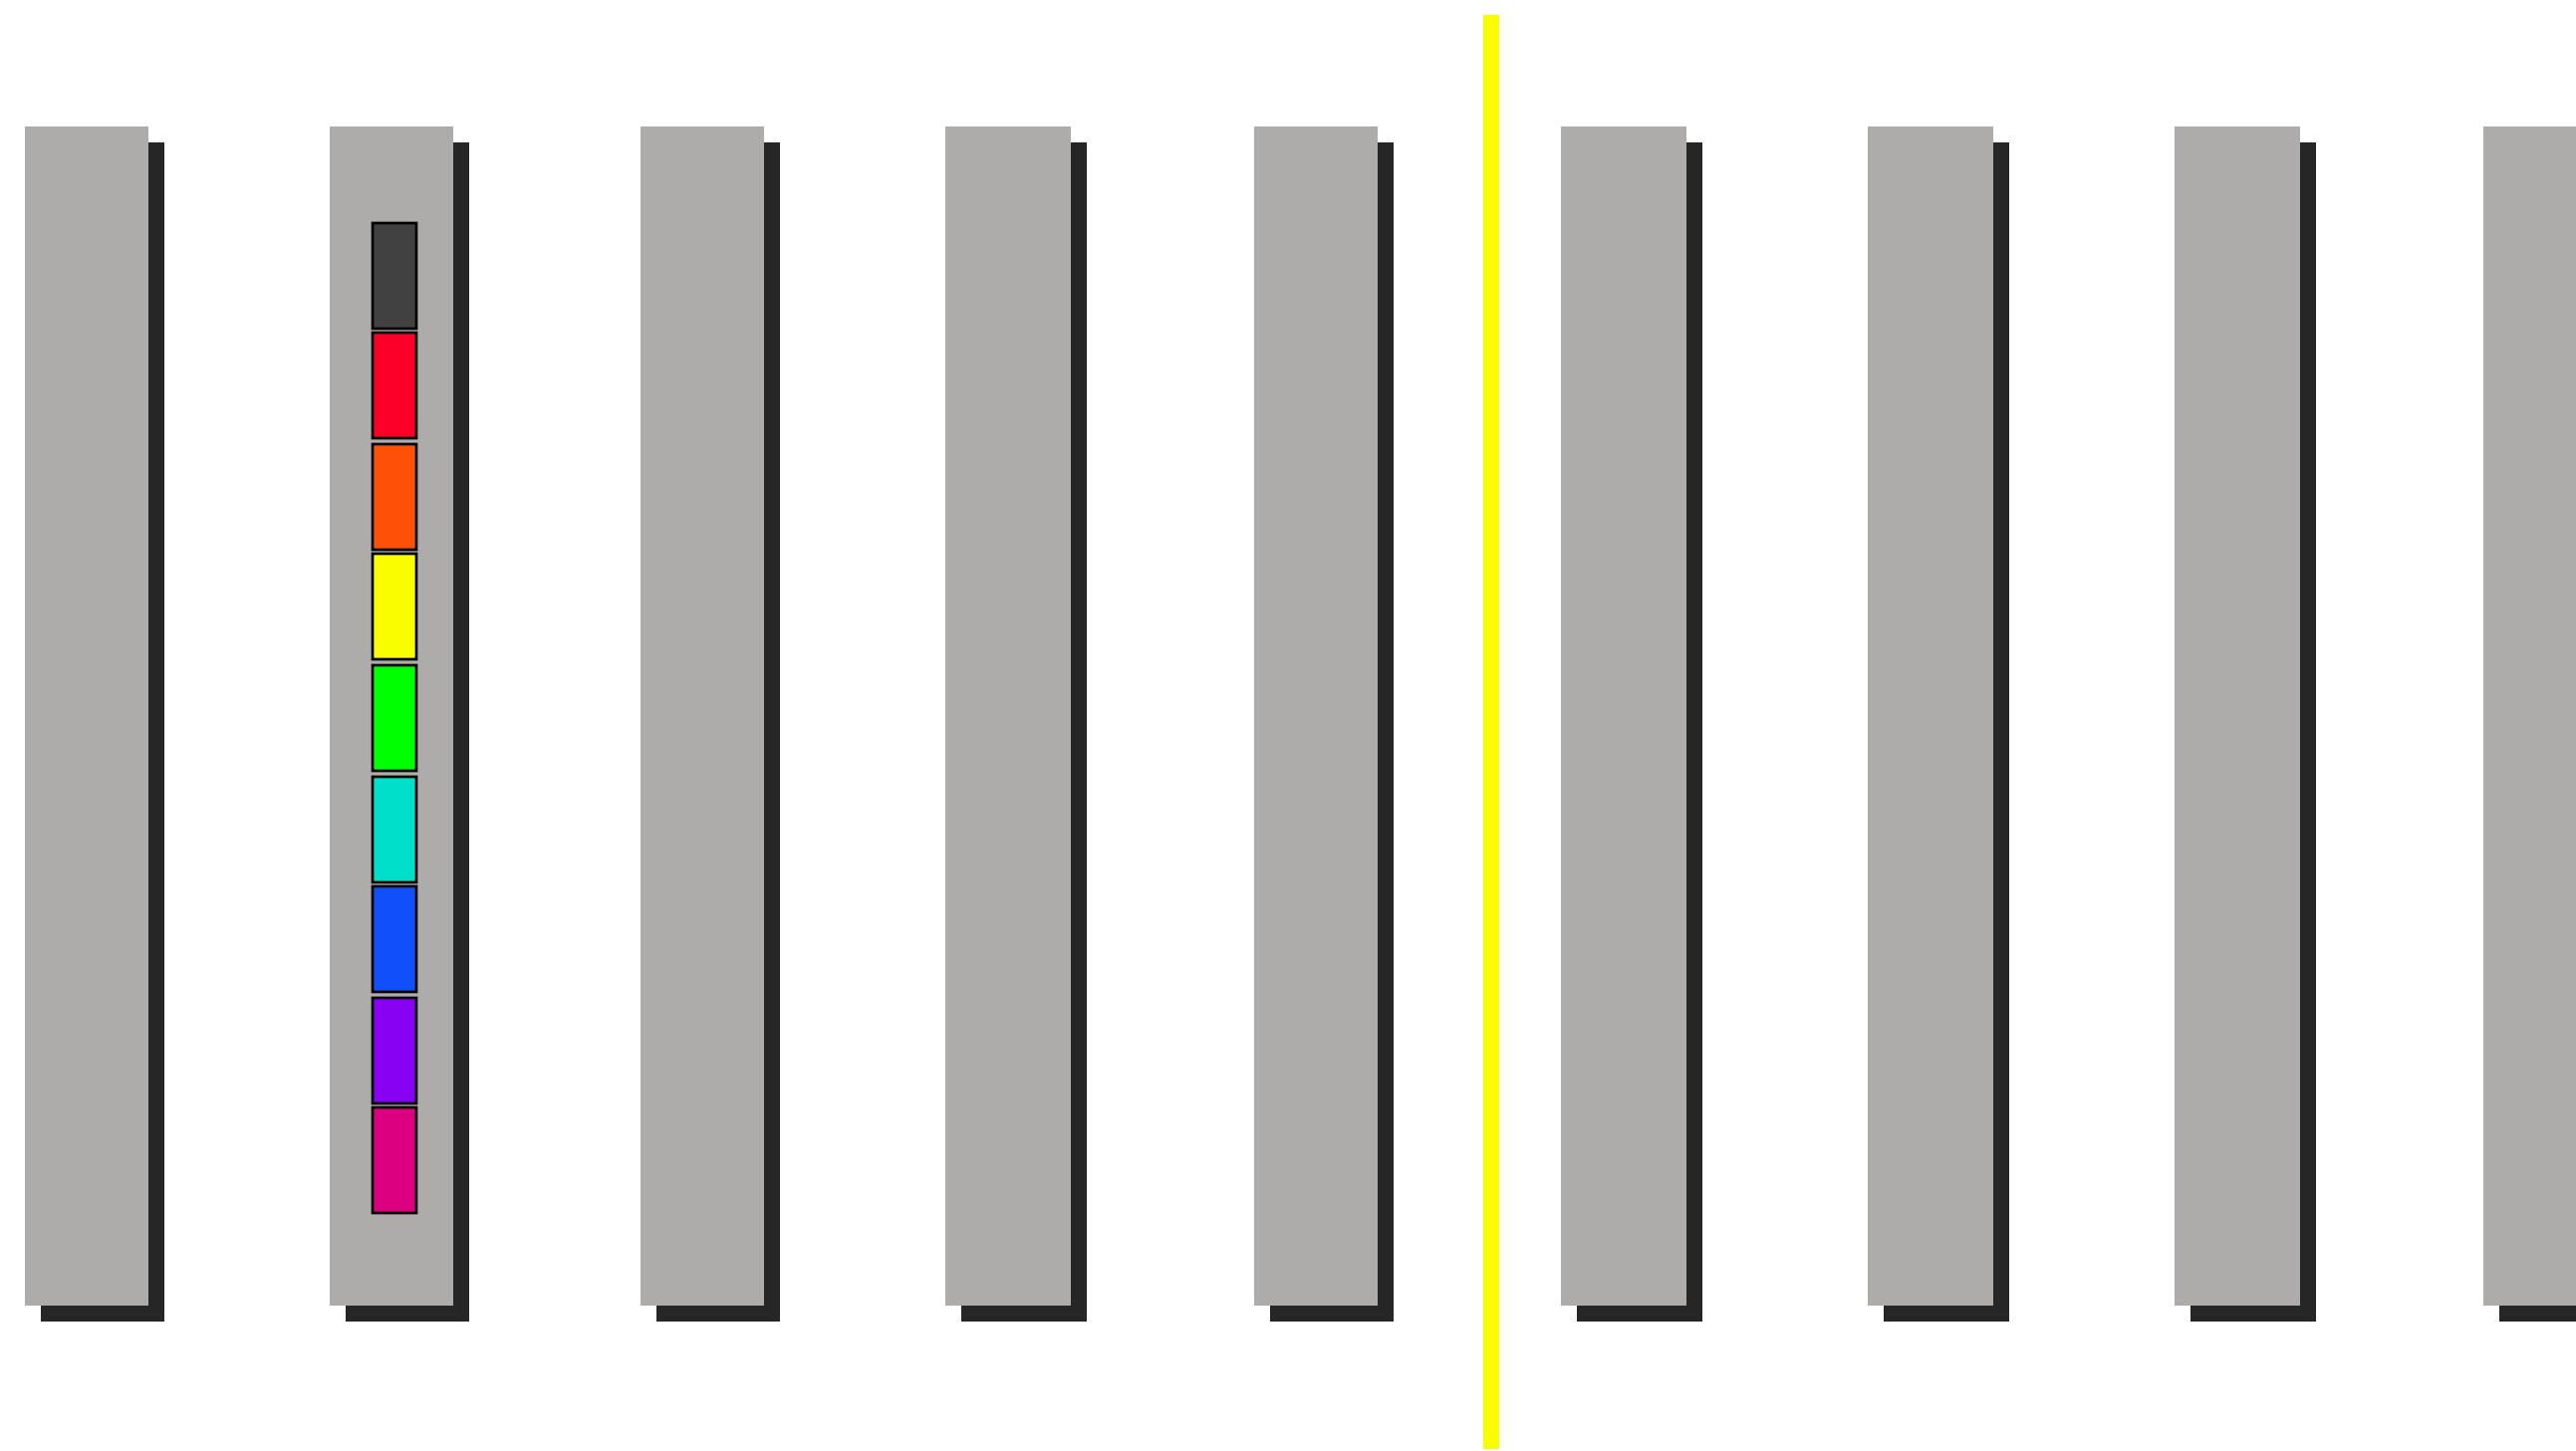
- Minimize the number of rounds needed for communication
- Minimal-spanning tree algorithm

General principles



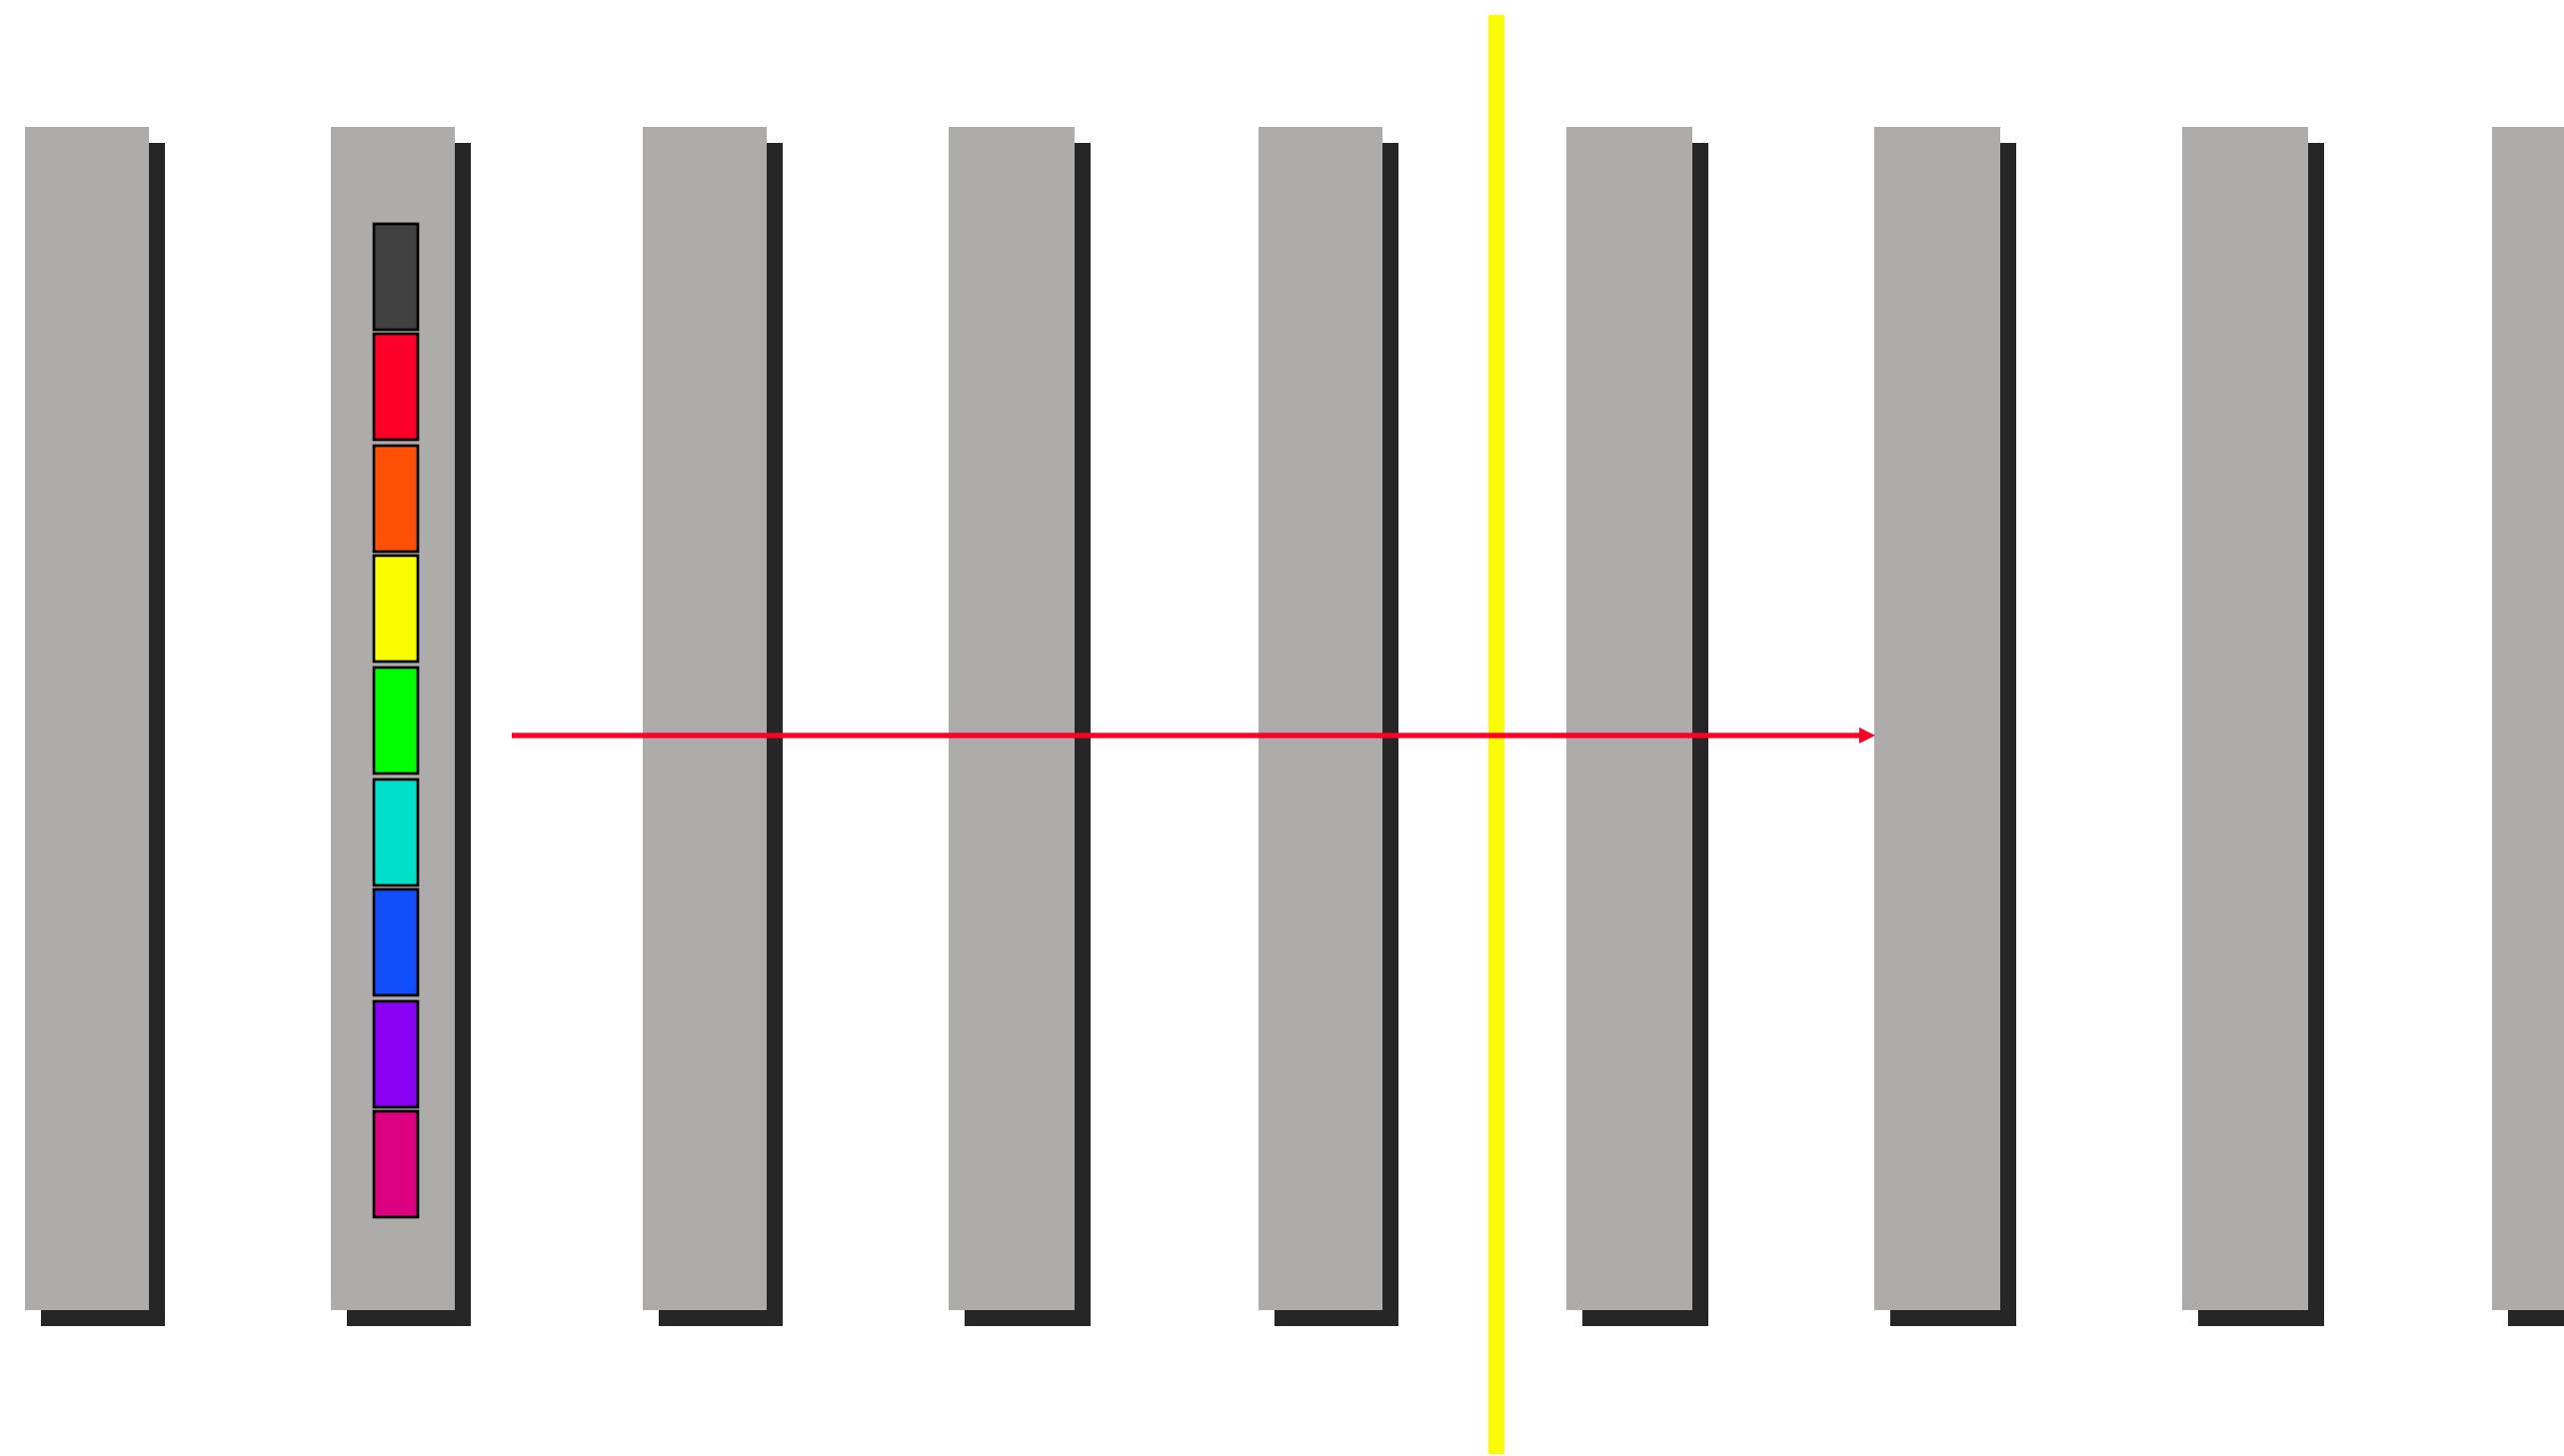
- message starts on one processor

General principles



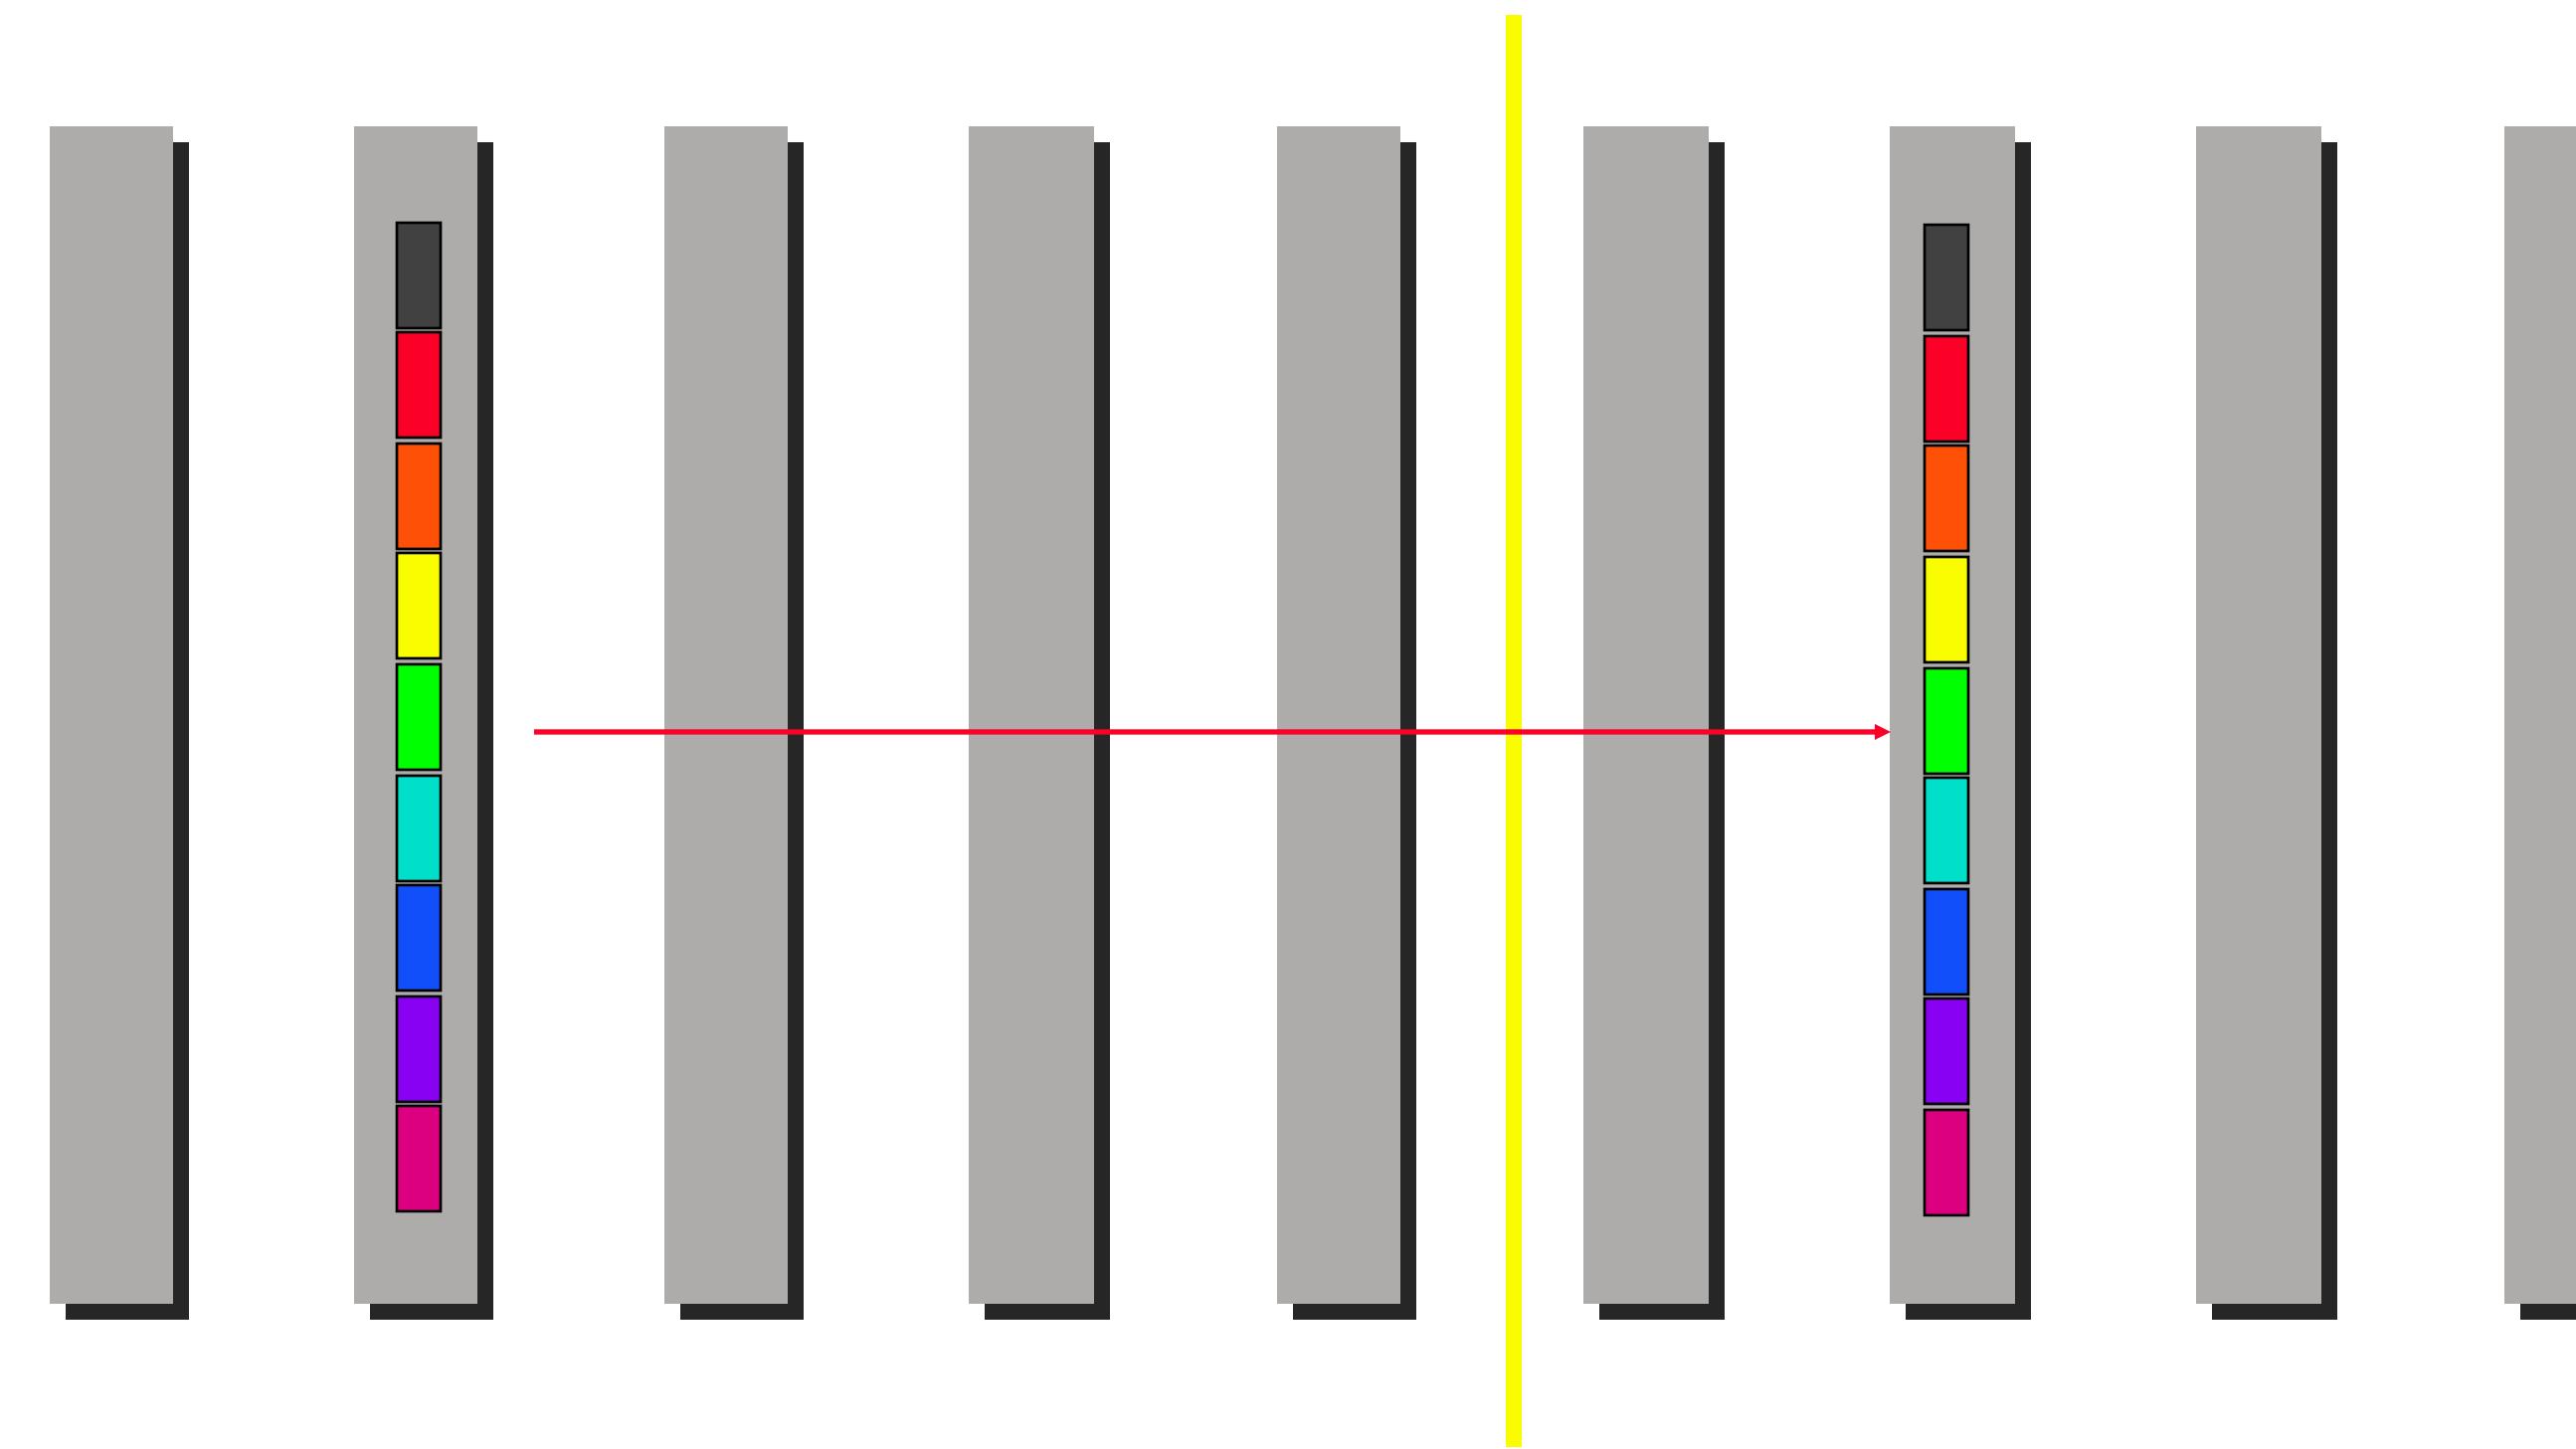
- divide logical linear array in half

General principles



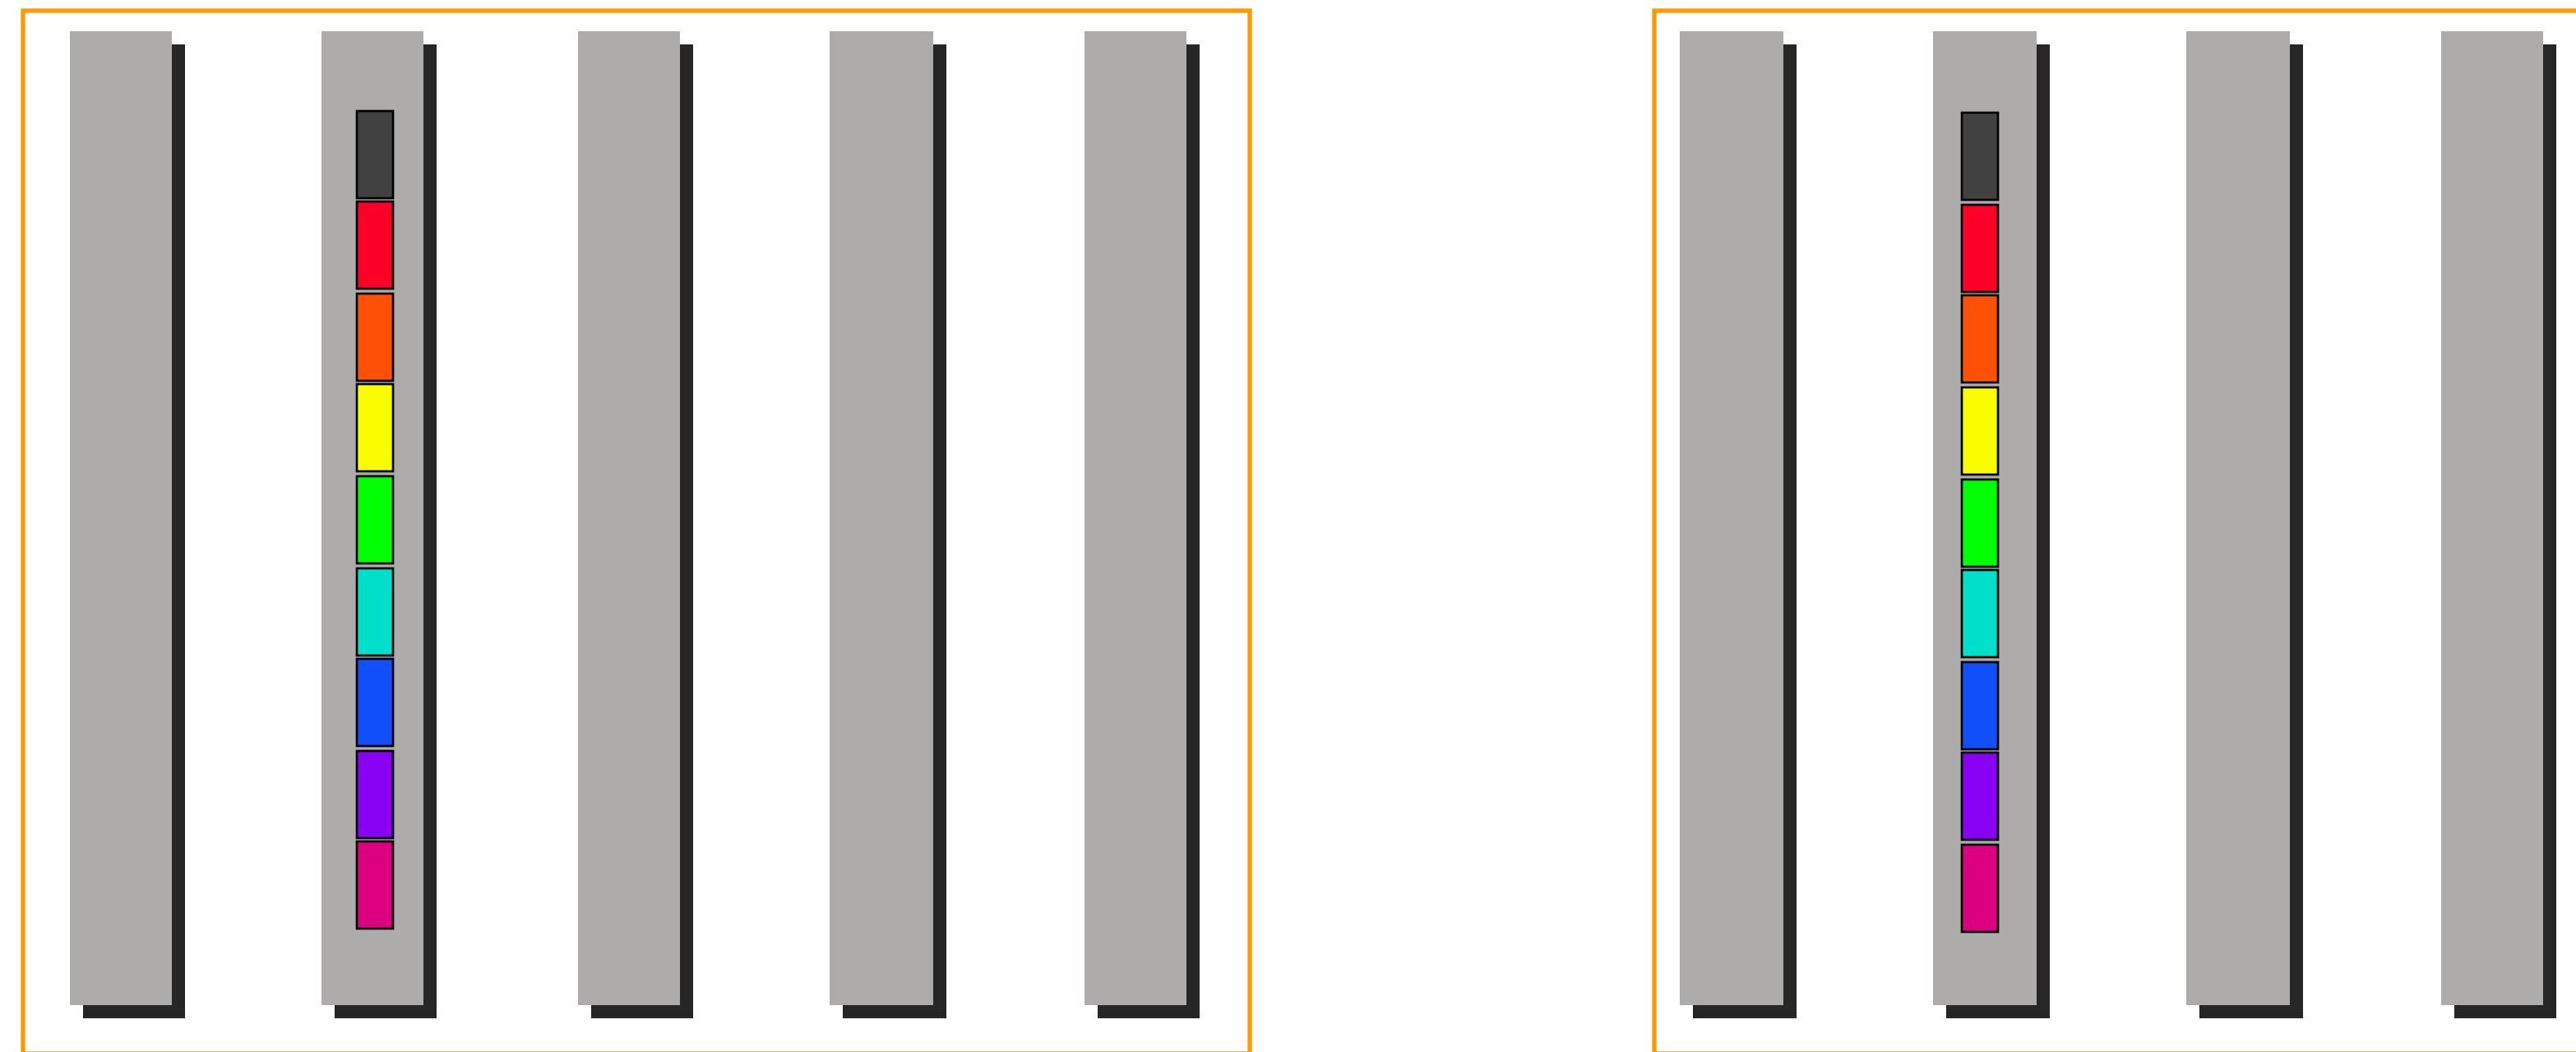
- send message to the half of the network that does not contain the current node (root) that holds the message

General principles



- send message to the half of the network that does not contain the current node (root) that holds the message

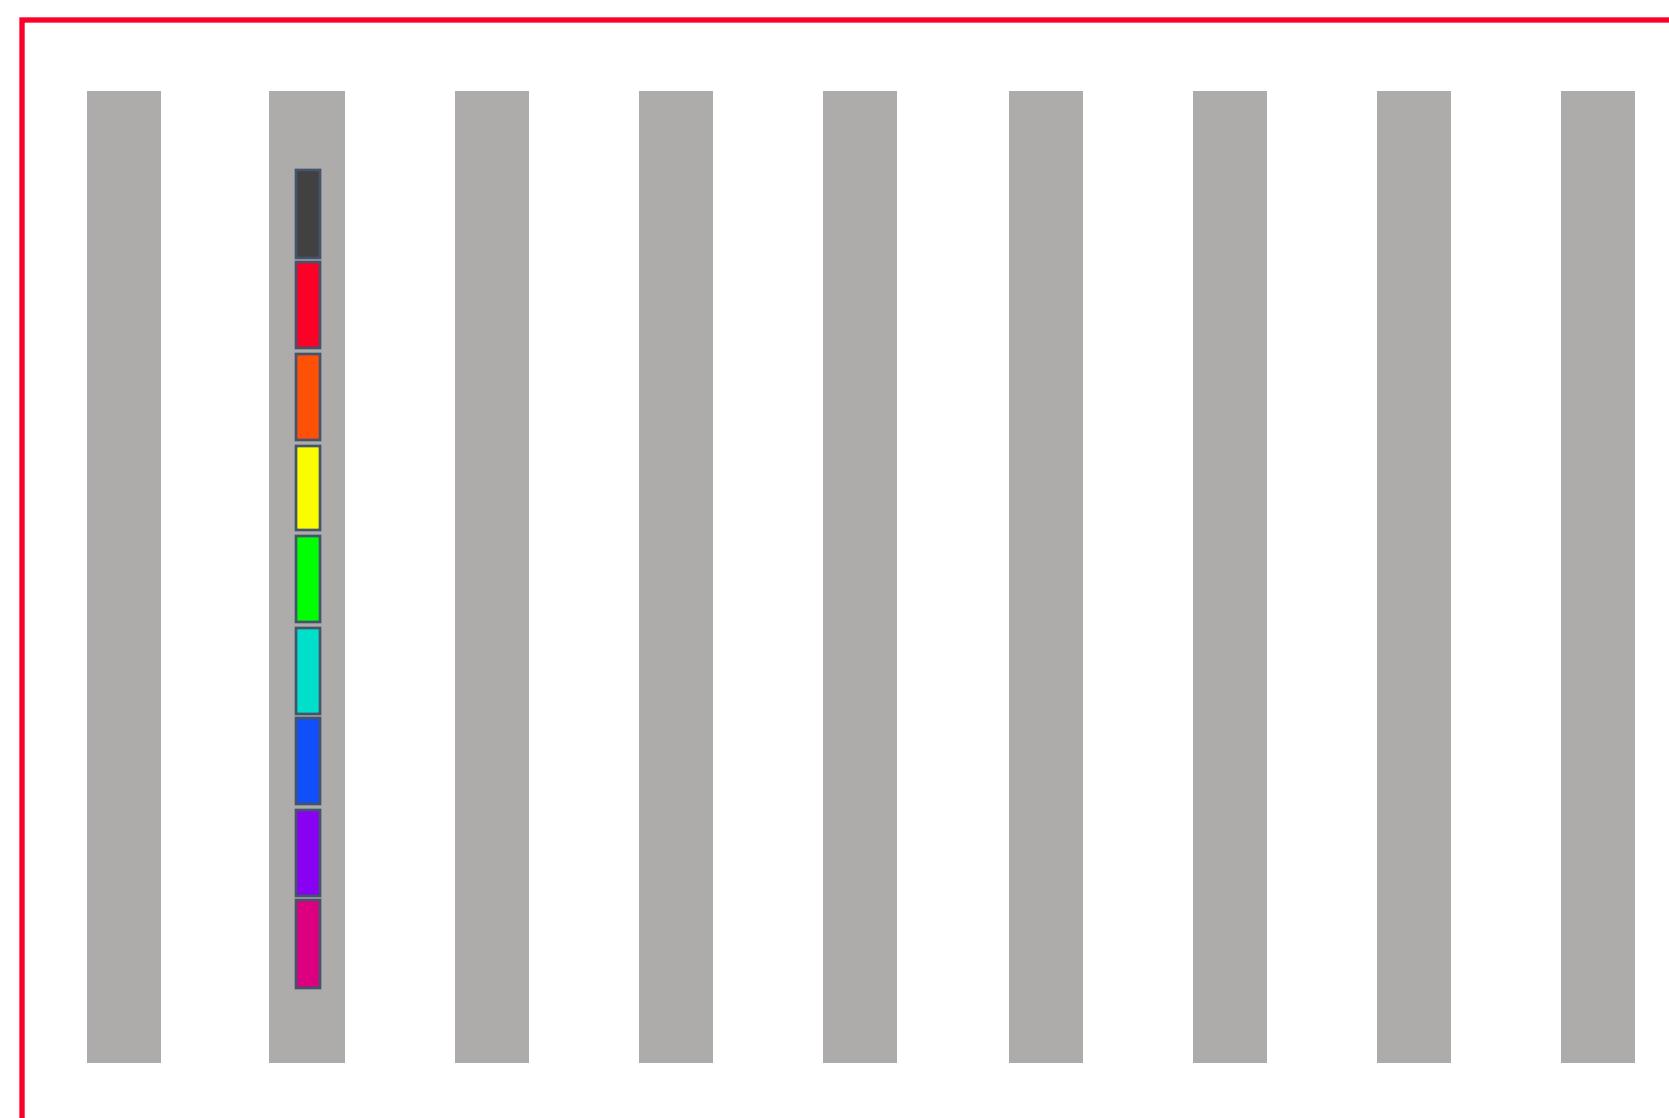
General principles



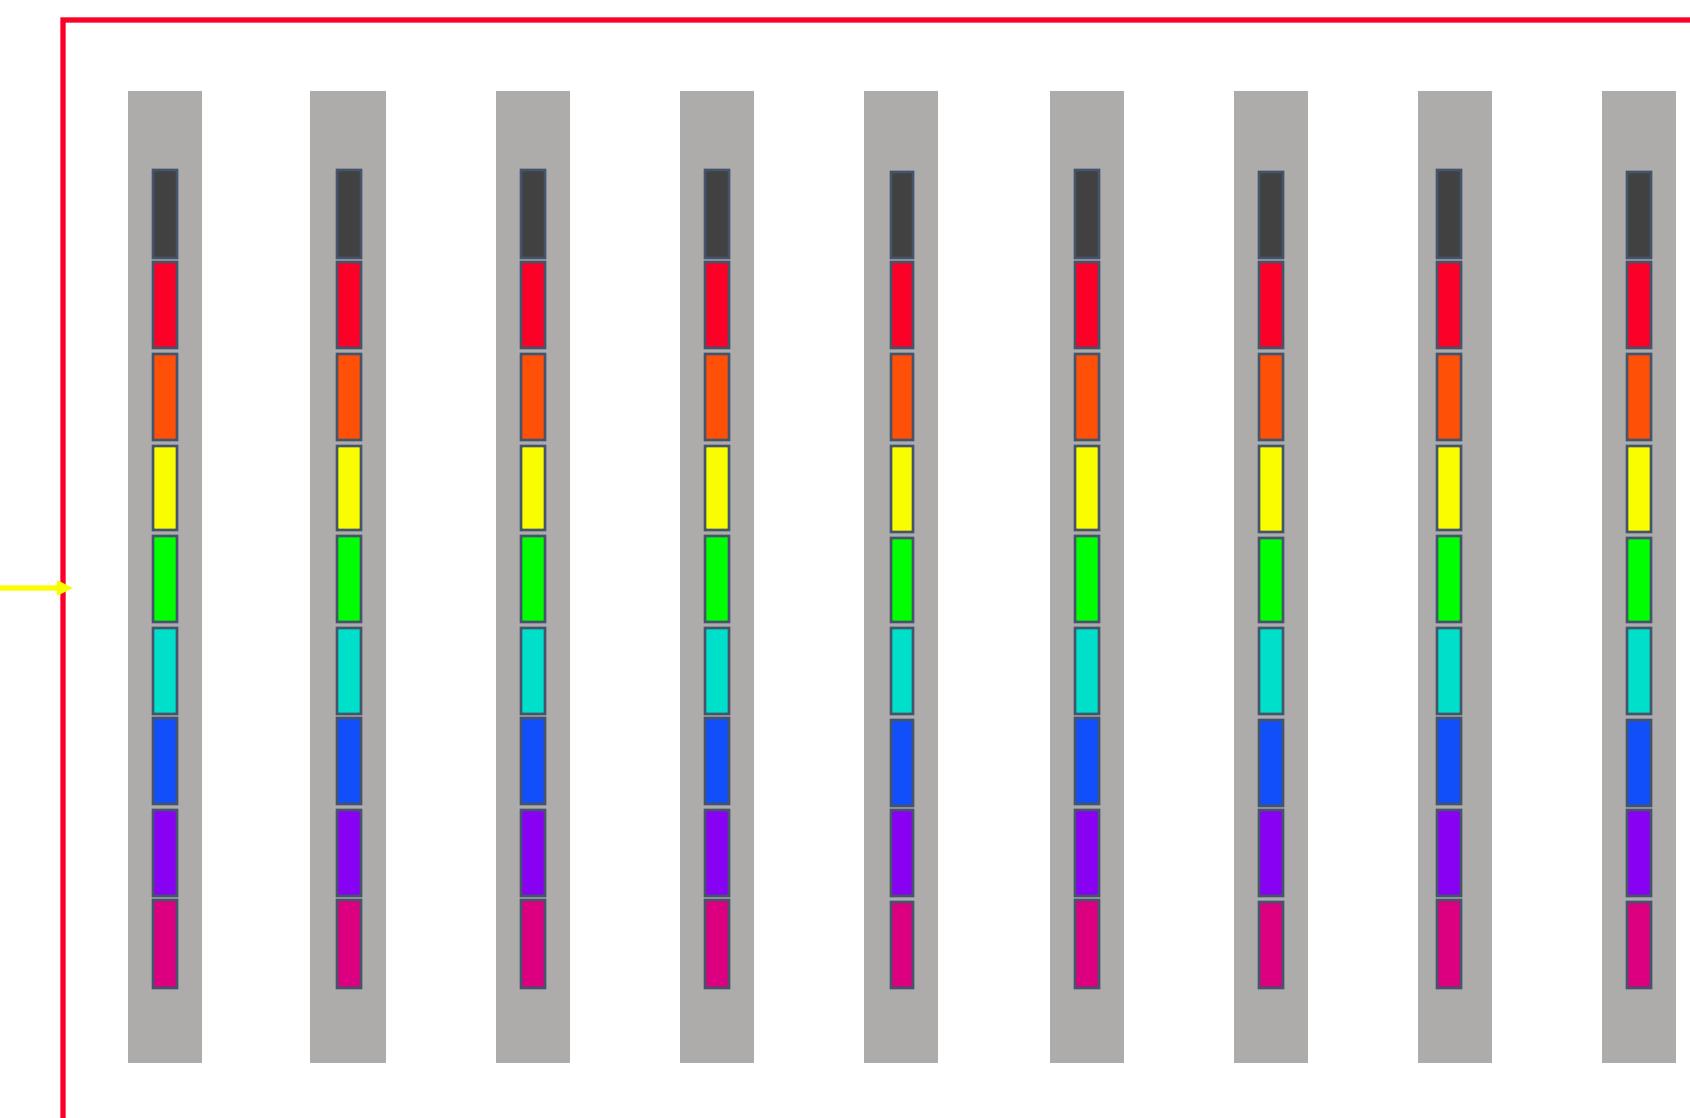
- continue recursively in each of the two halves

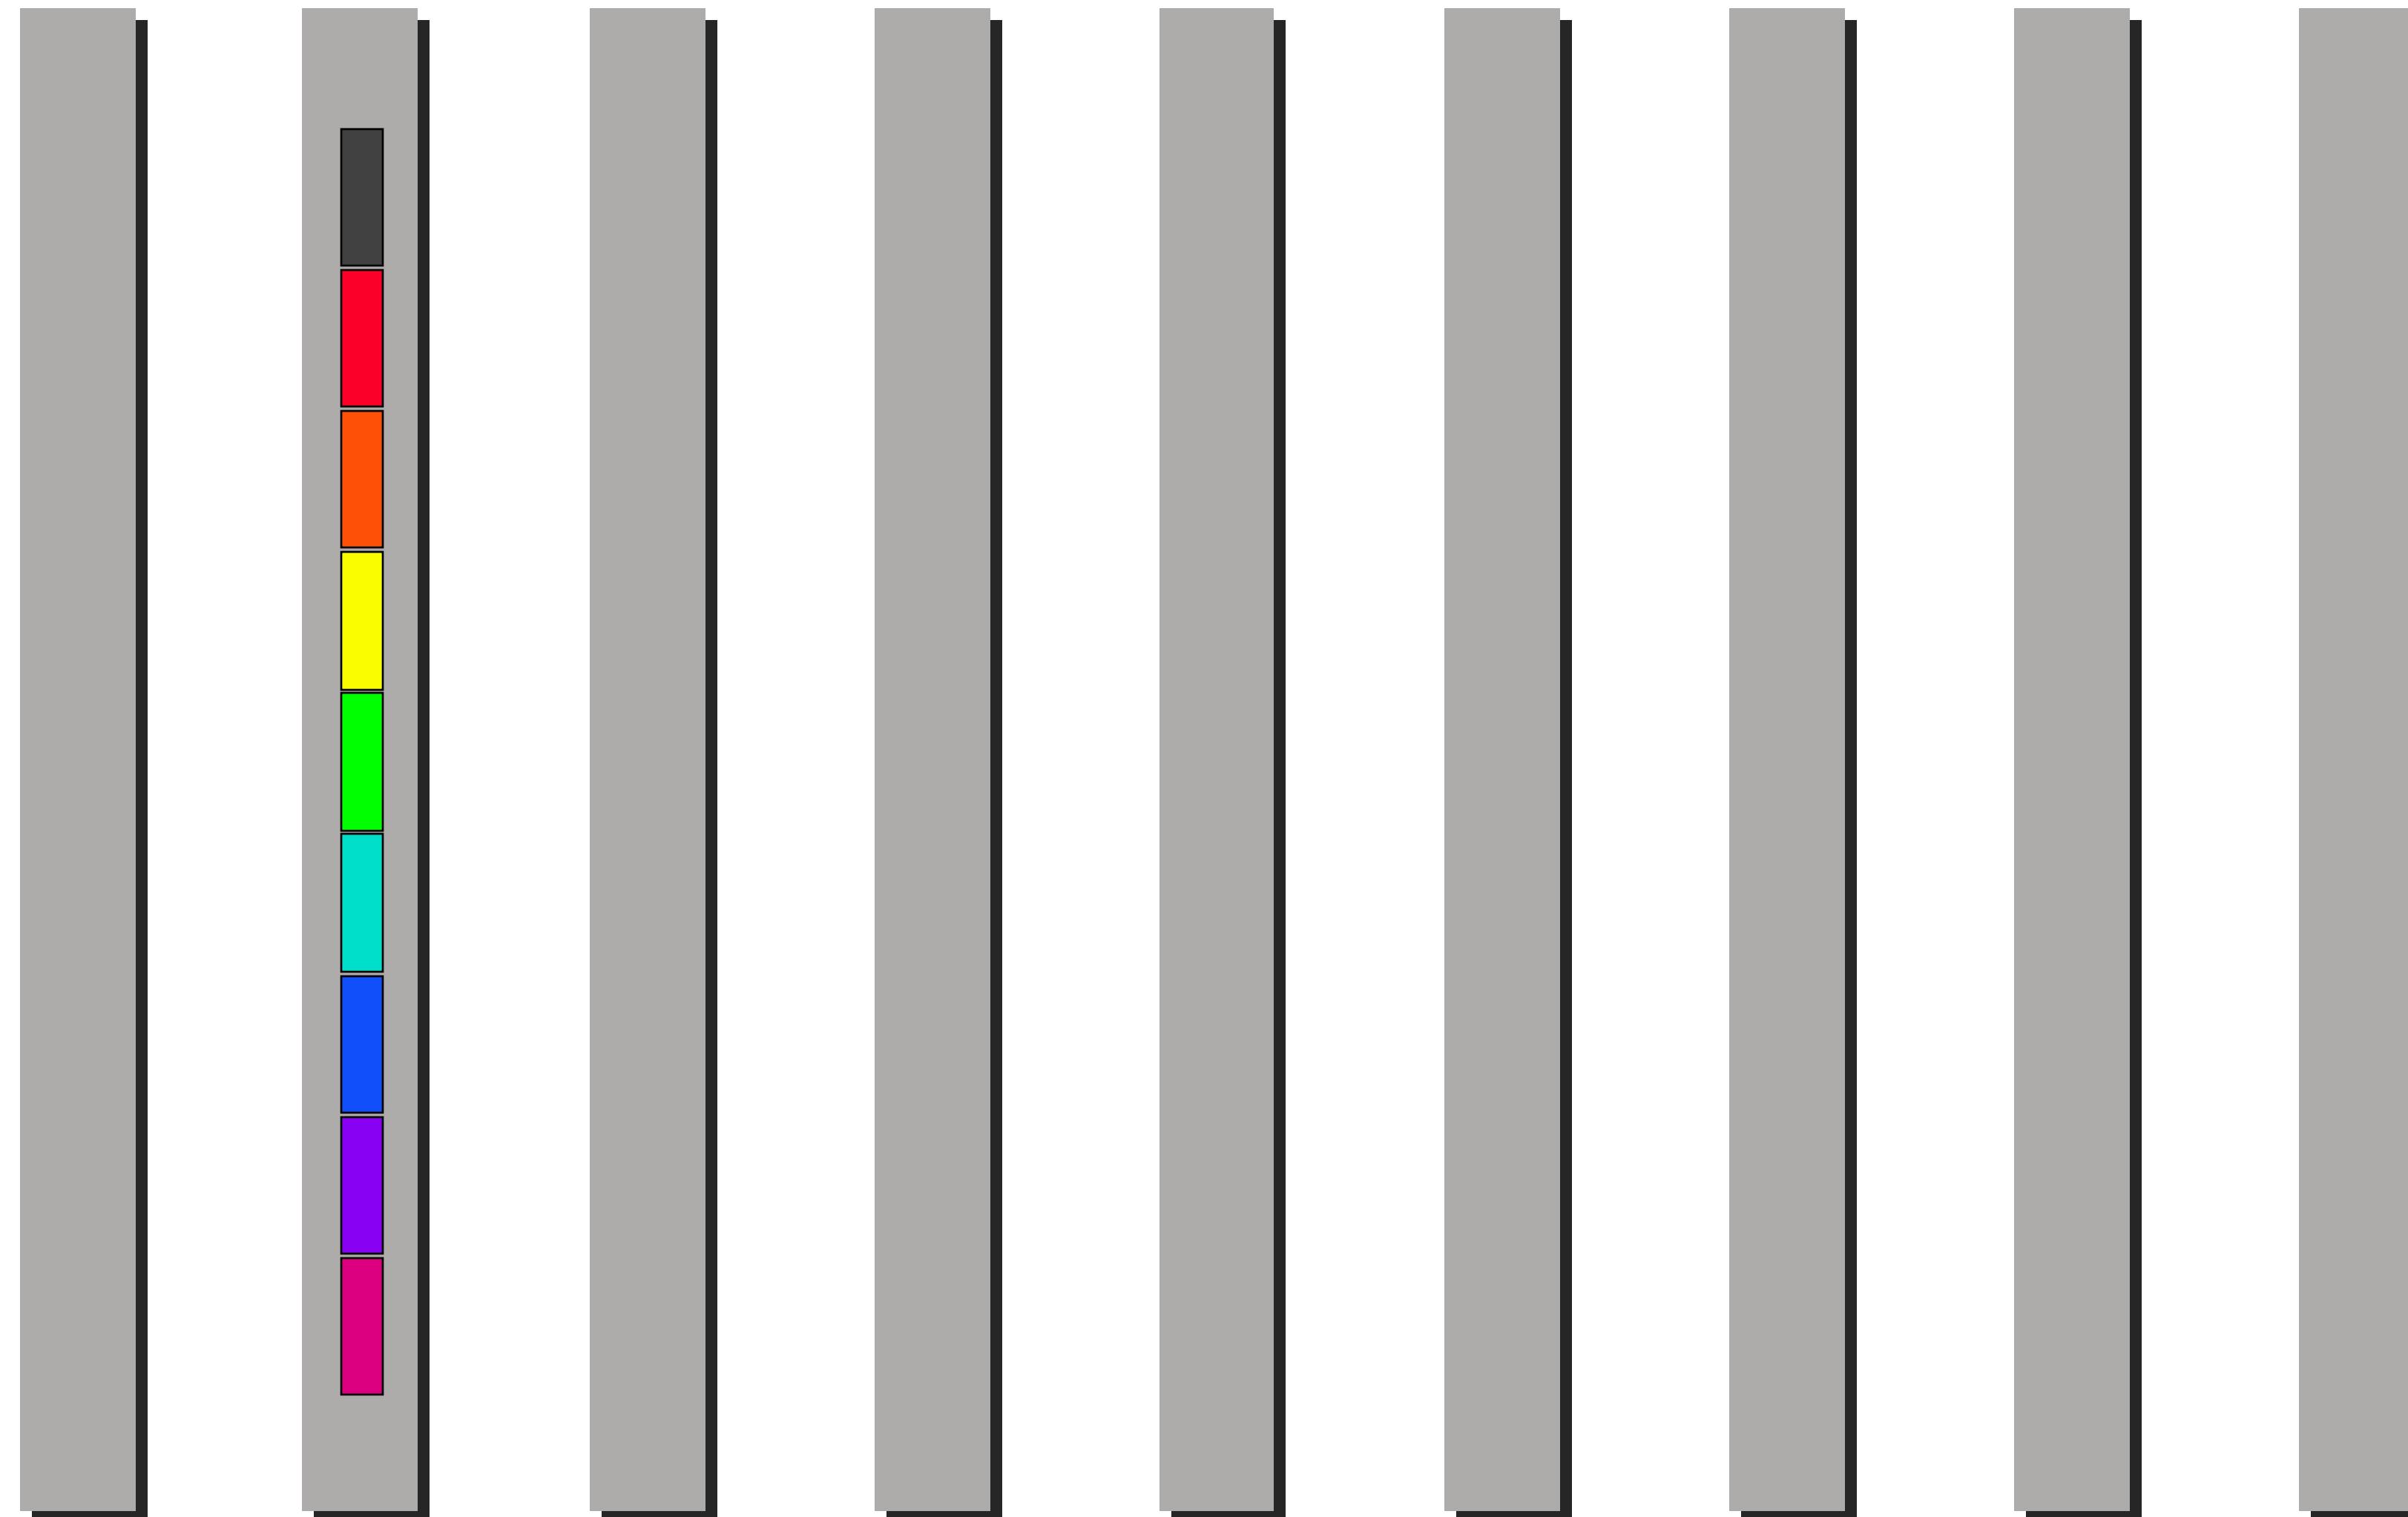
Broadcast

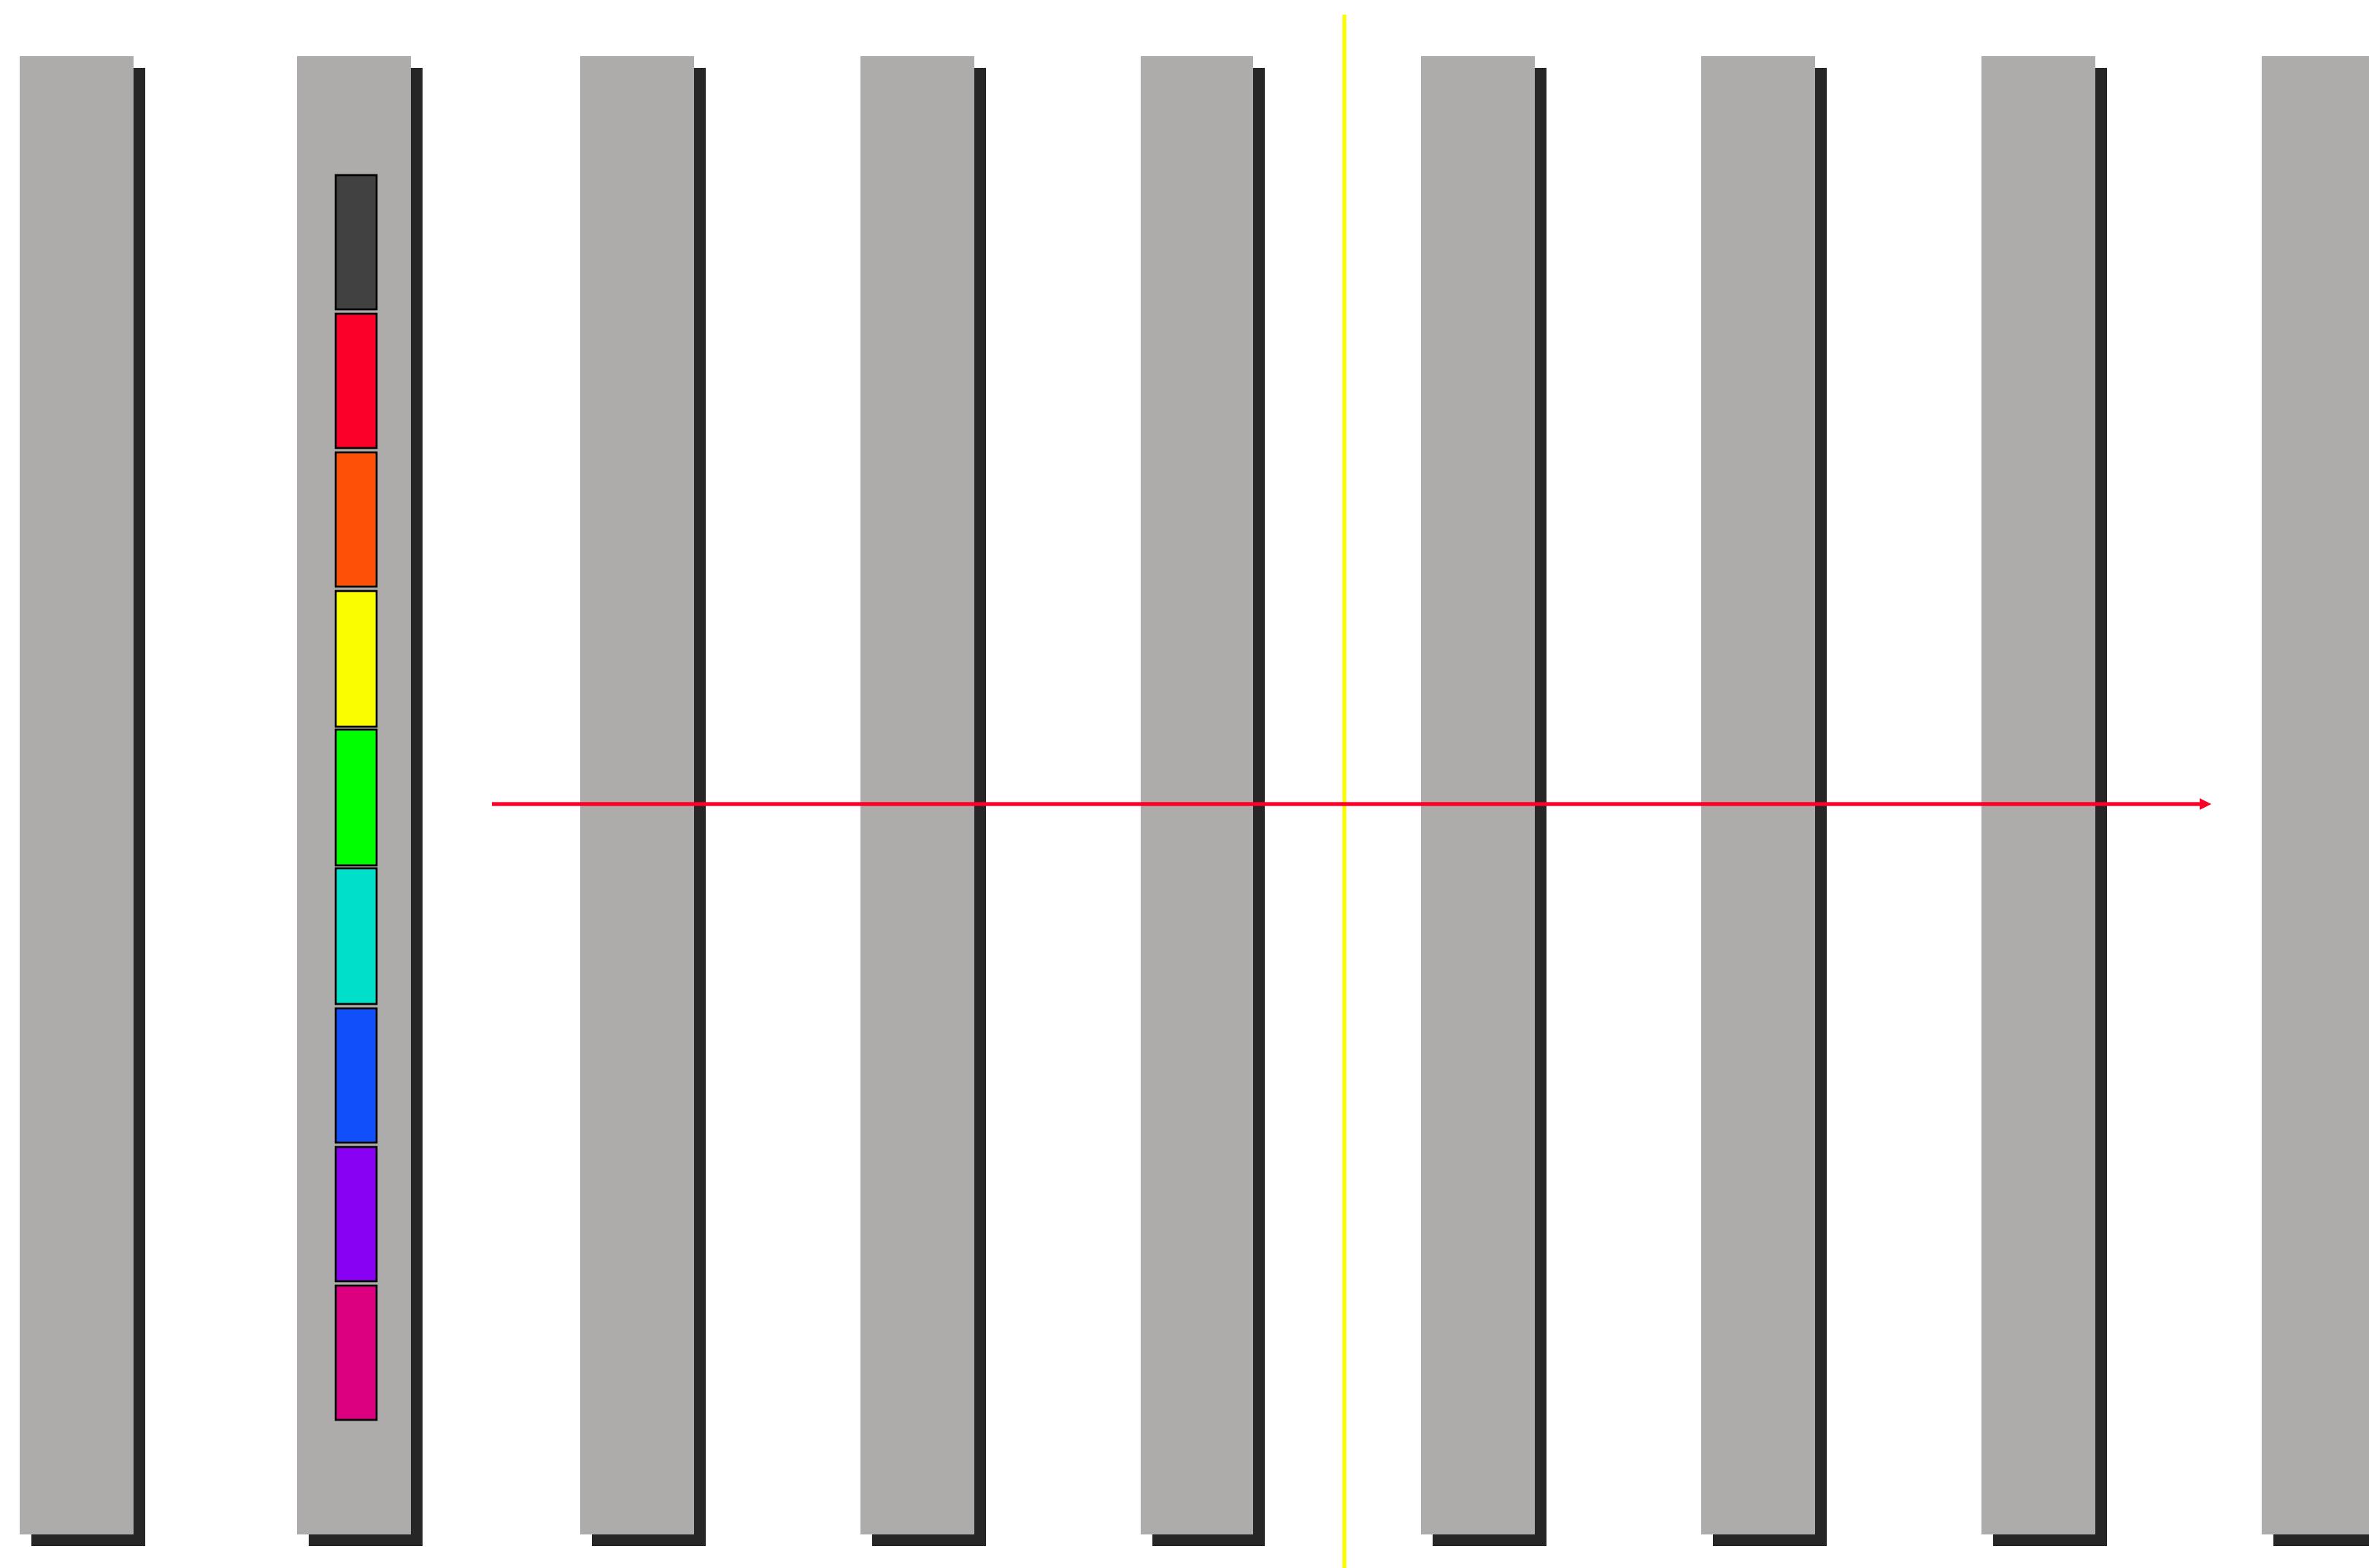
Before

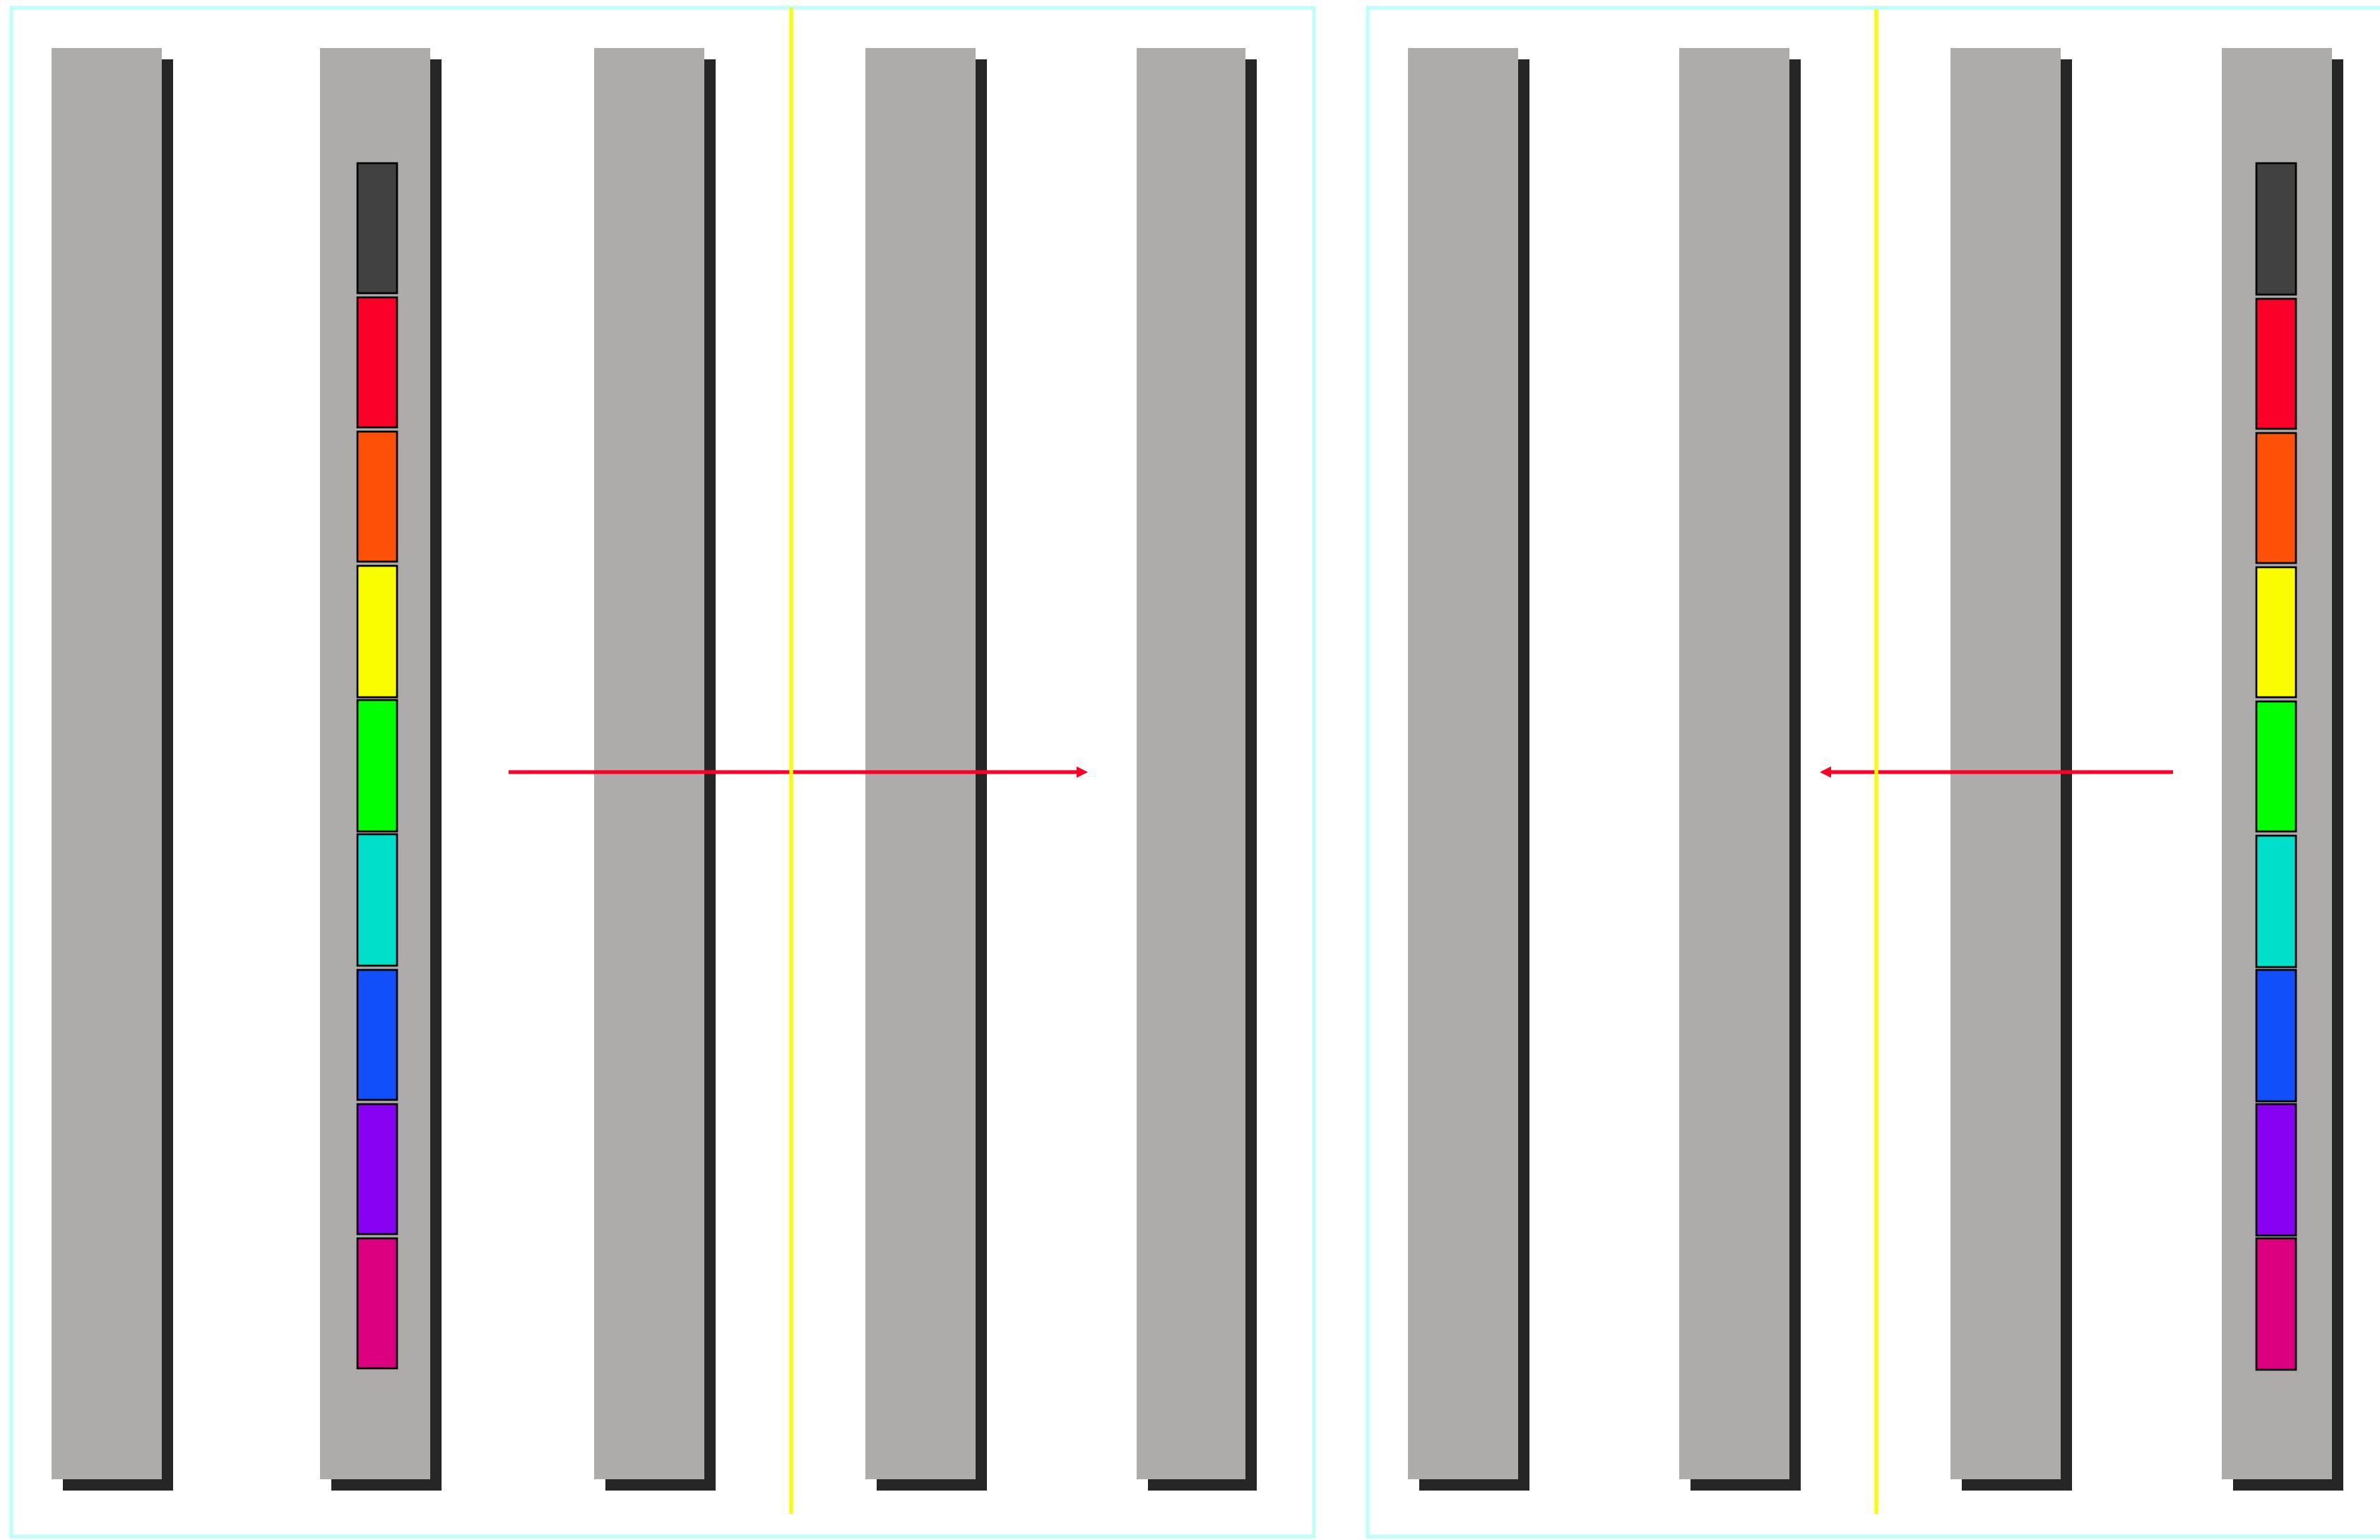


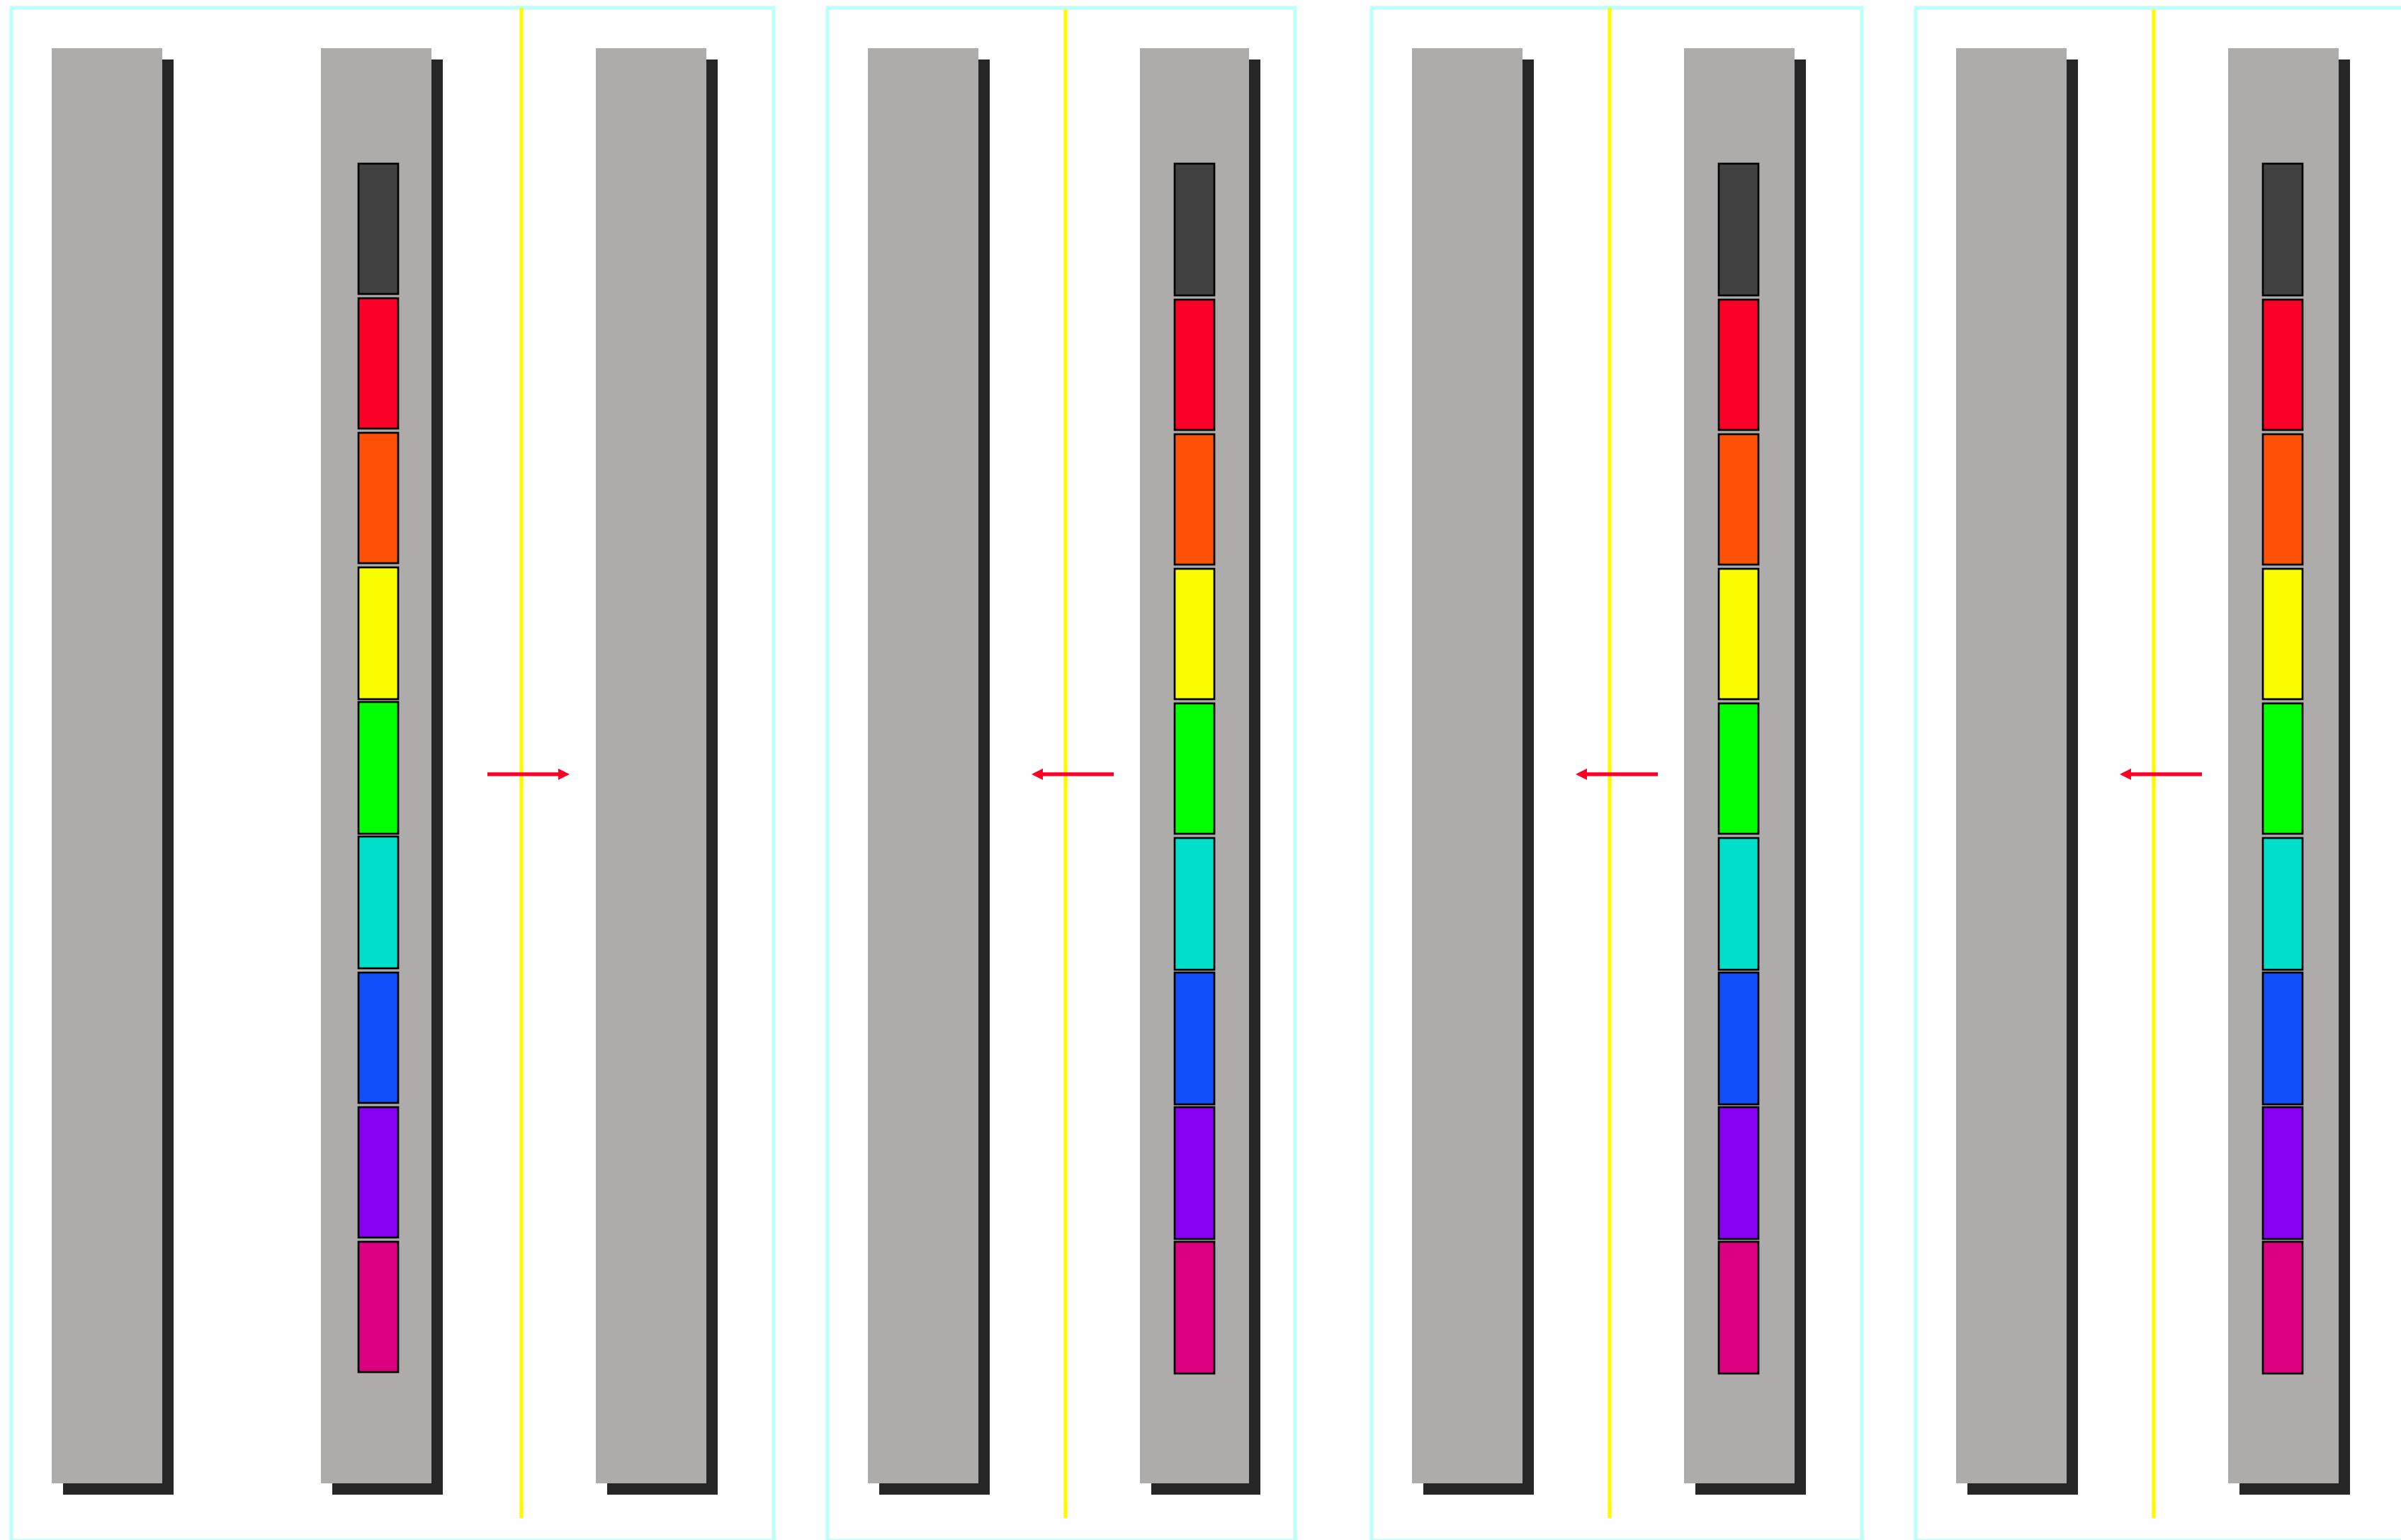
After

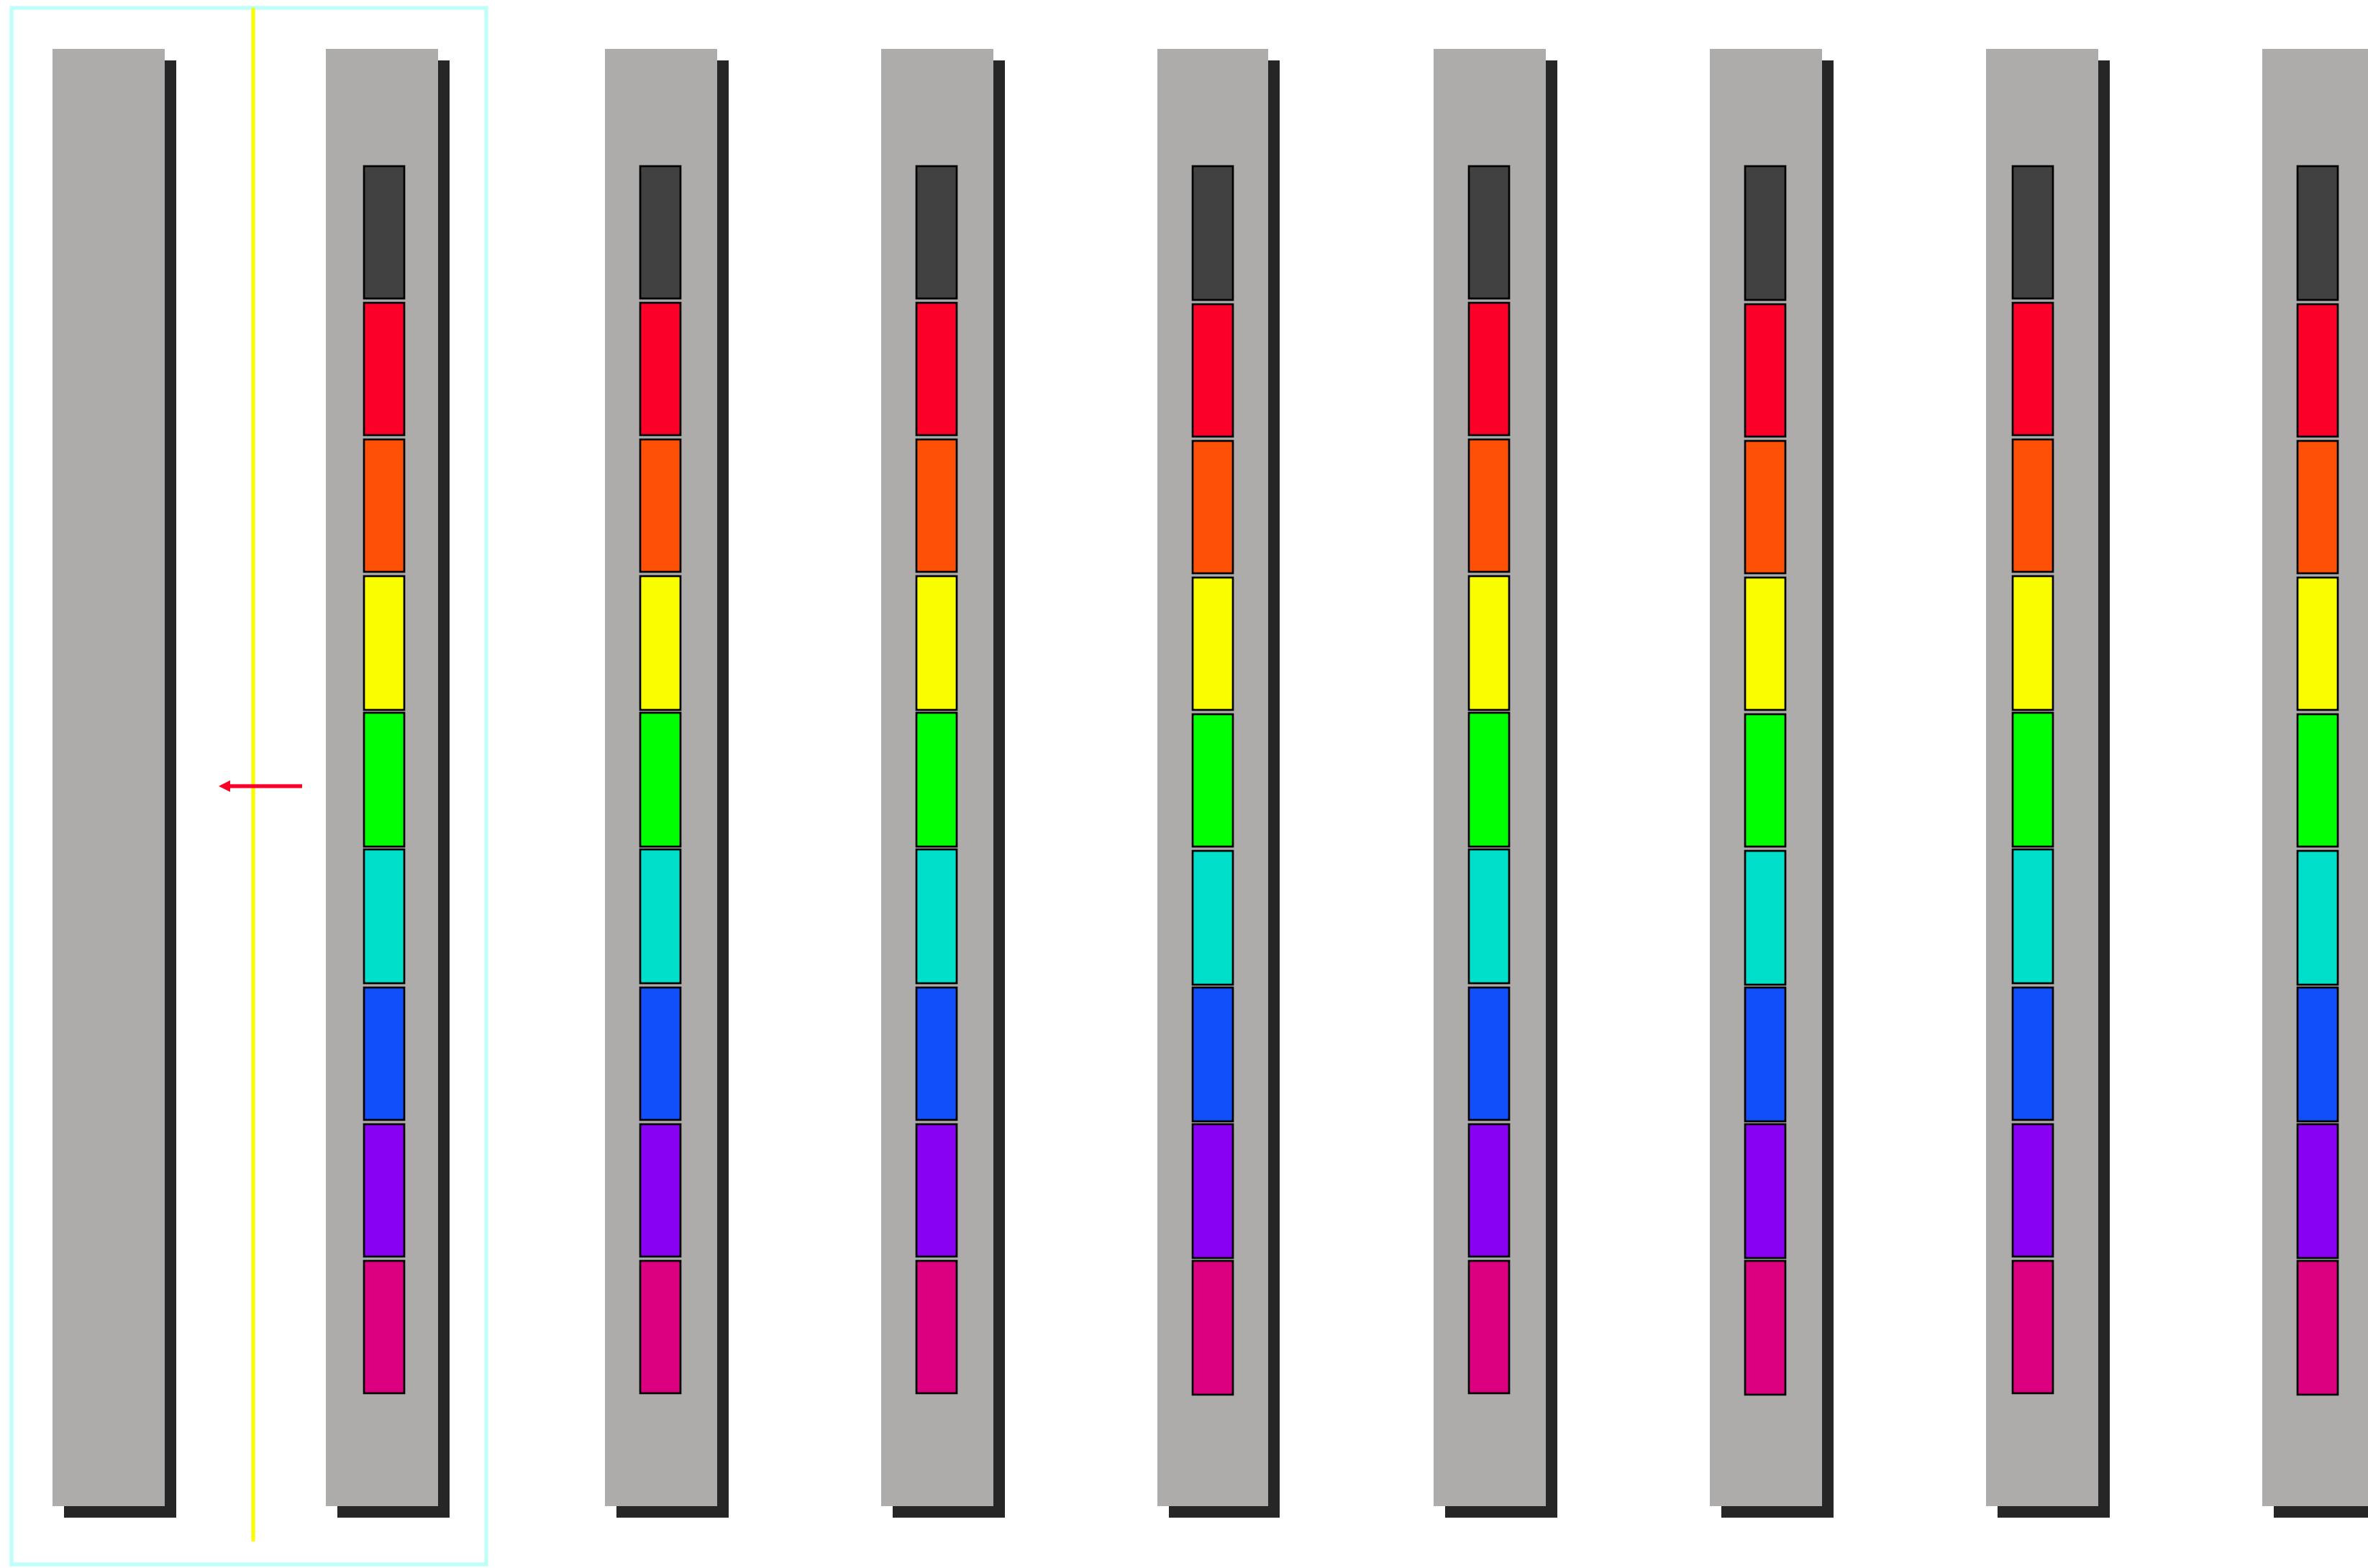


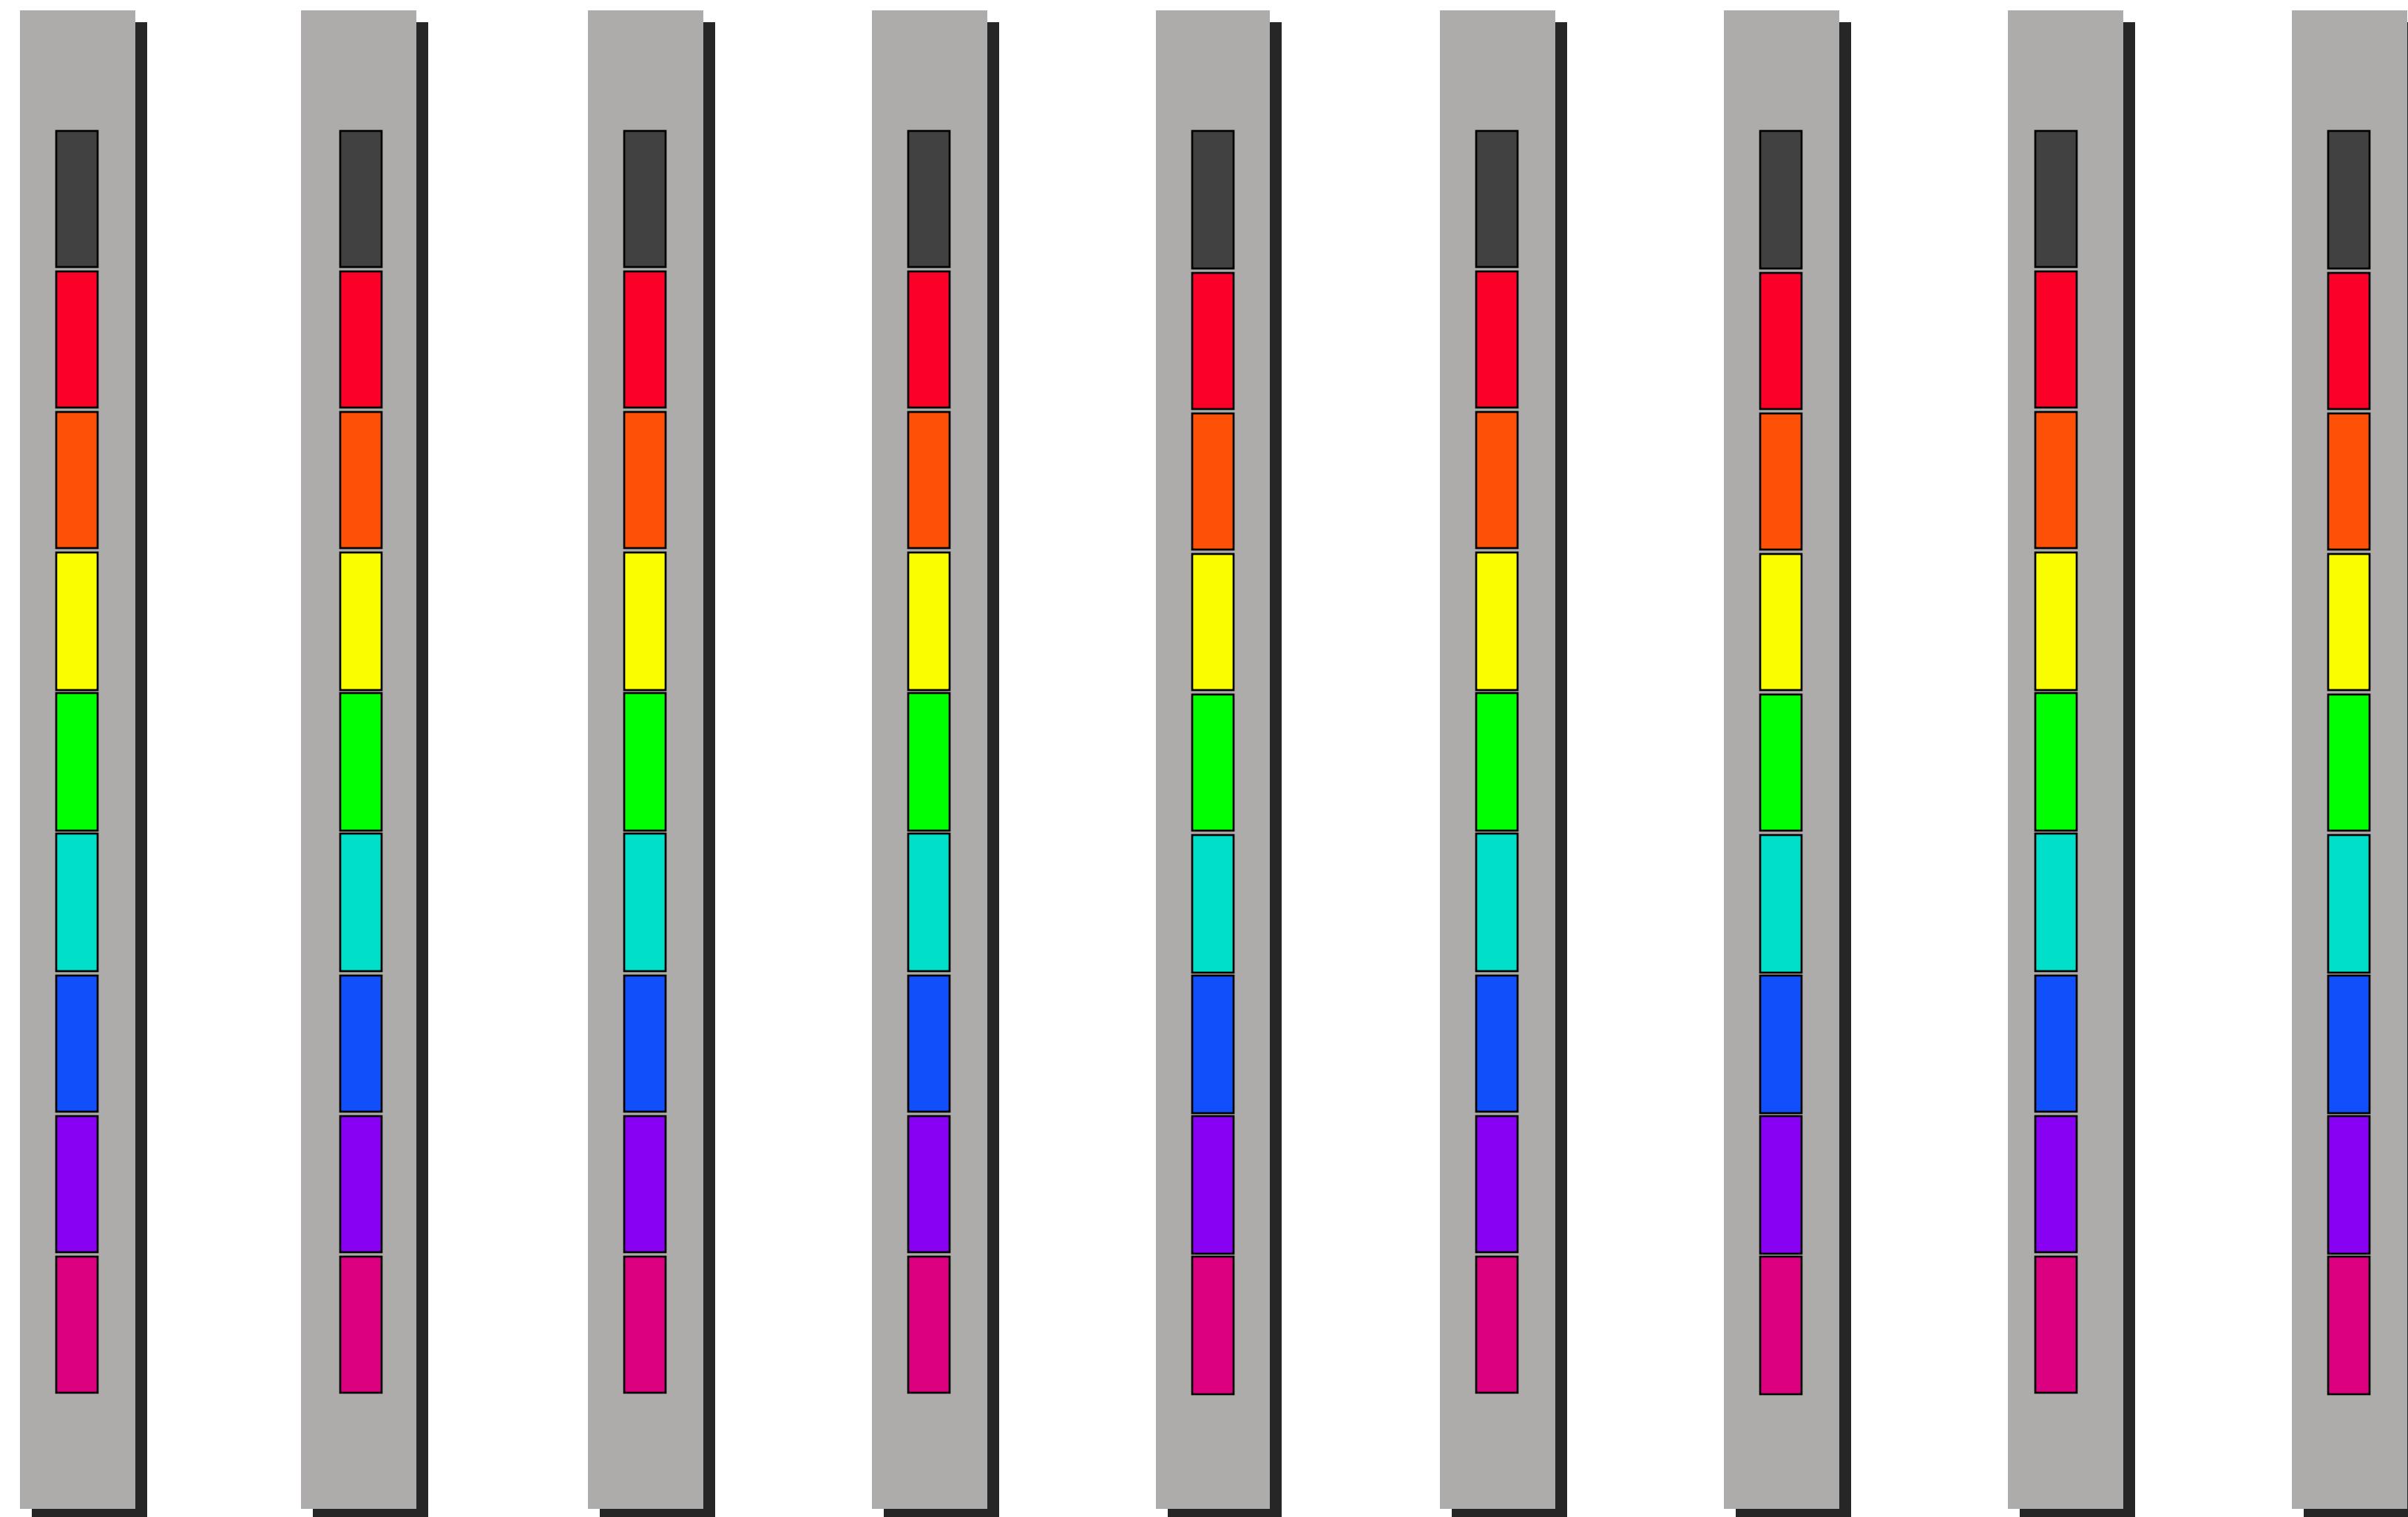






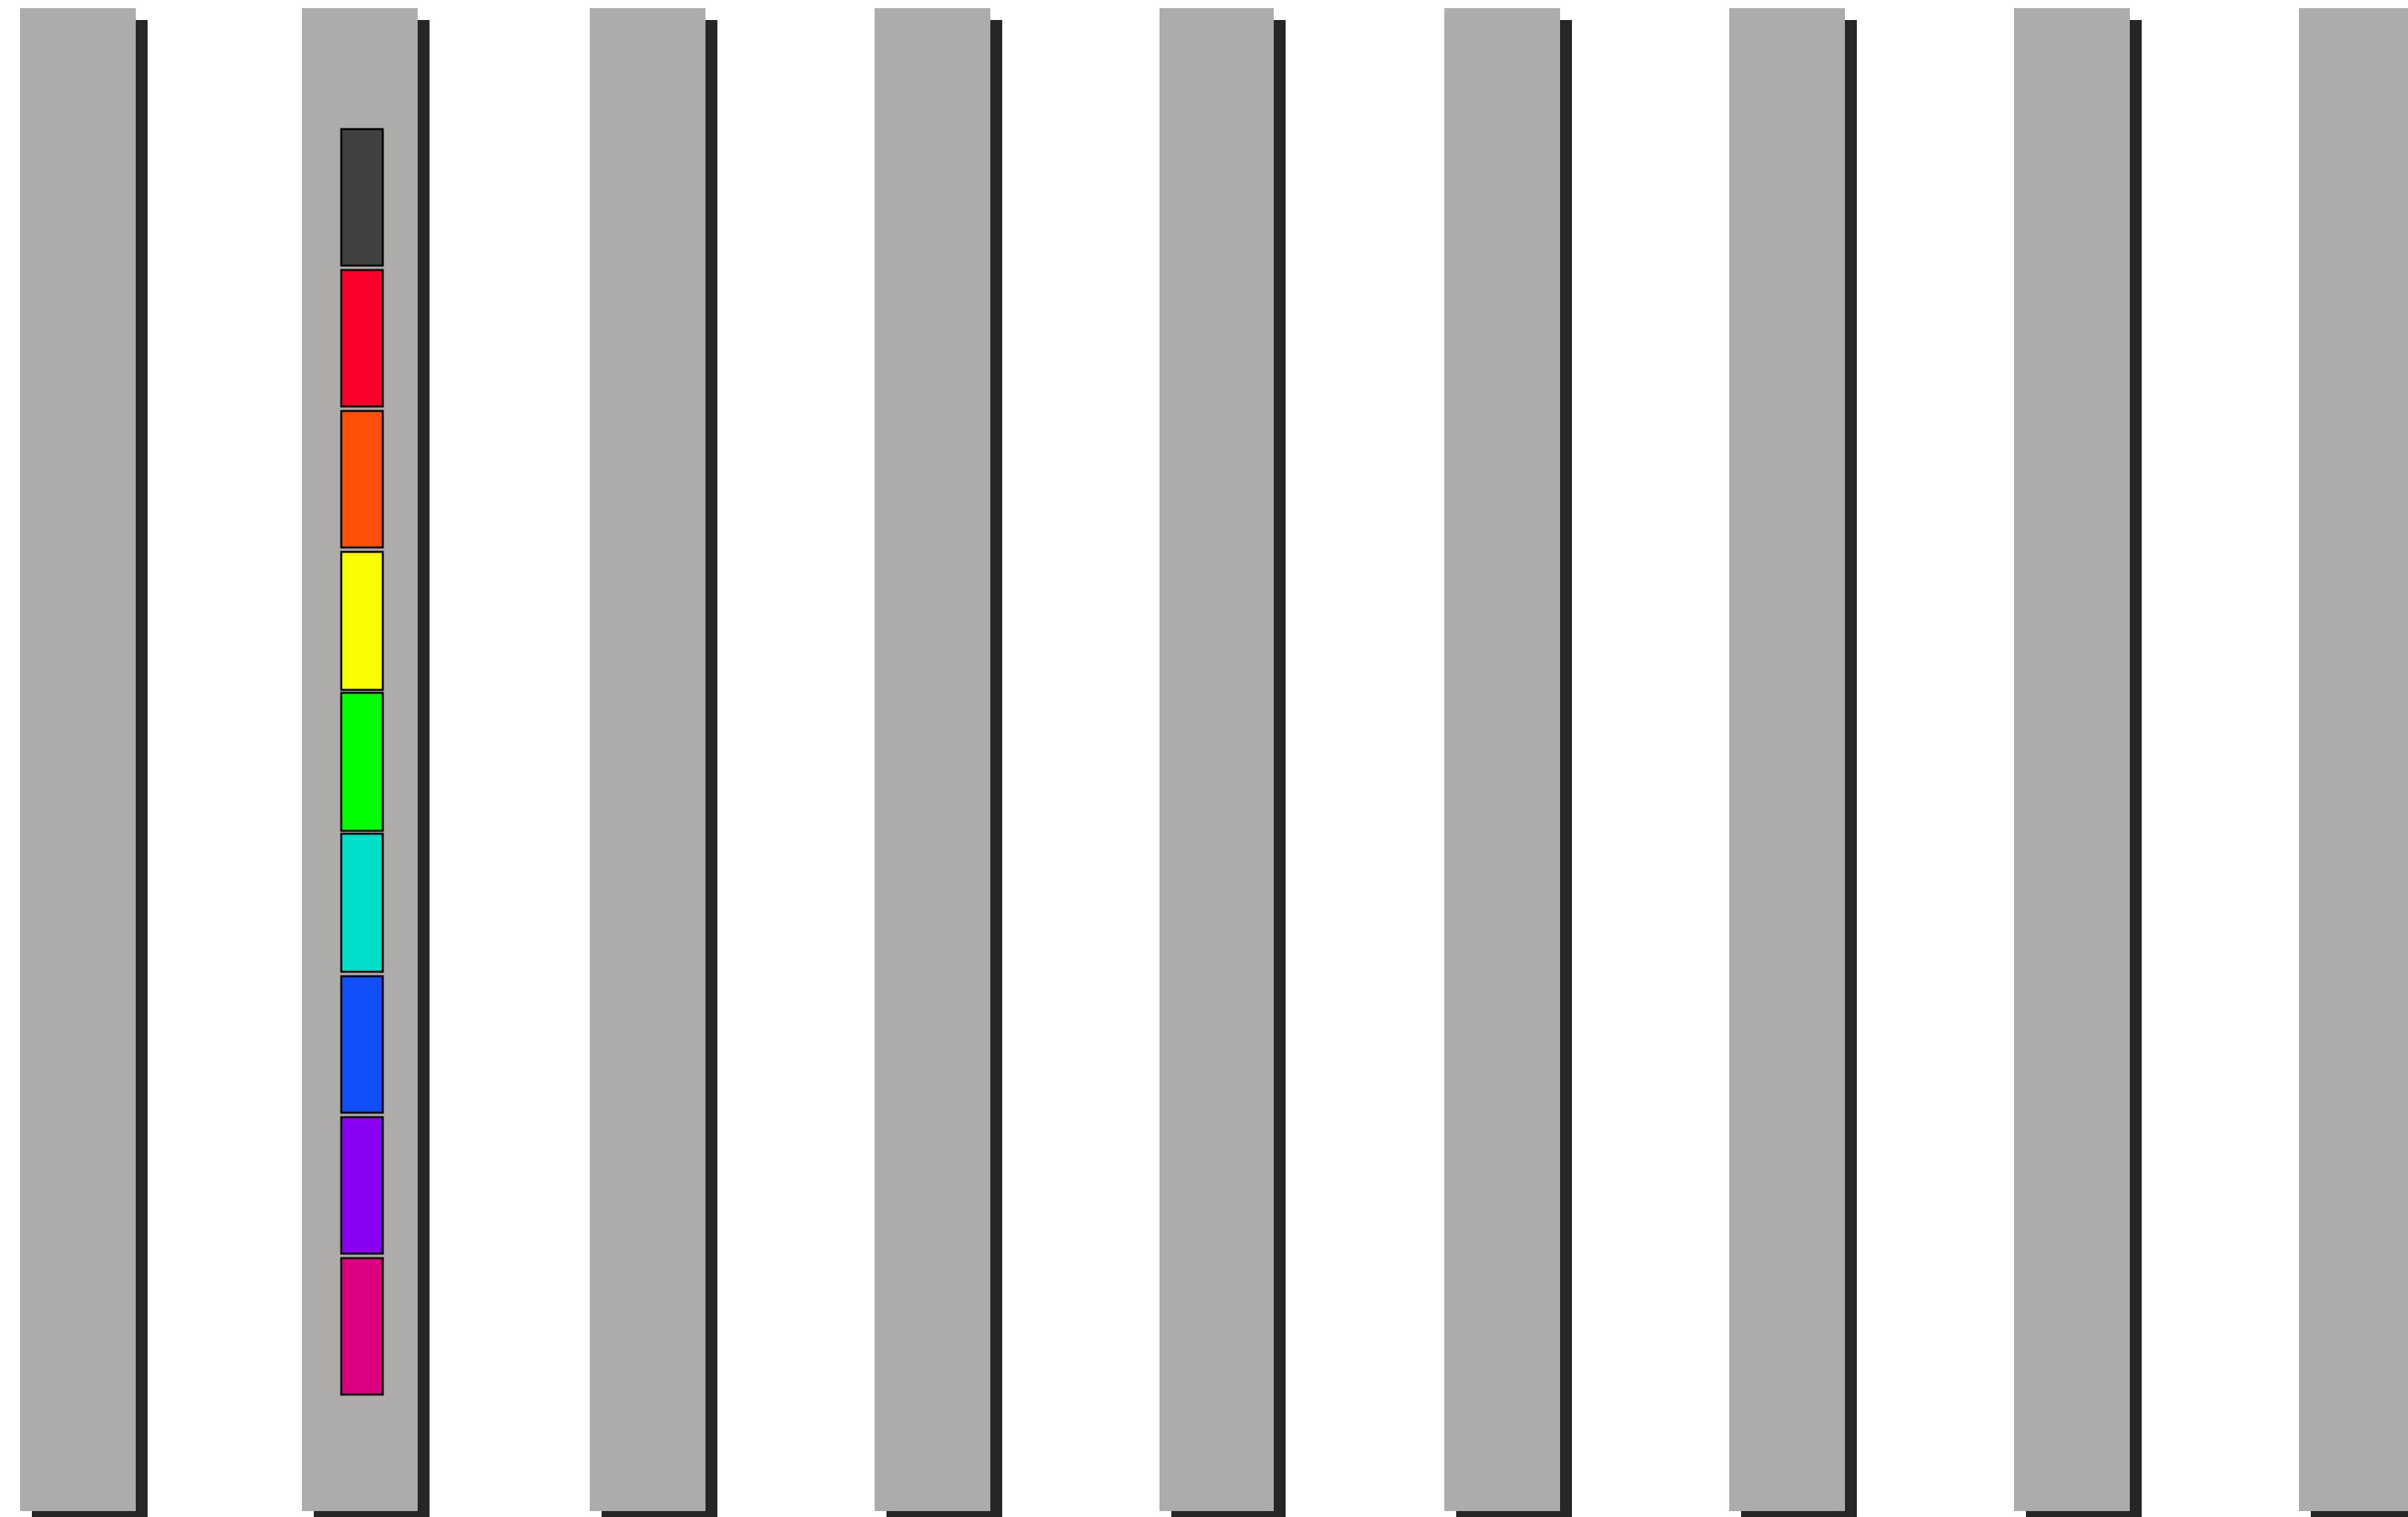


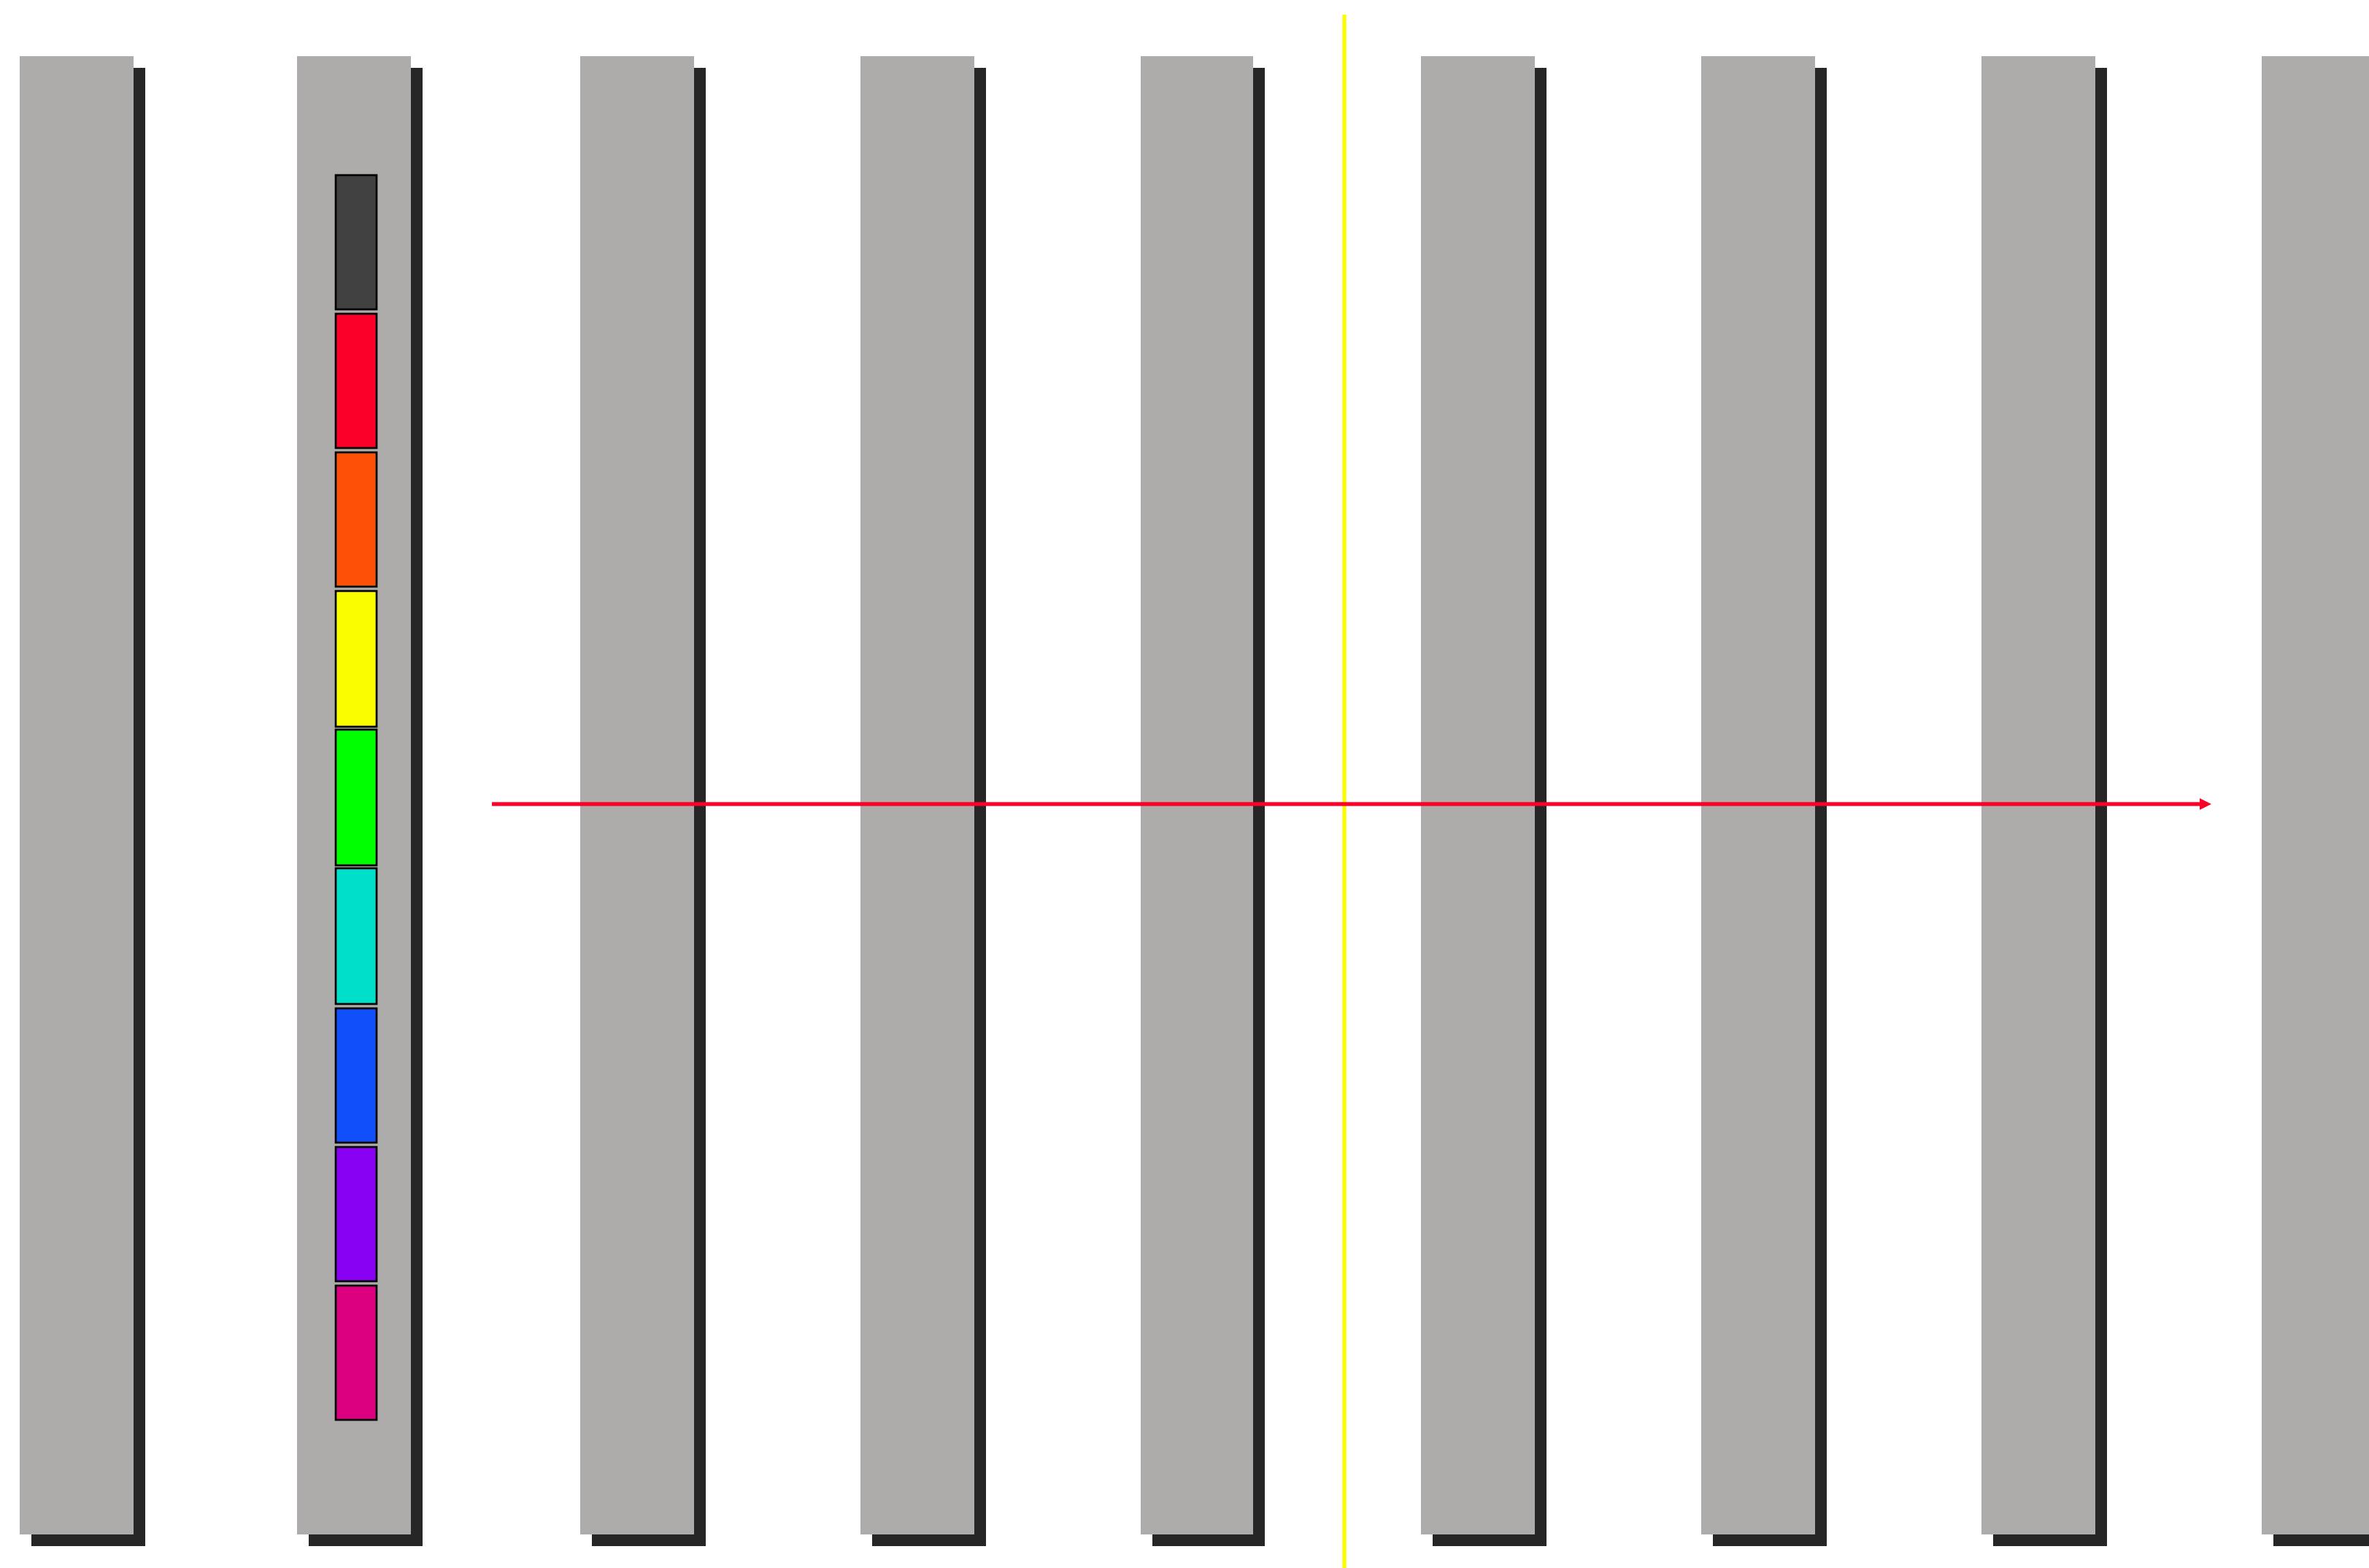


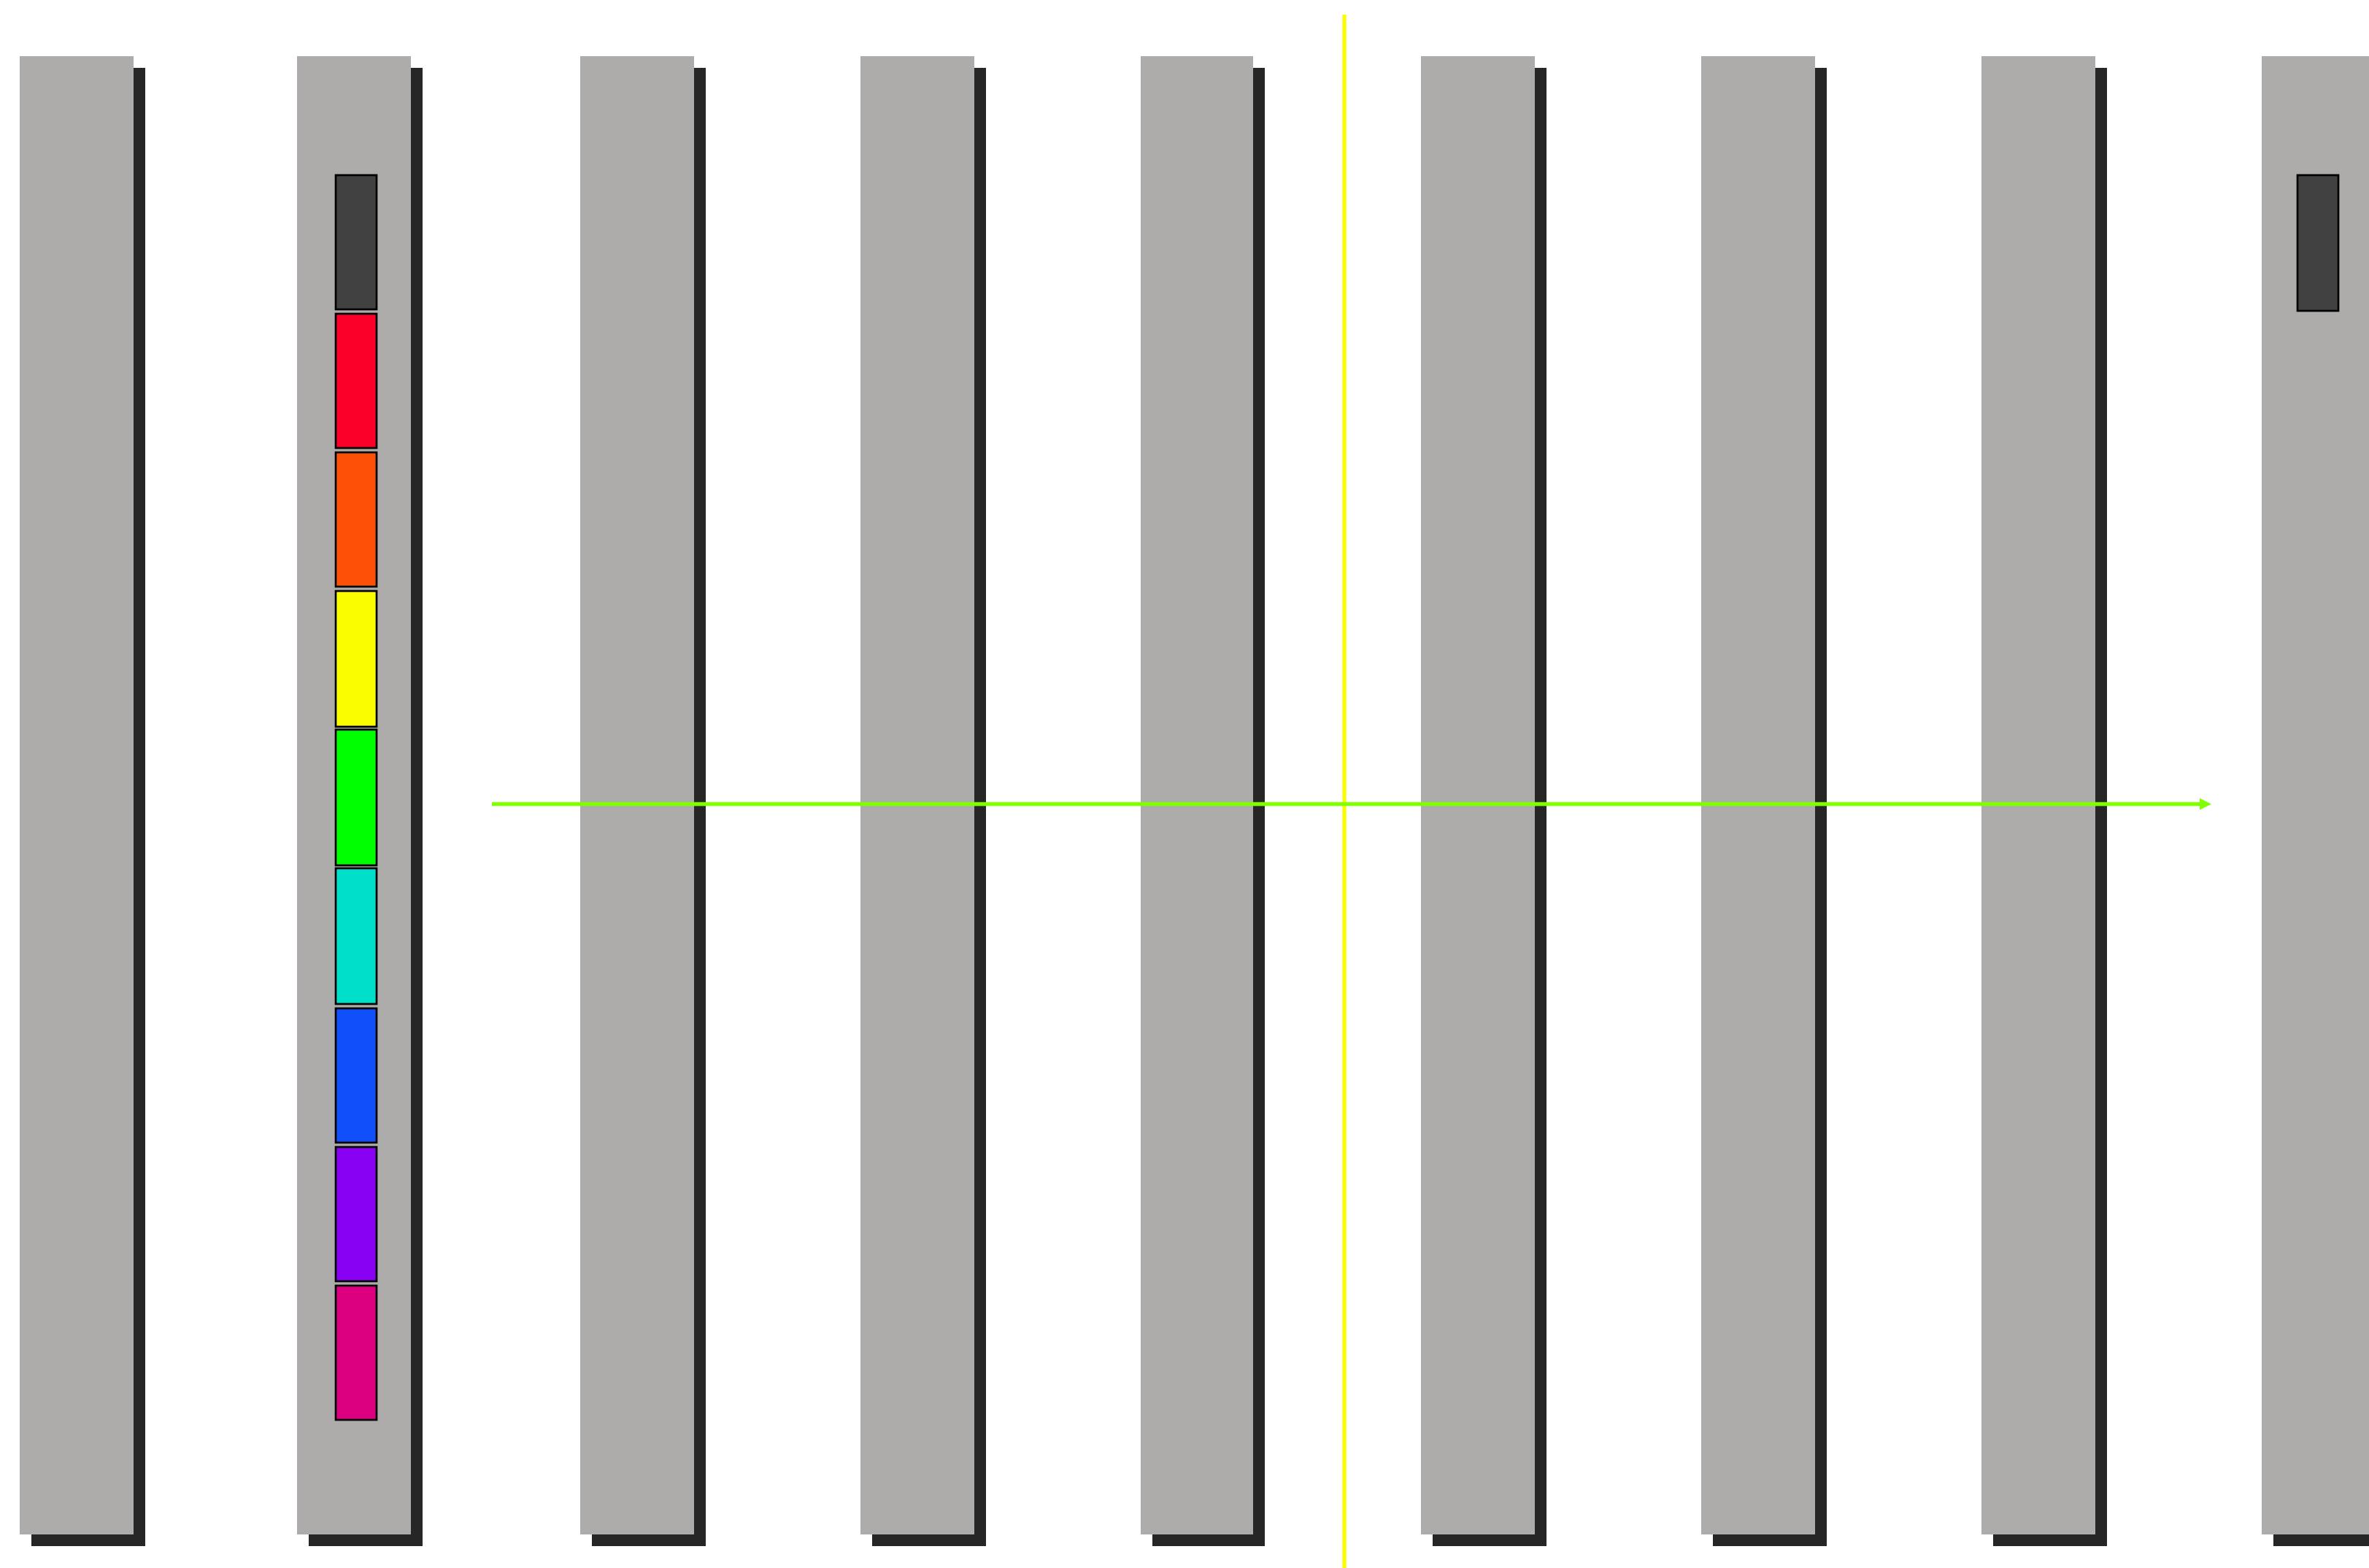


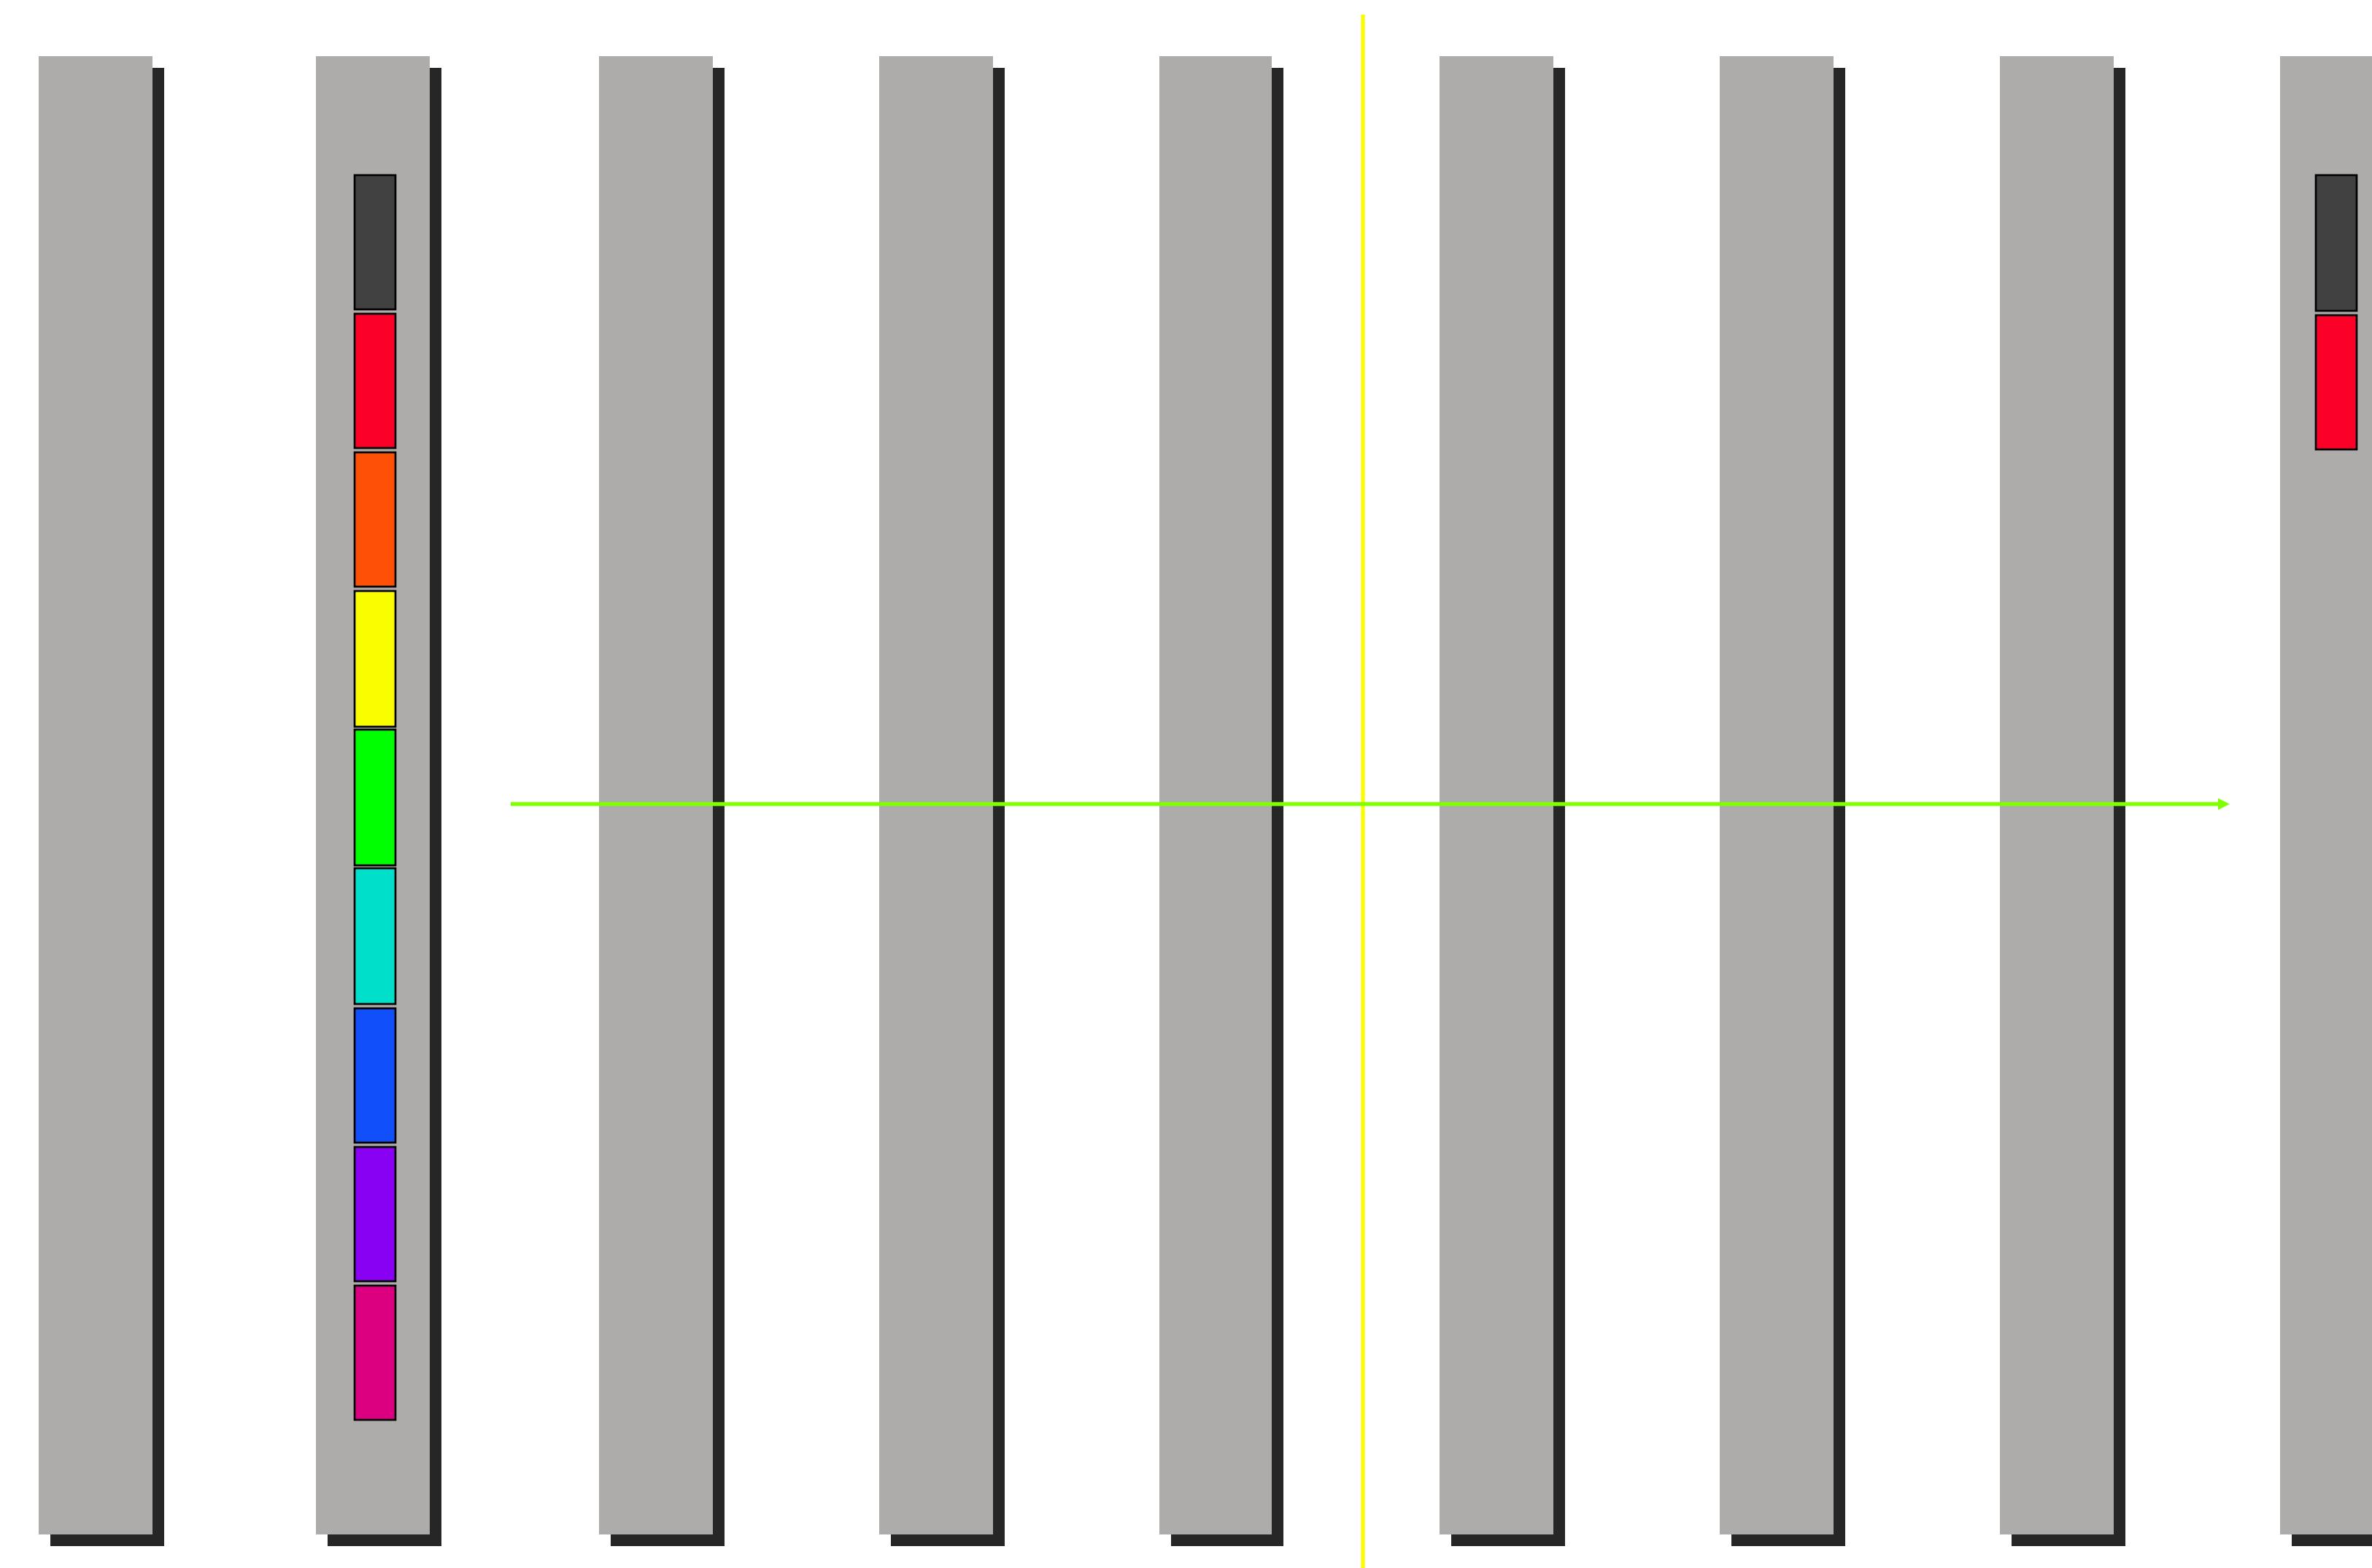
Let us view this more closely

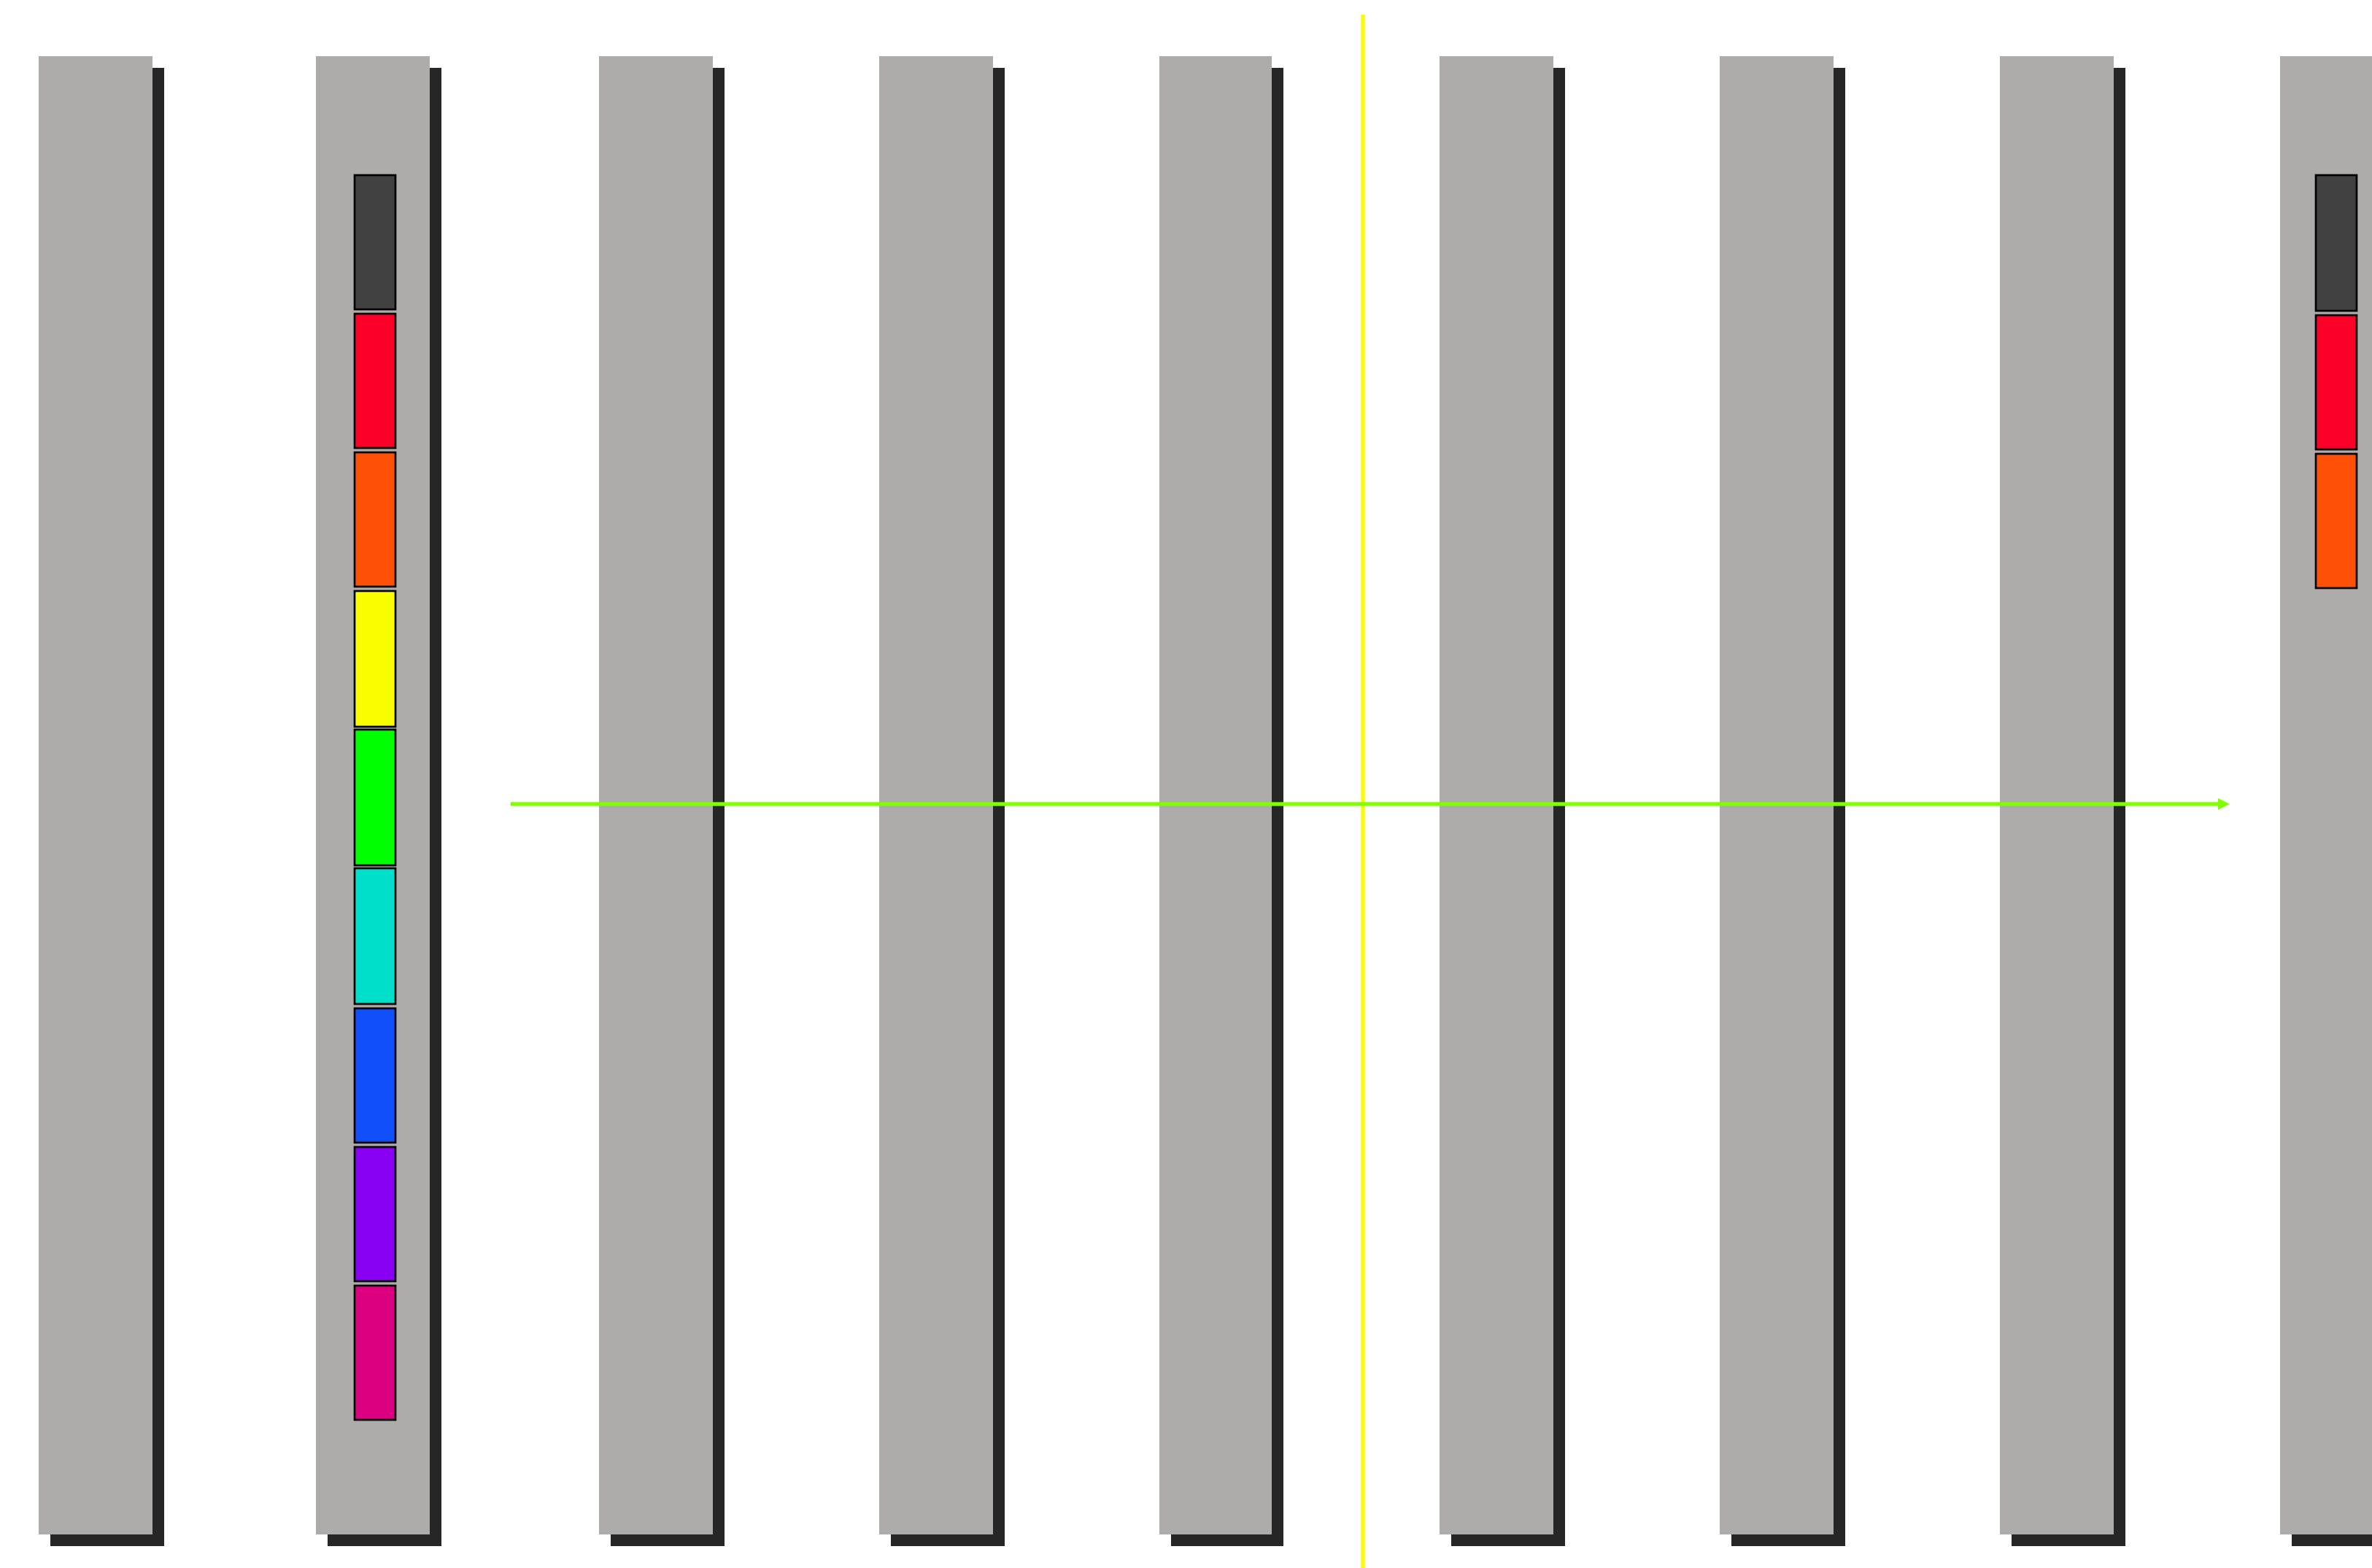
- Red arrows indicate startup of communication (leading to latency, α)
- Green arrows indicate packets in transit (leading to a bandwidth related cost proportional to β and the length of the packet)

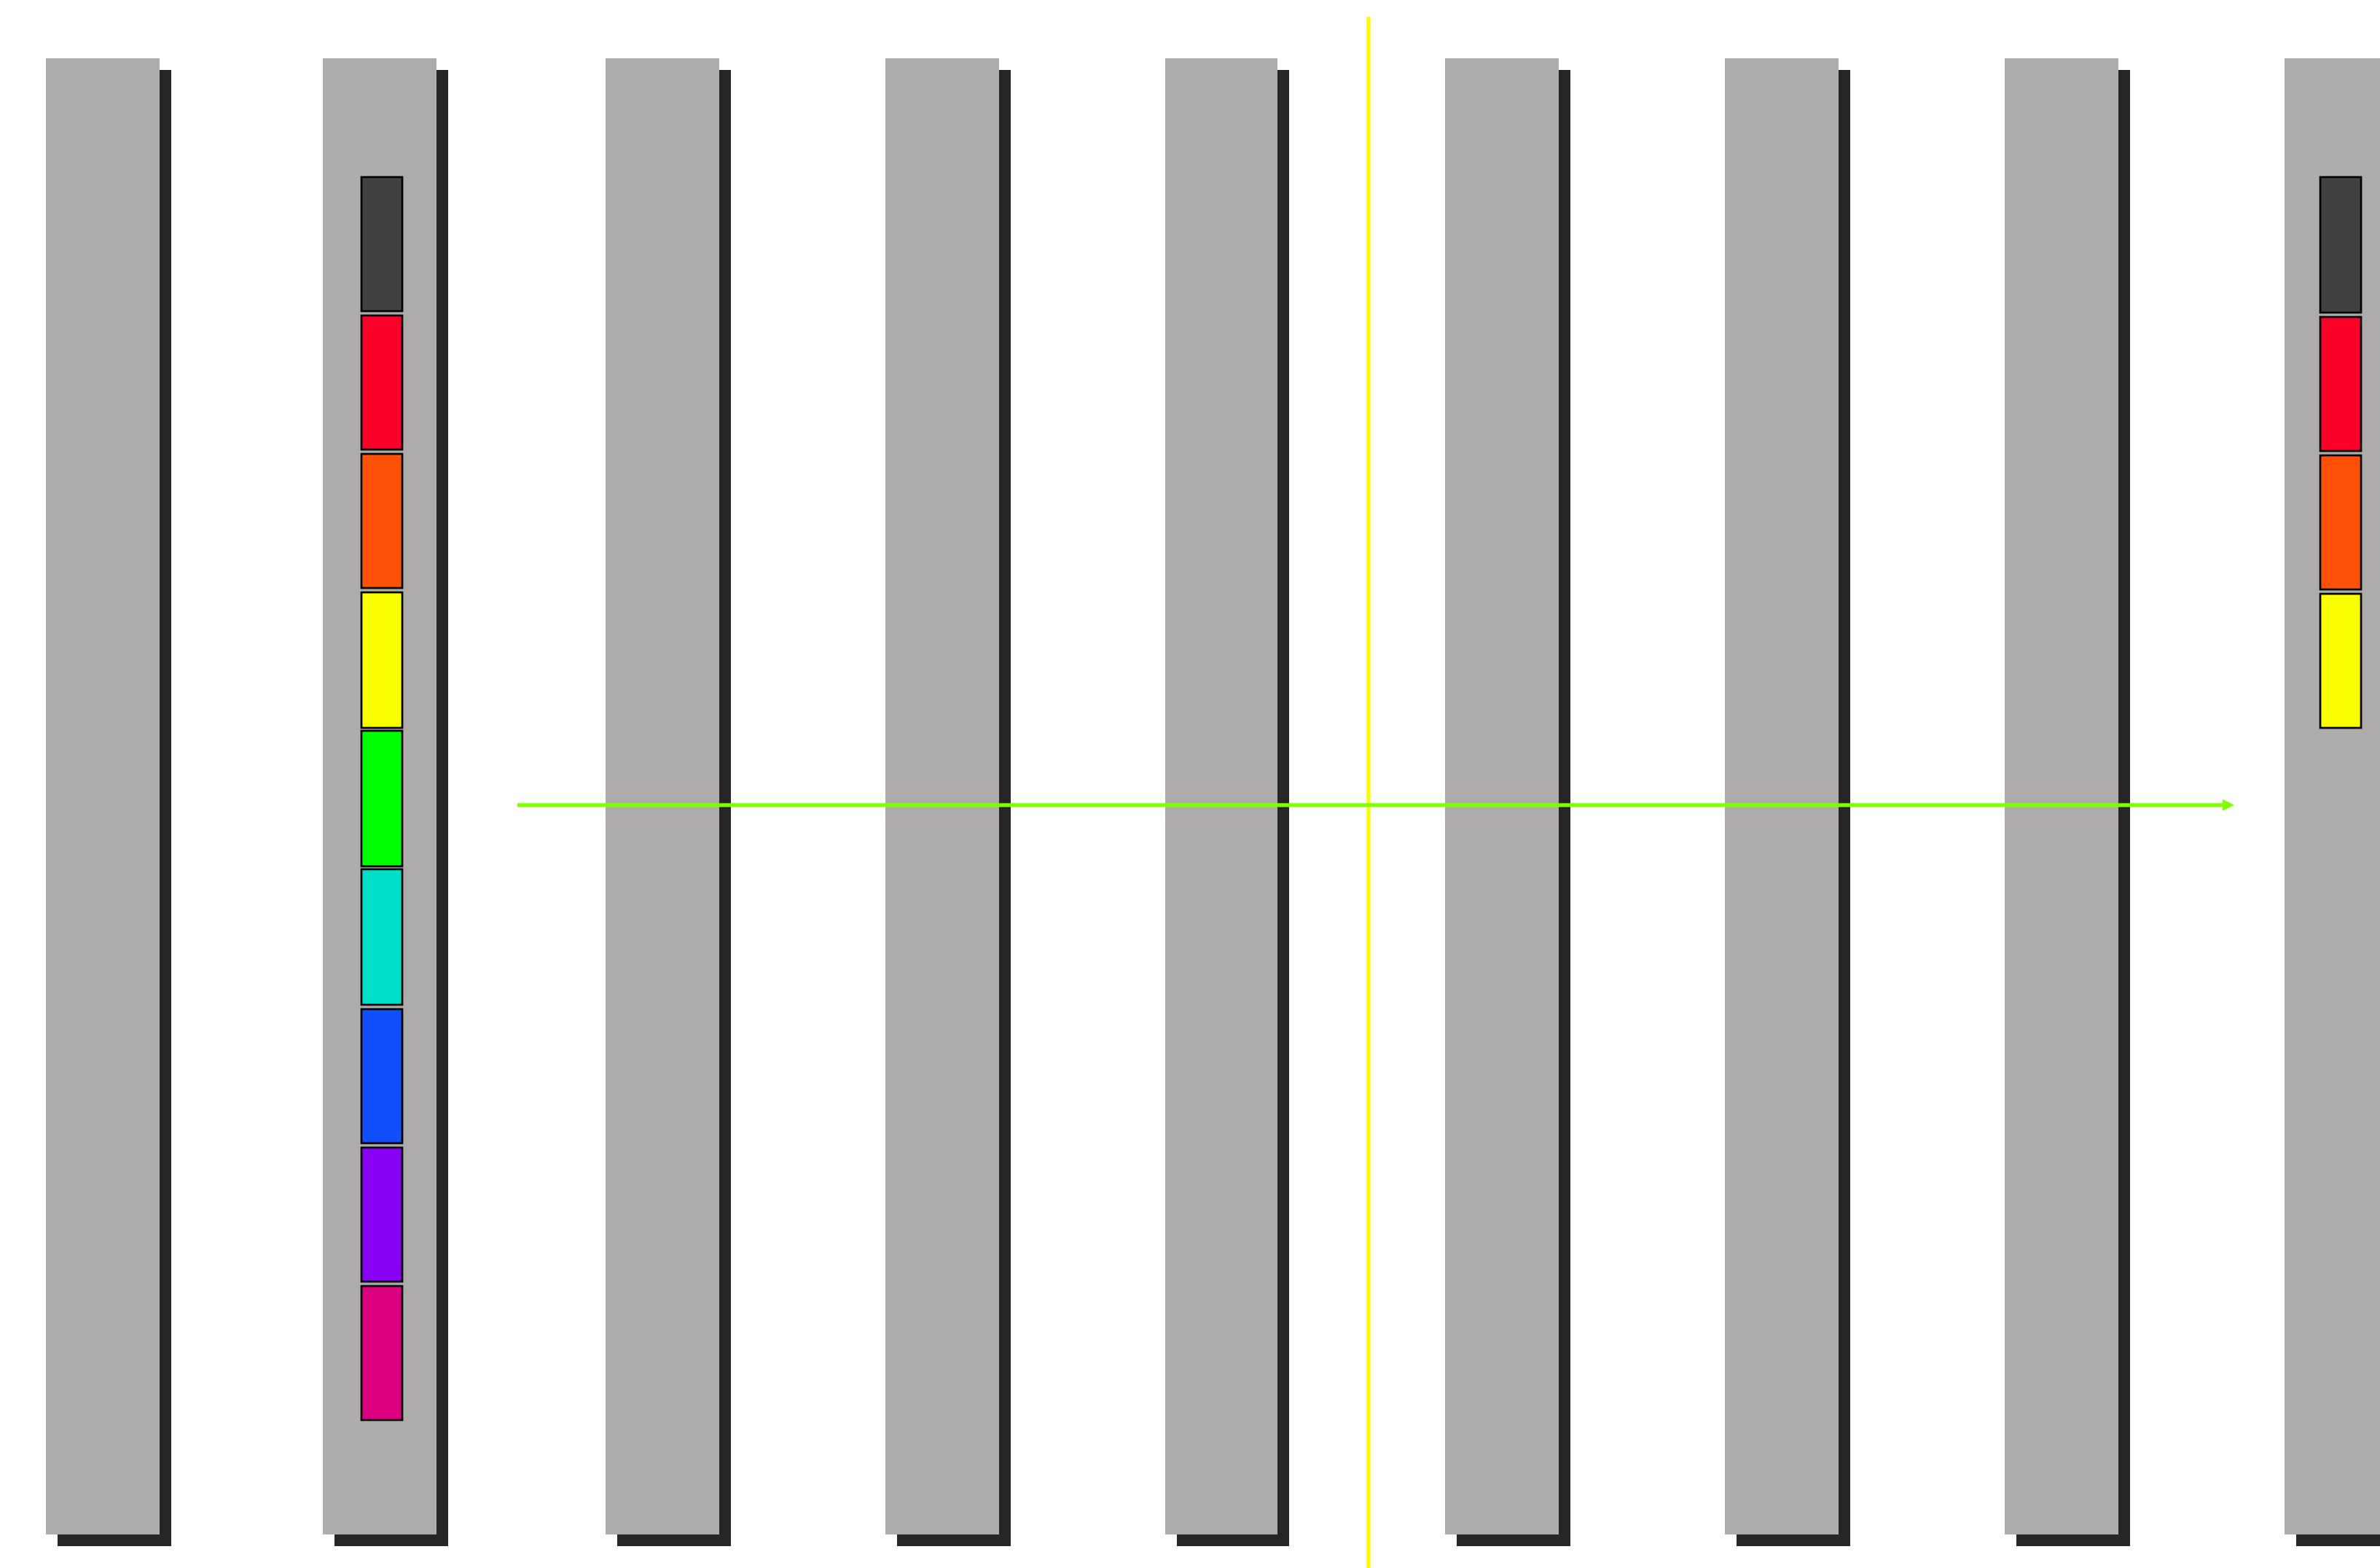


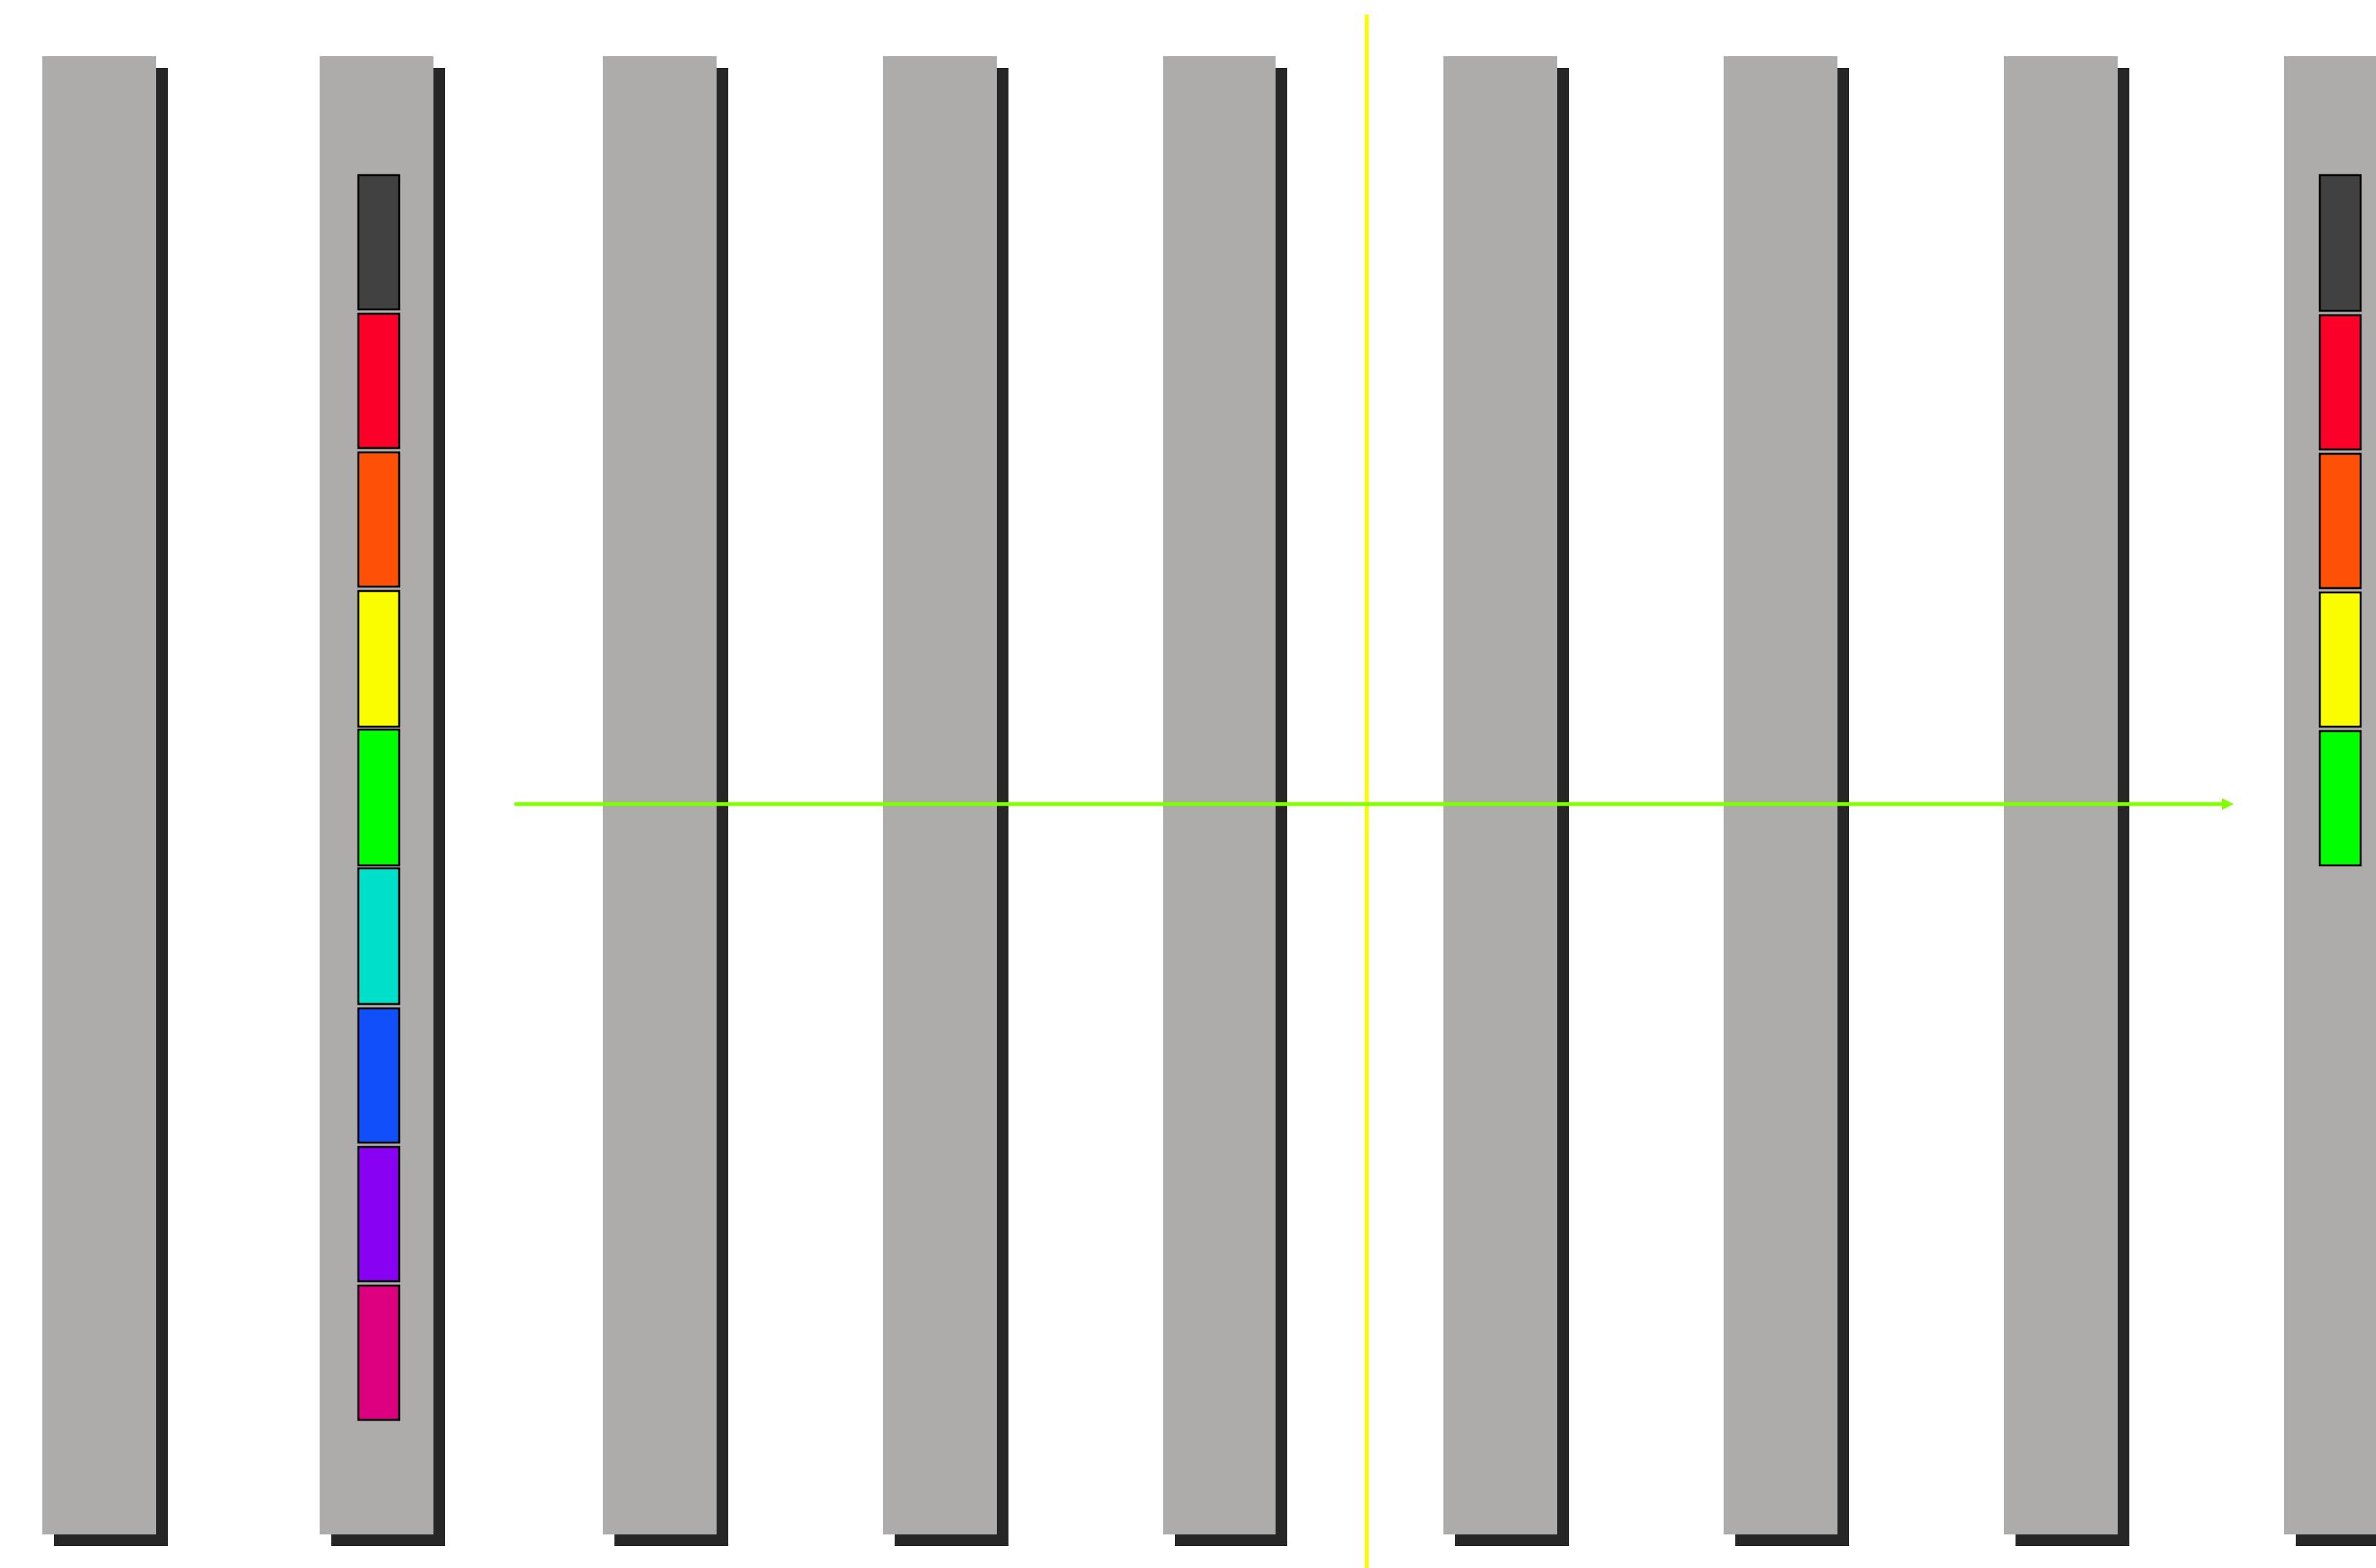


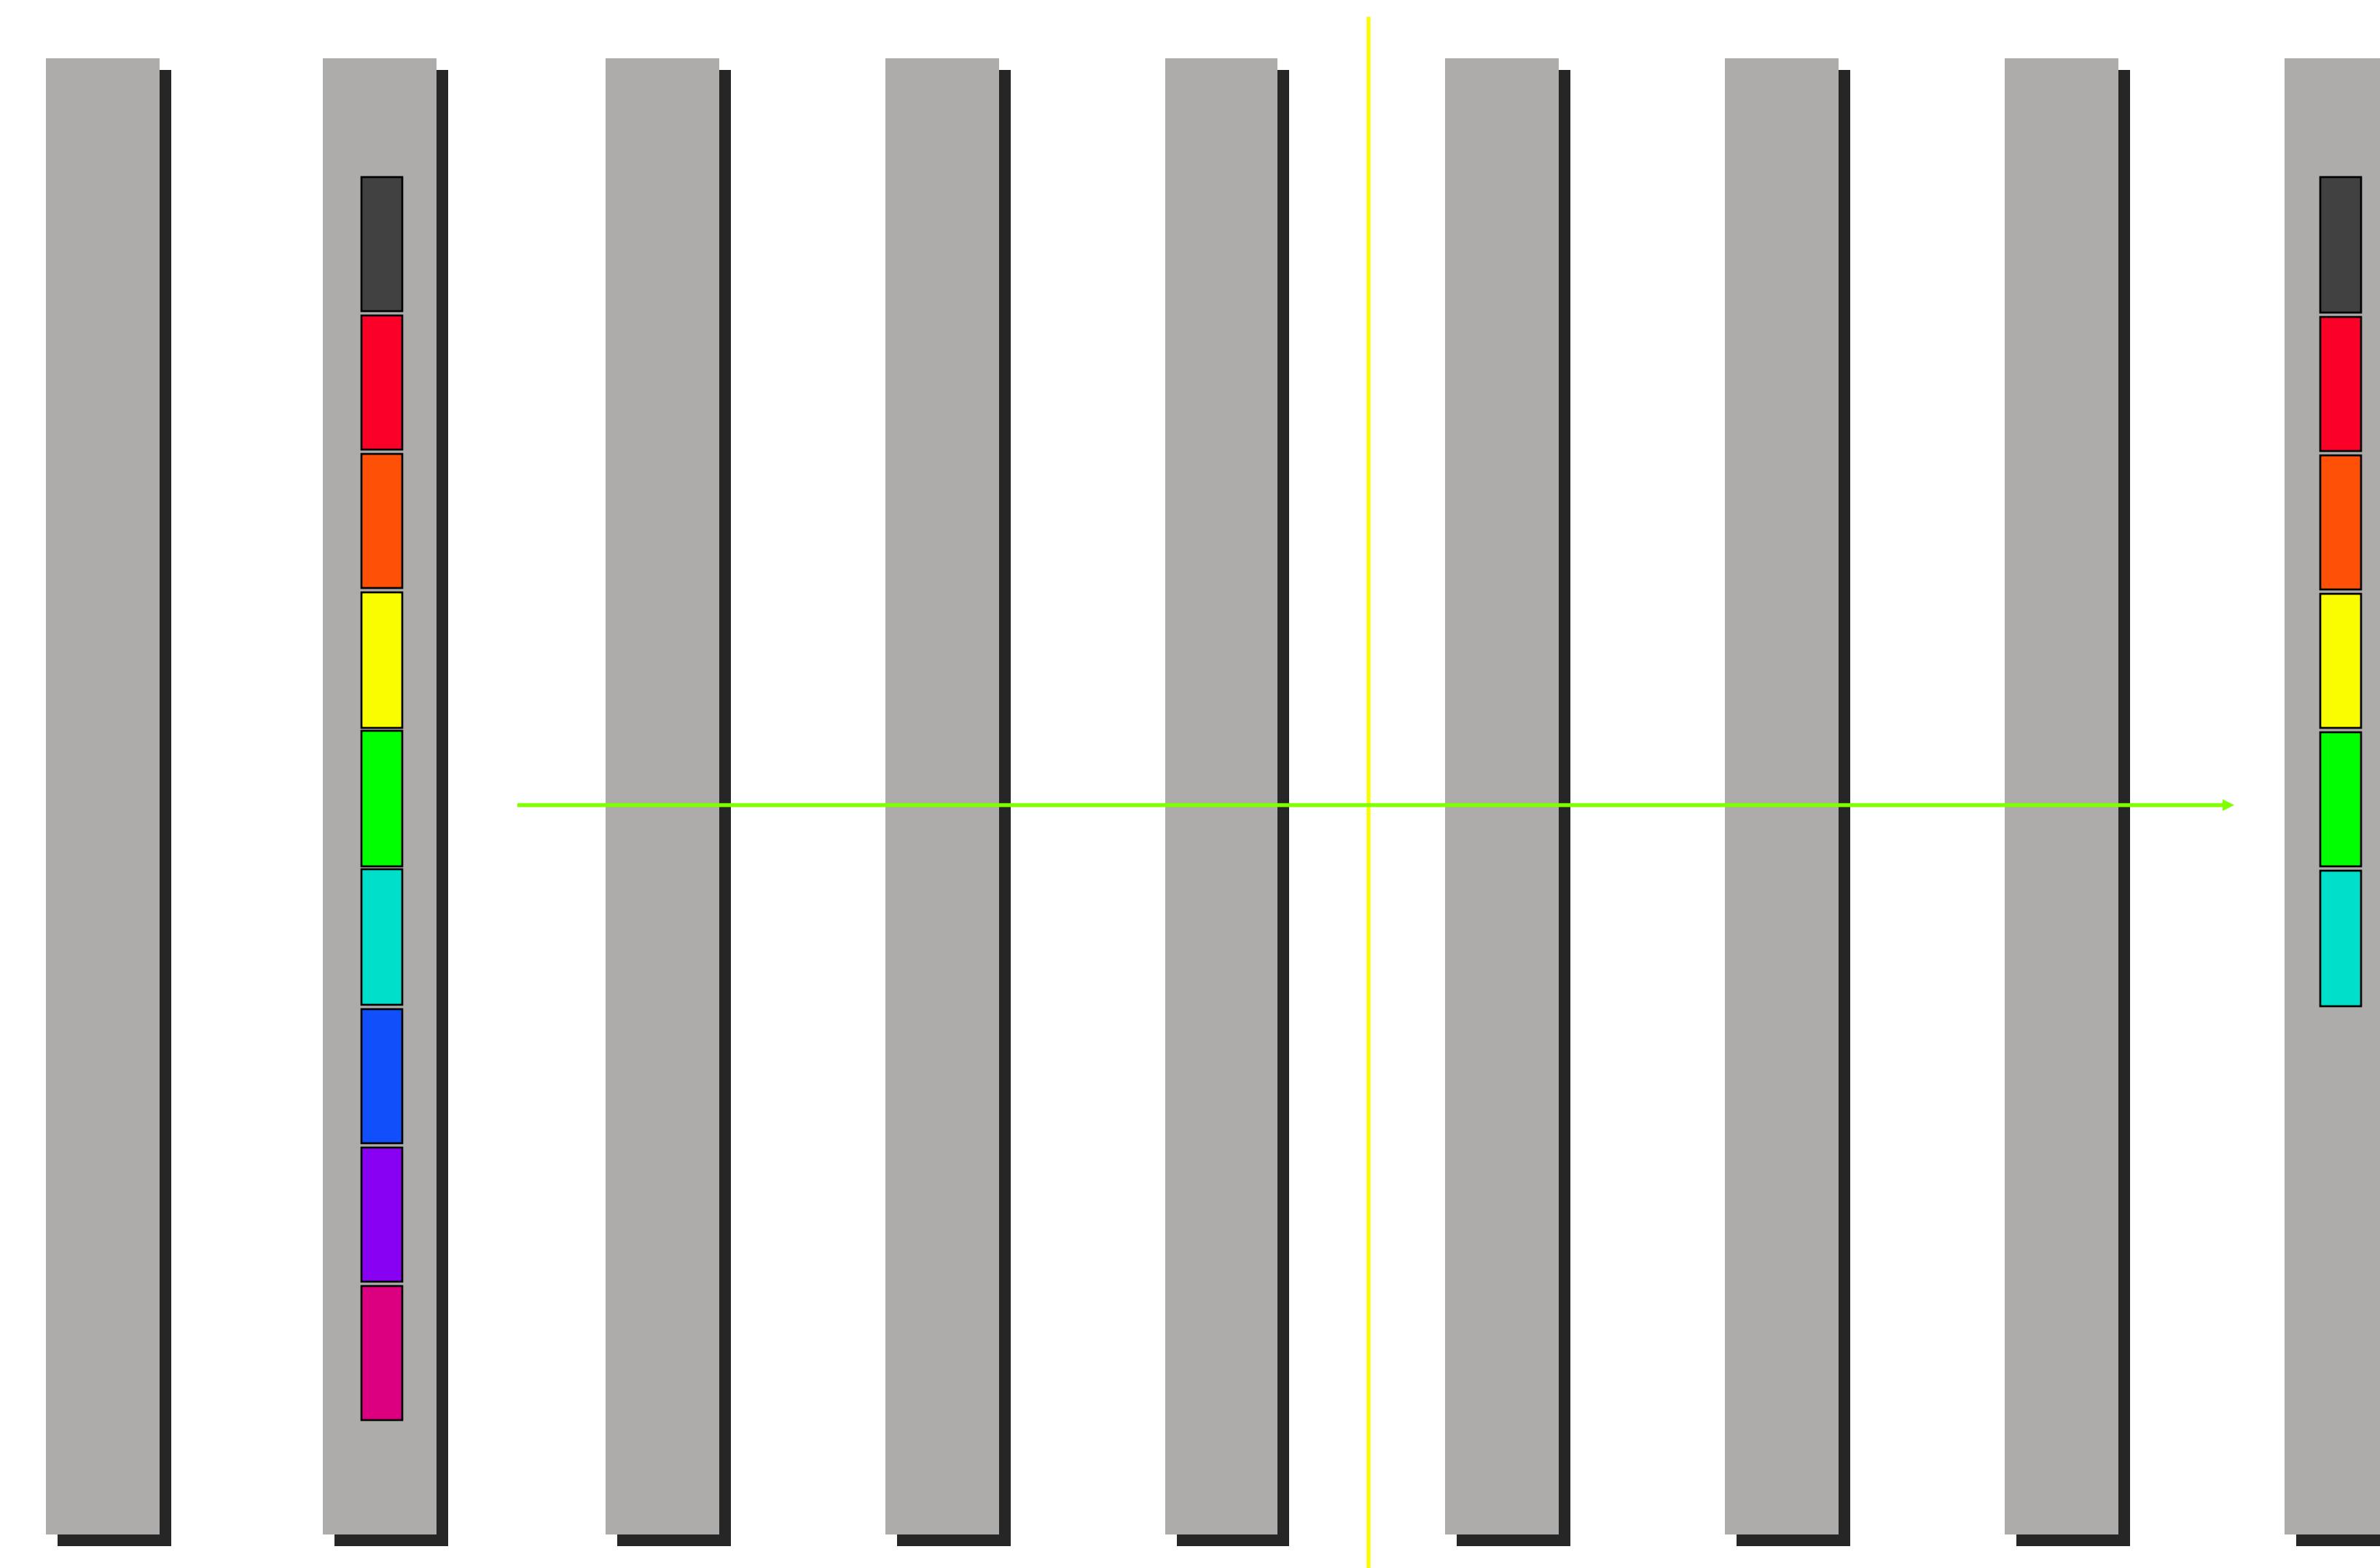


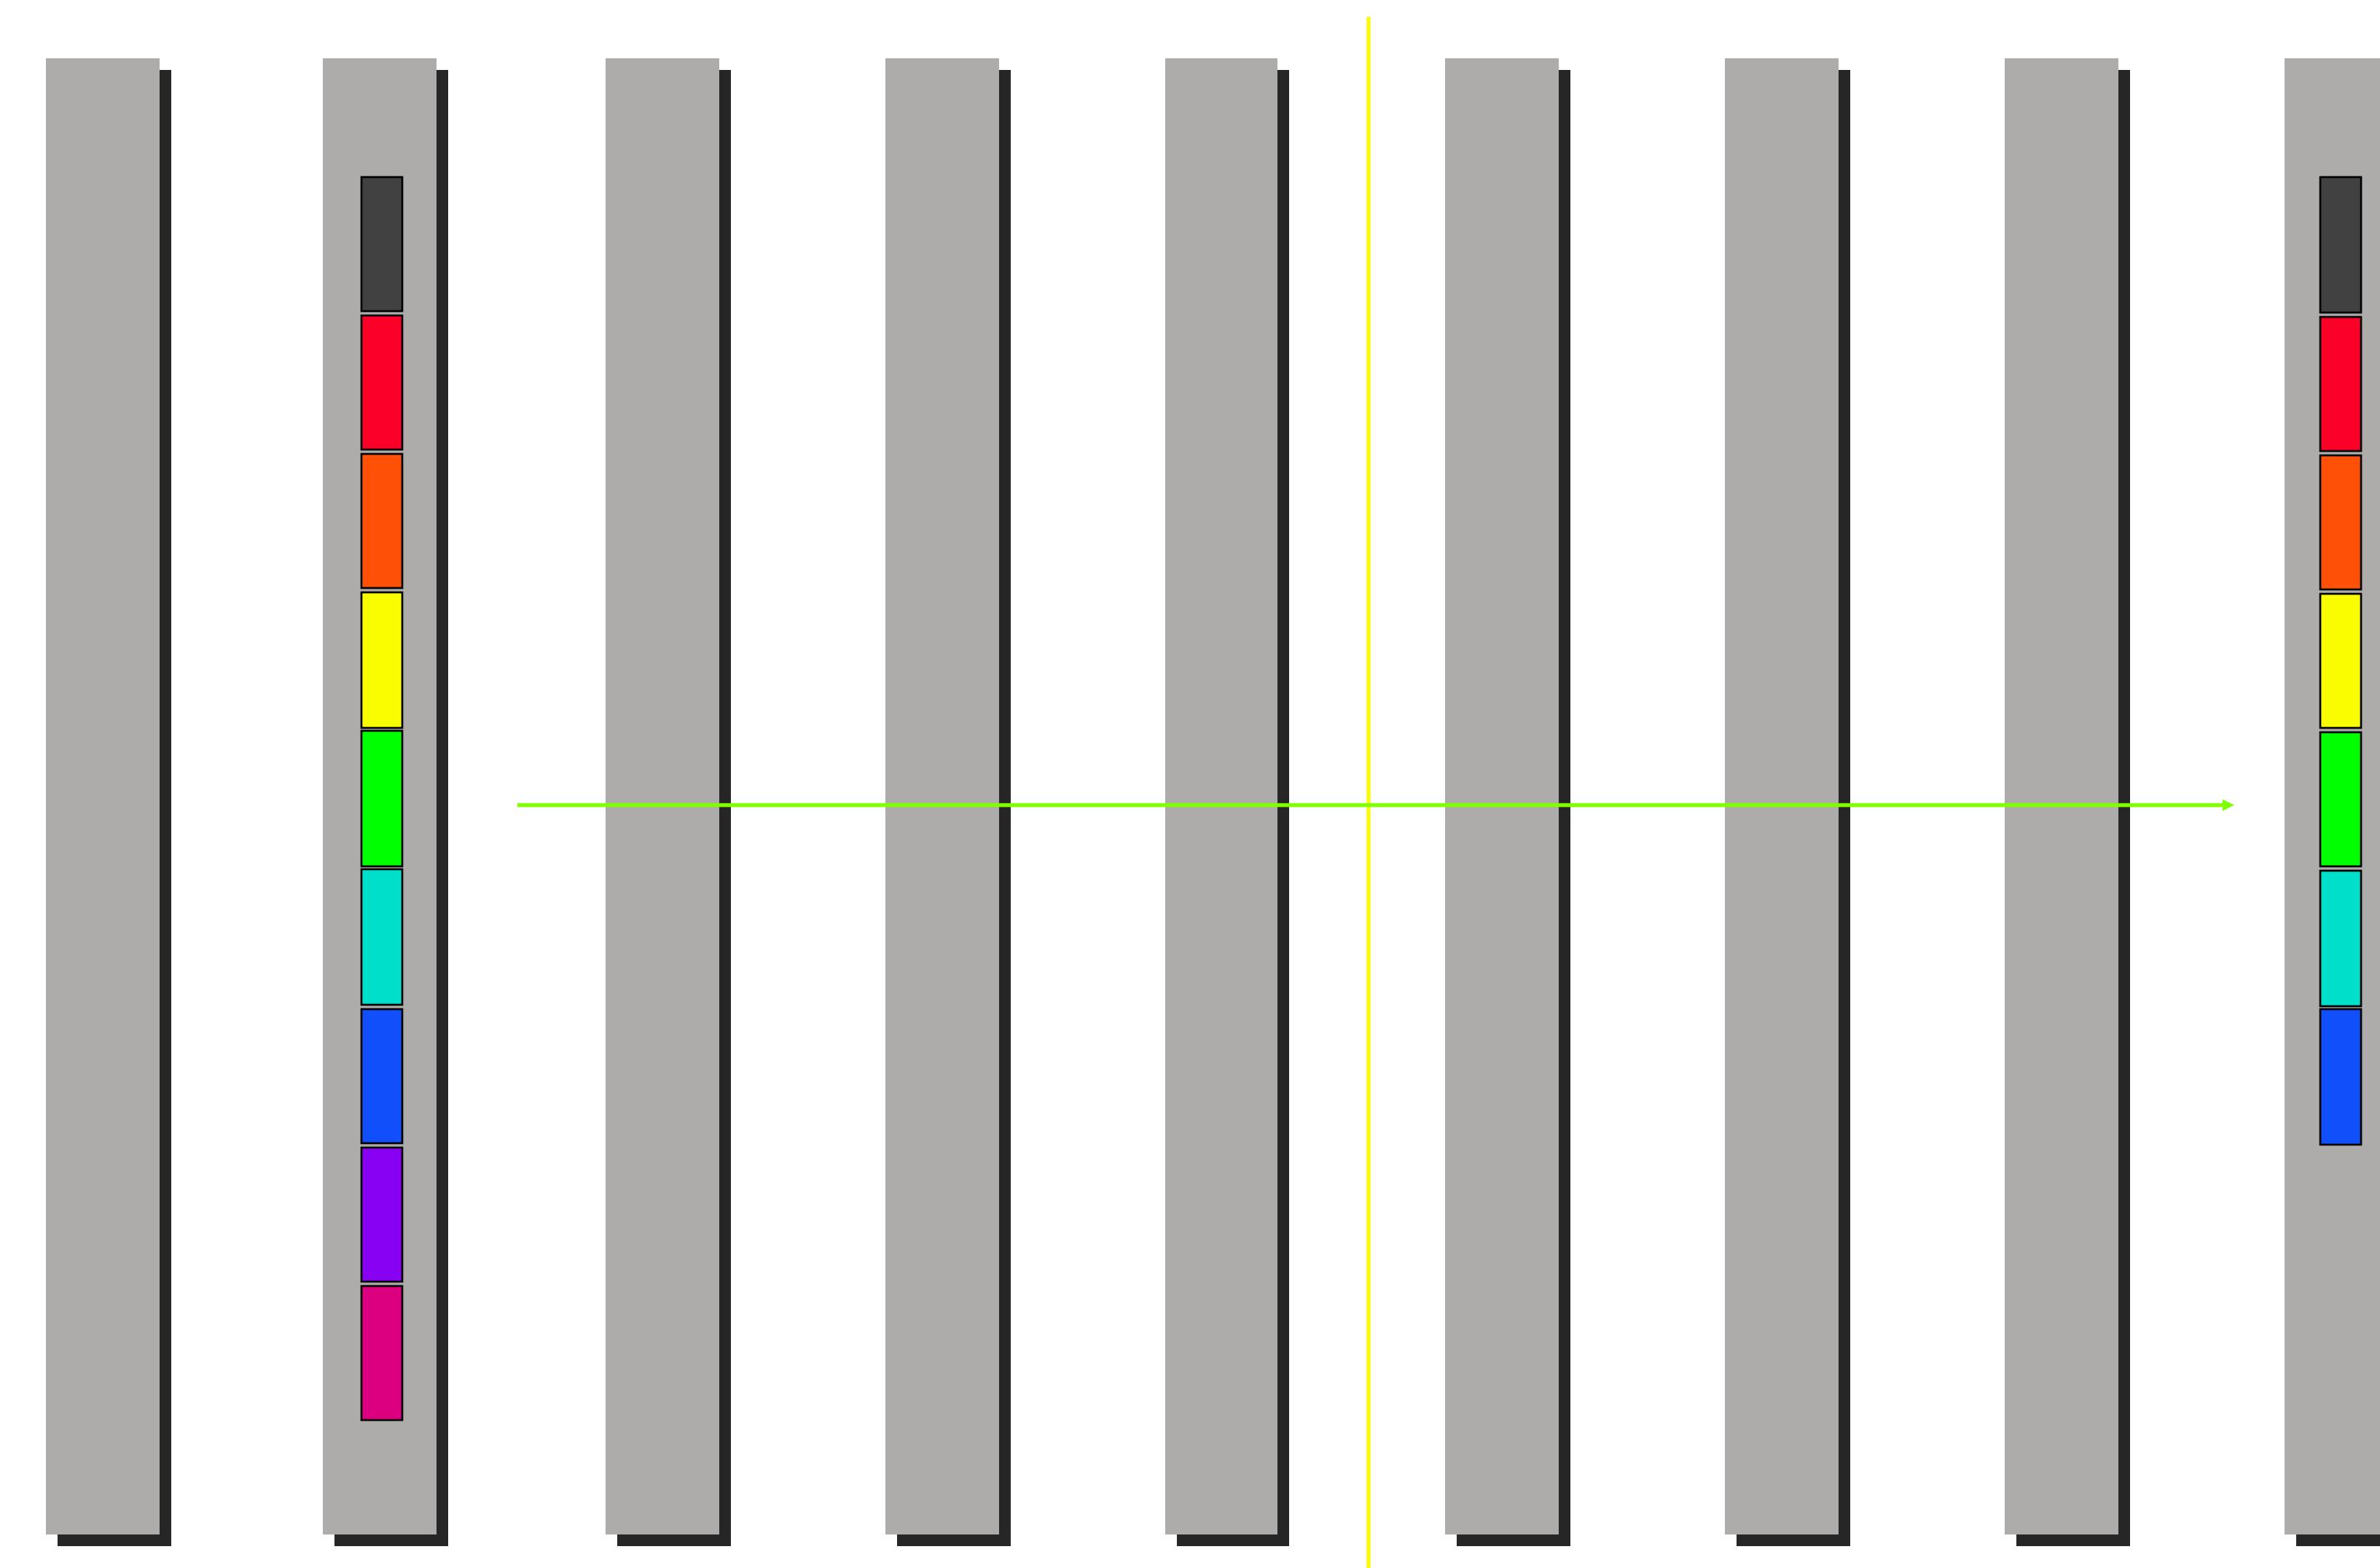


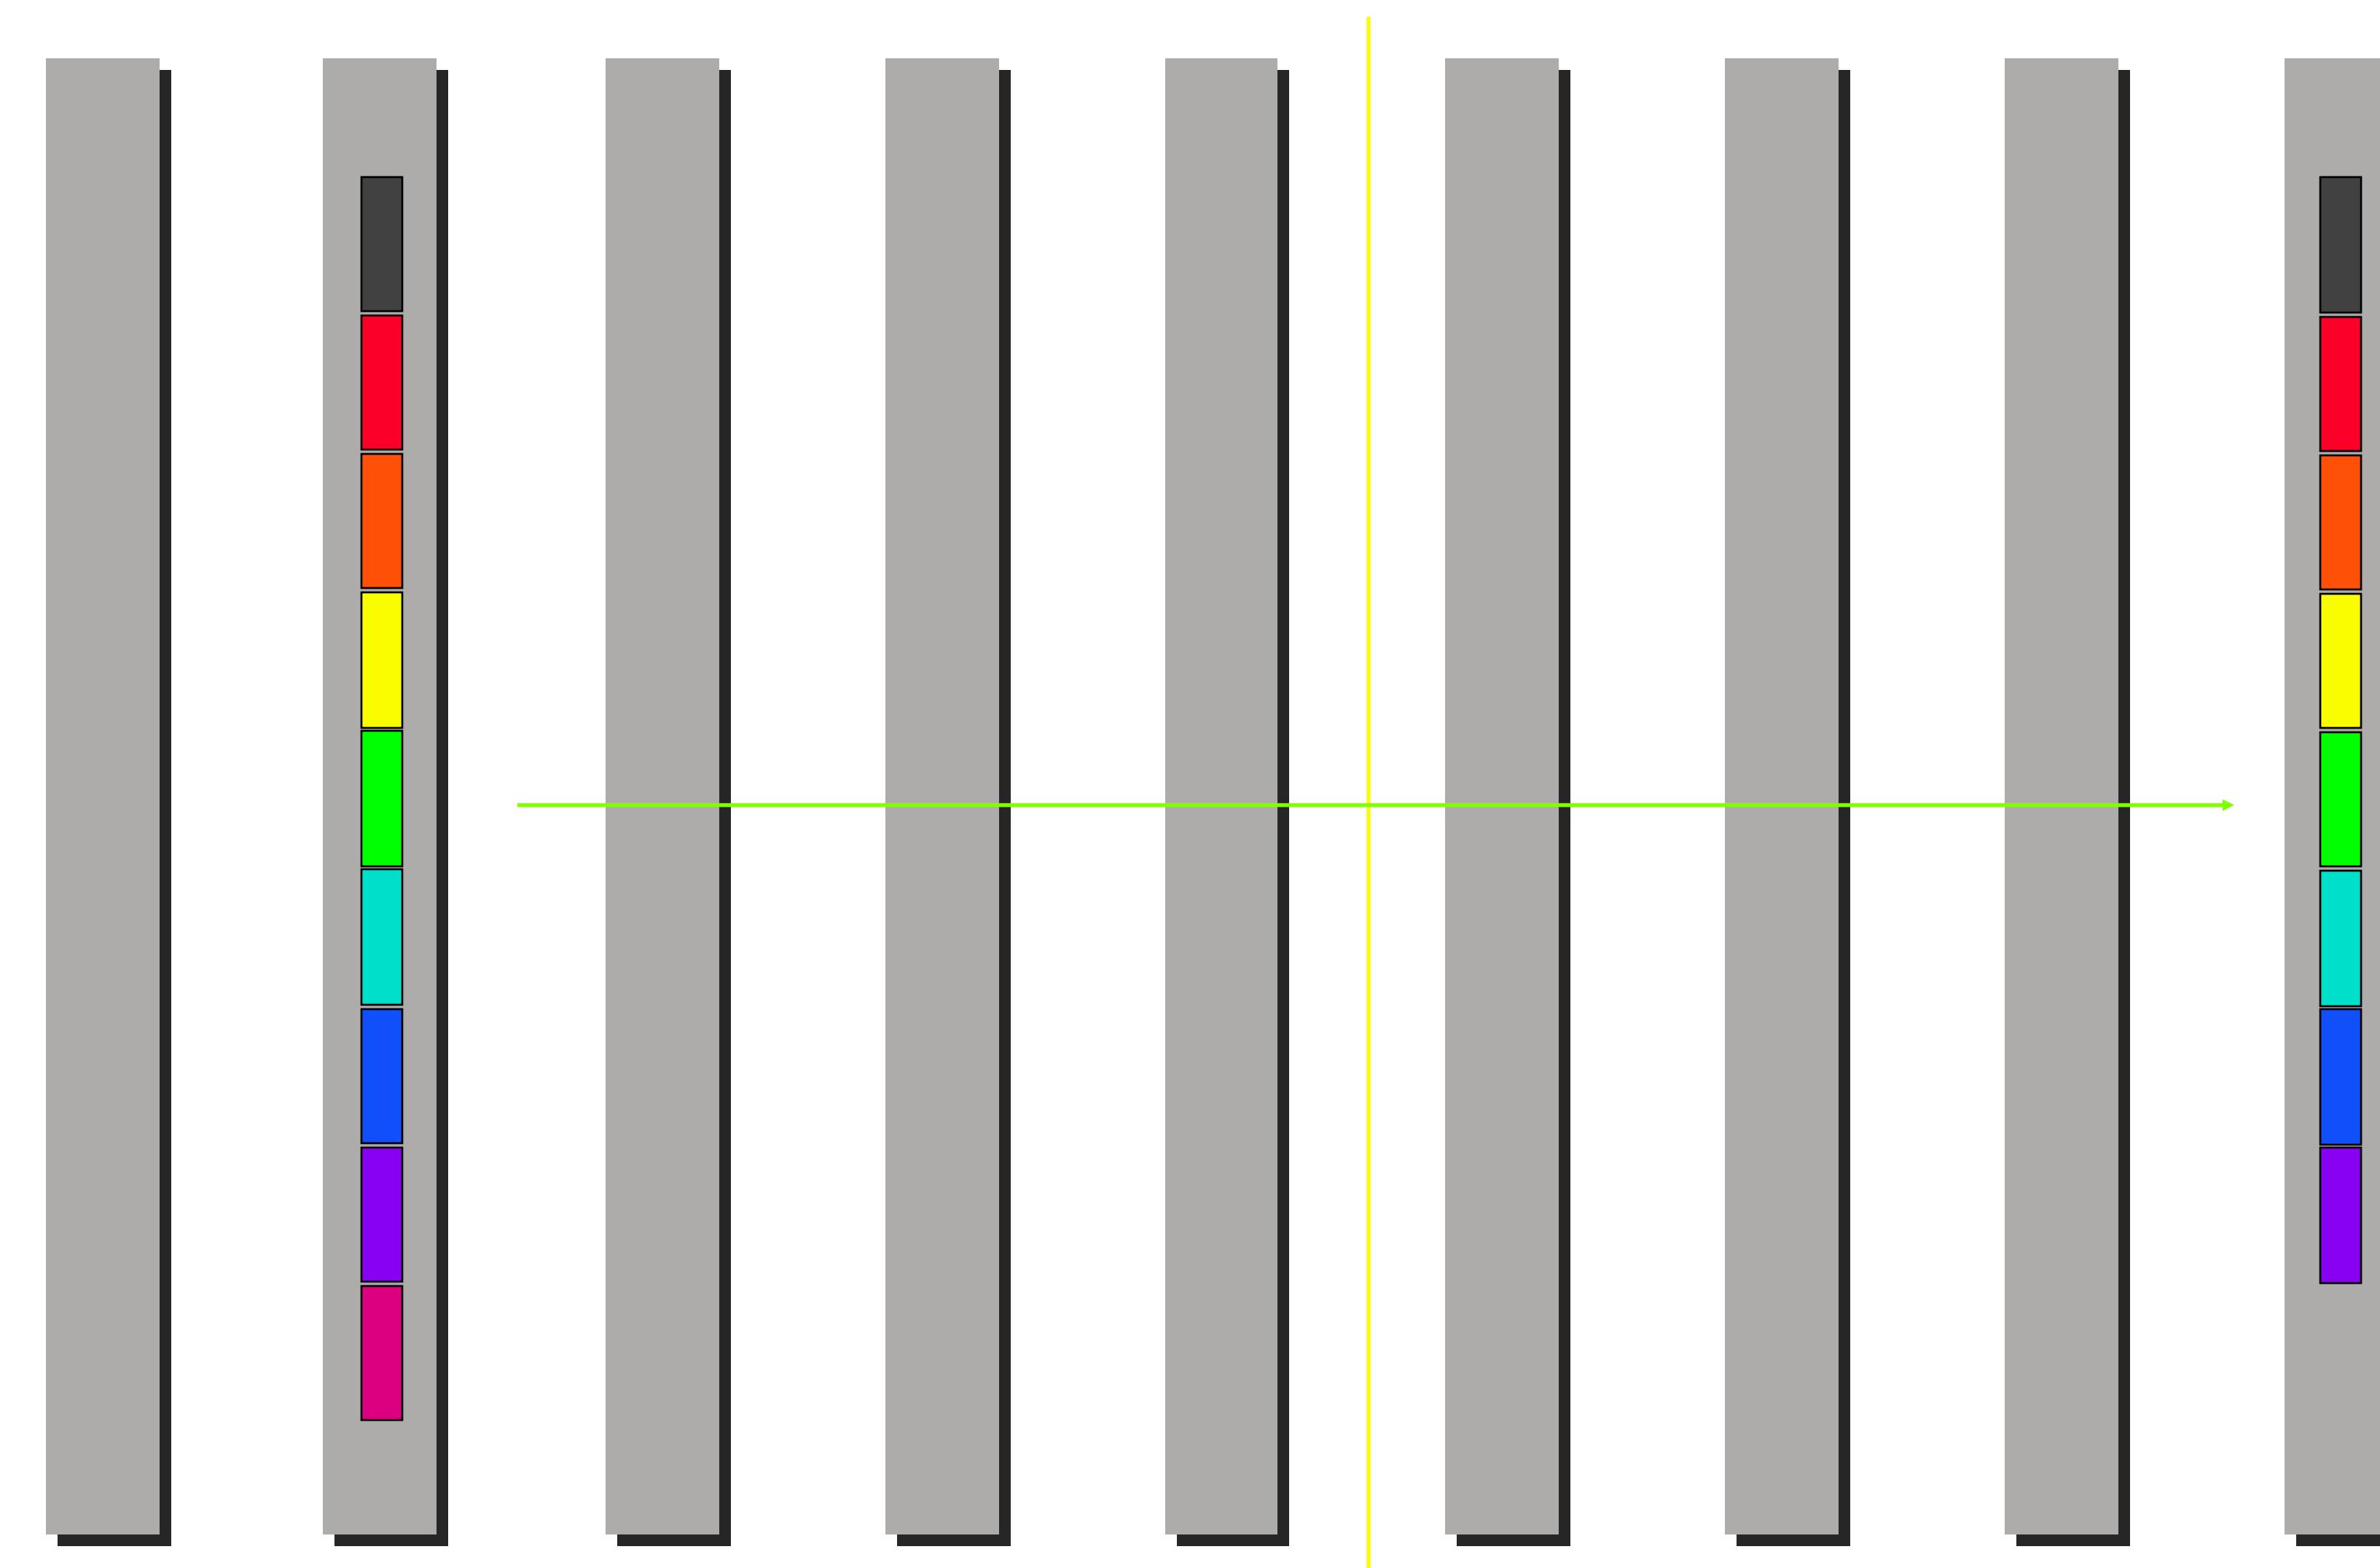


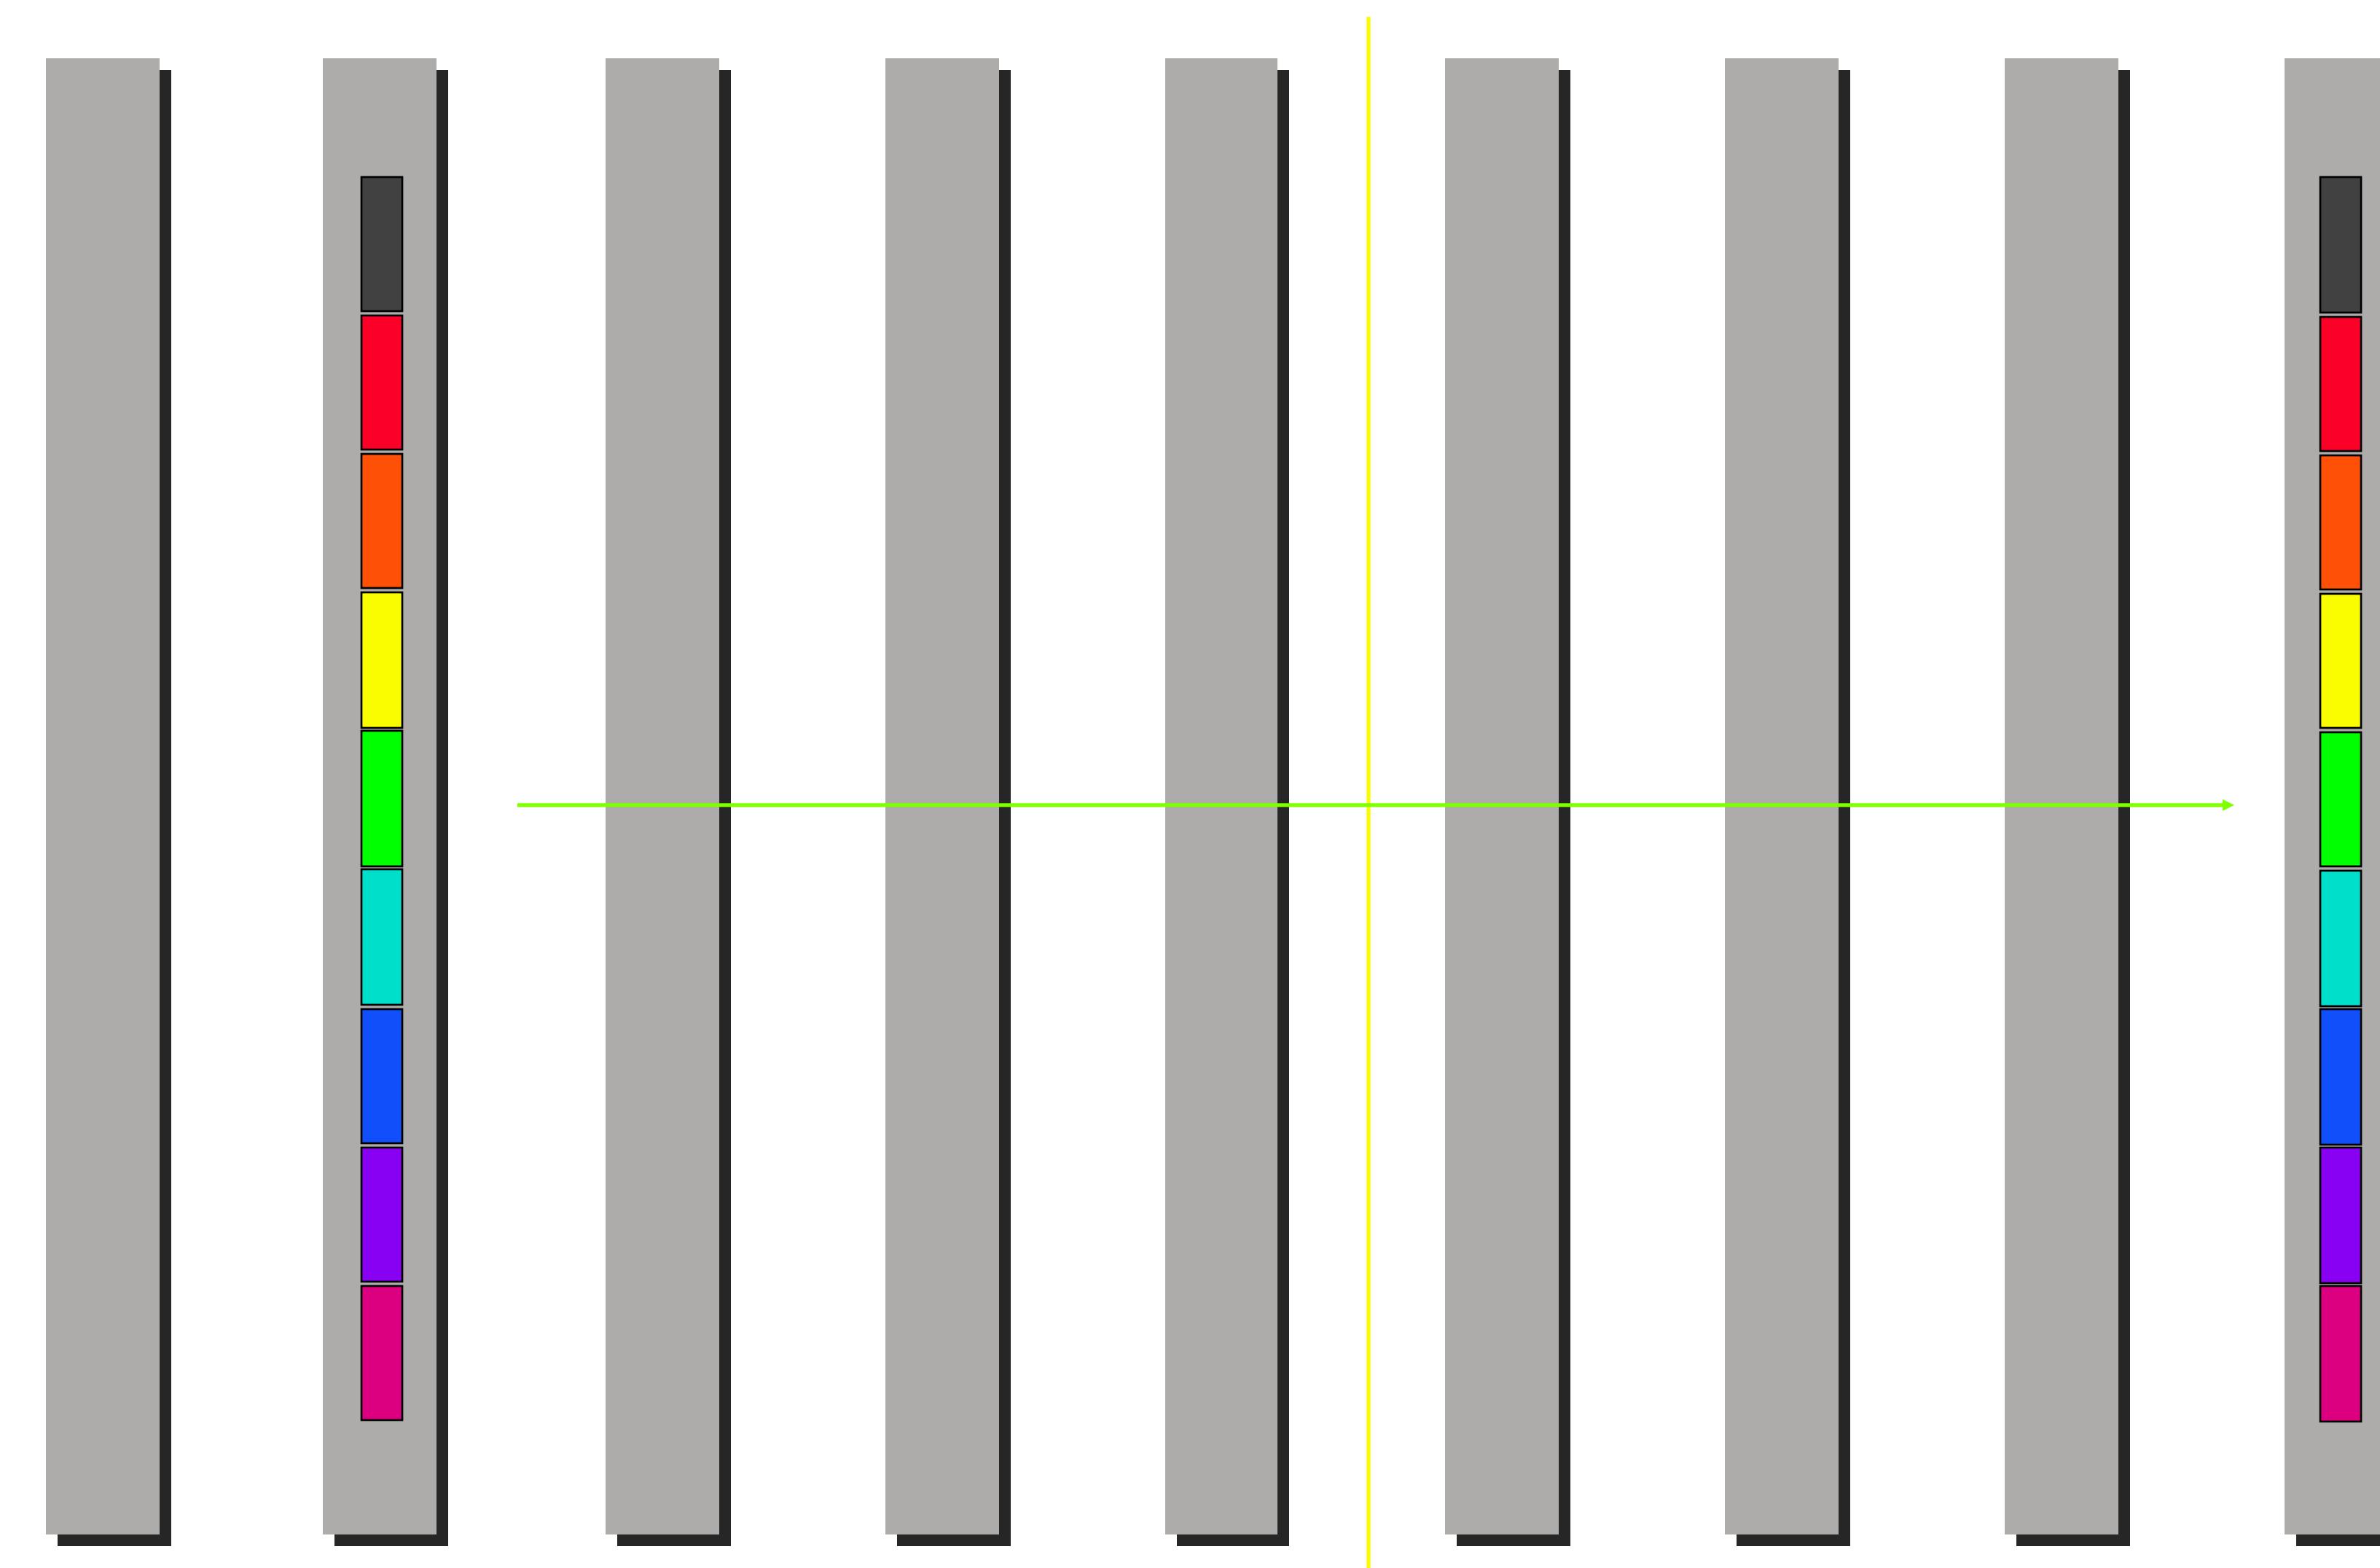


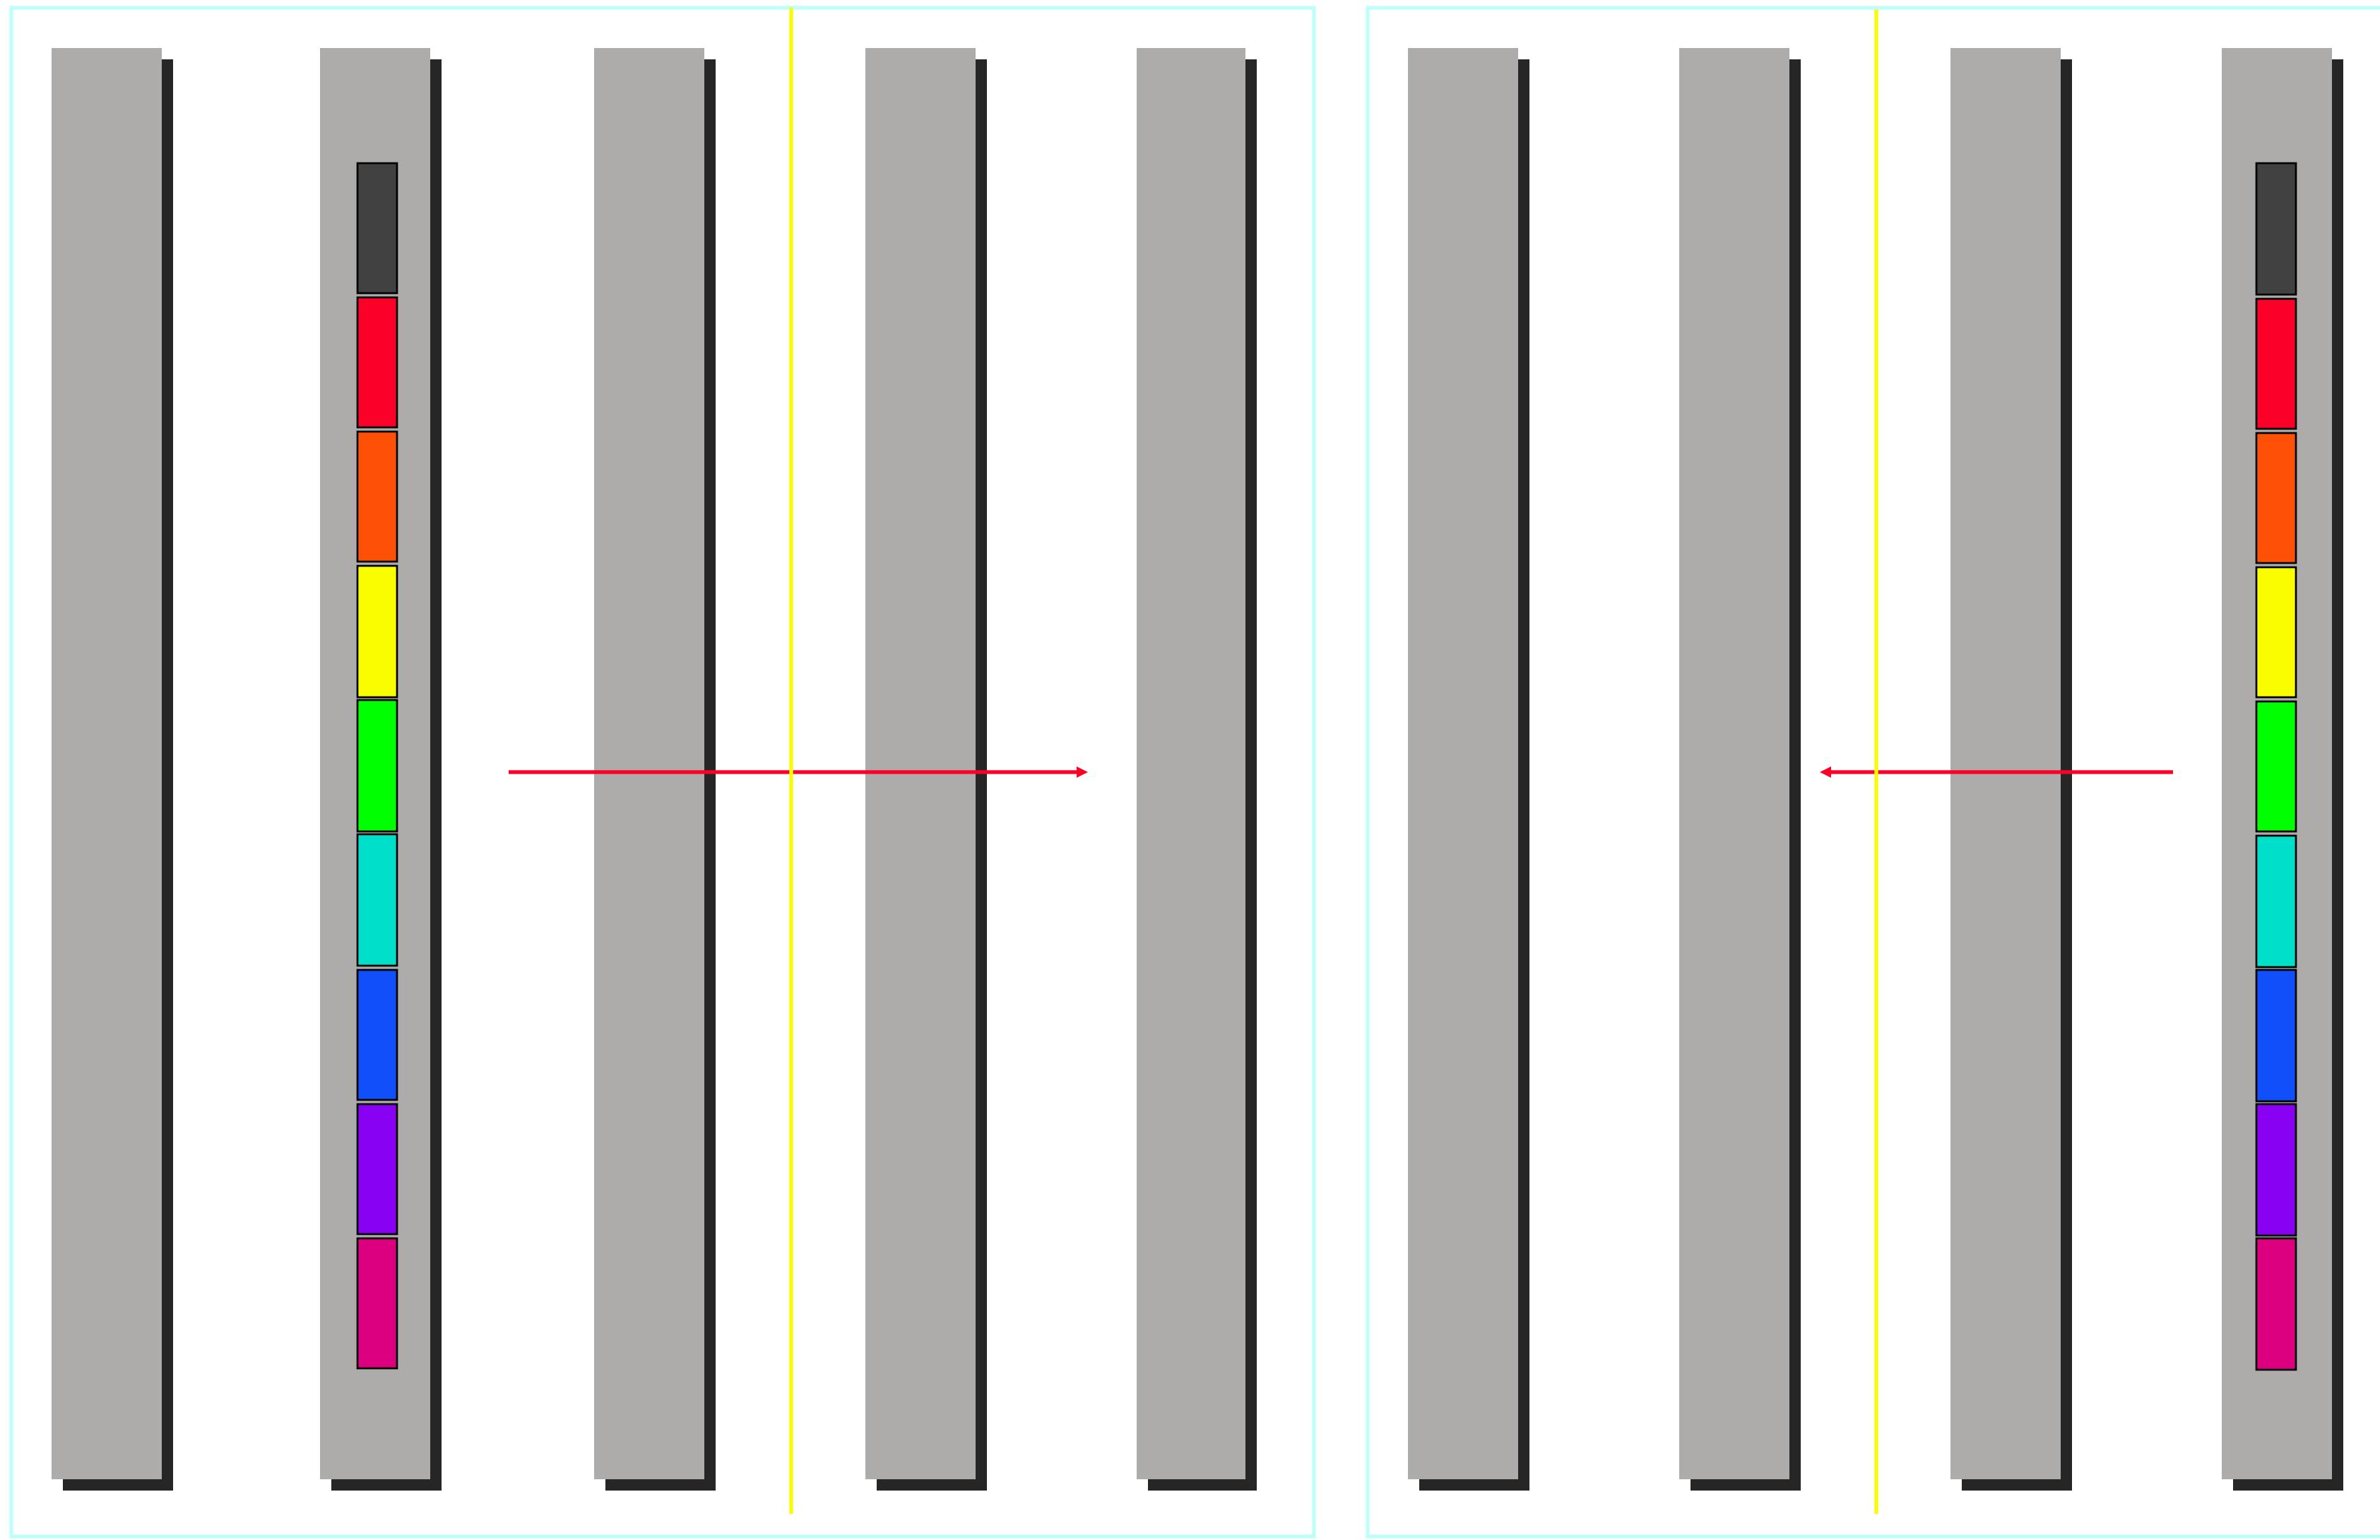


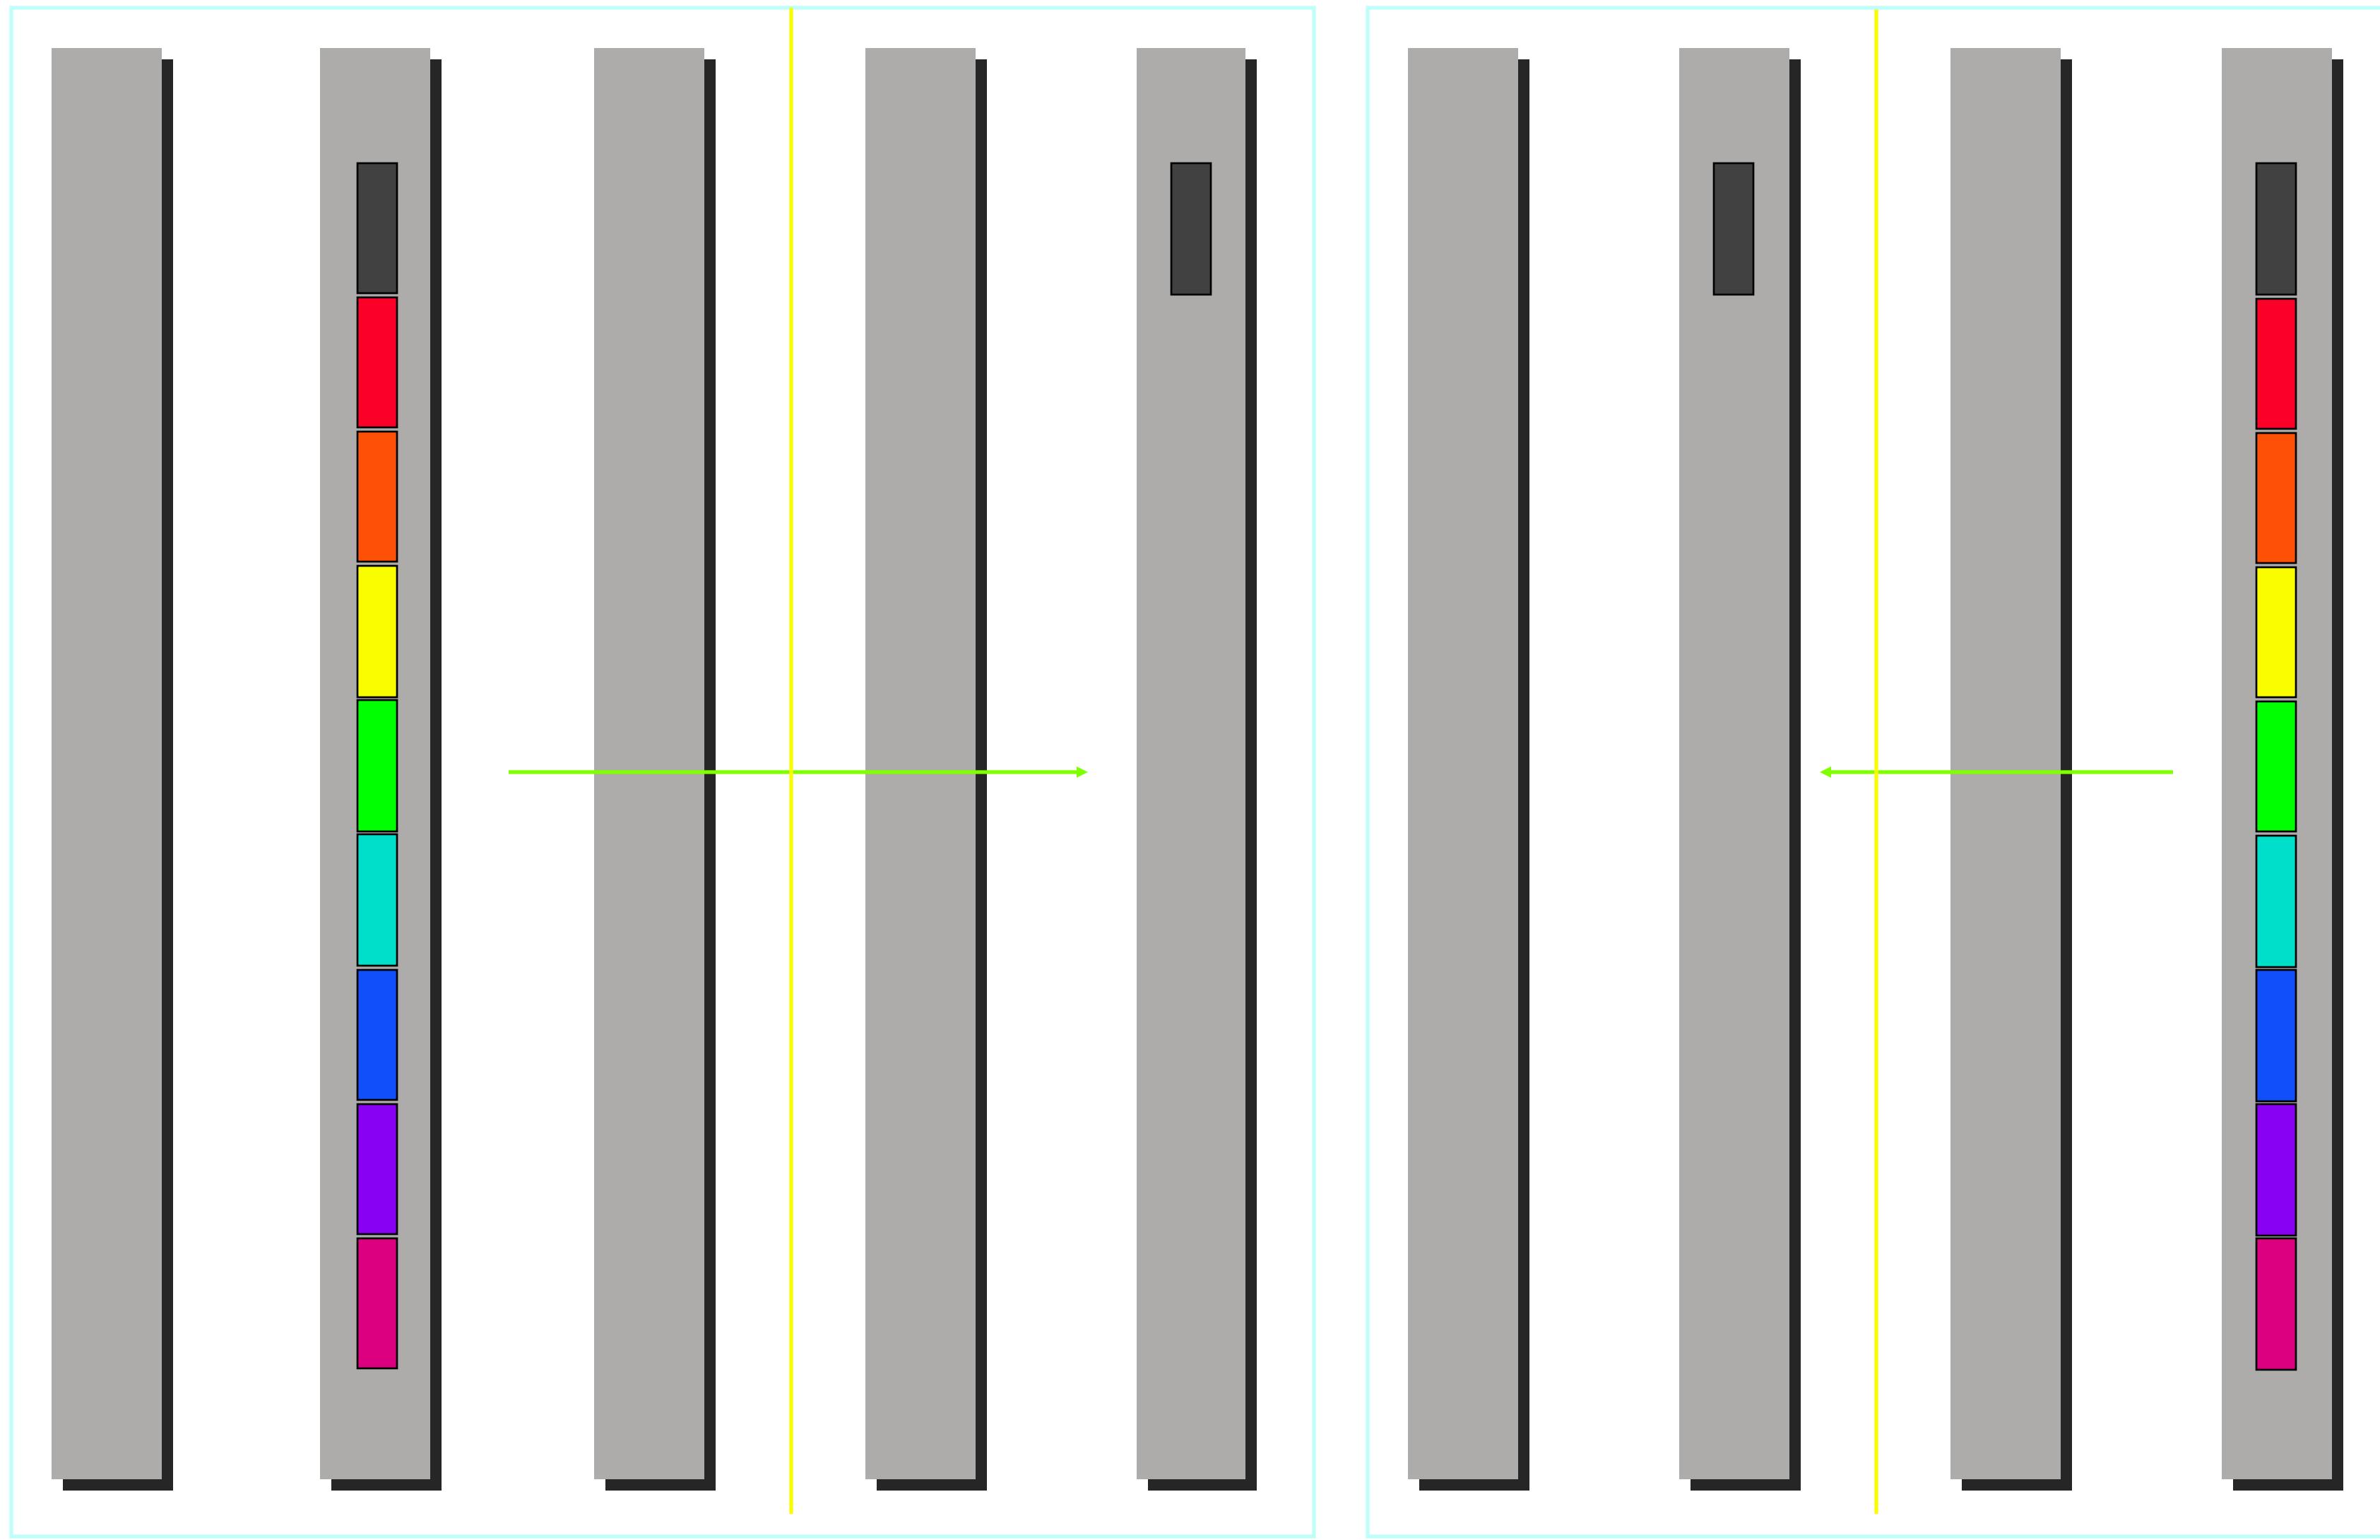


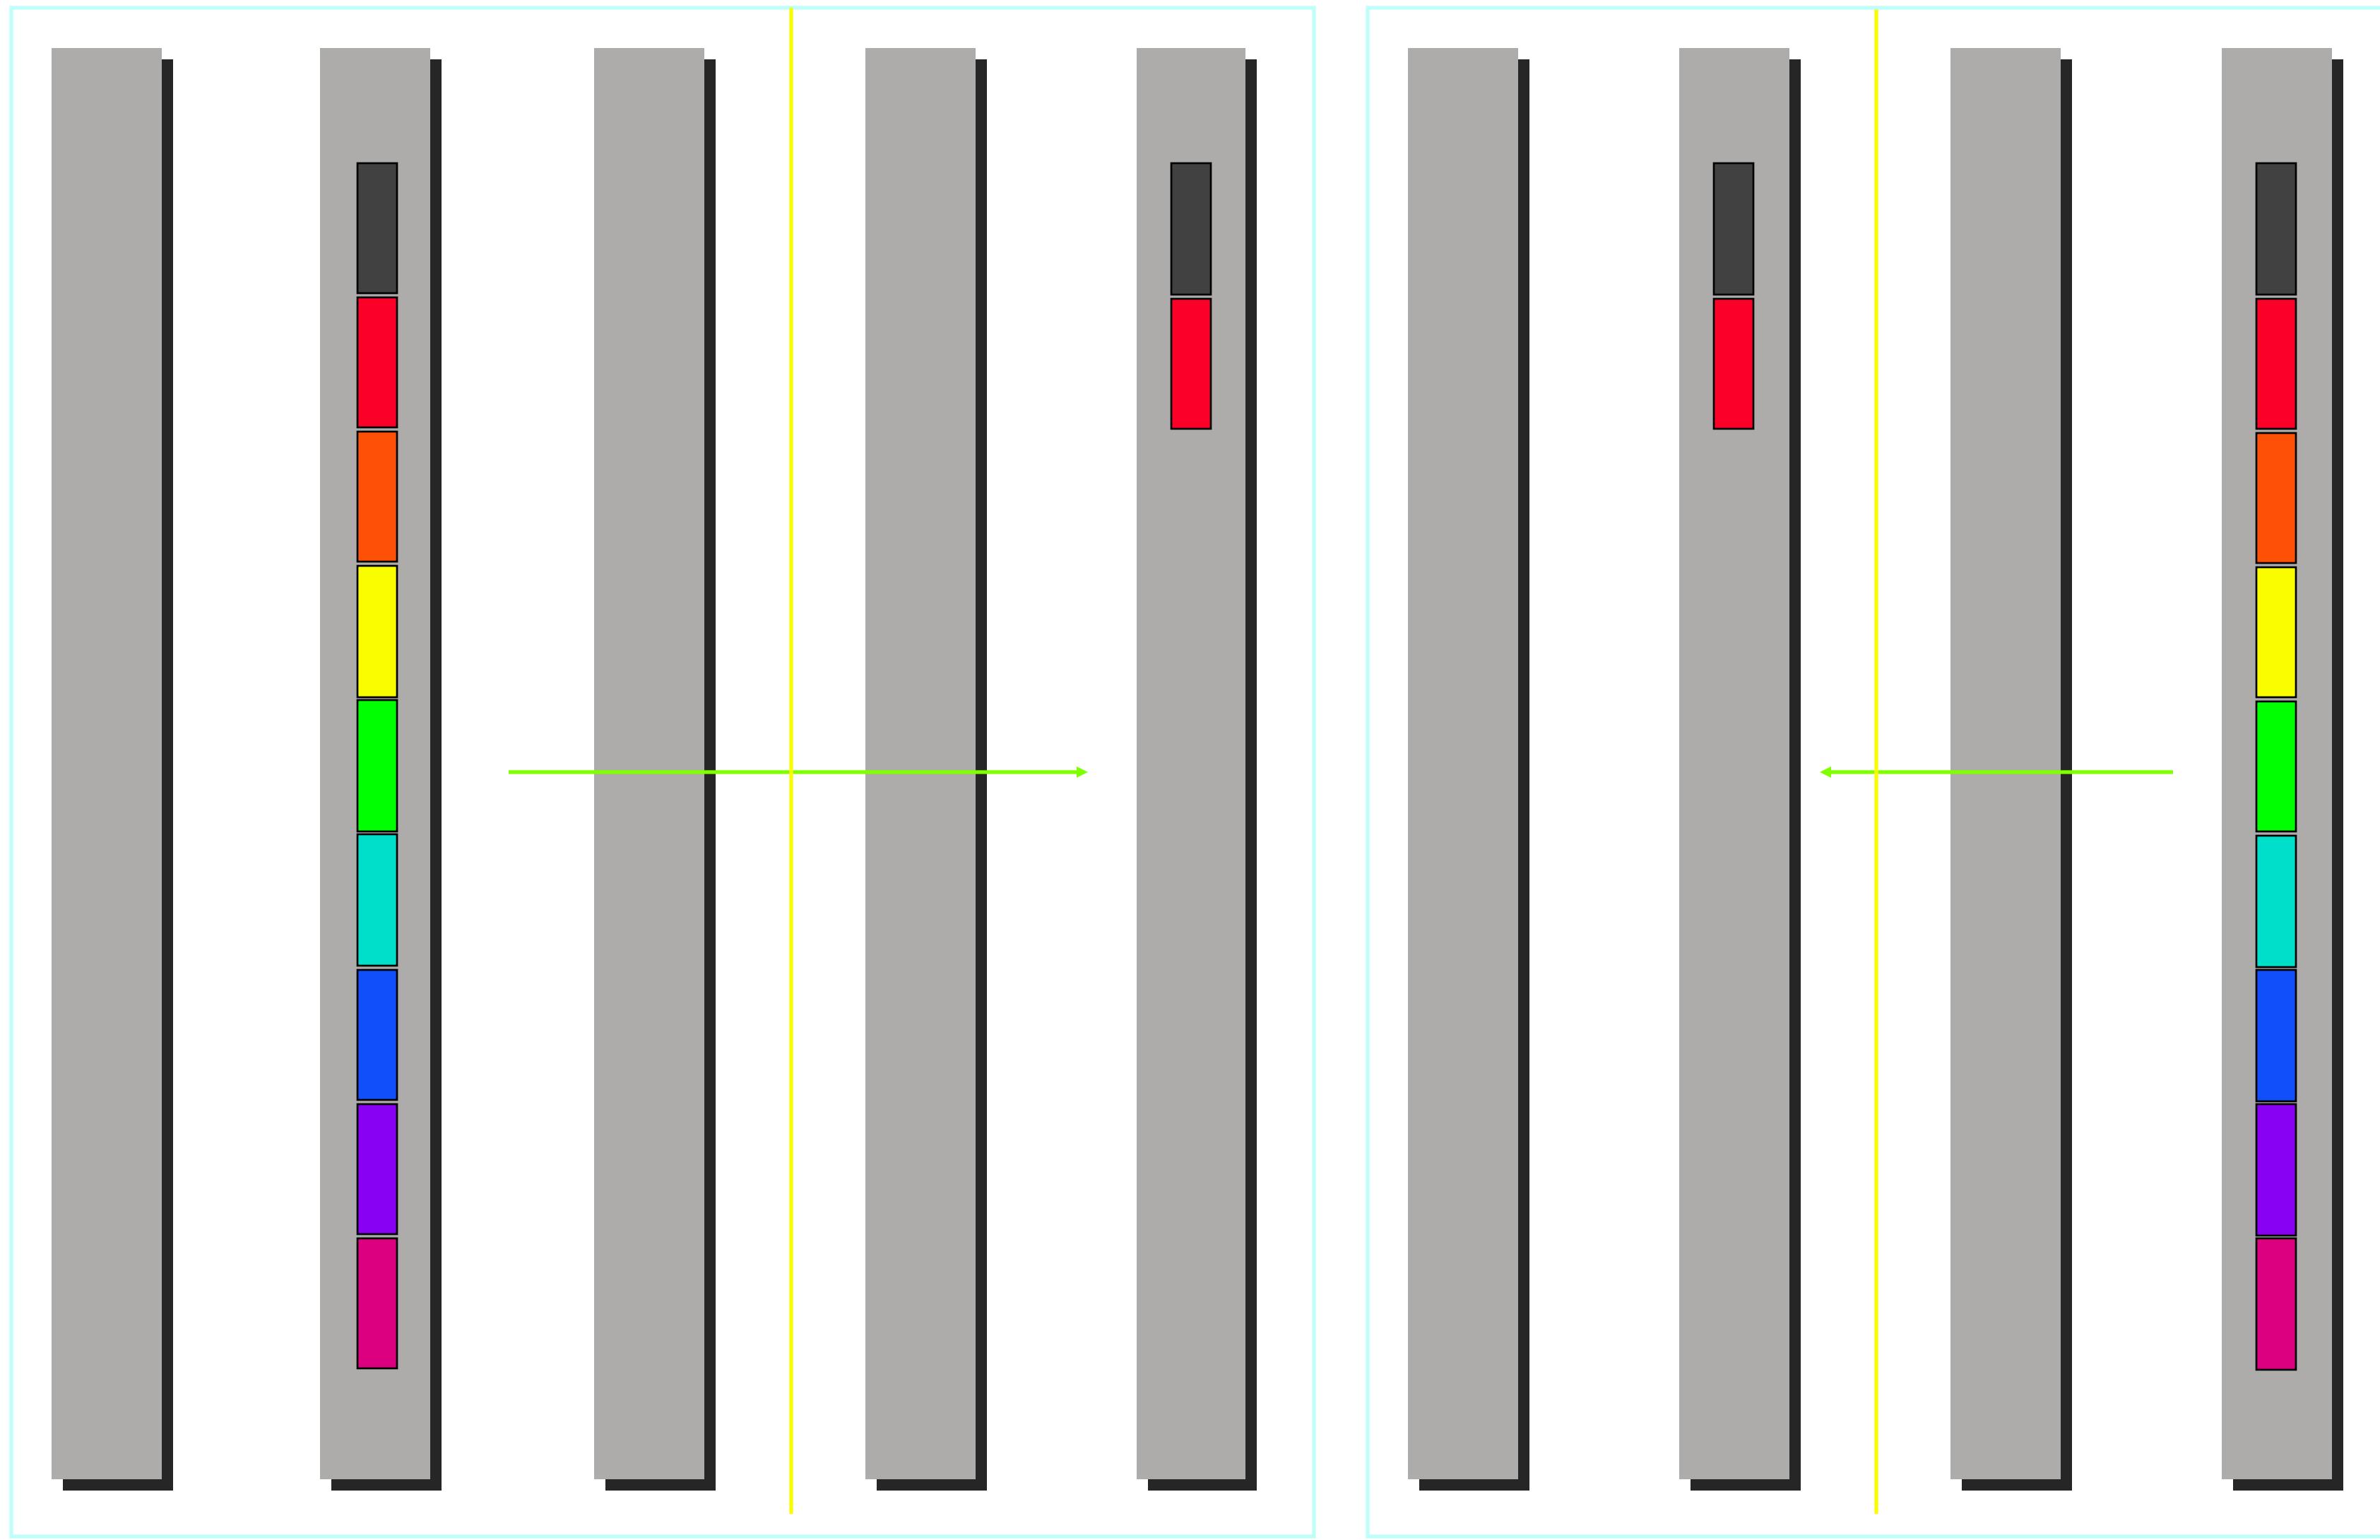


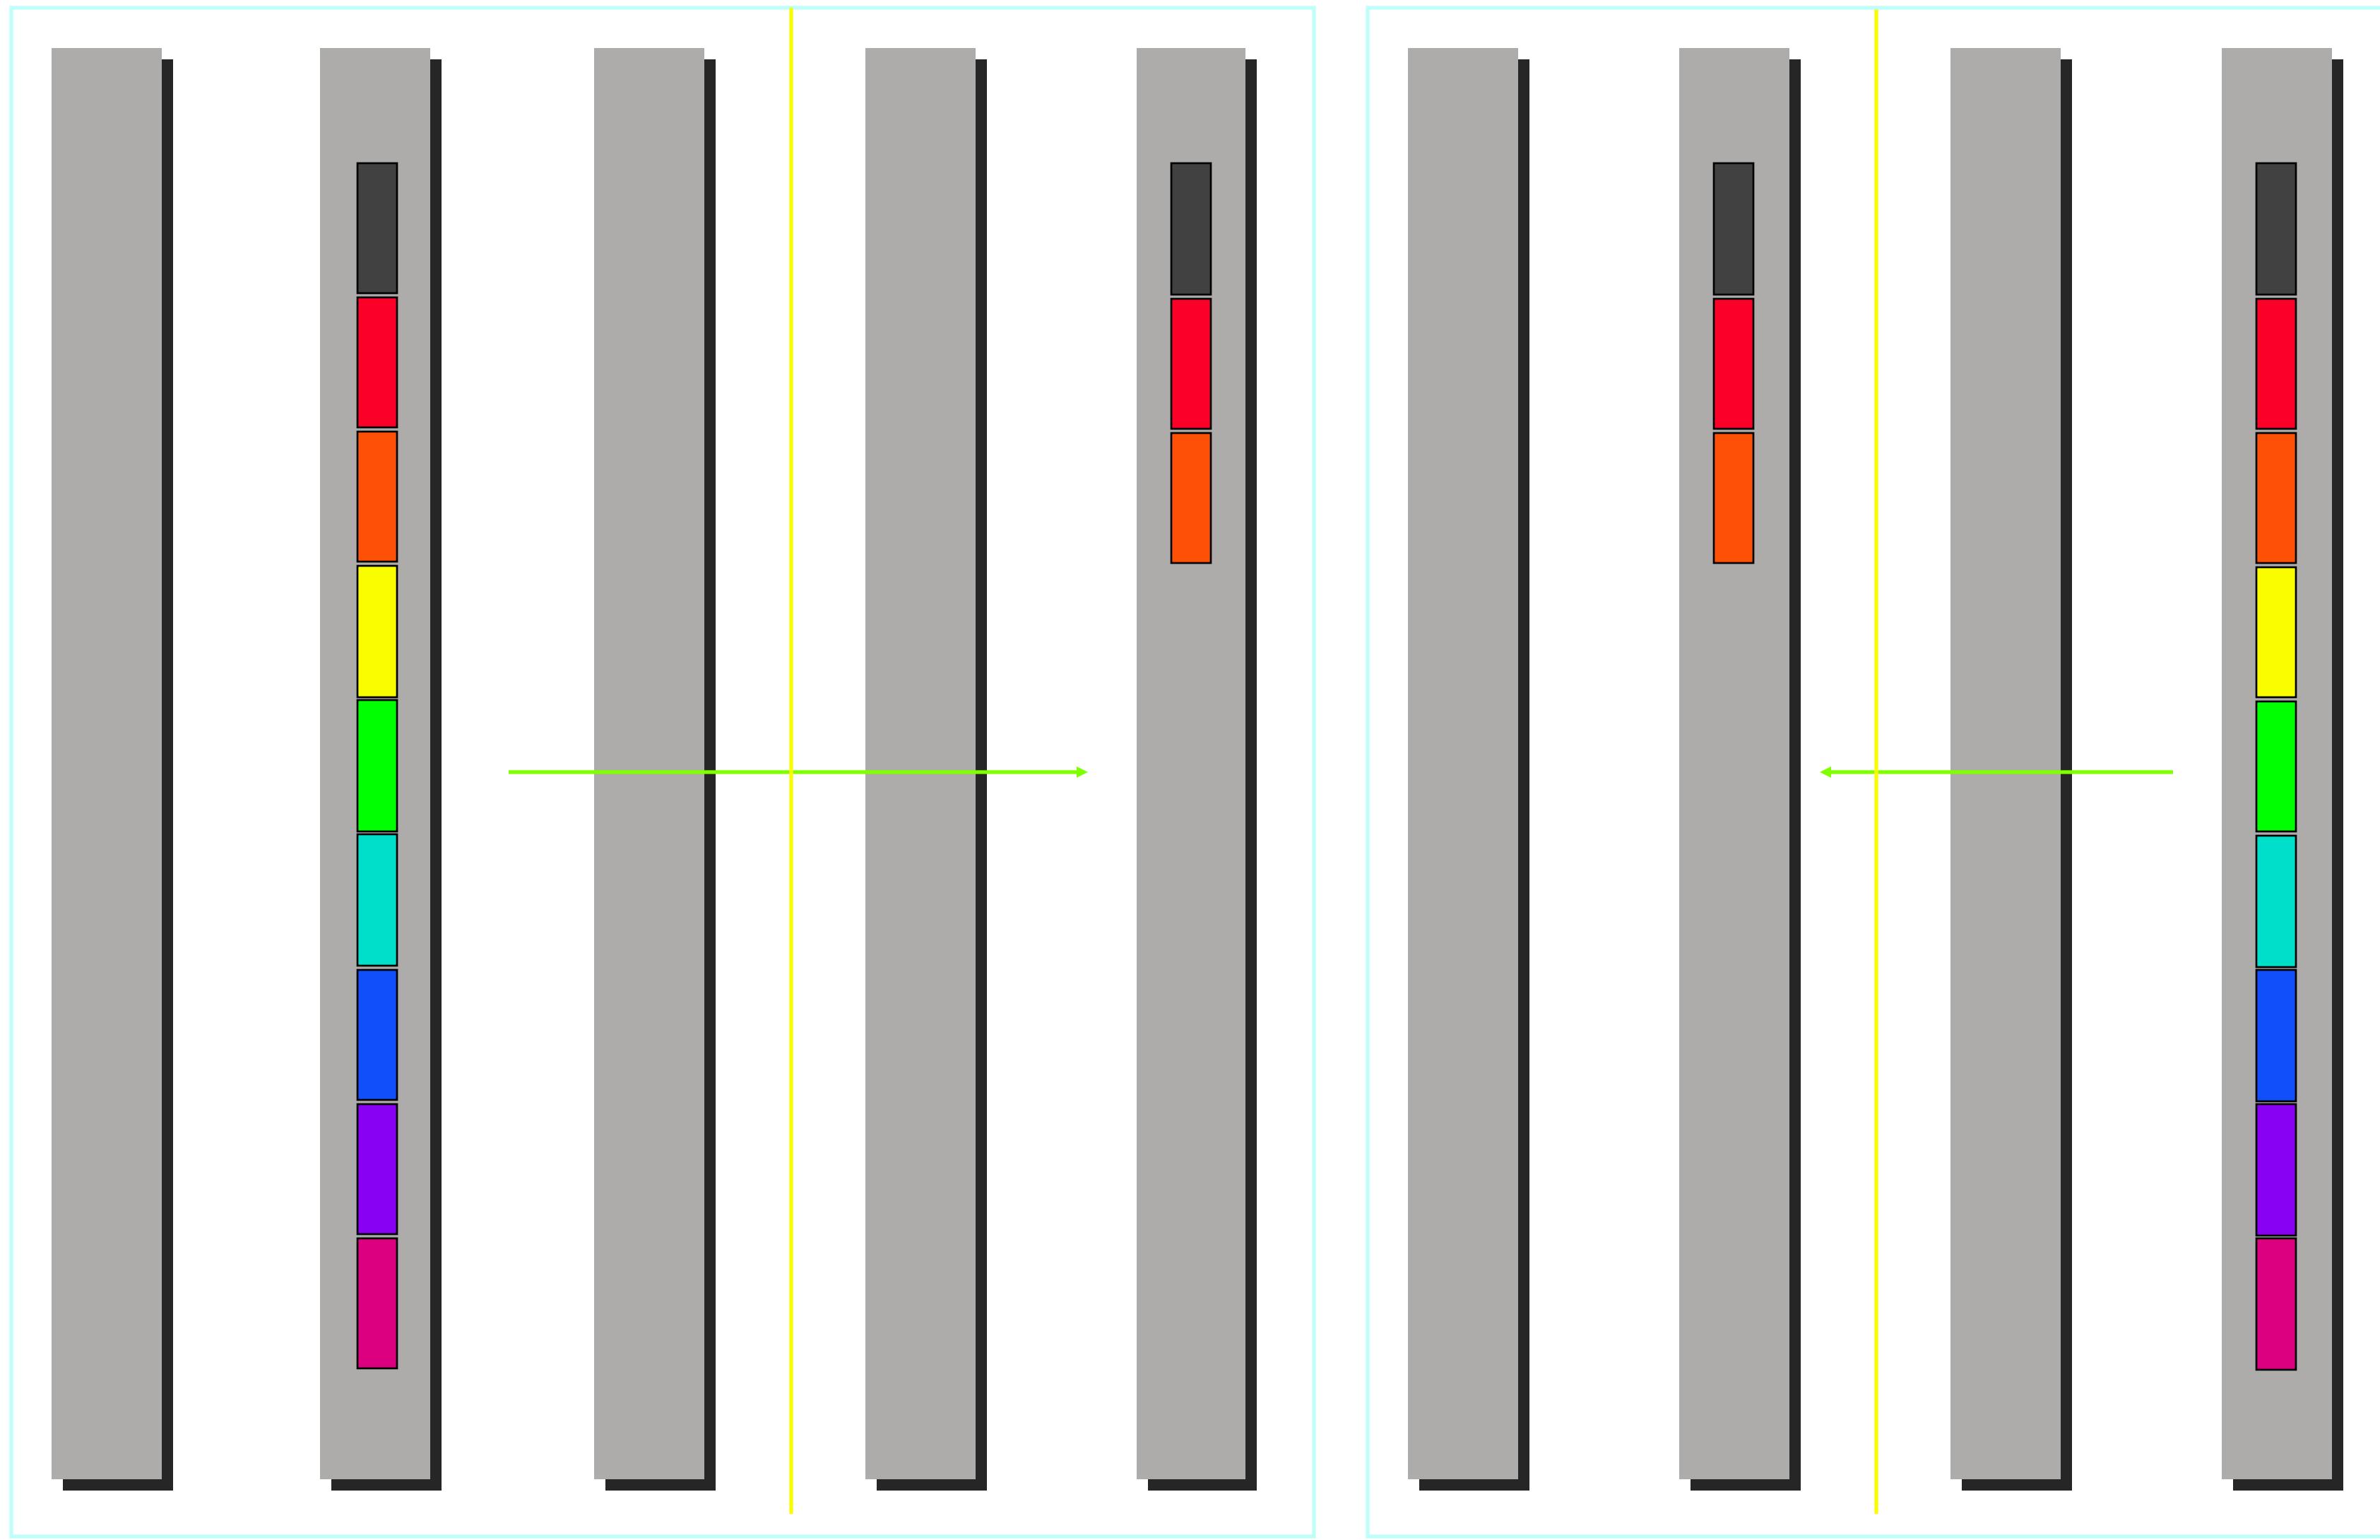


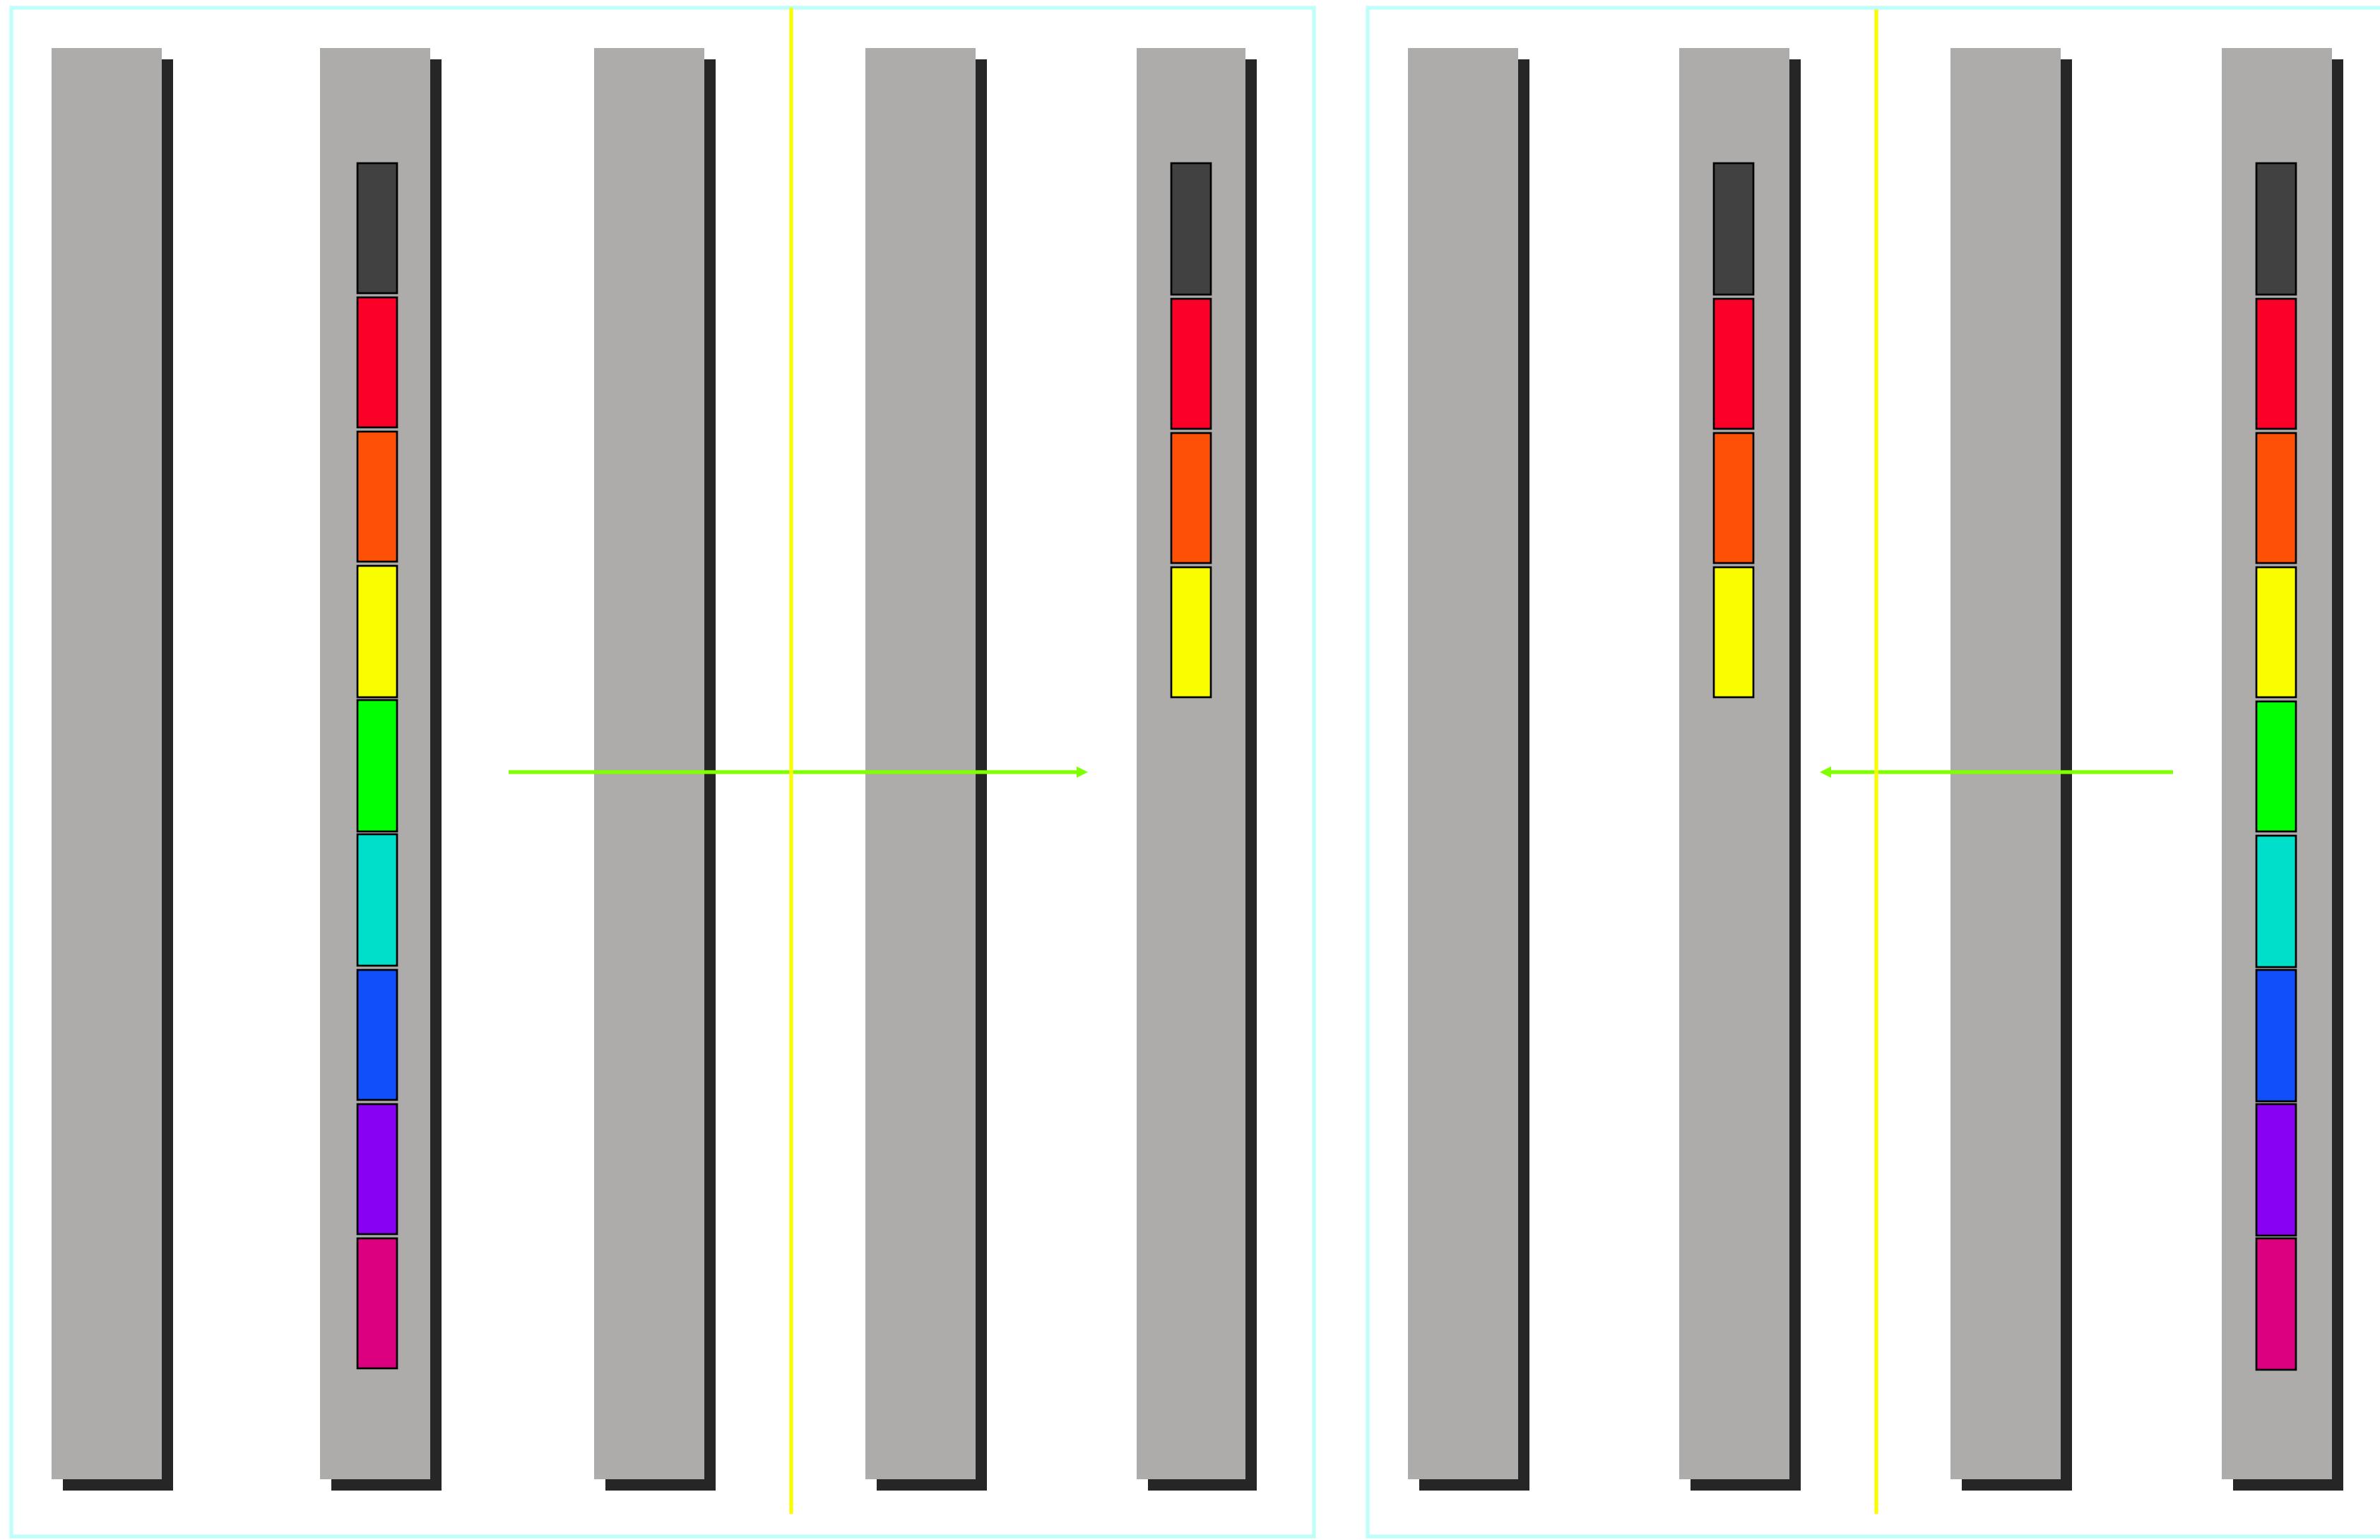


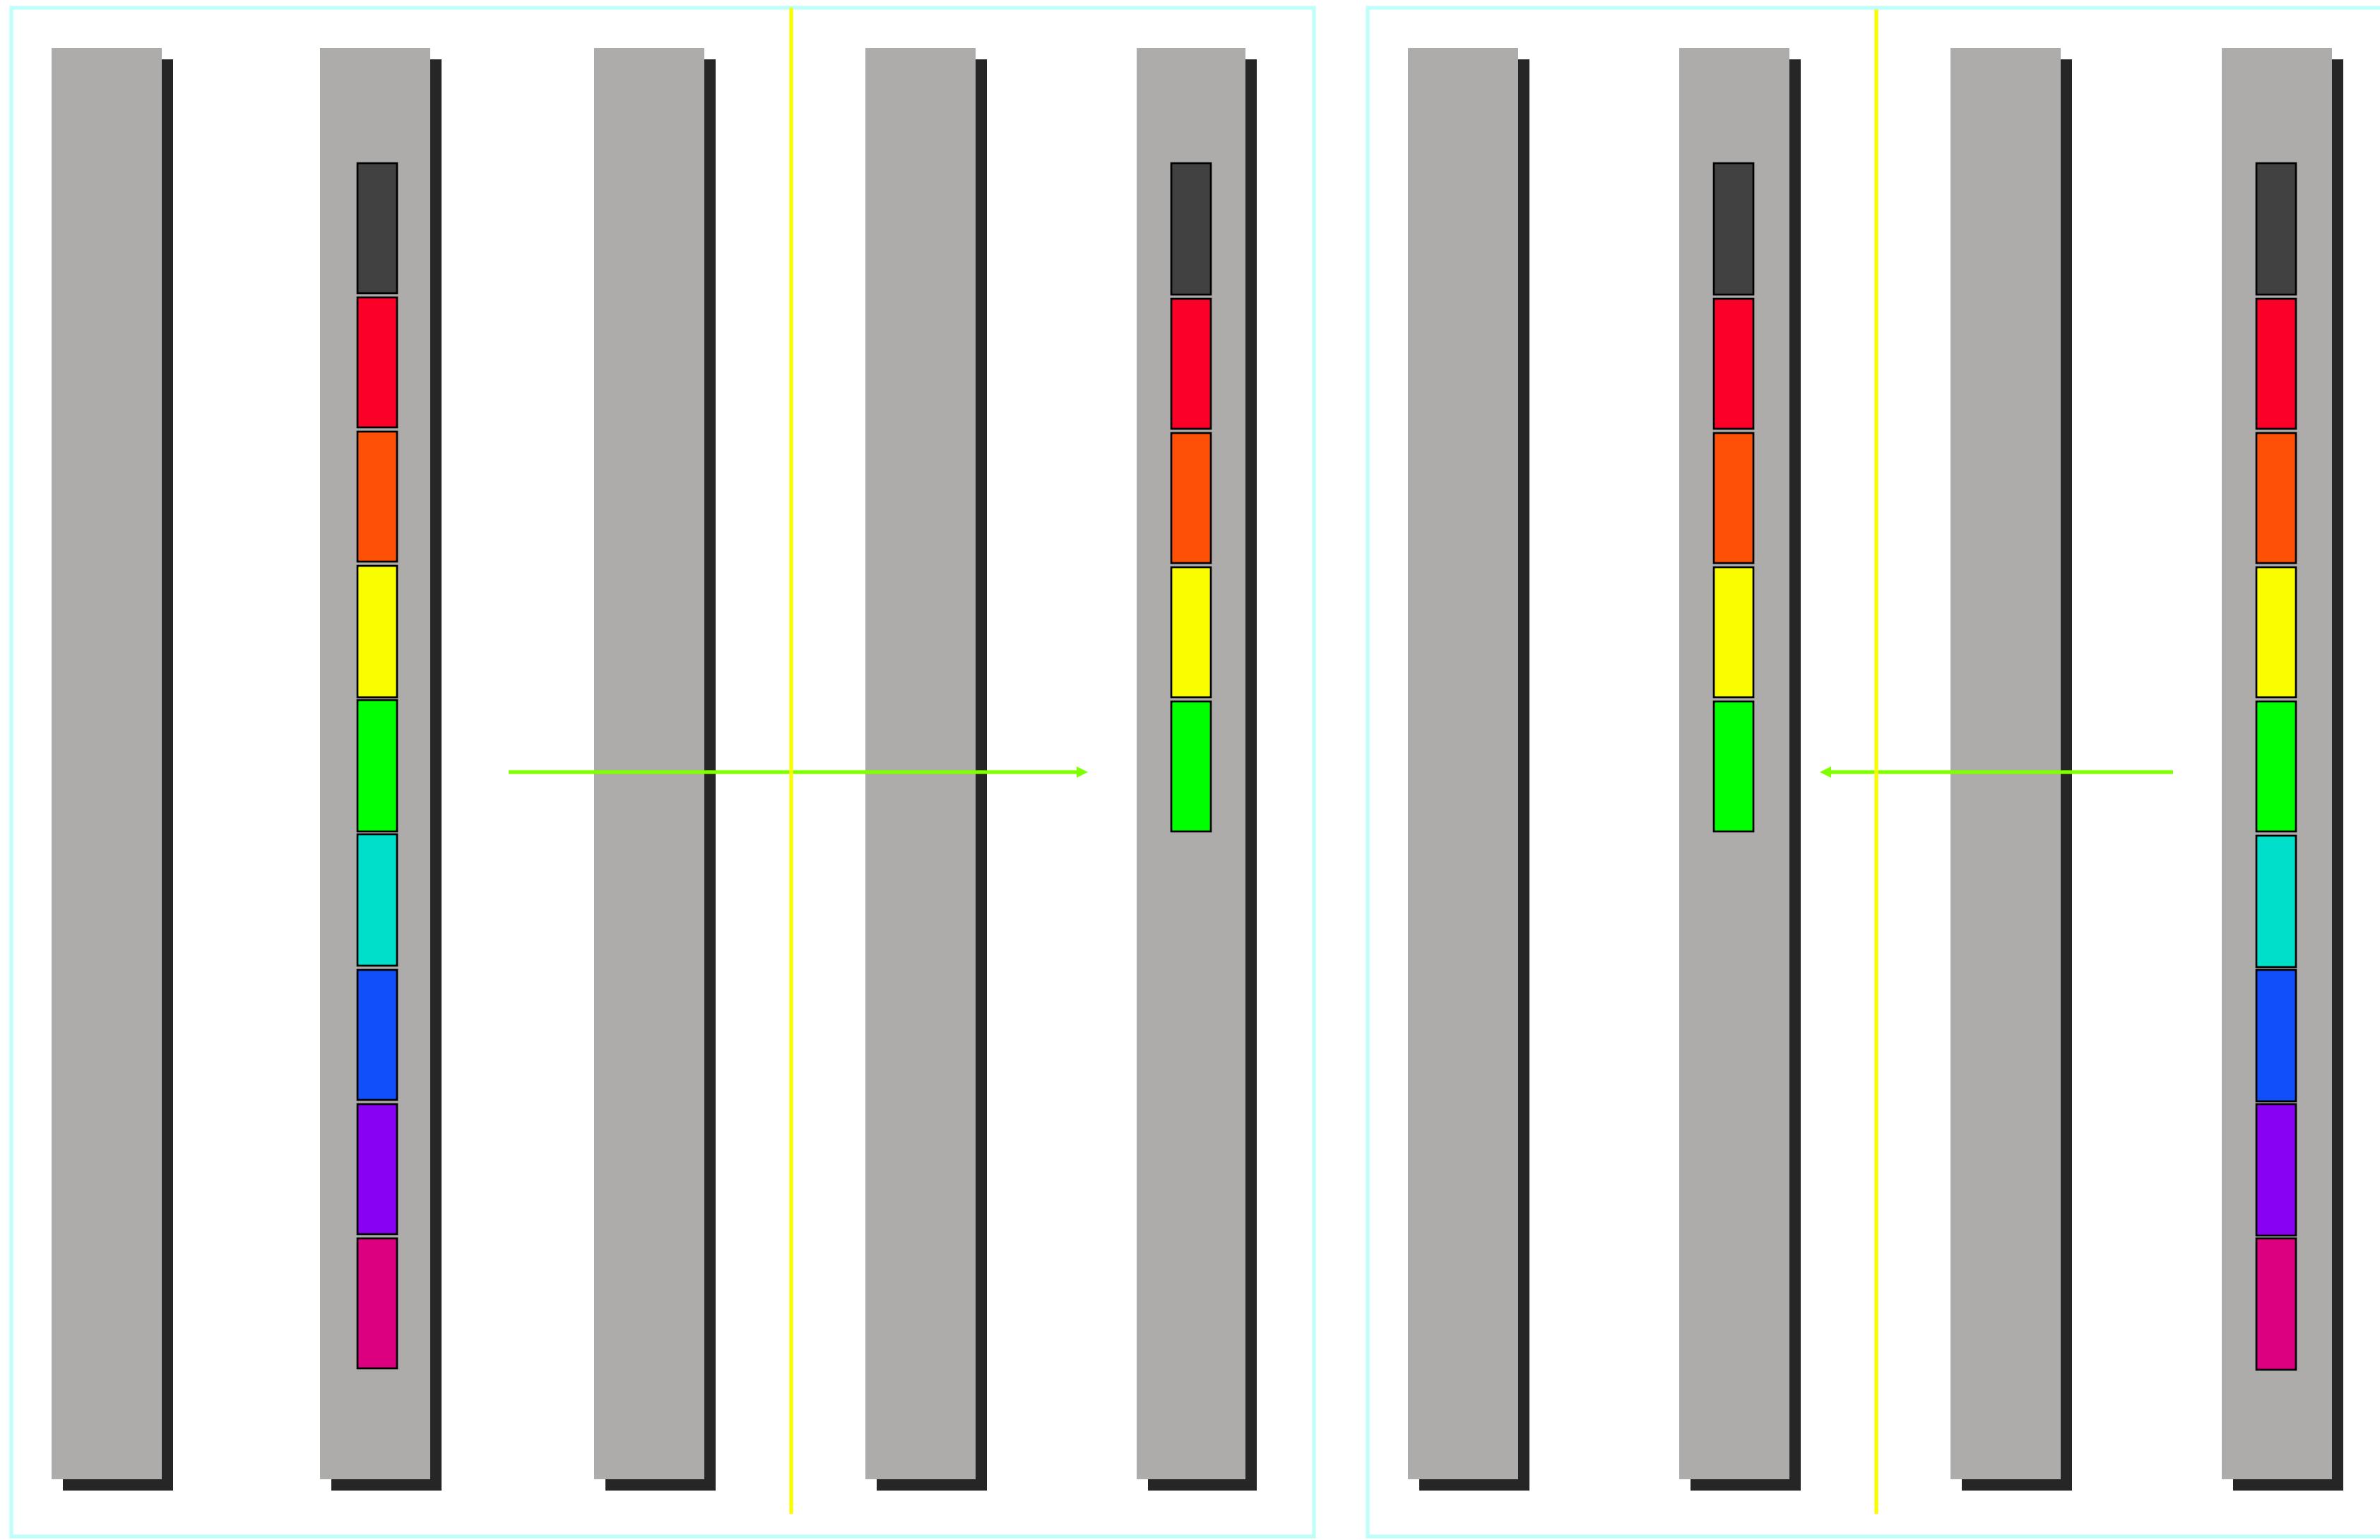


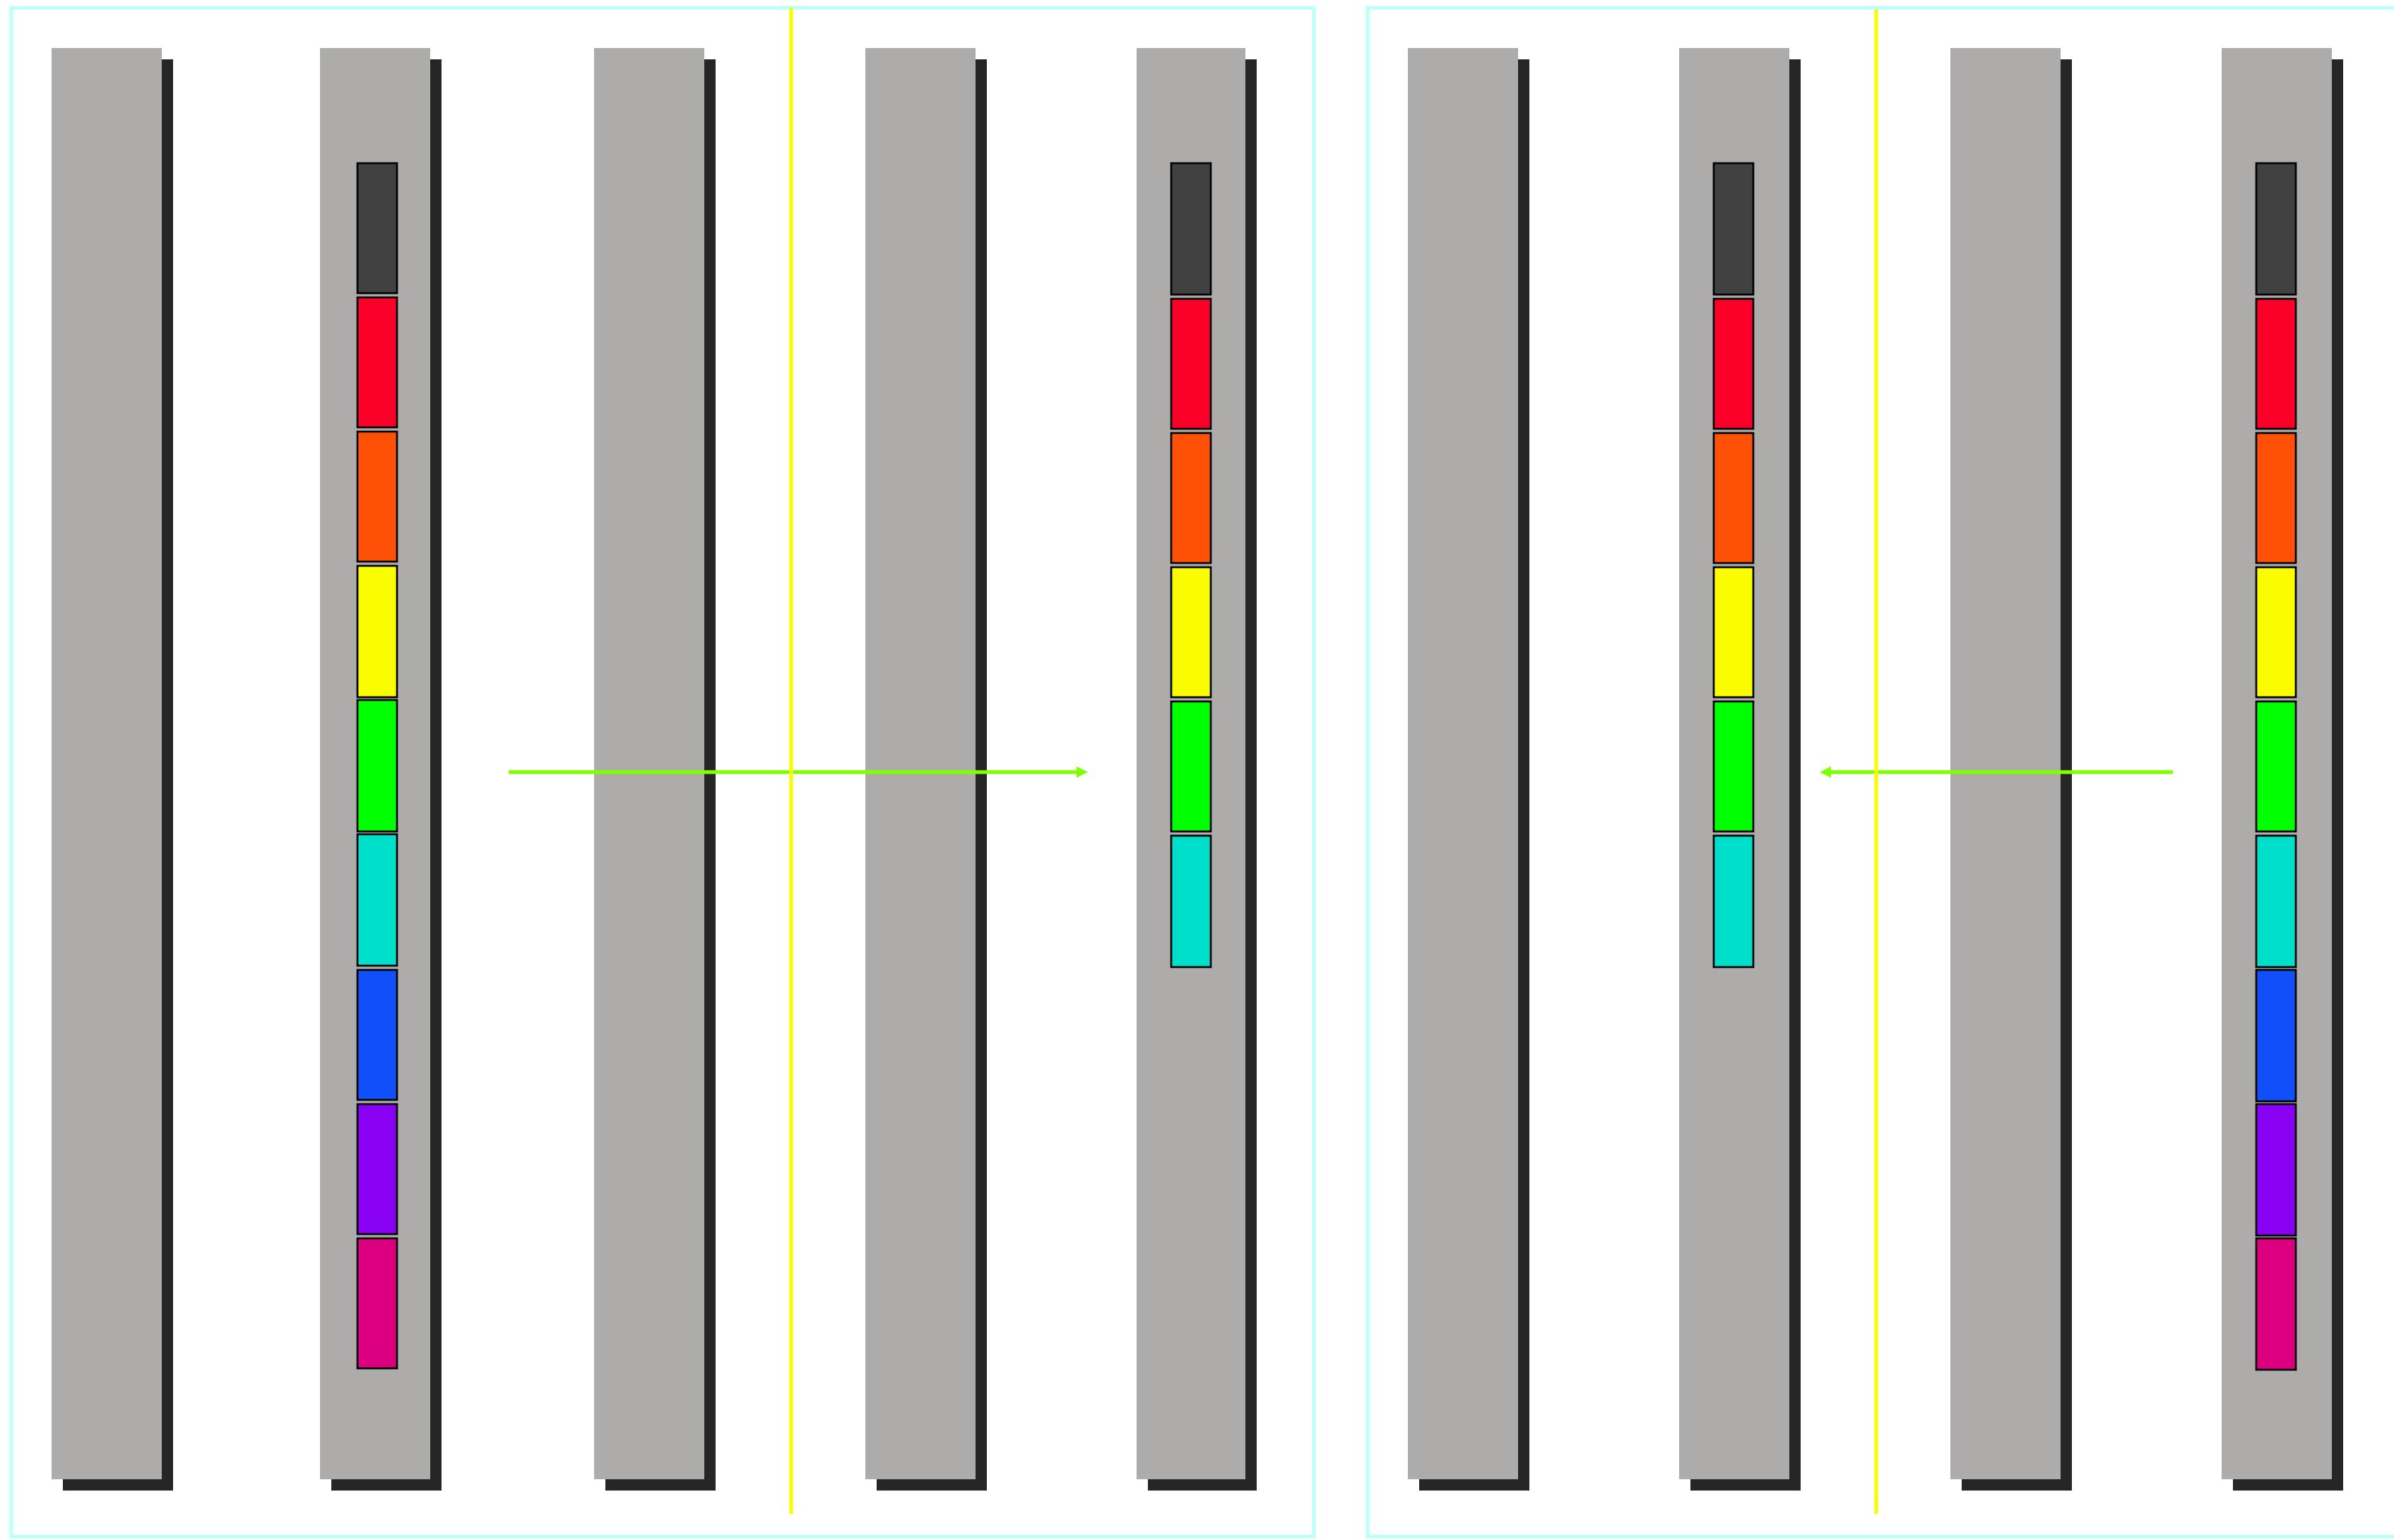


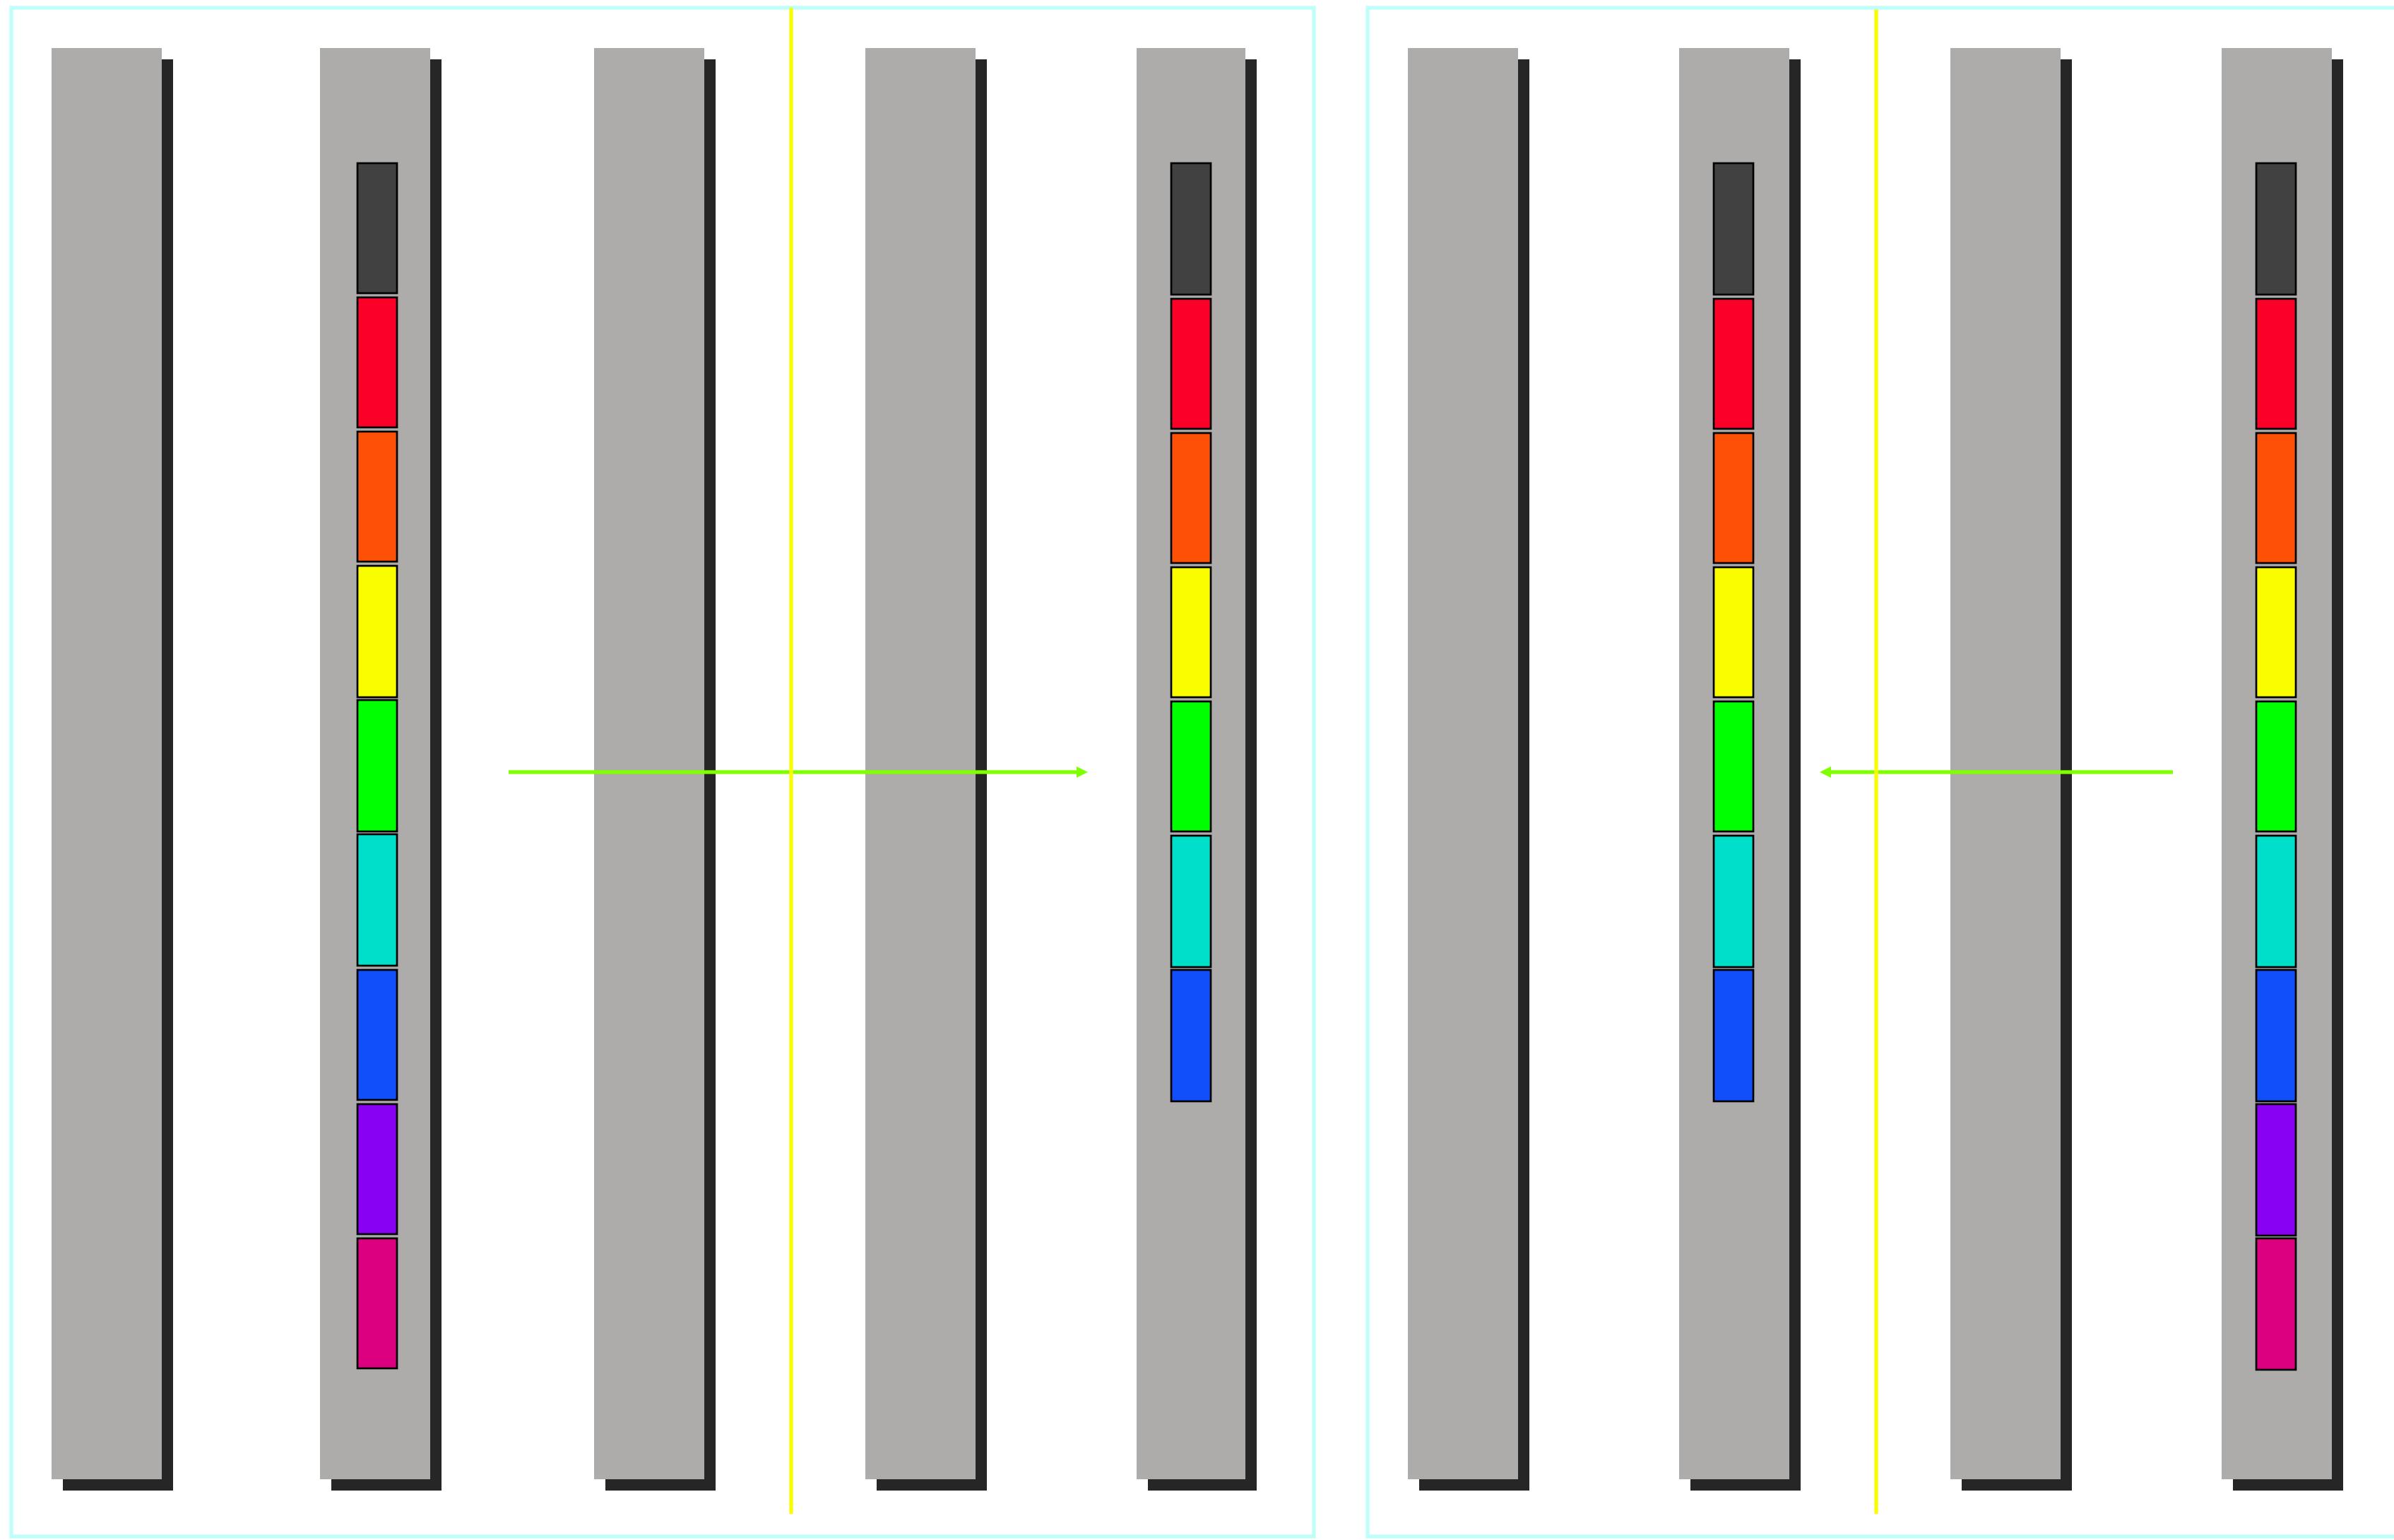


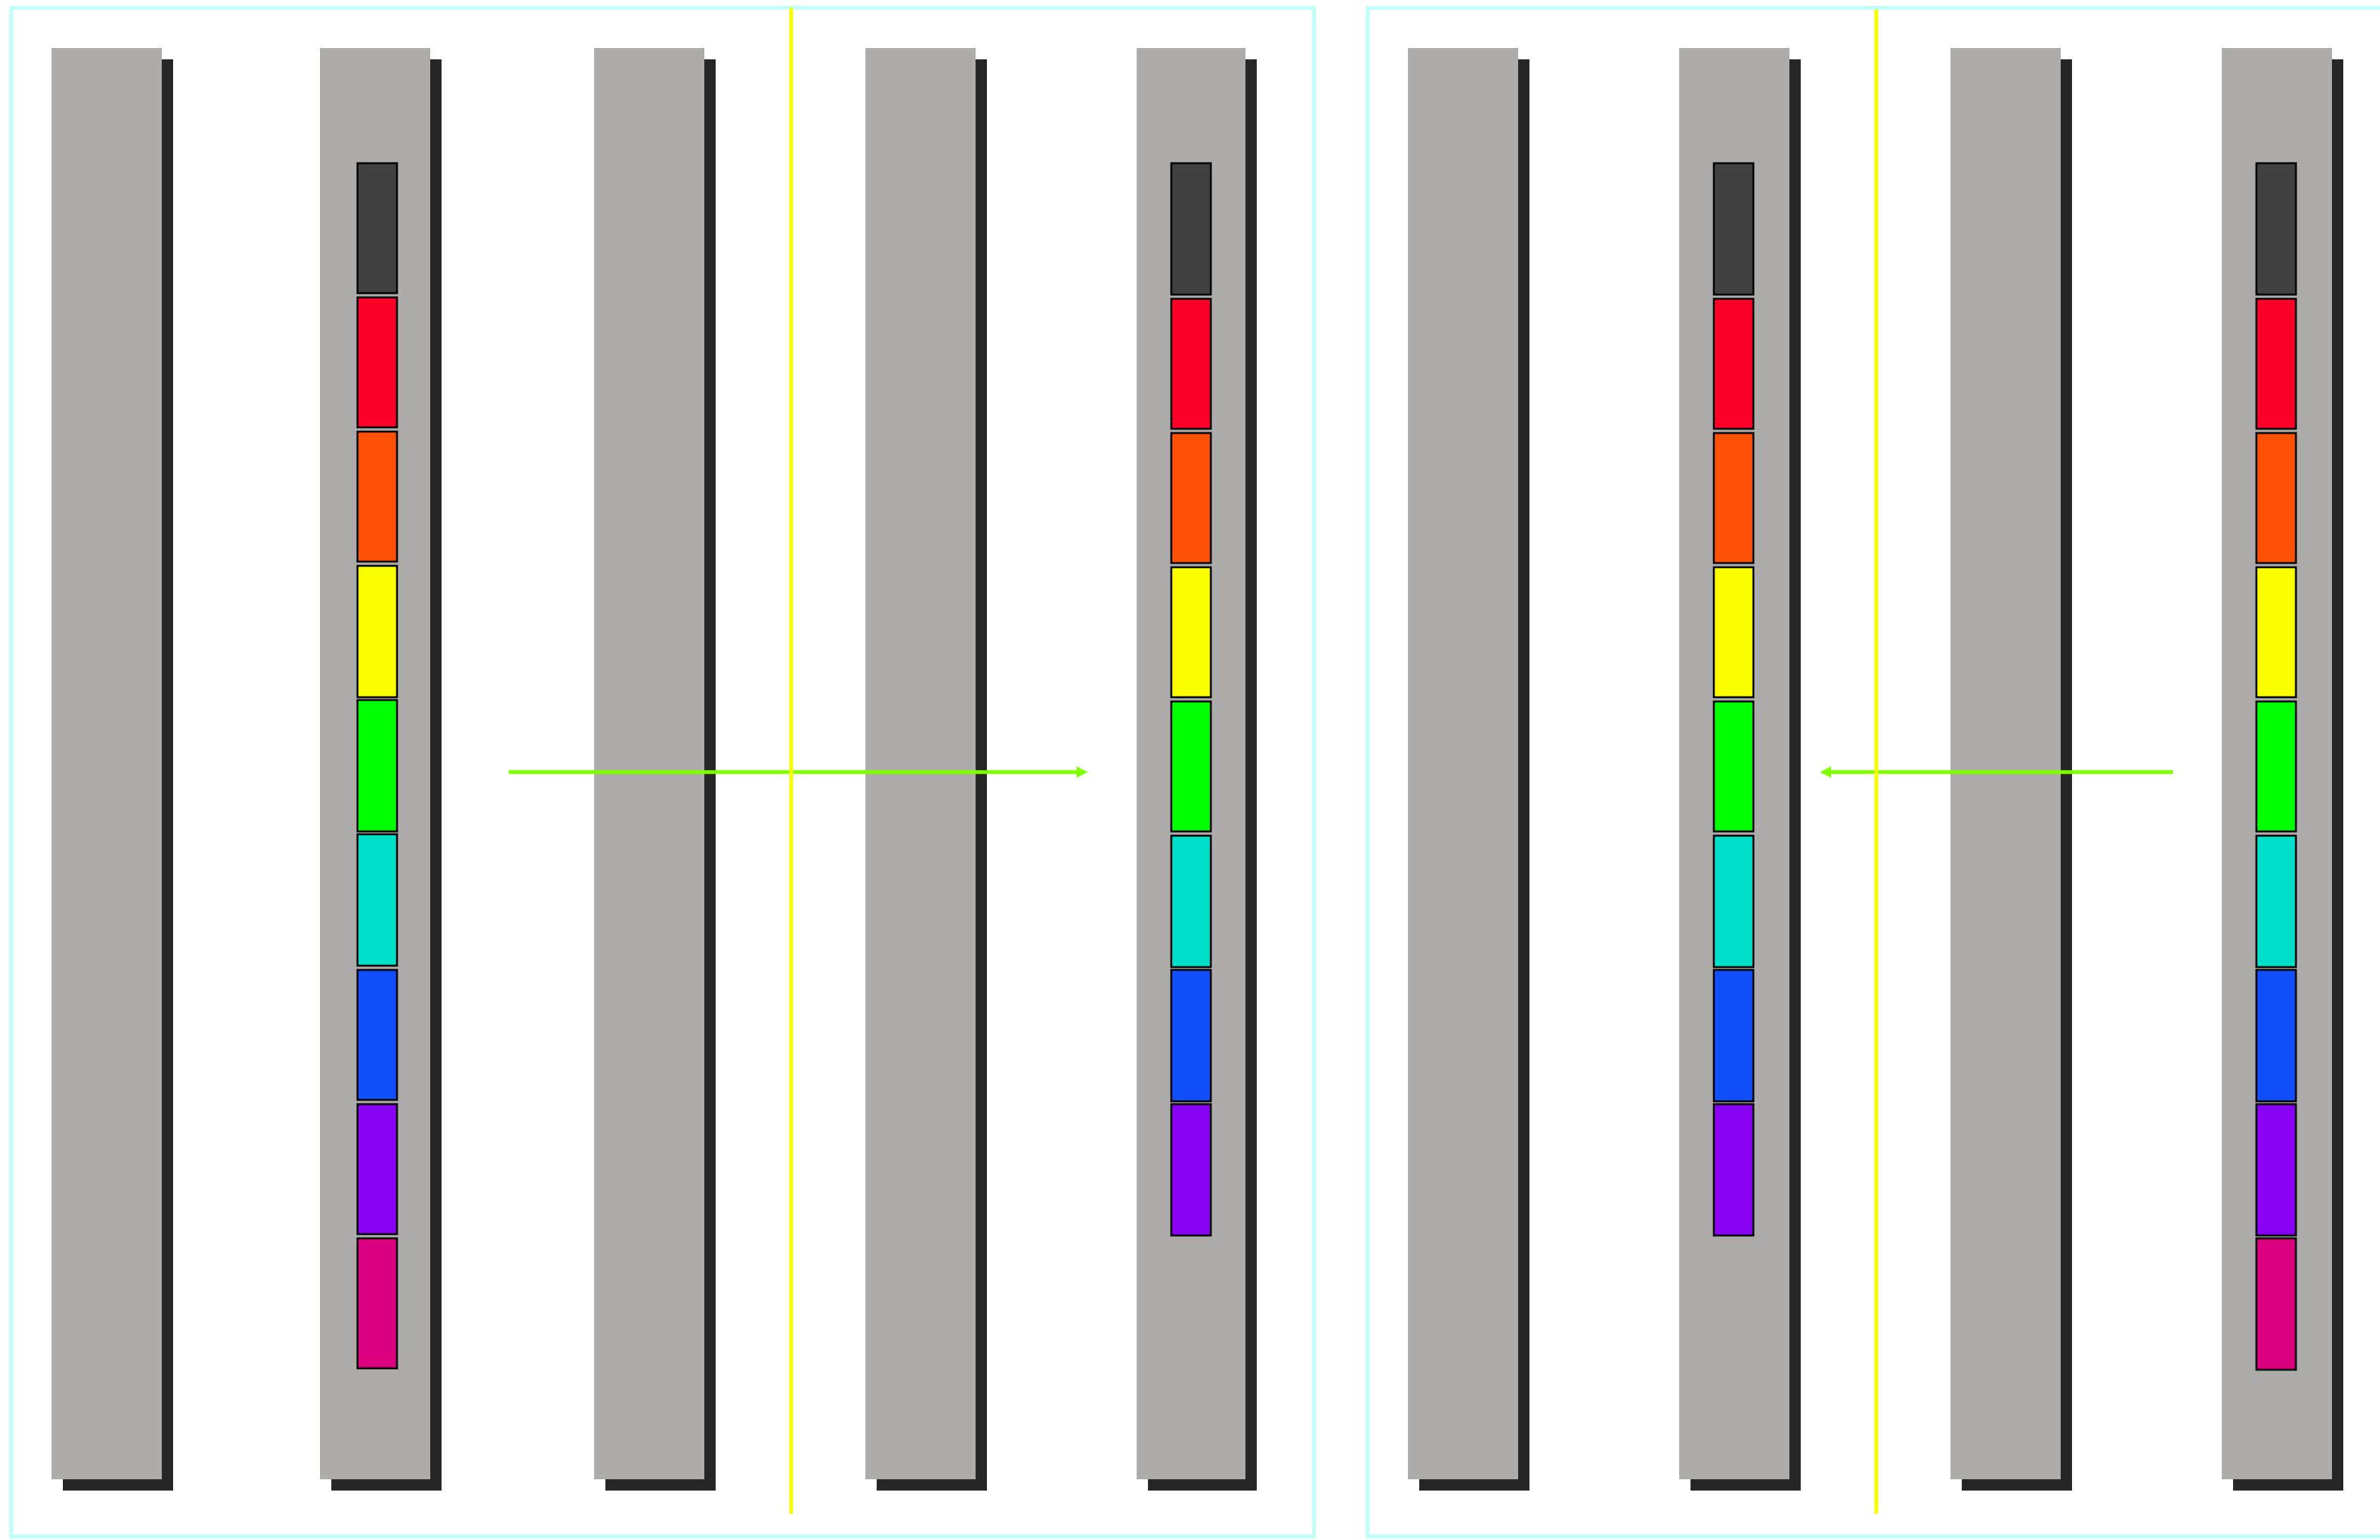


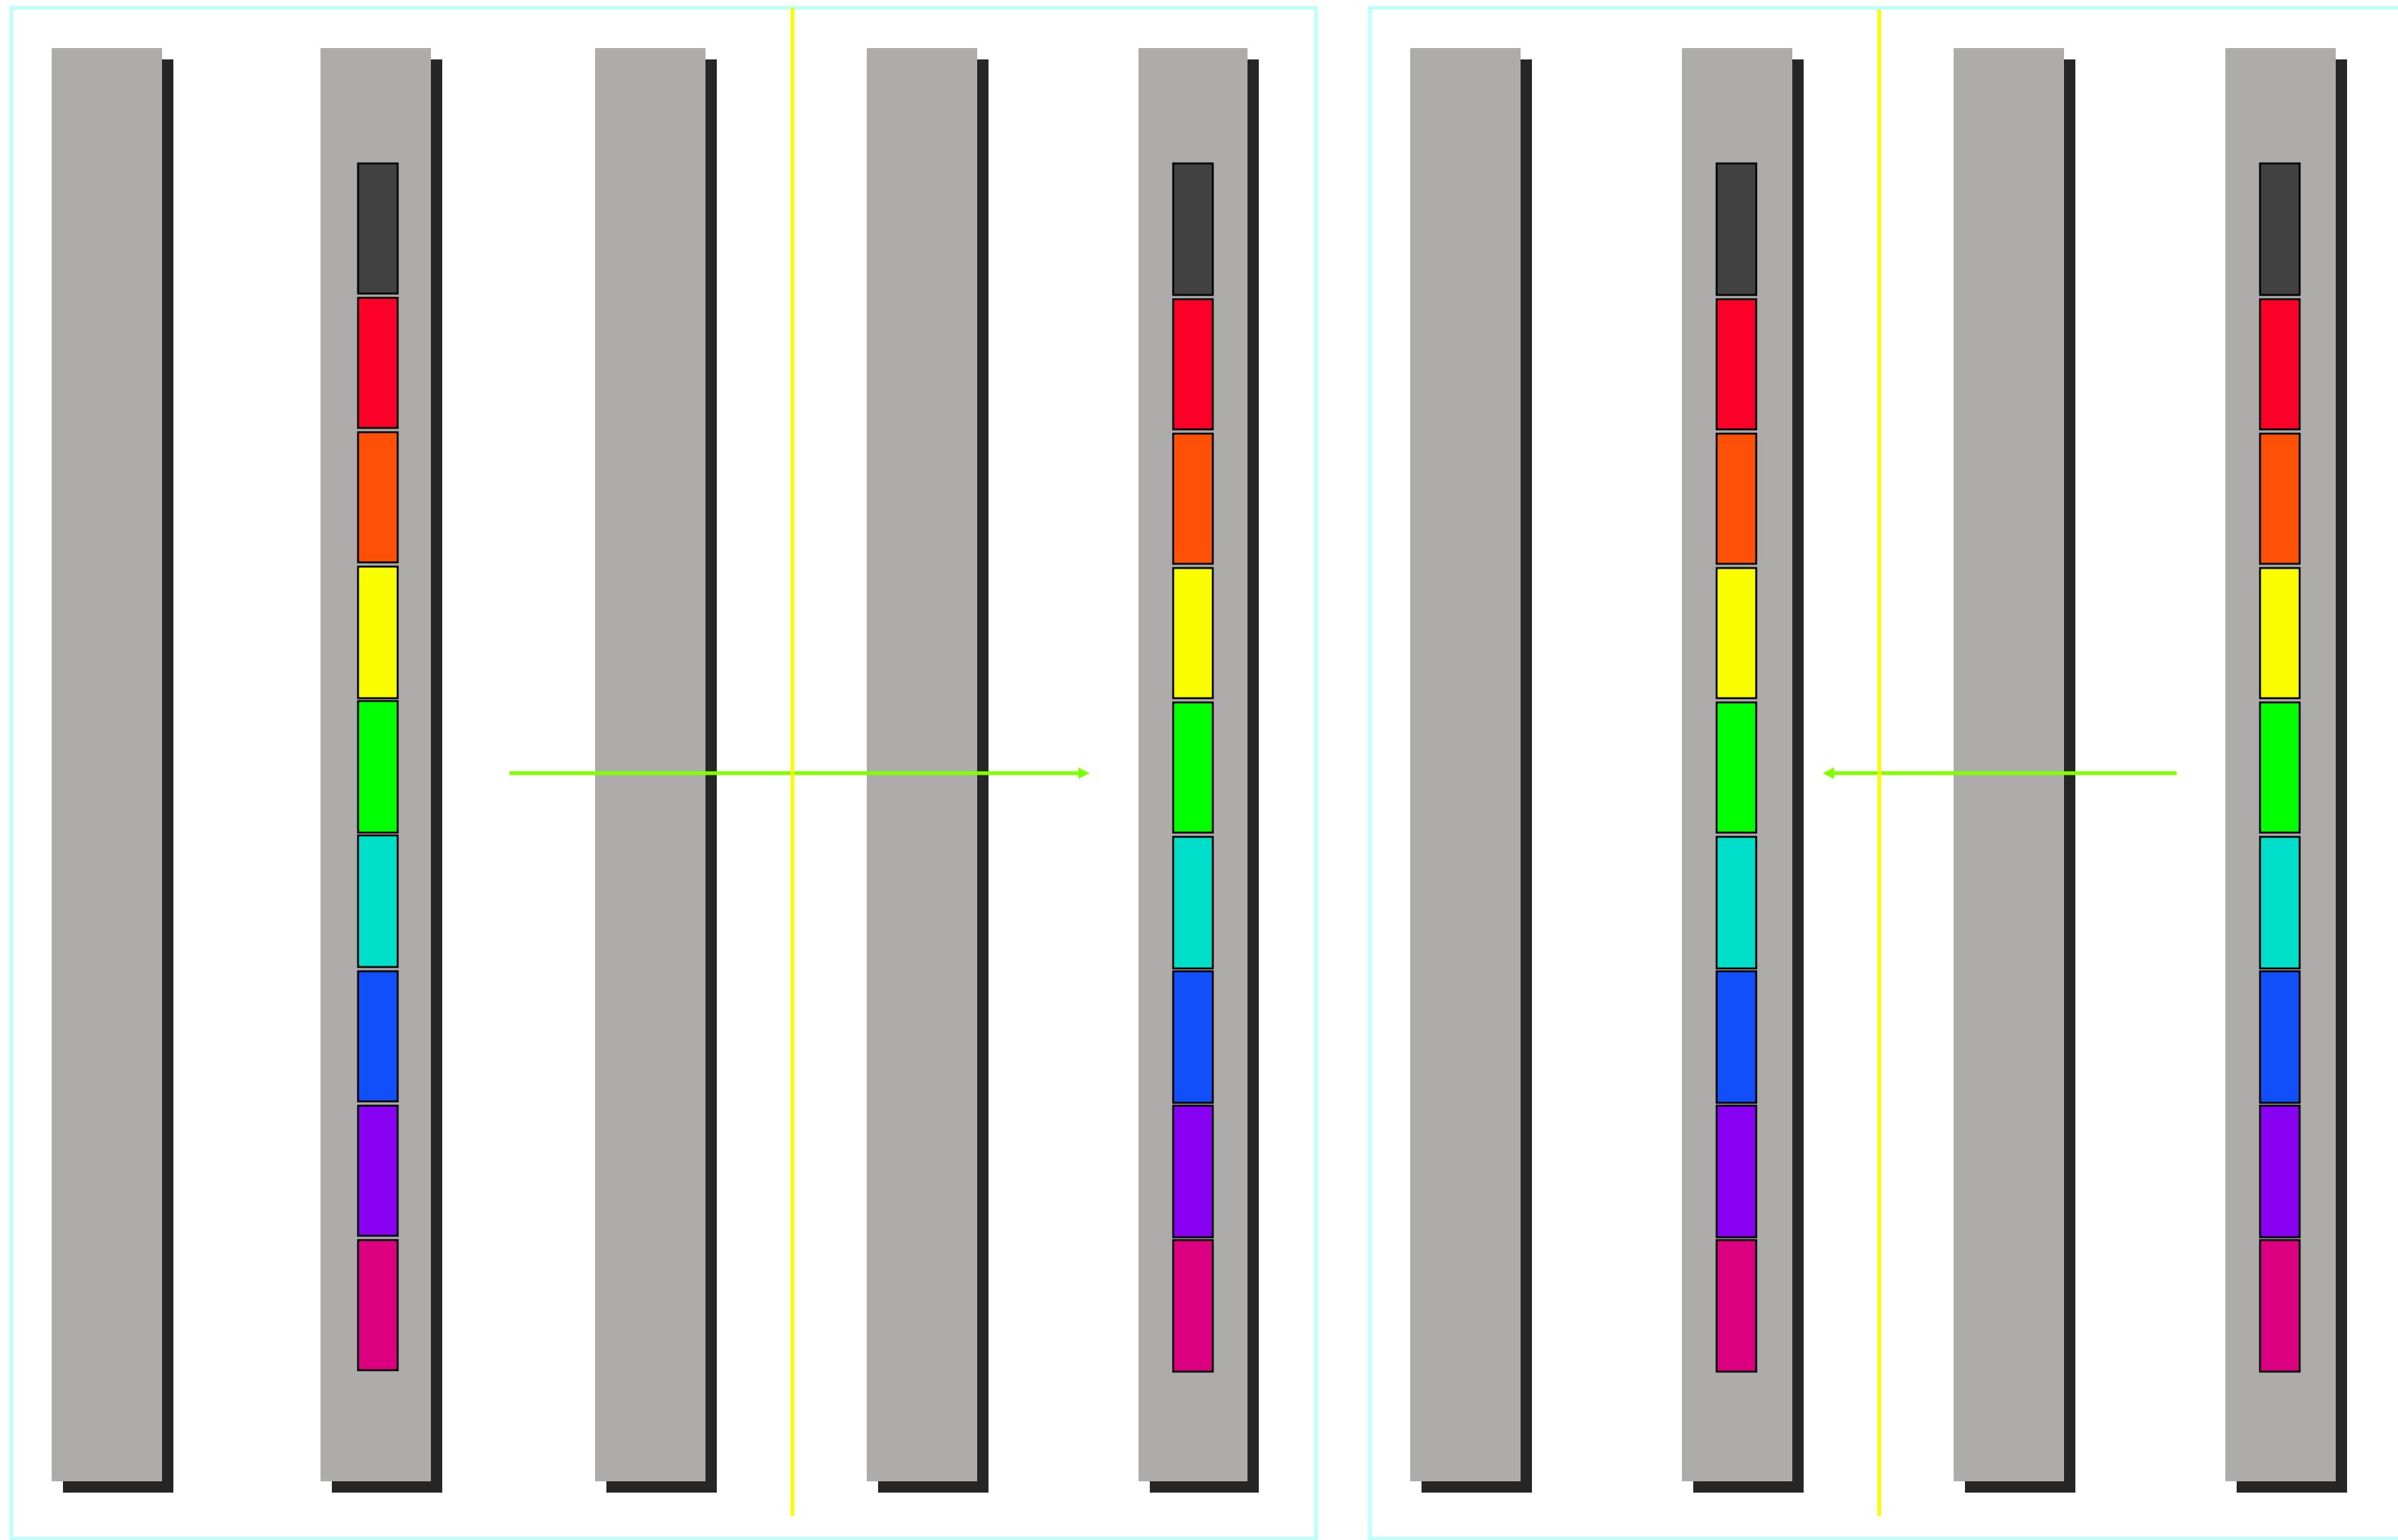


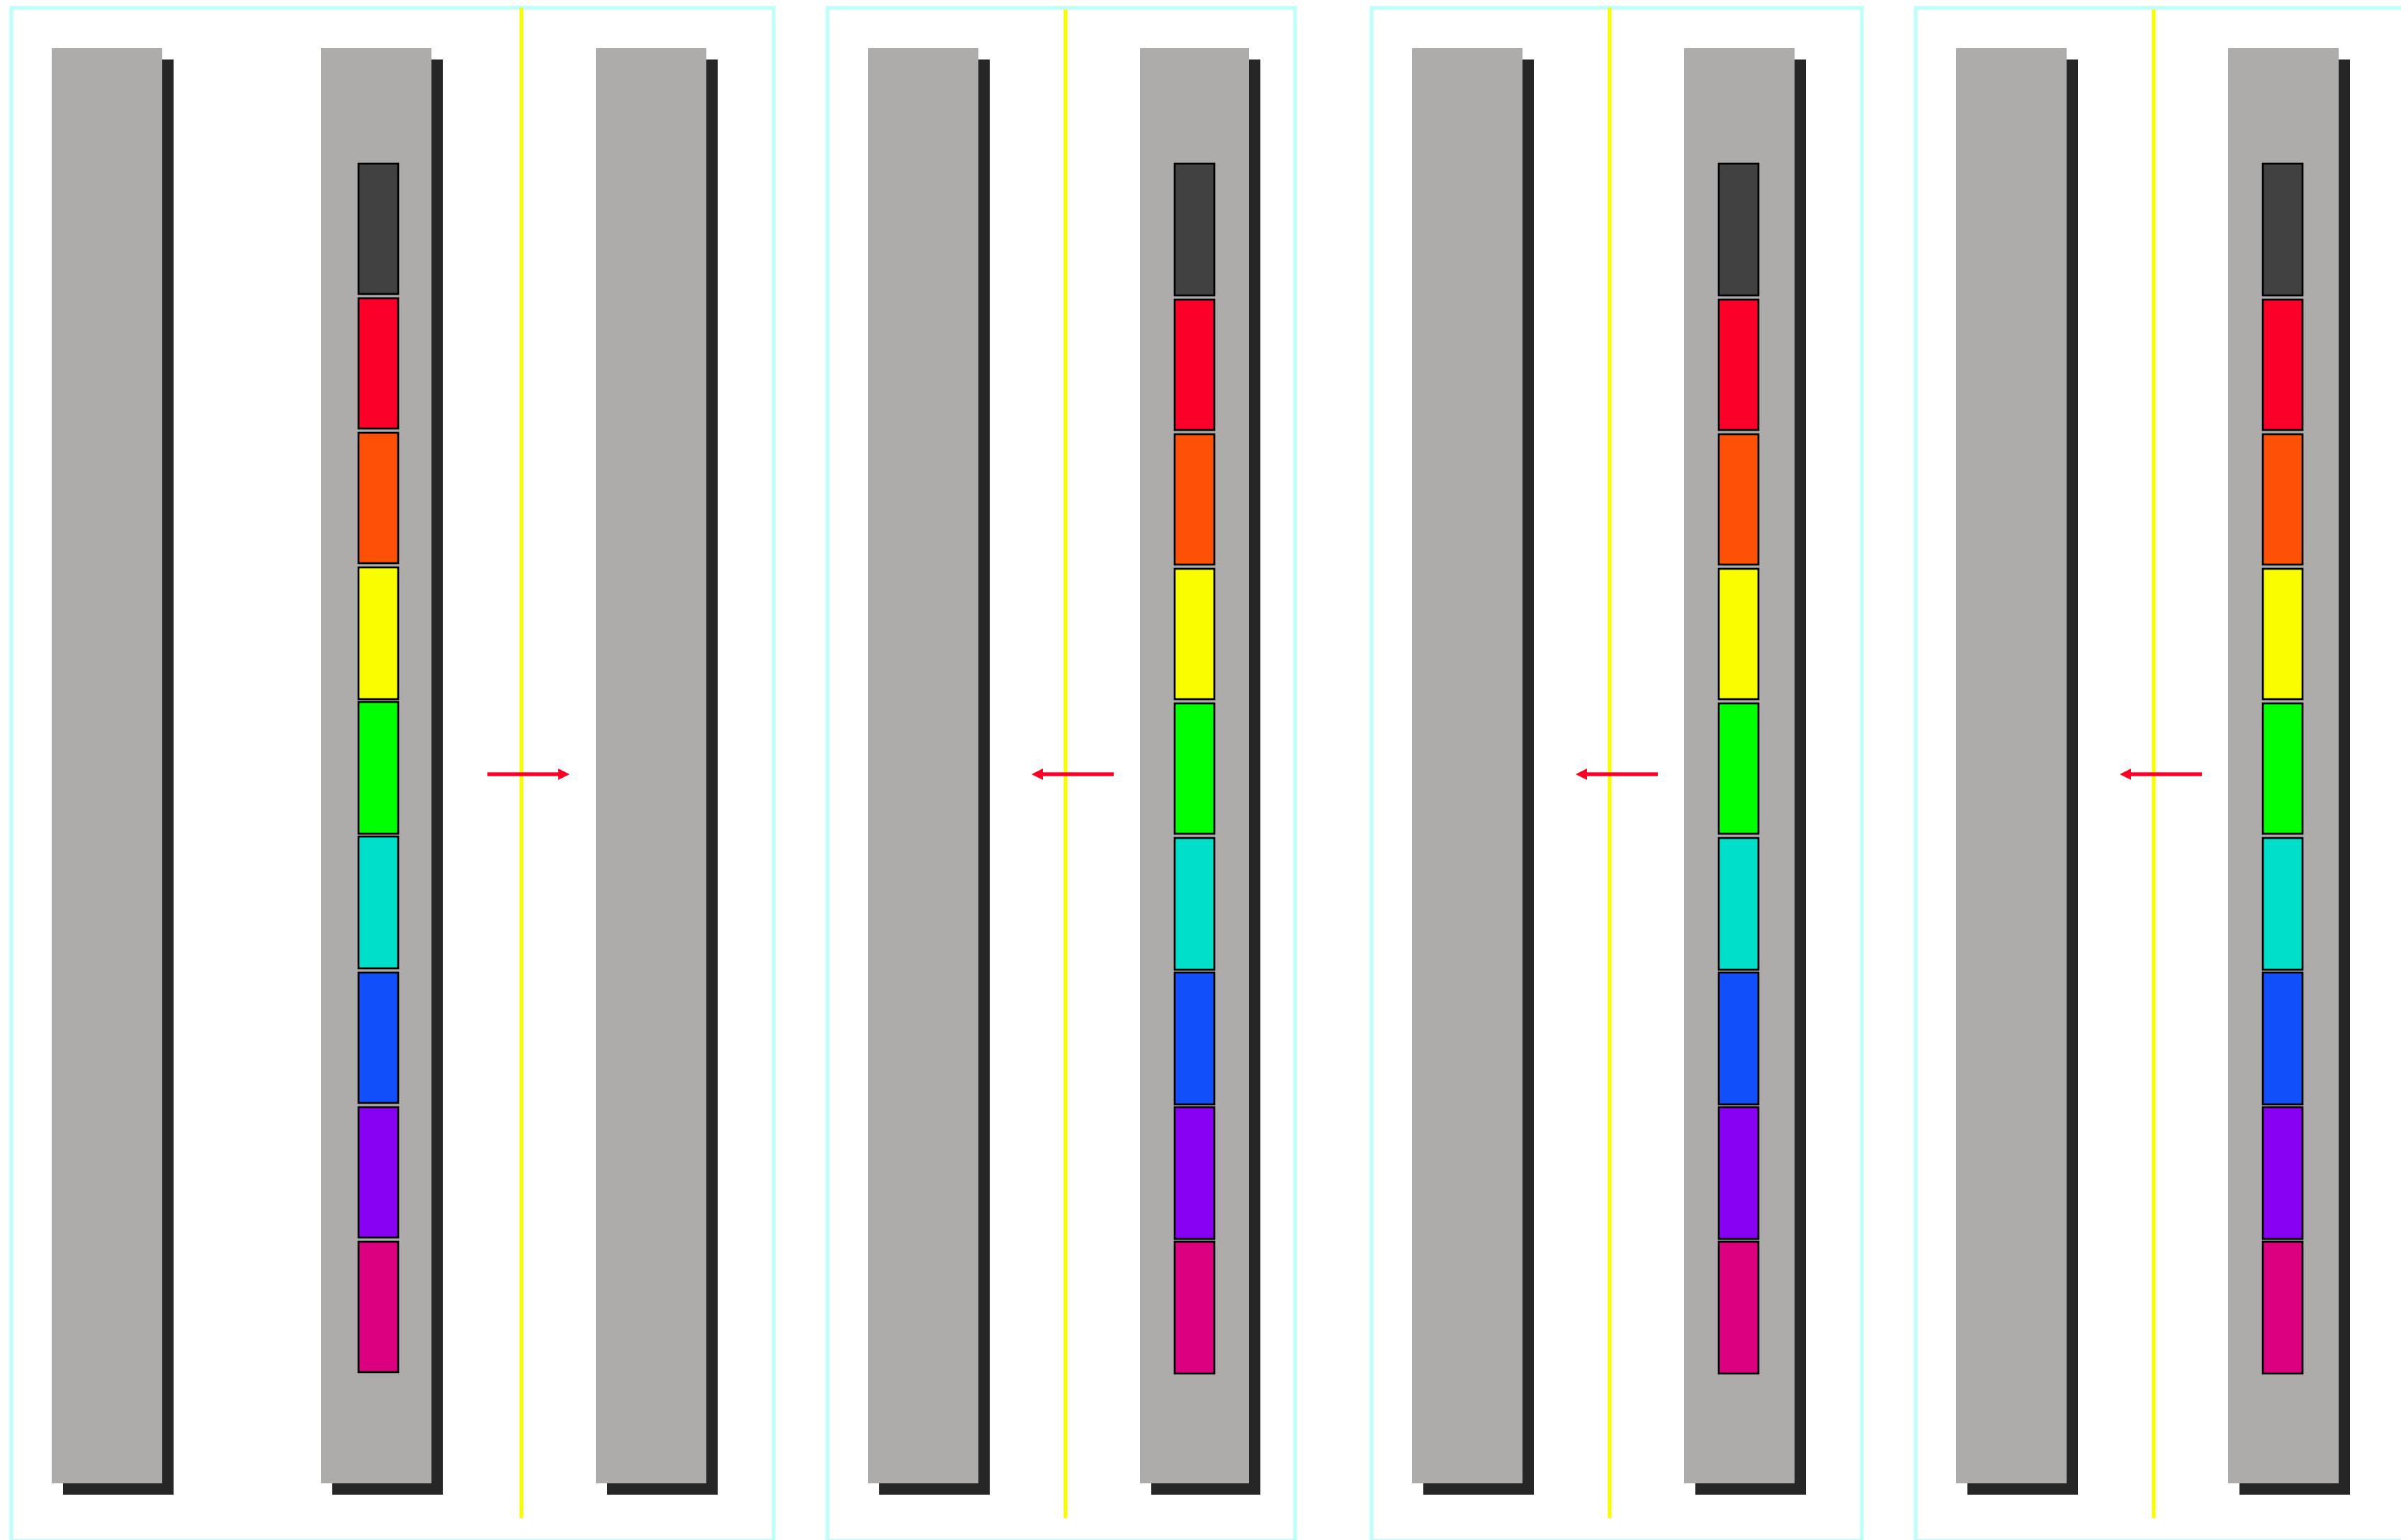


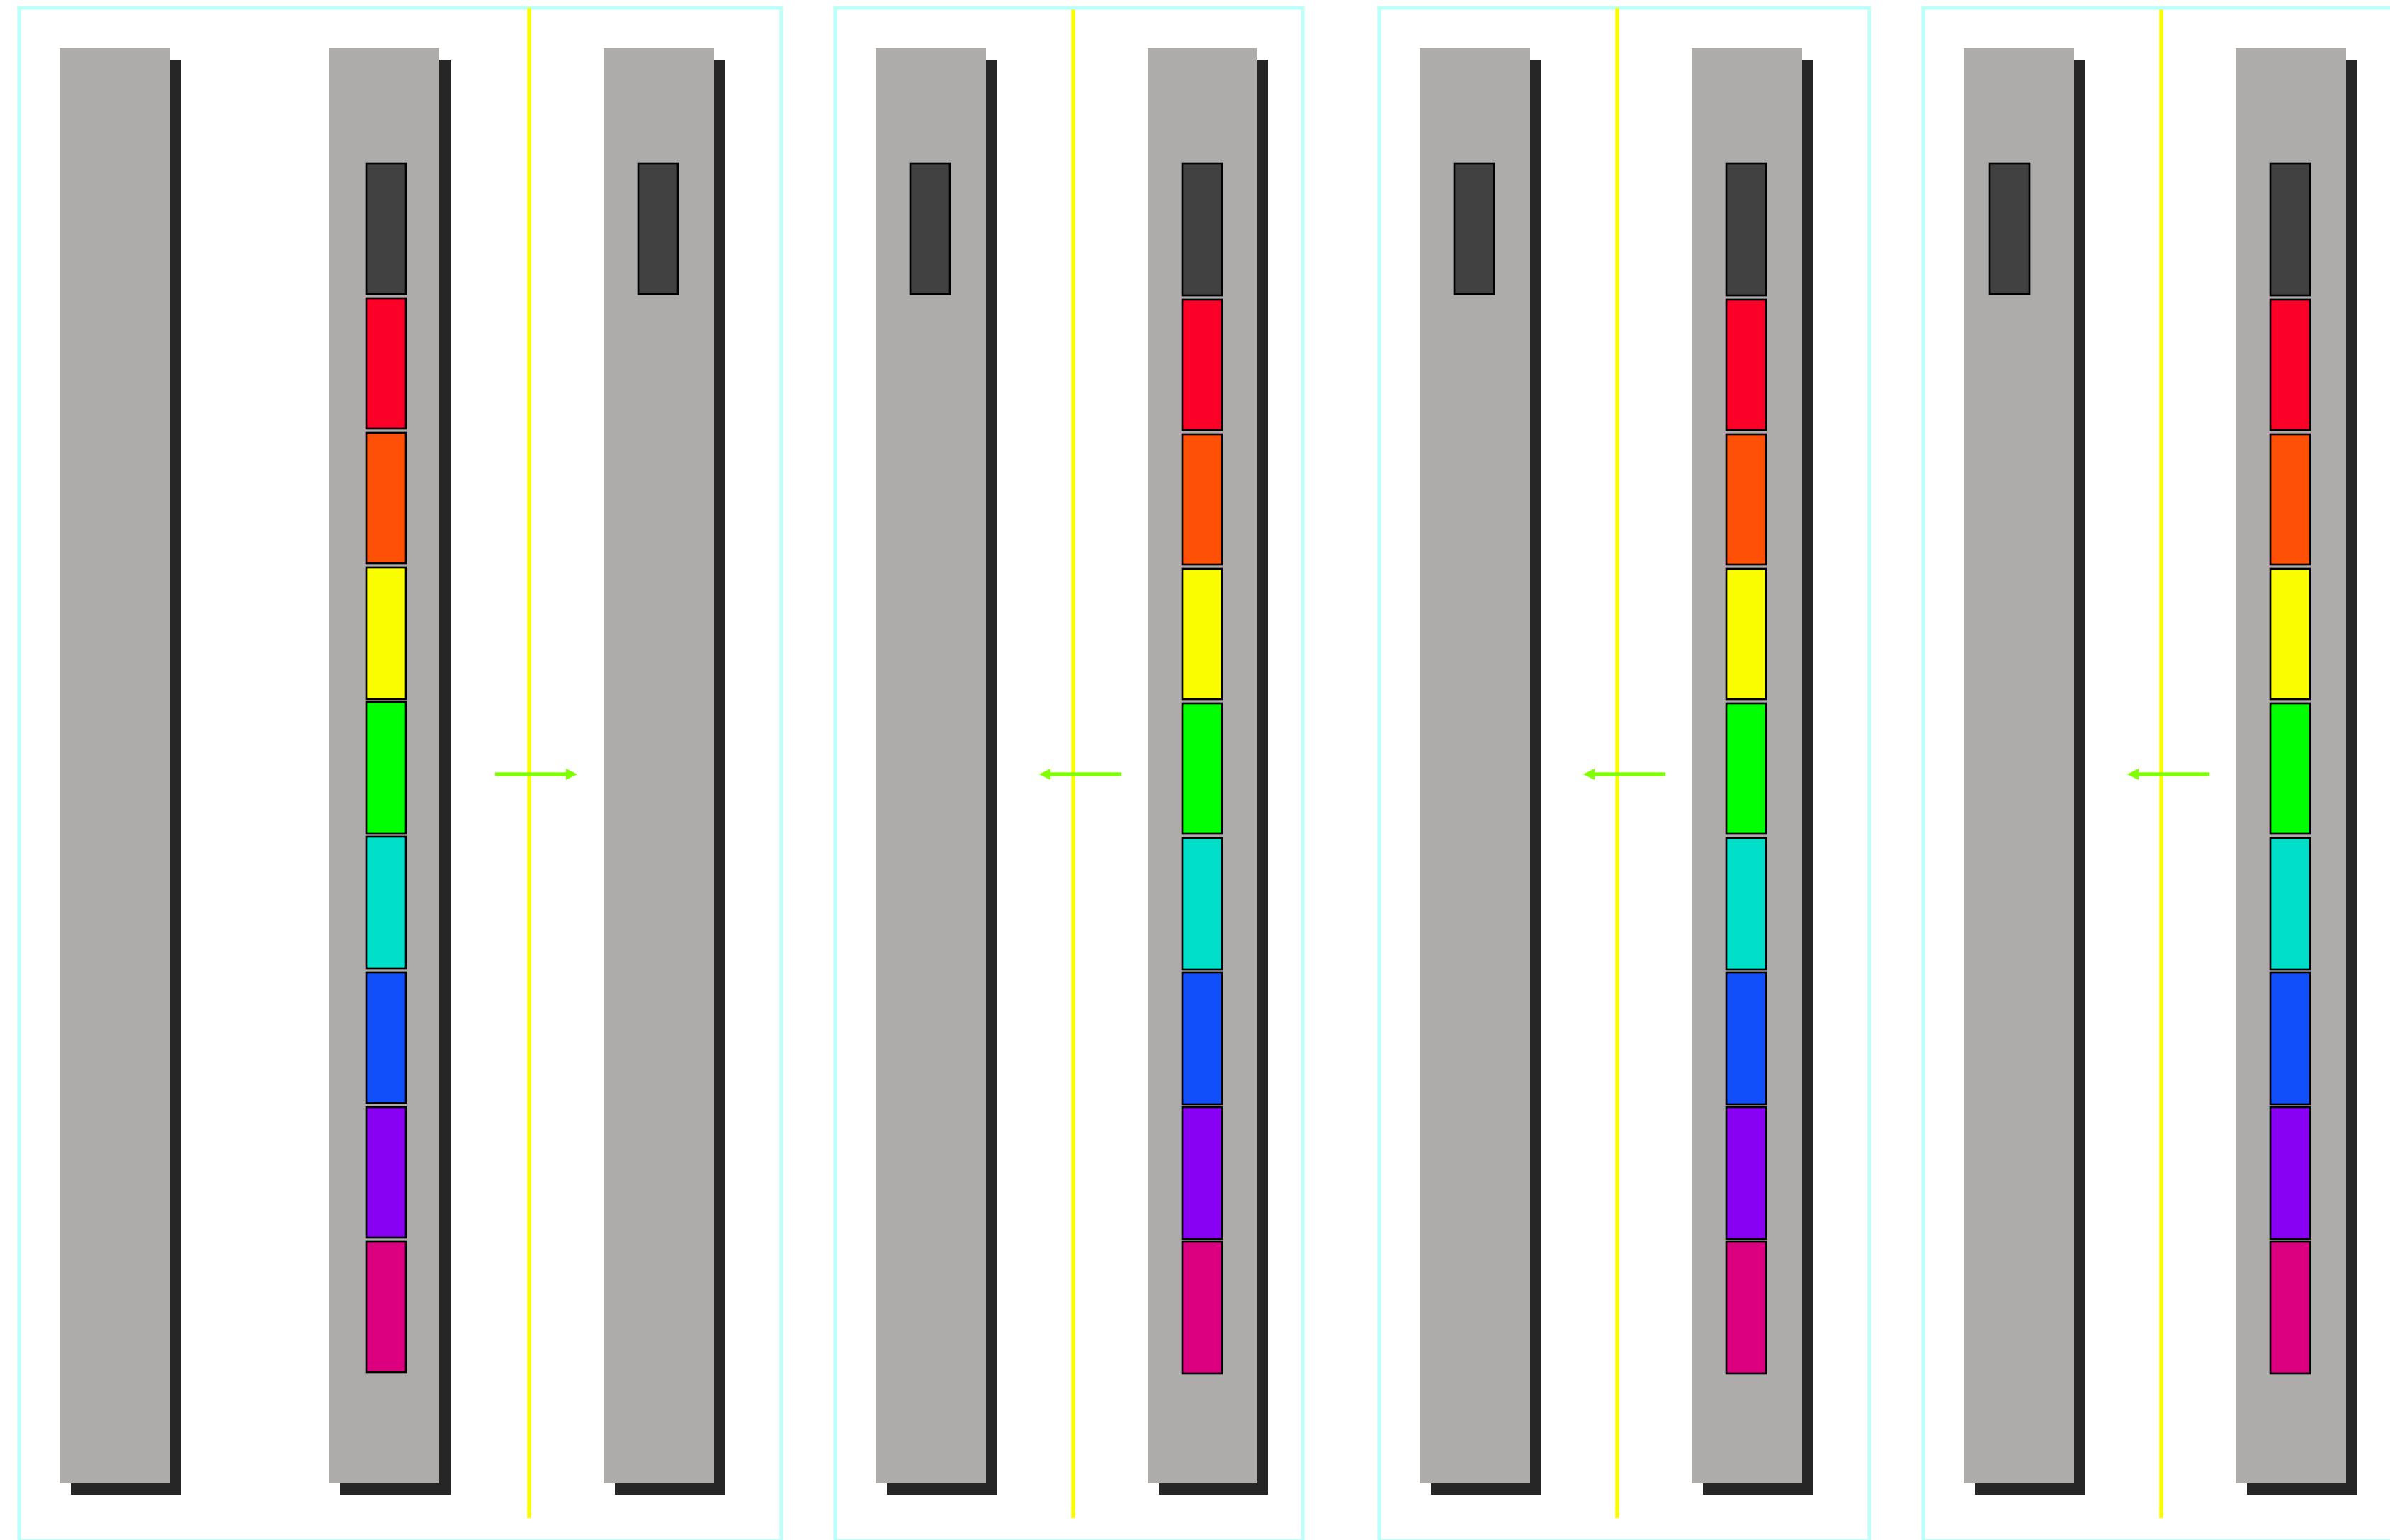


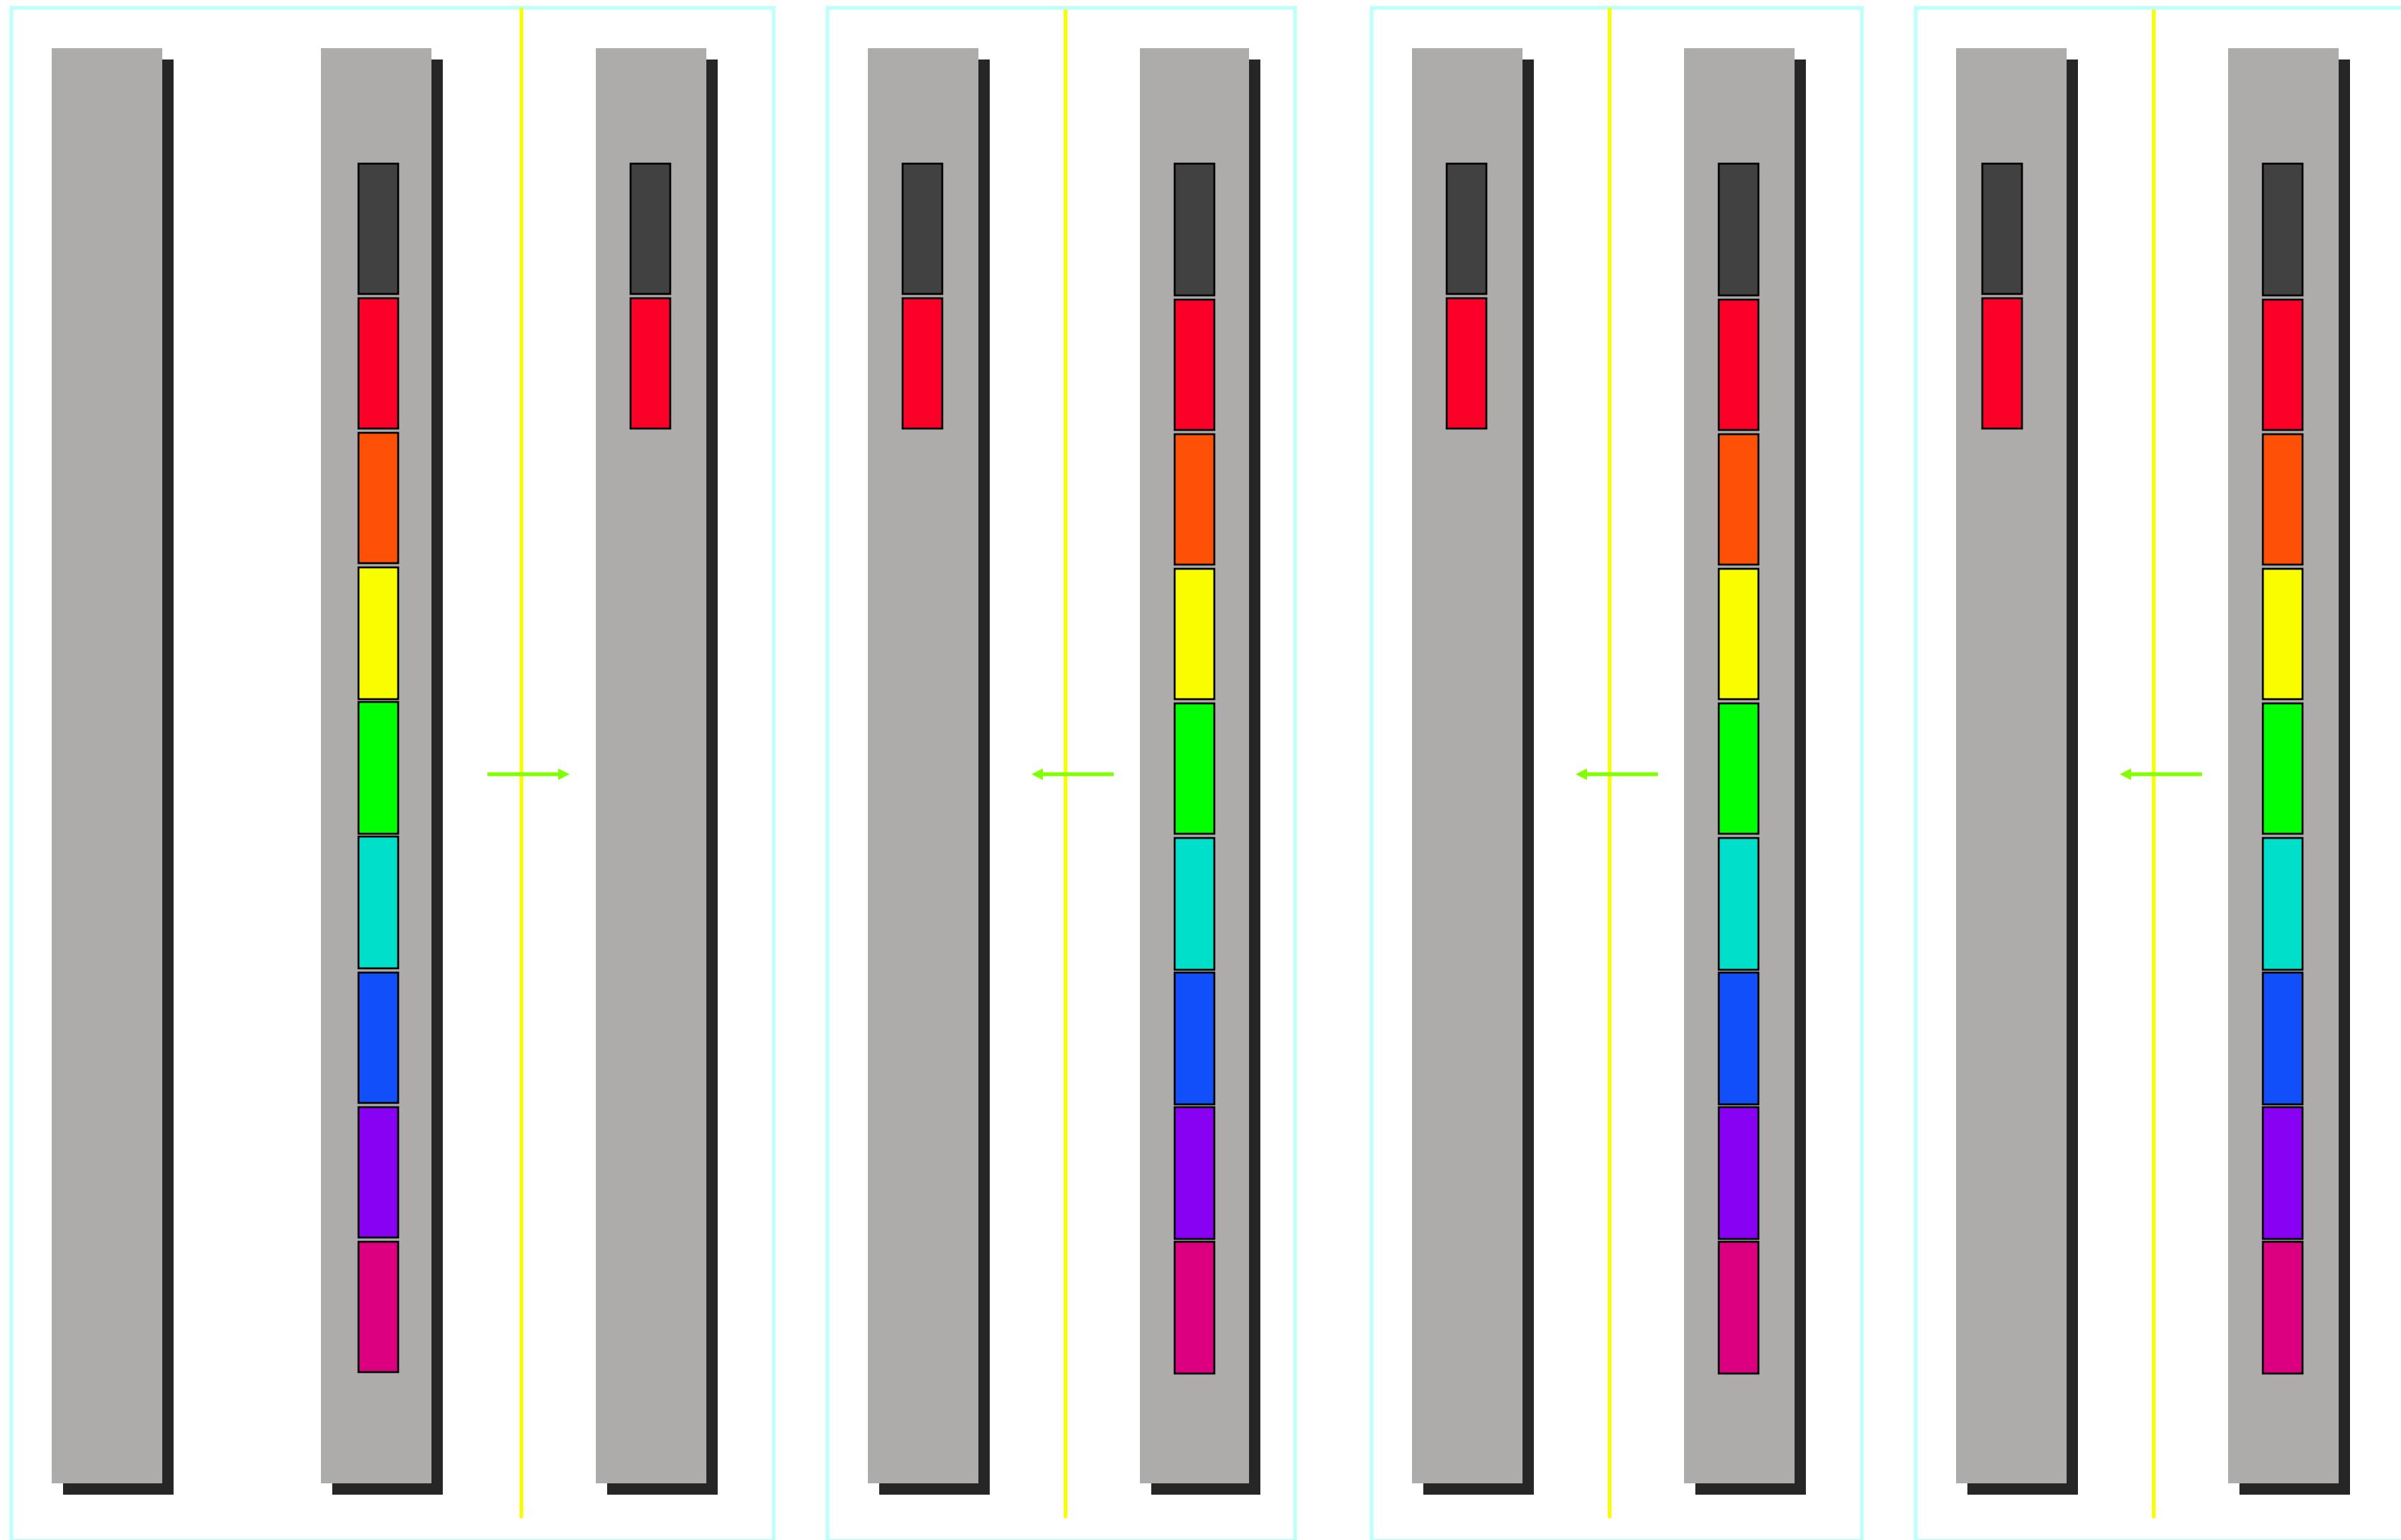


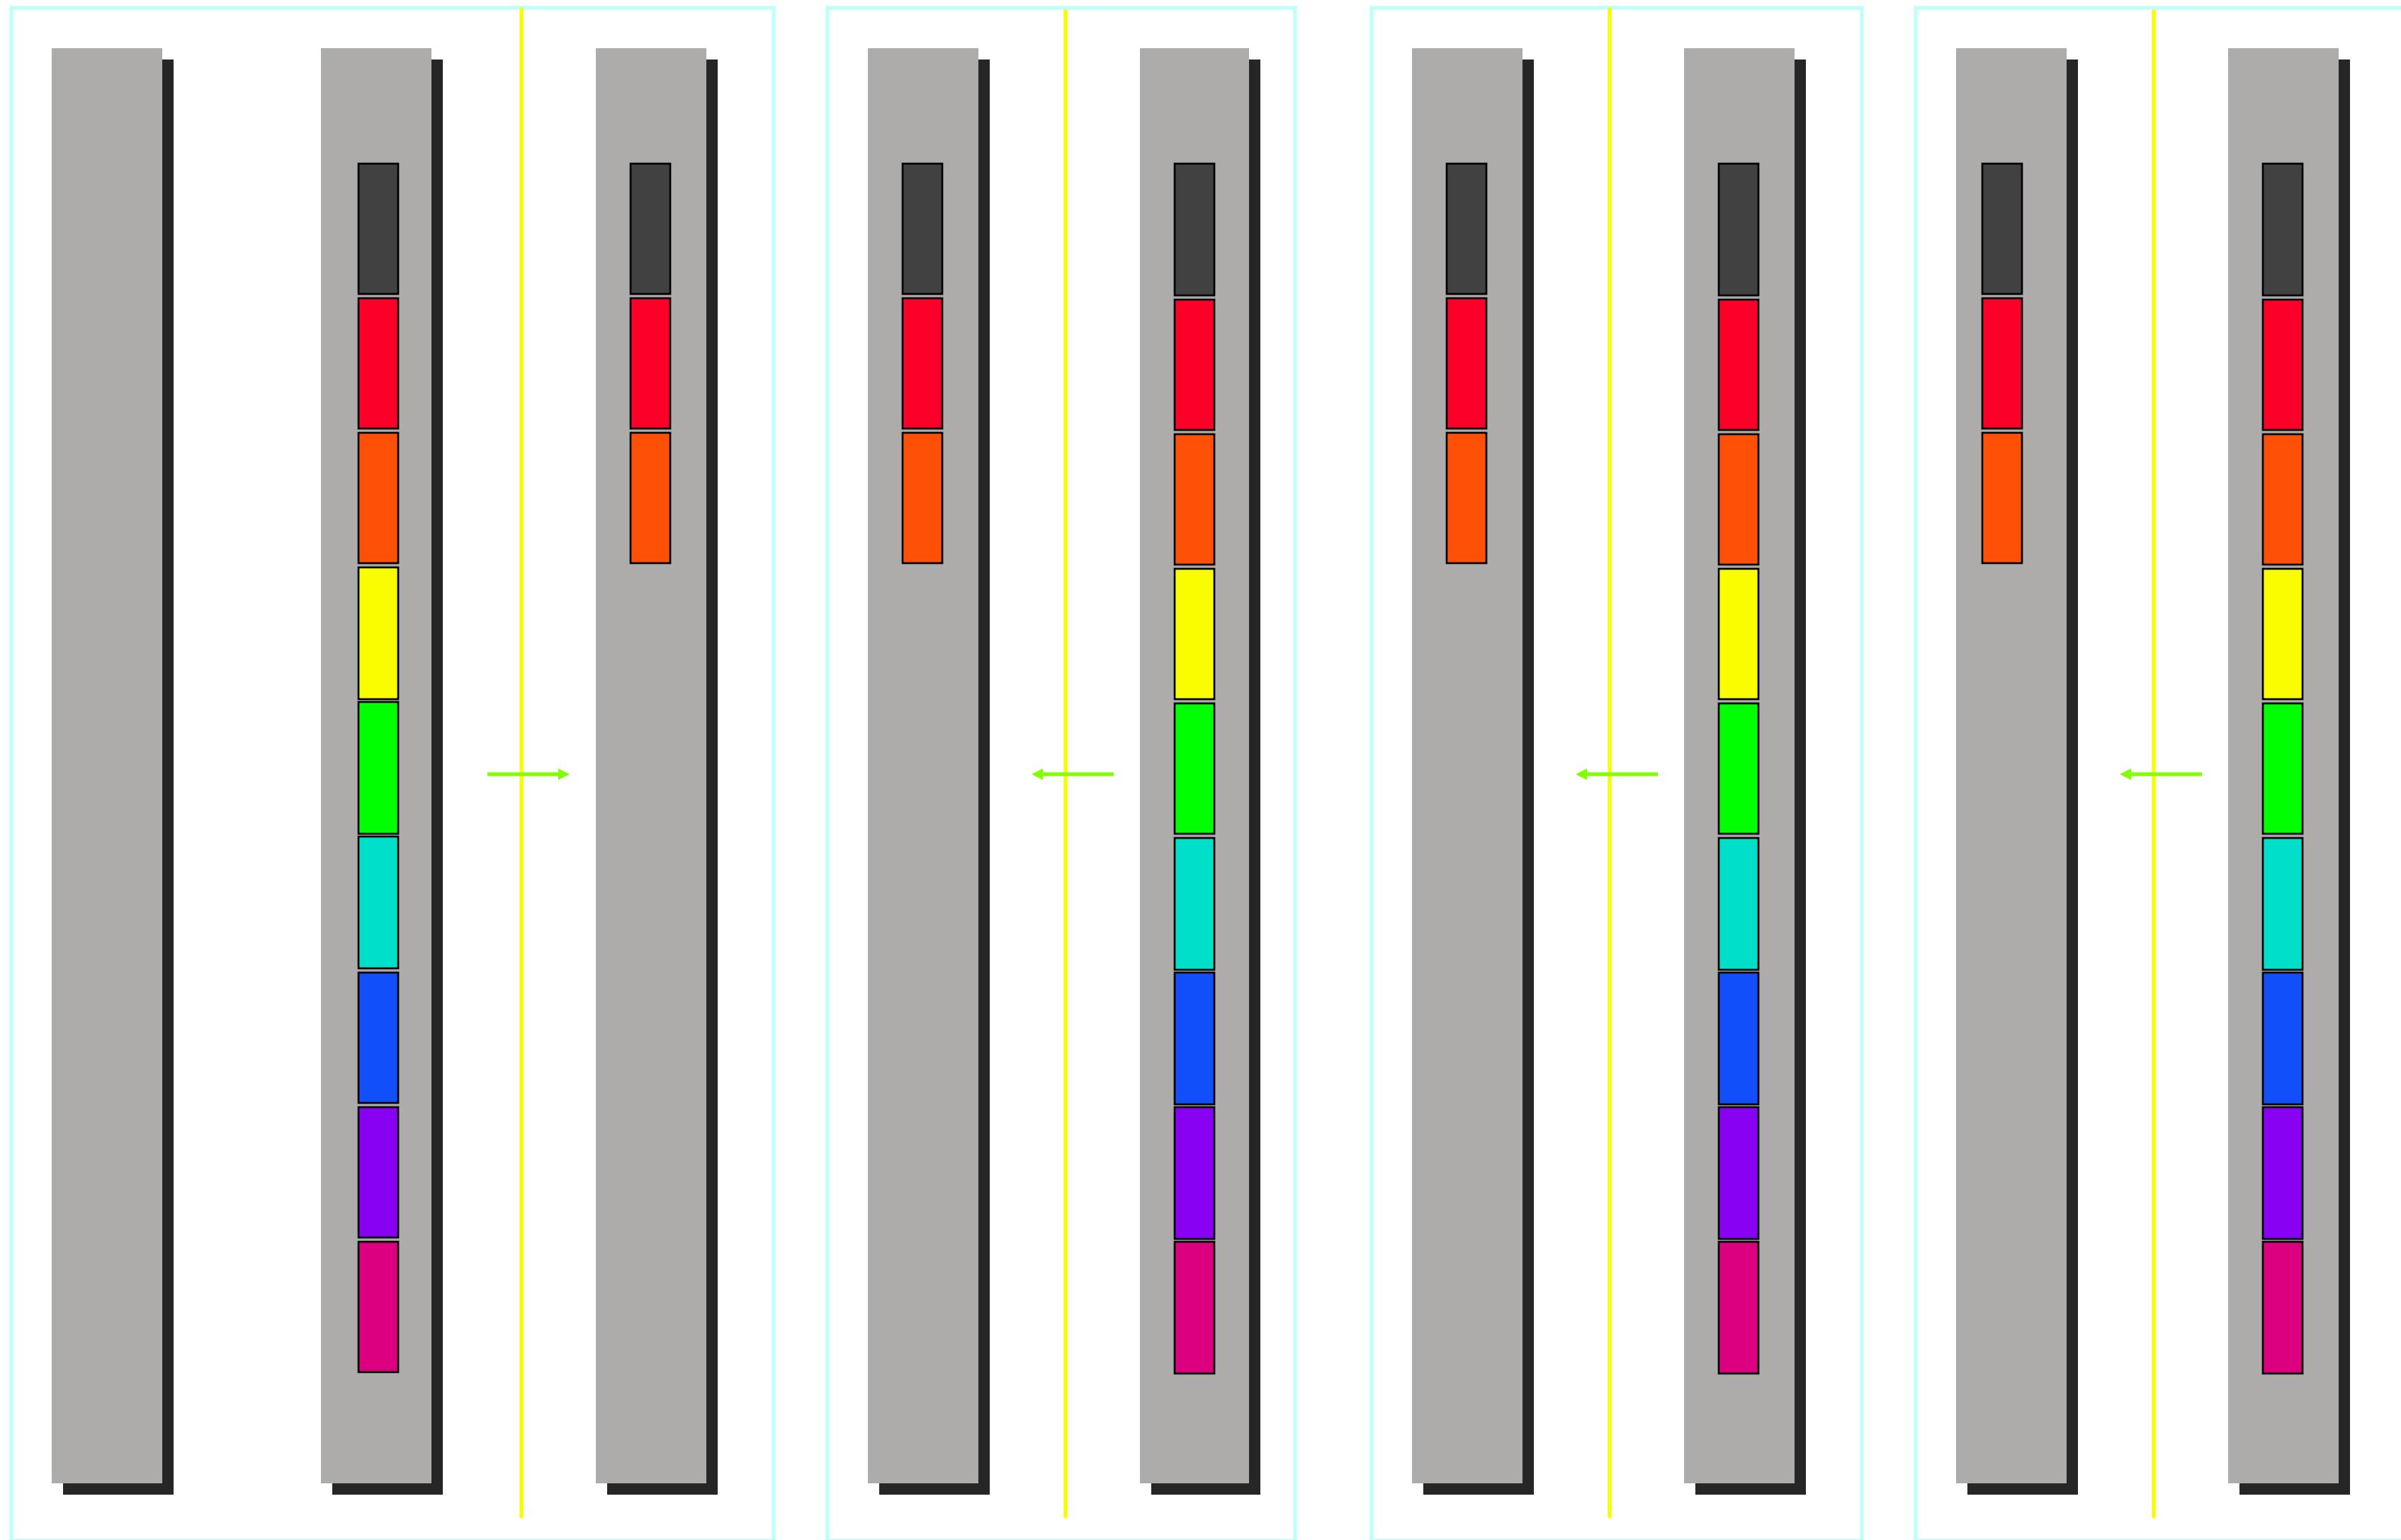


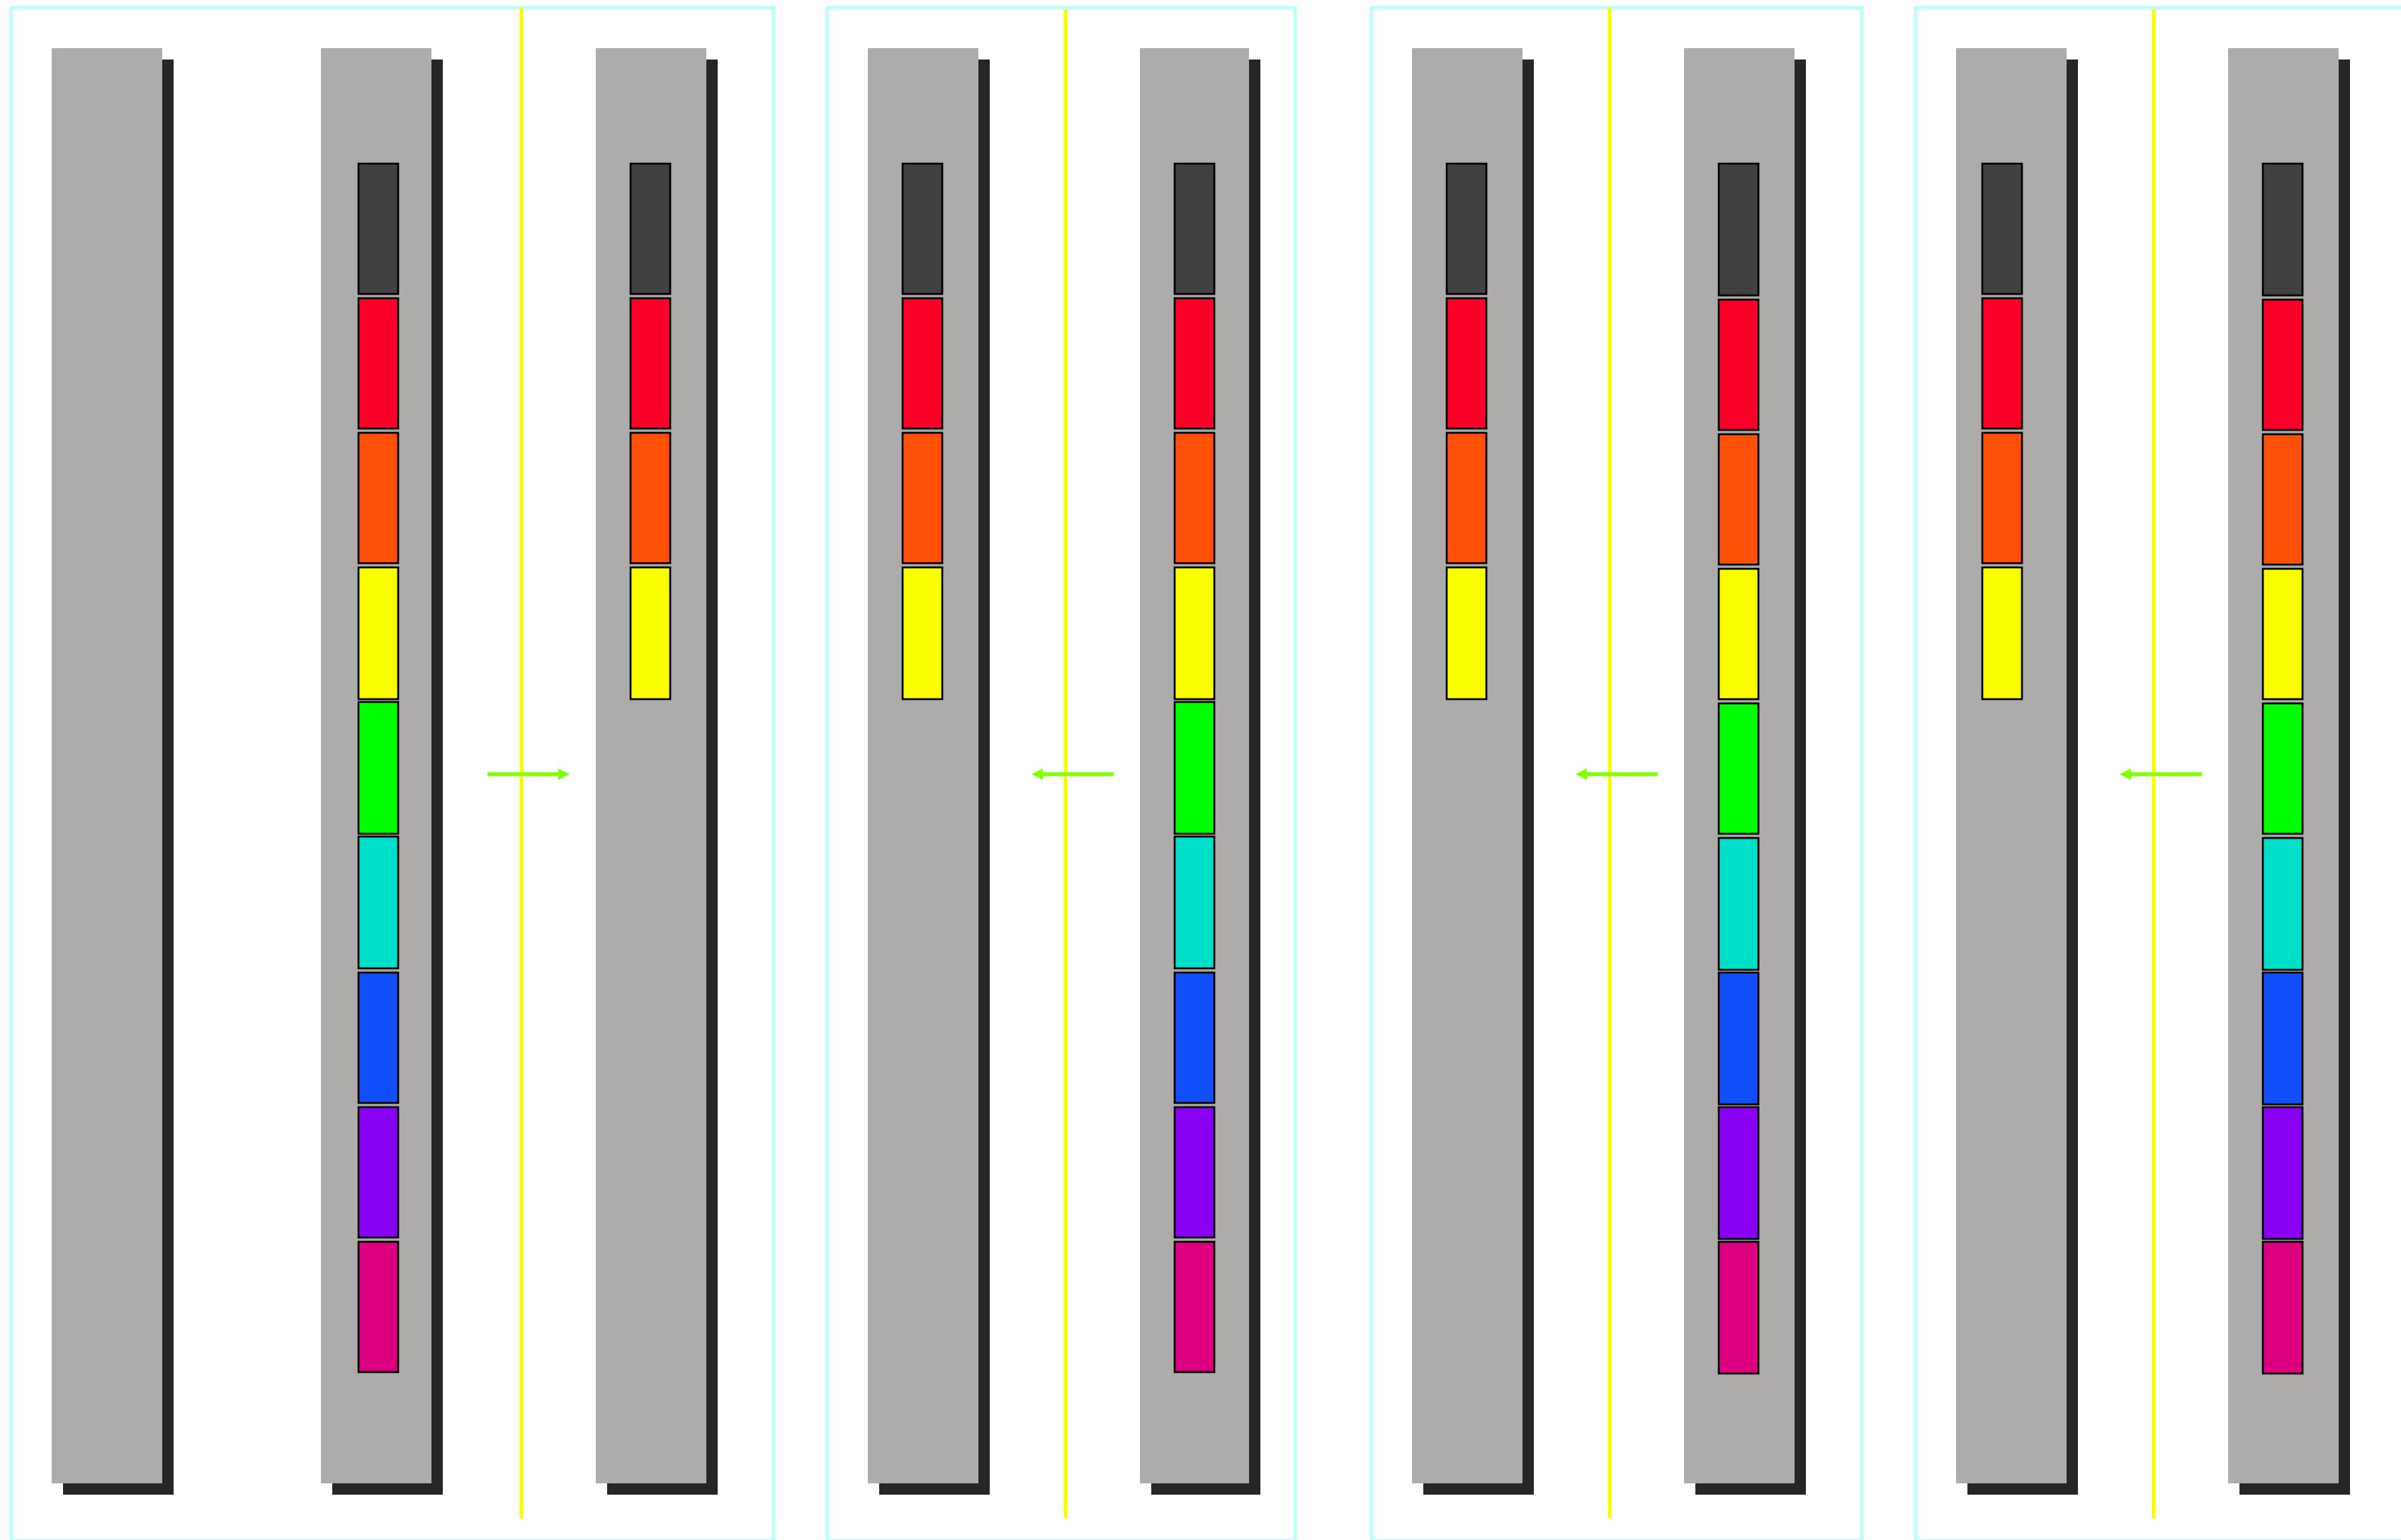


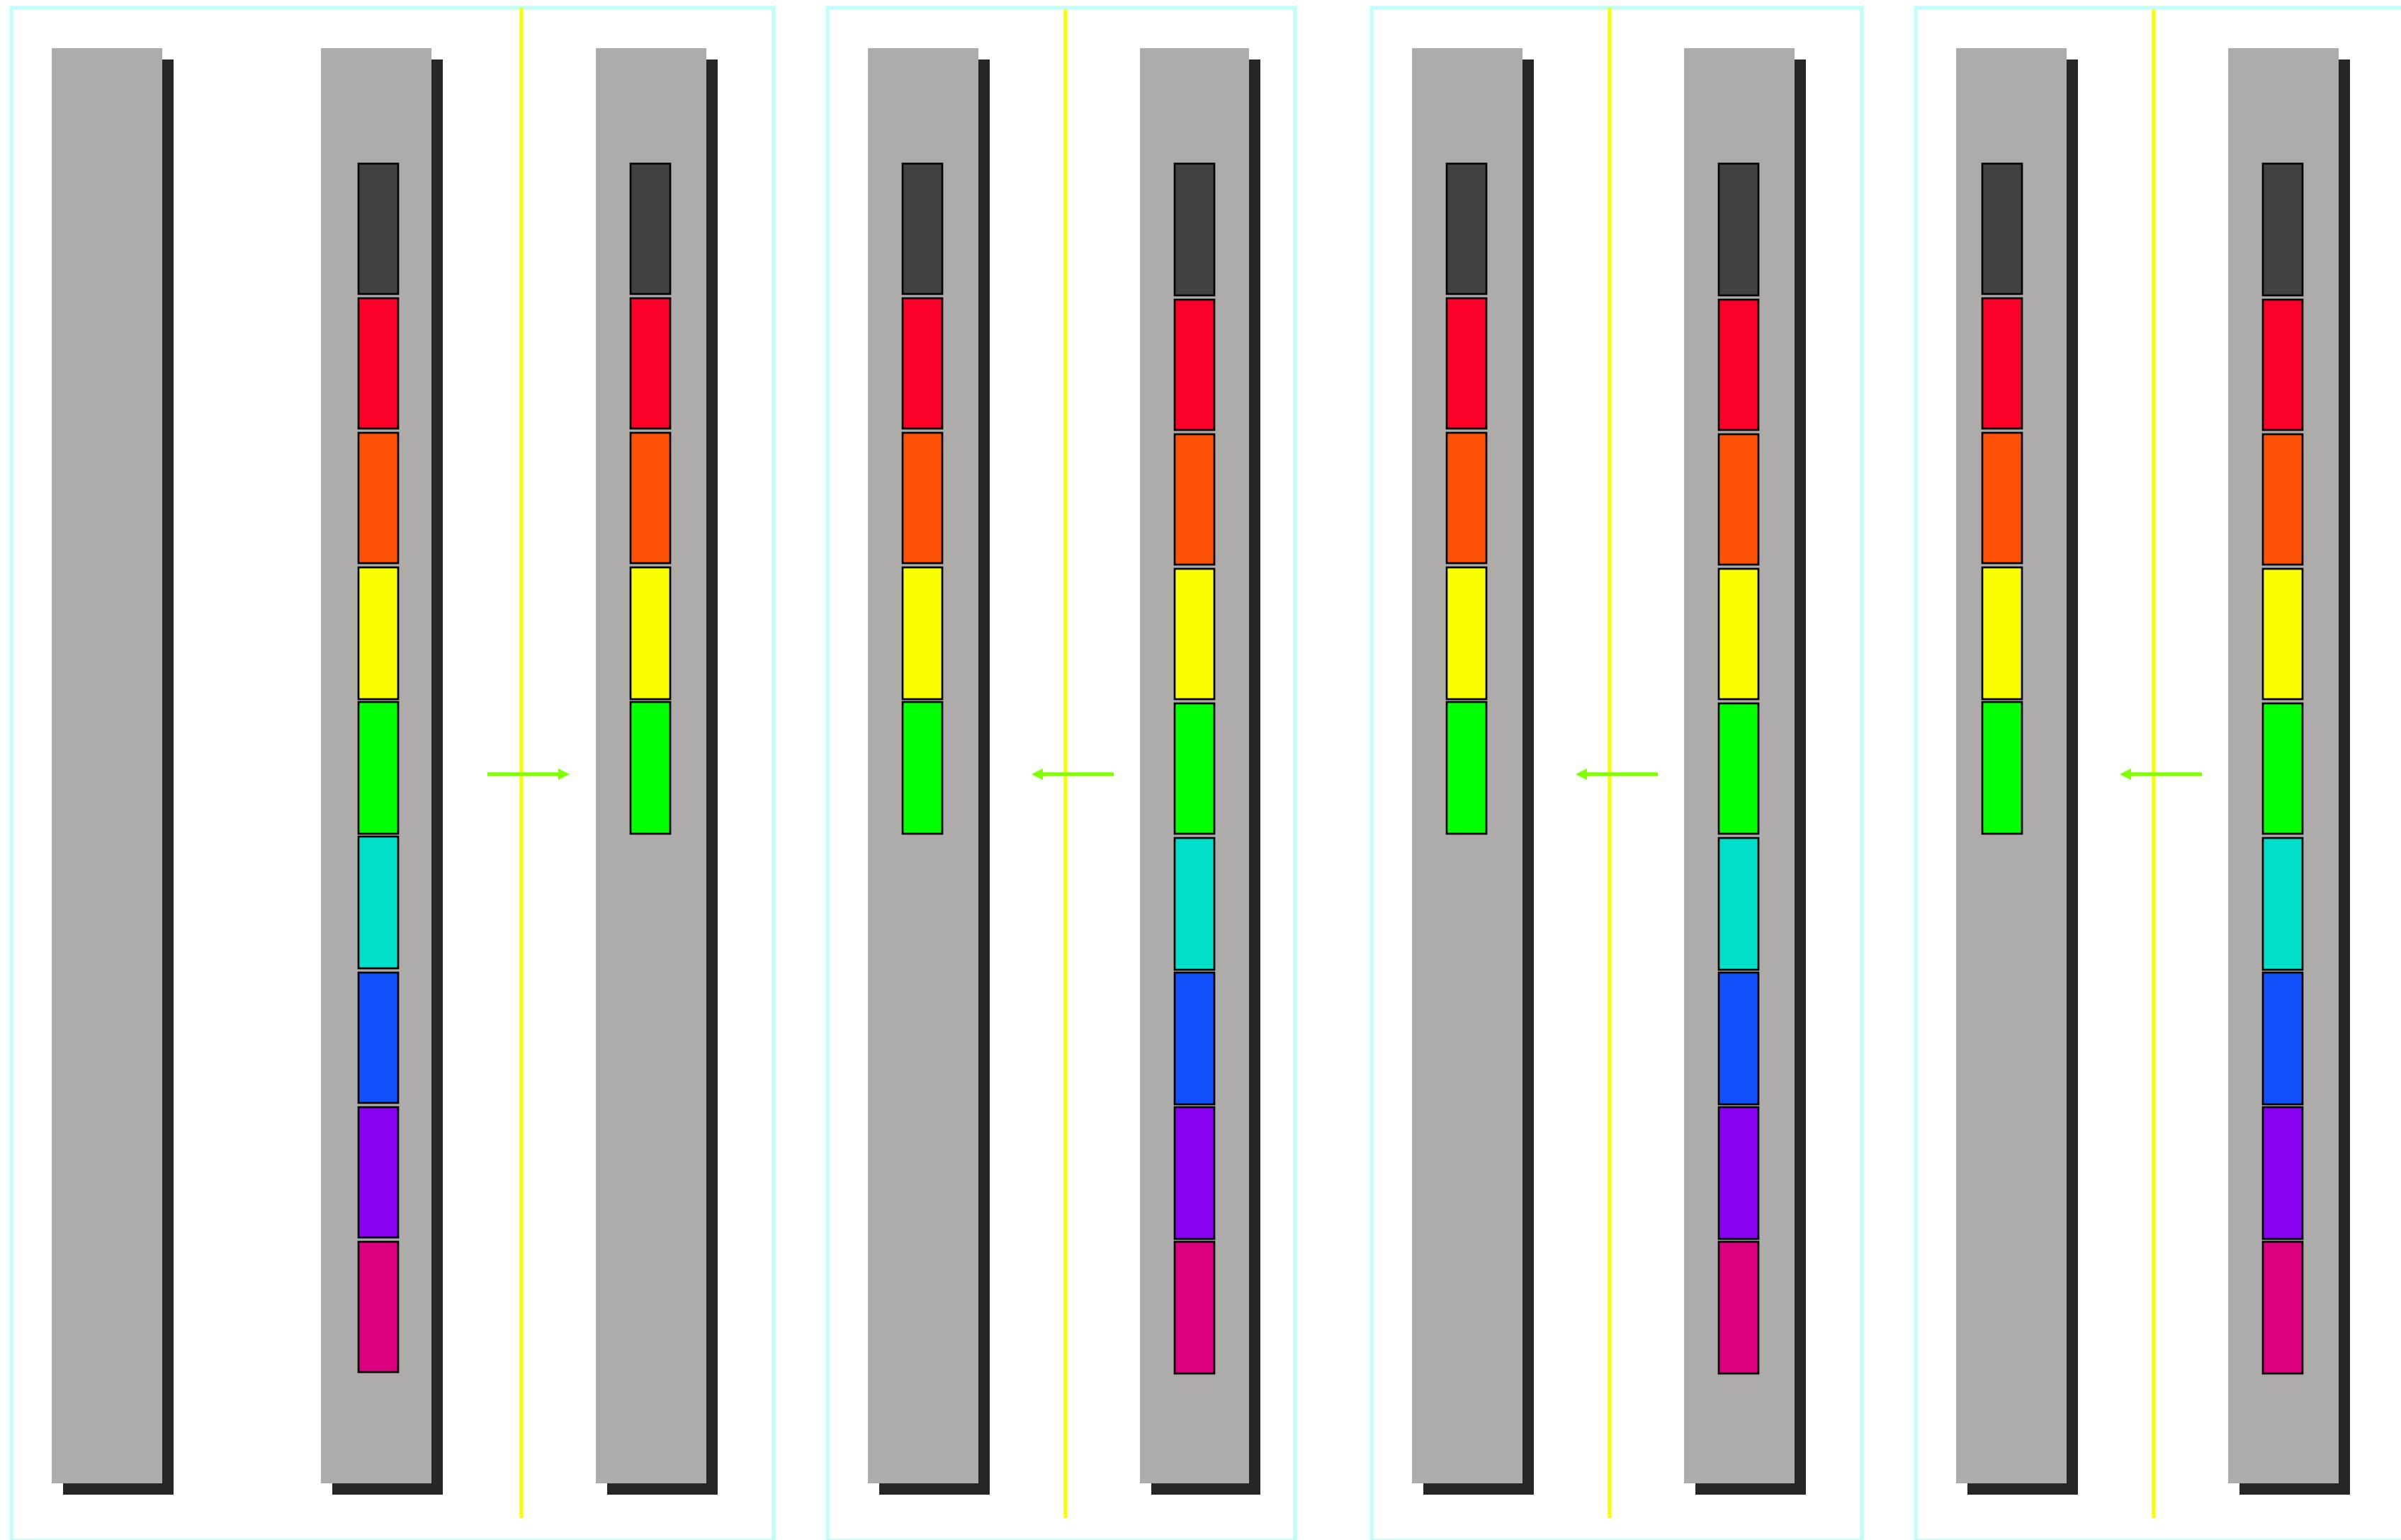


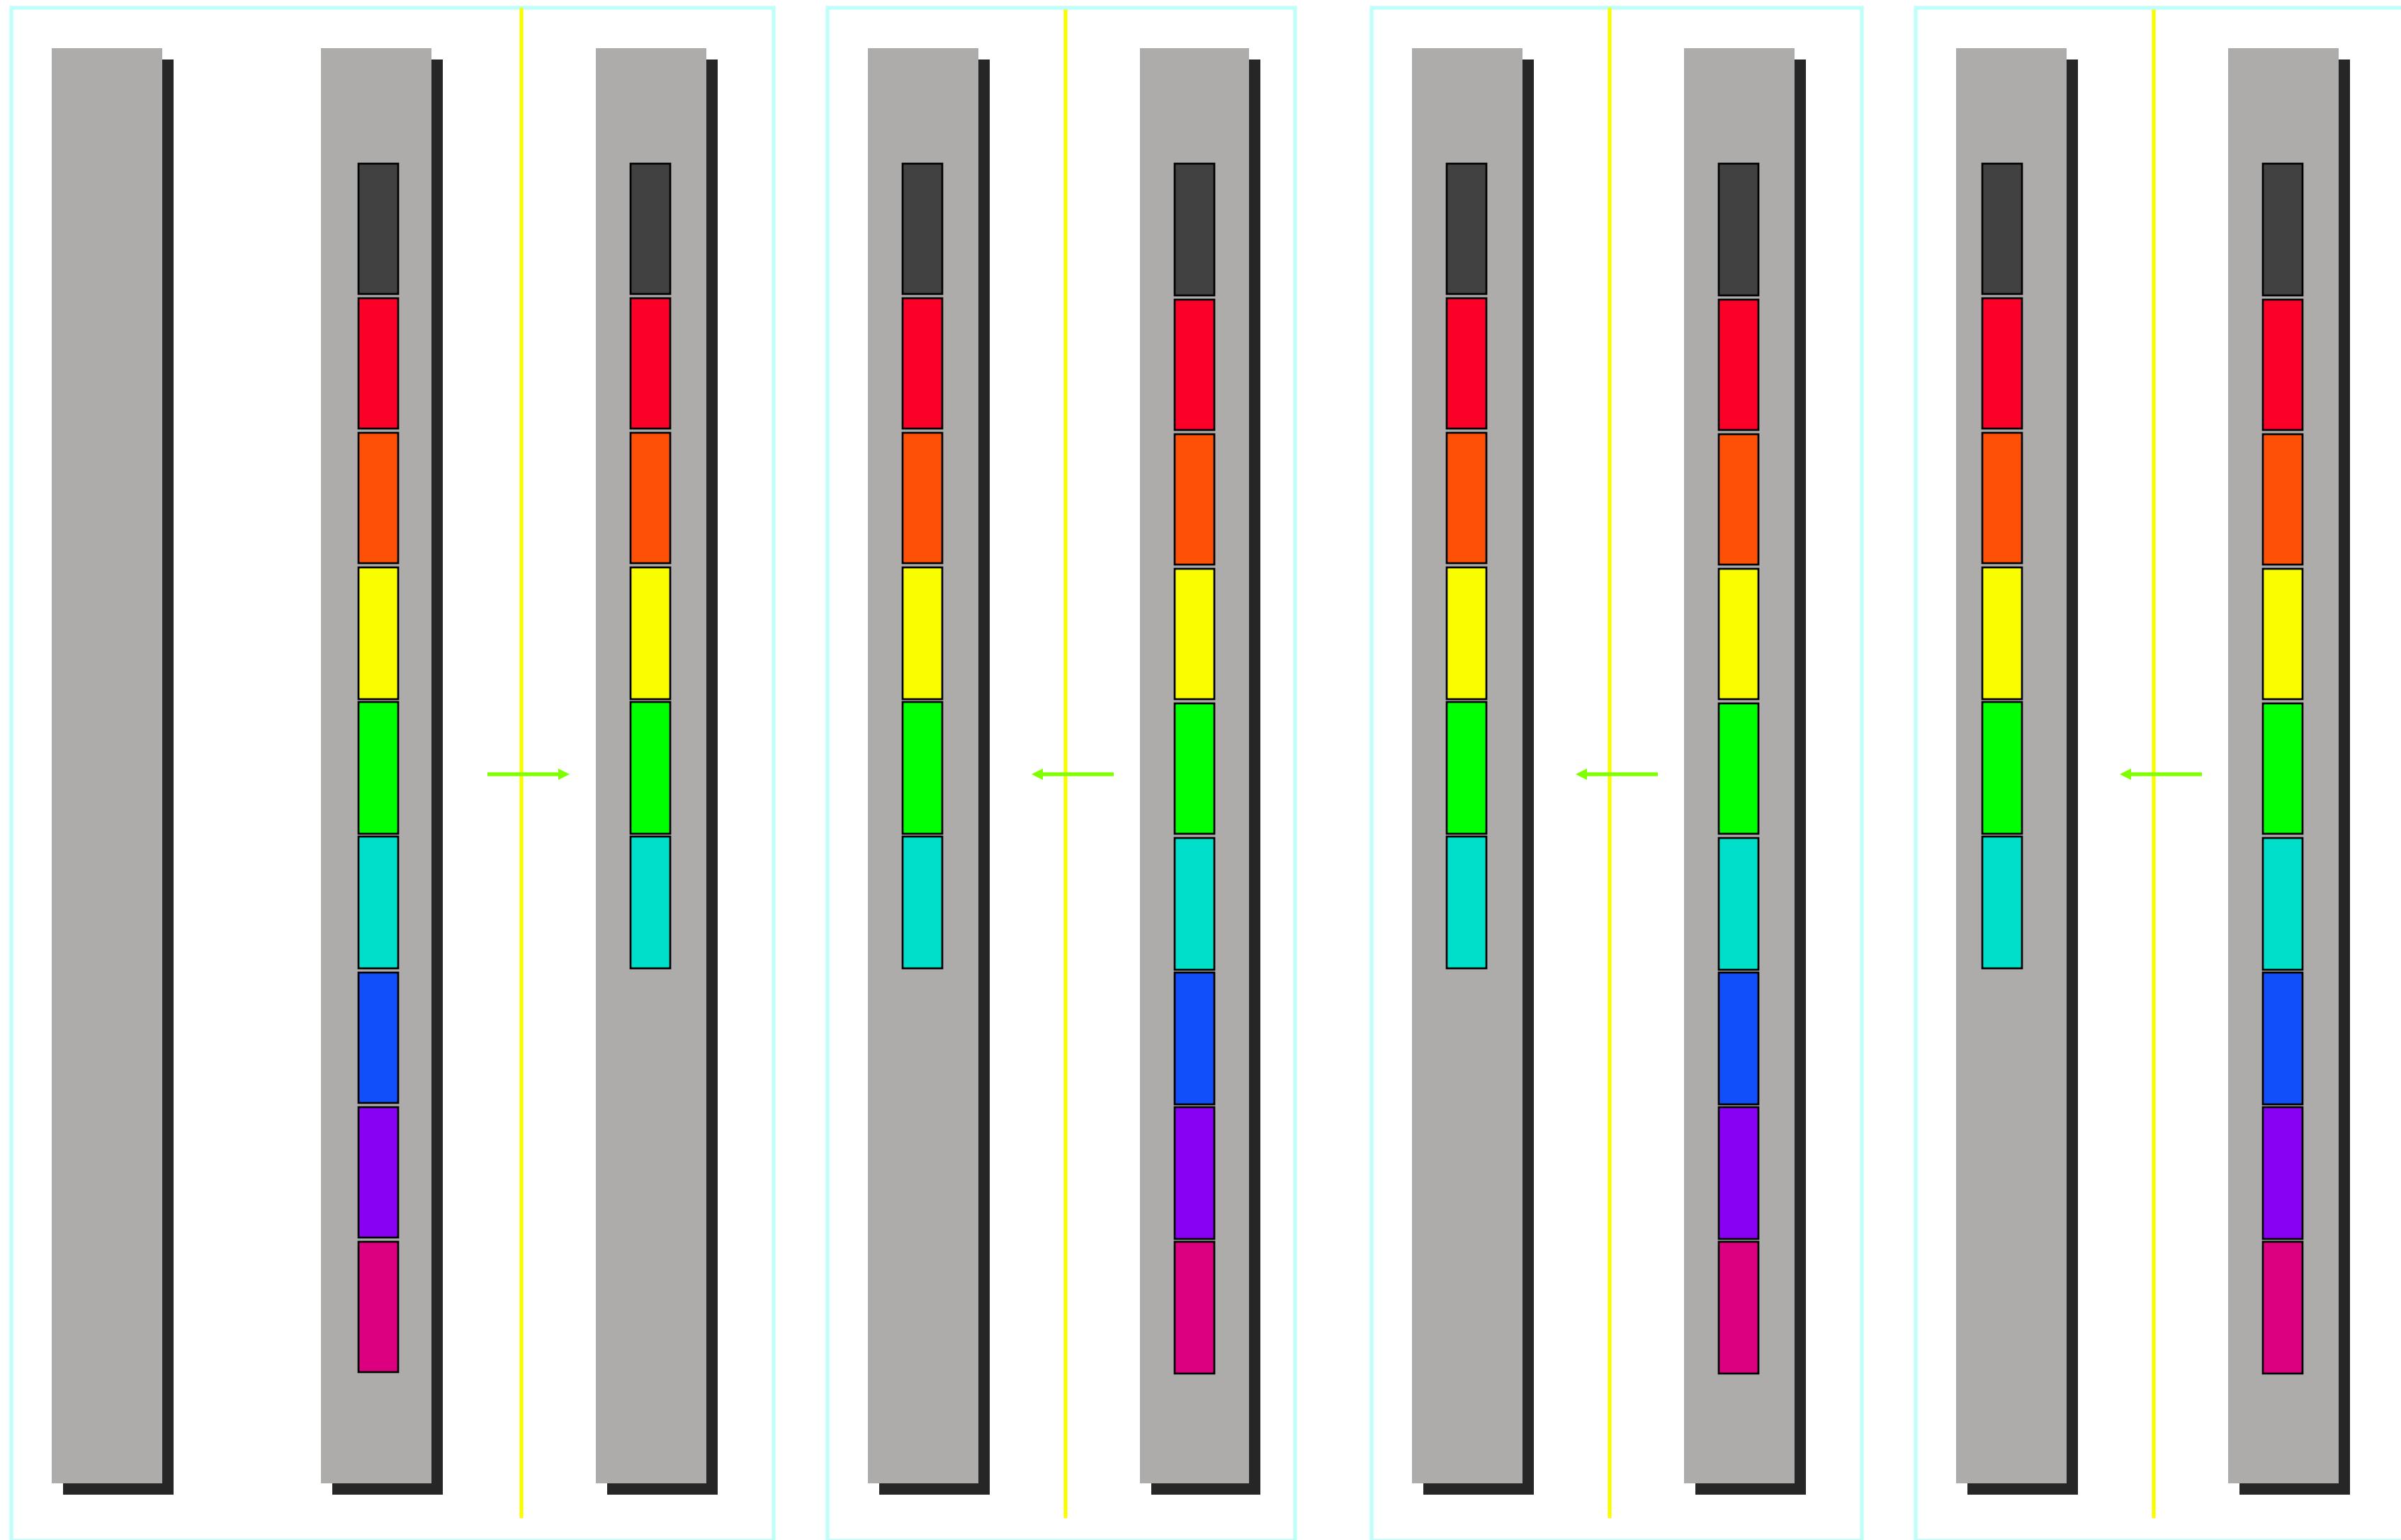


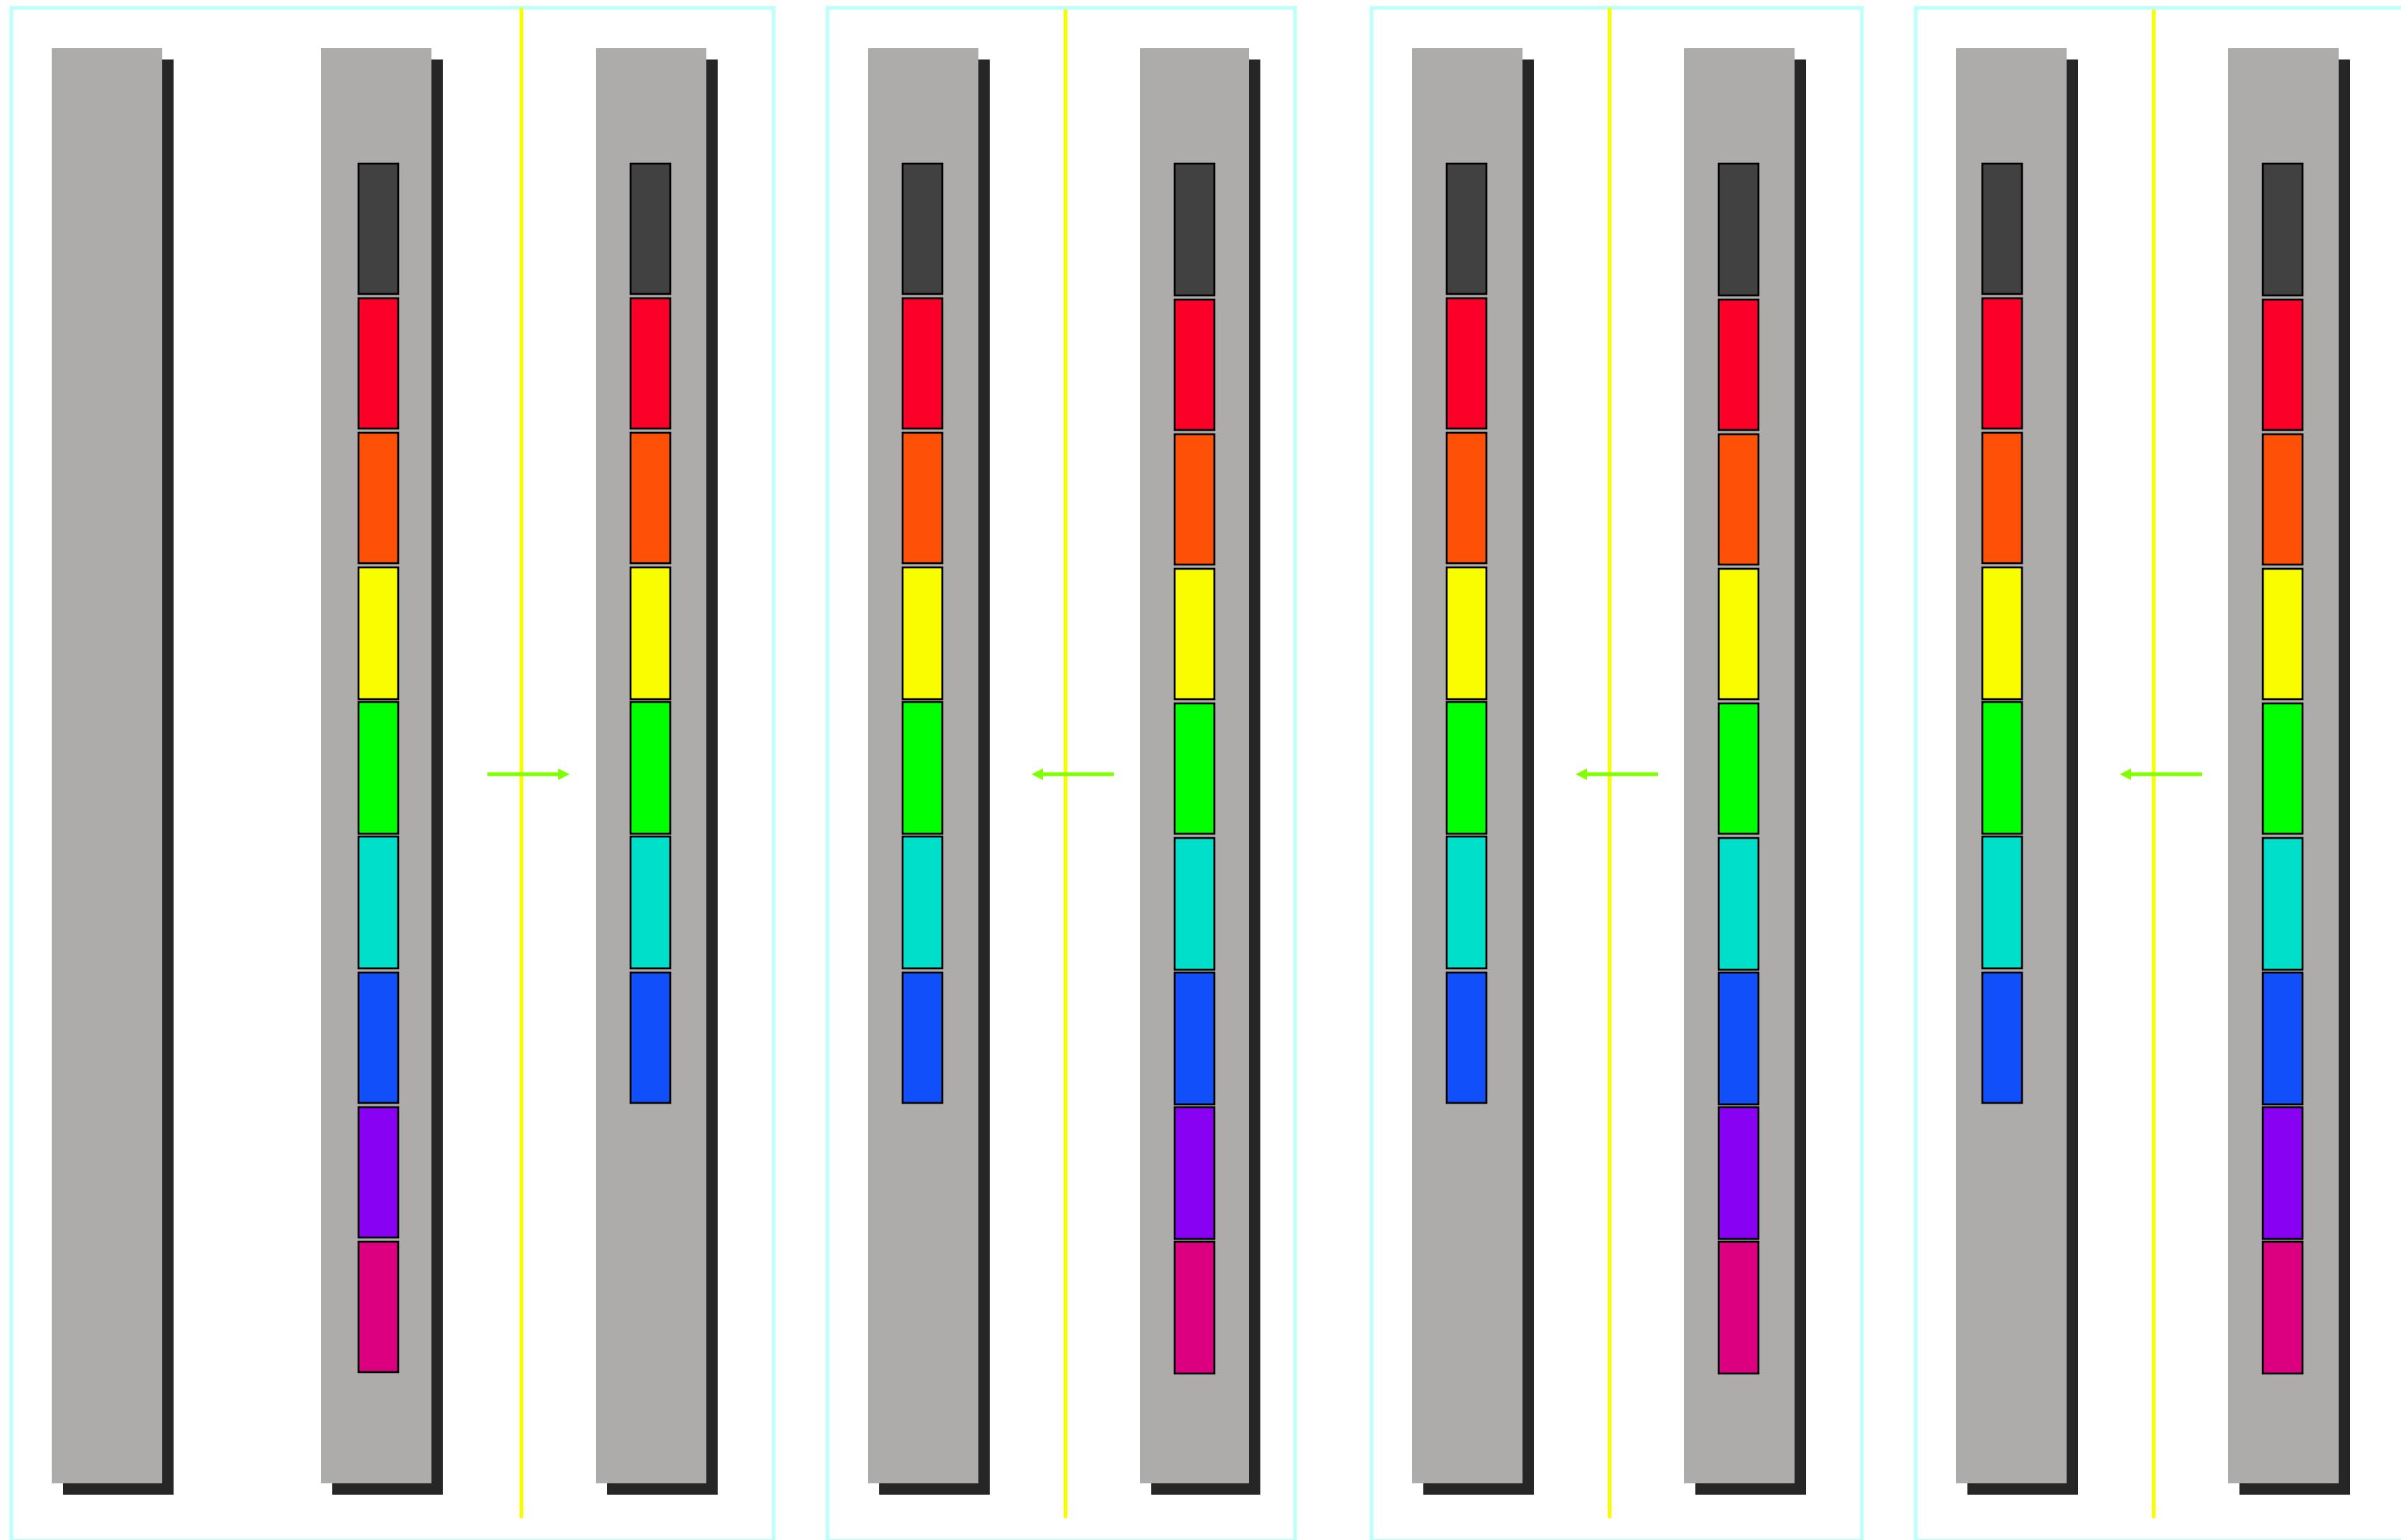


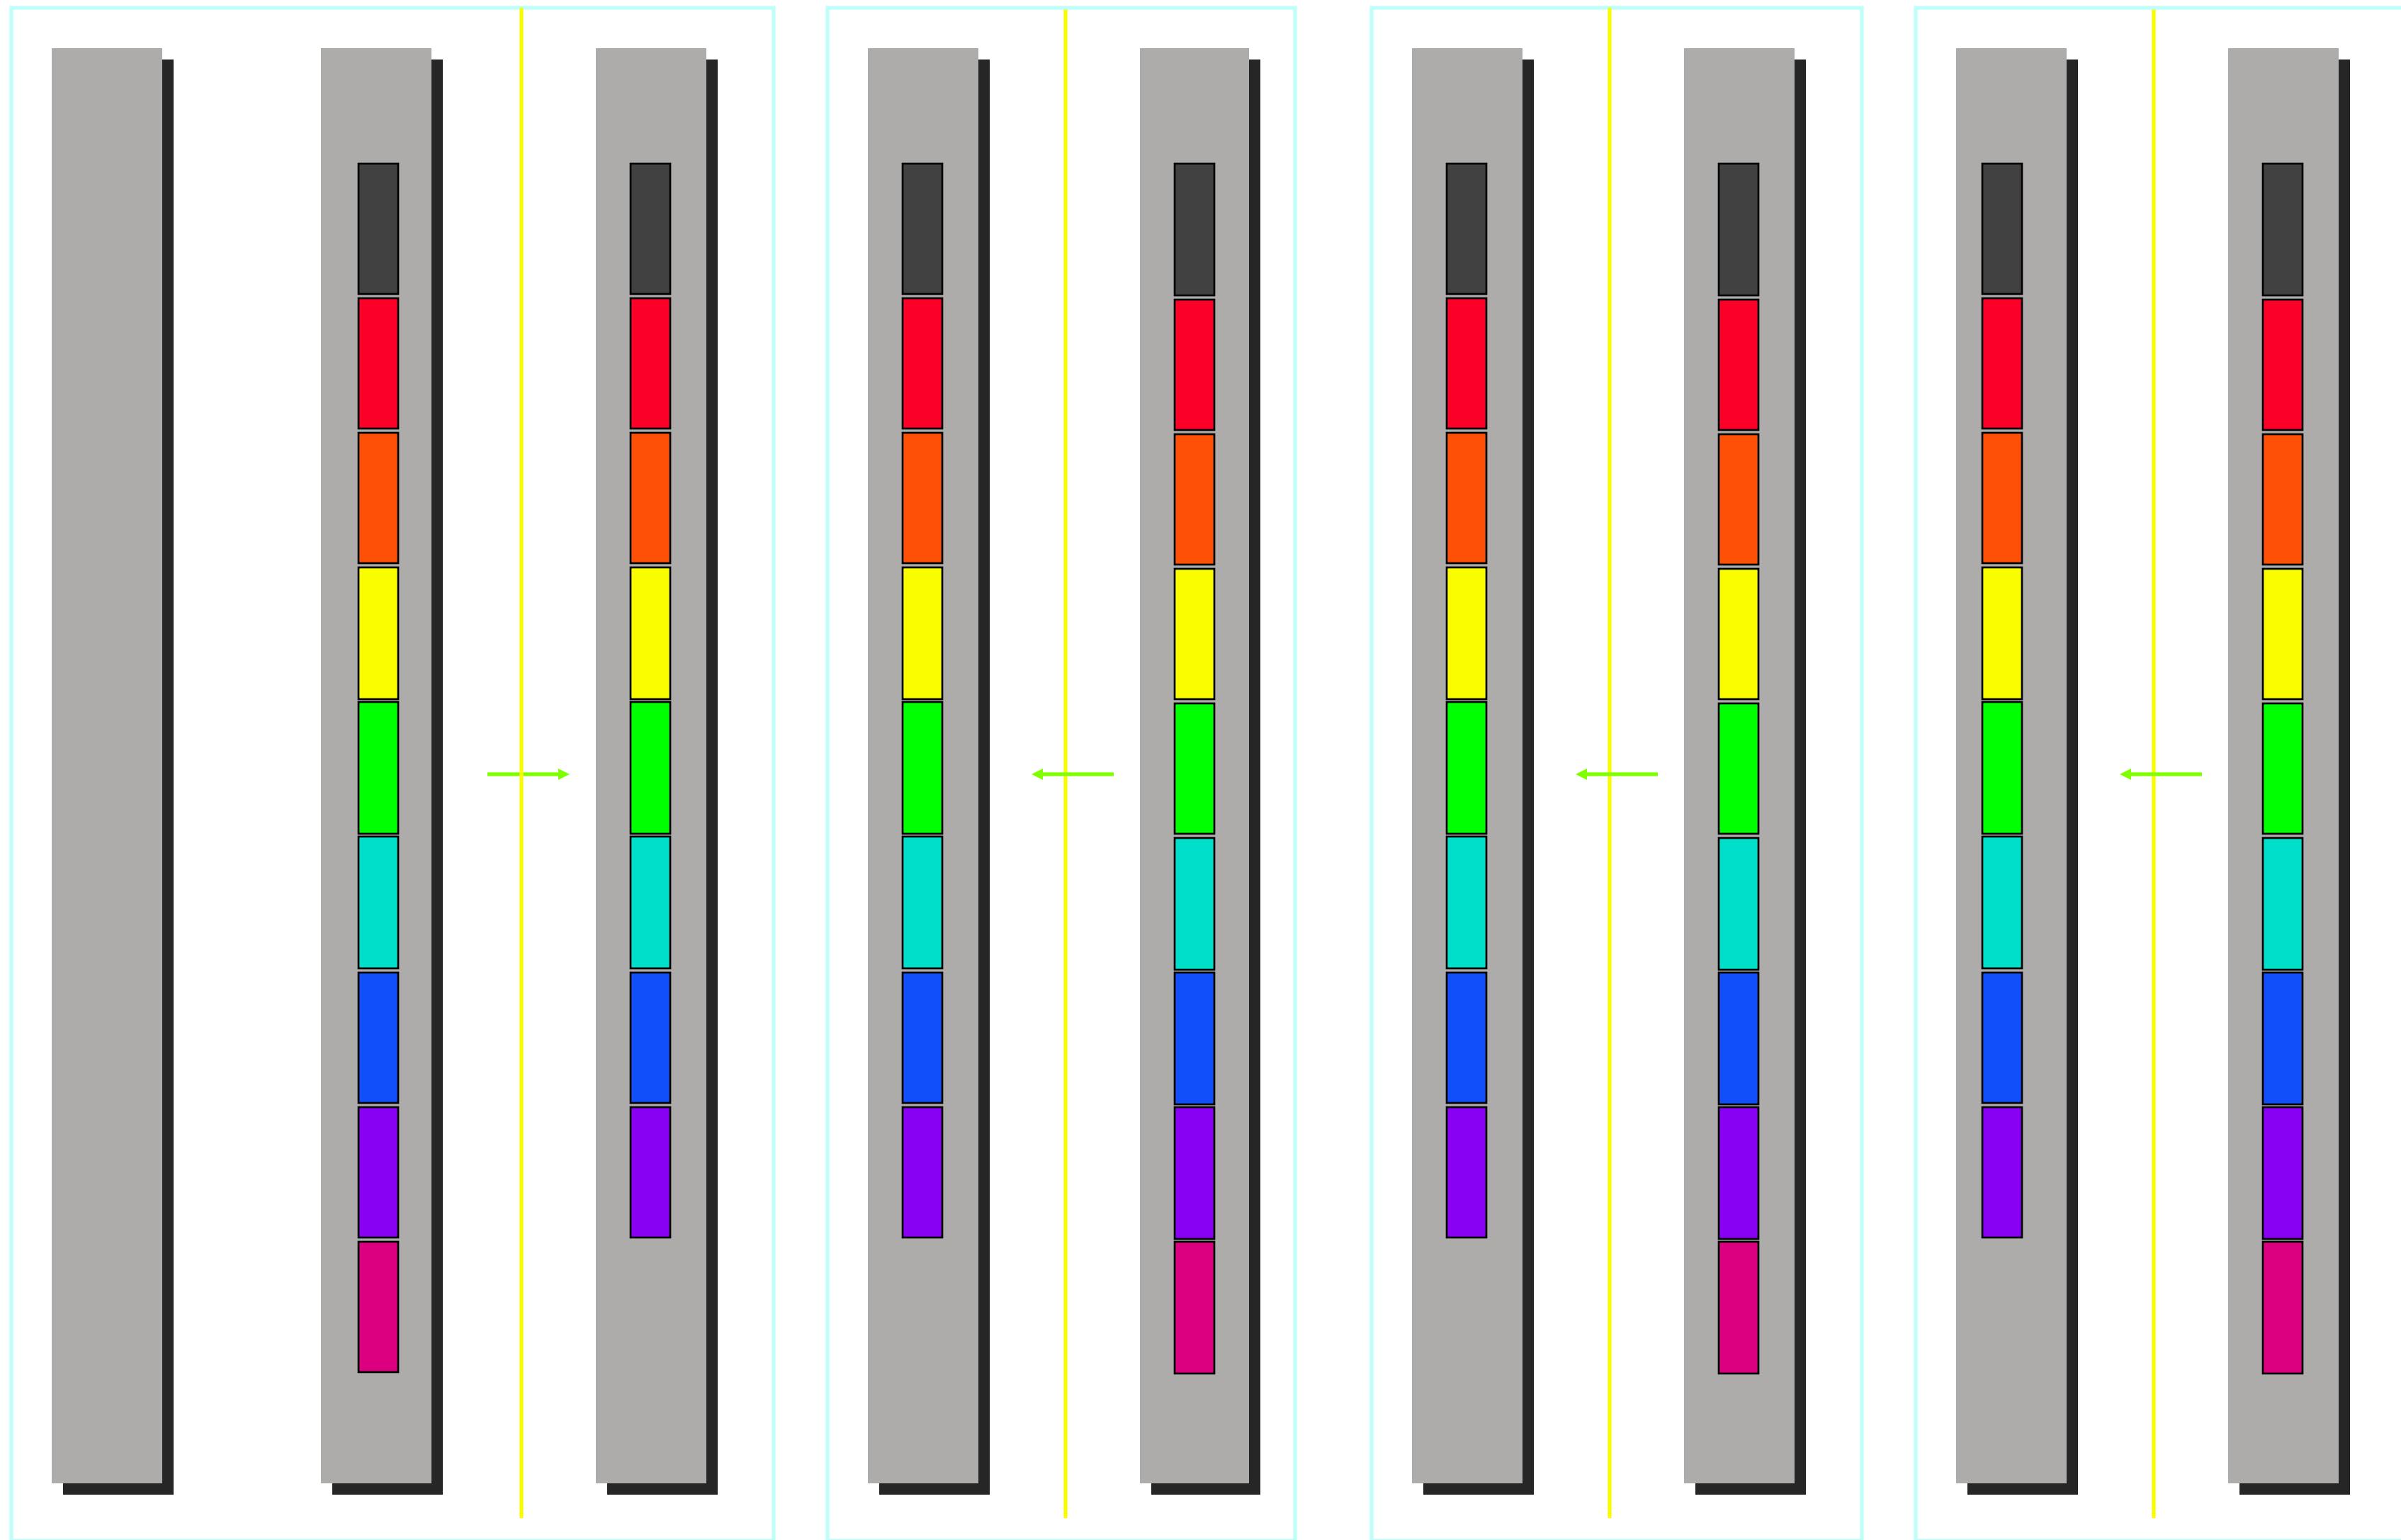


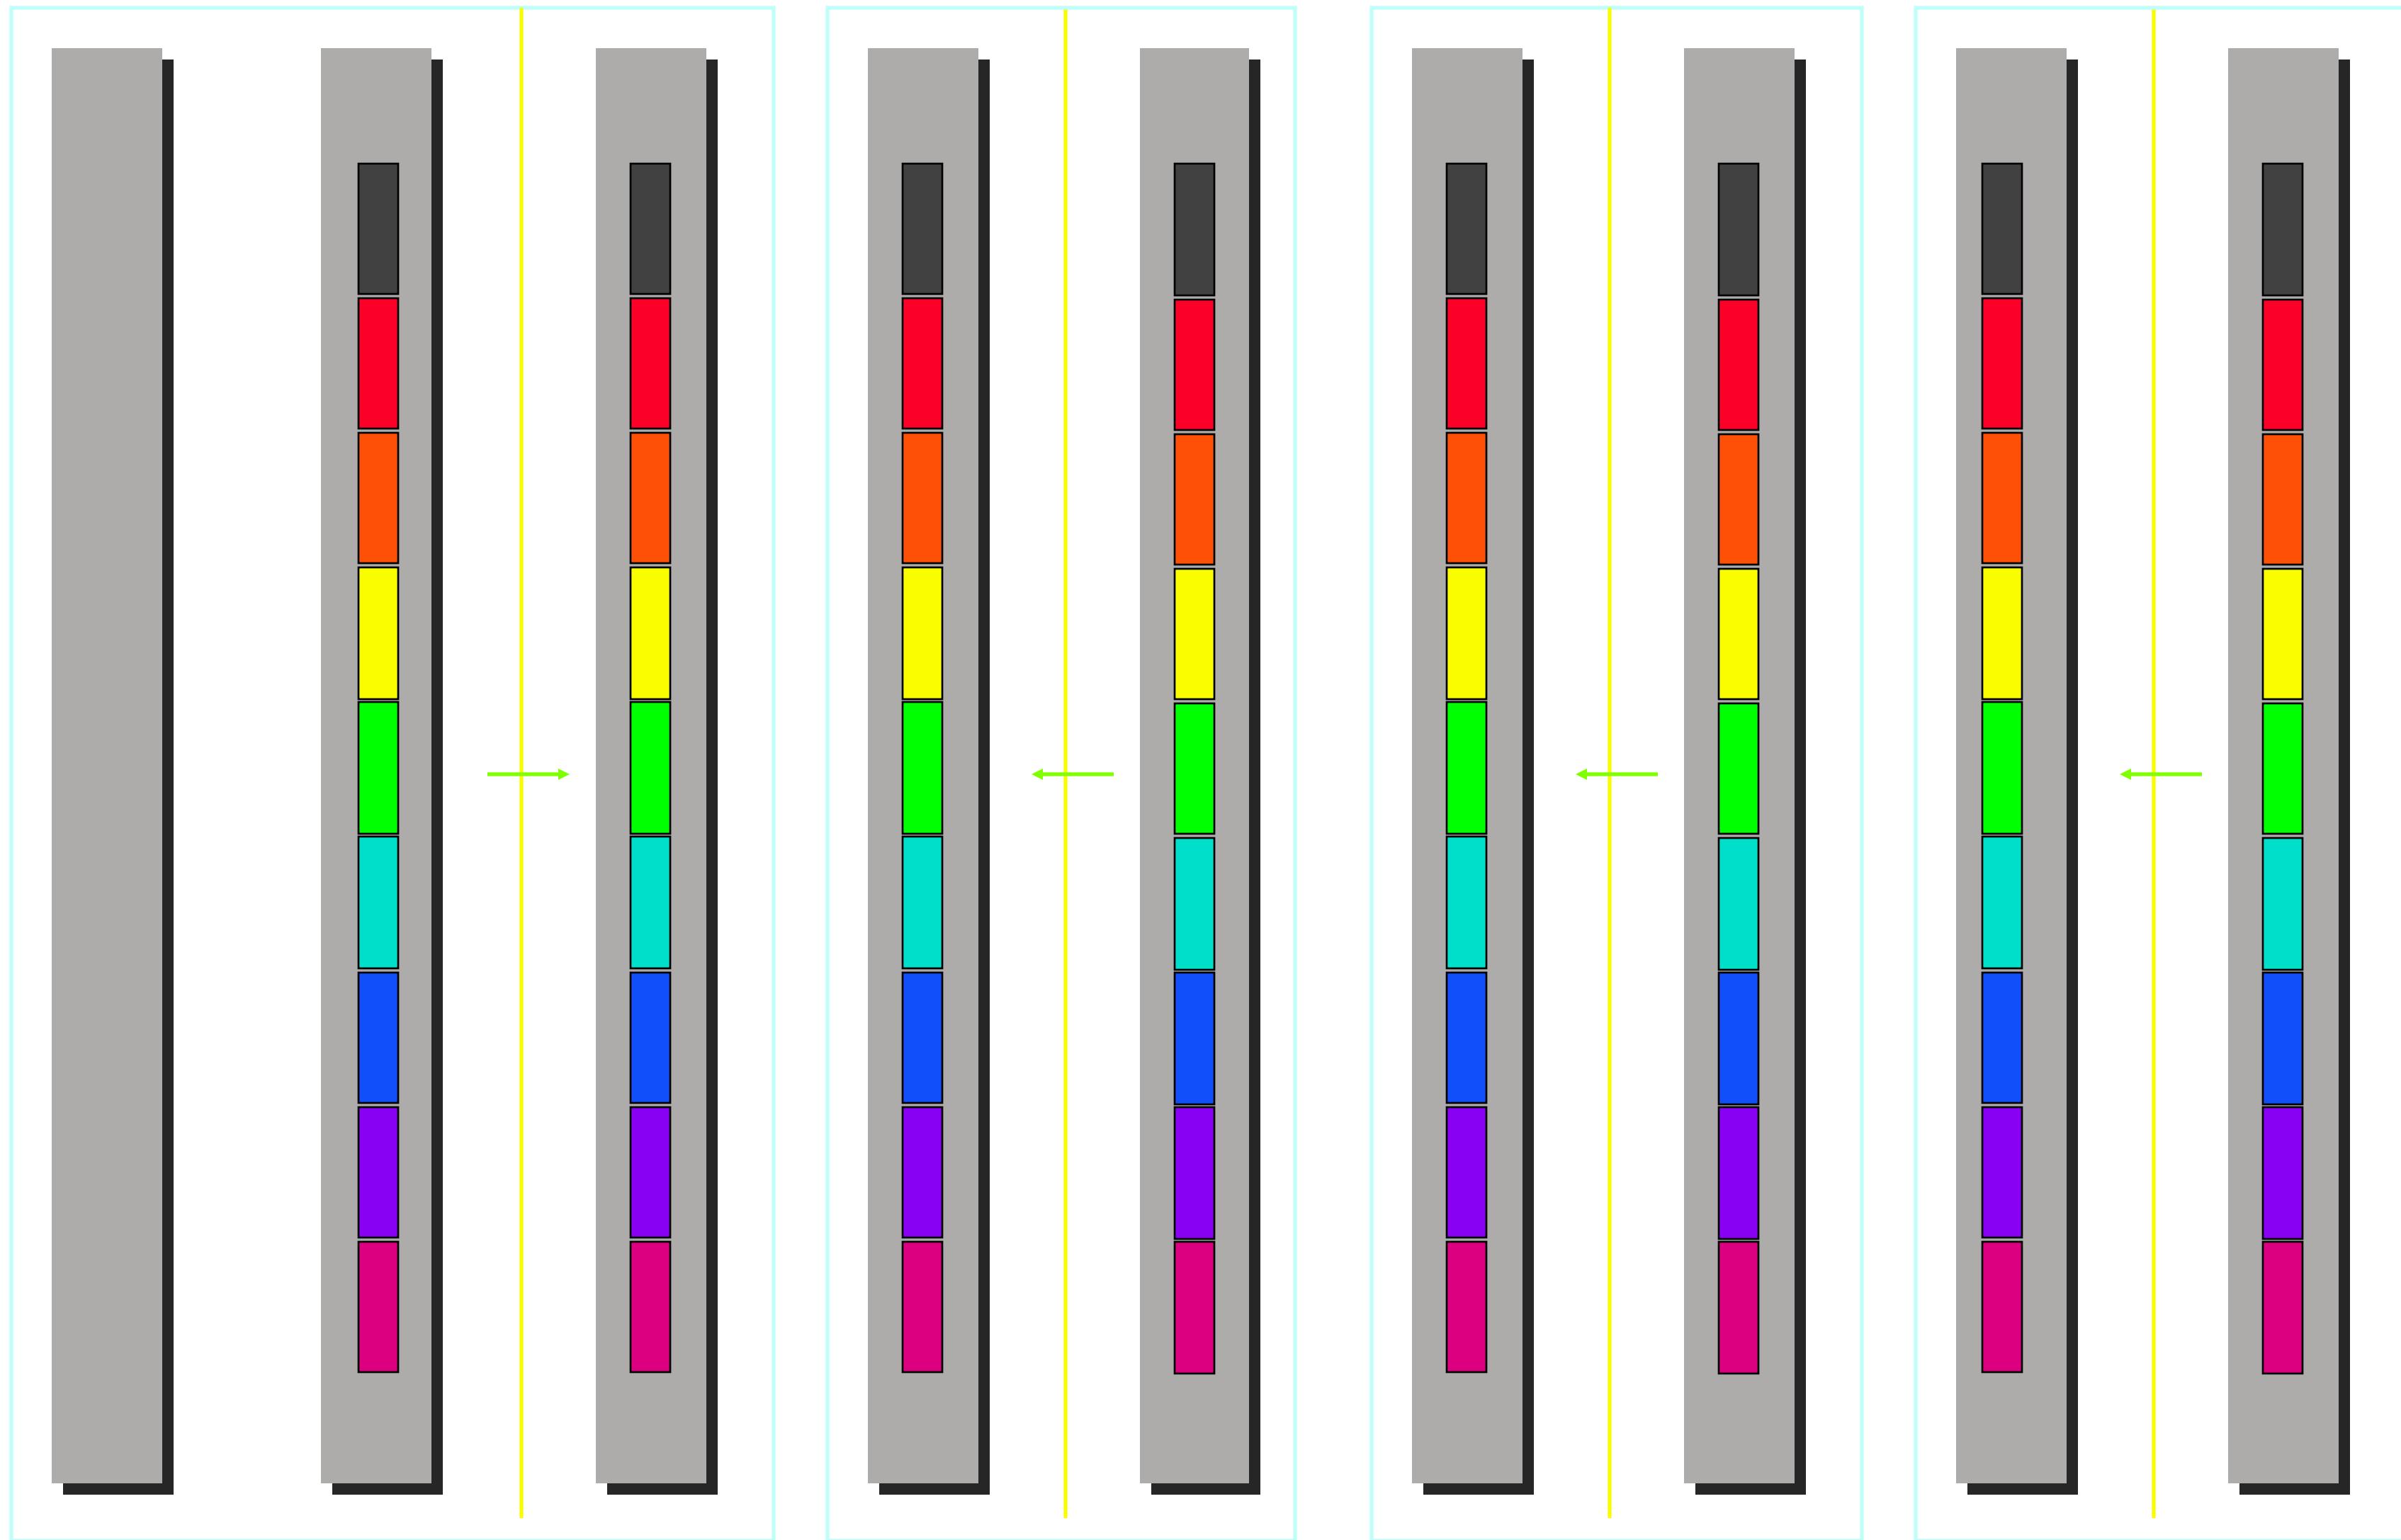


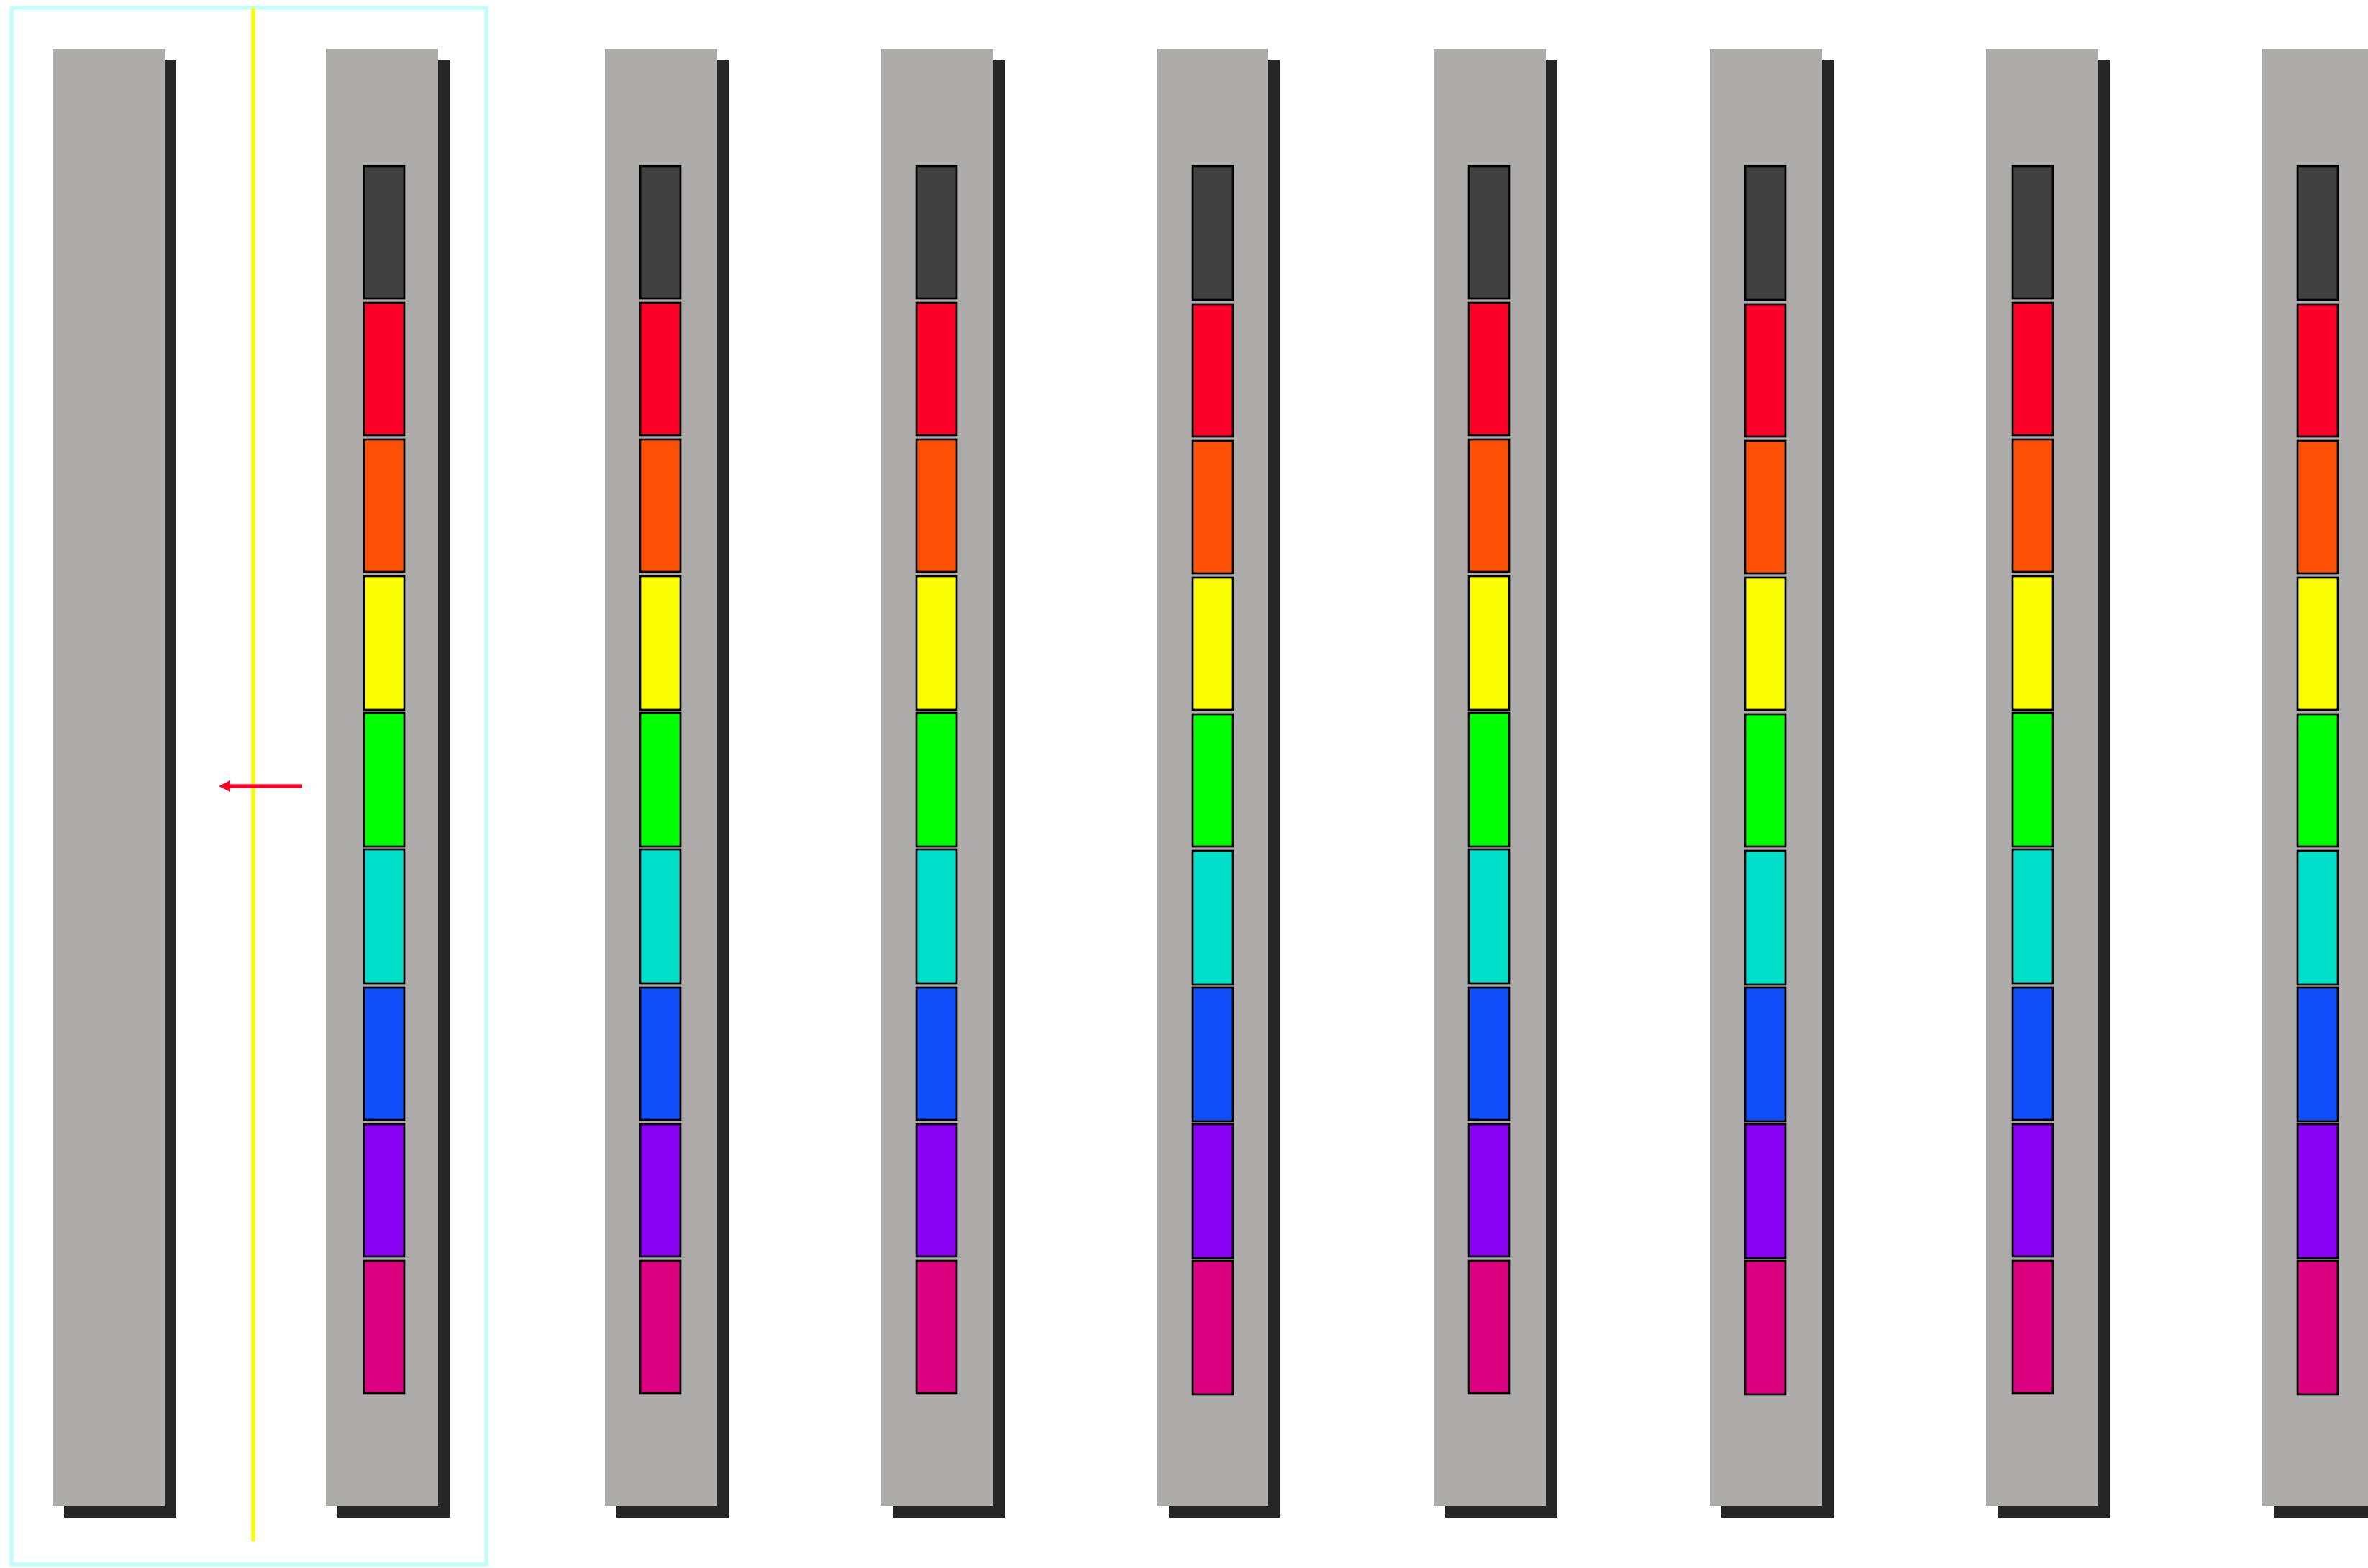


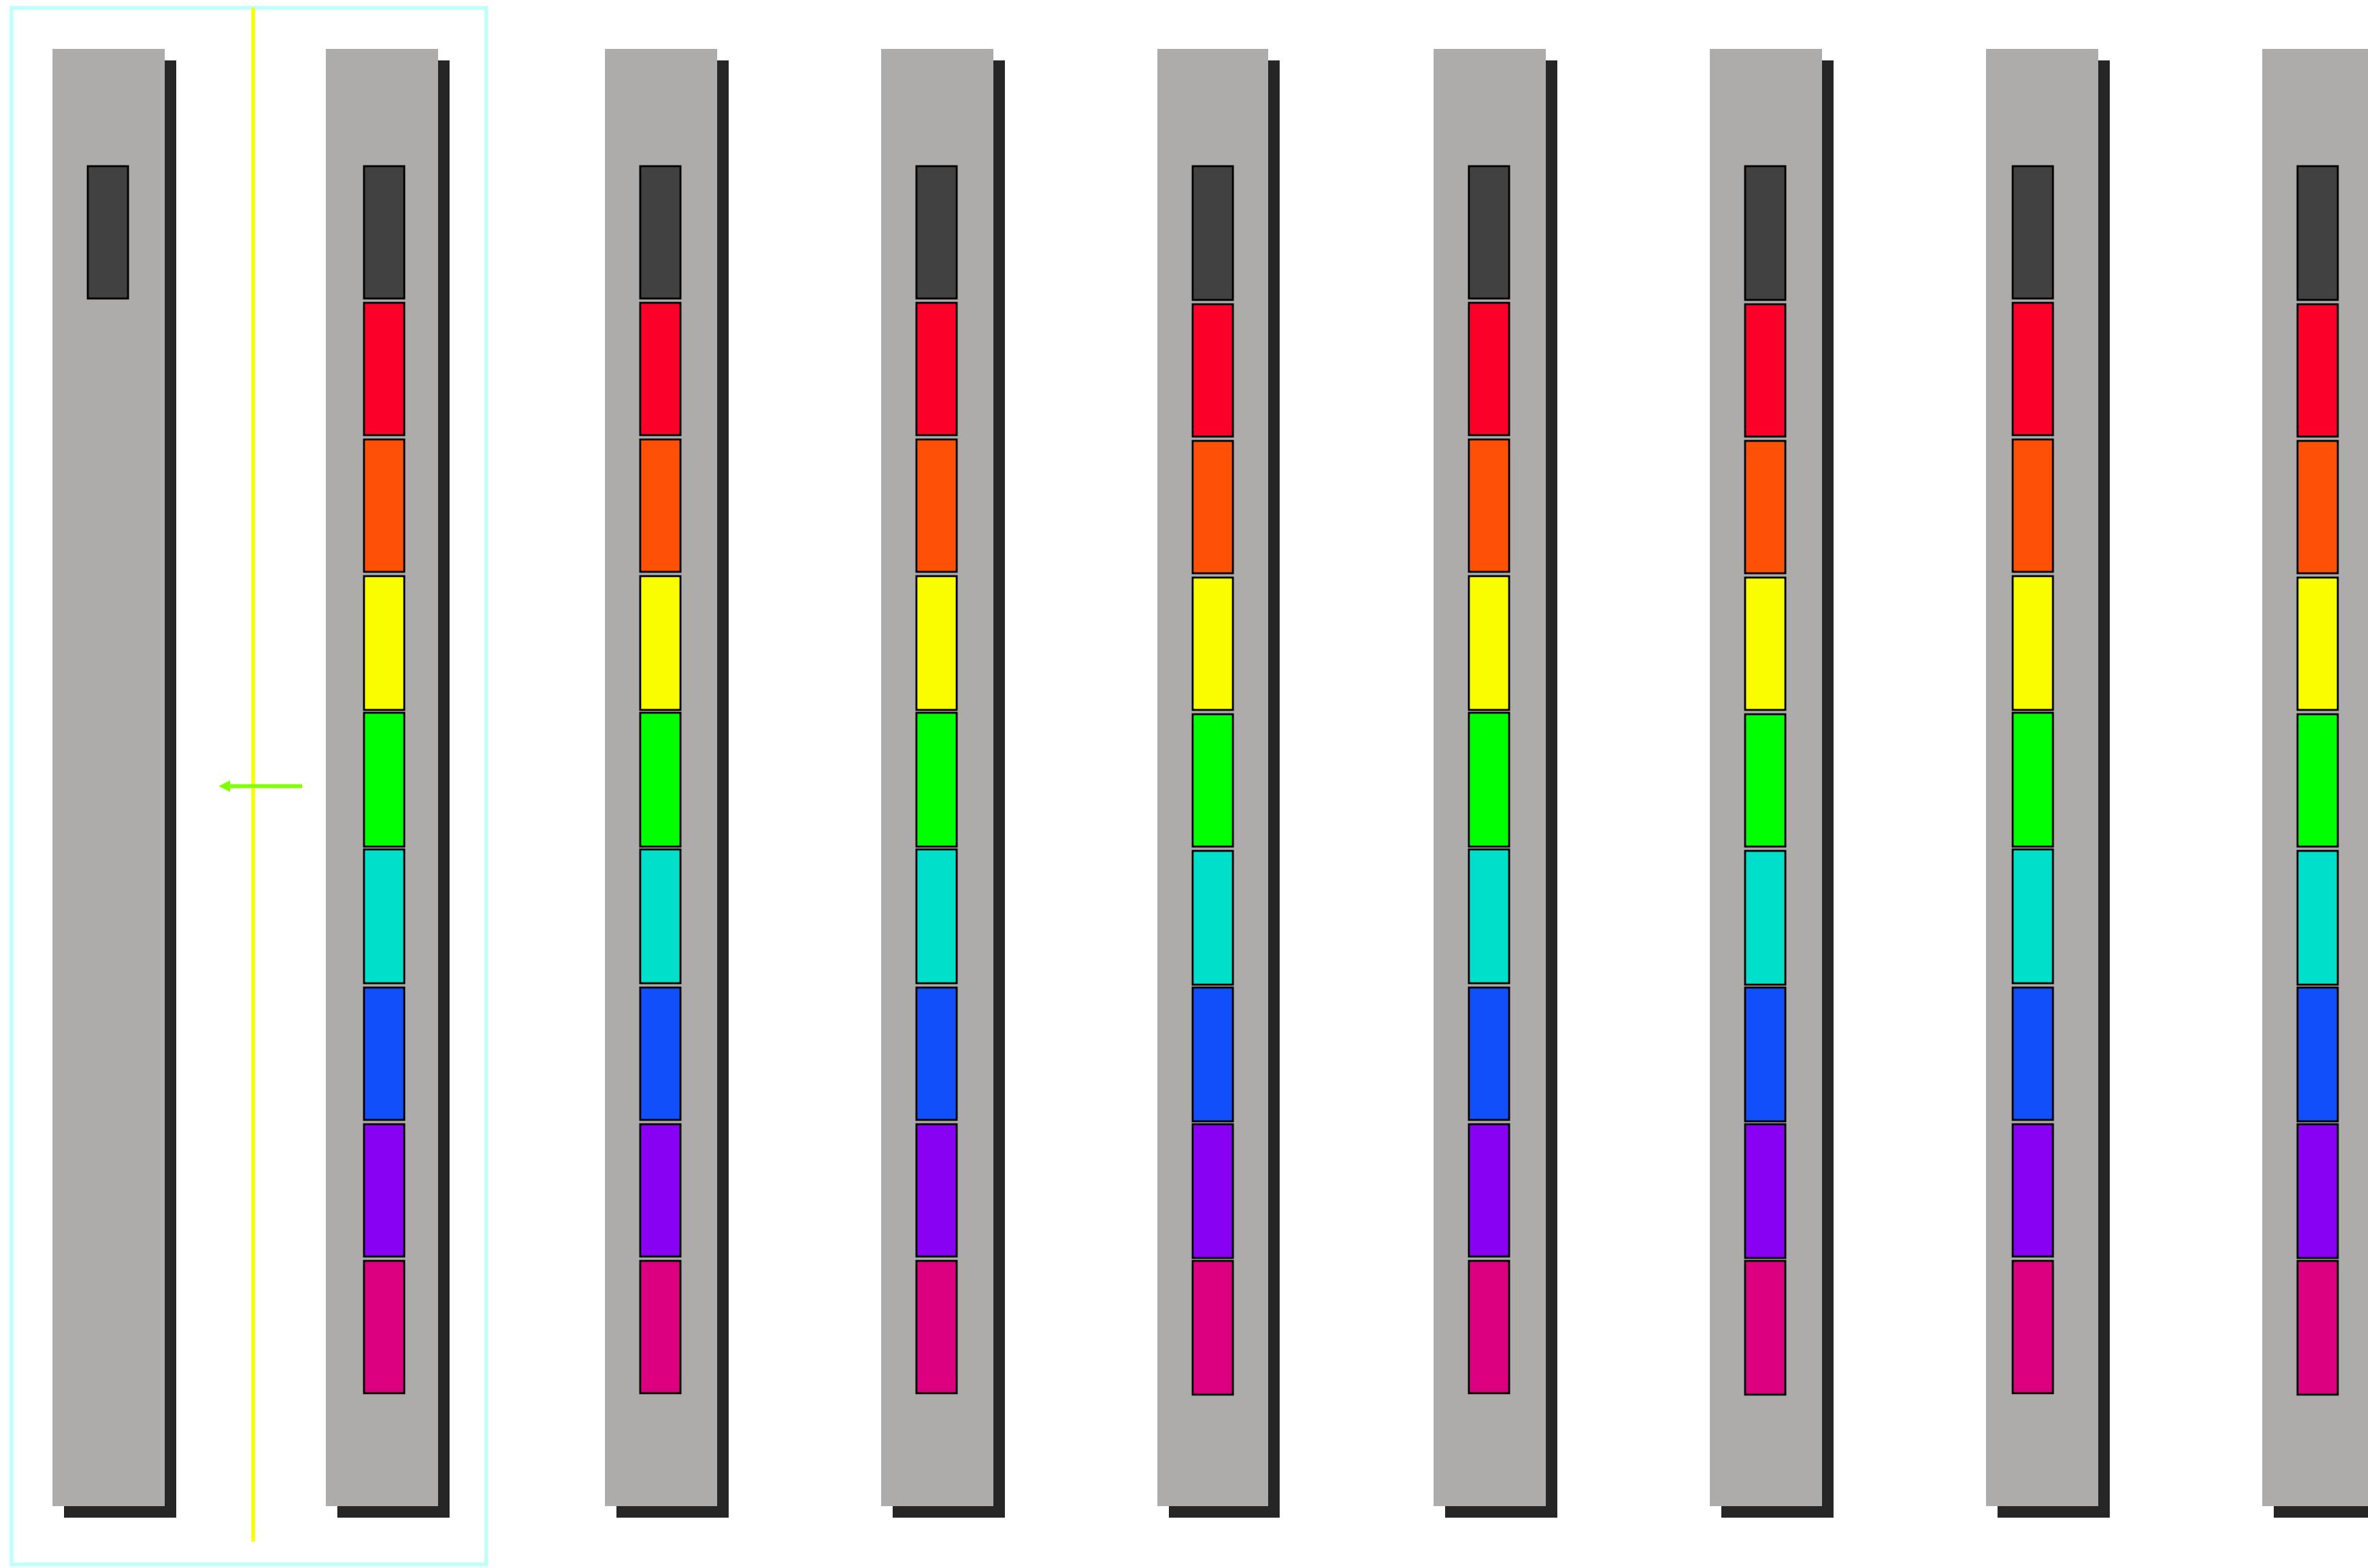


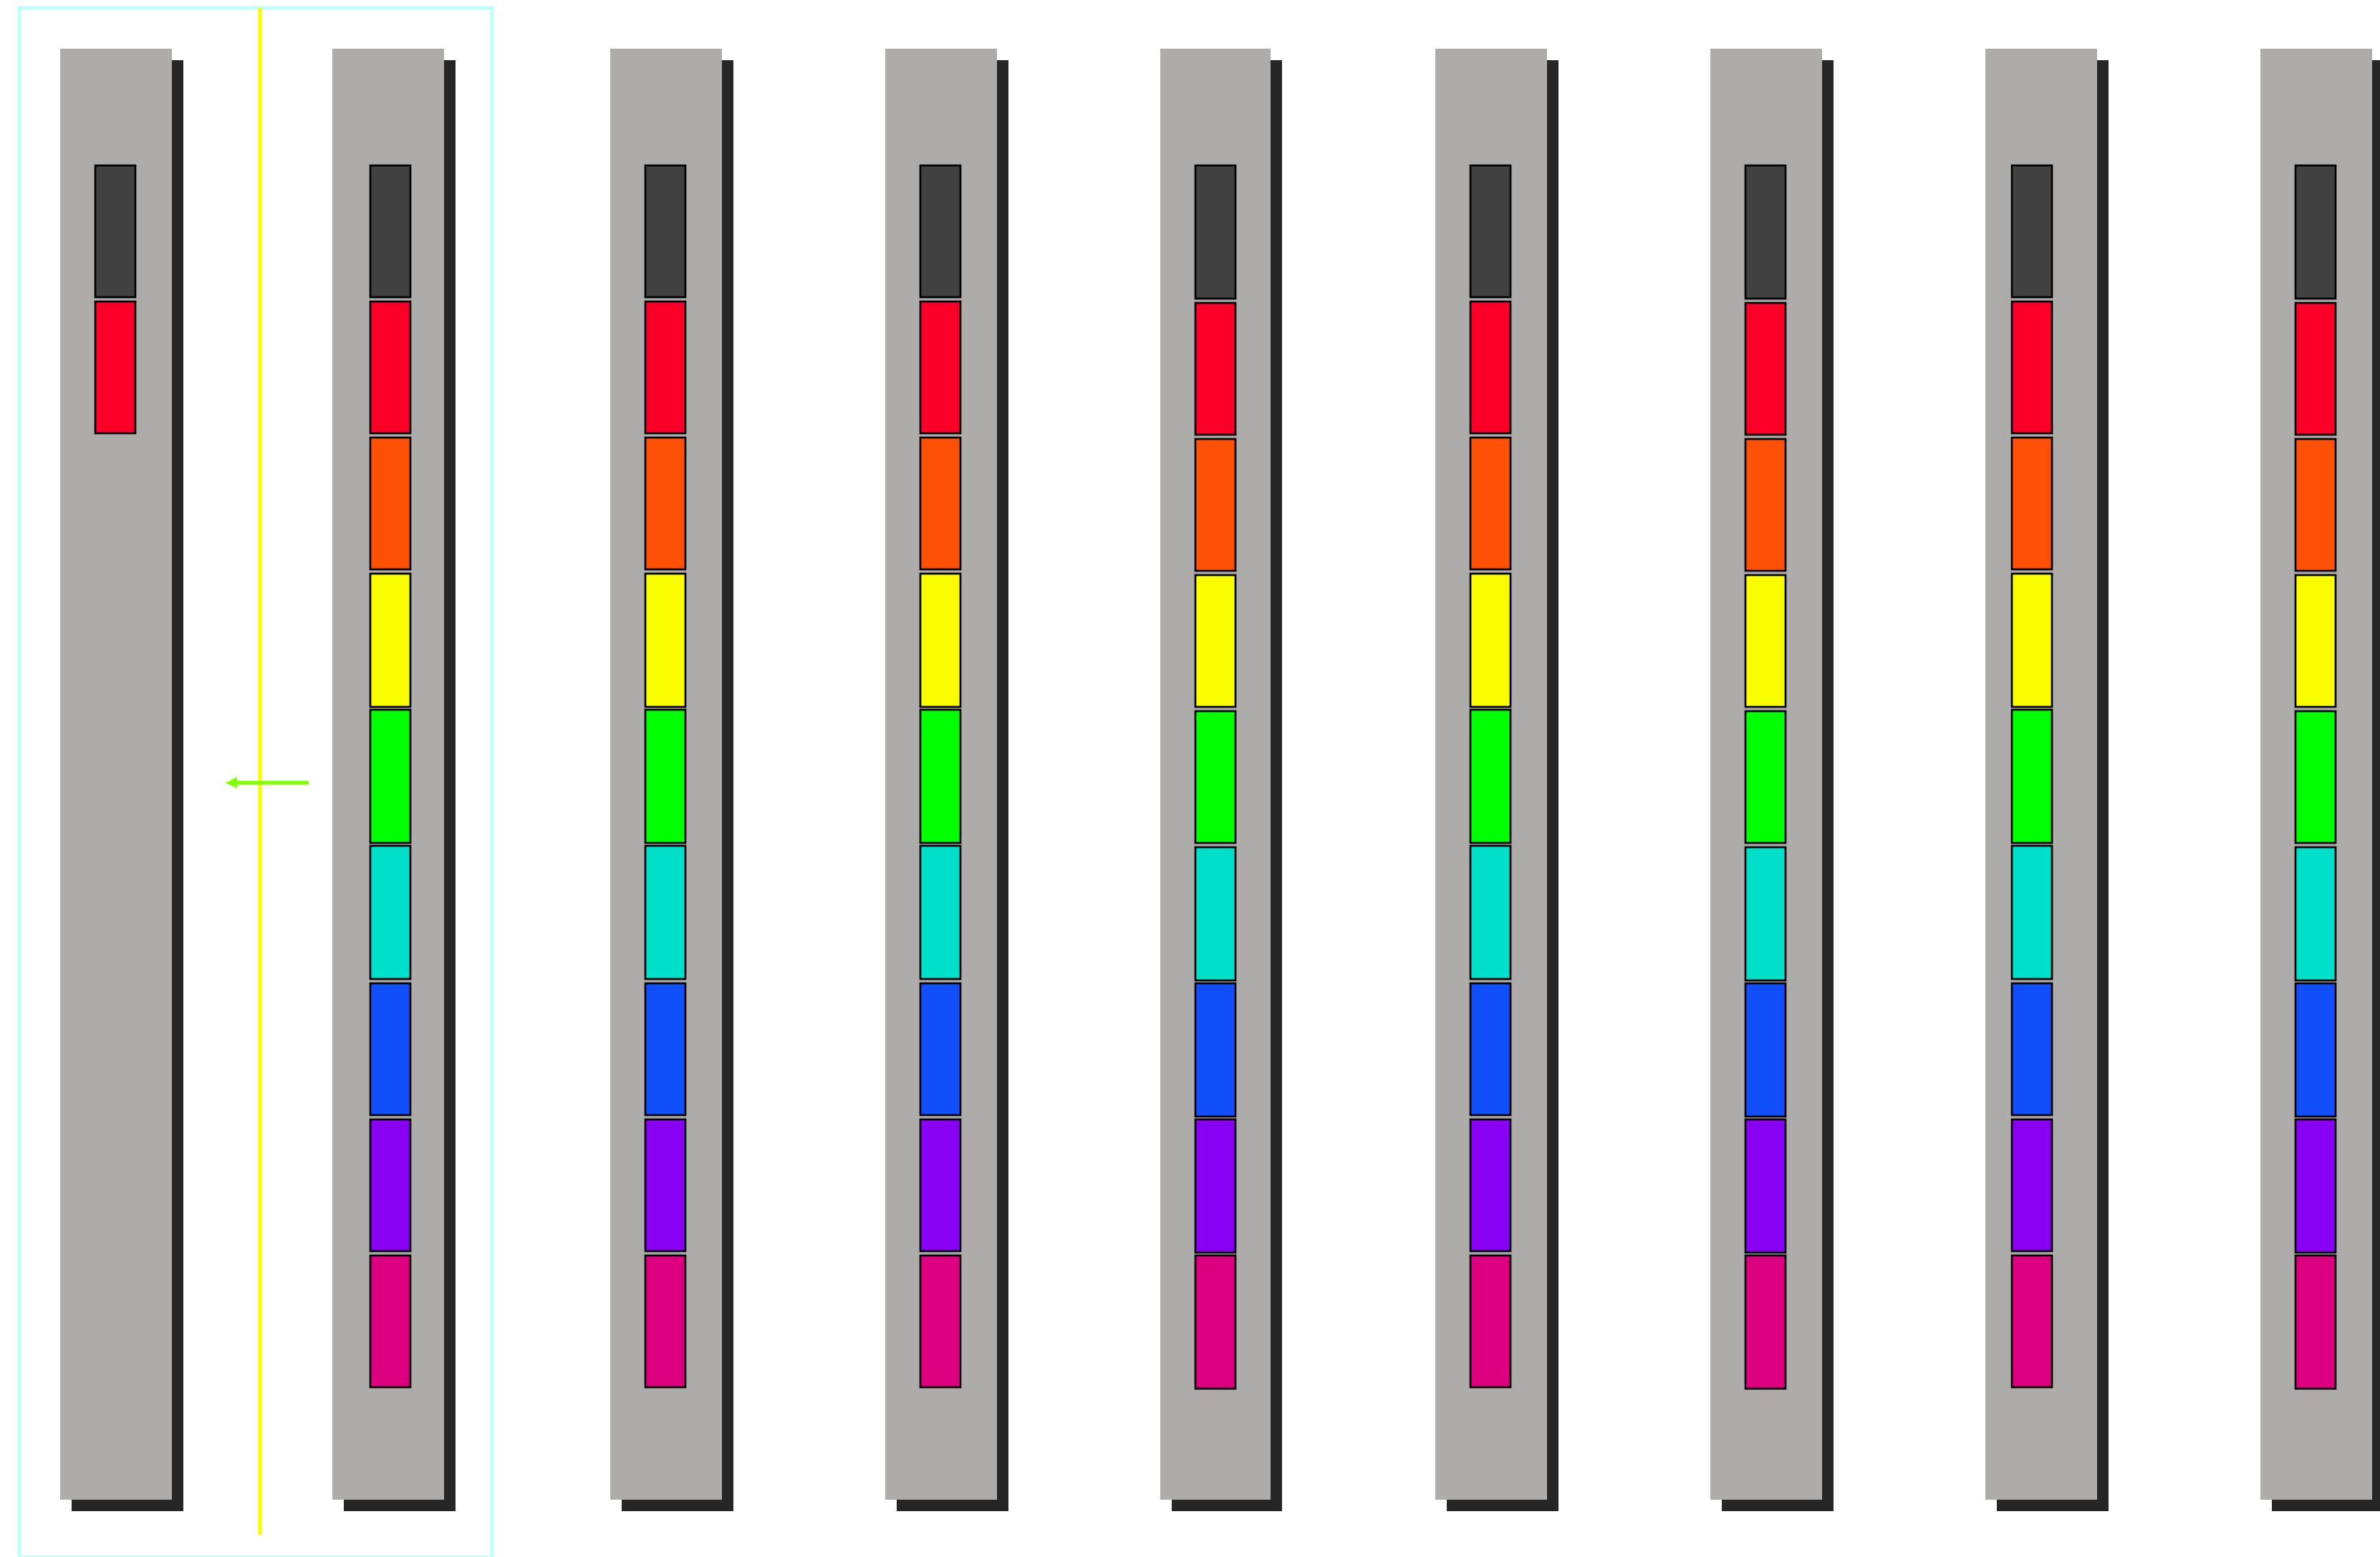


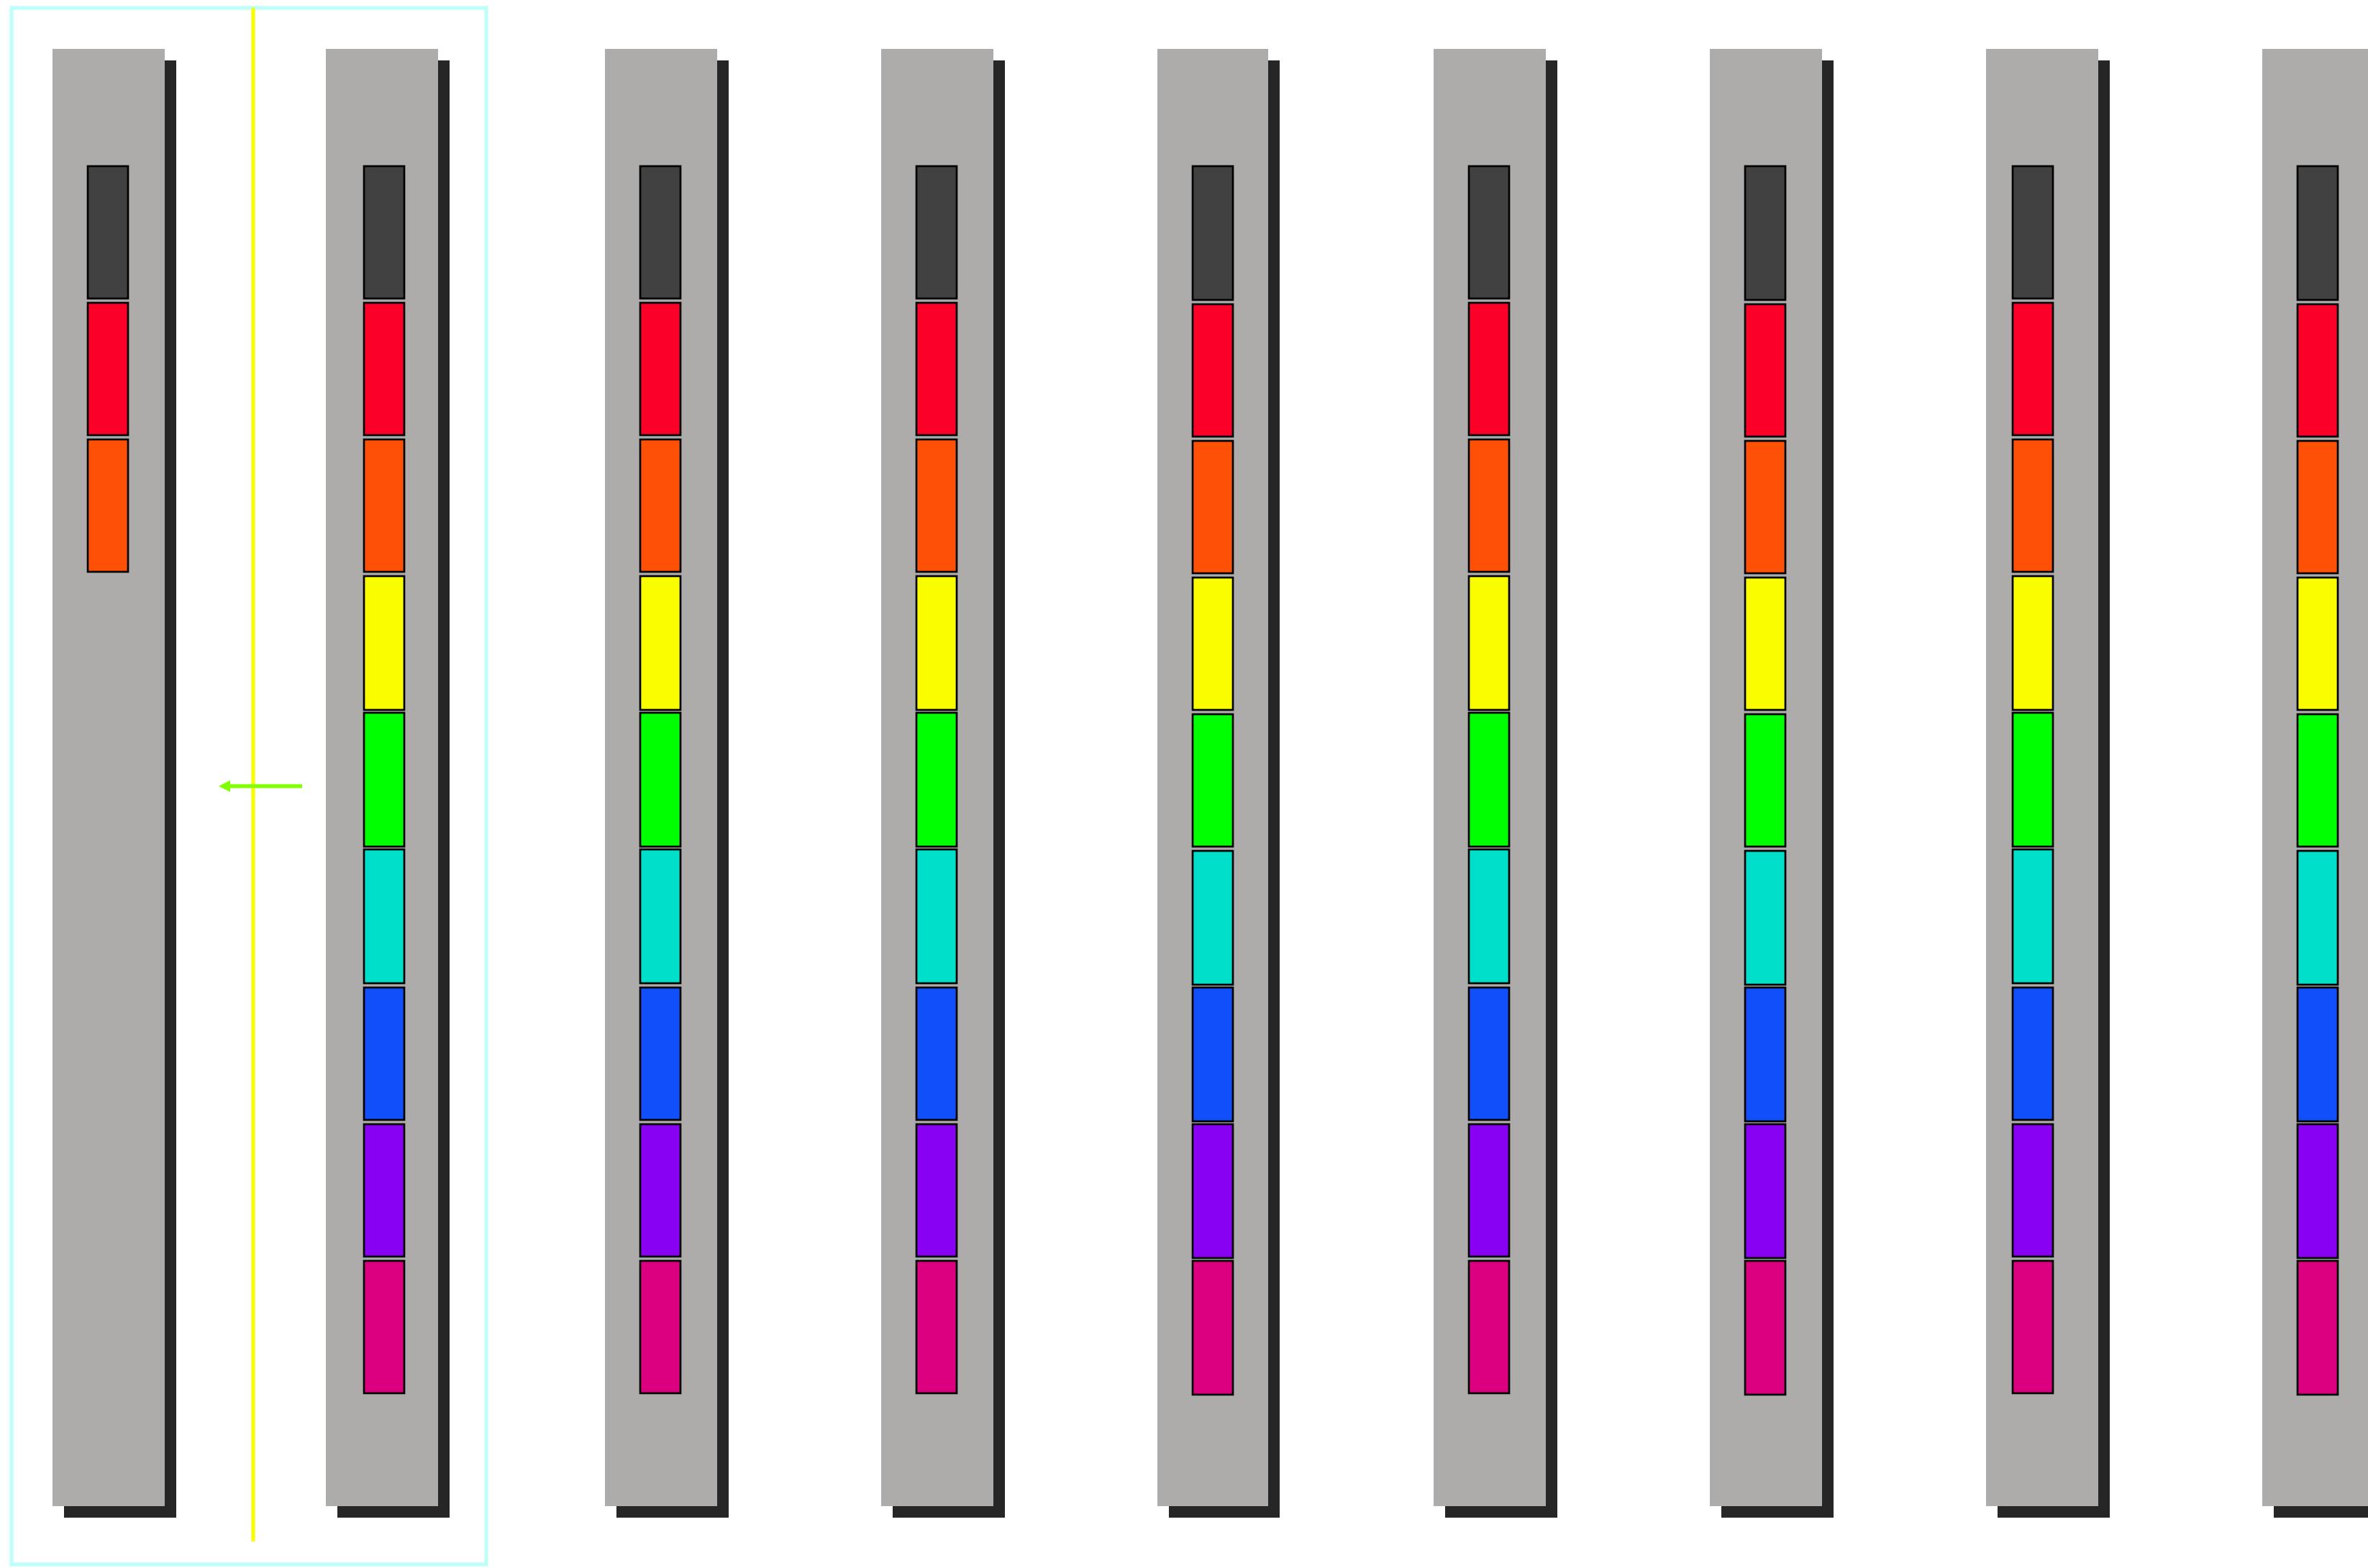


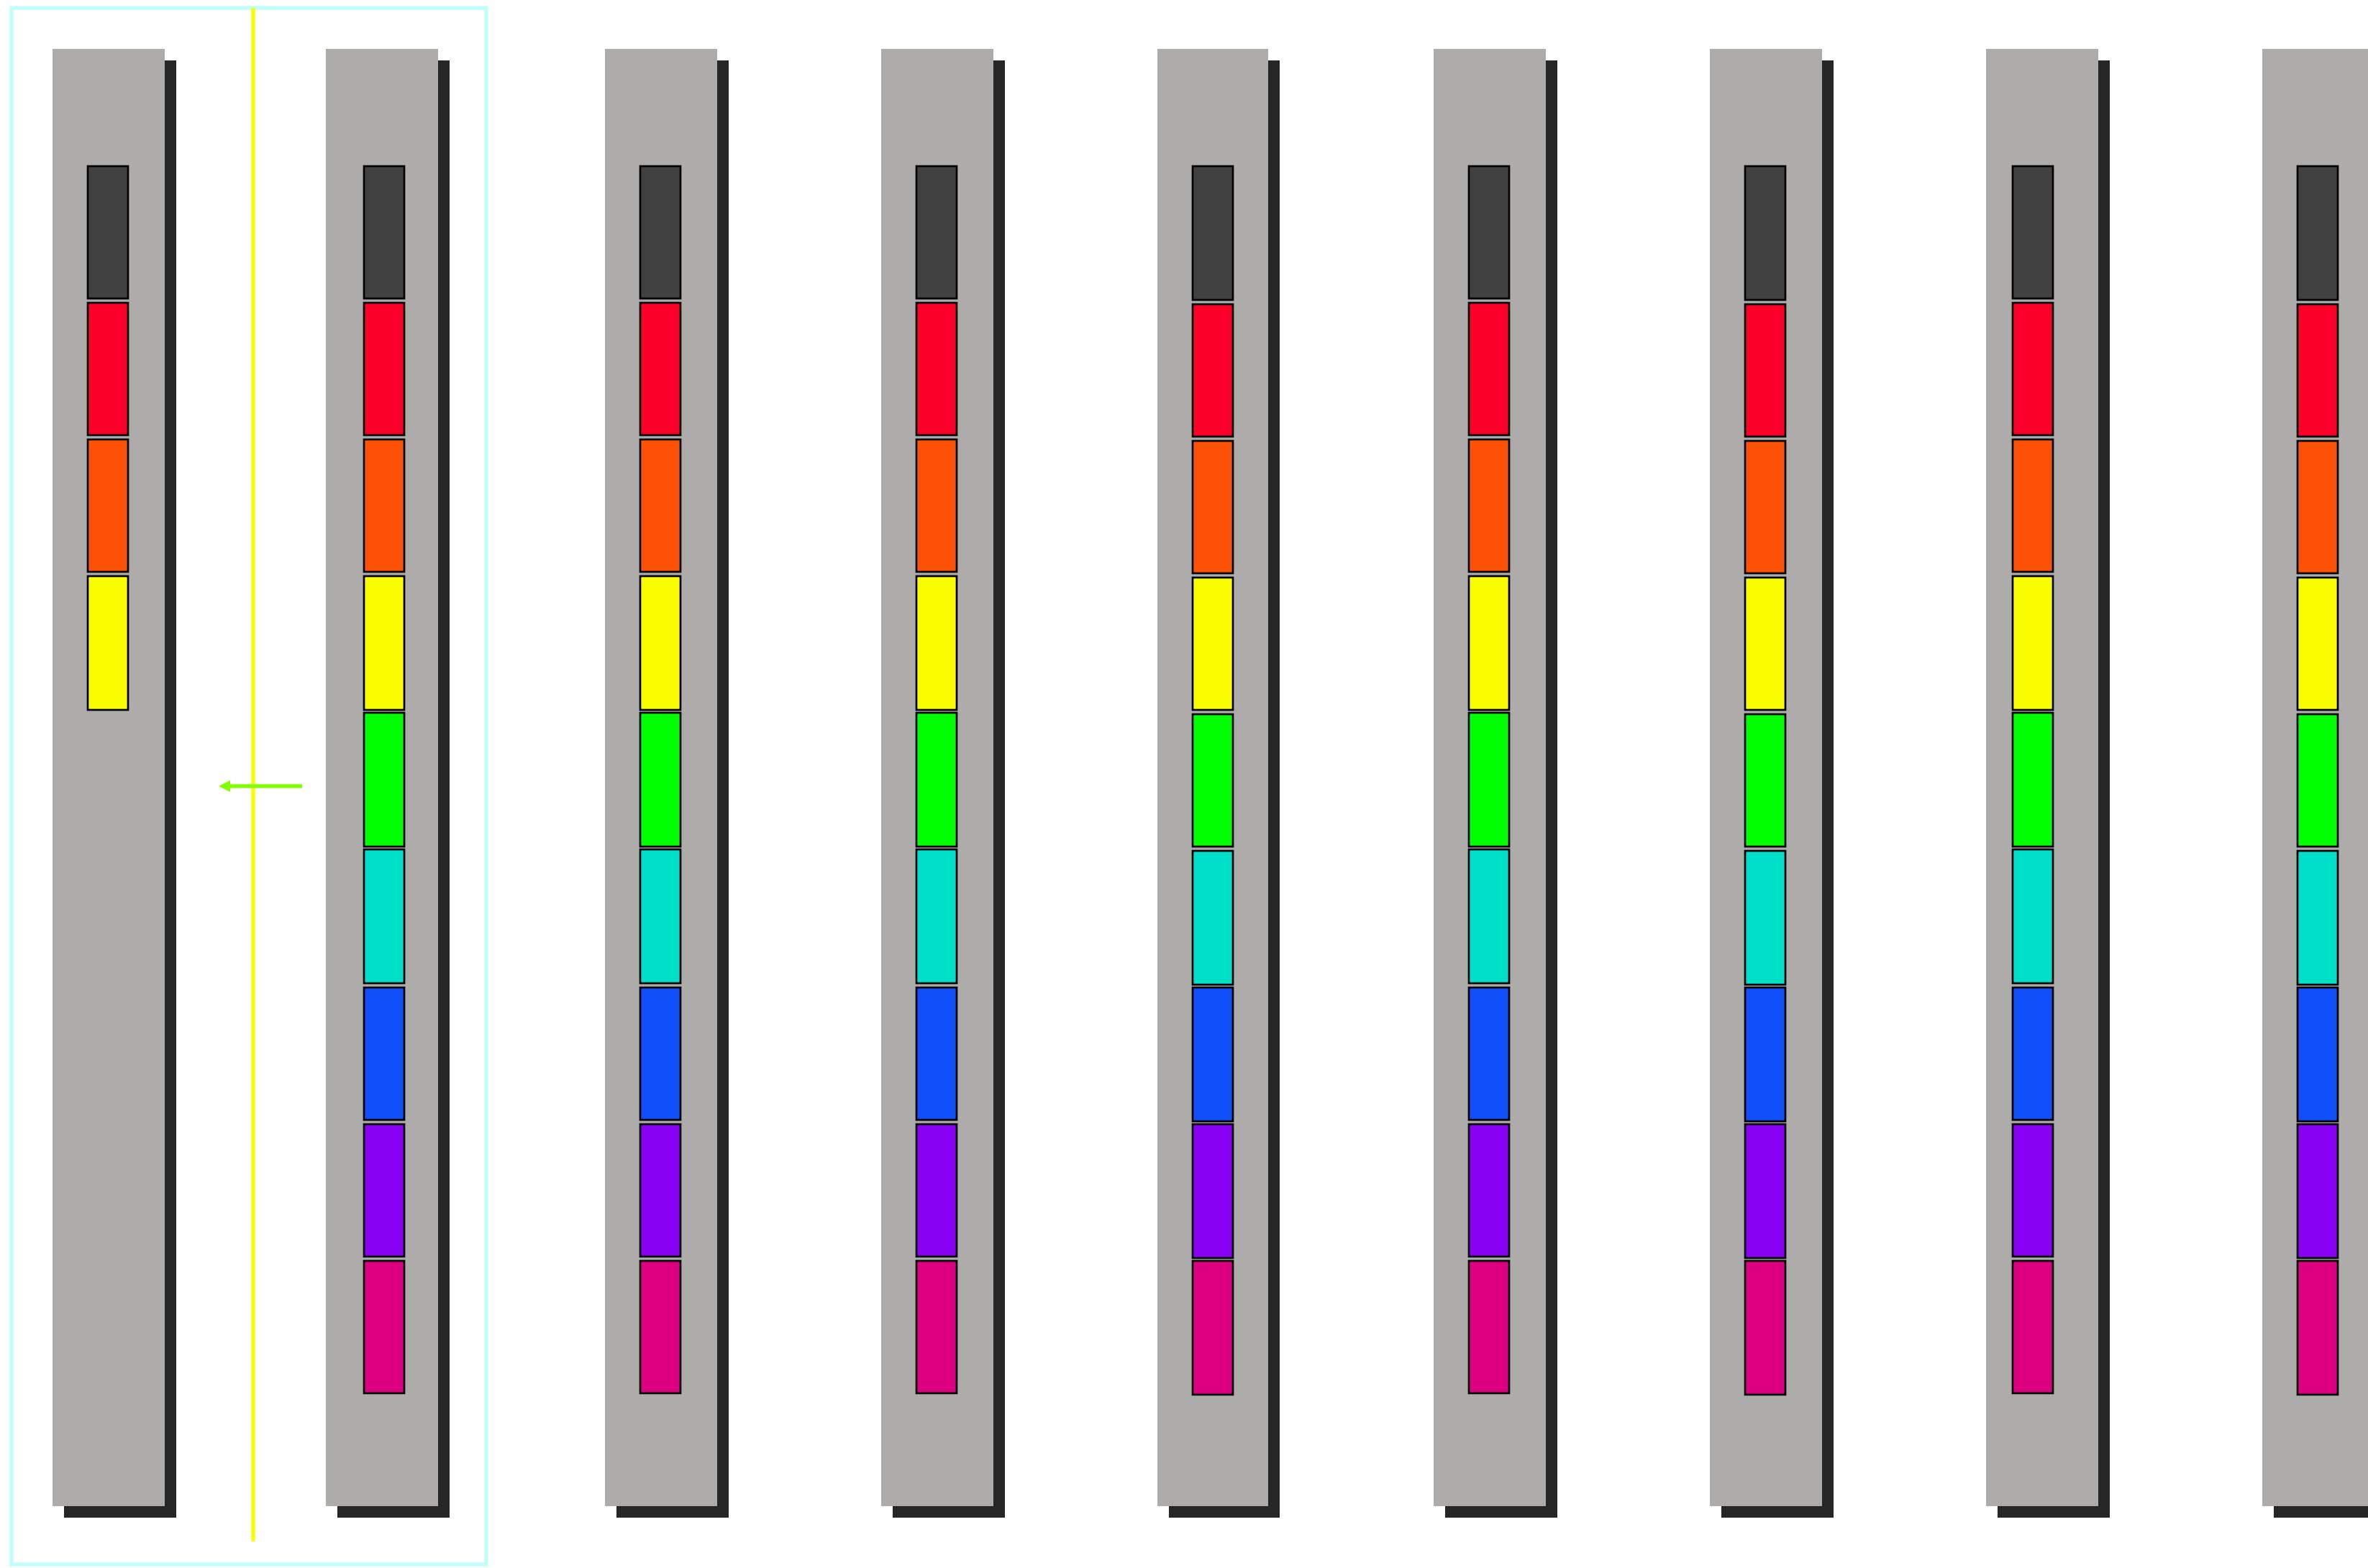


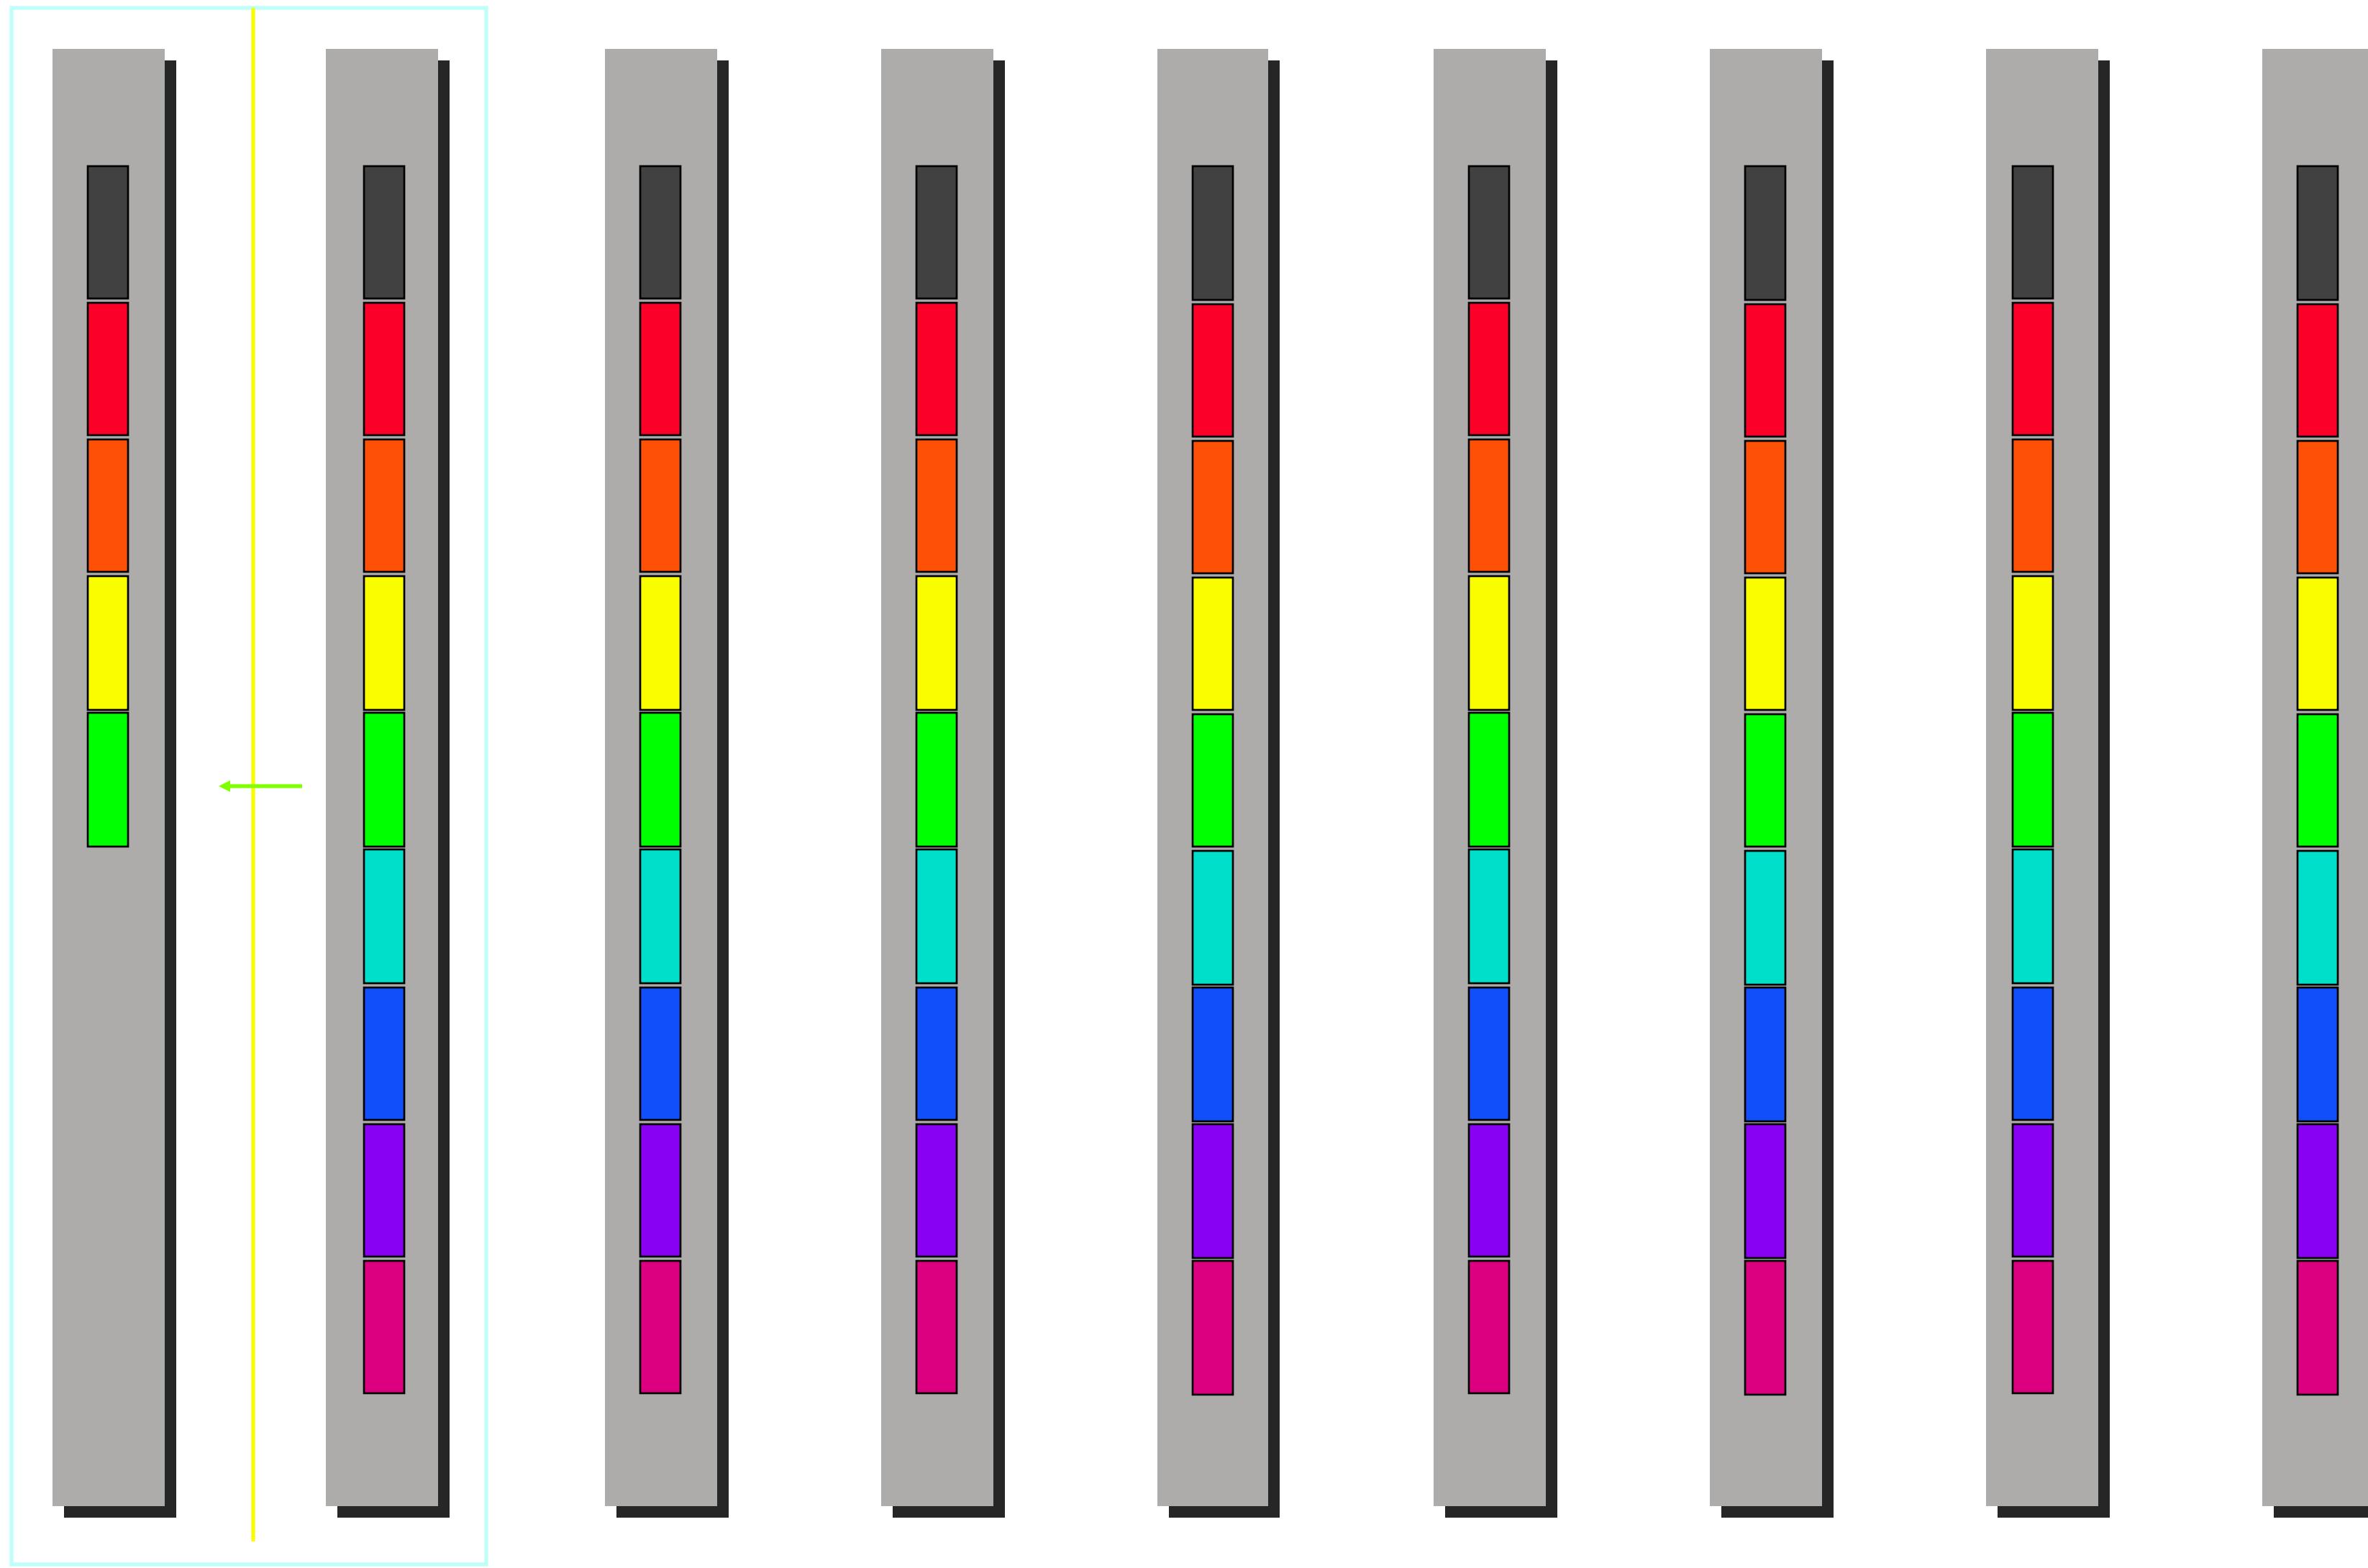


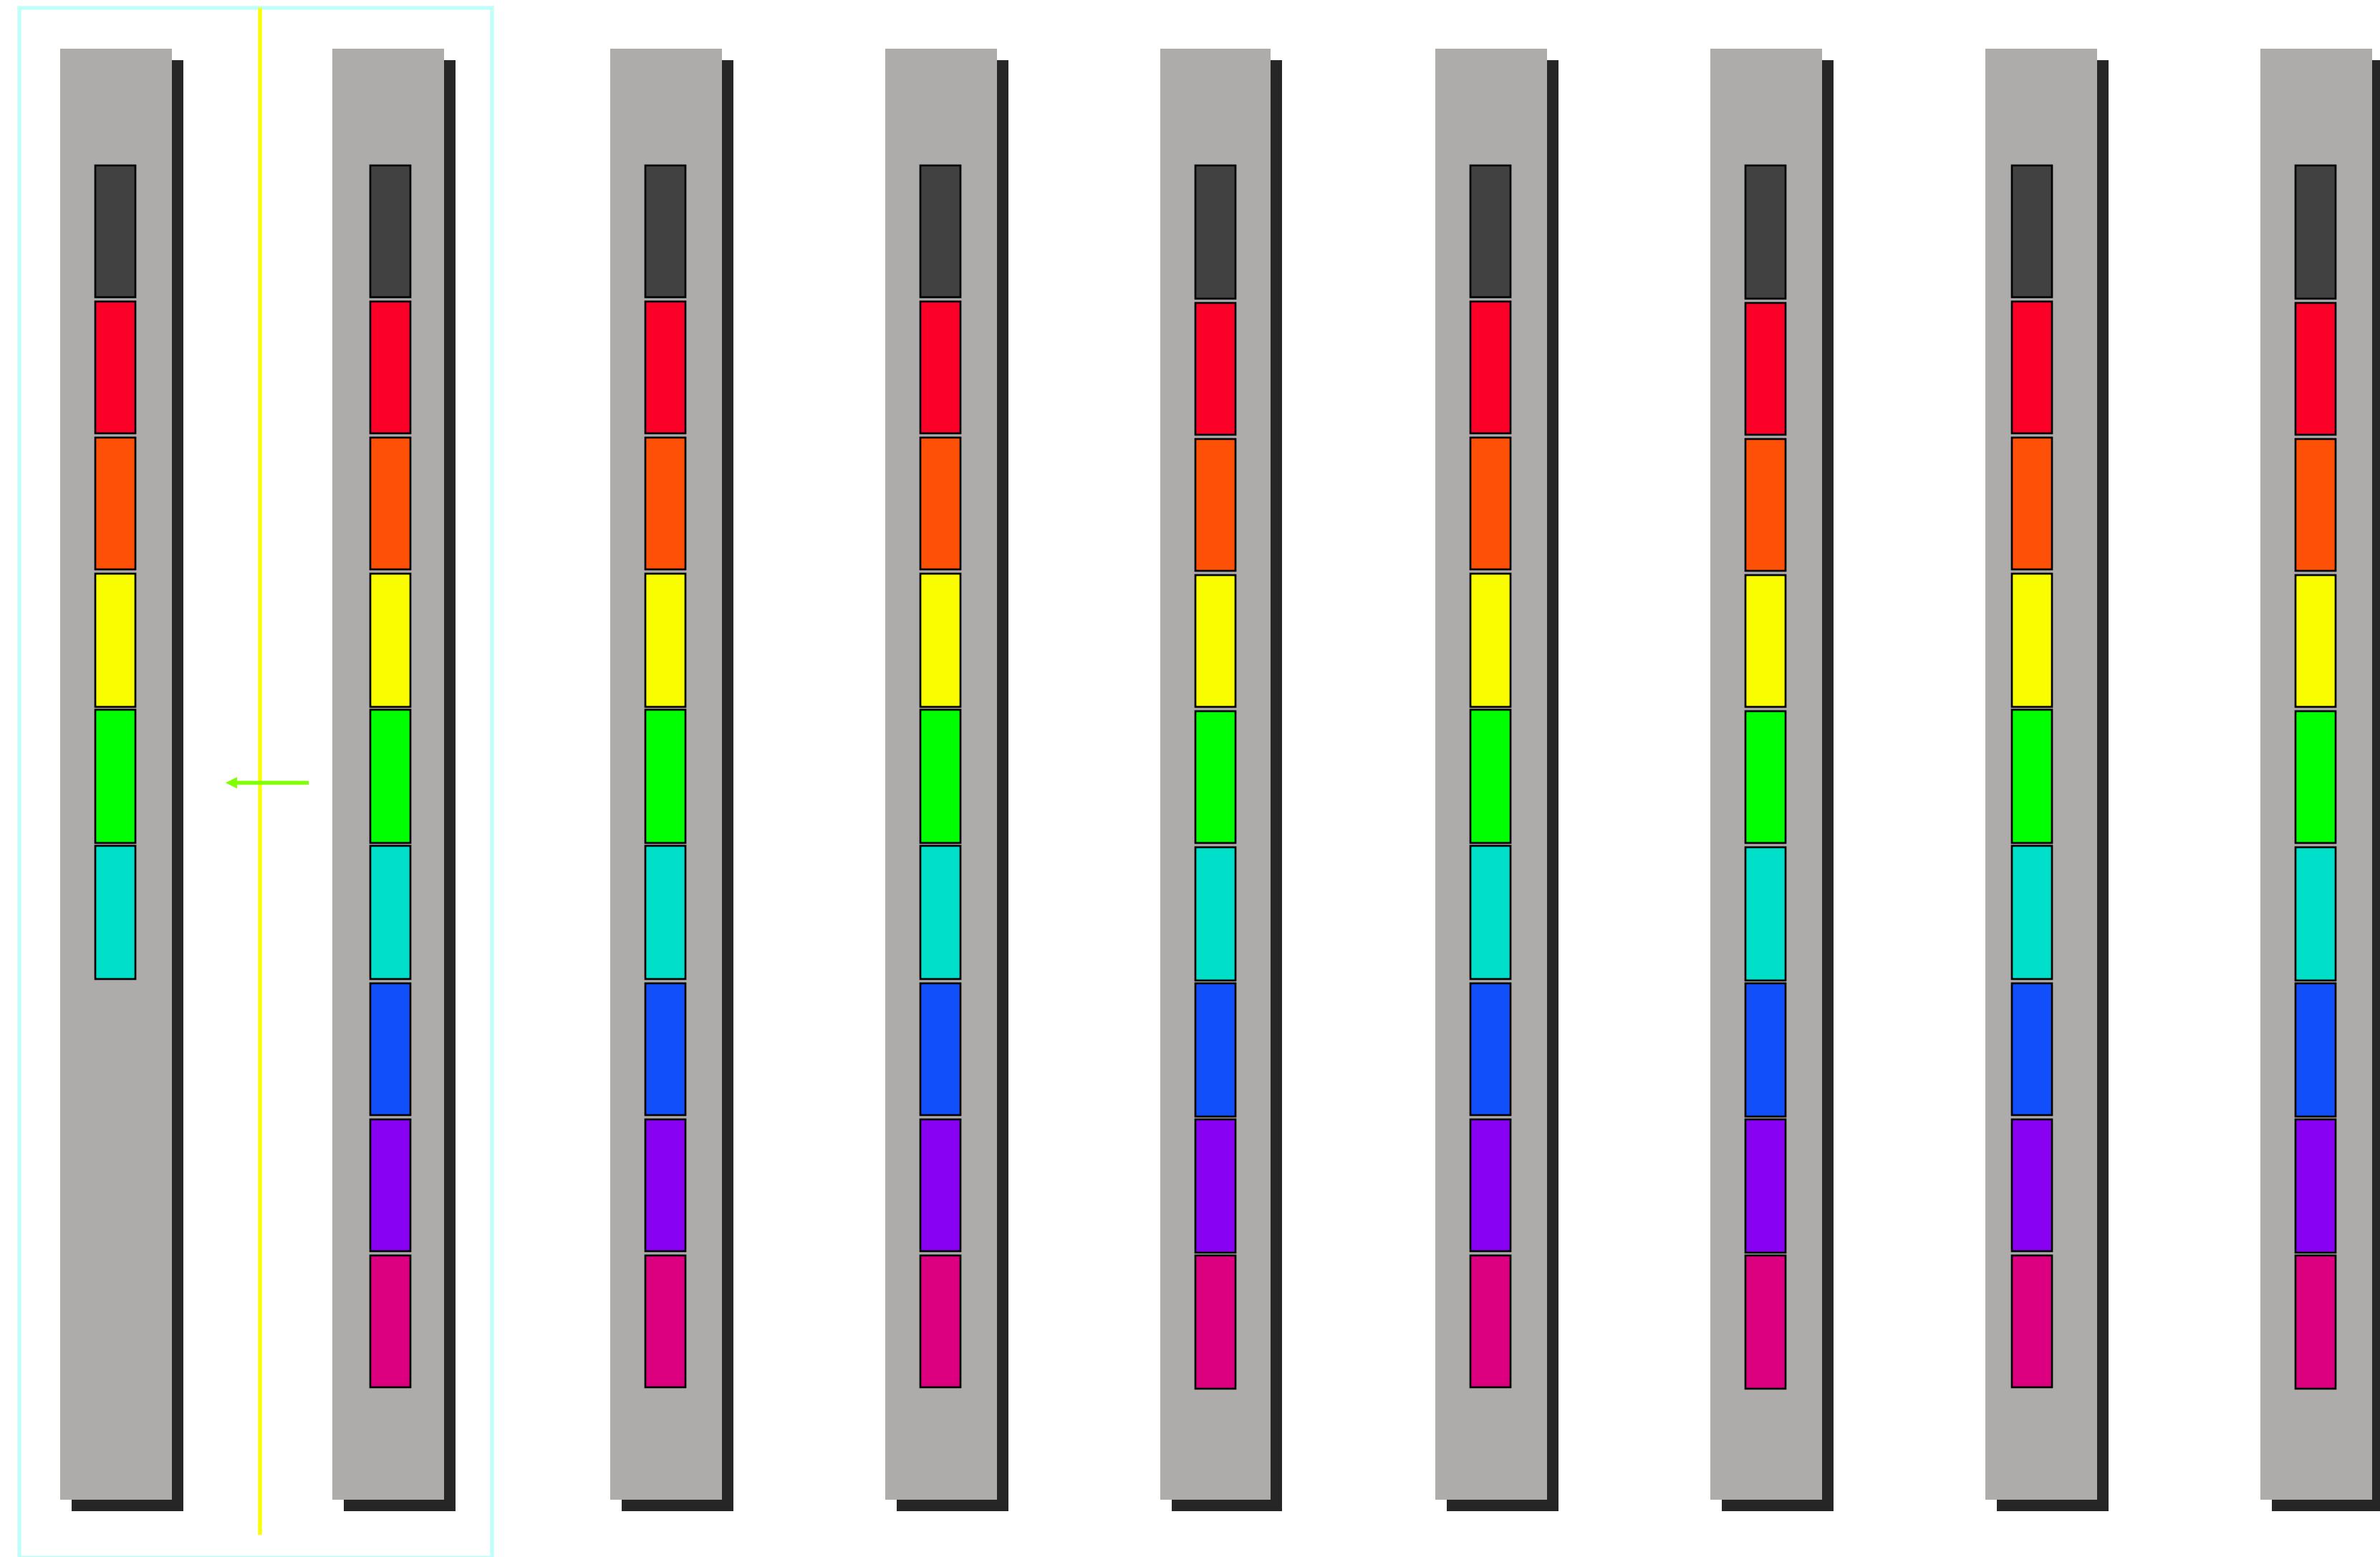


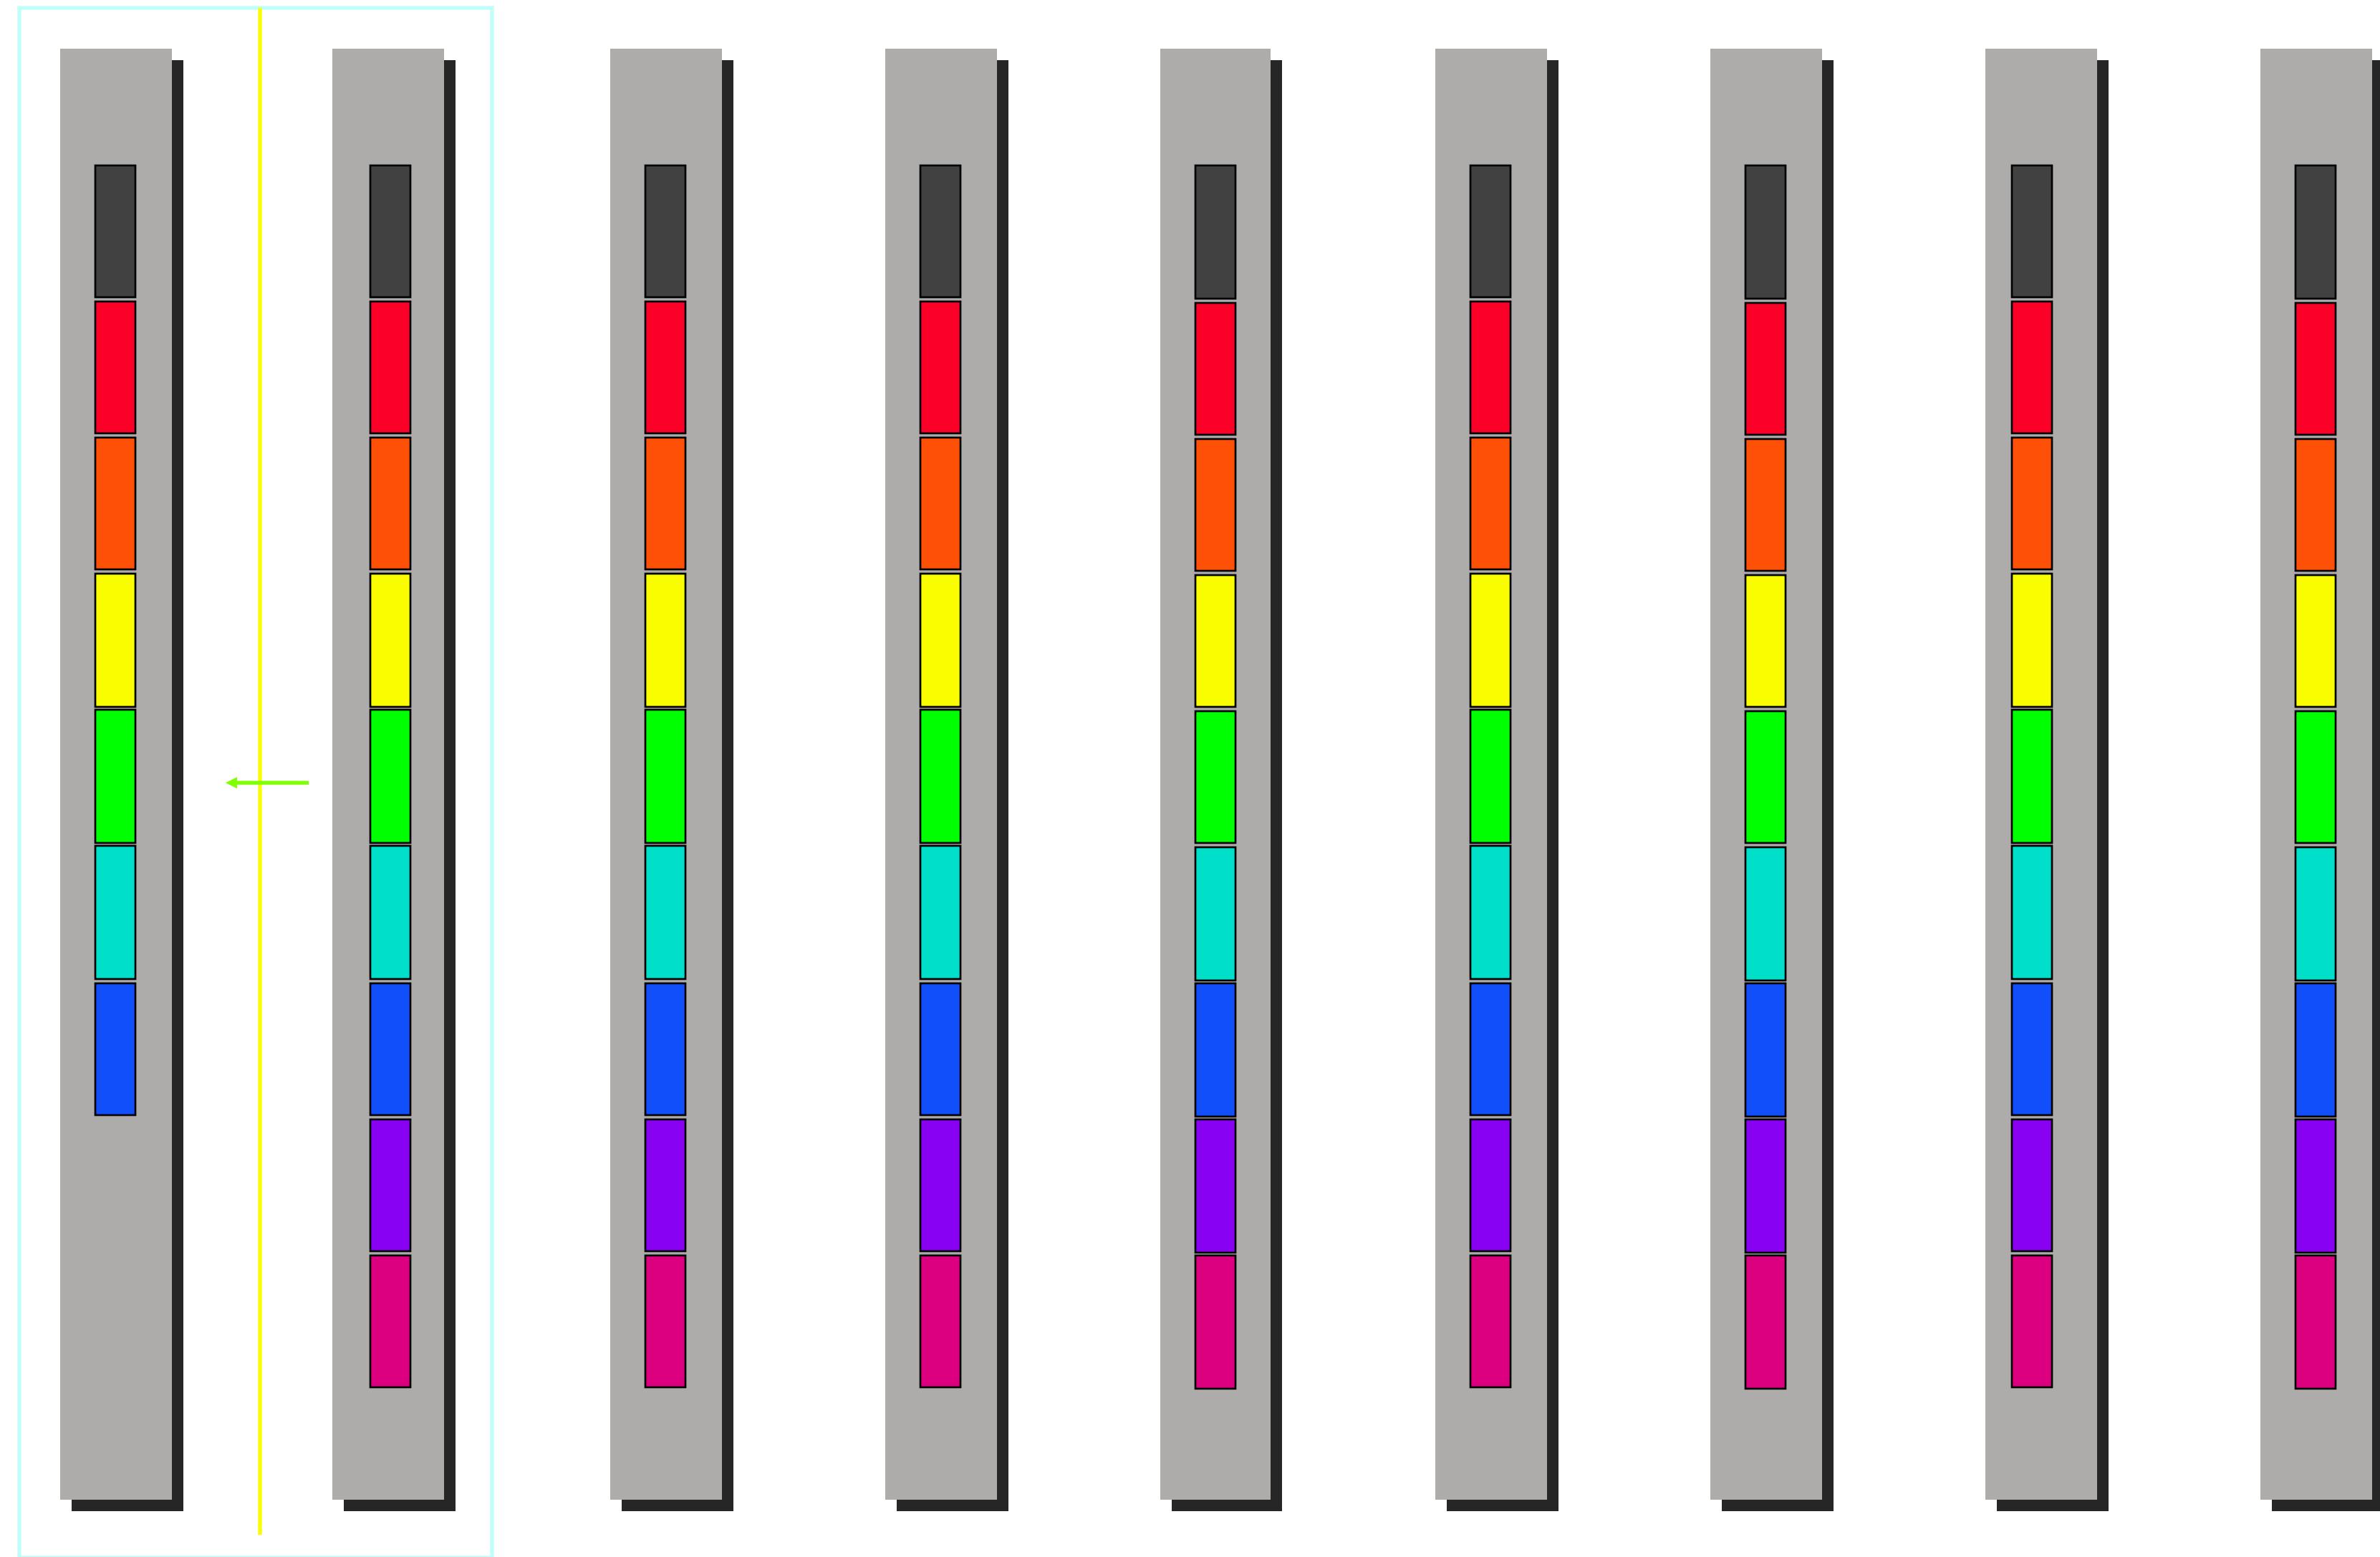


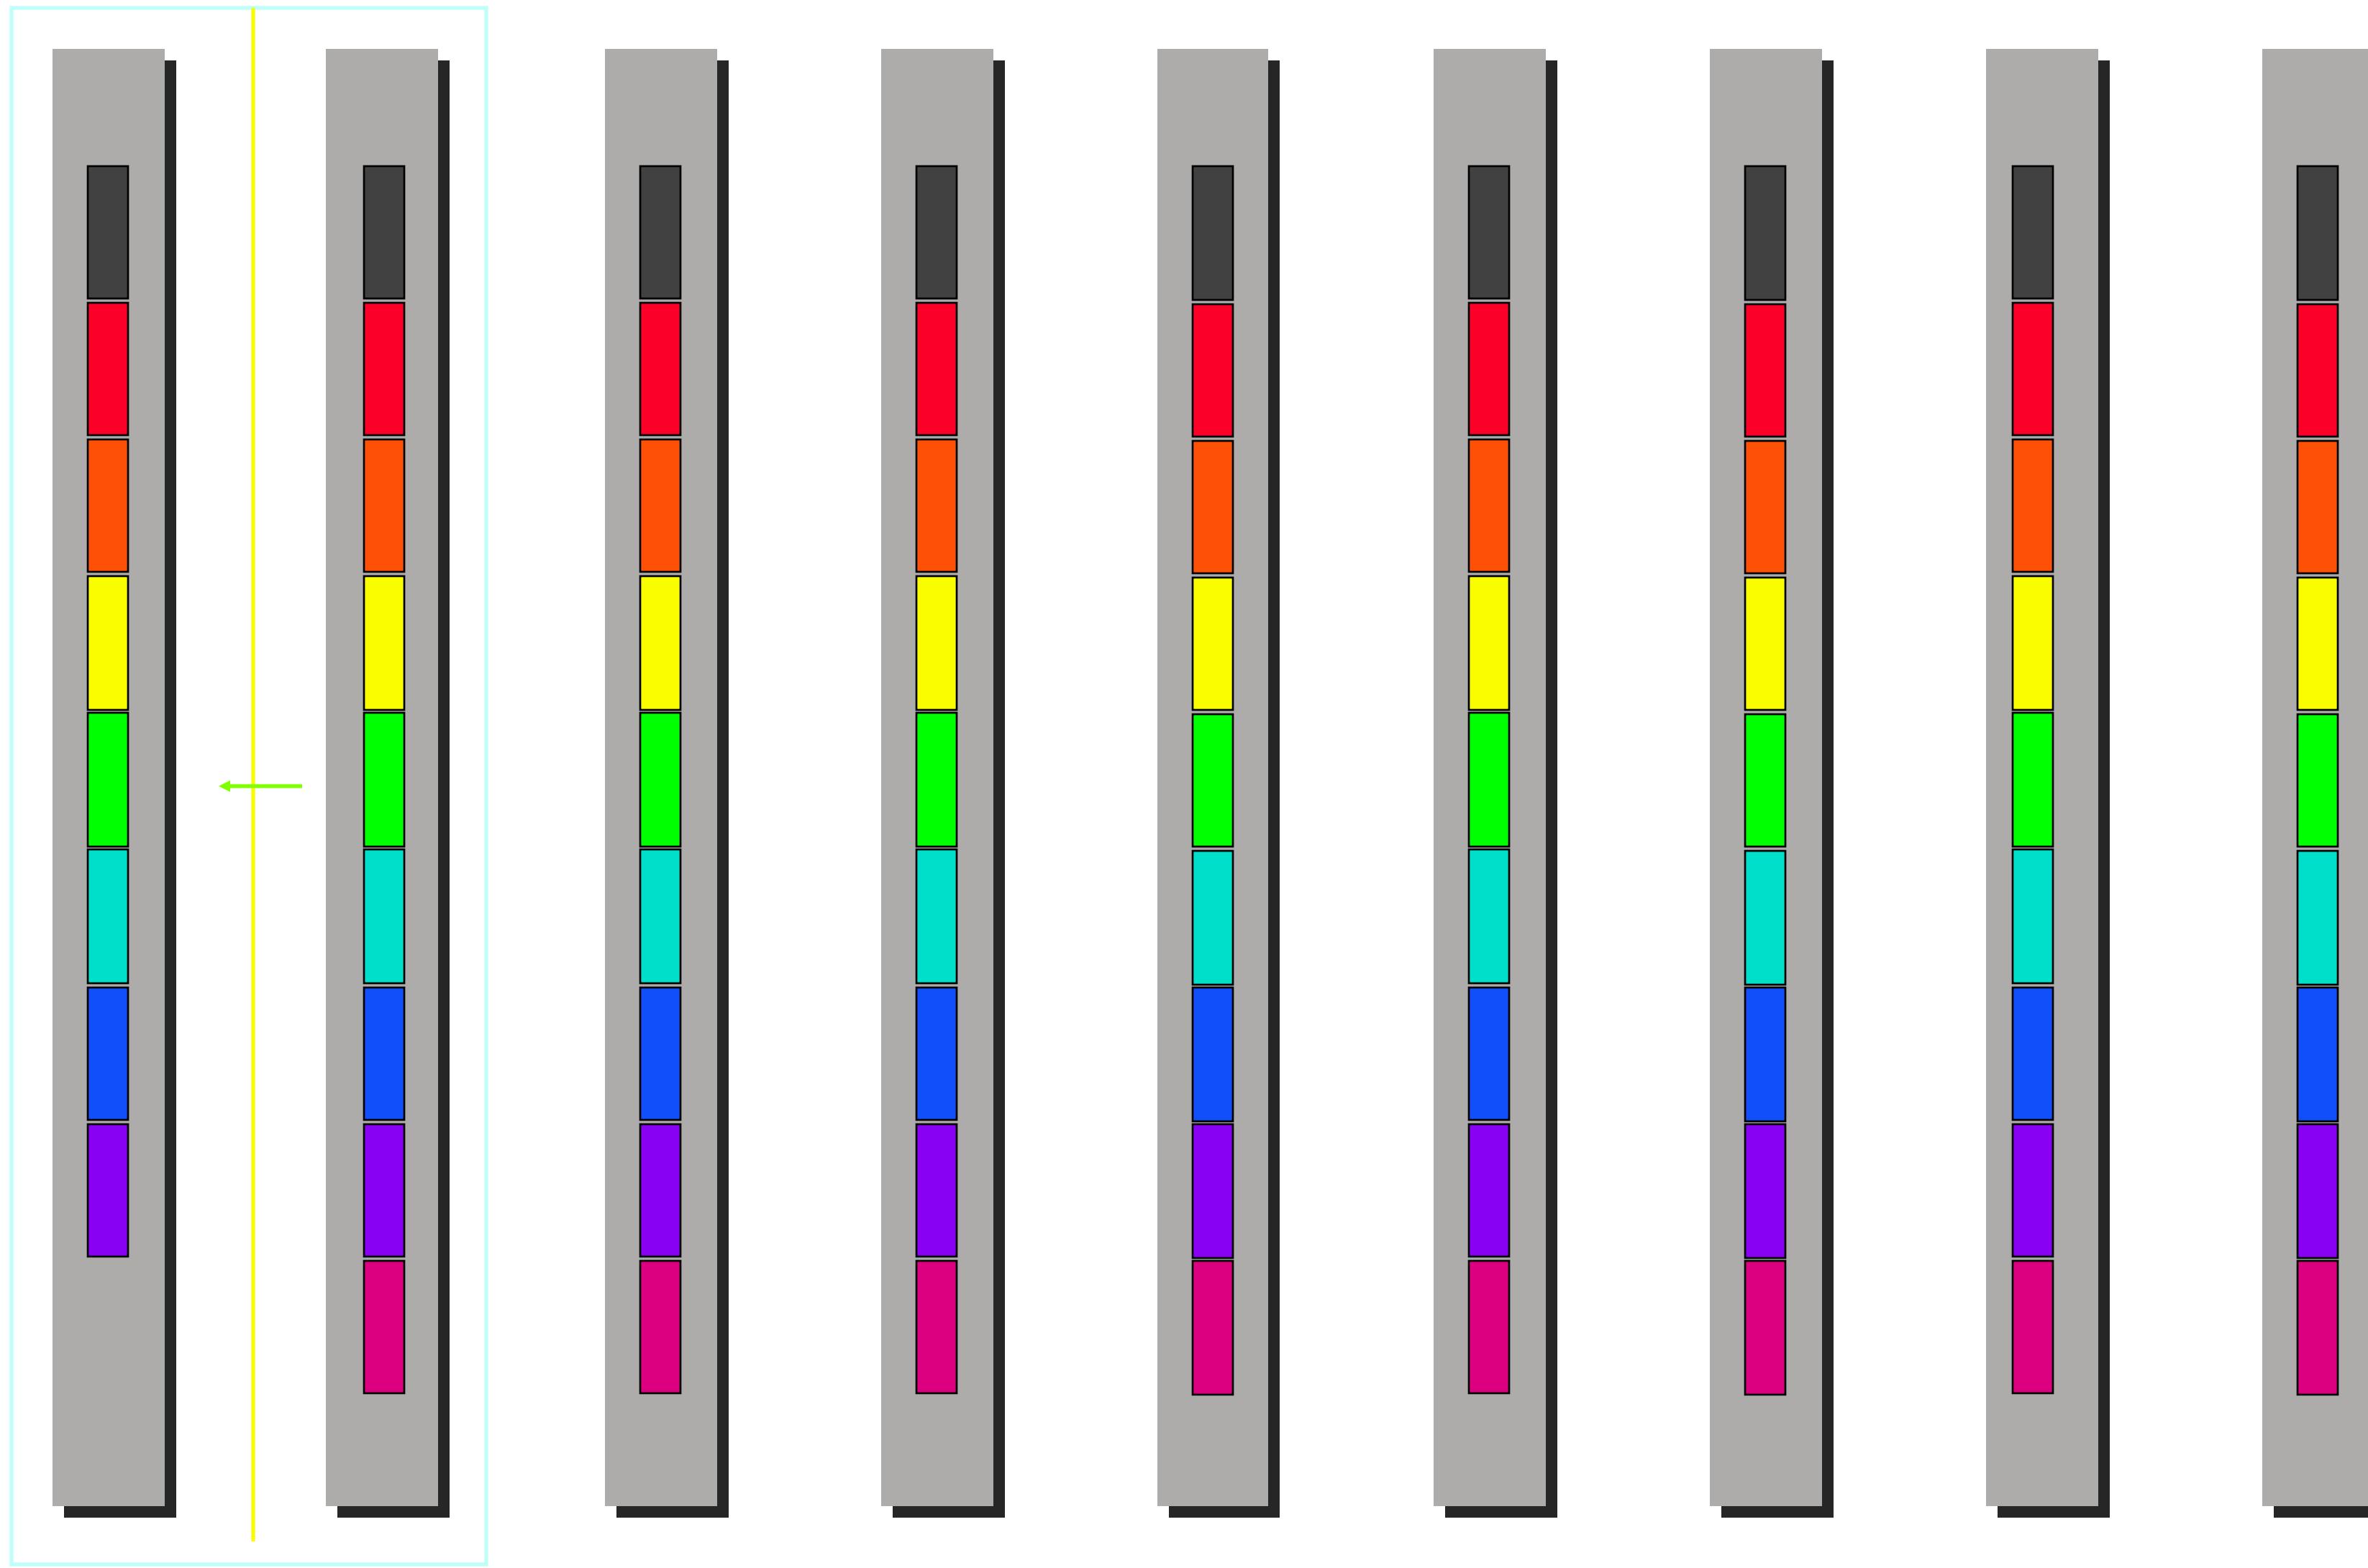


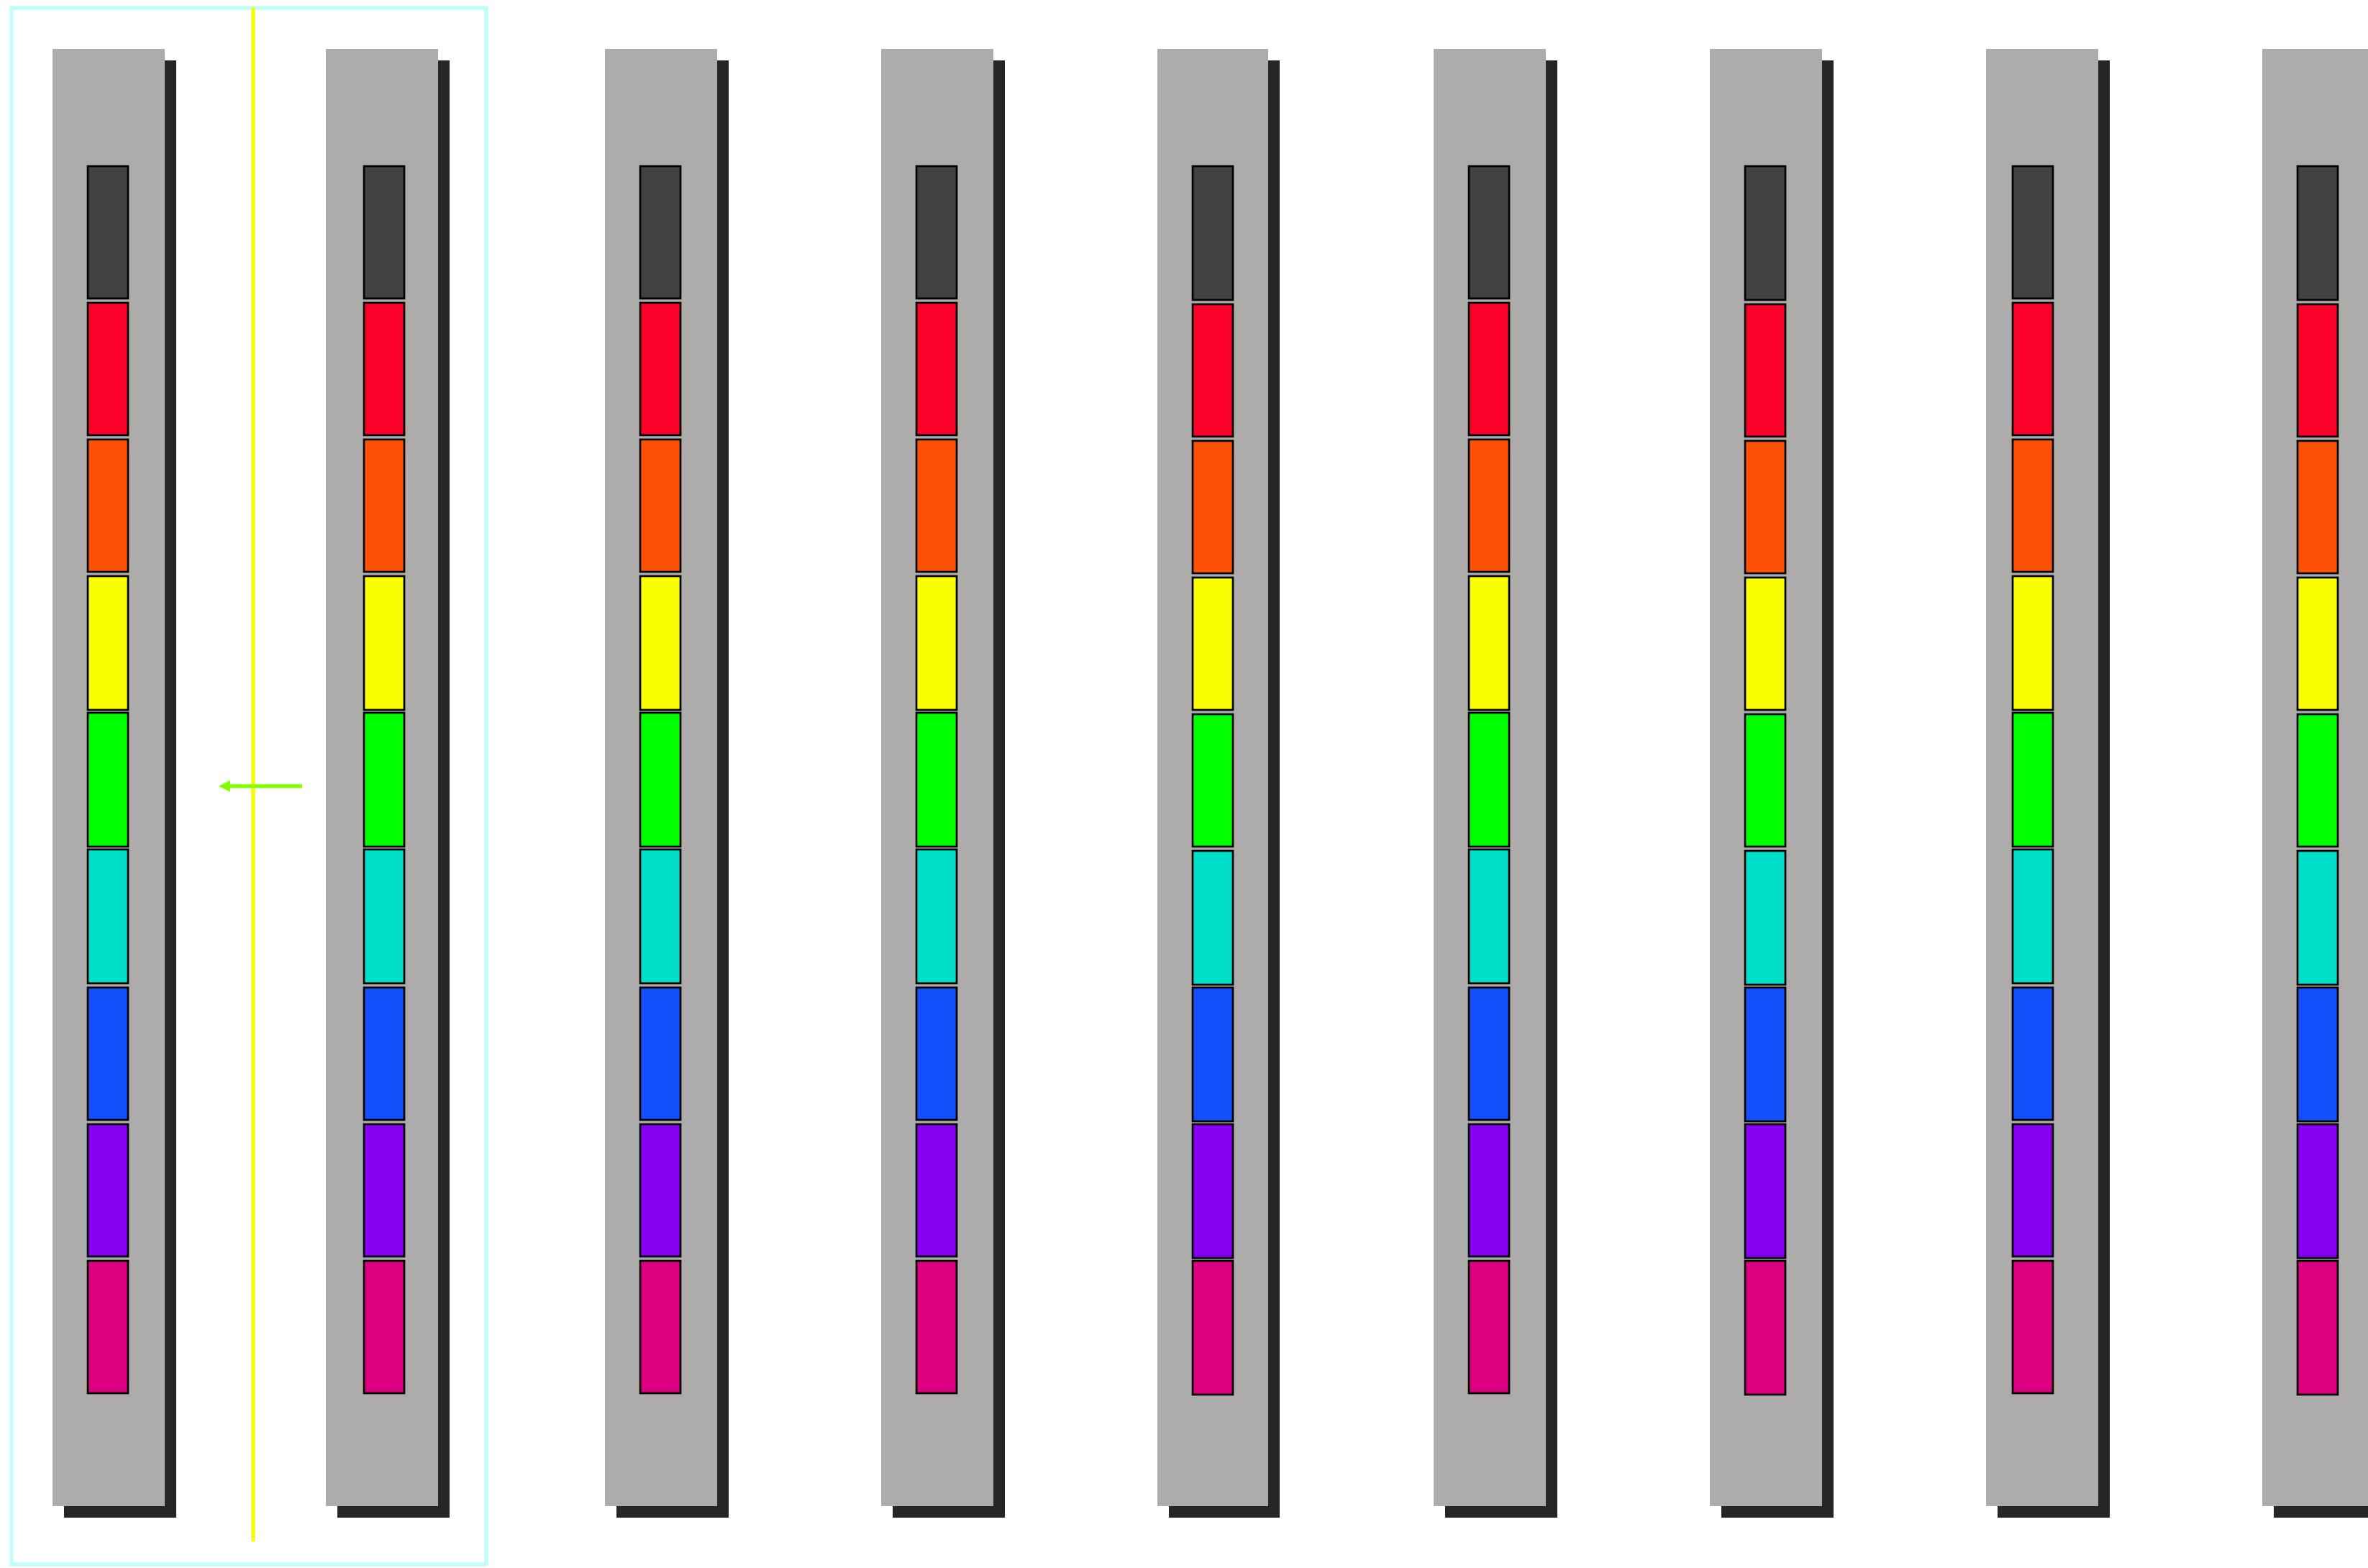


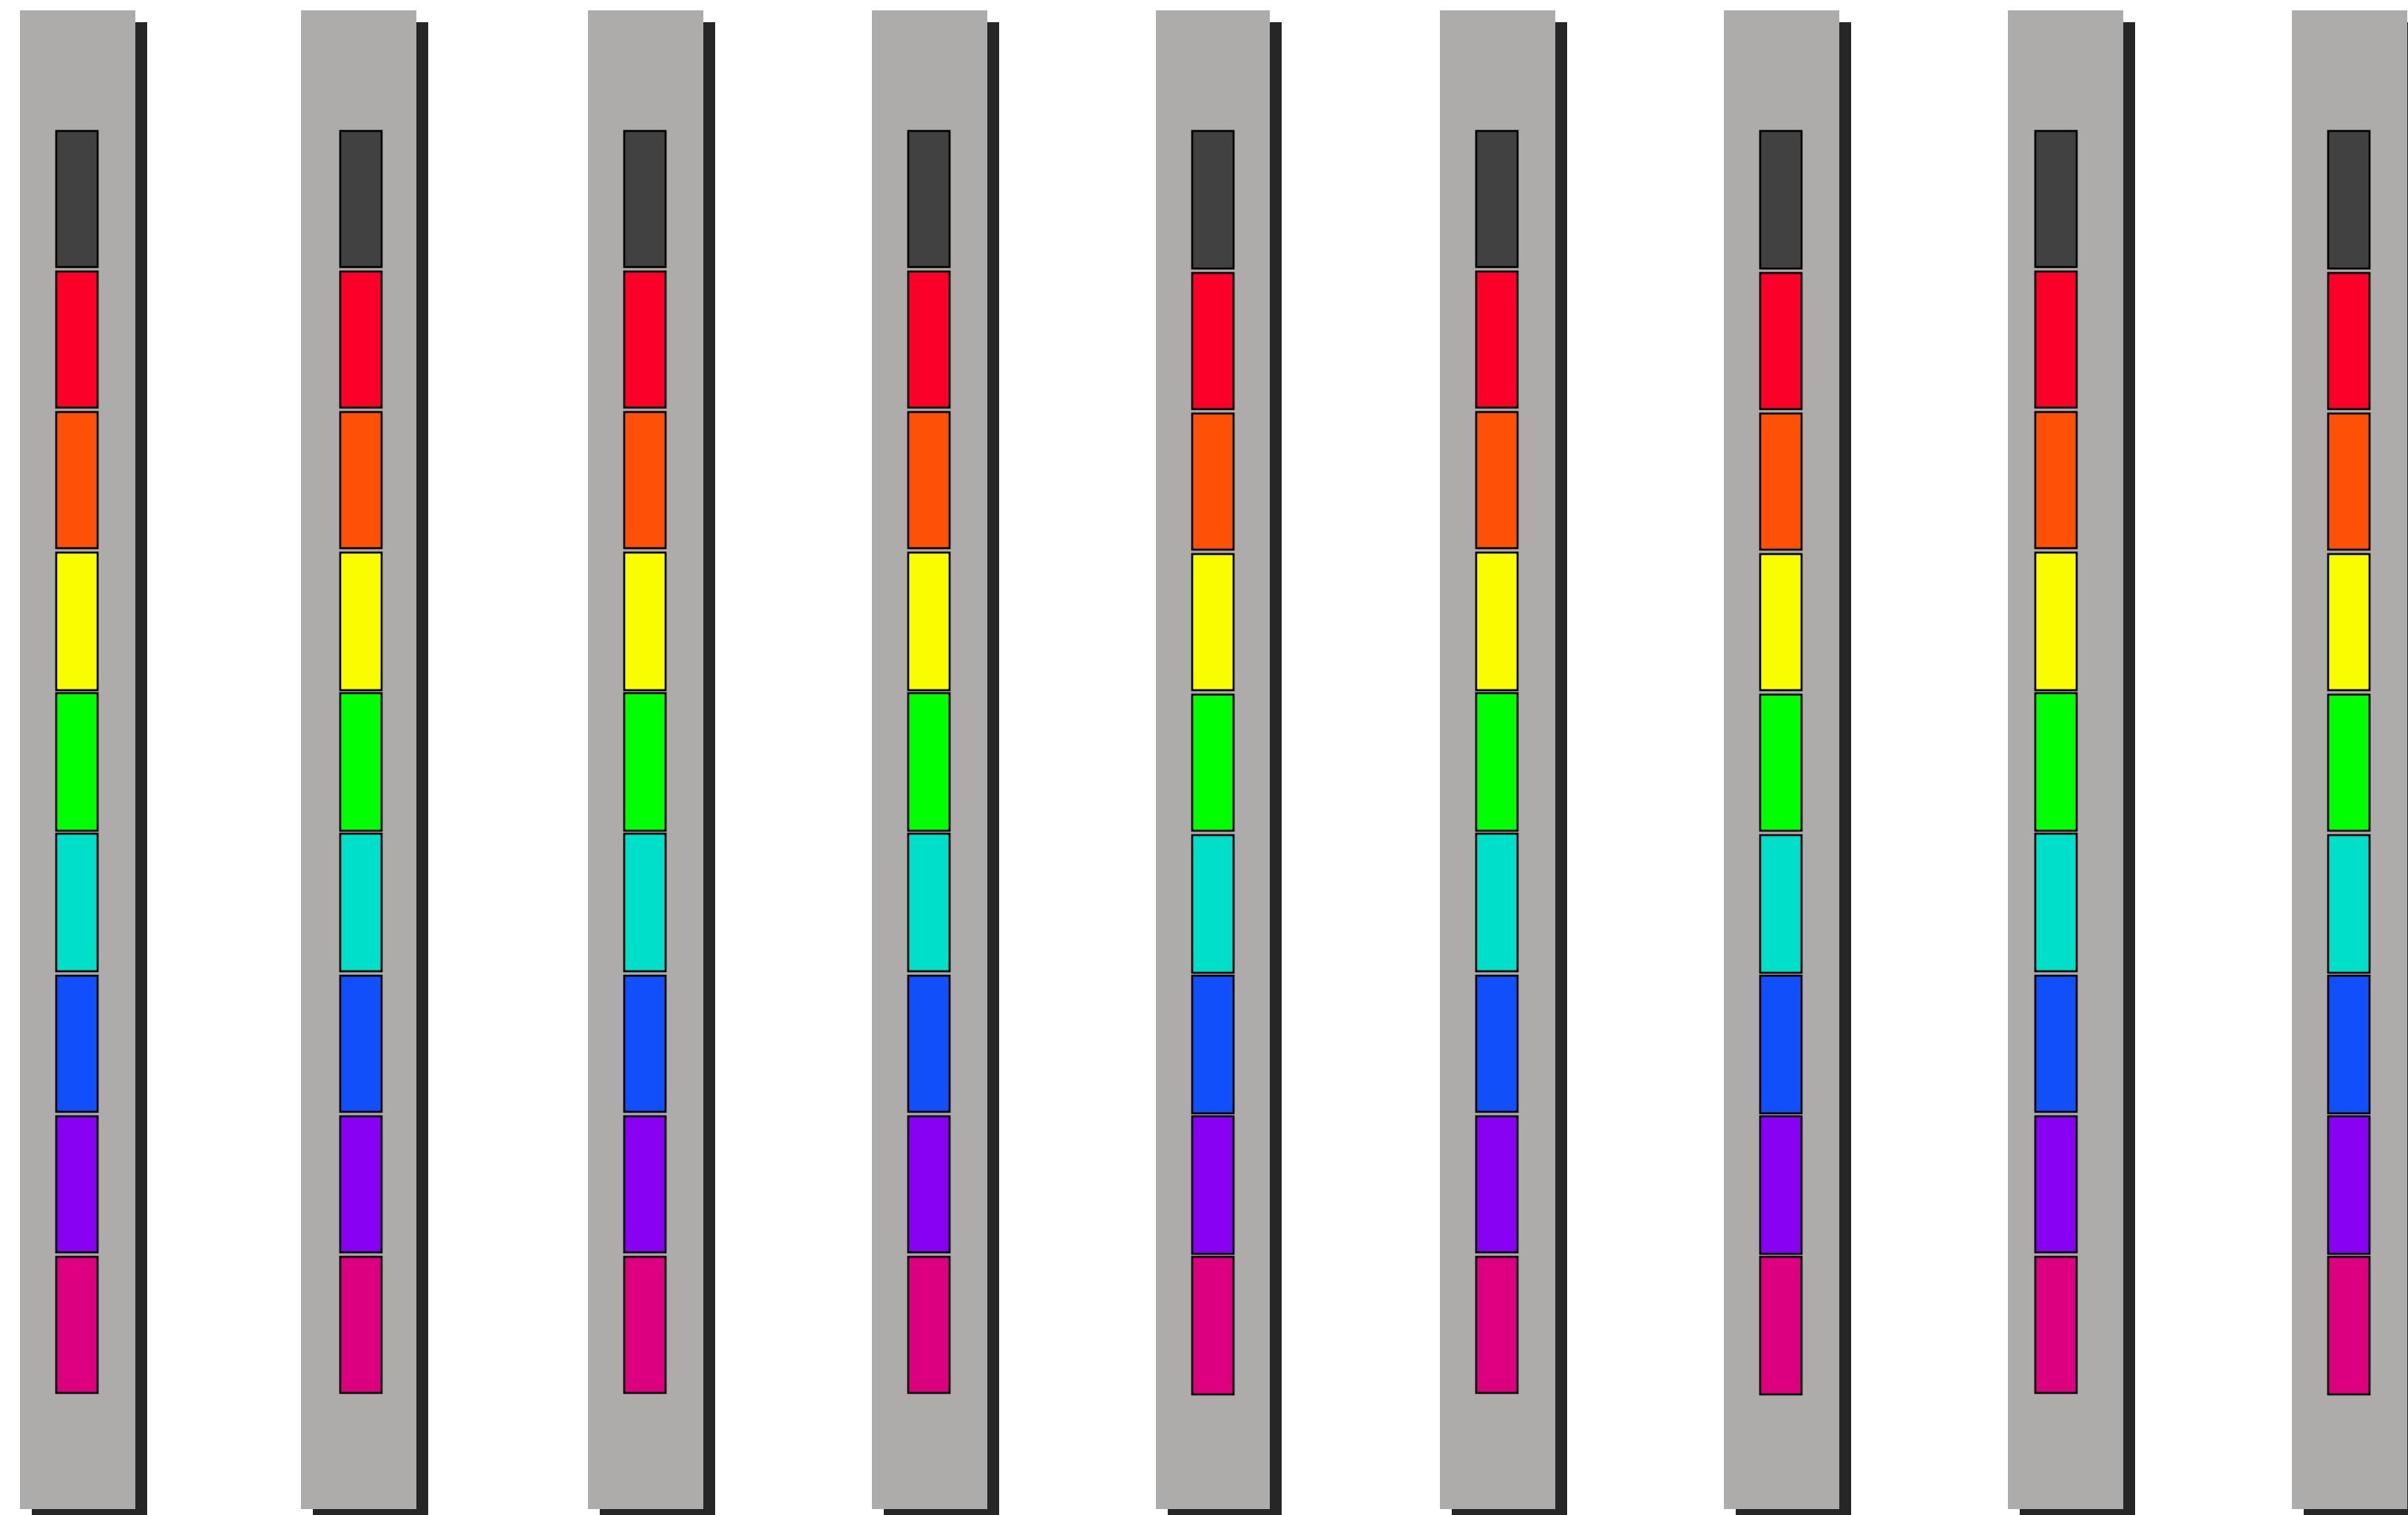








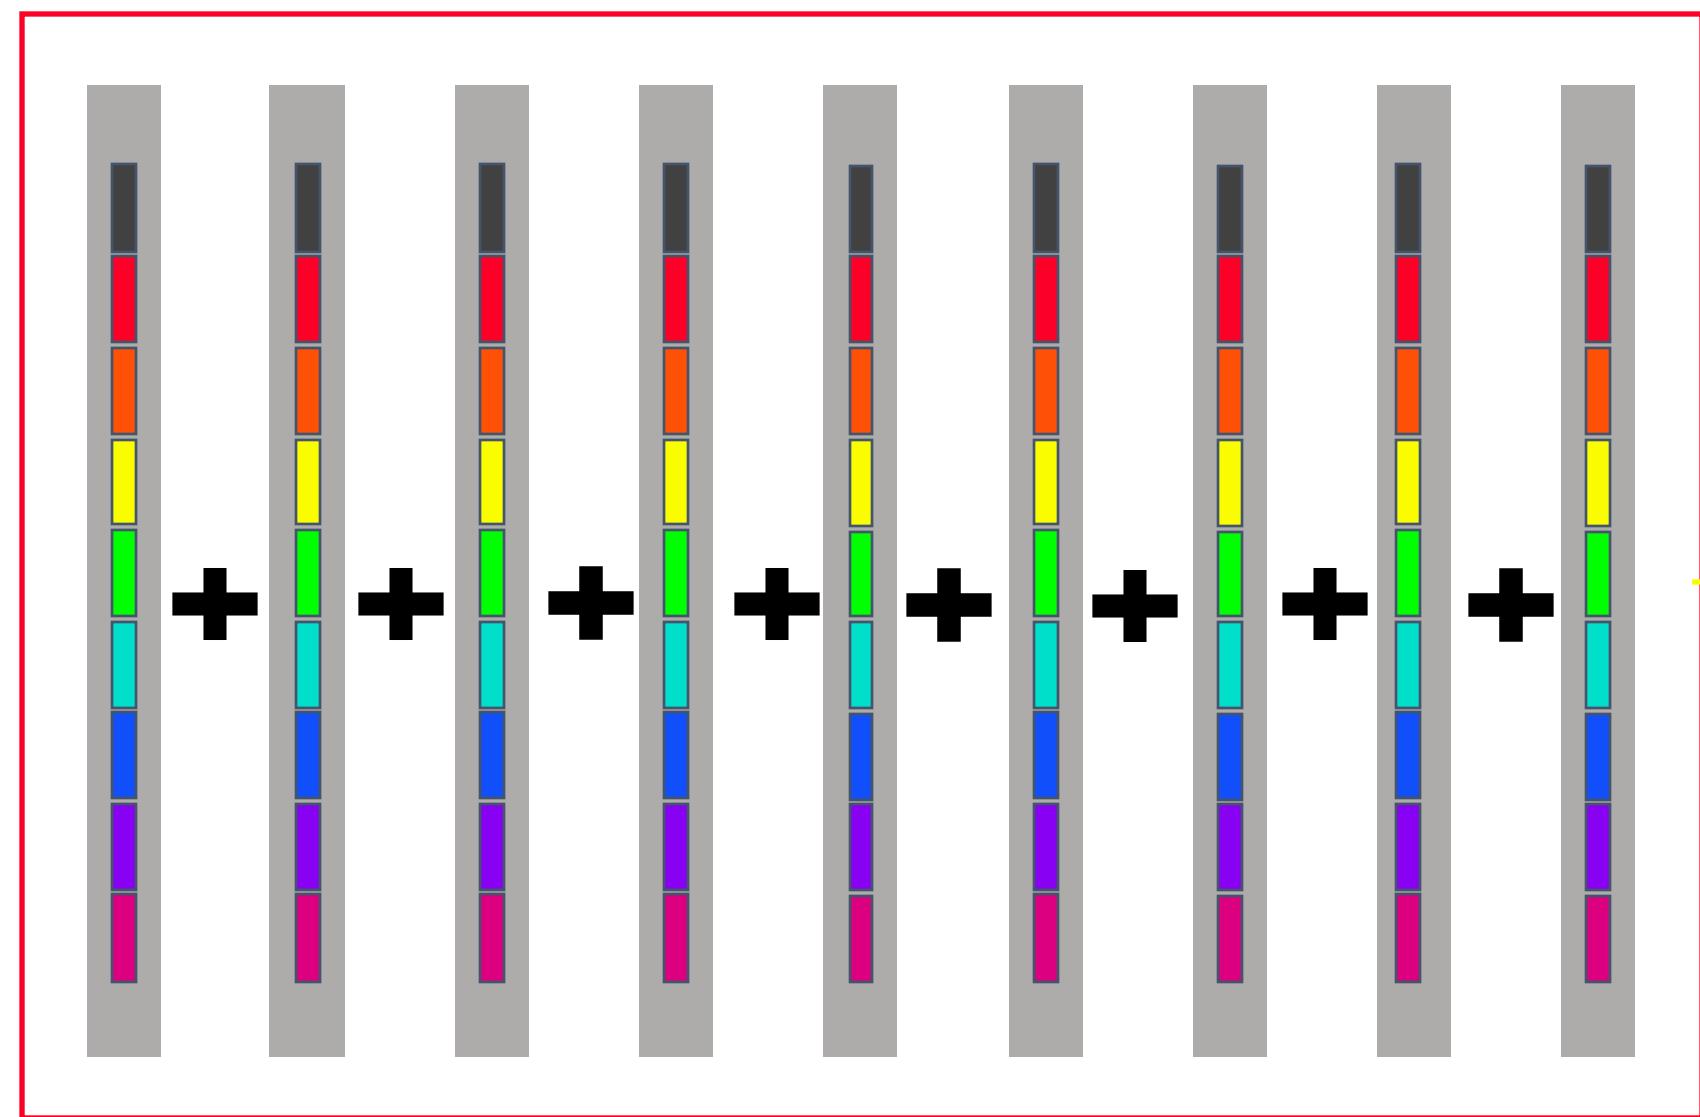




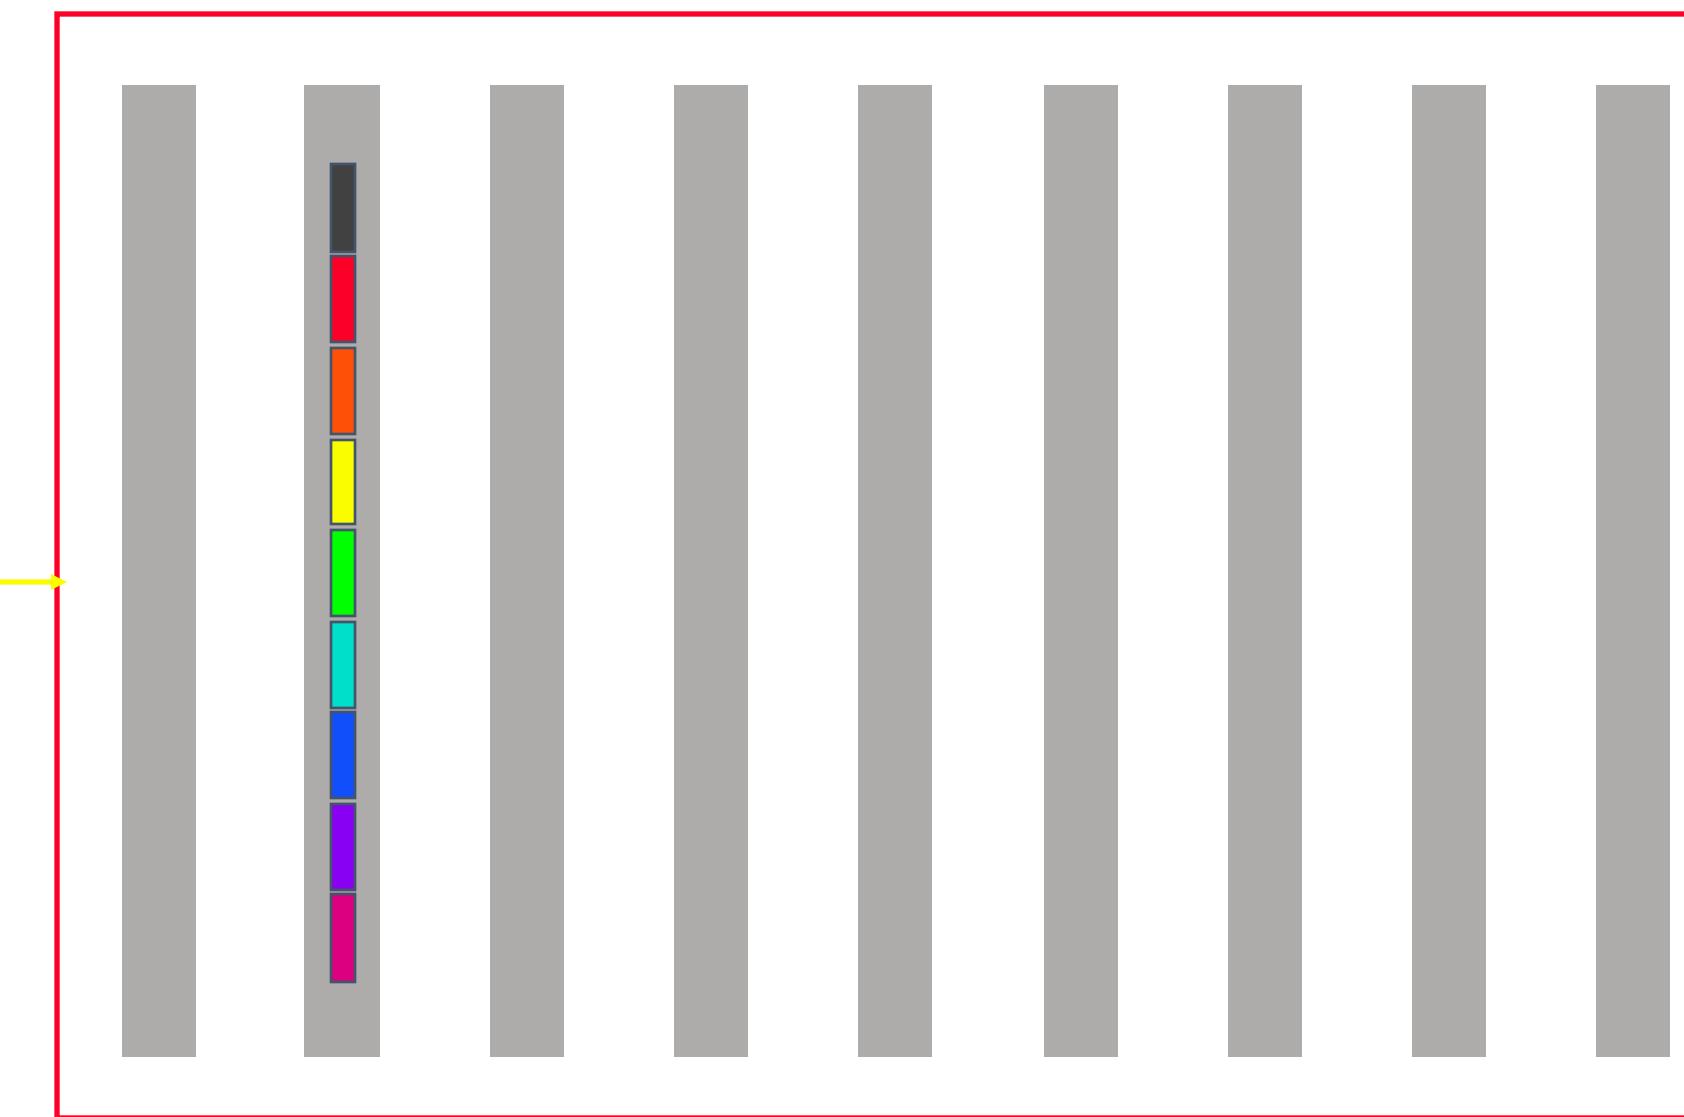
Cost of minimum spanning tree broadcast

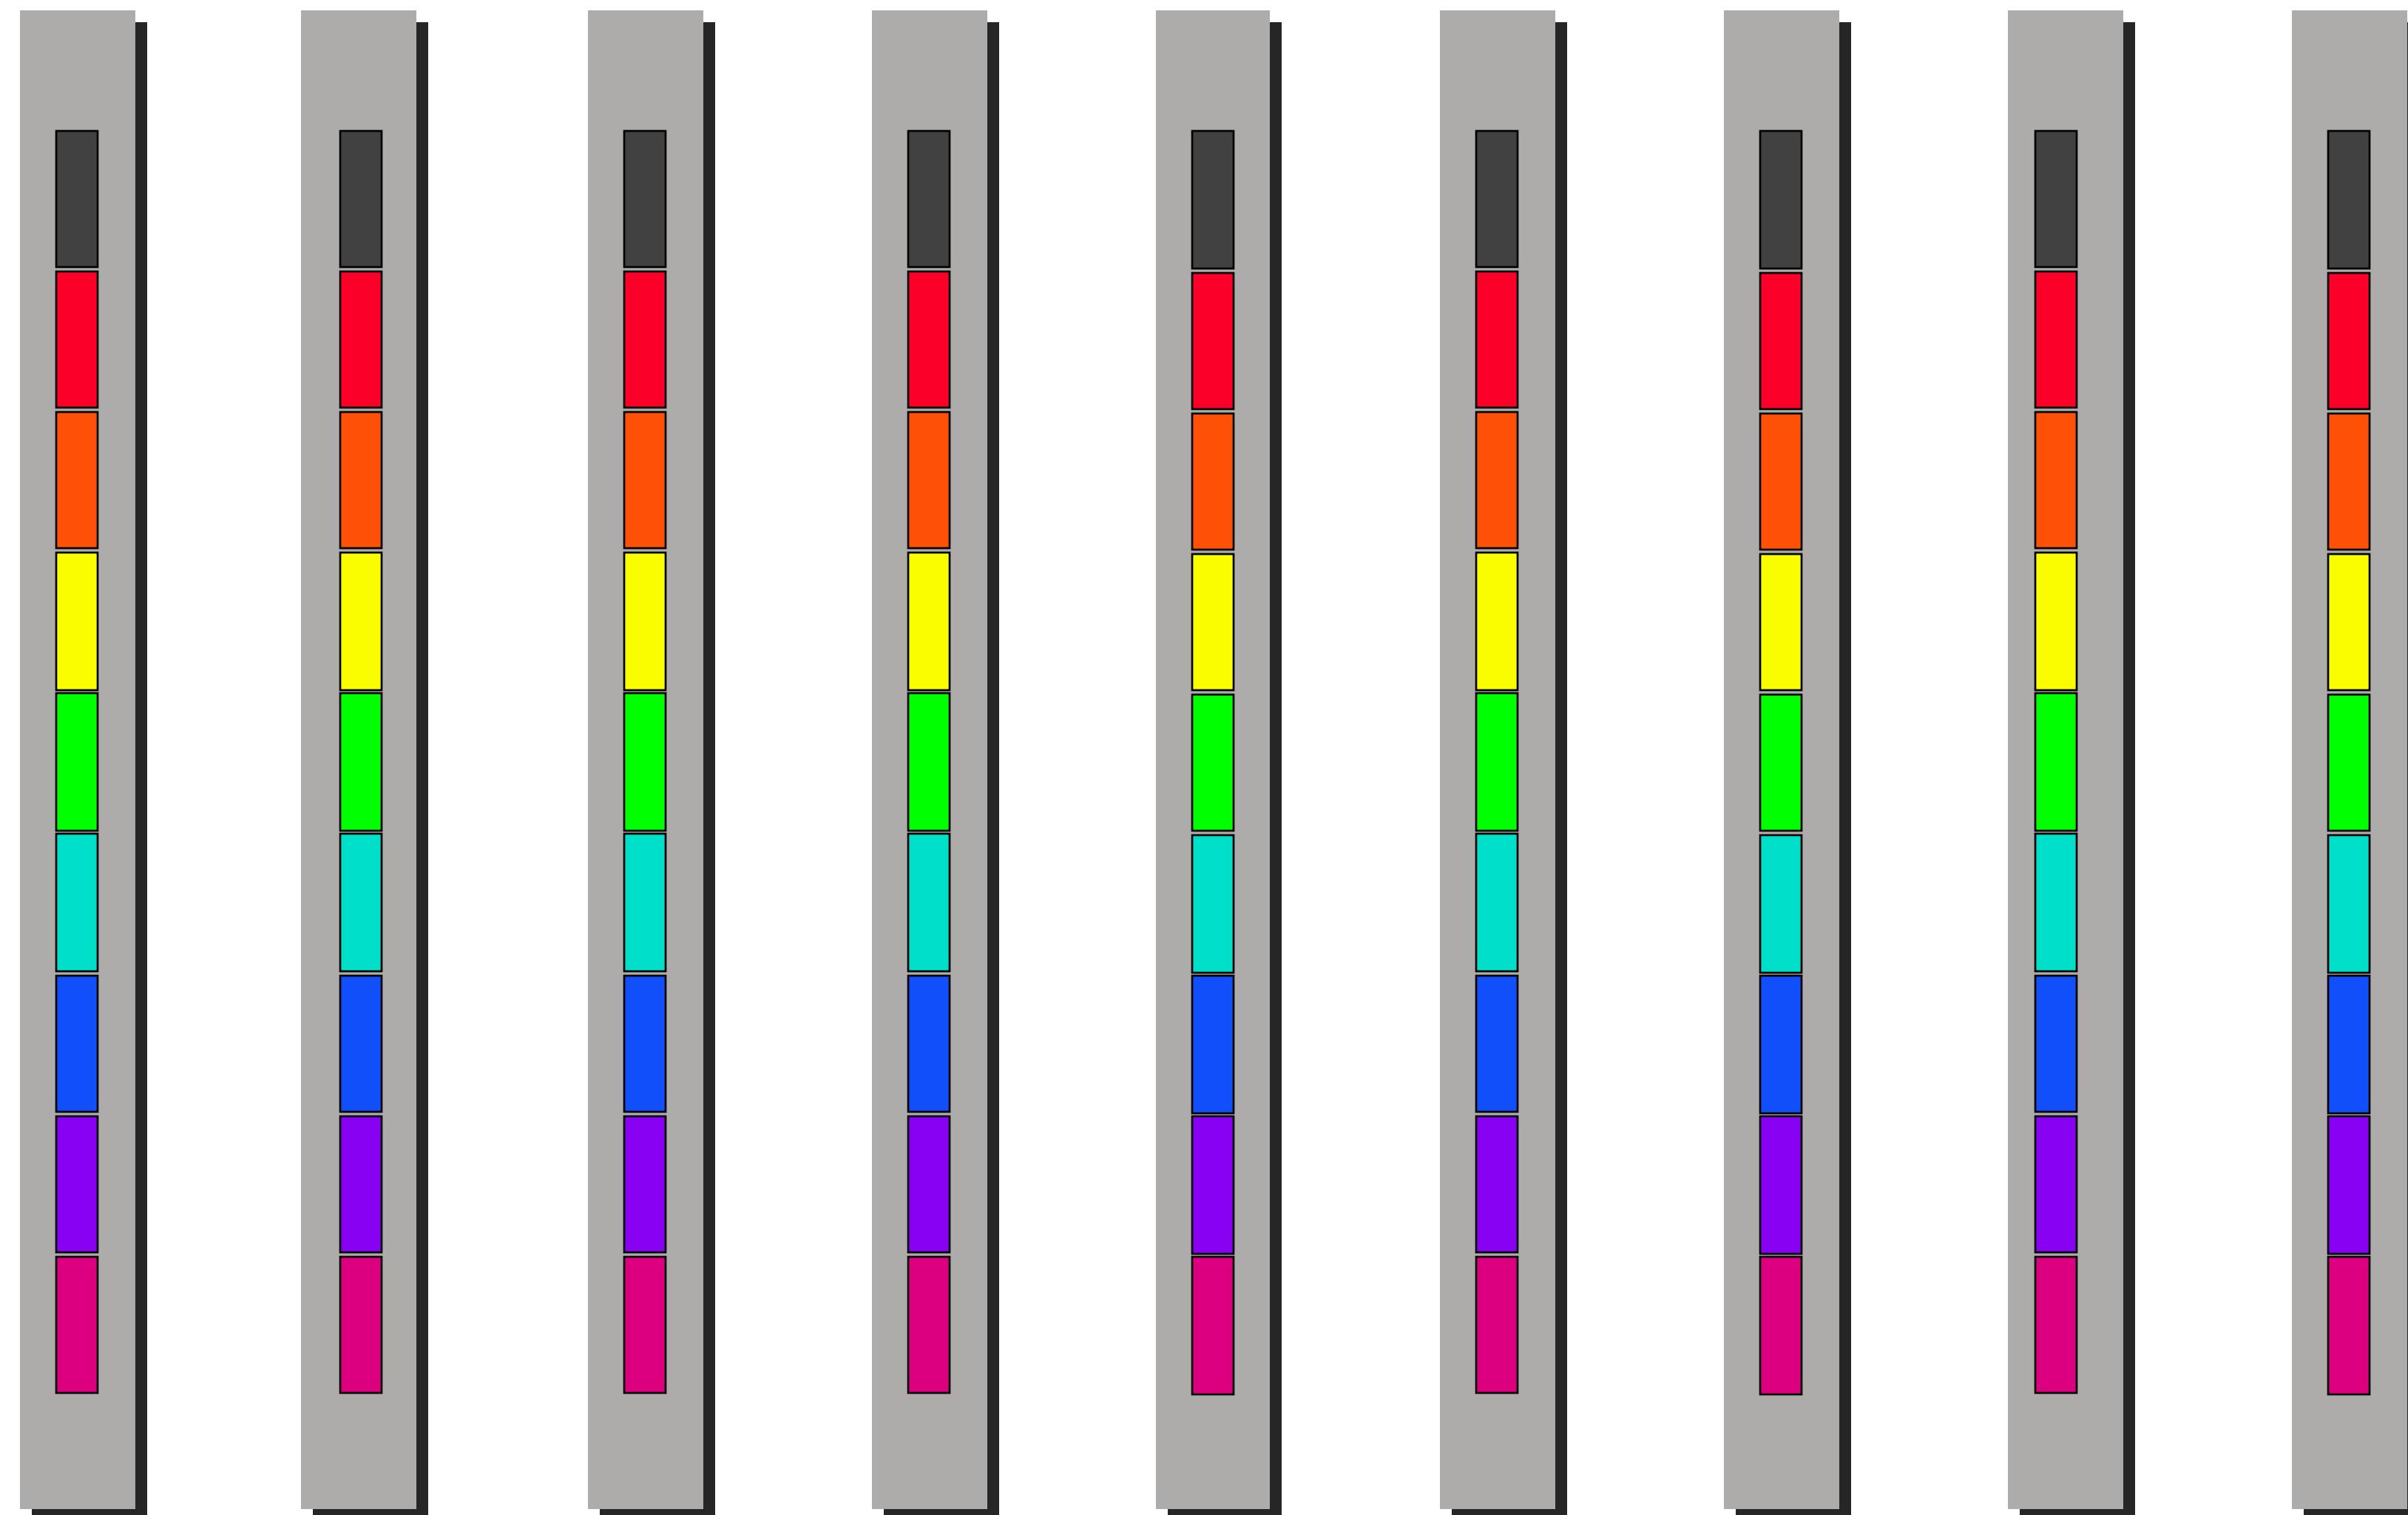
Reduce(-to-one)

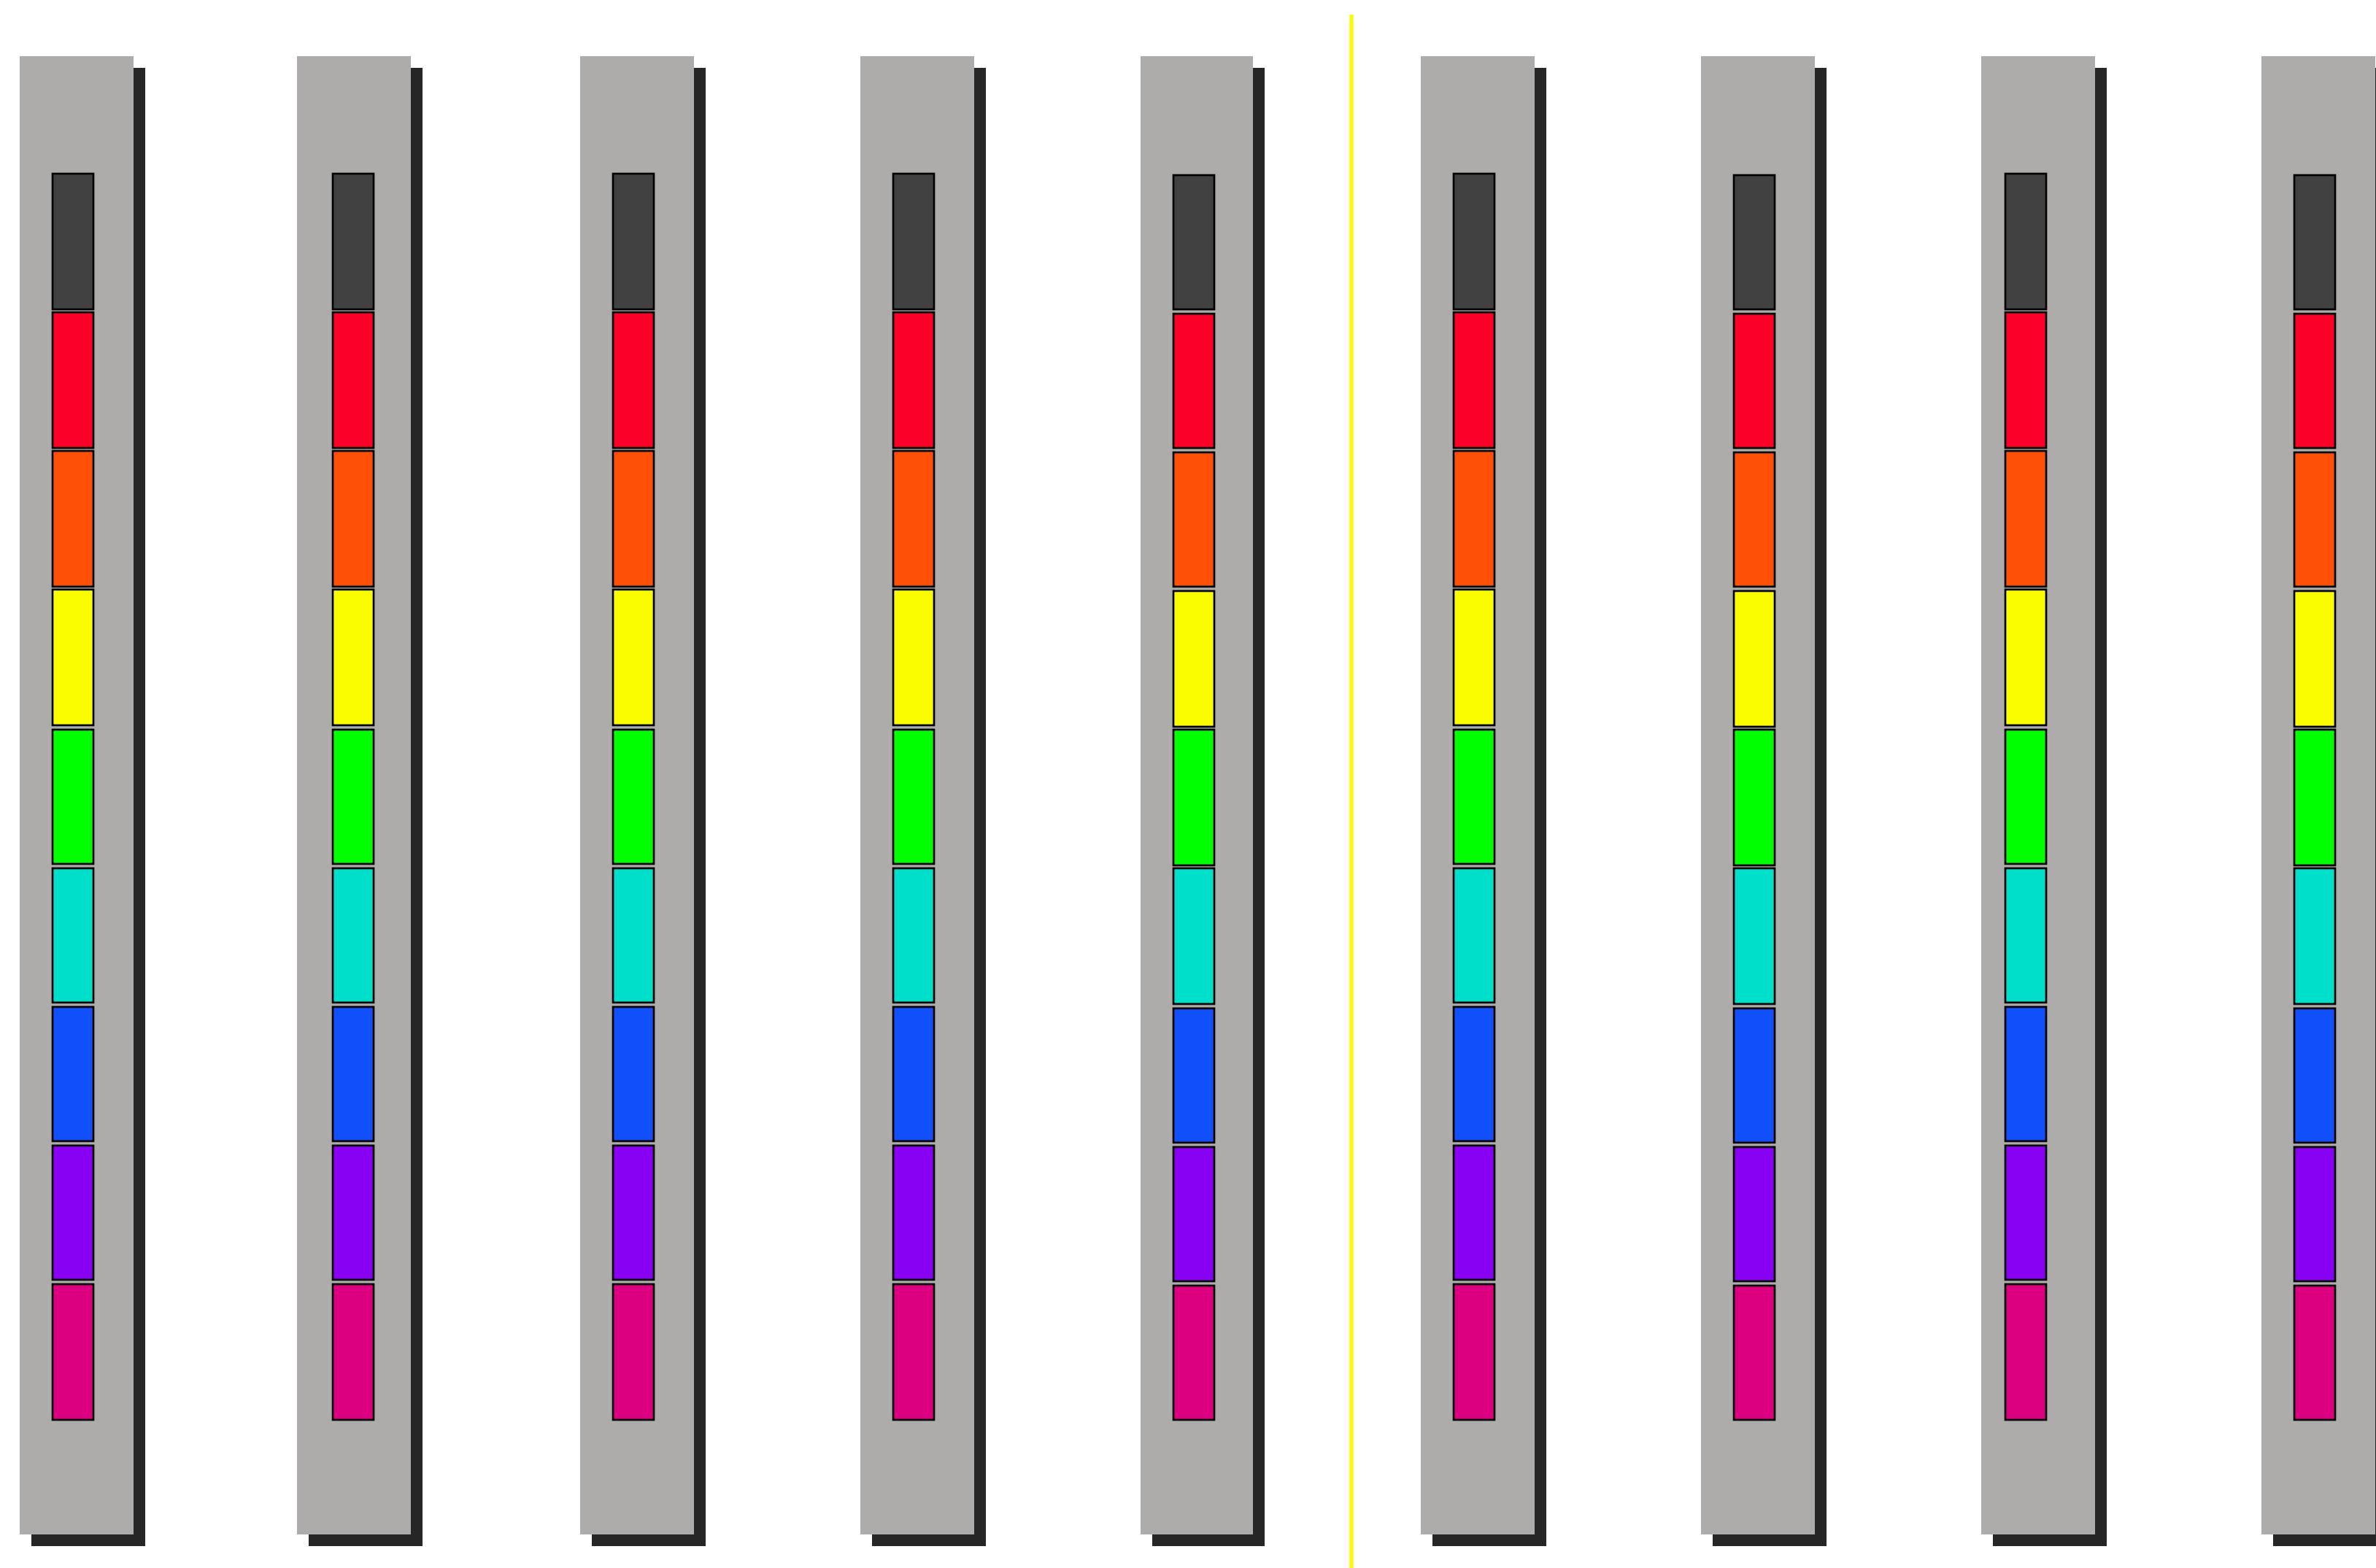
Before

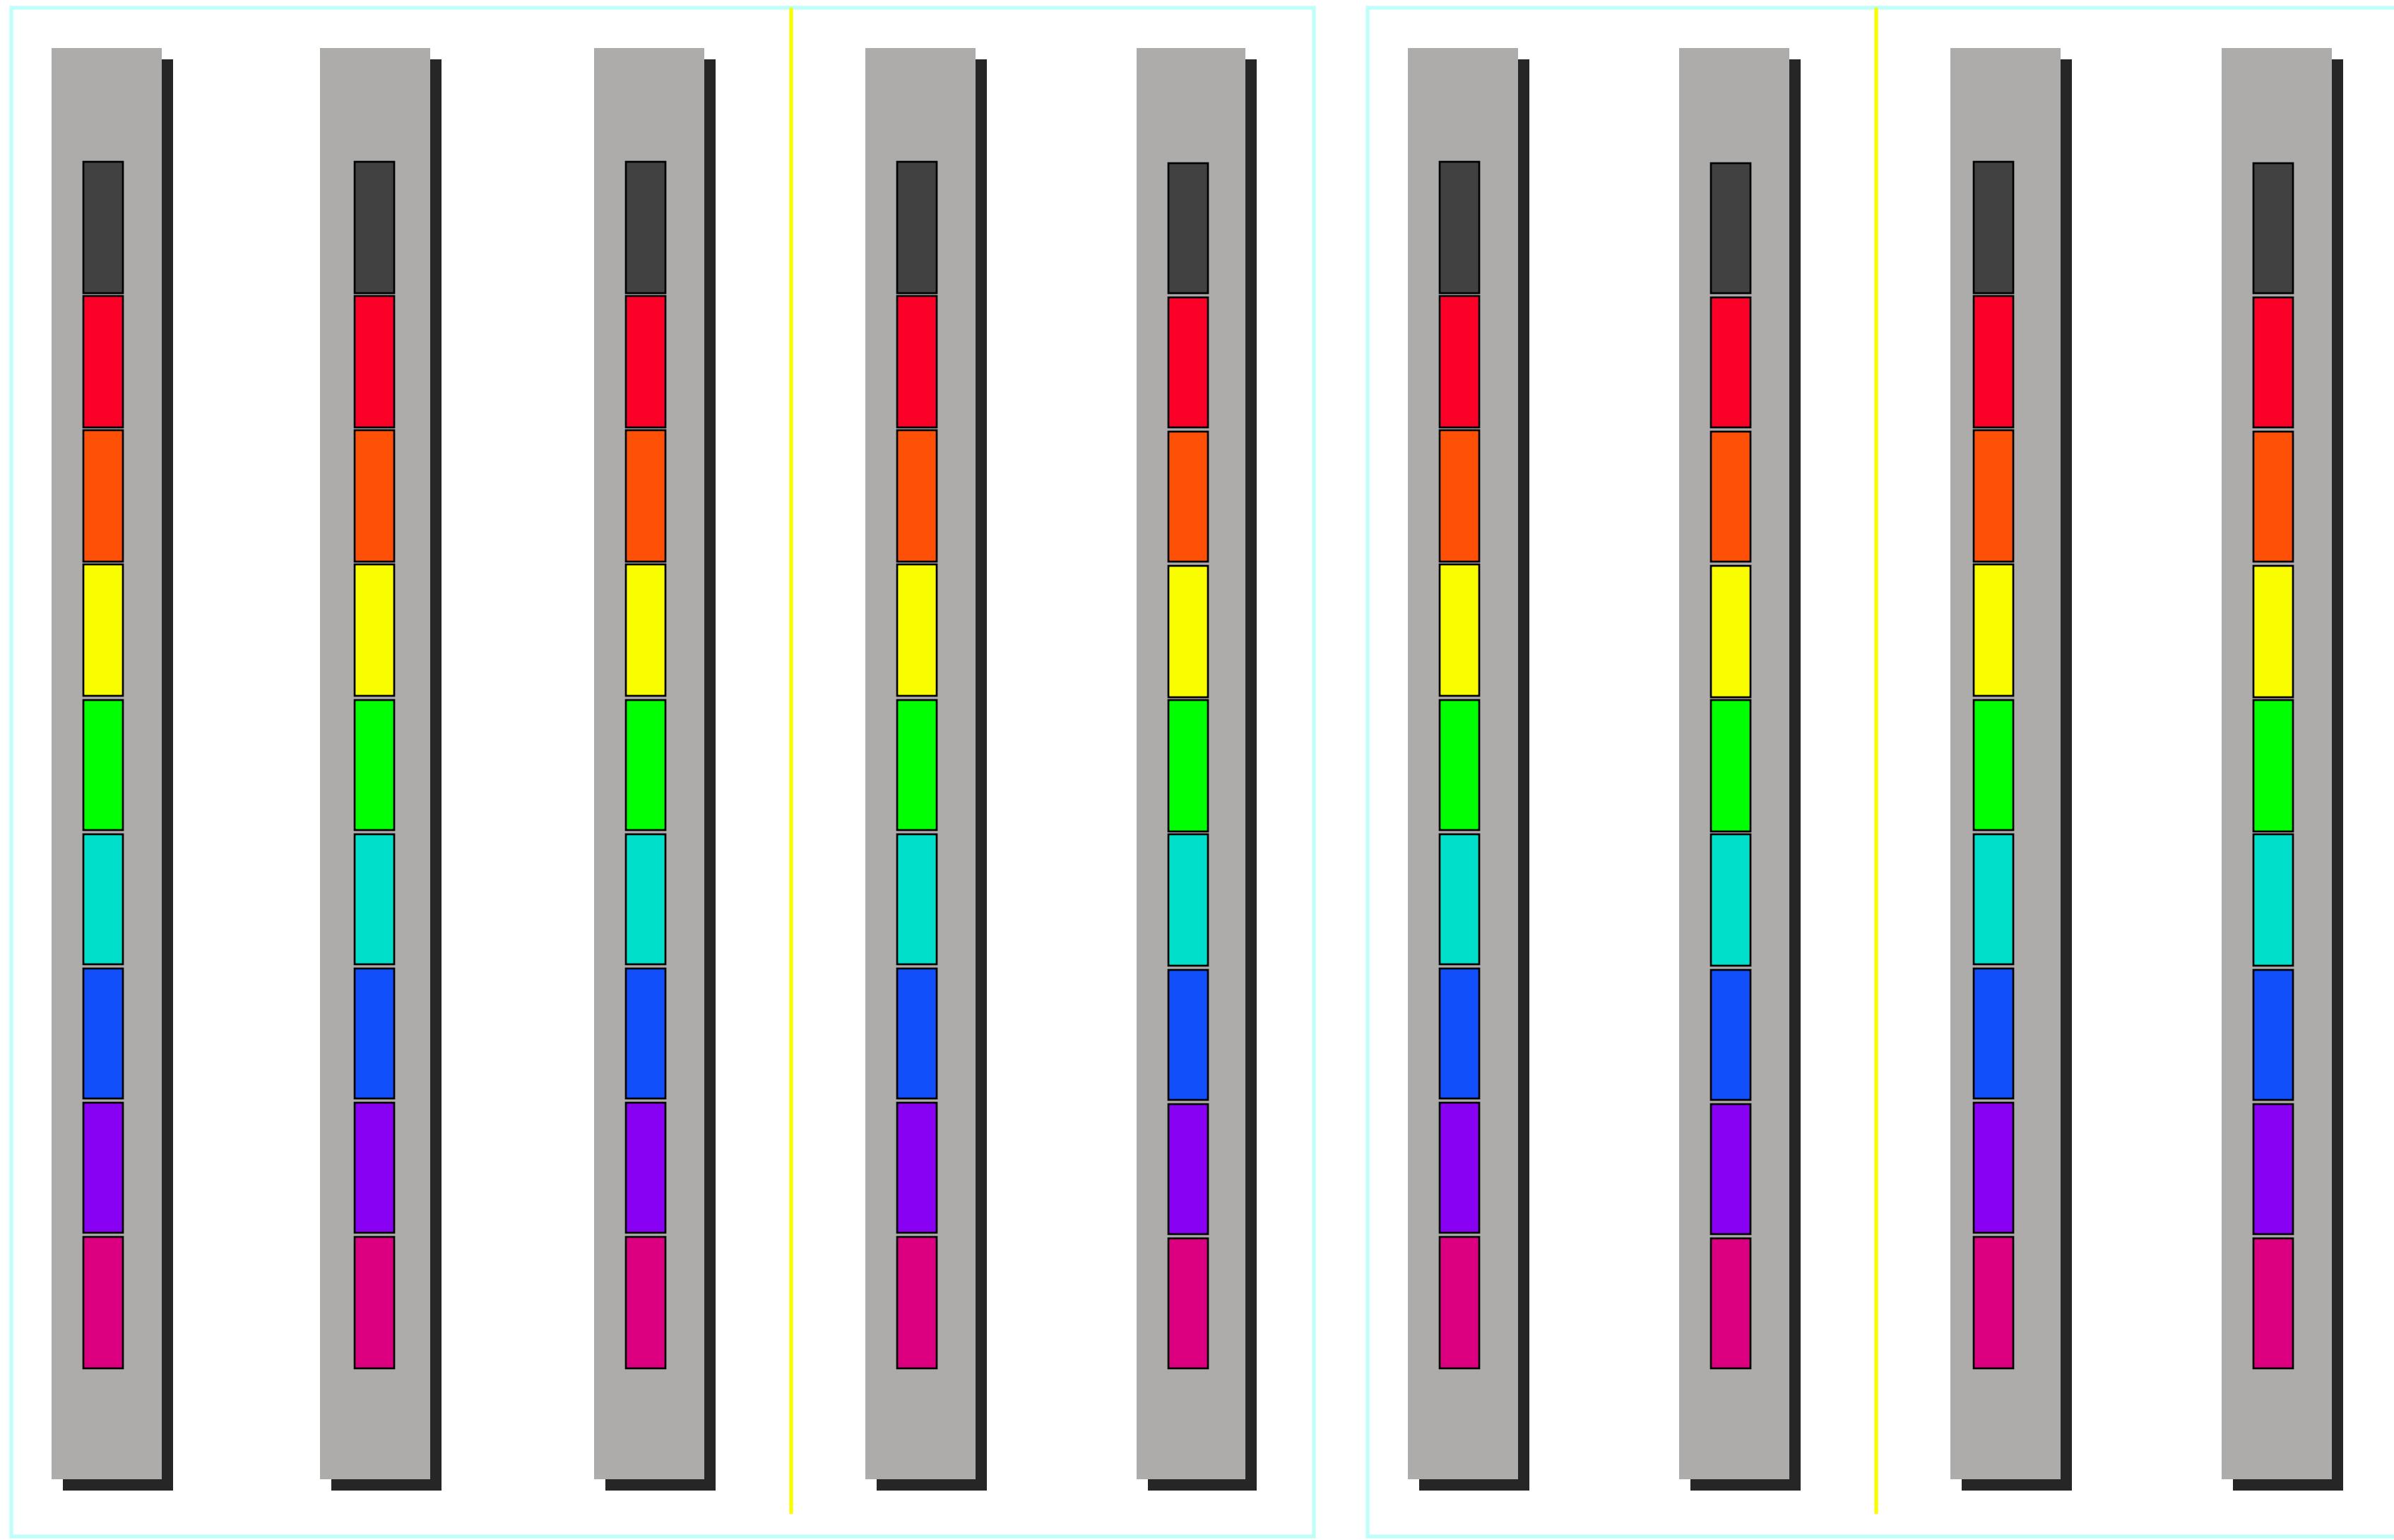


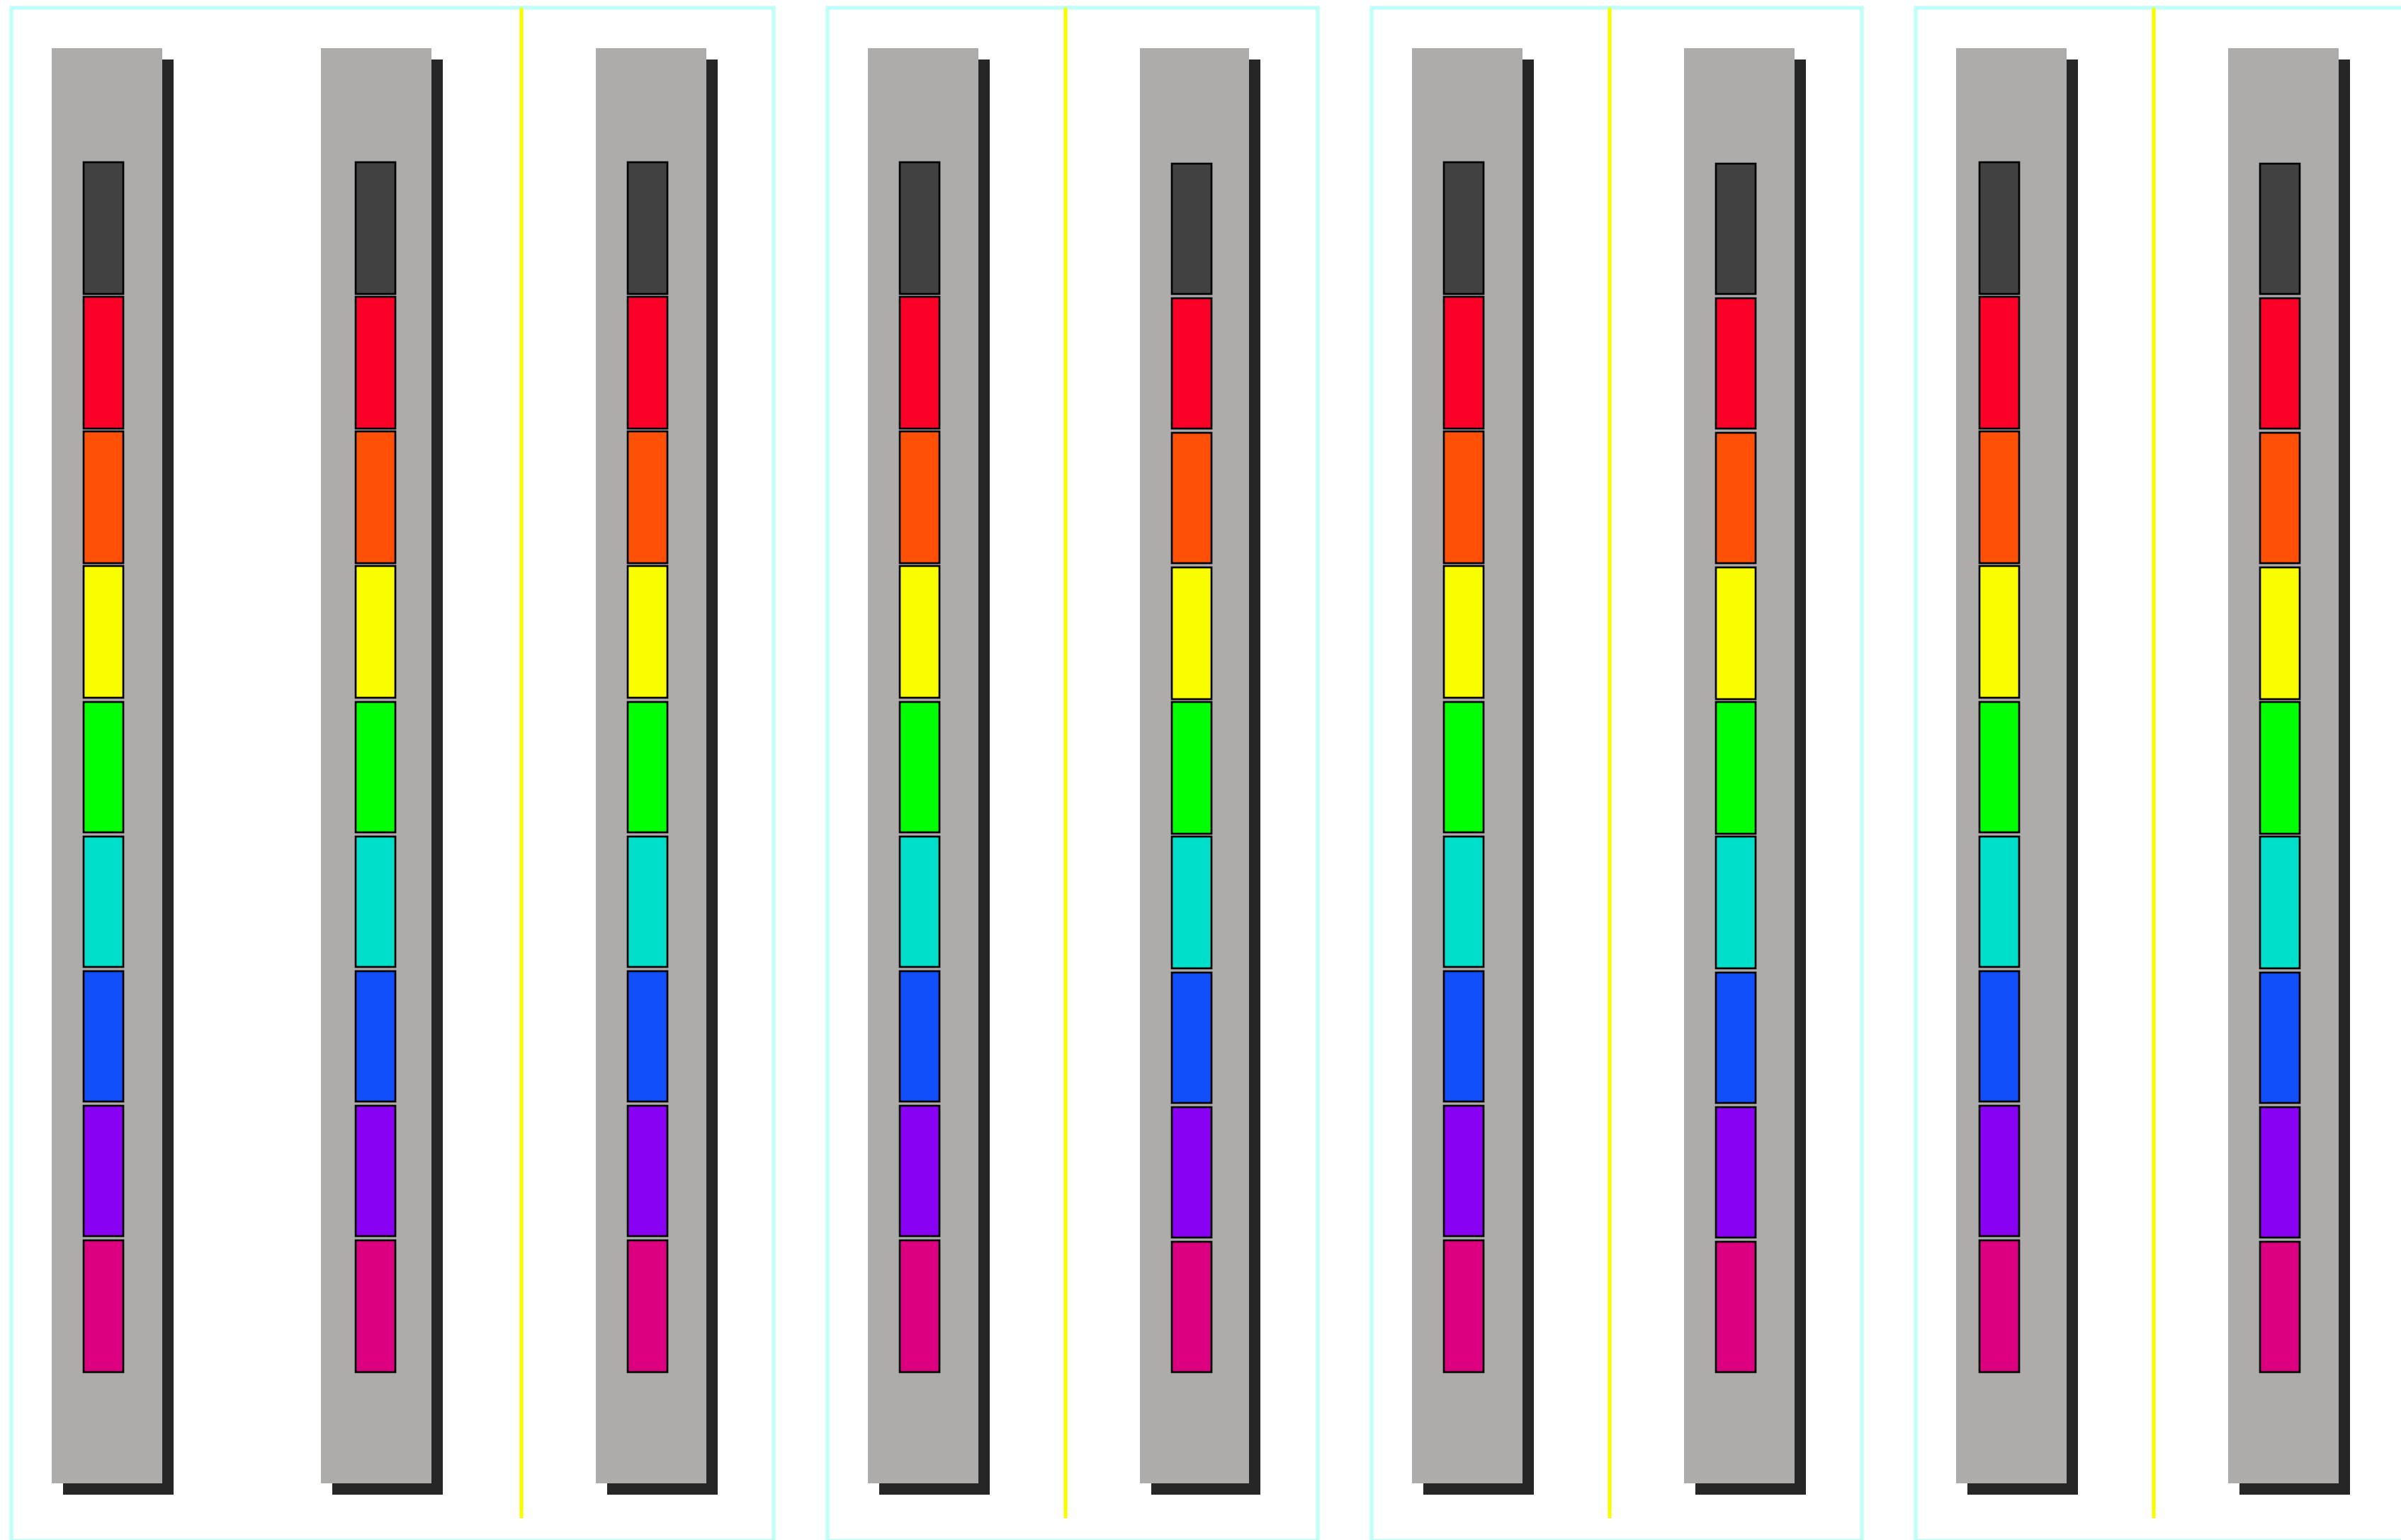
After

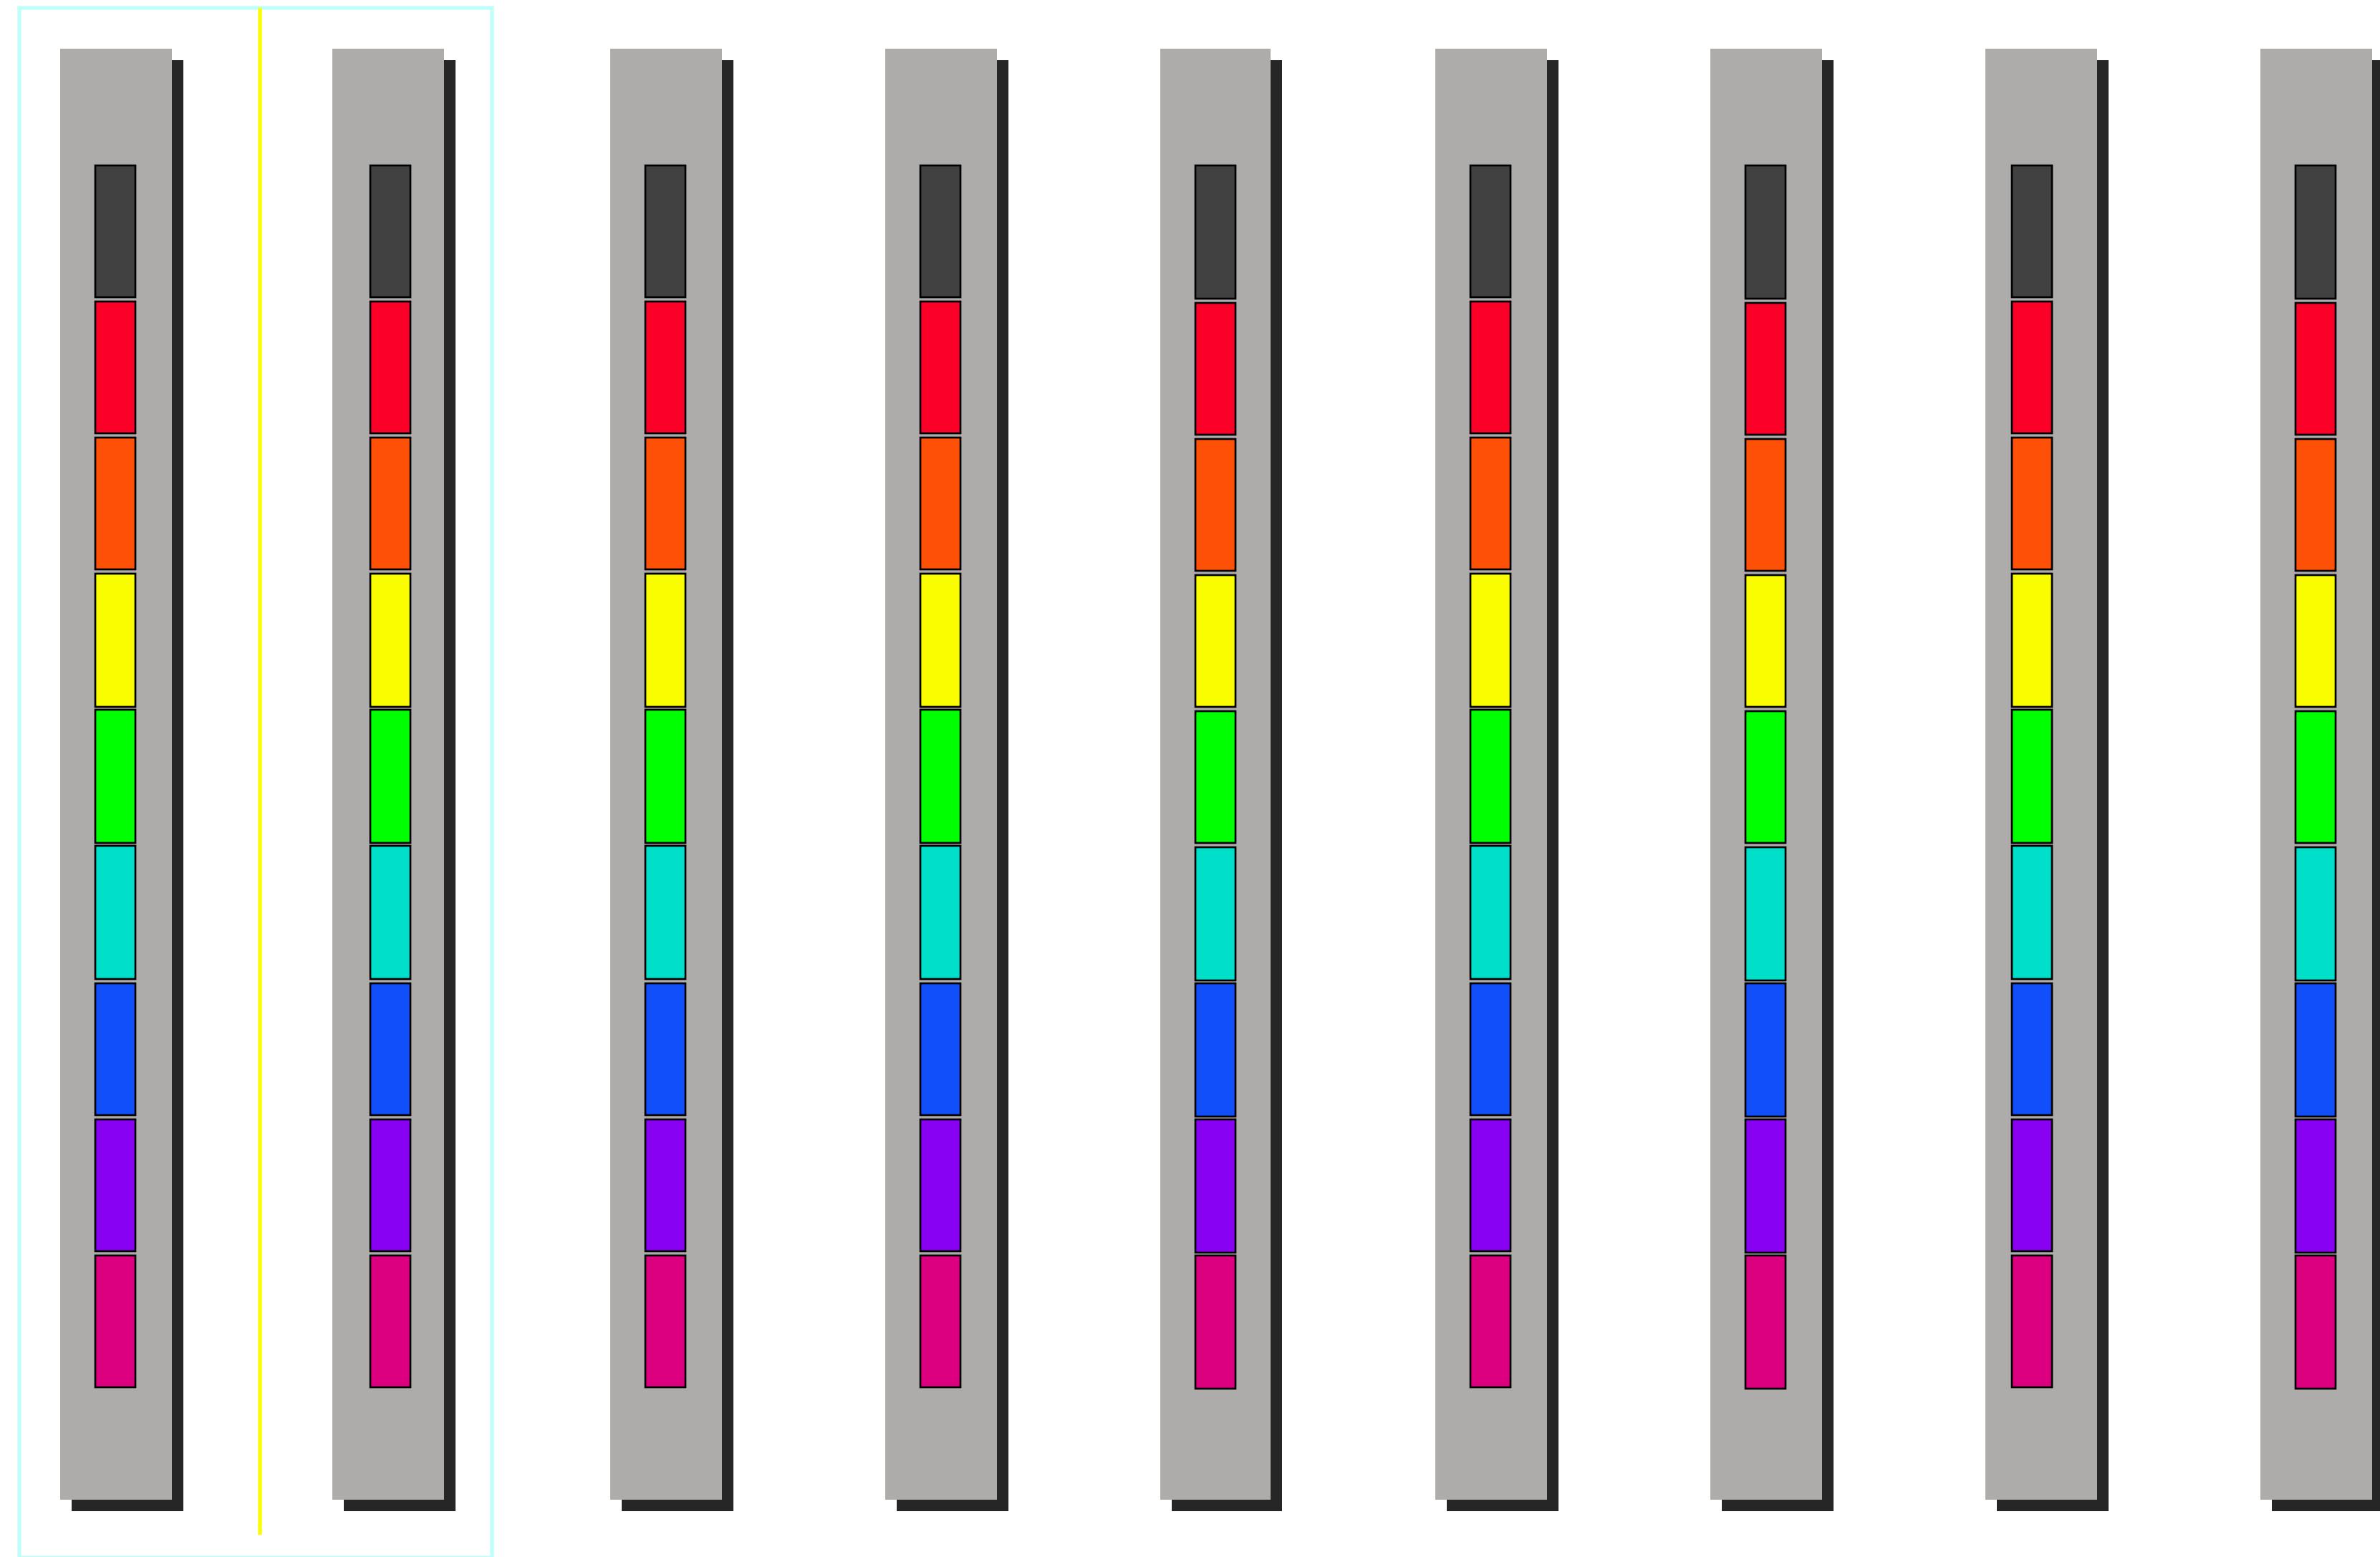


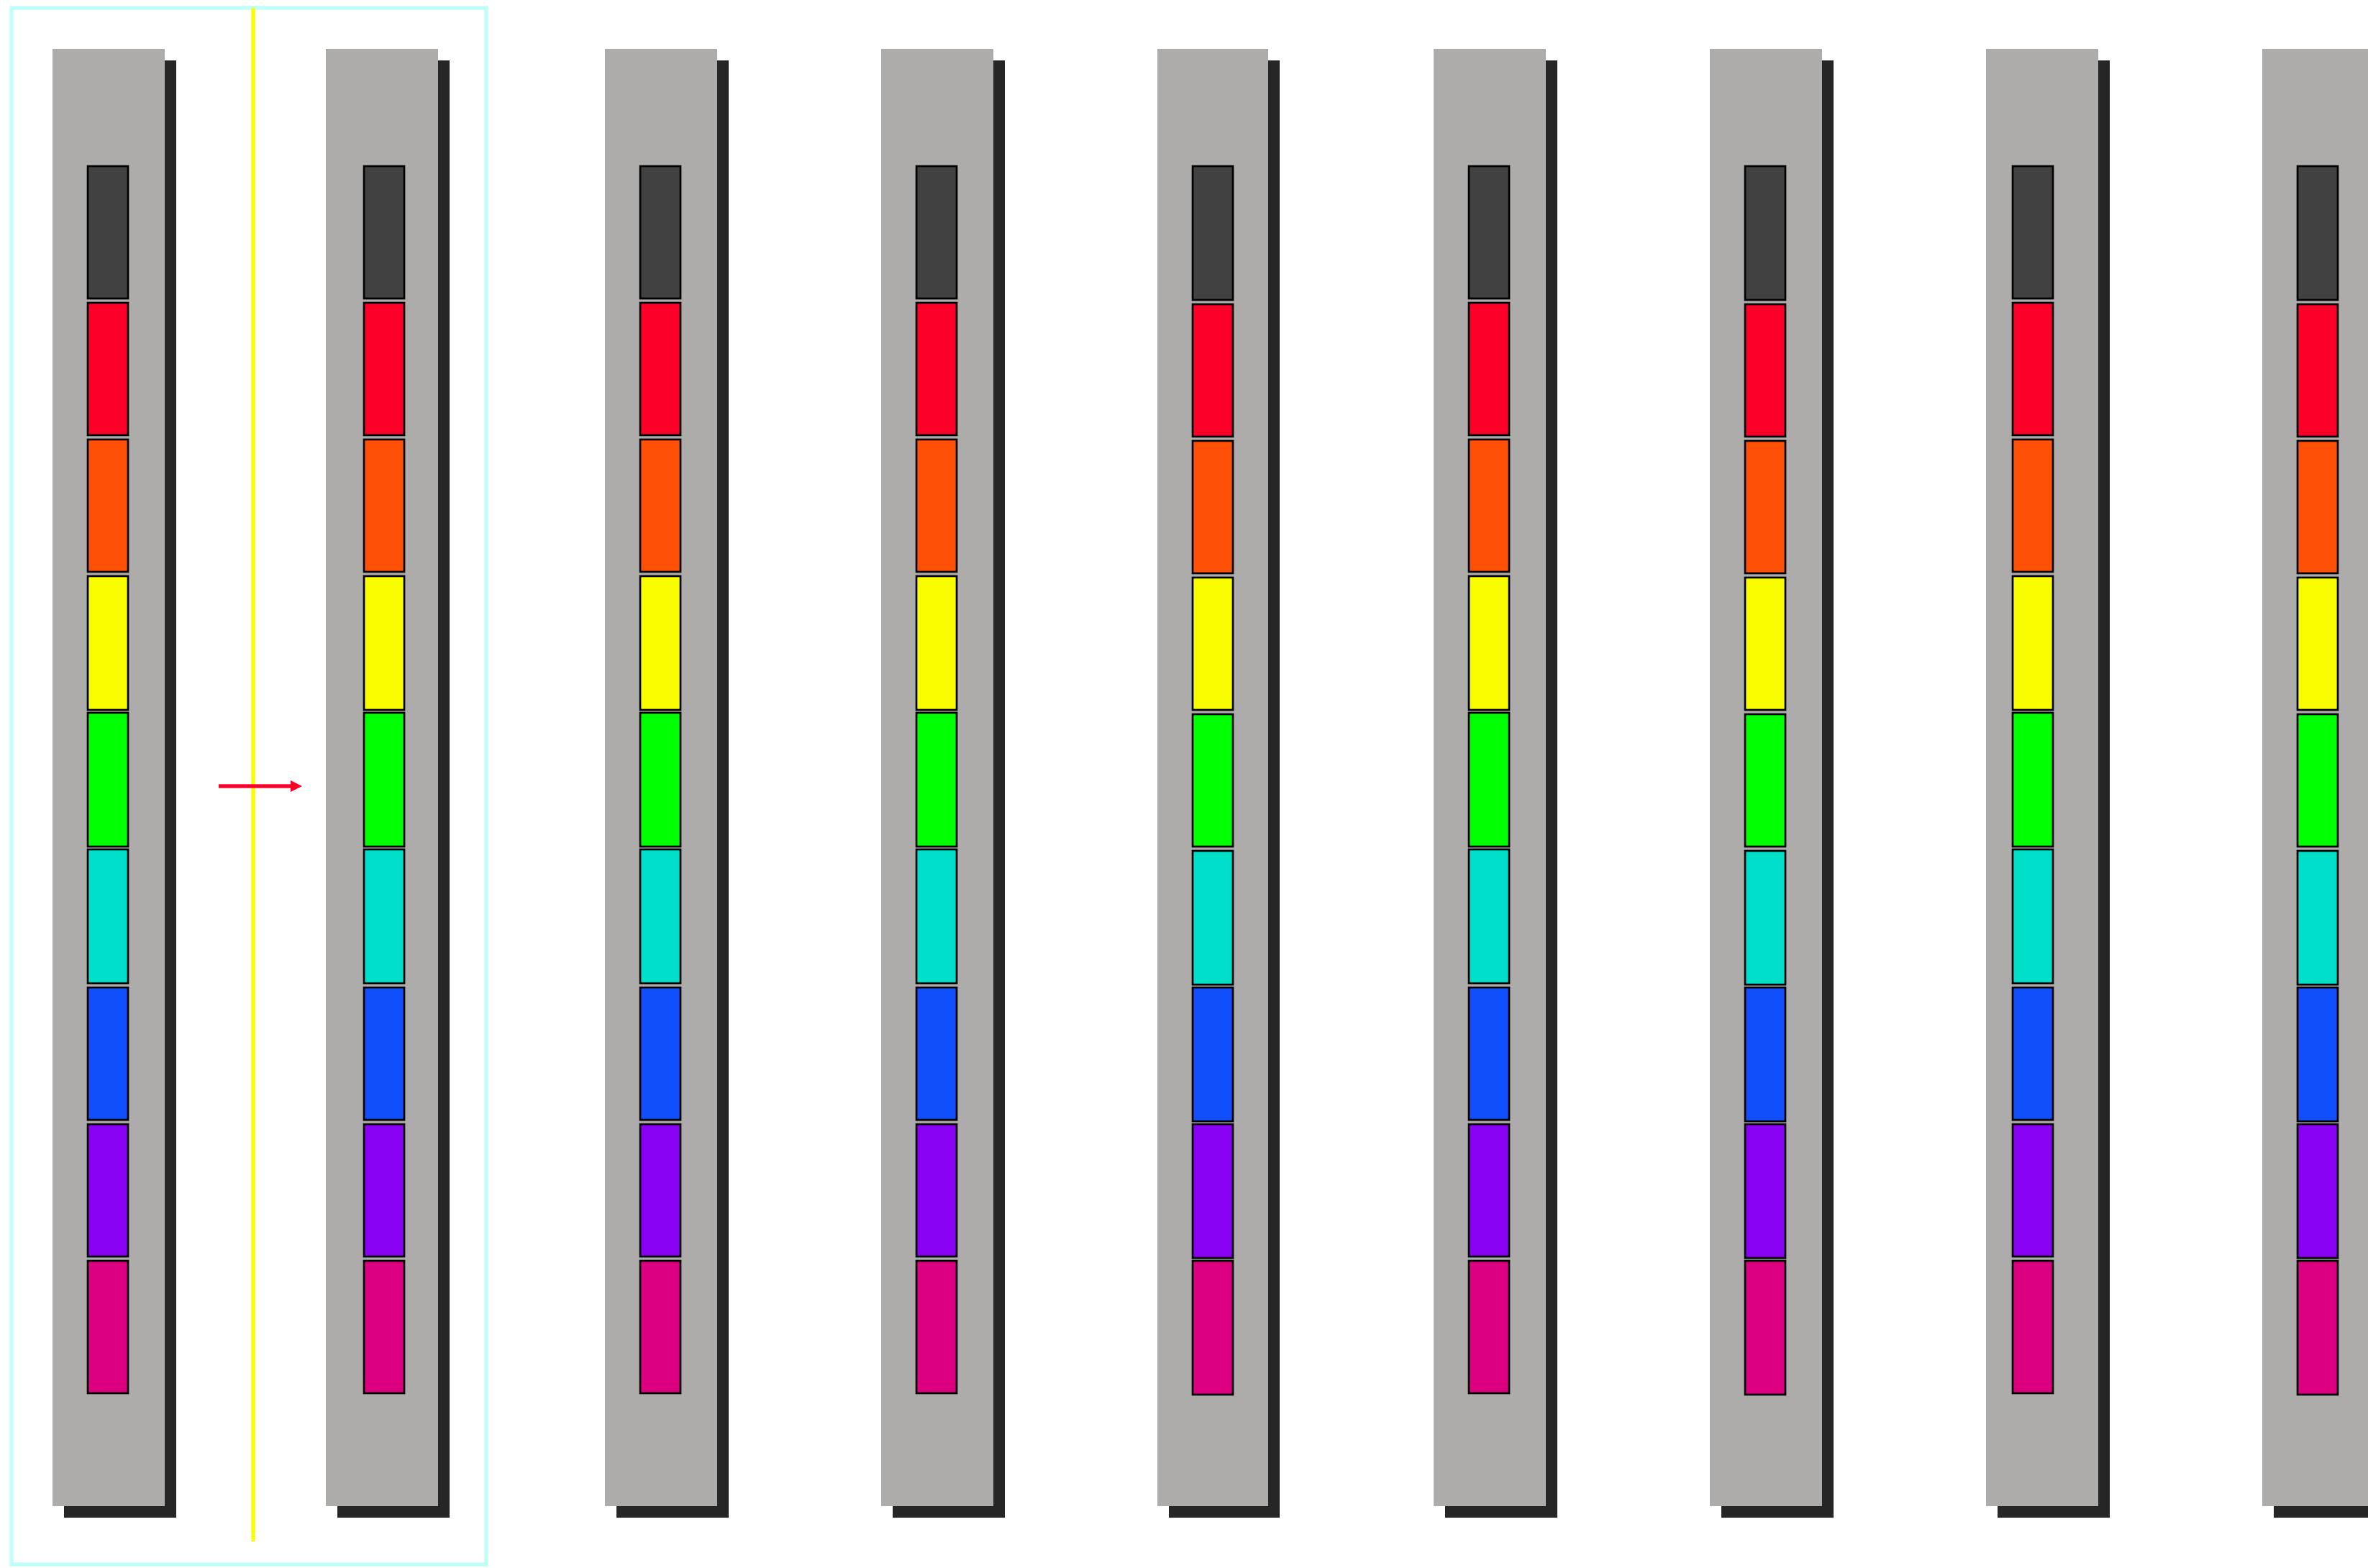


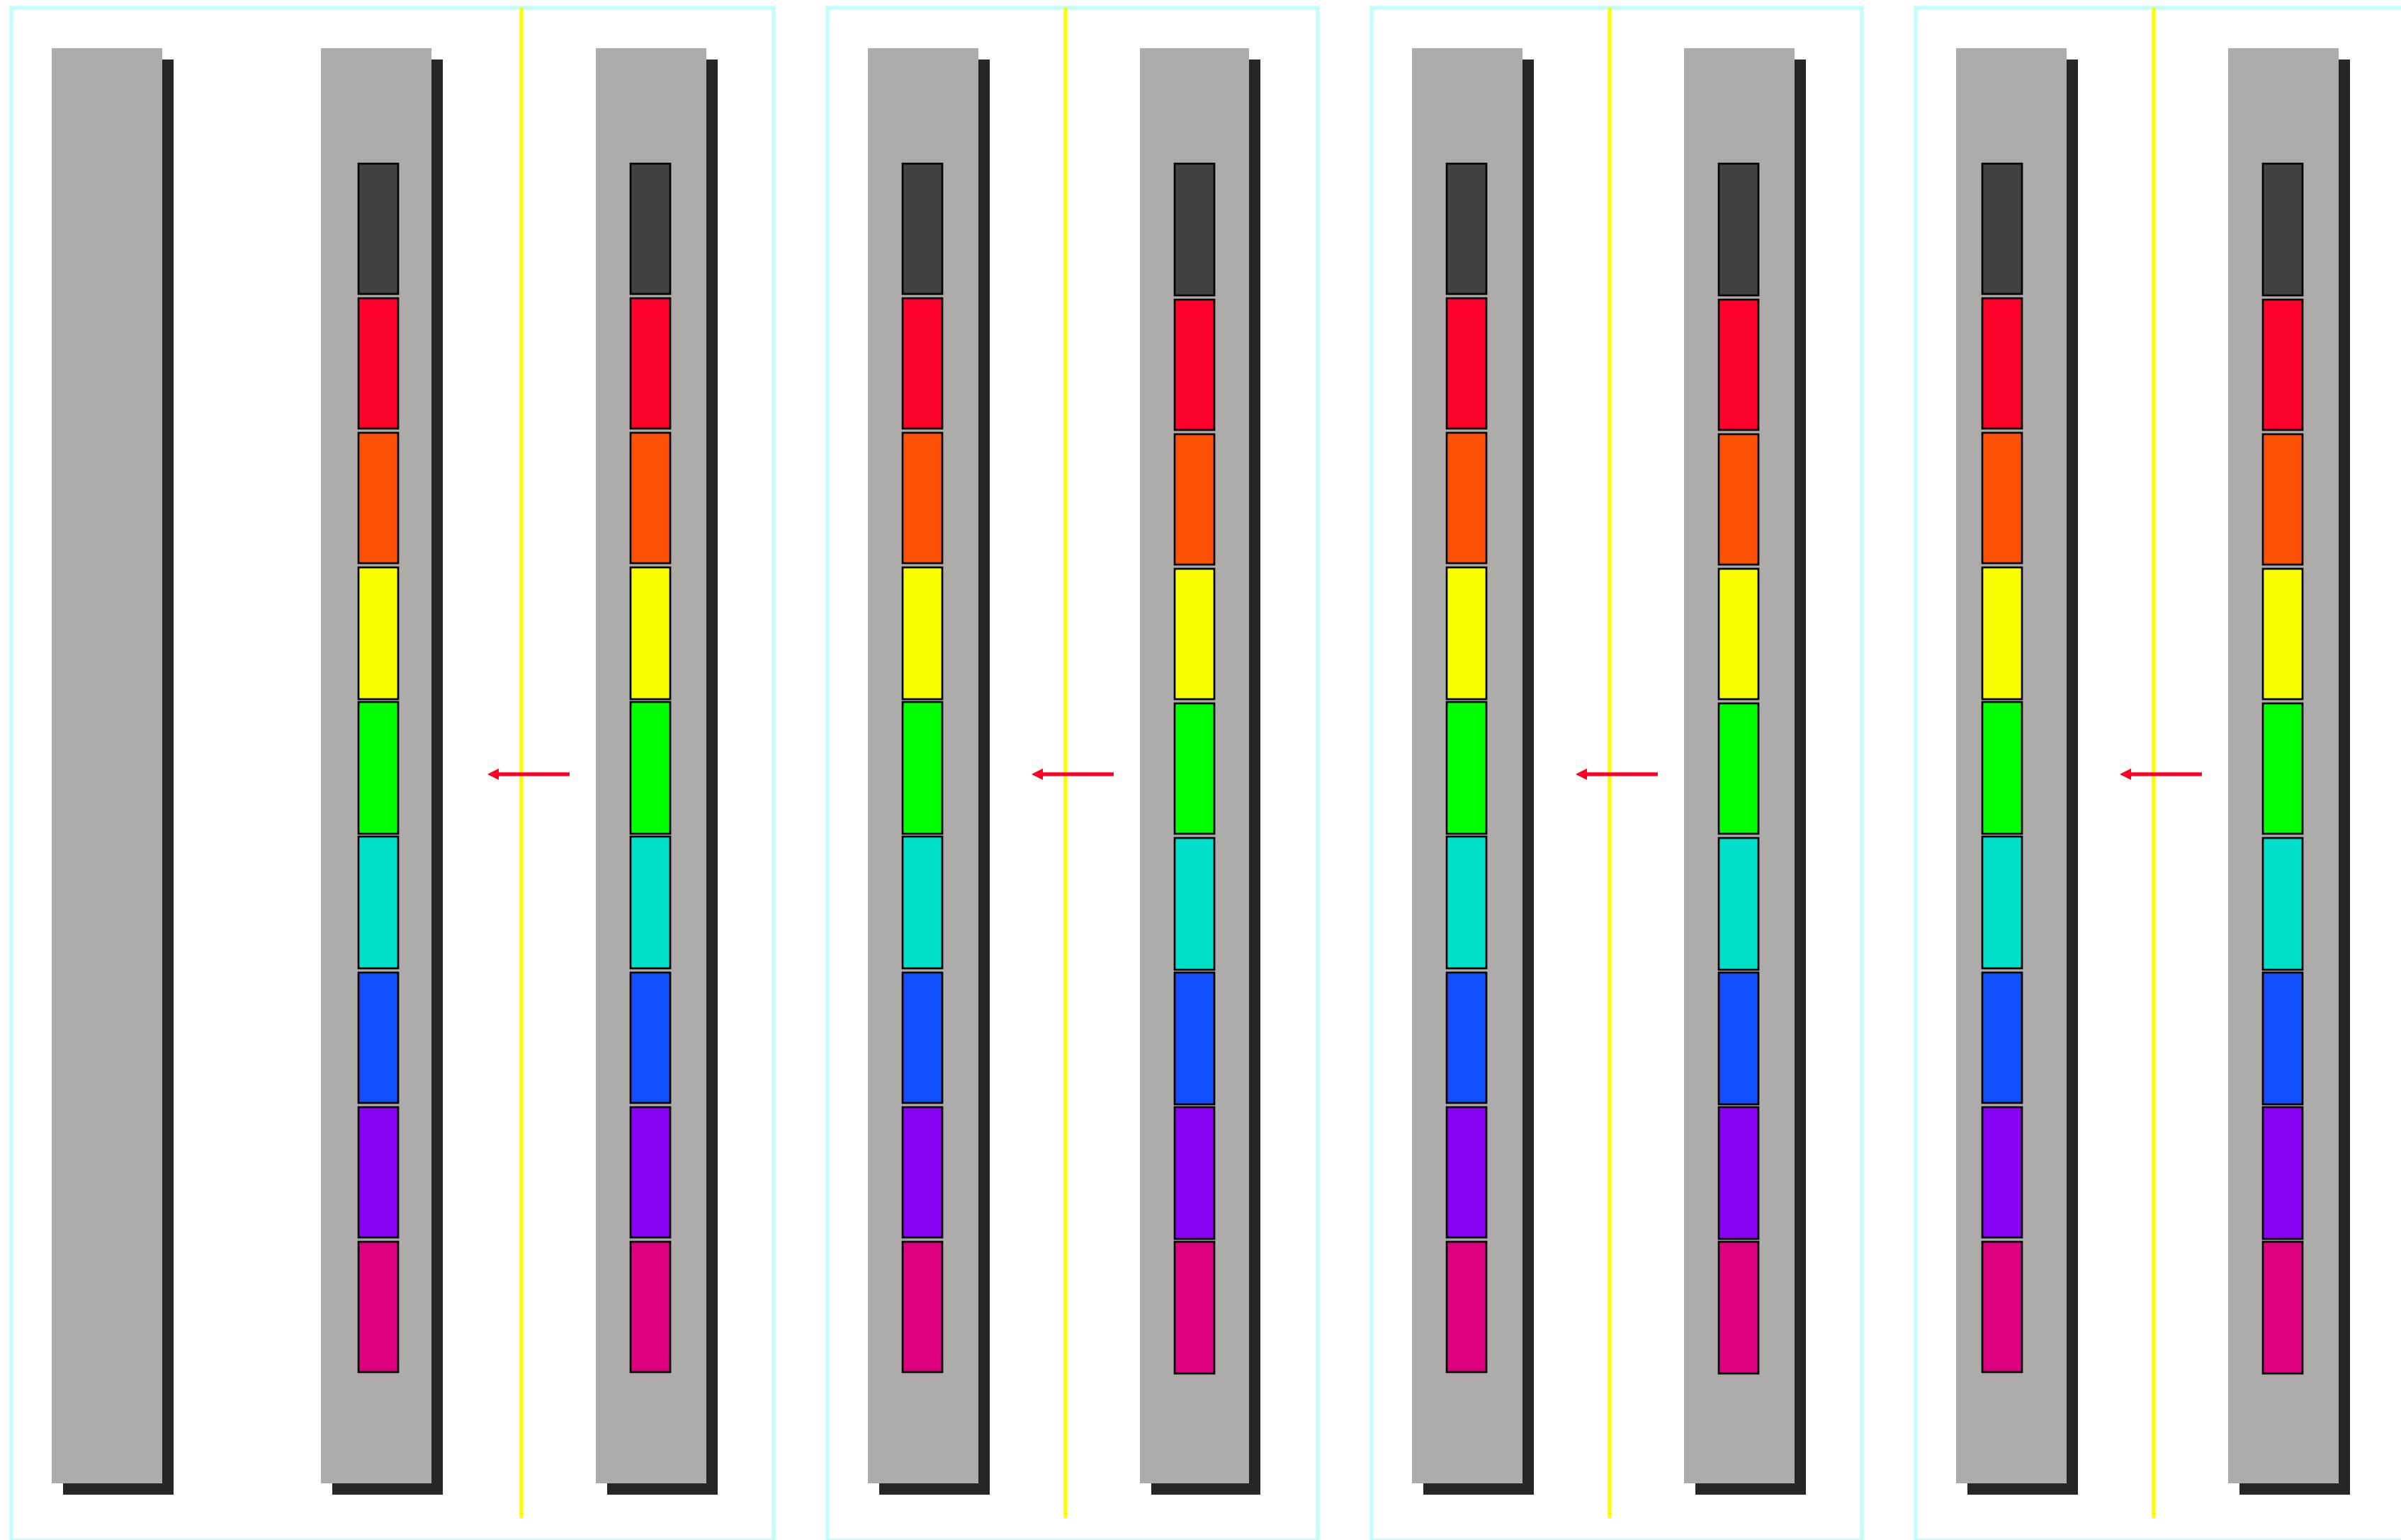


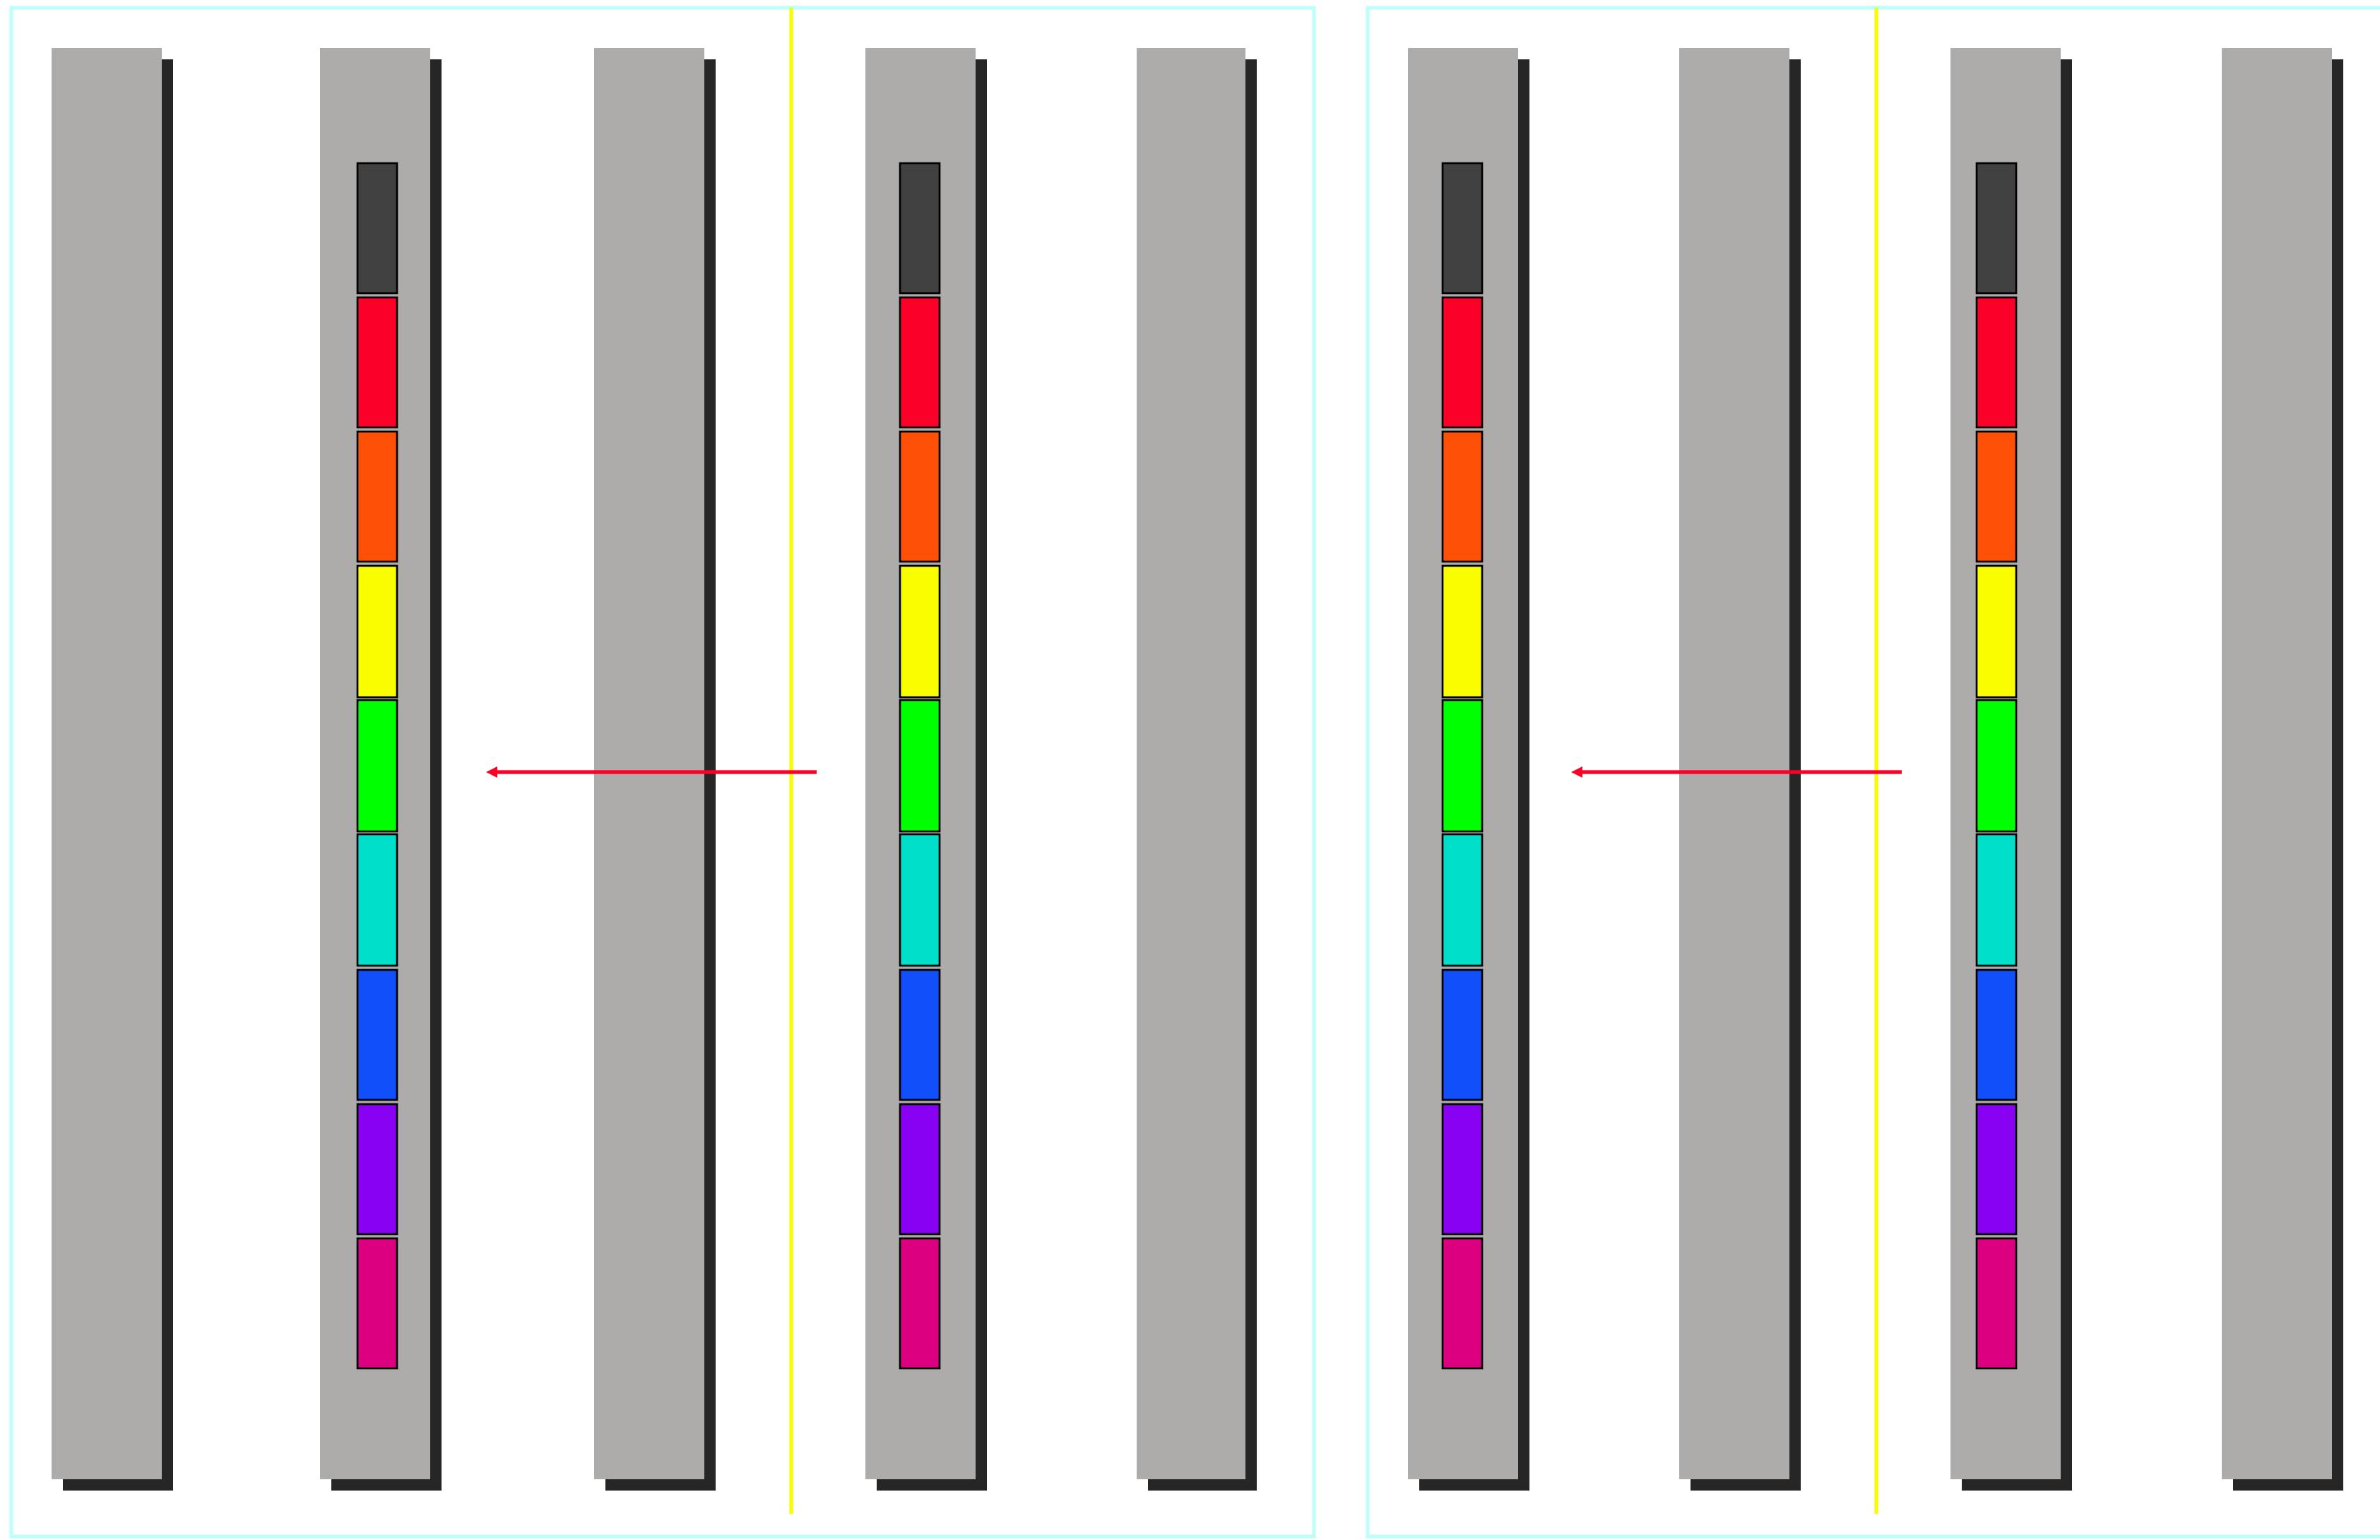


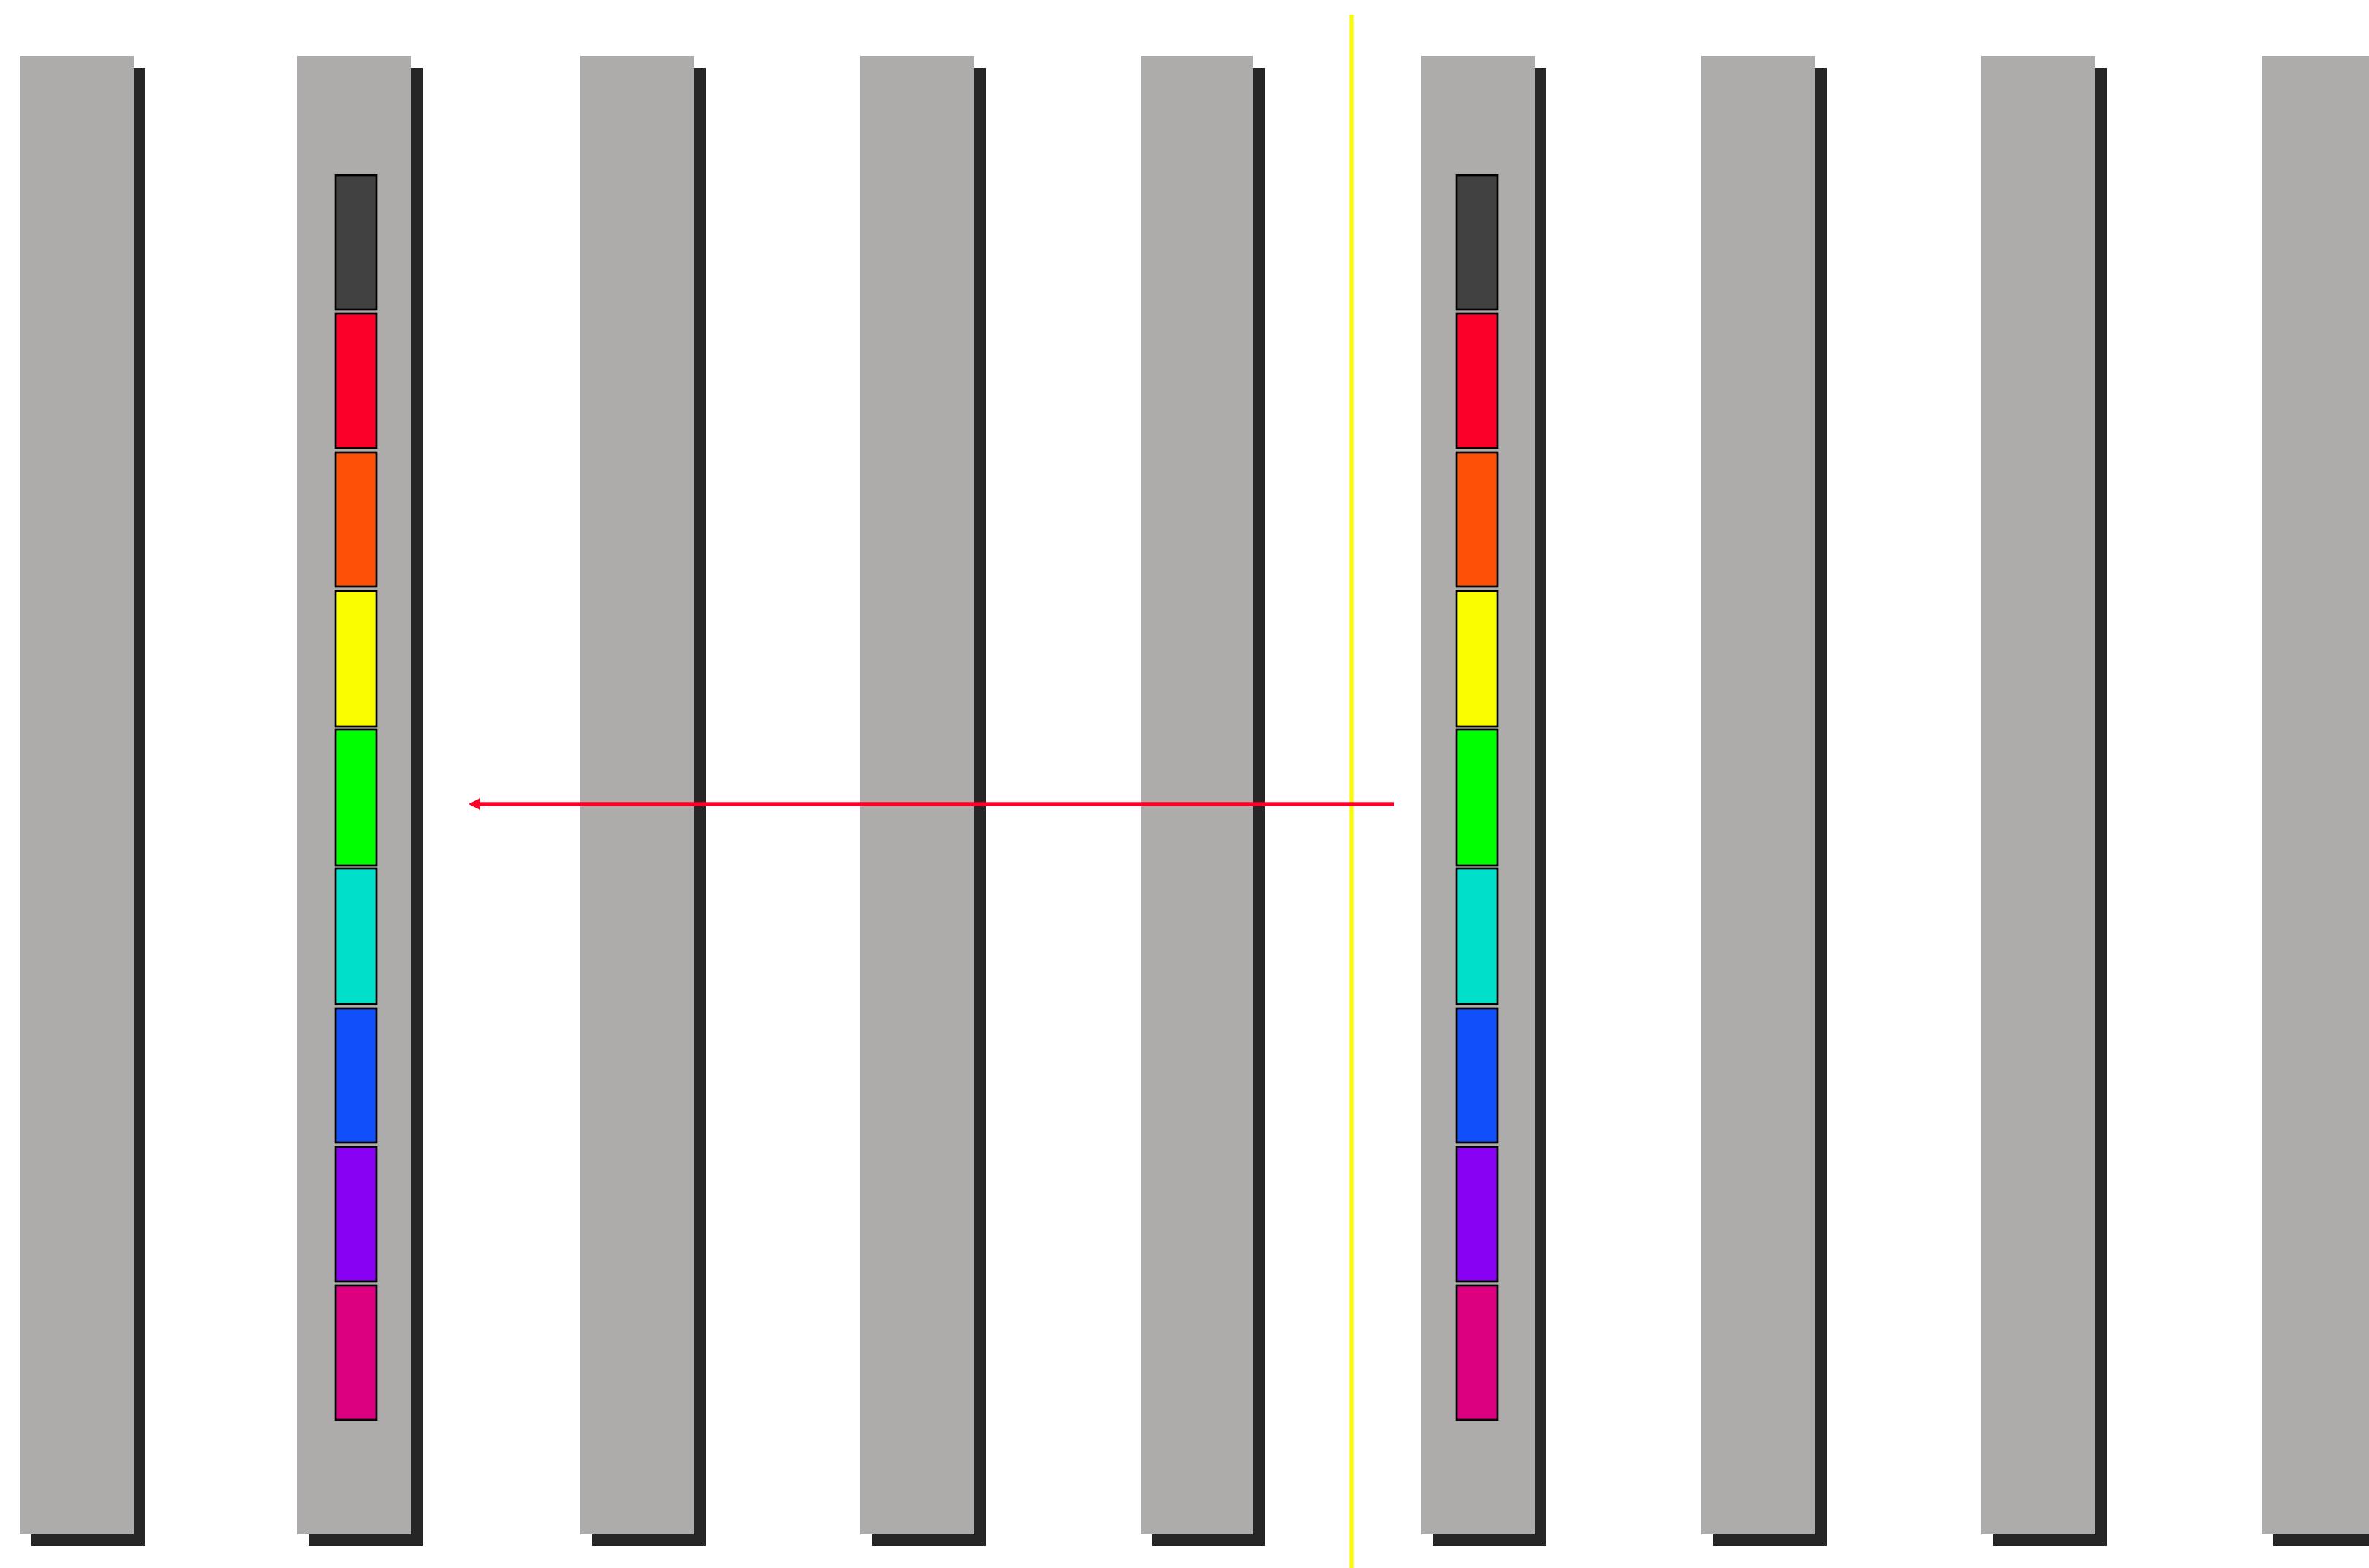










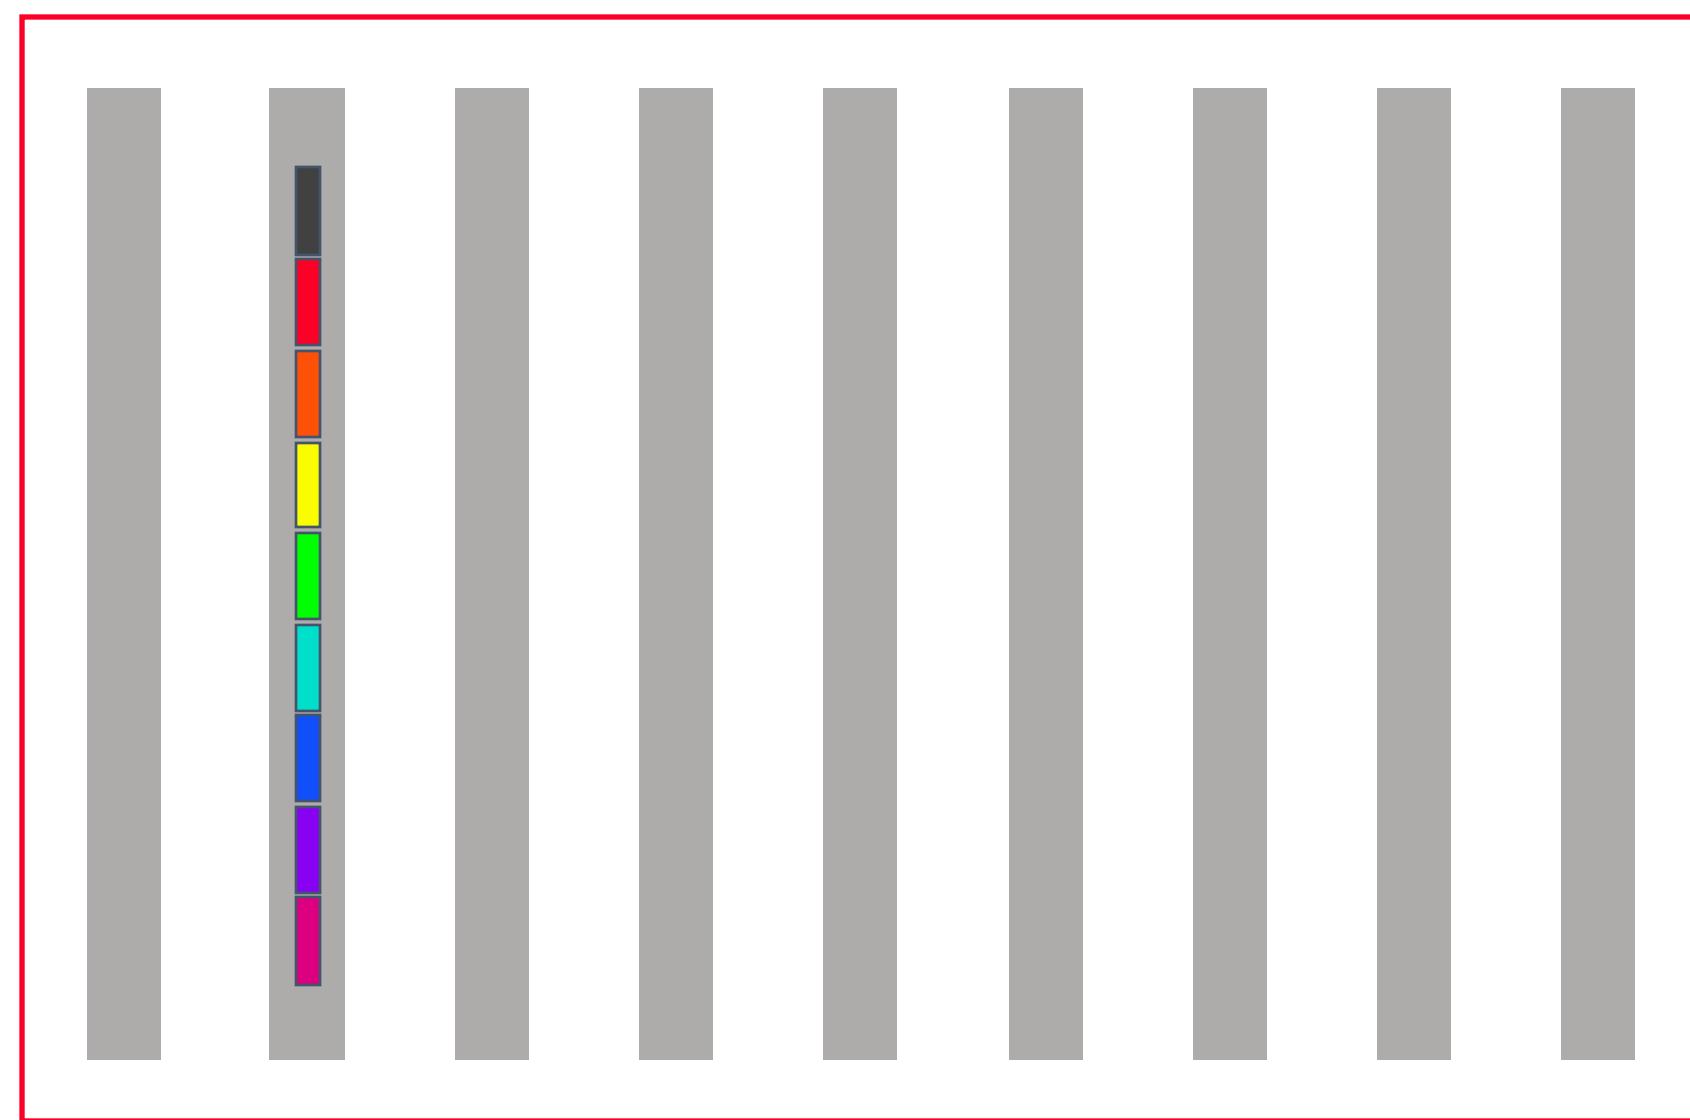


Cost of minimum spanning tree reduce(-to-one)

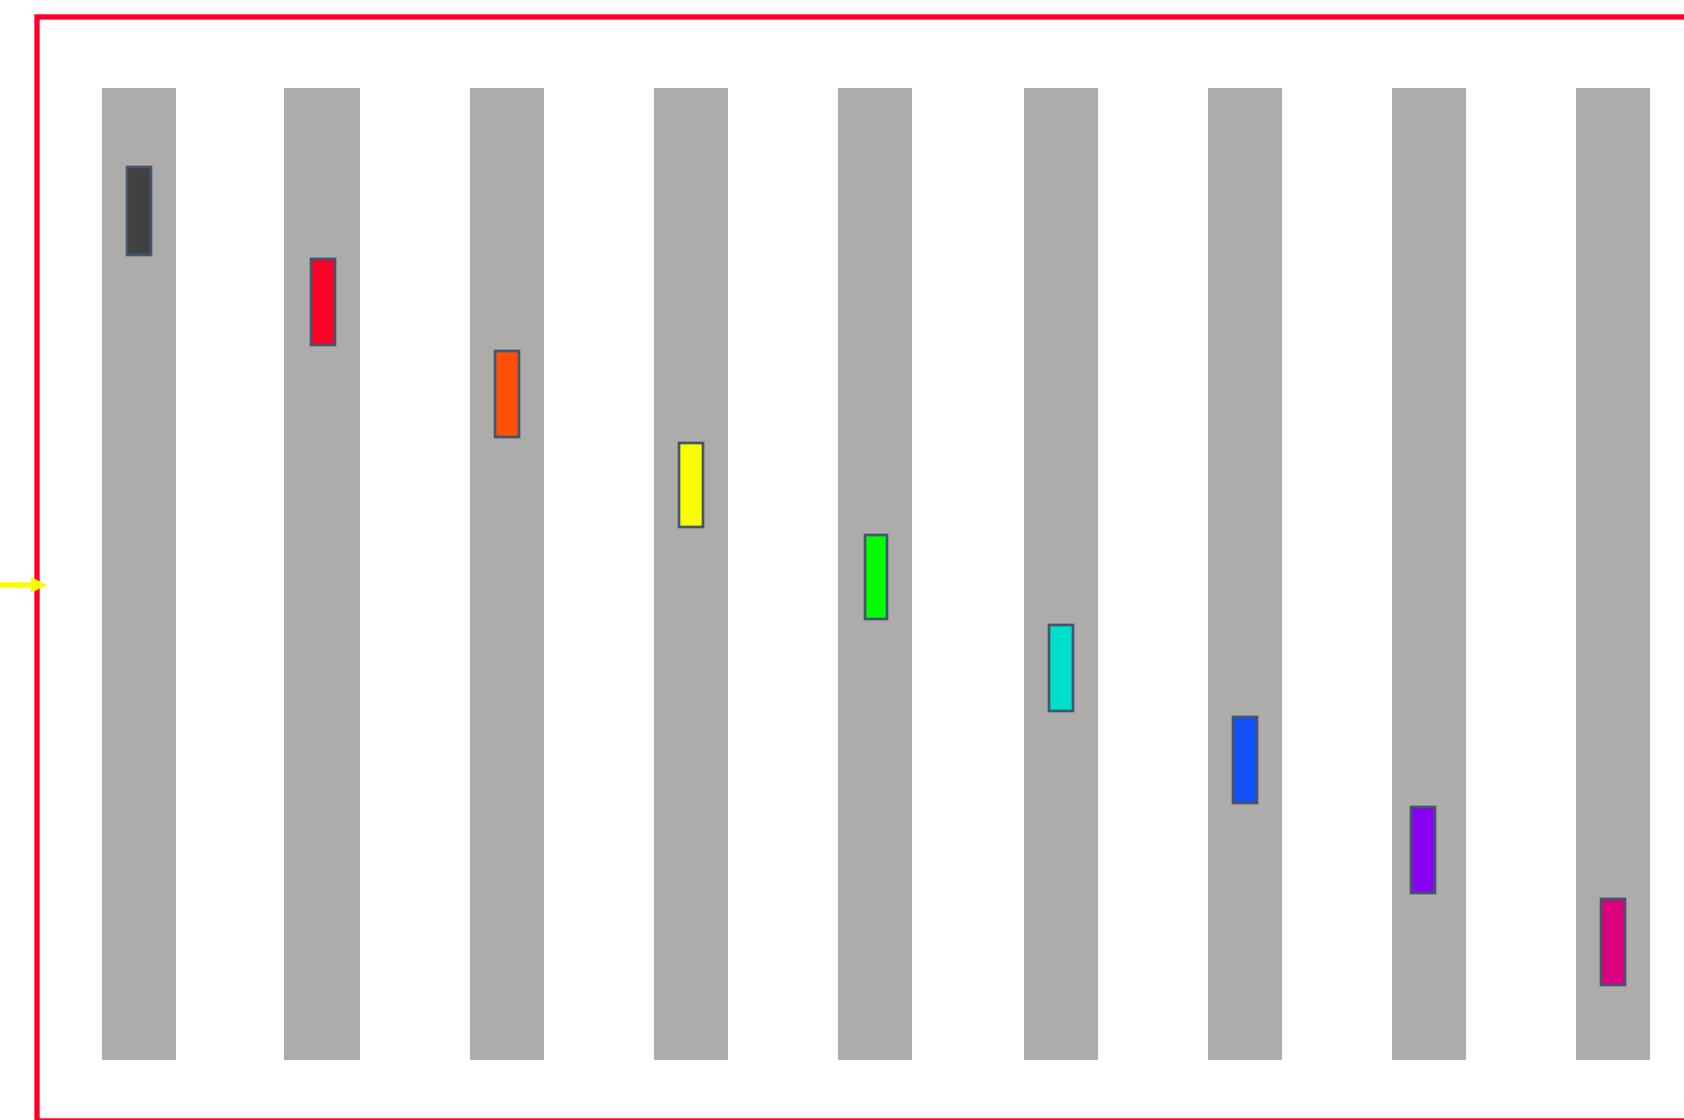
$$\lceil \log(p) \rceil (\alpha + n\beta + n\gamma)$$

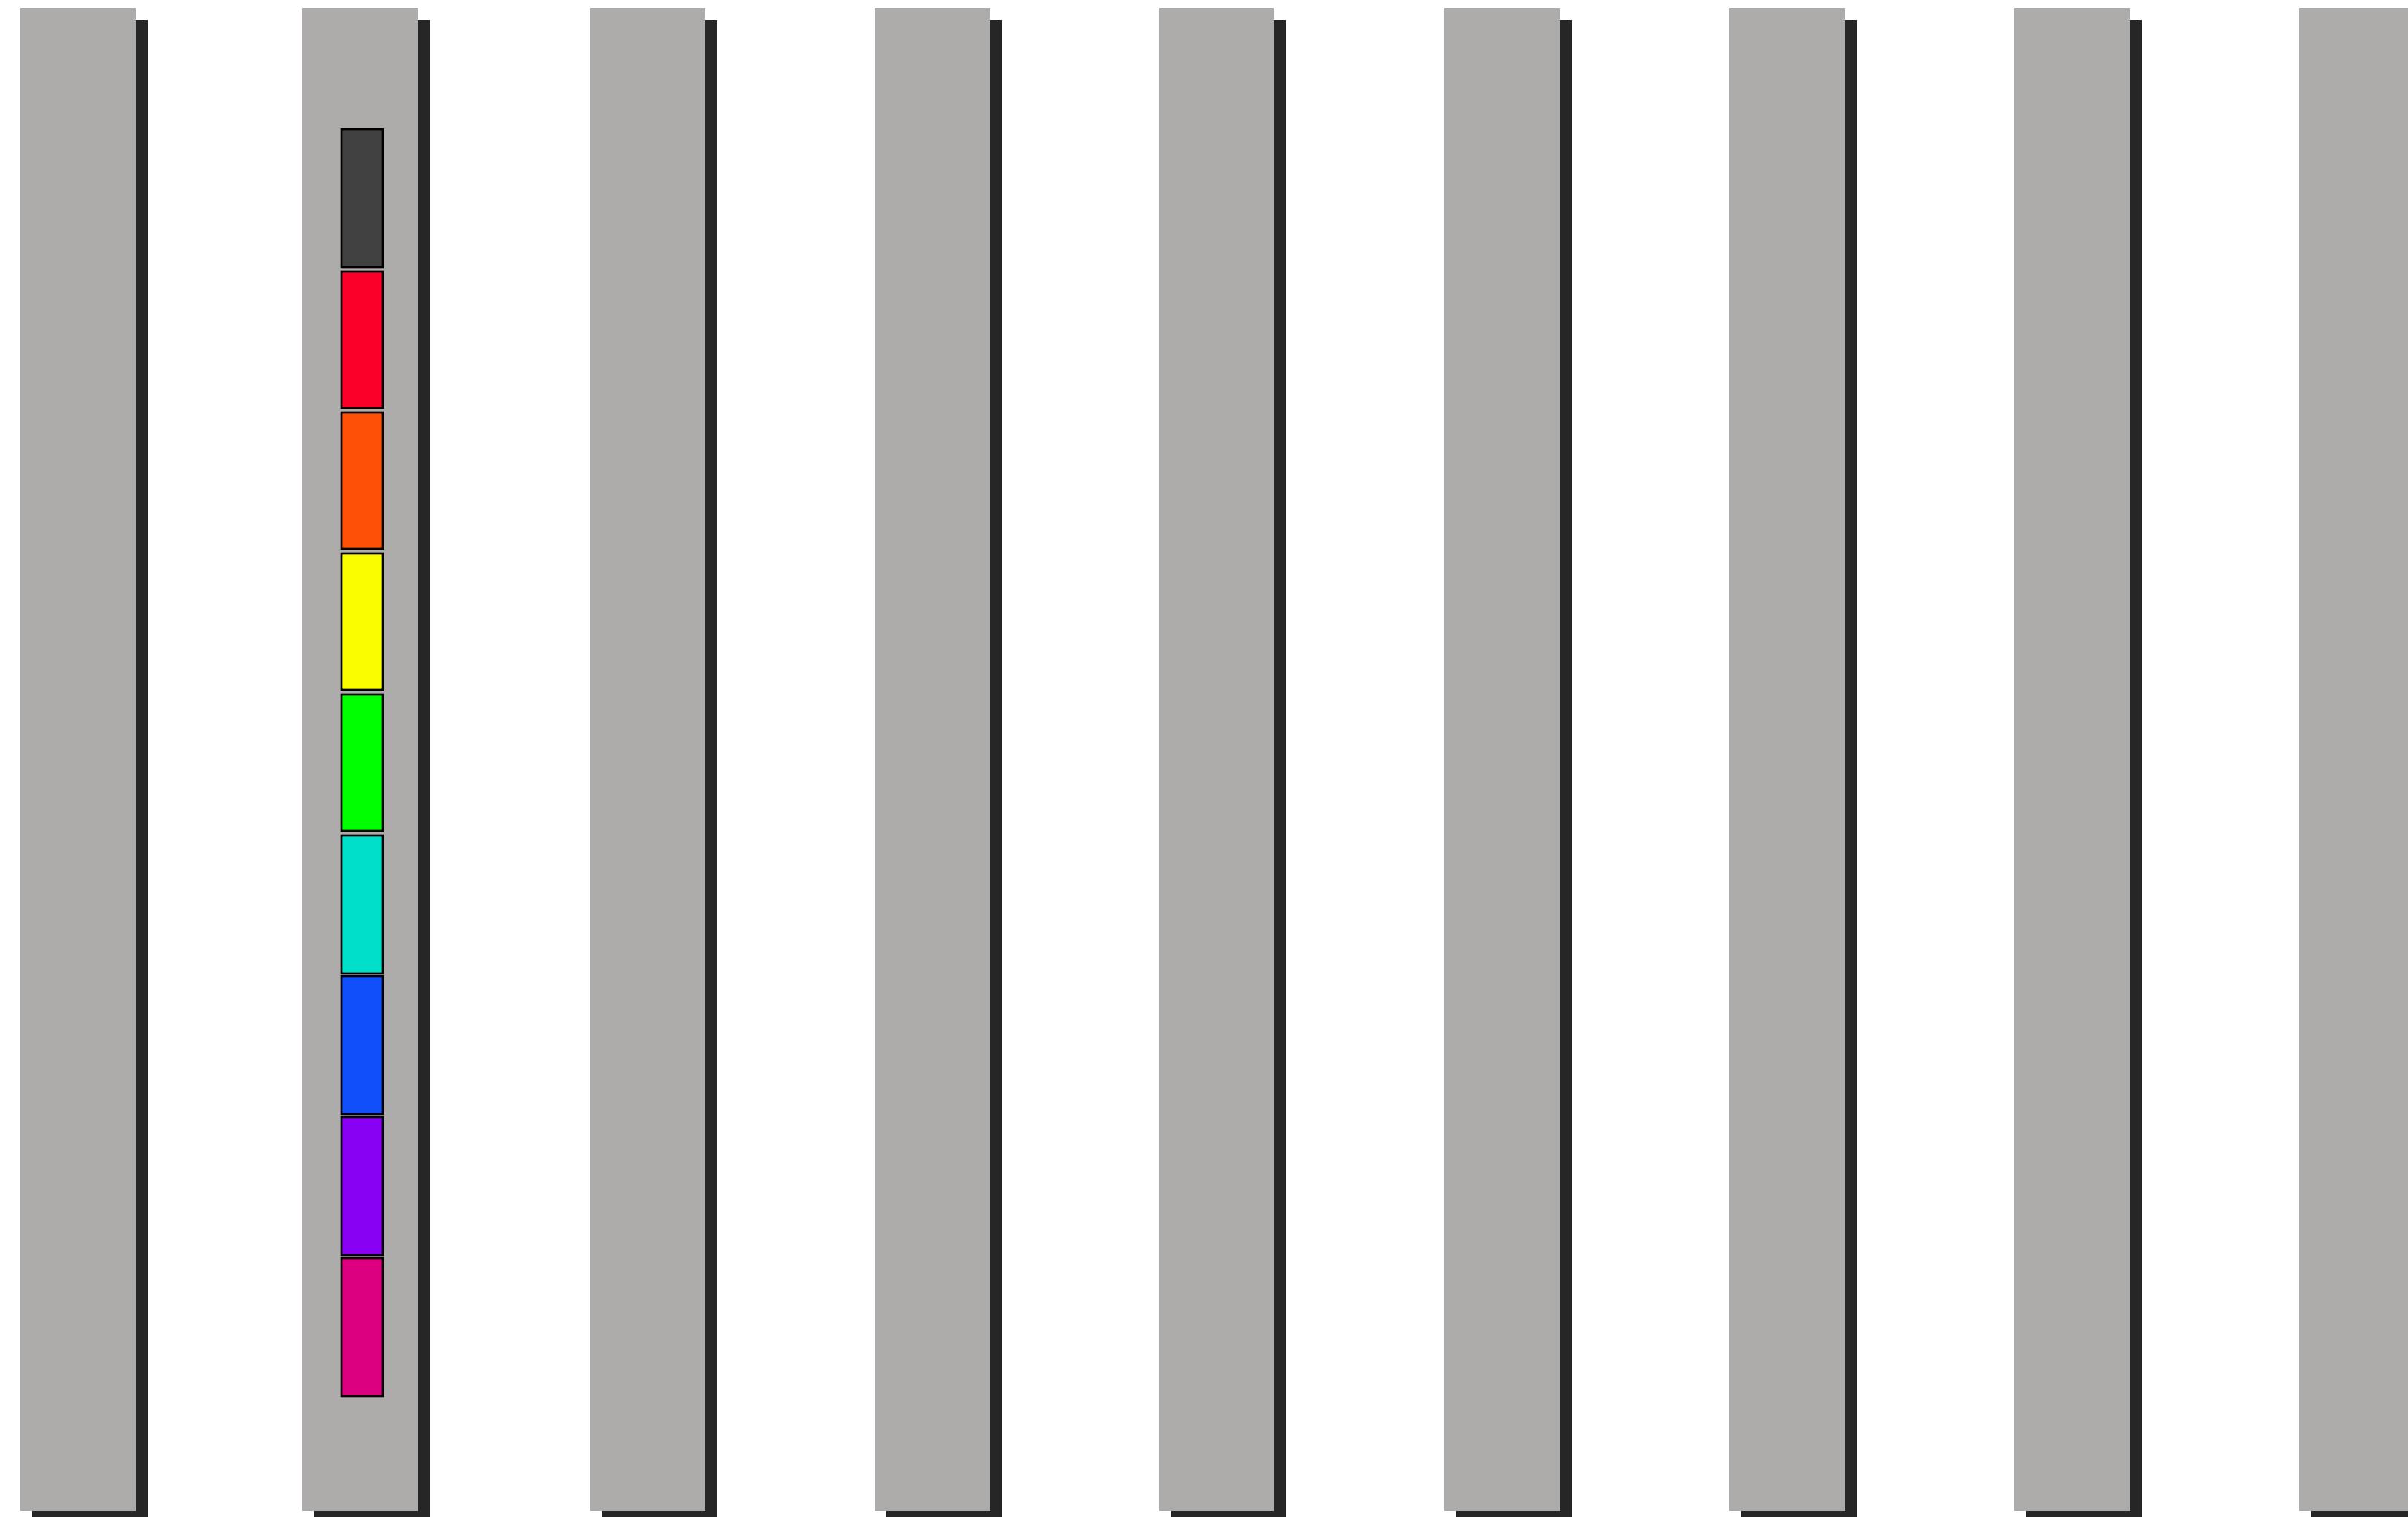
Scatter

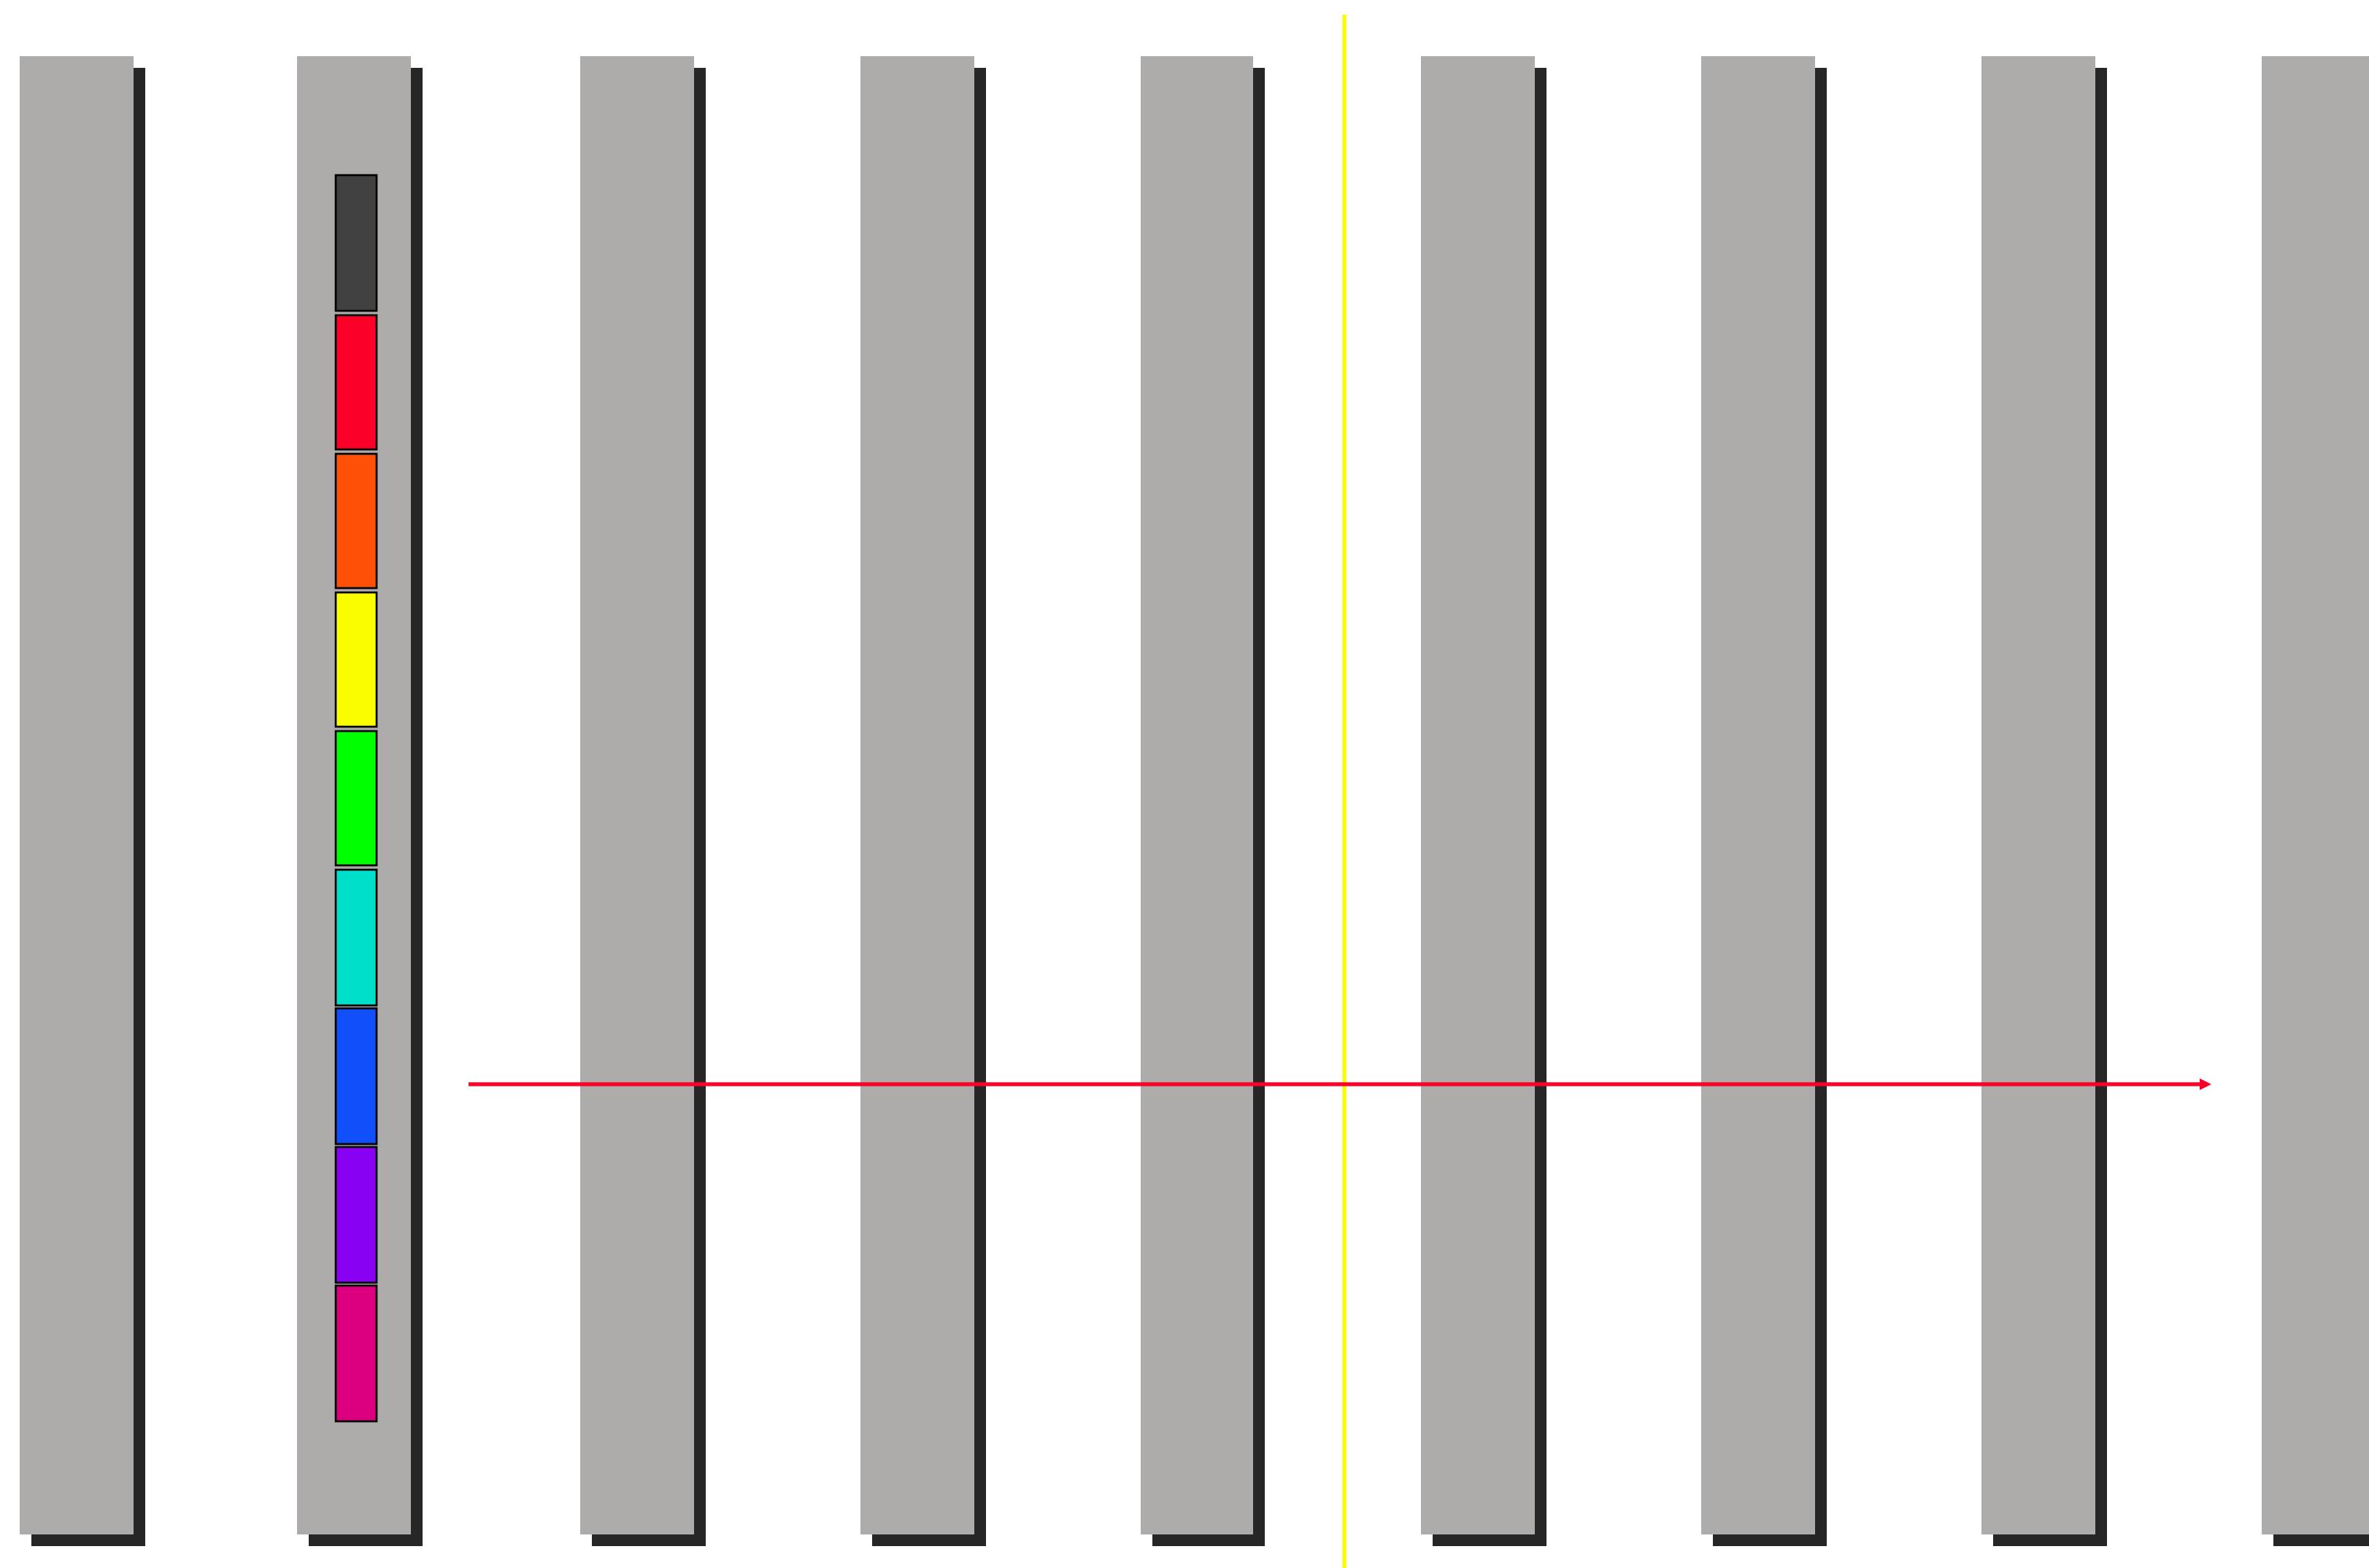
Before

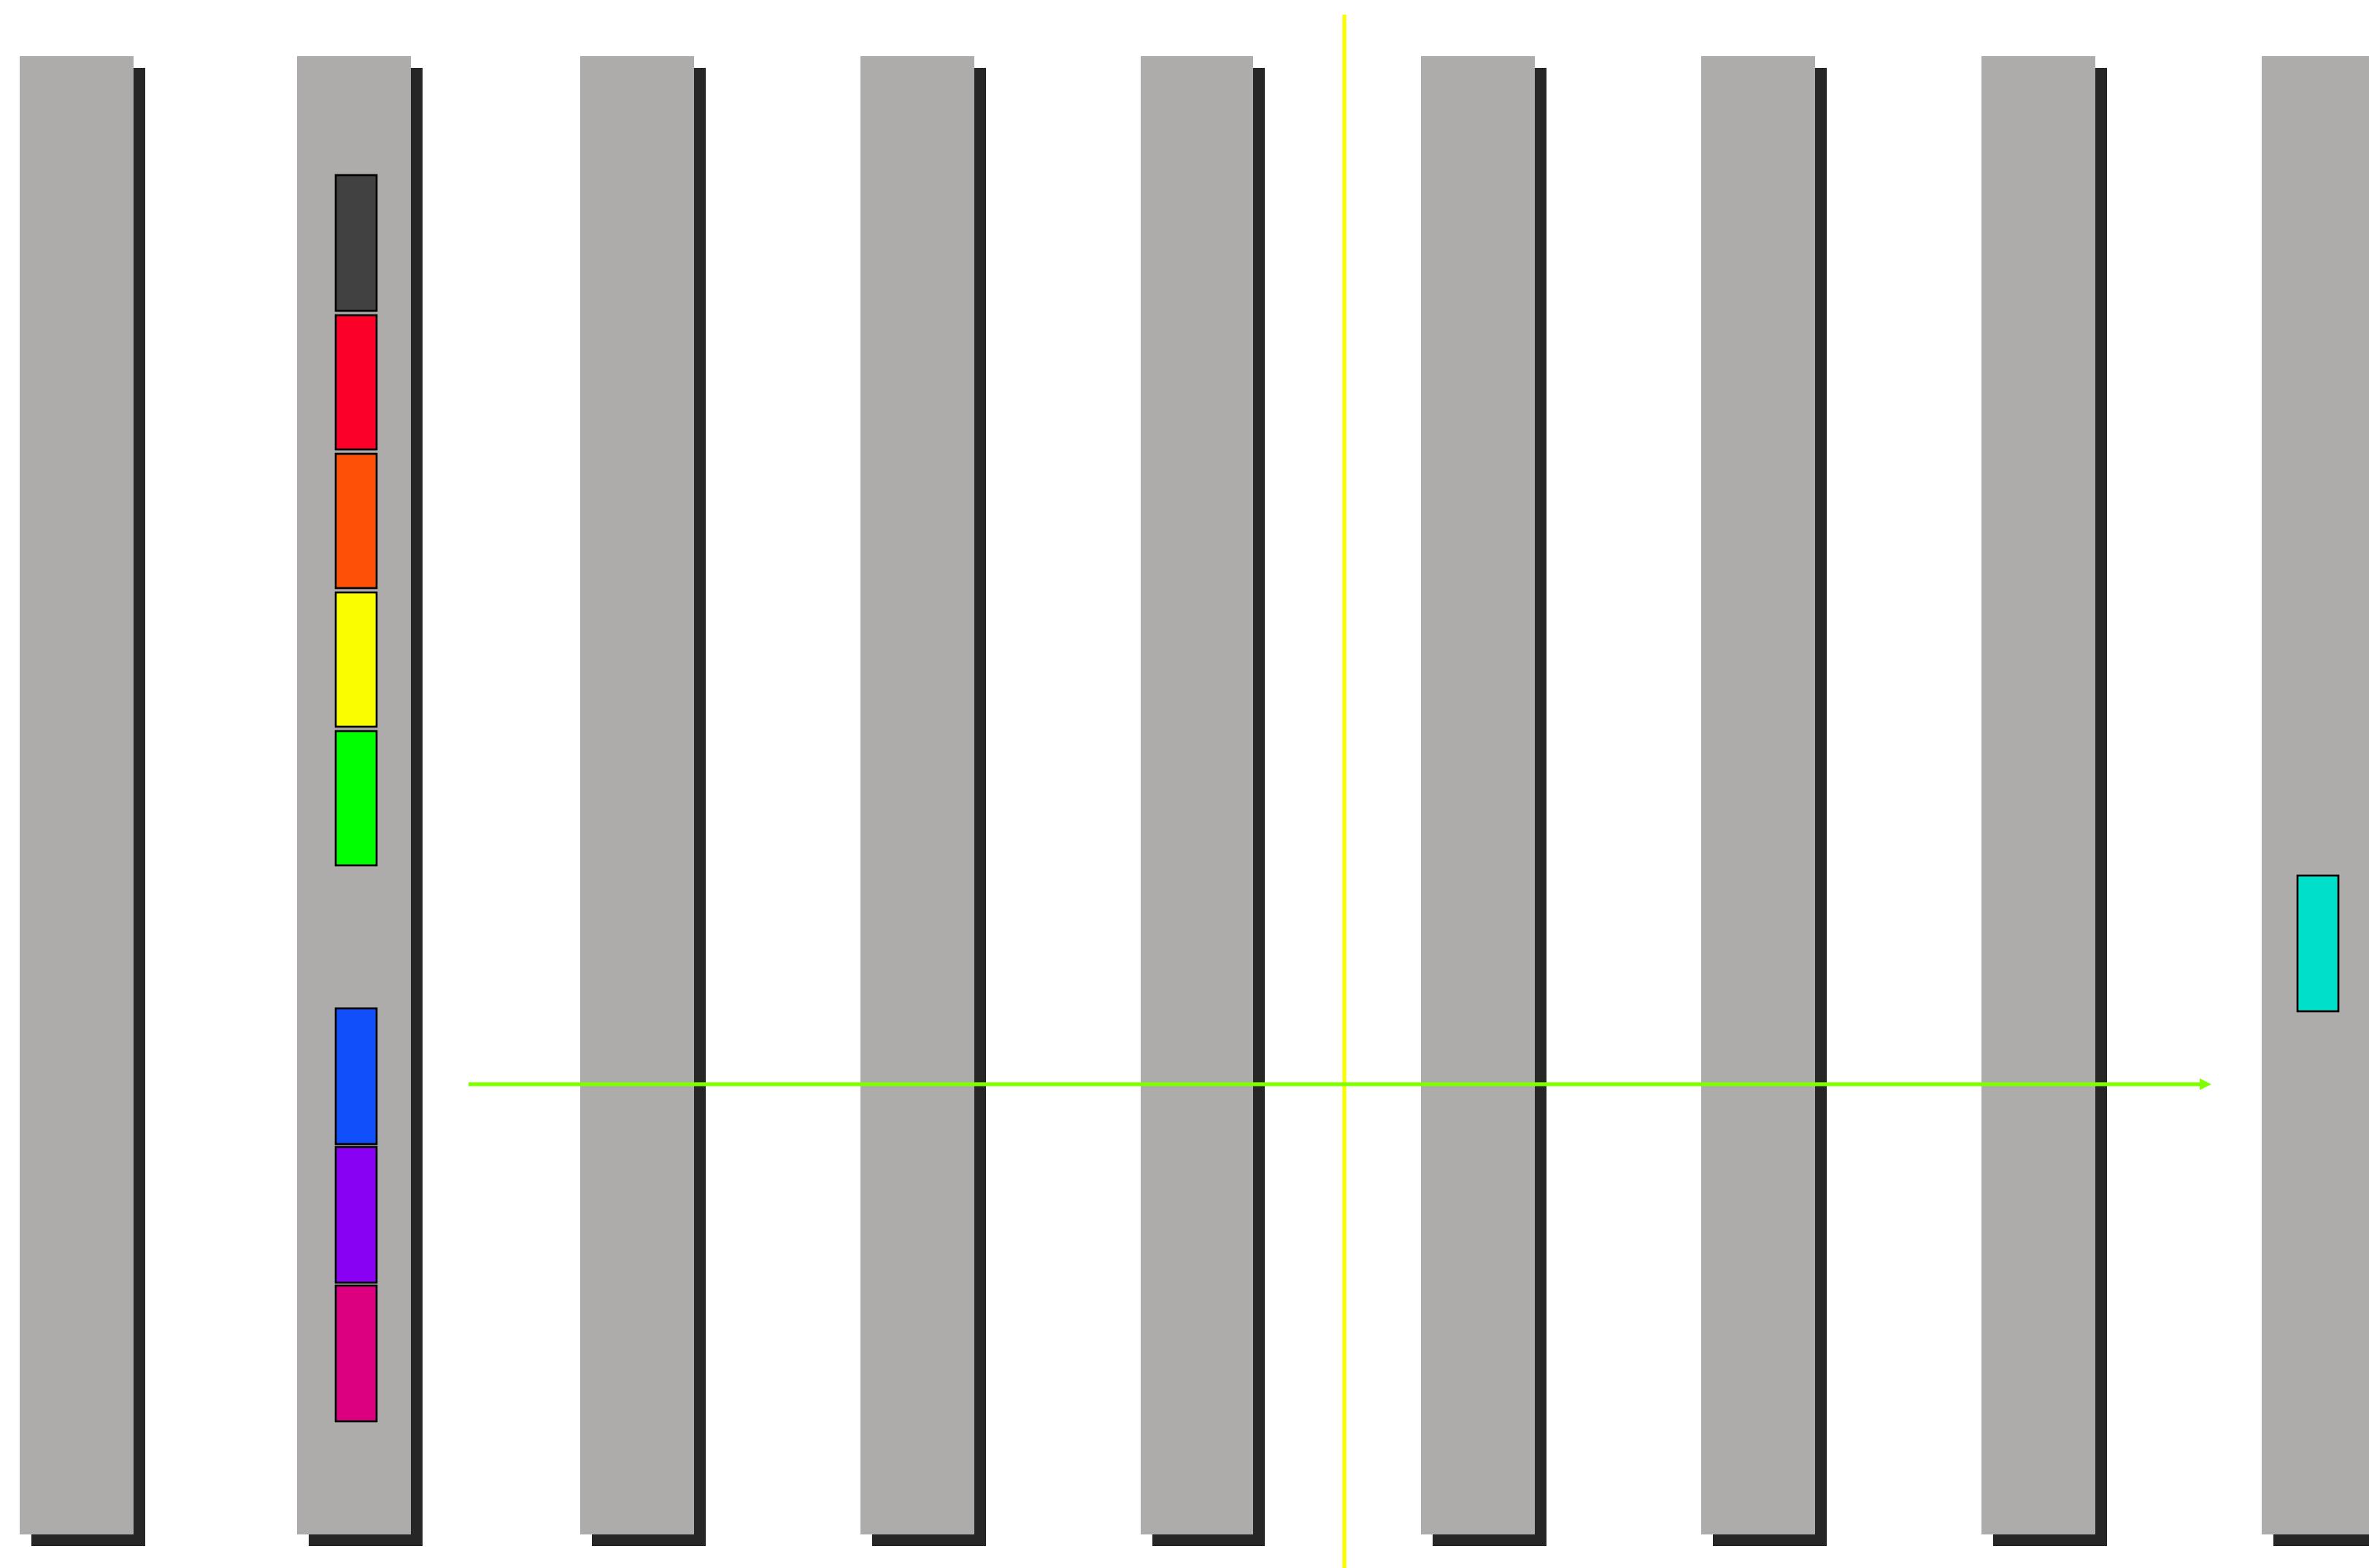


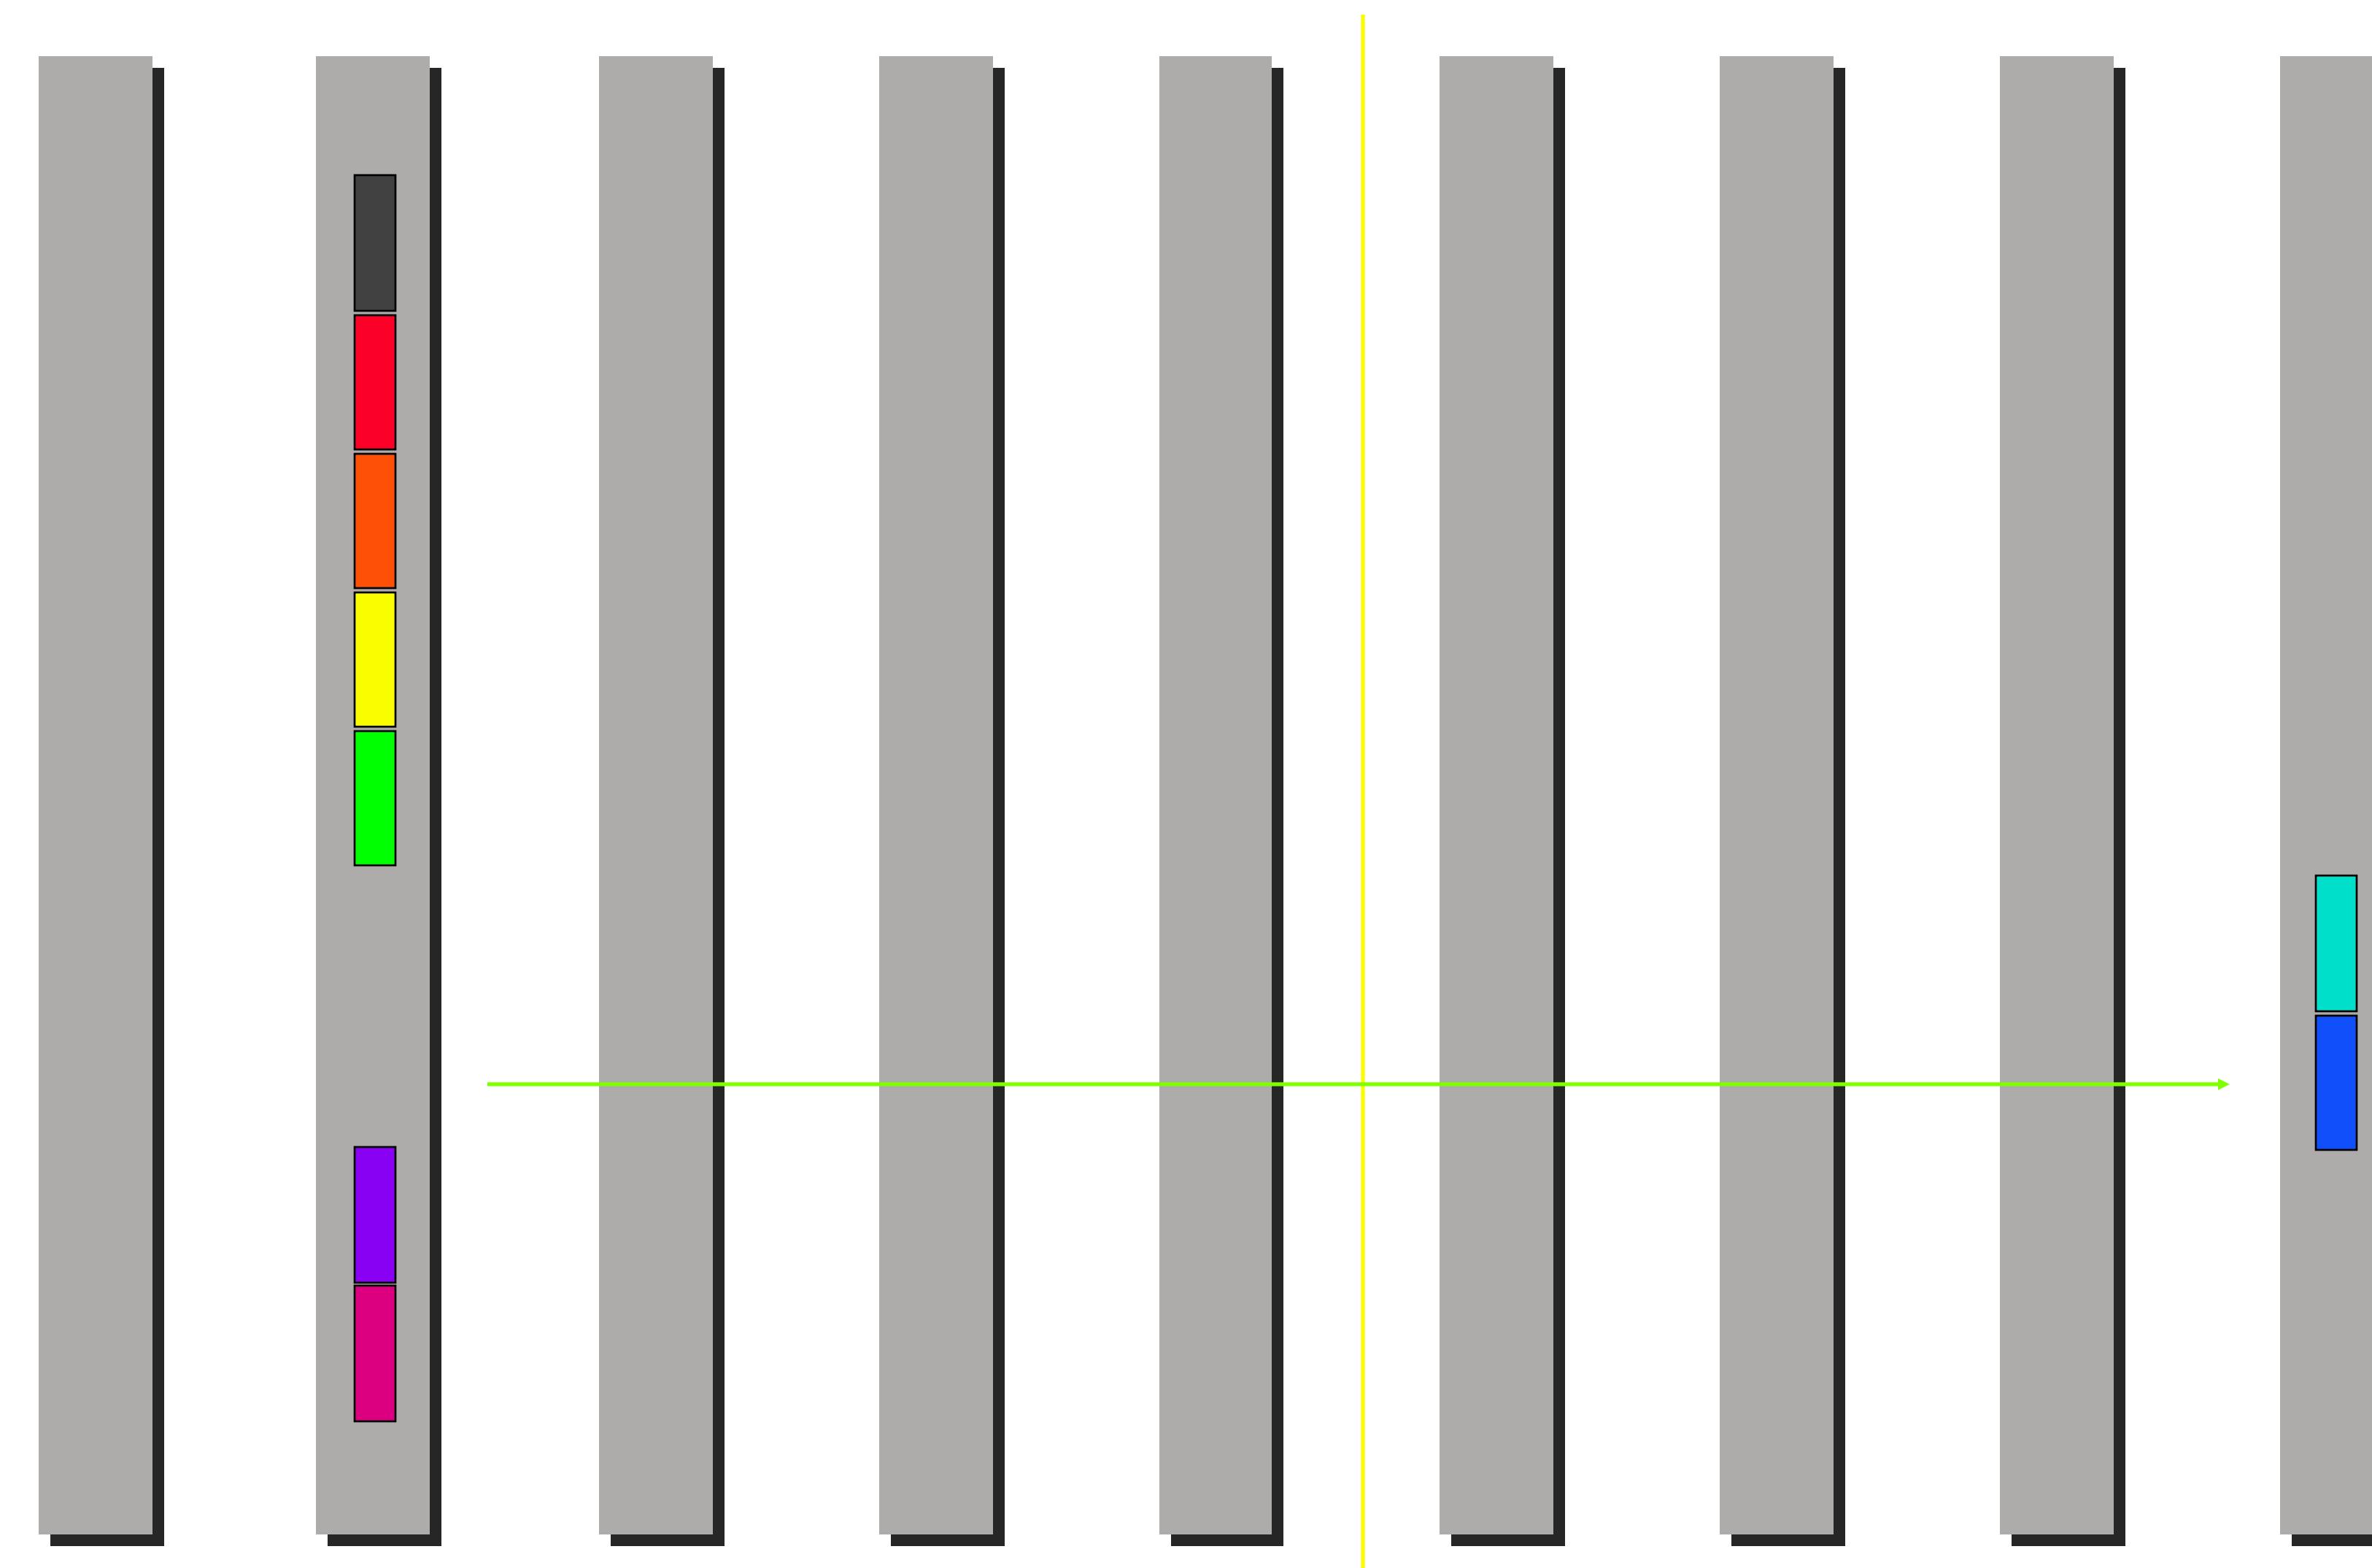
After

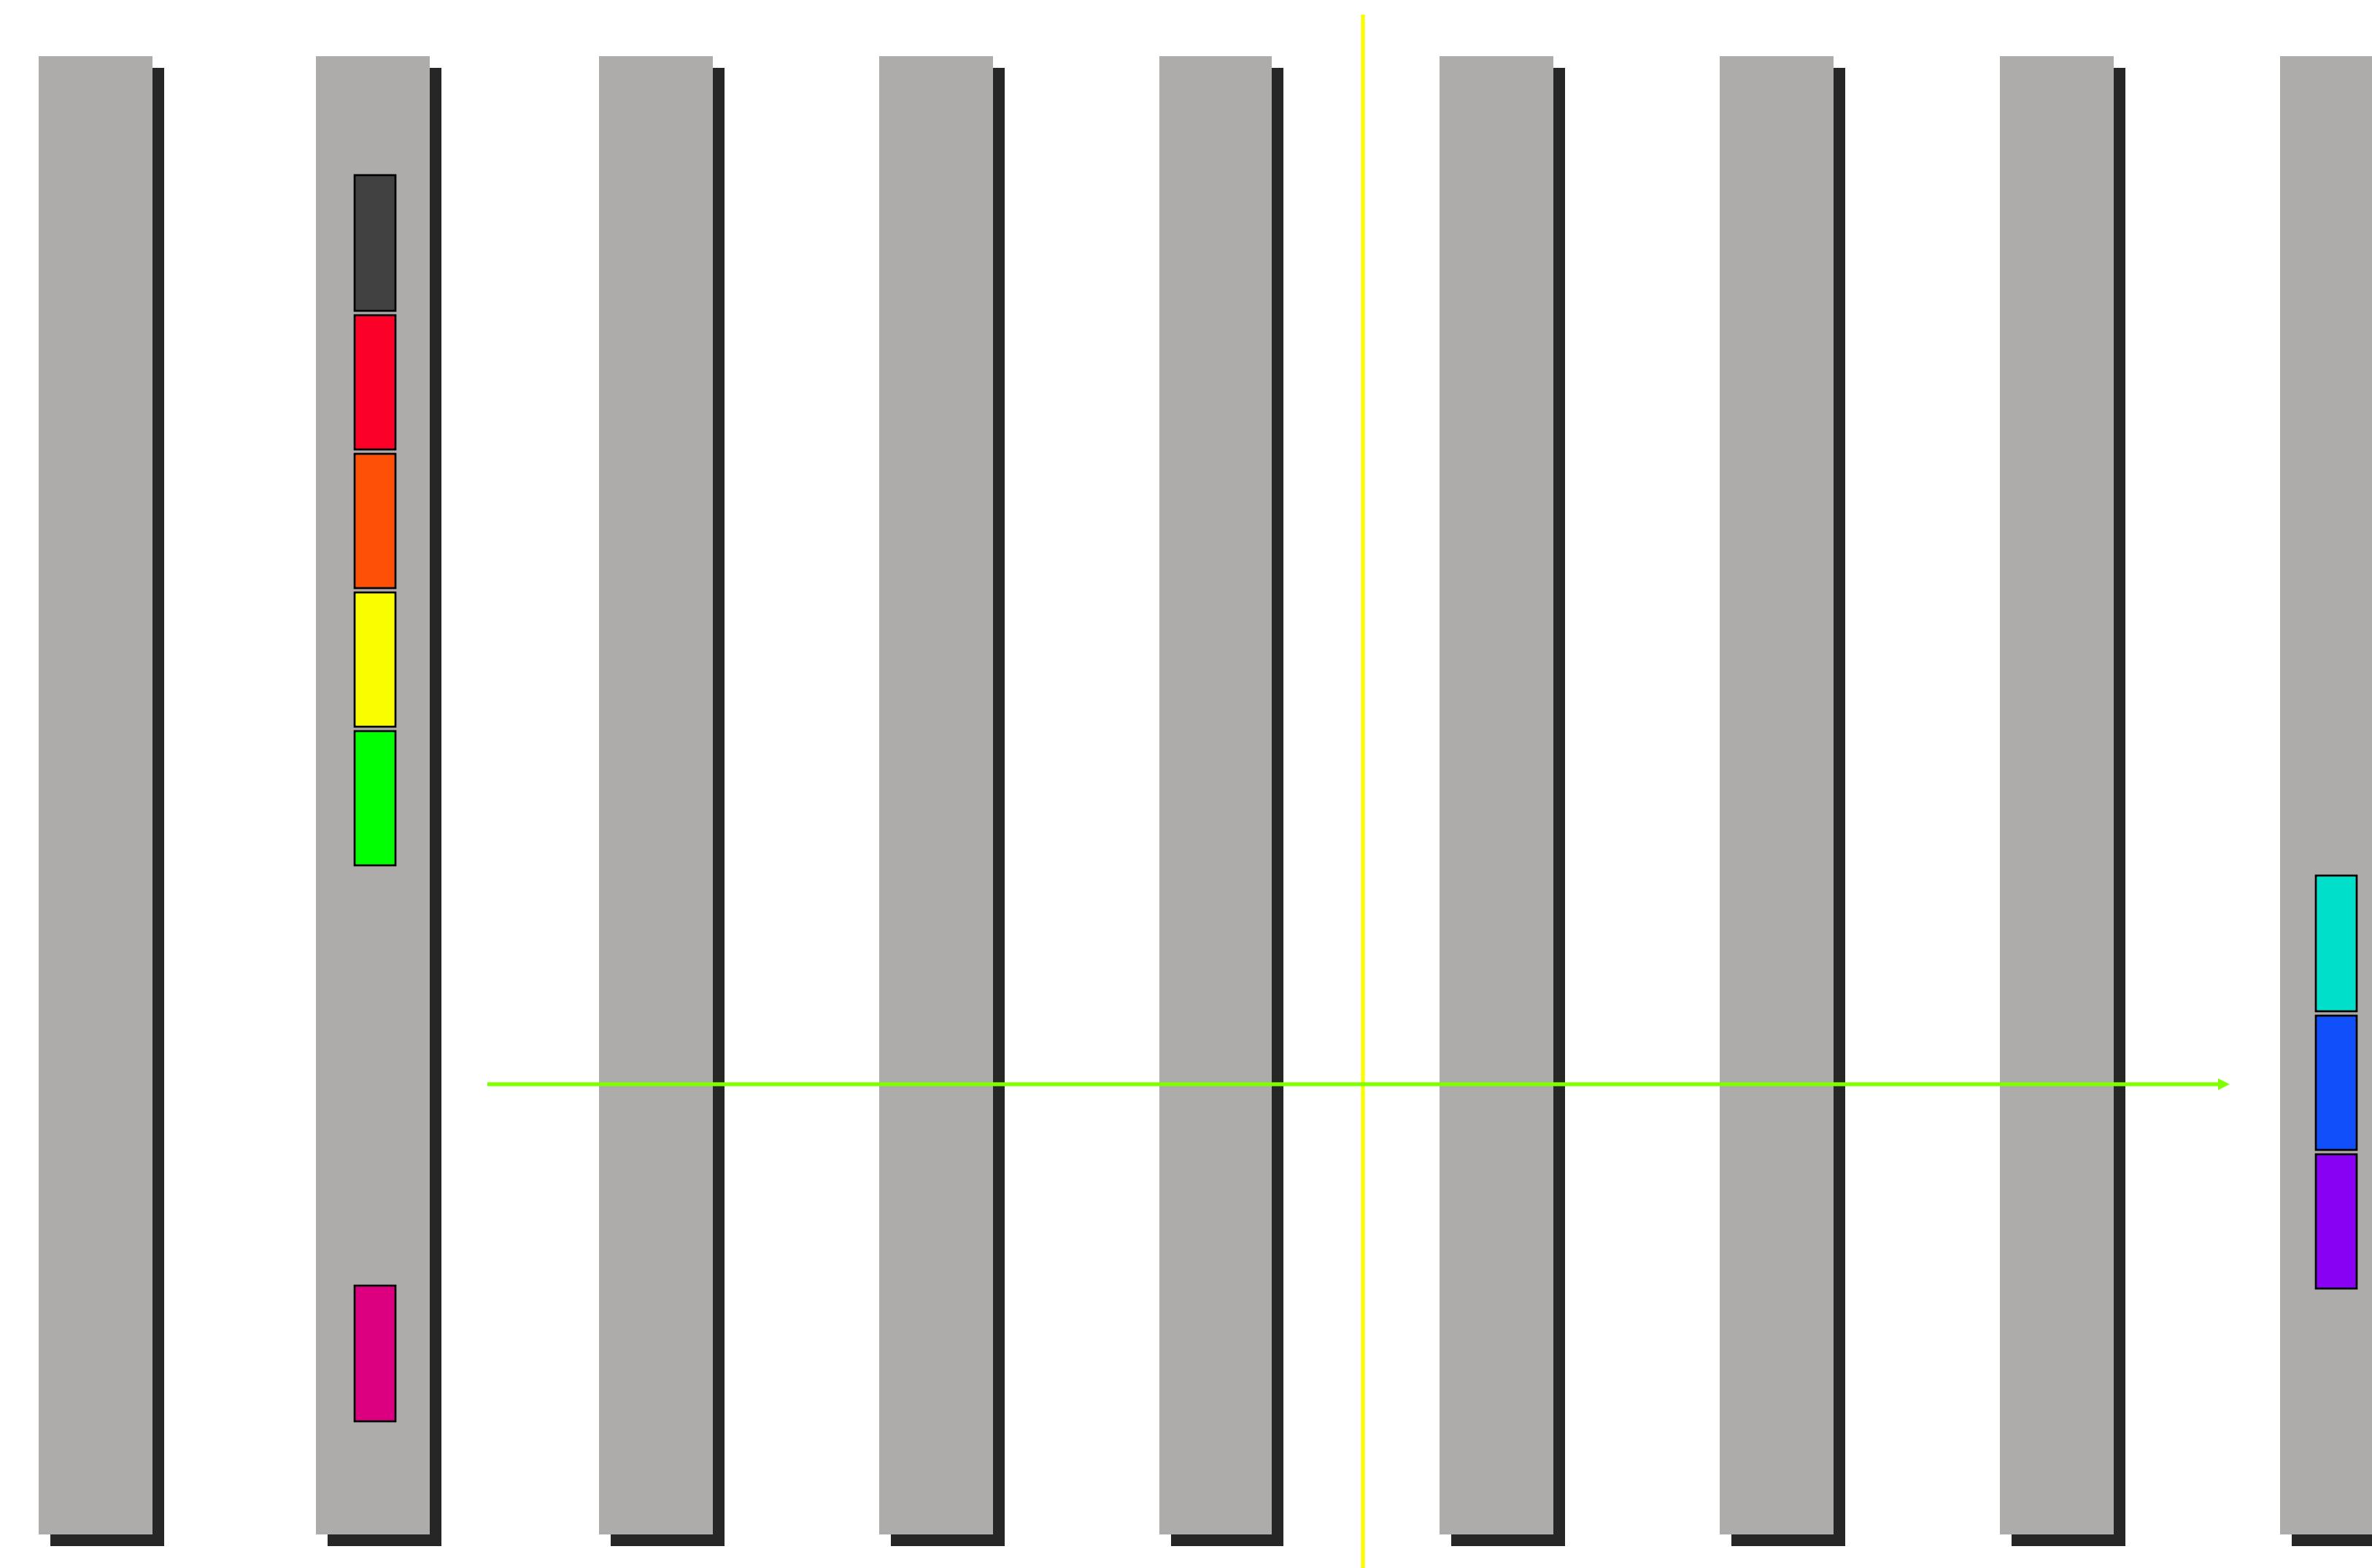


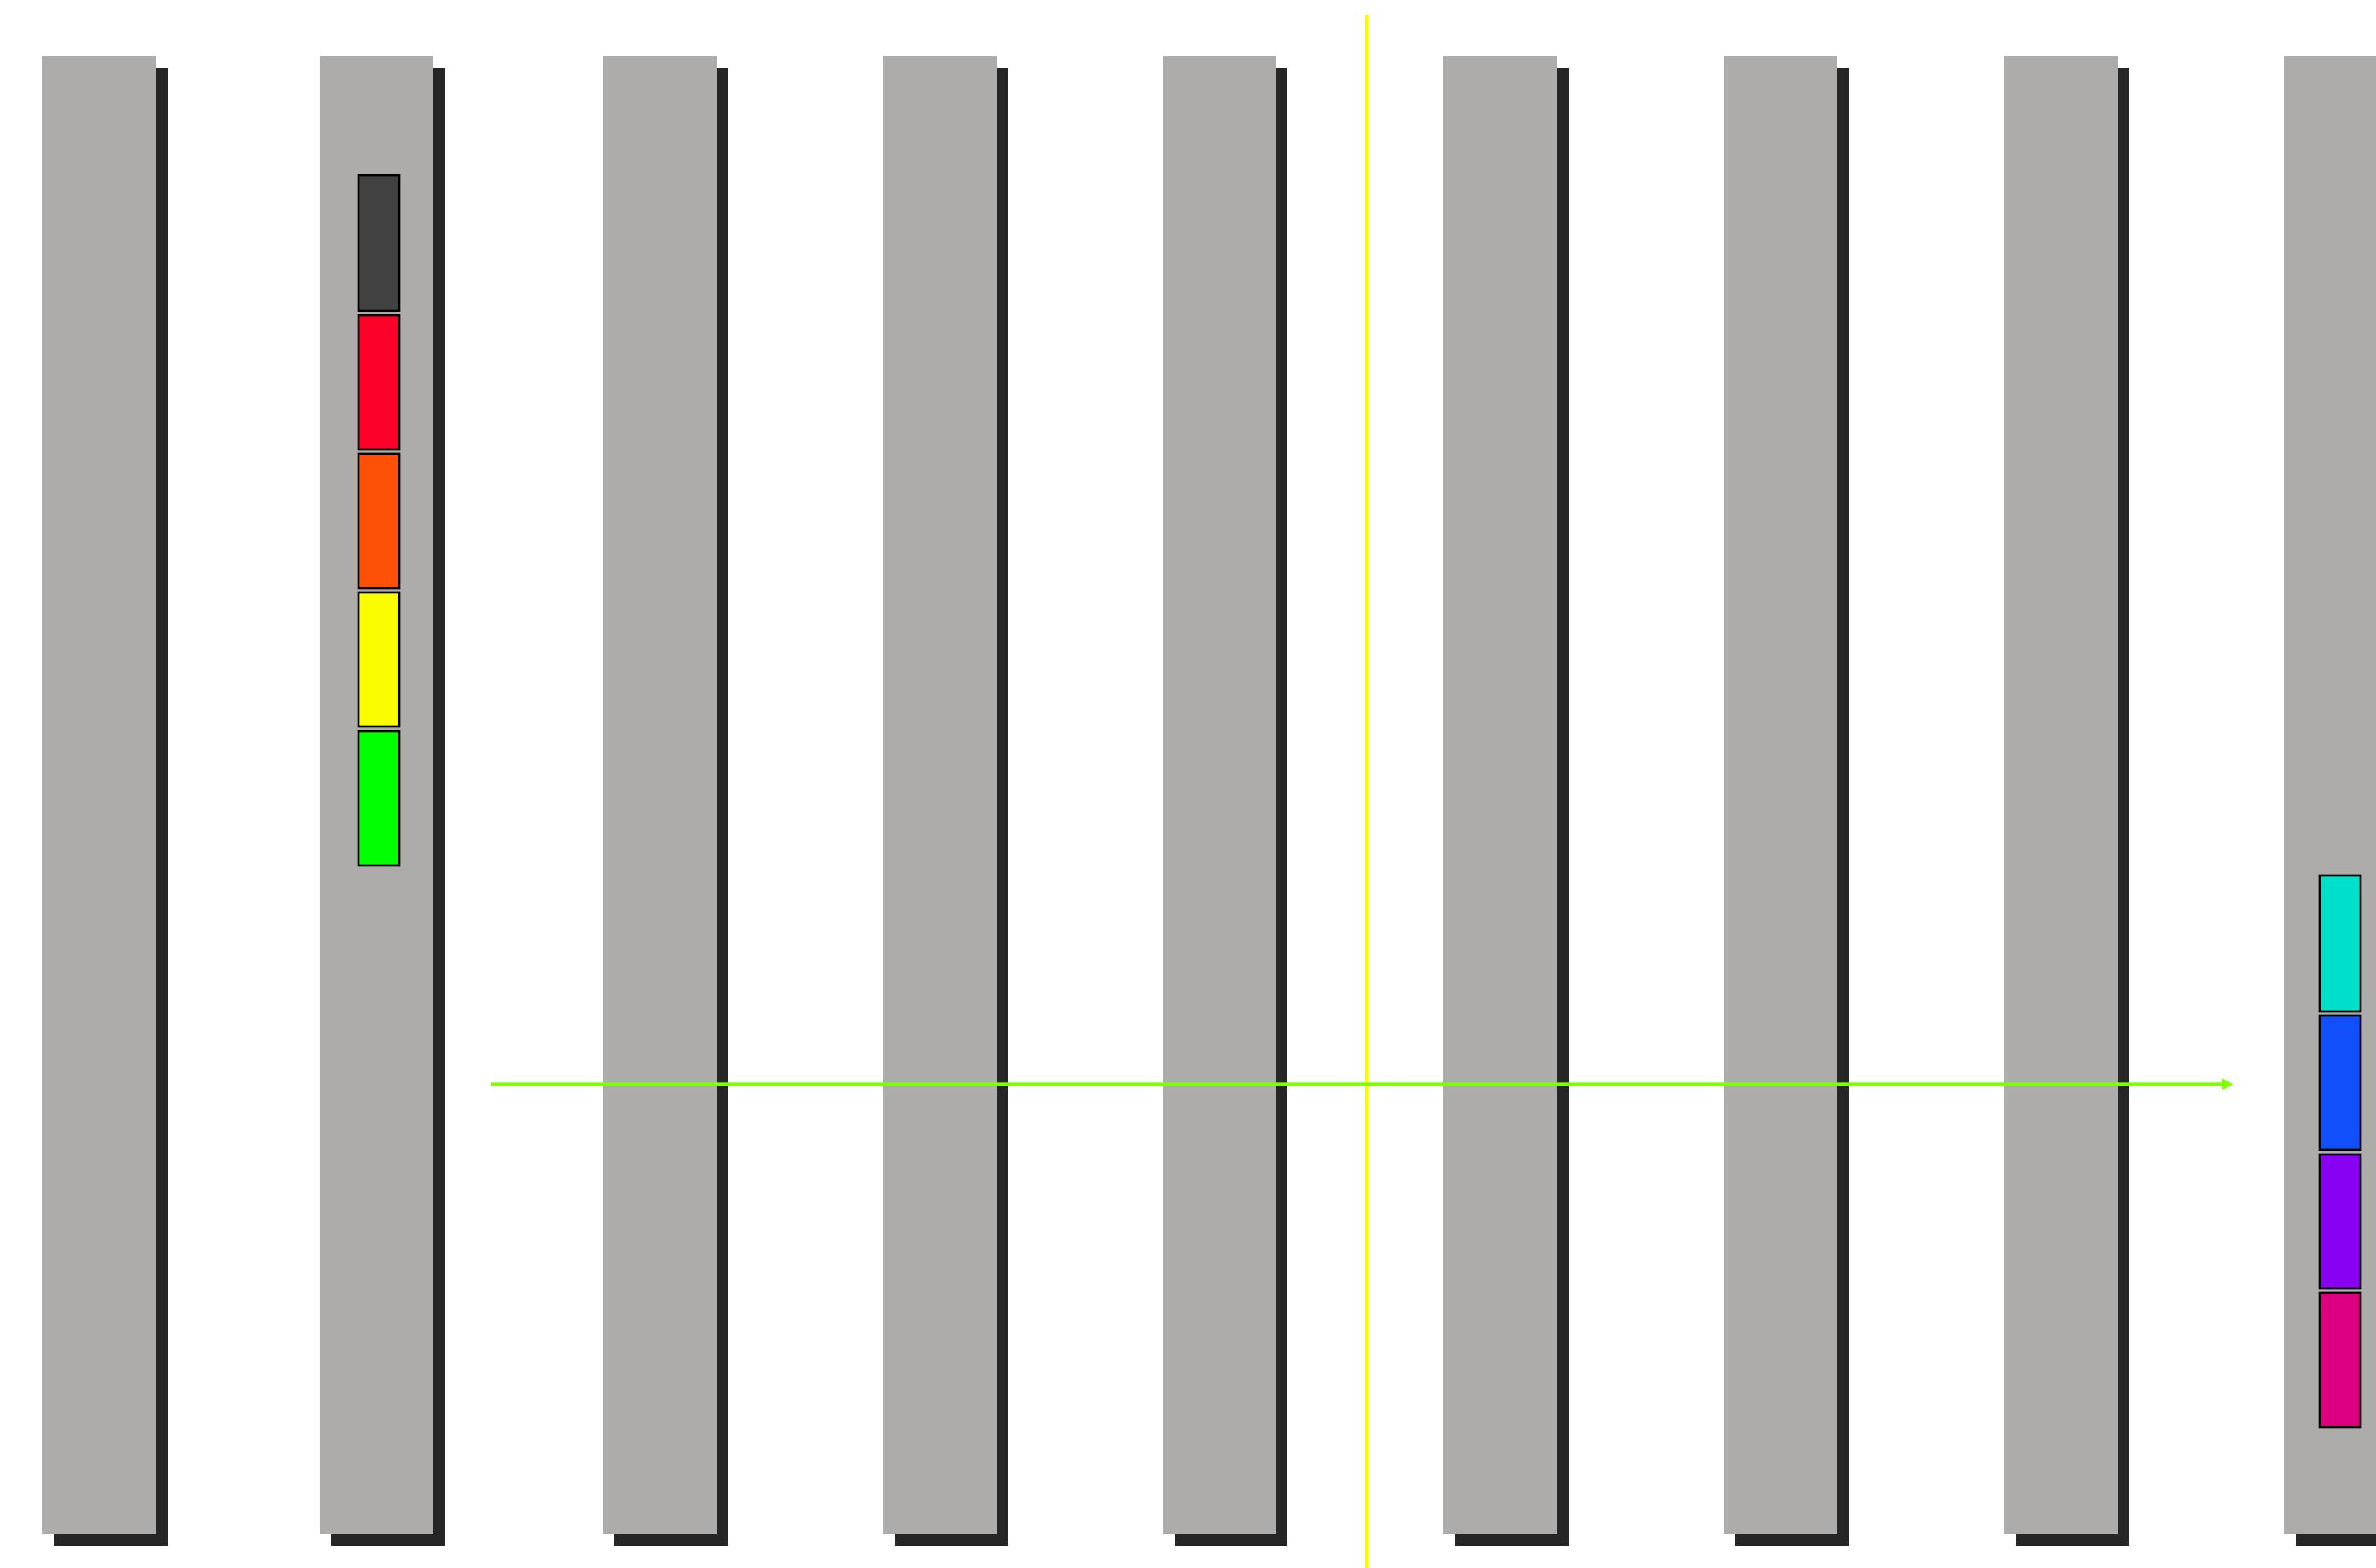


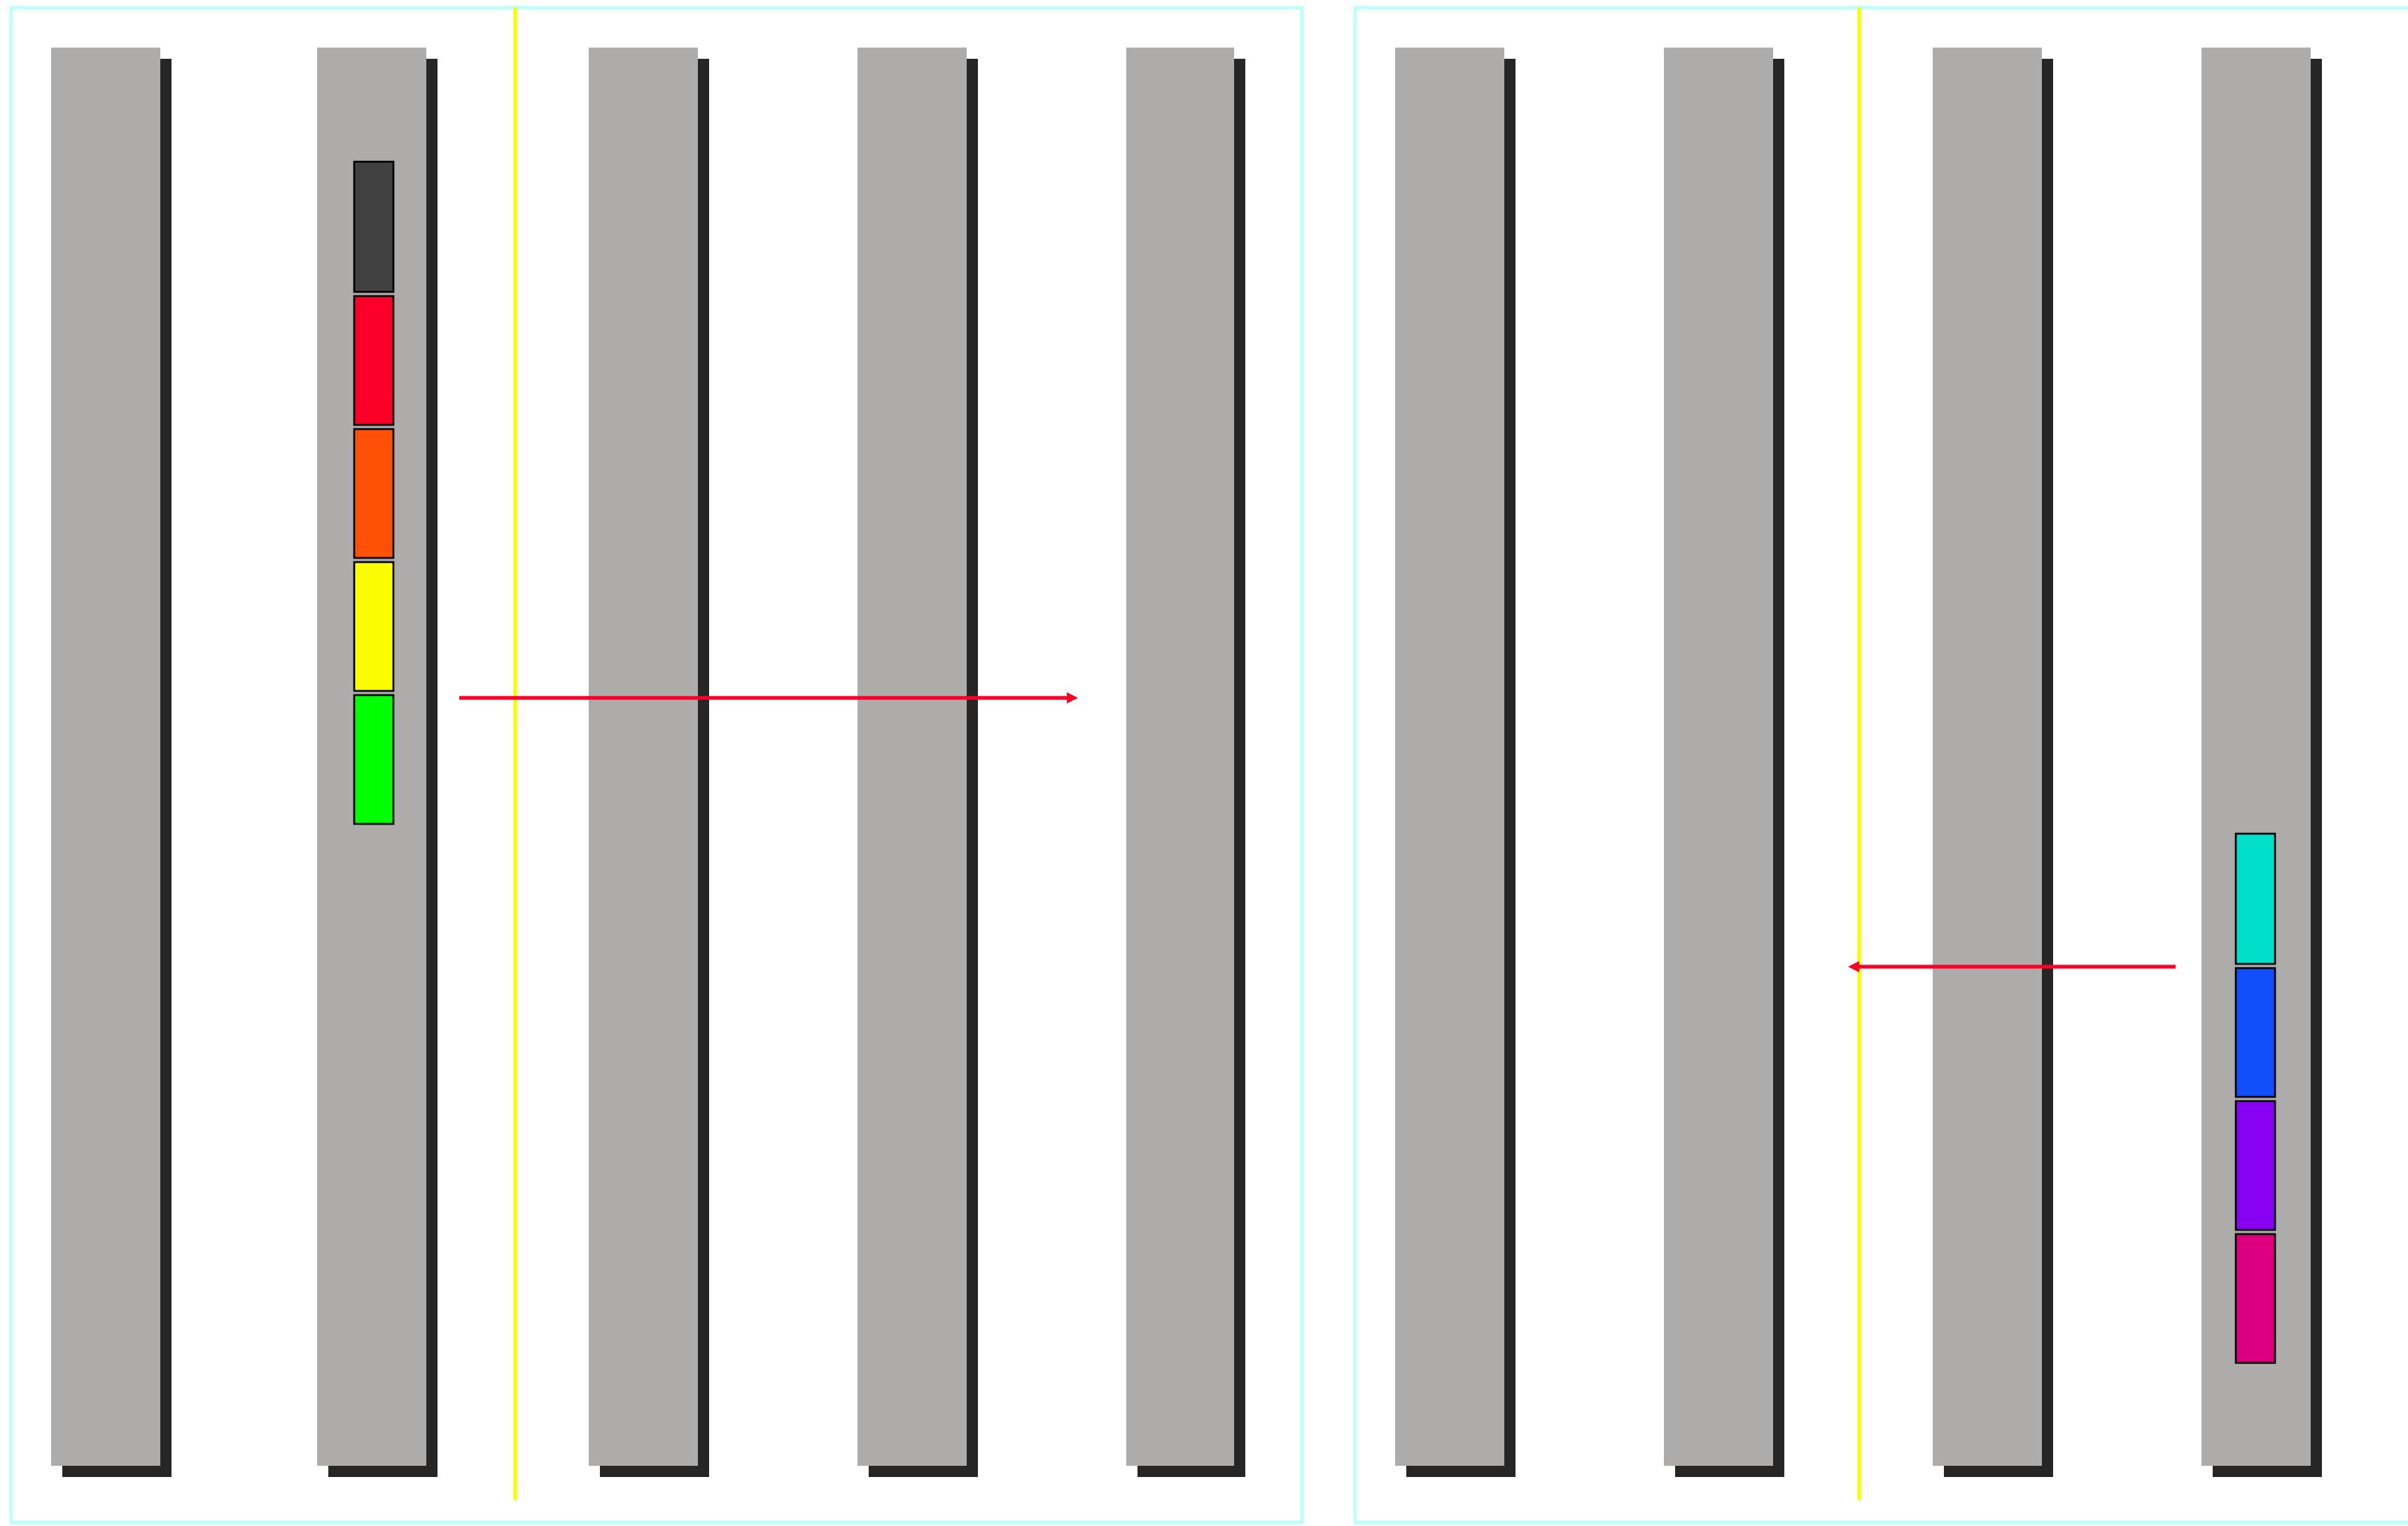


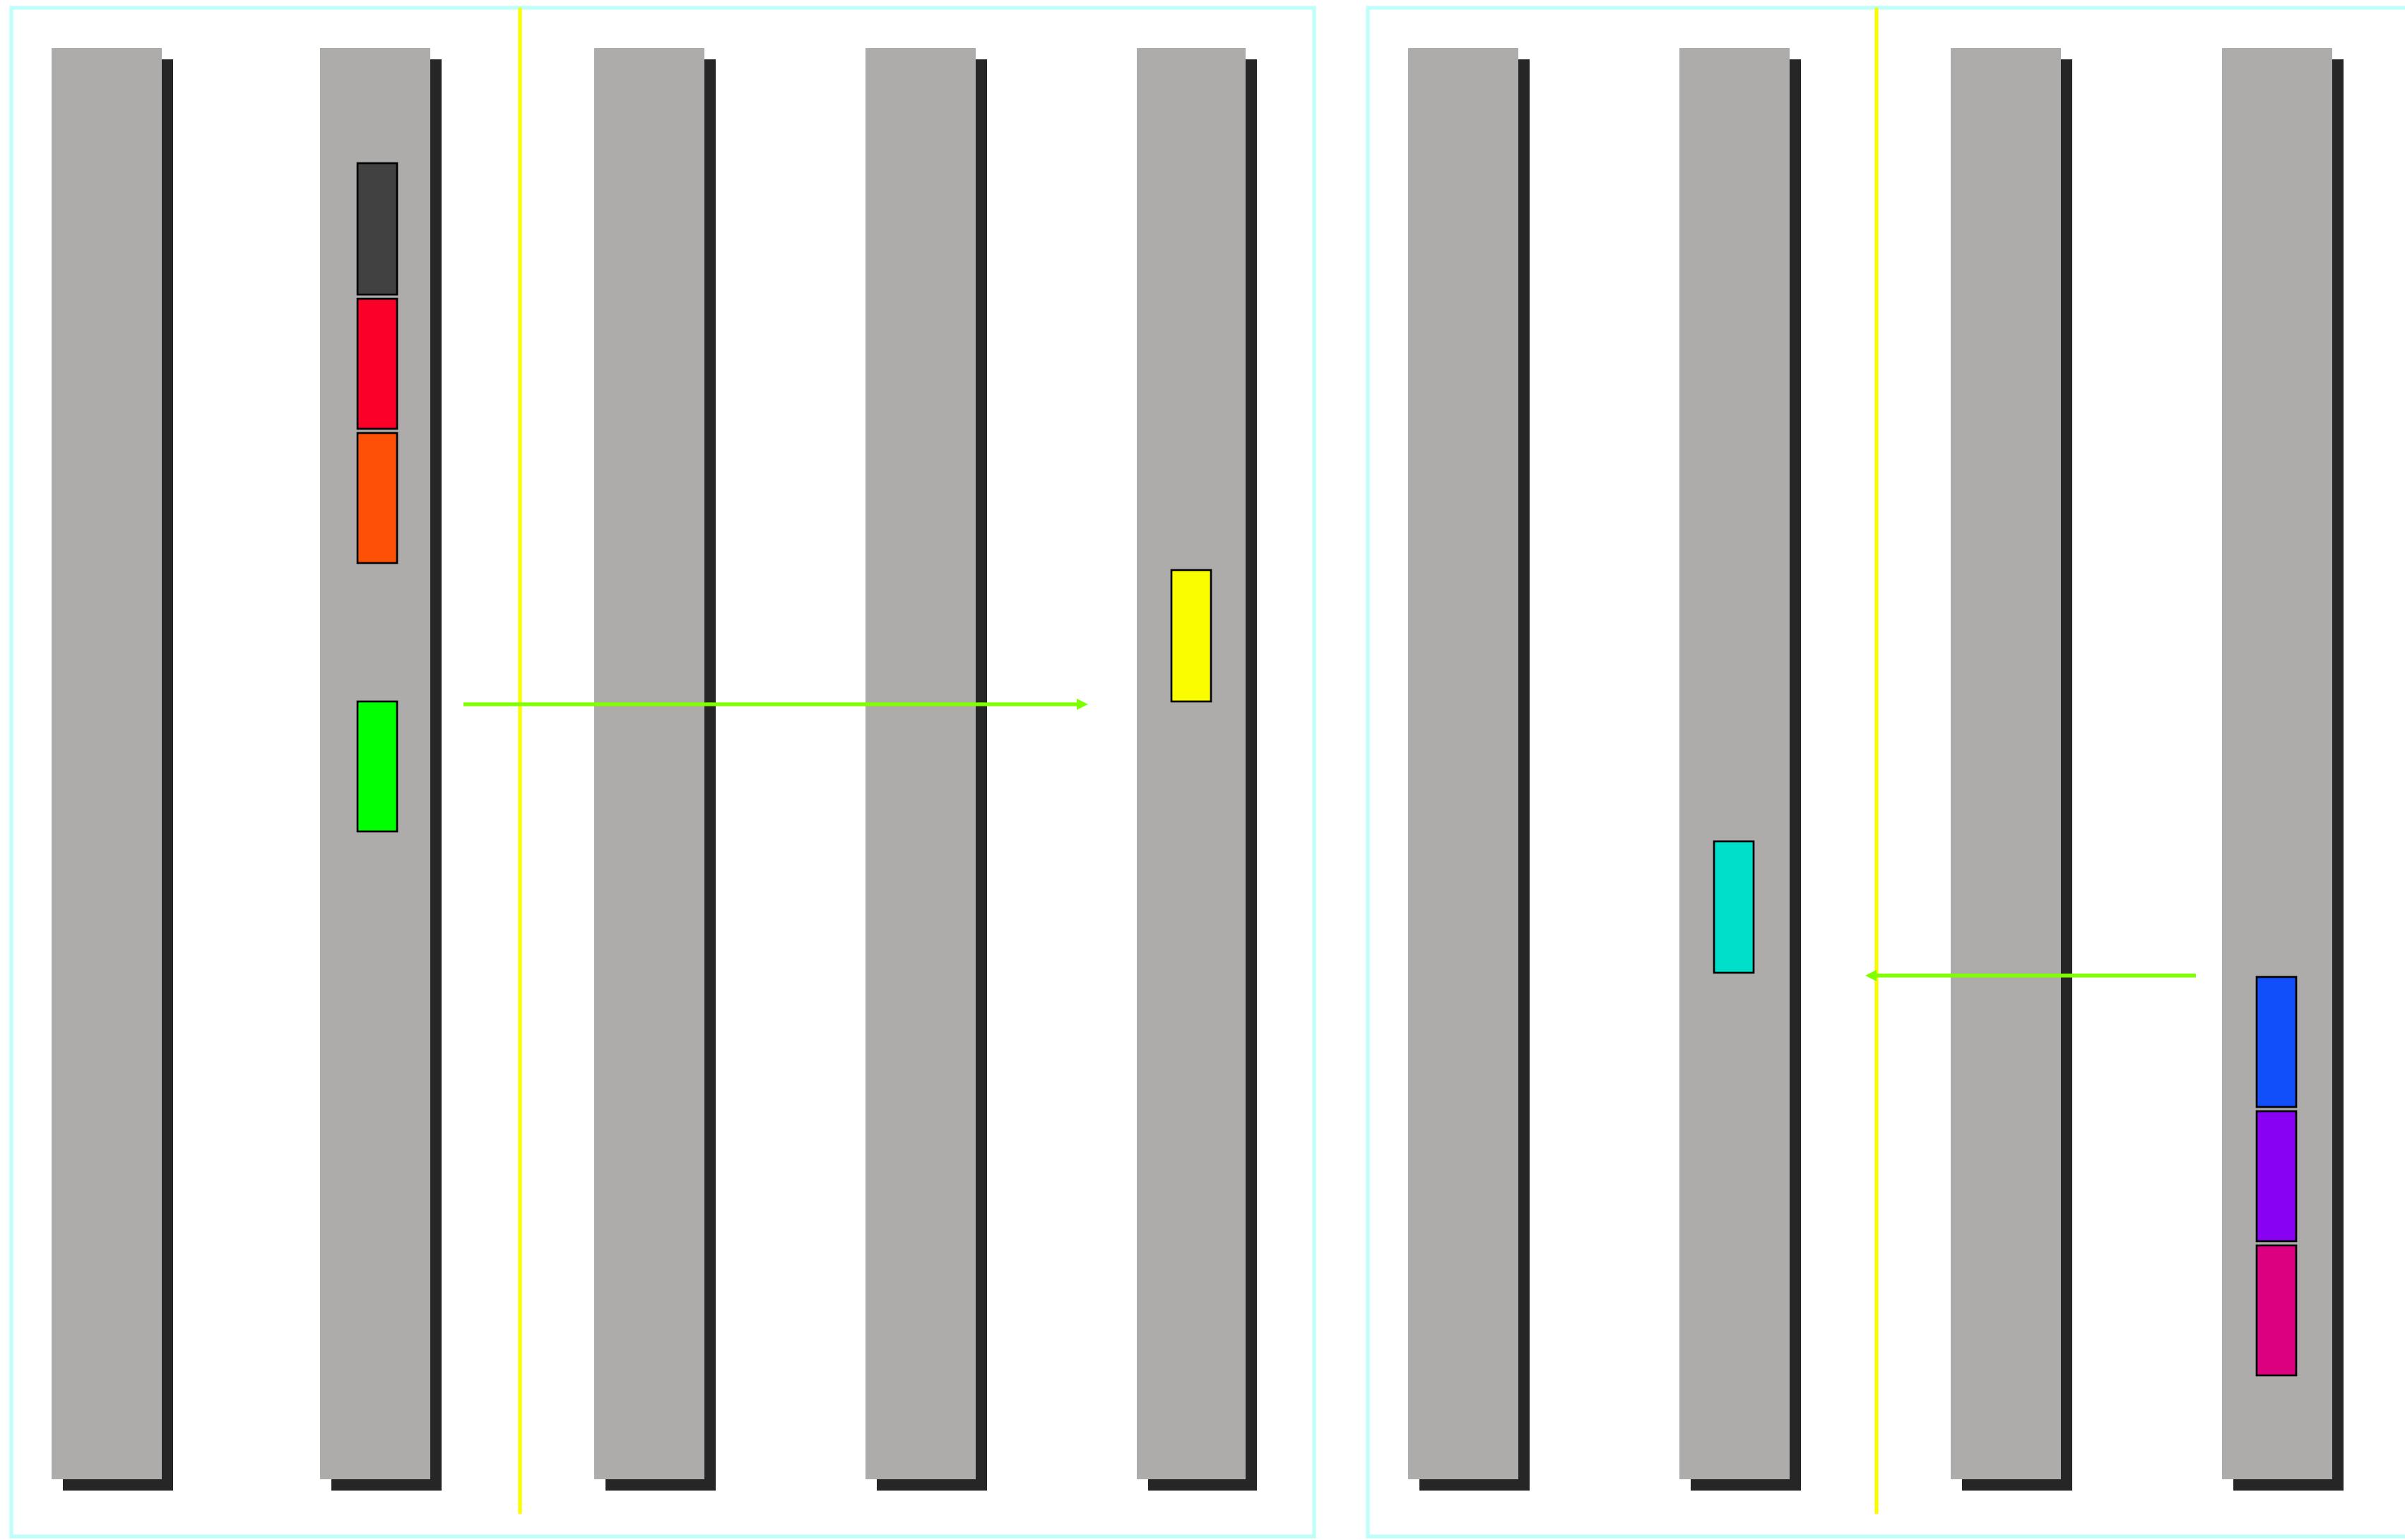


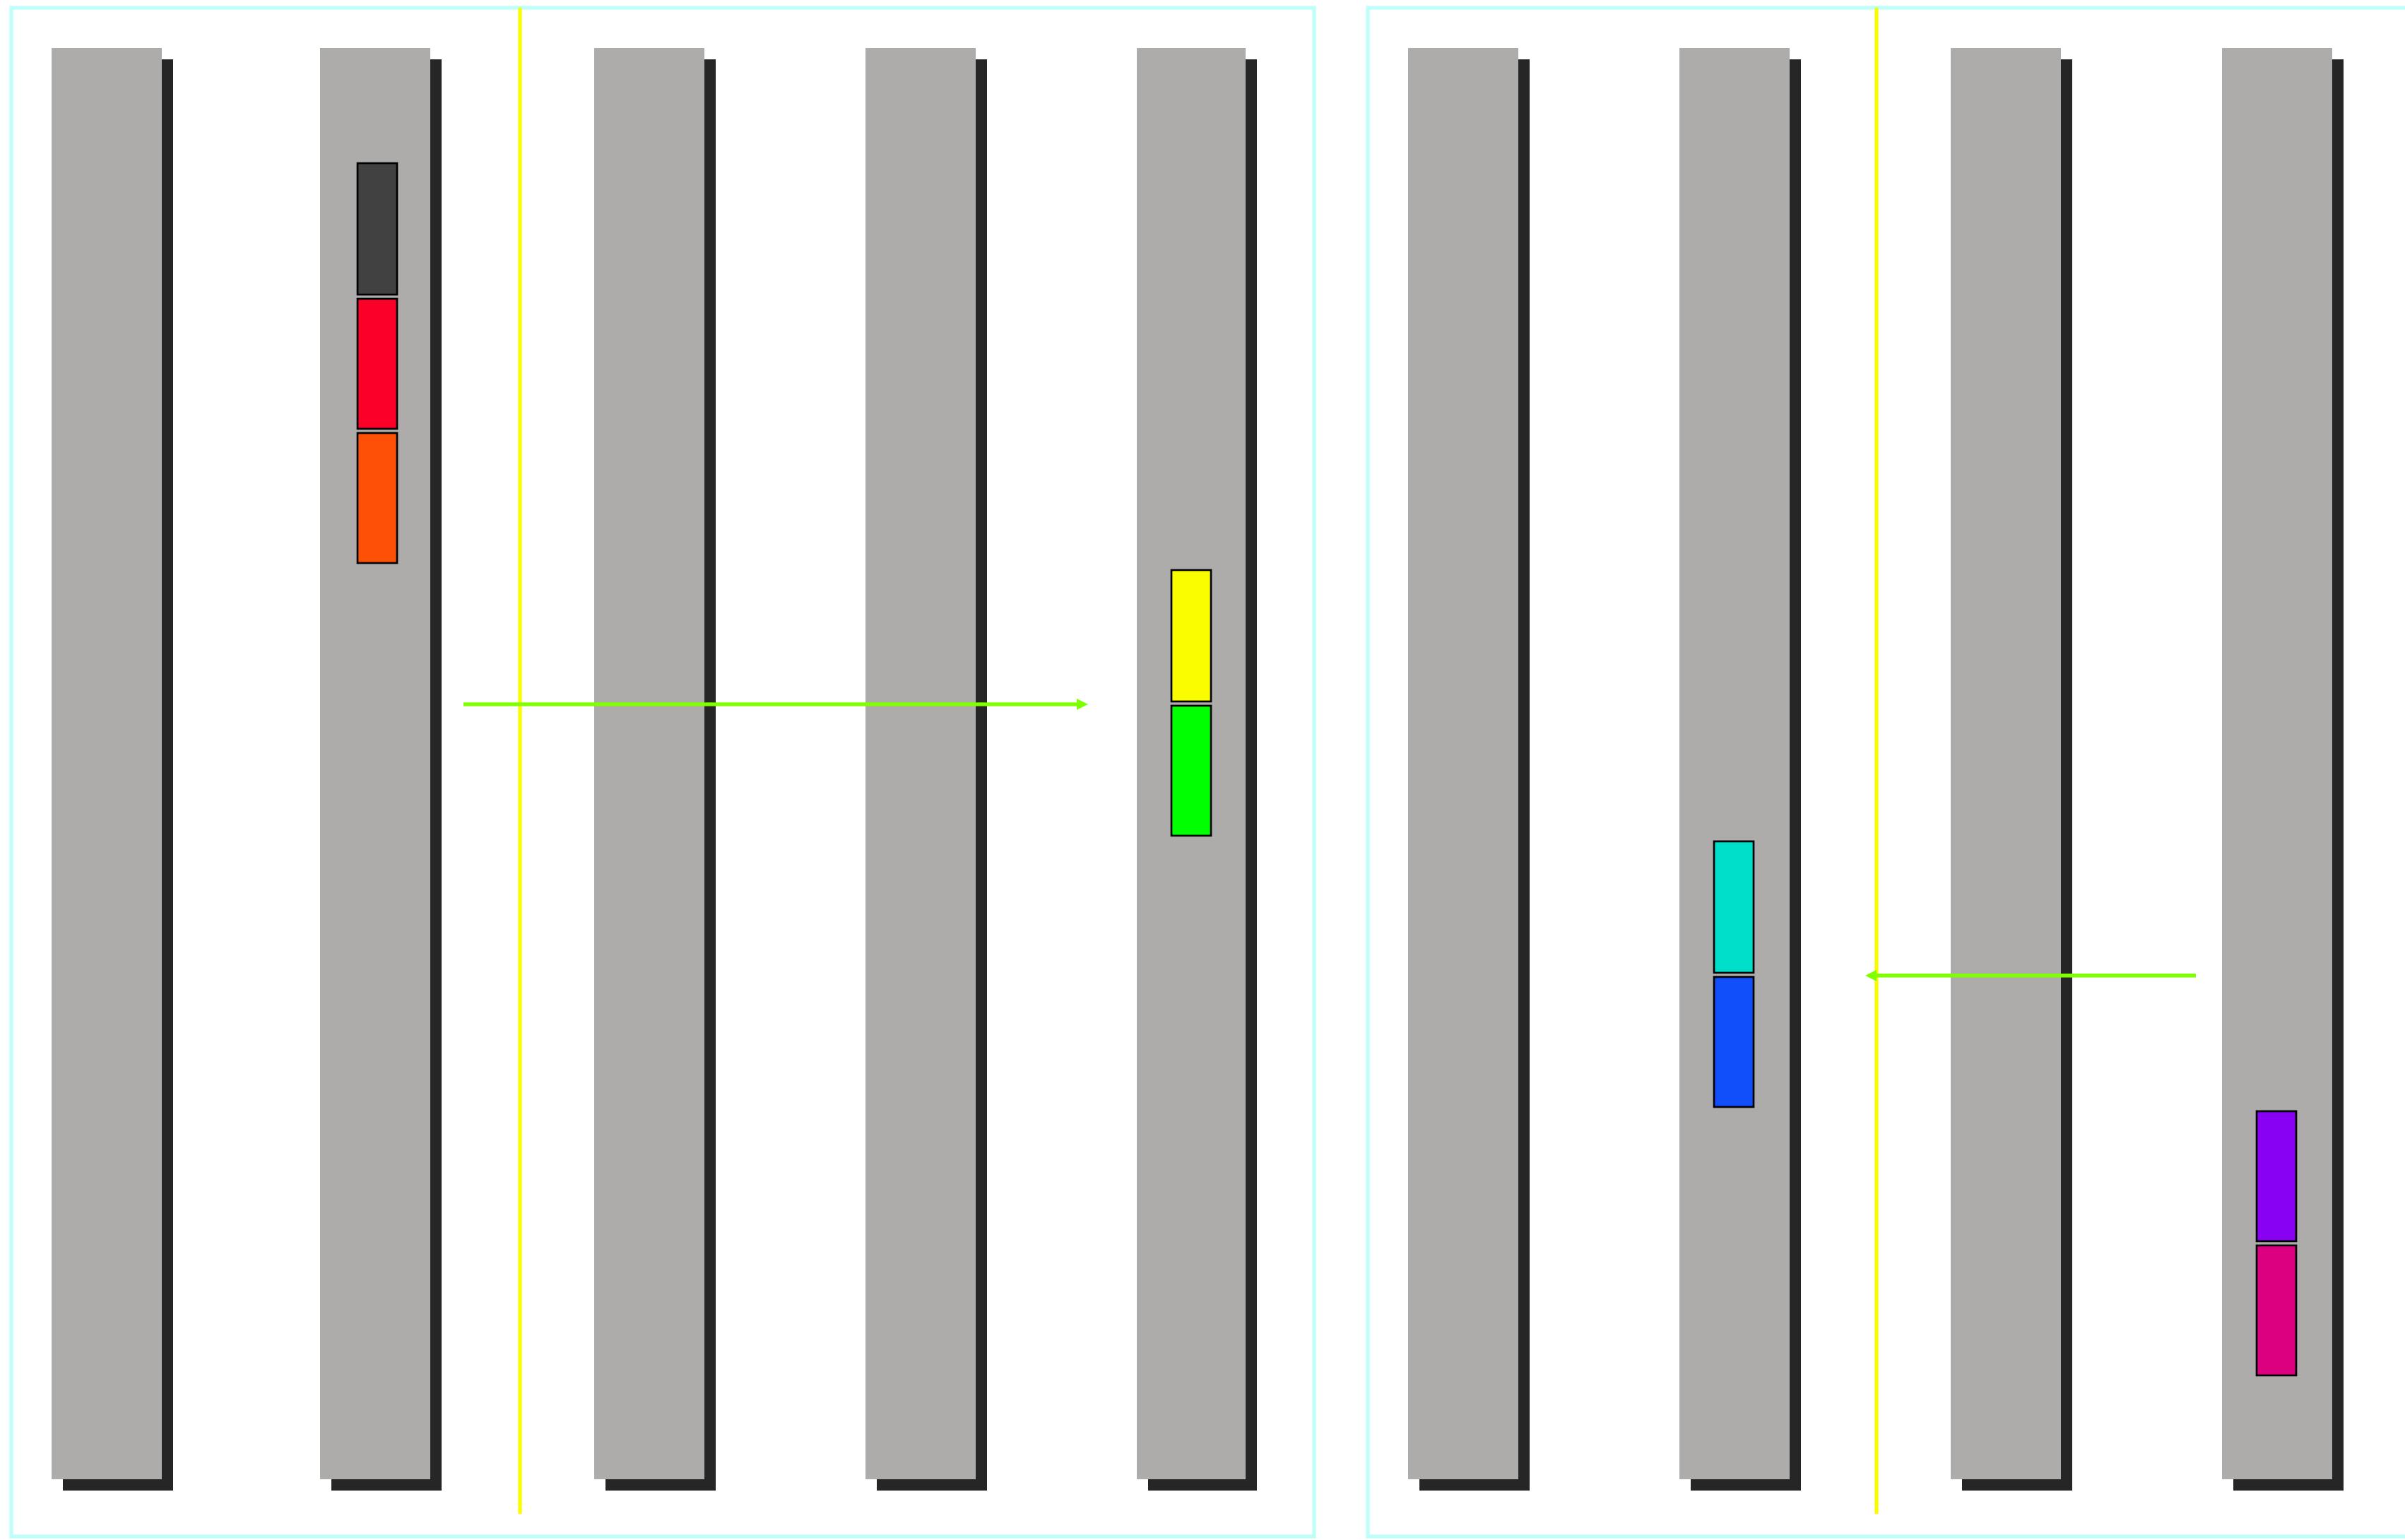


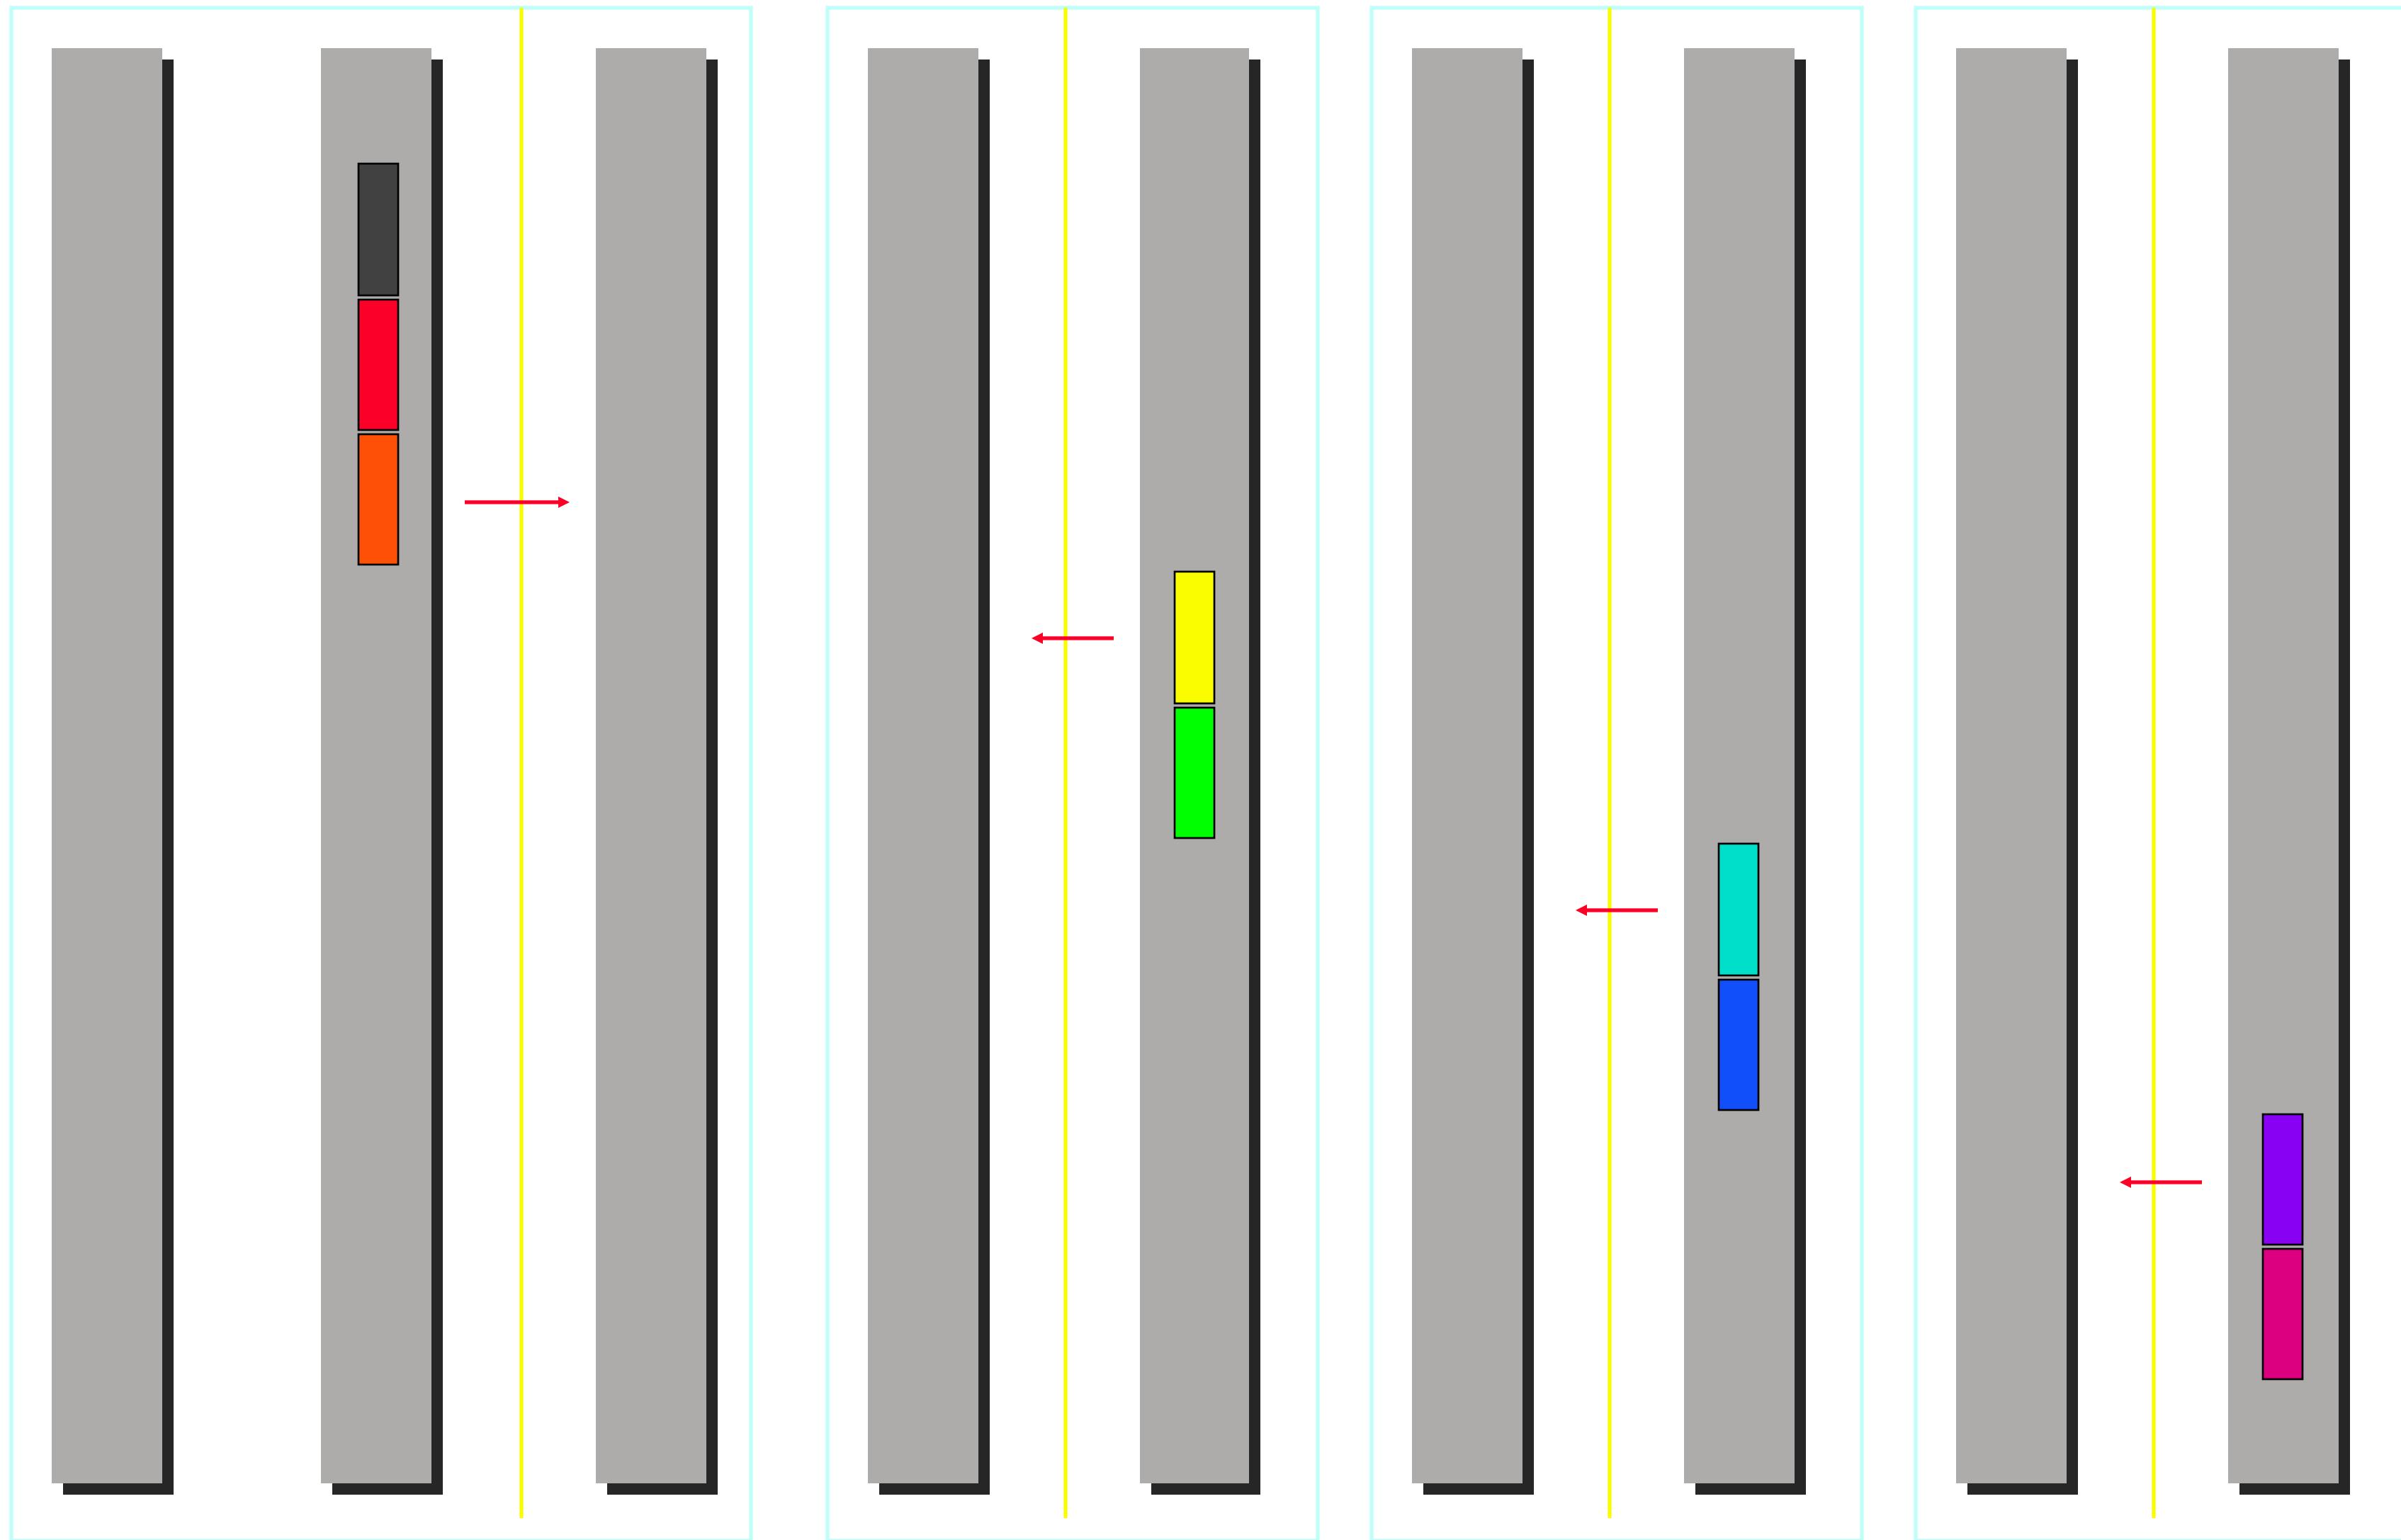


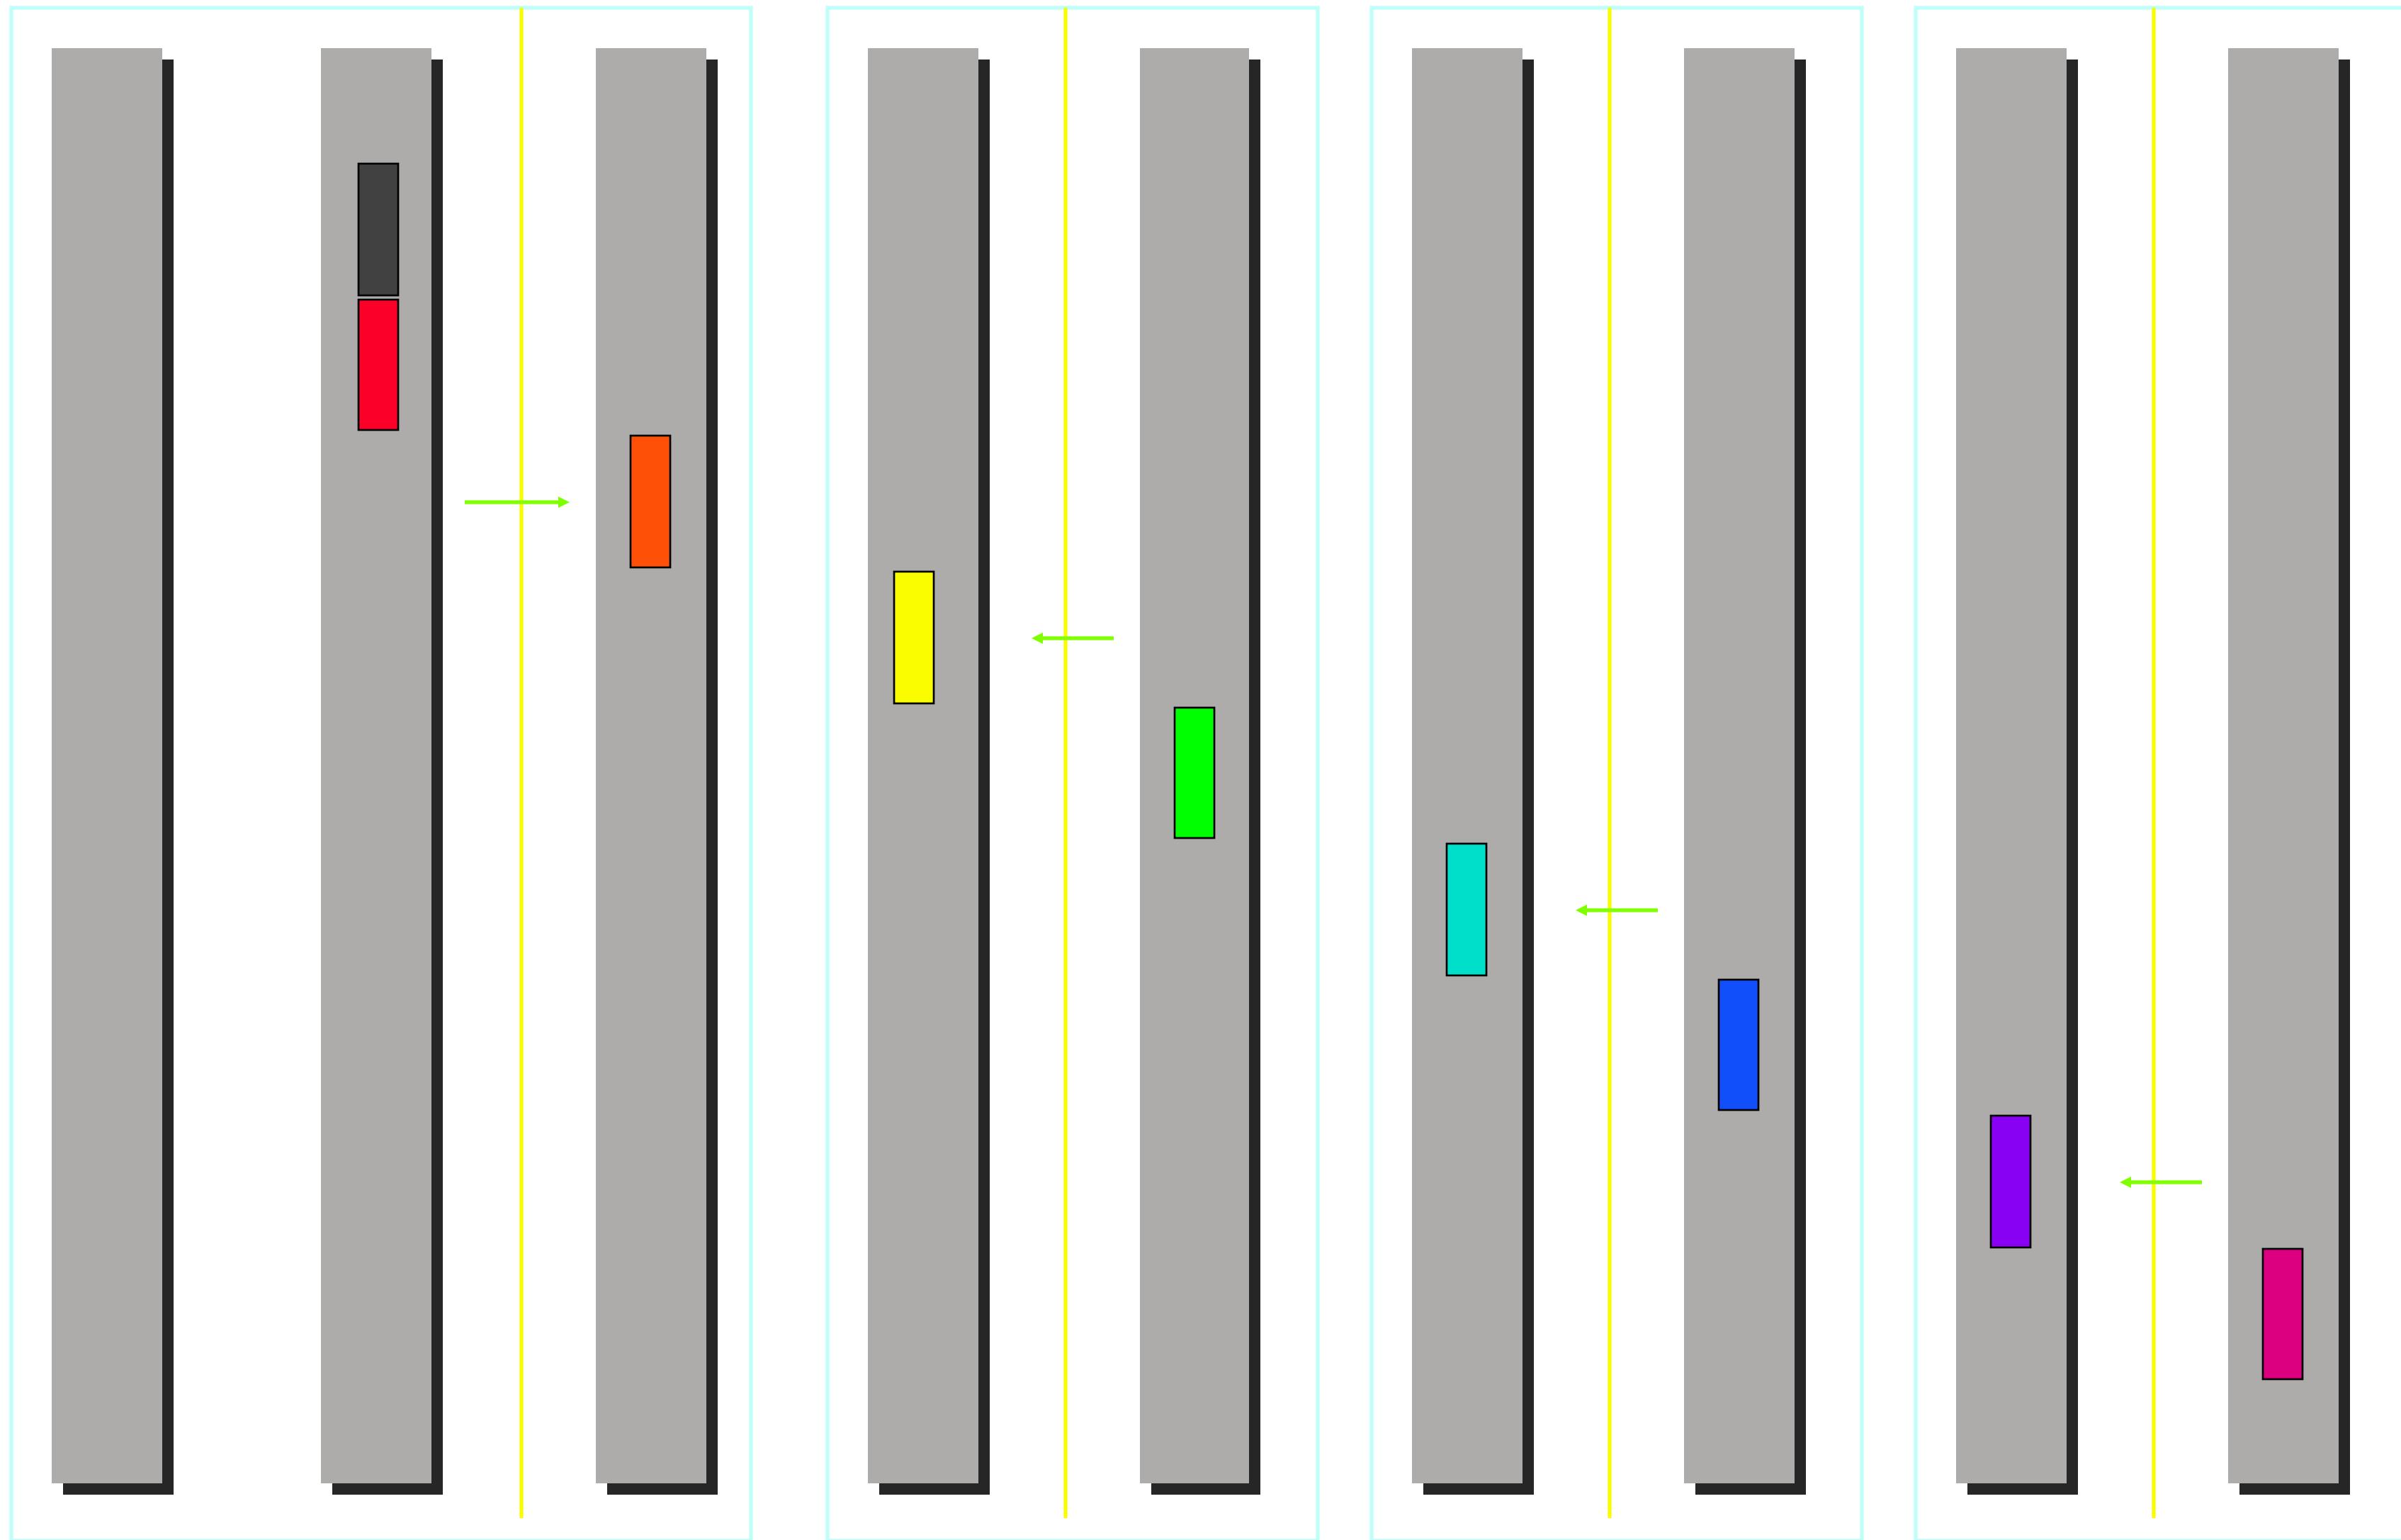


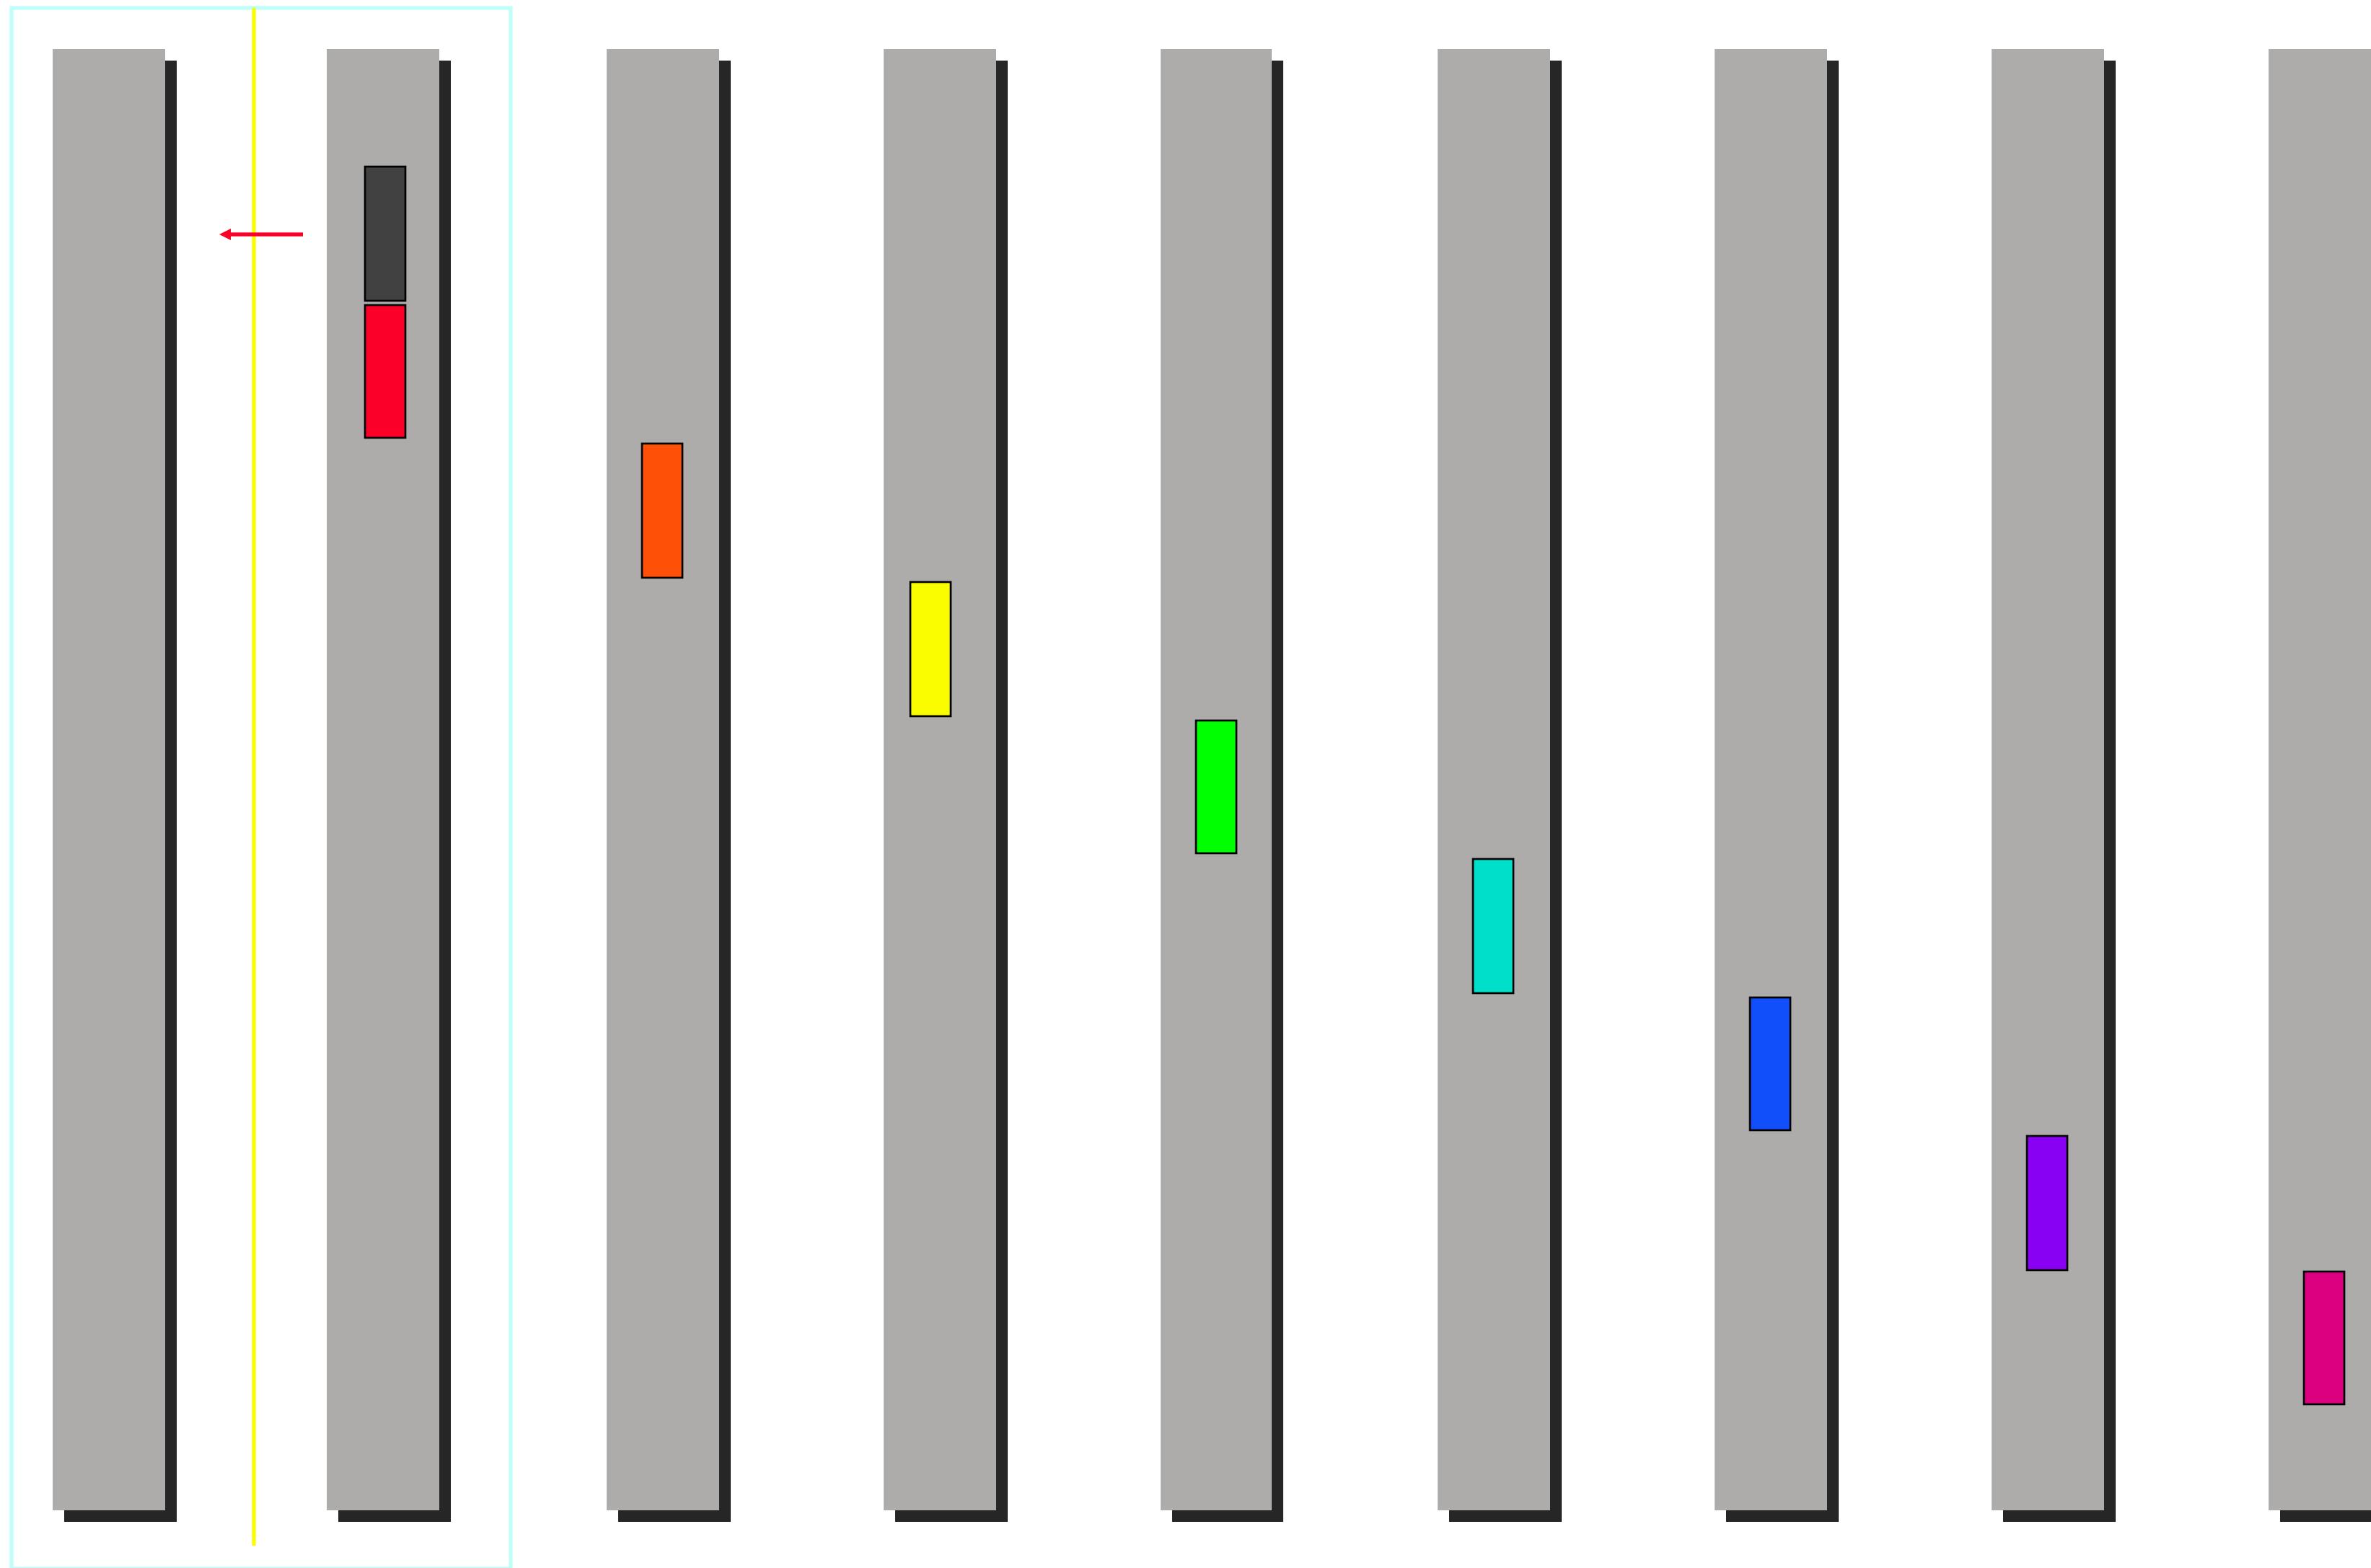


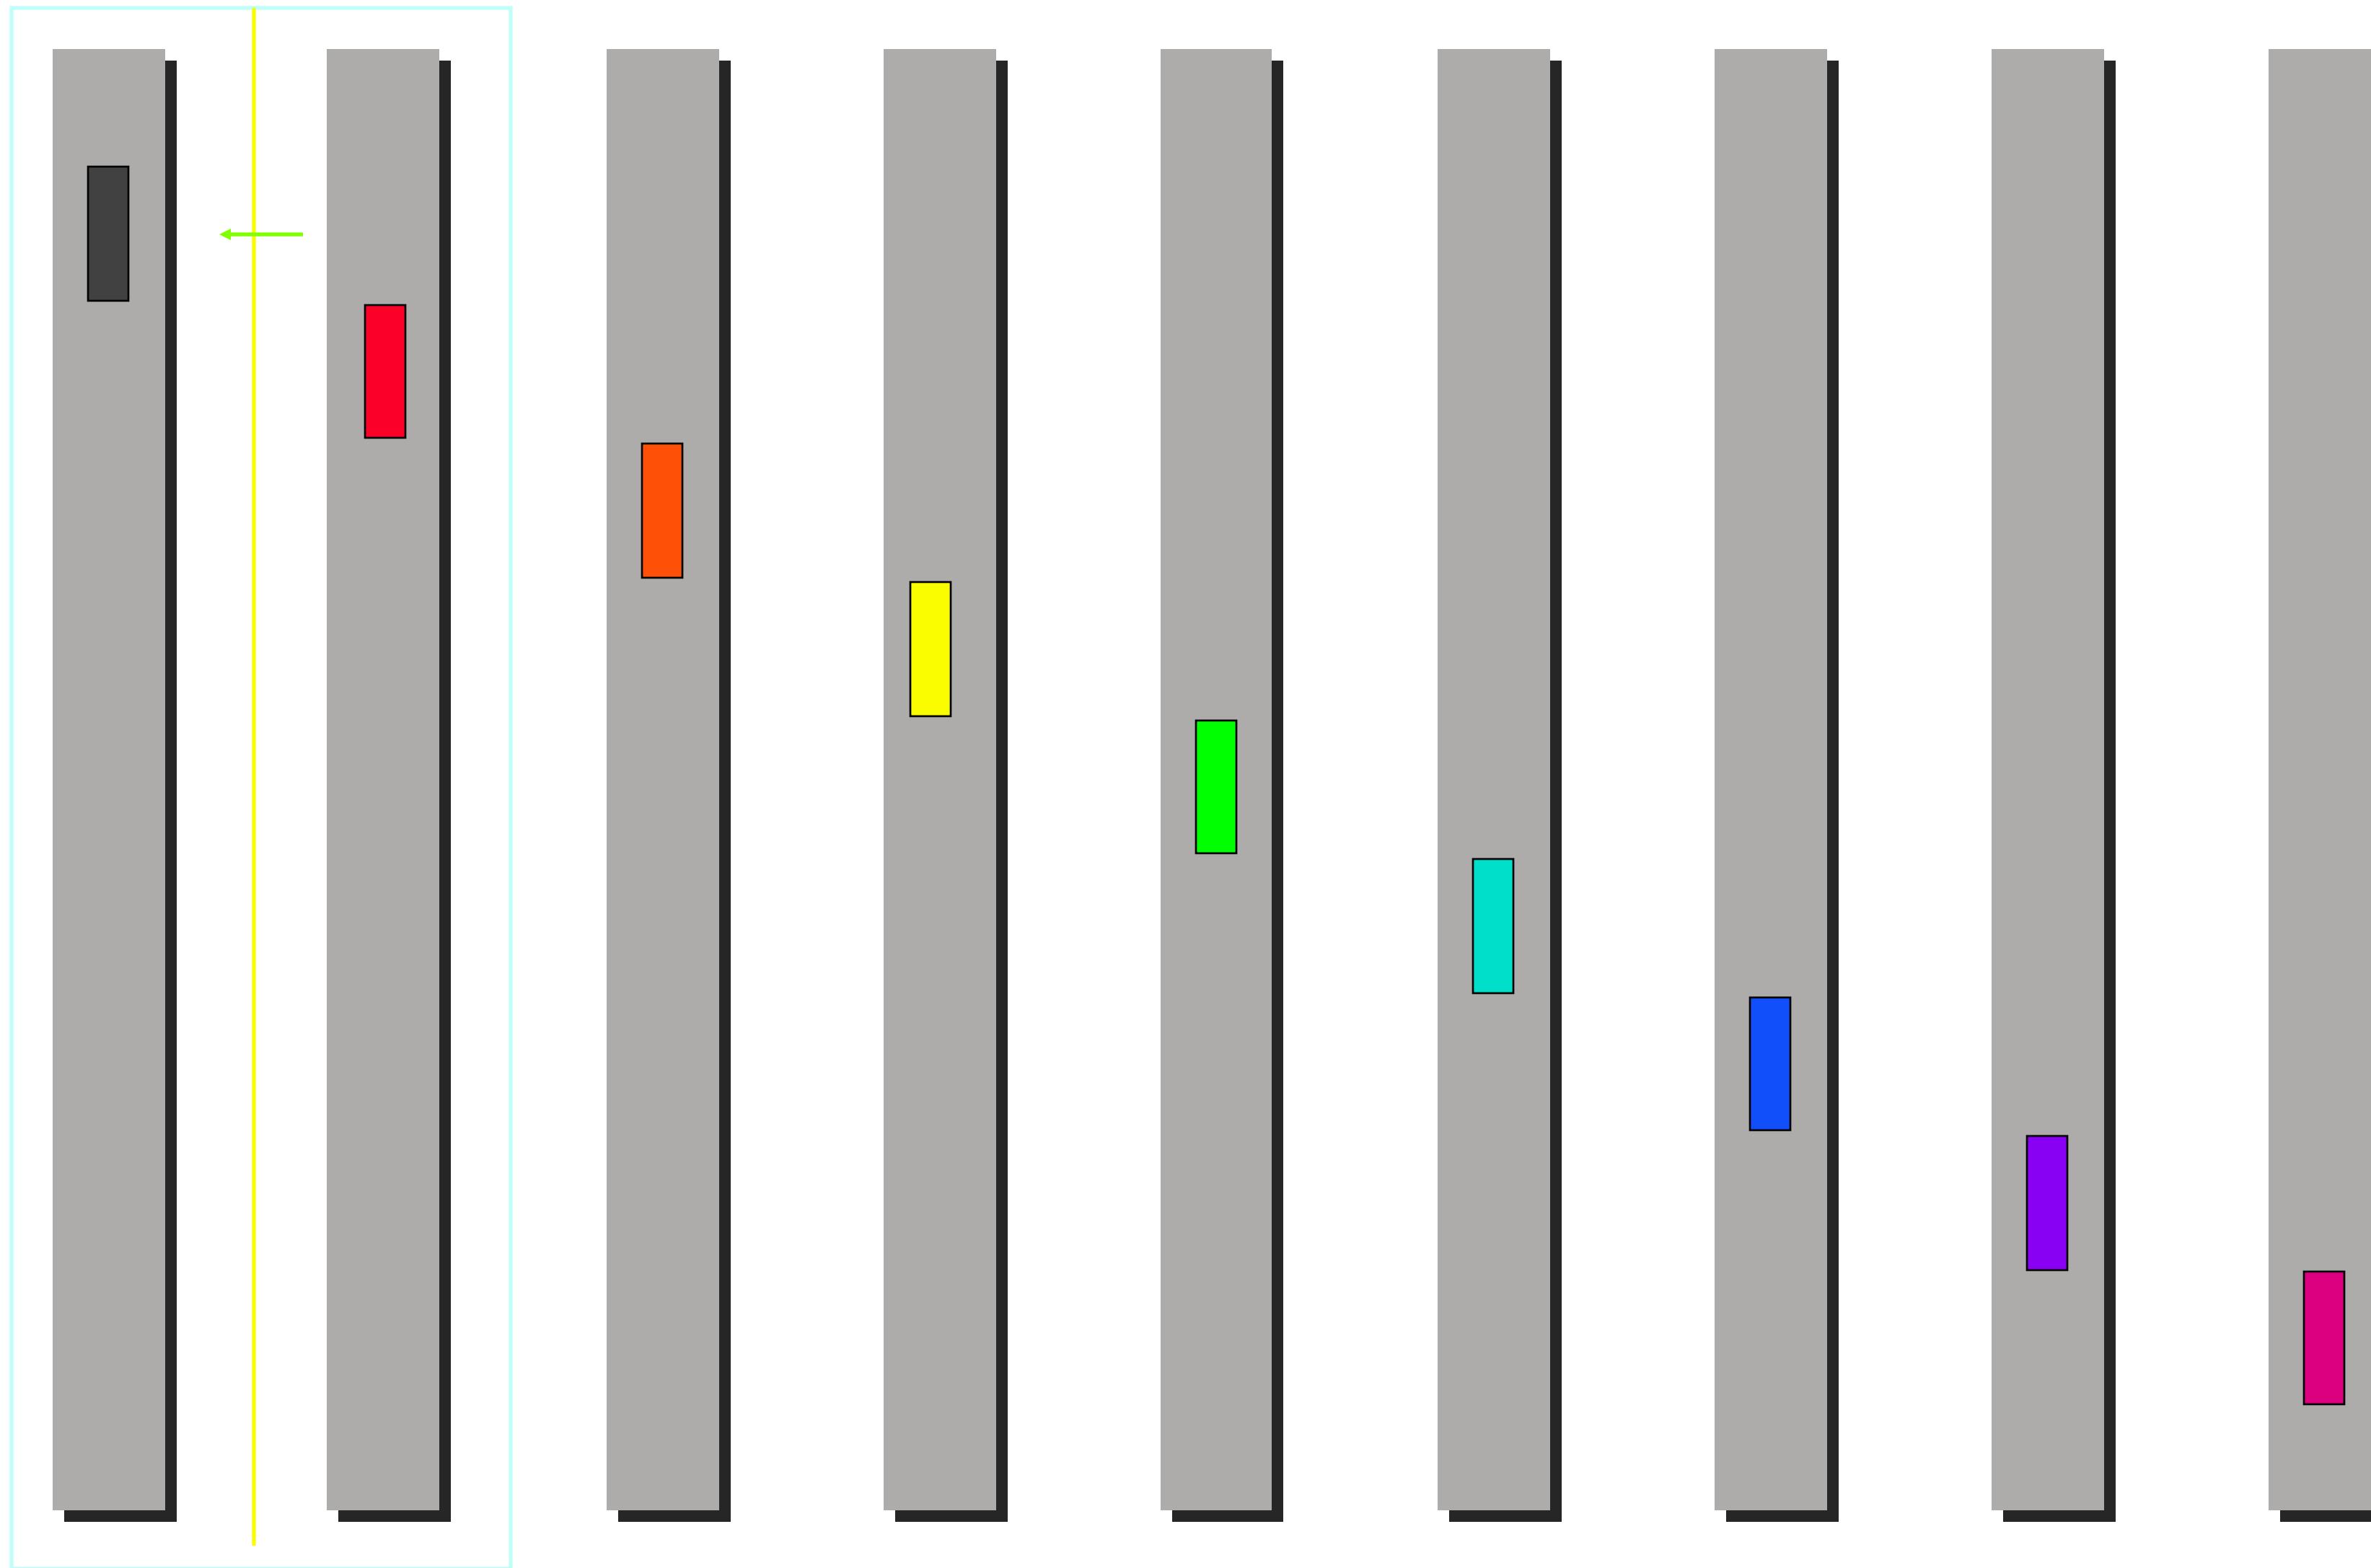


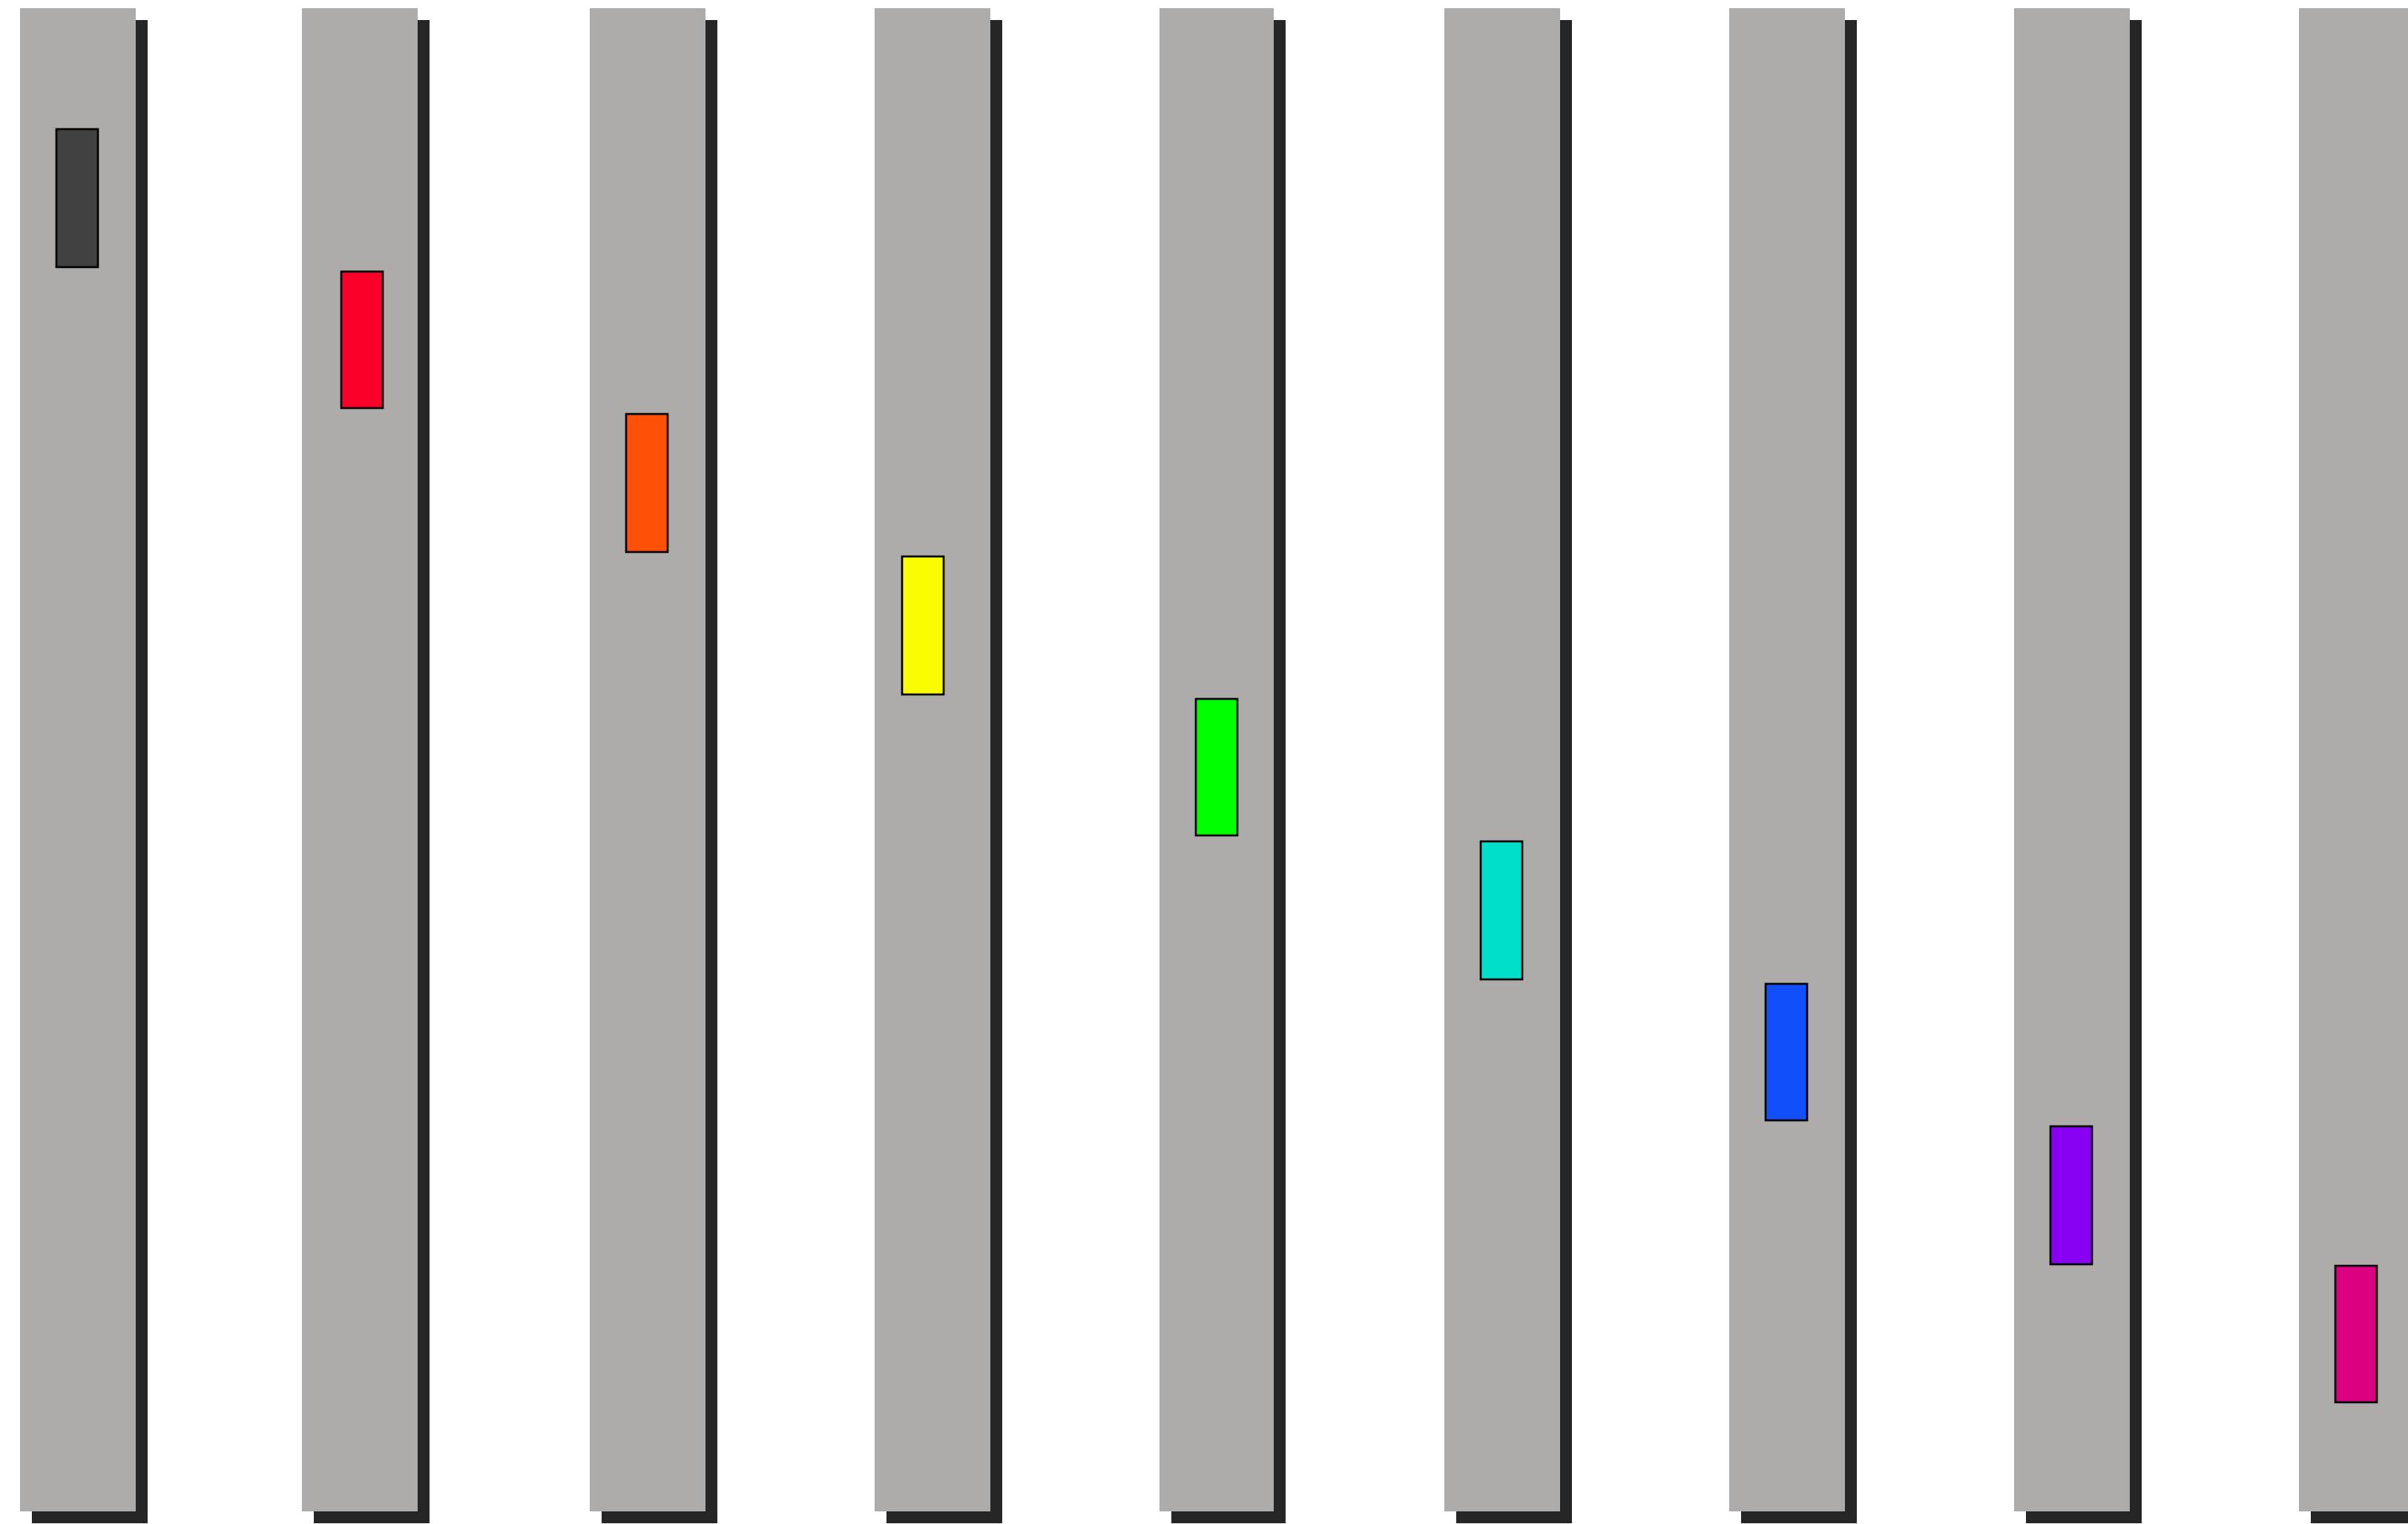












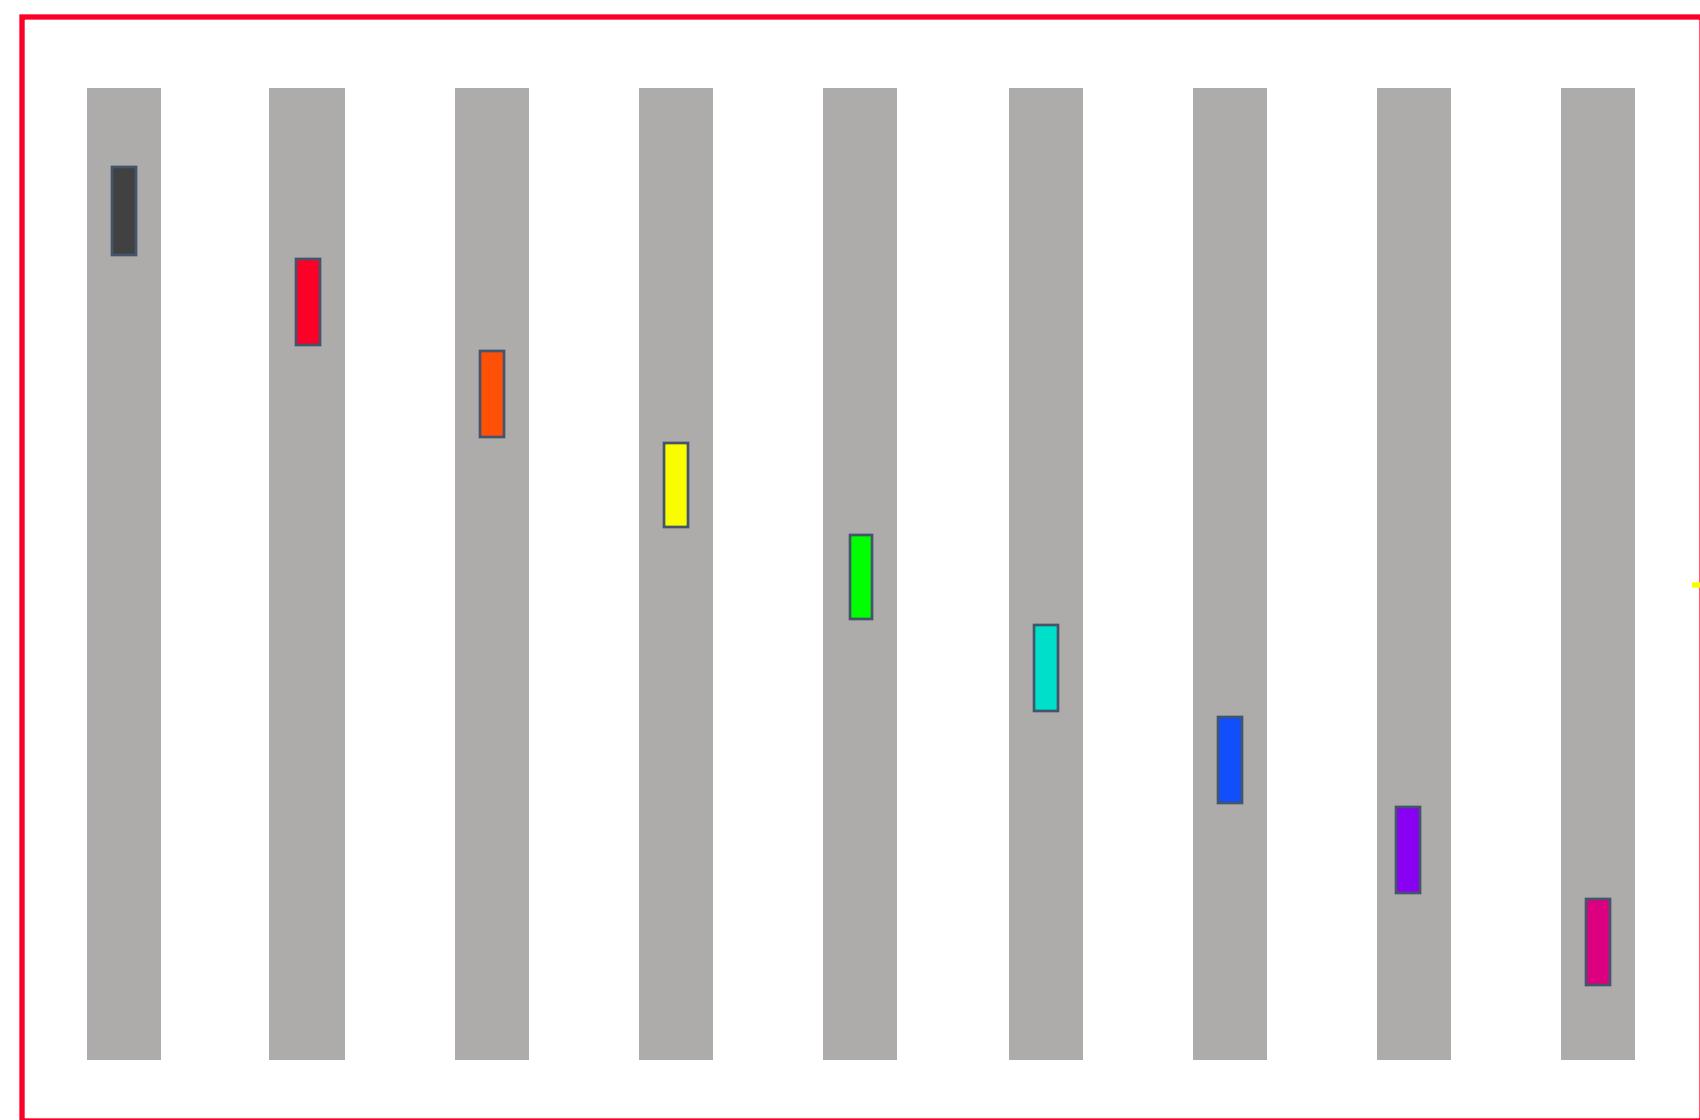
Cost of minimum spanning tree scatter

- Assumption: power of two number of nodes

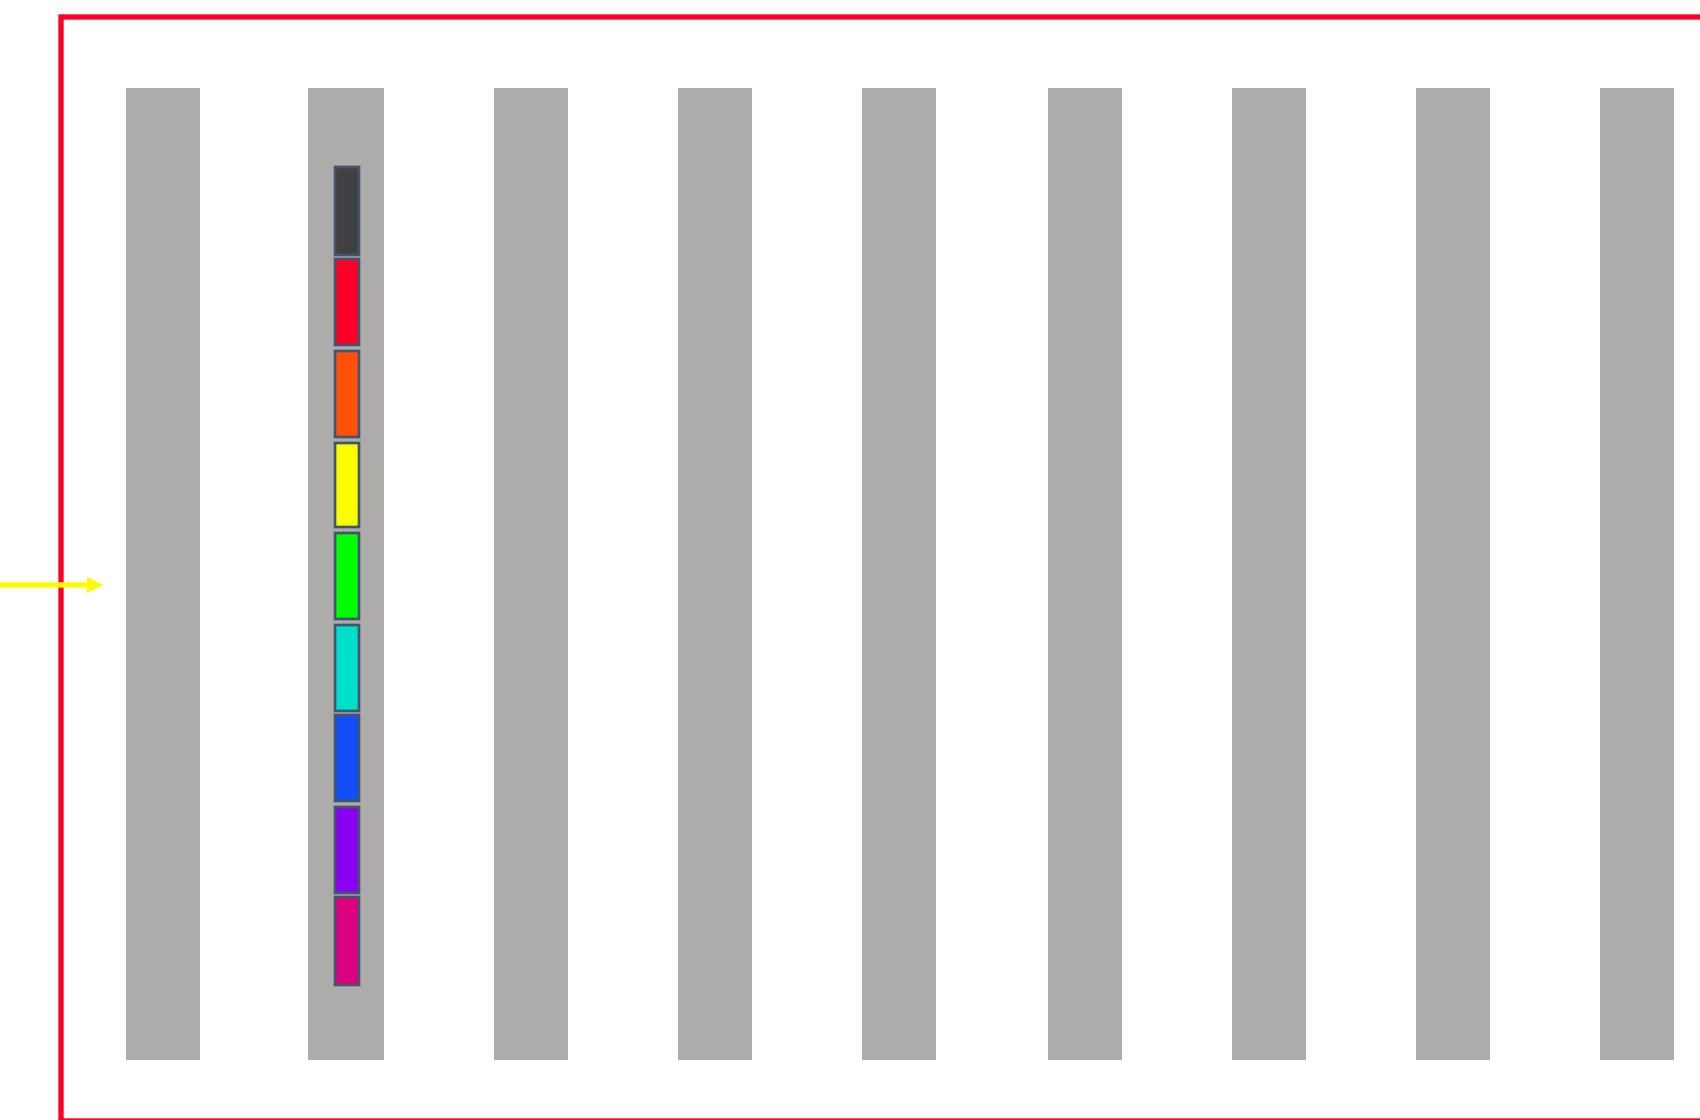
$$\begin{aligned} & \sum_{k=1}^{\log(p)} \left(\alpha + \frac{n}{2^k} \beta \right) \\ &= \\ & \log(p) - \alpha + \frac{p-1}{p} n \beta \end{aligned}$$

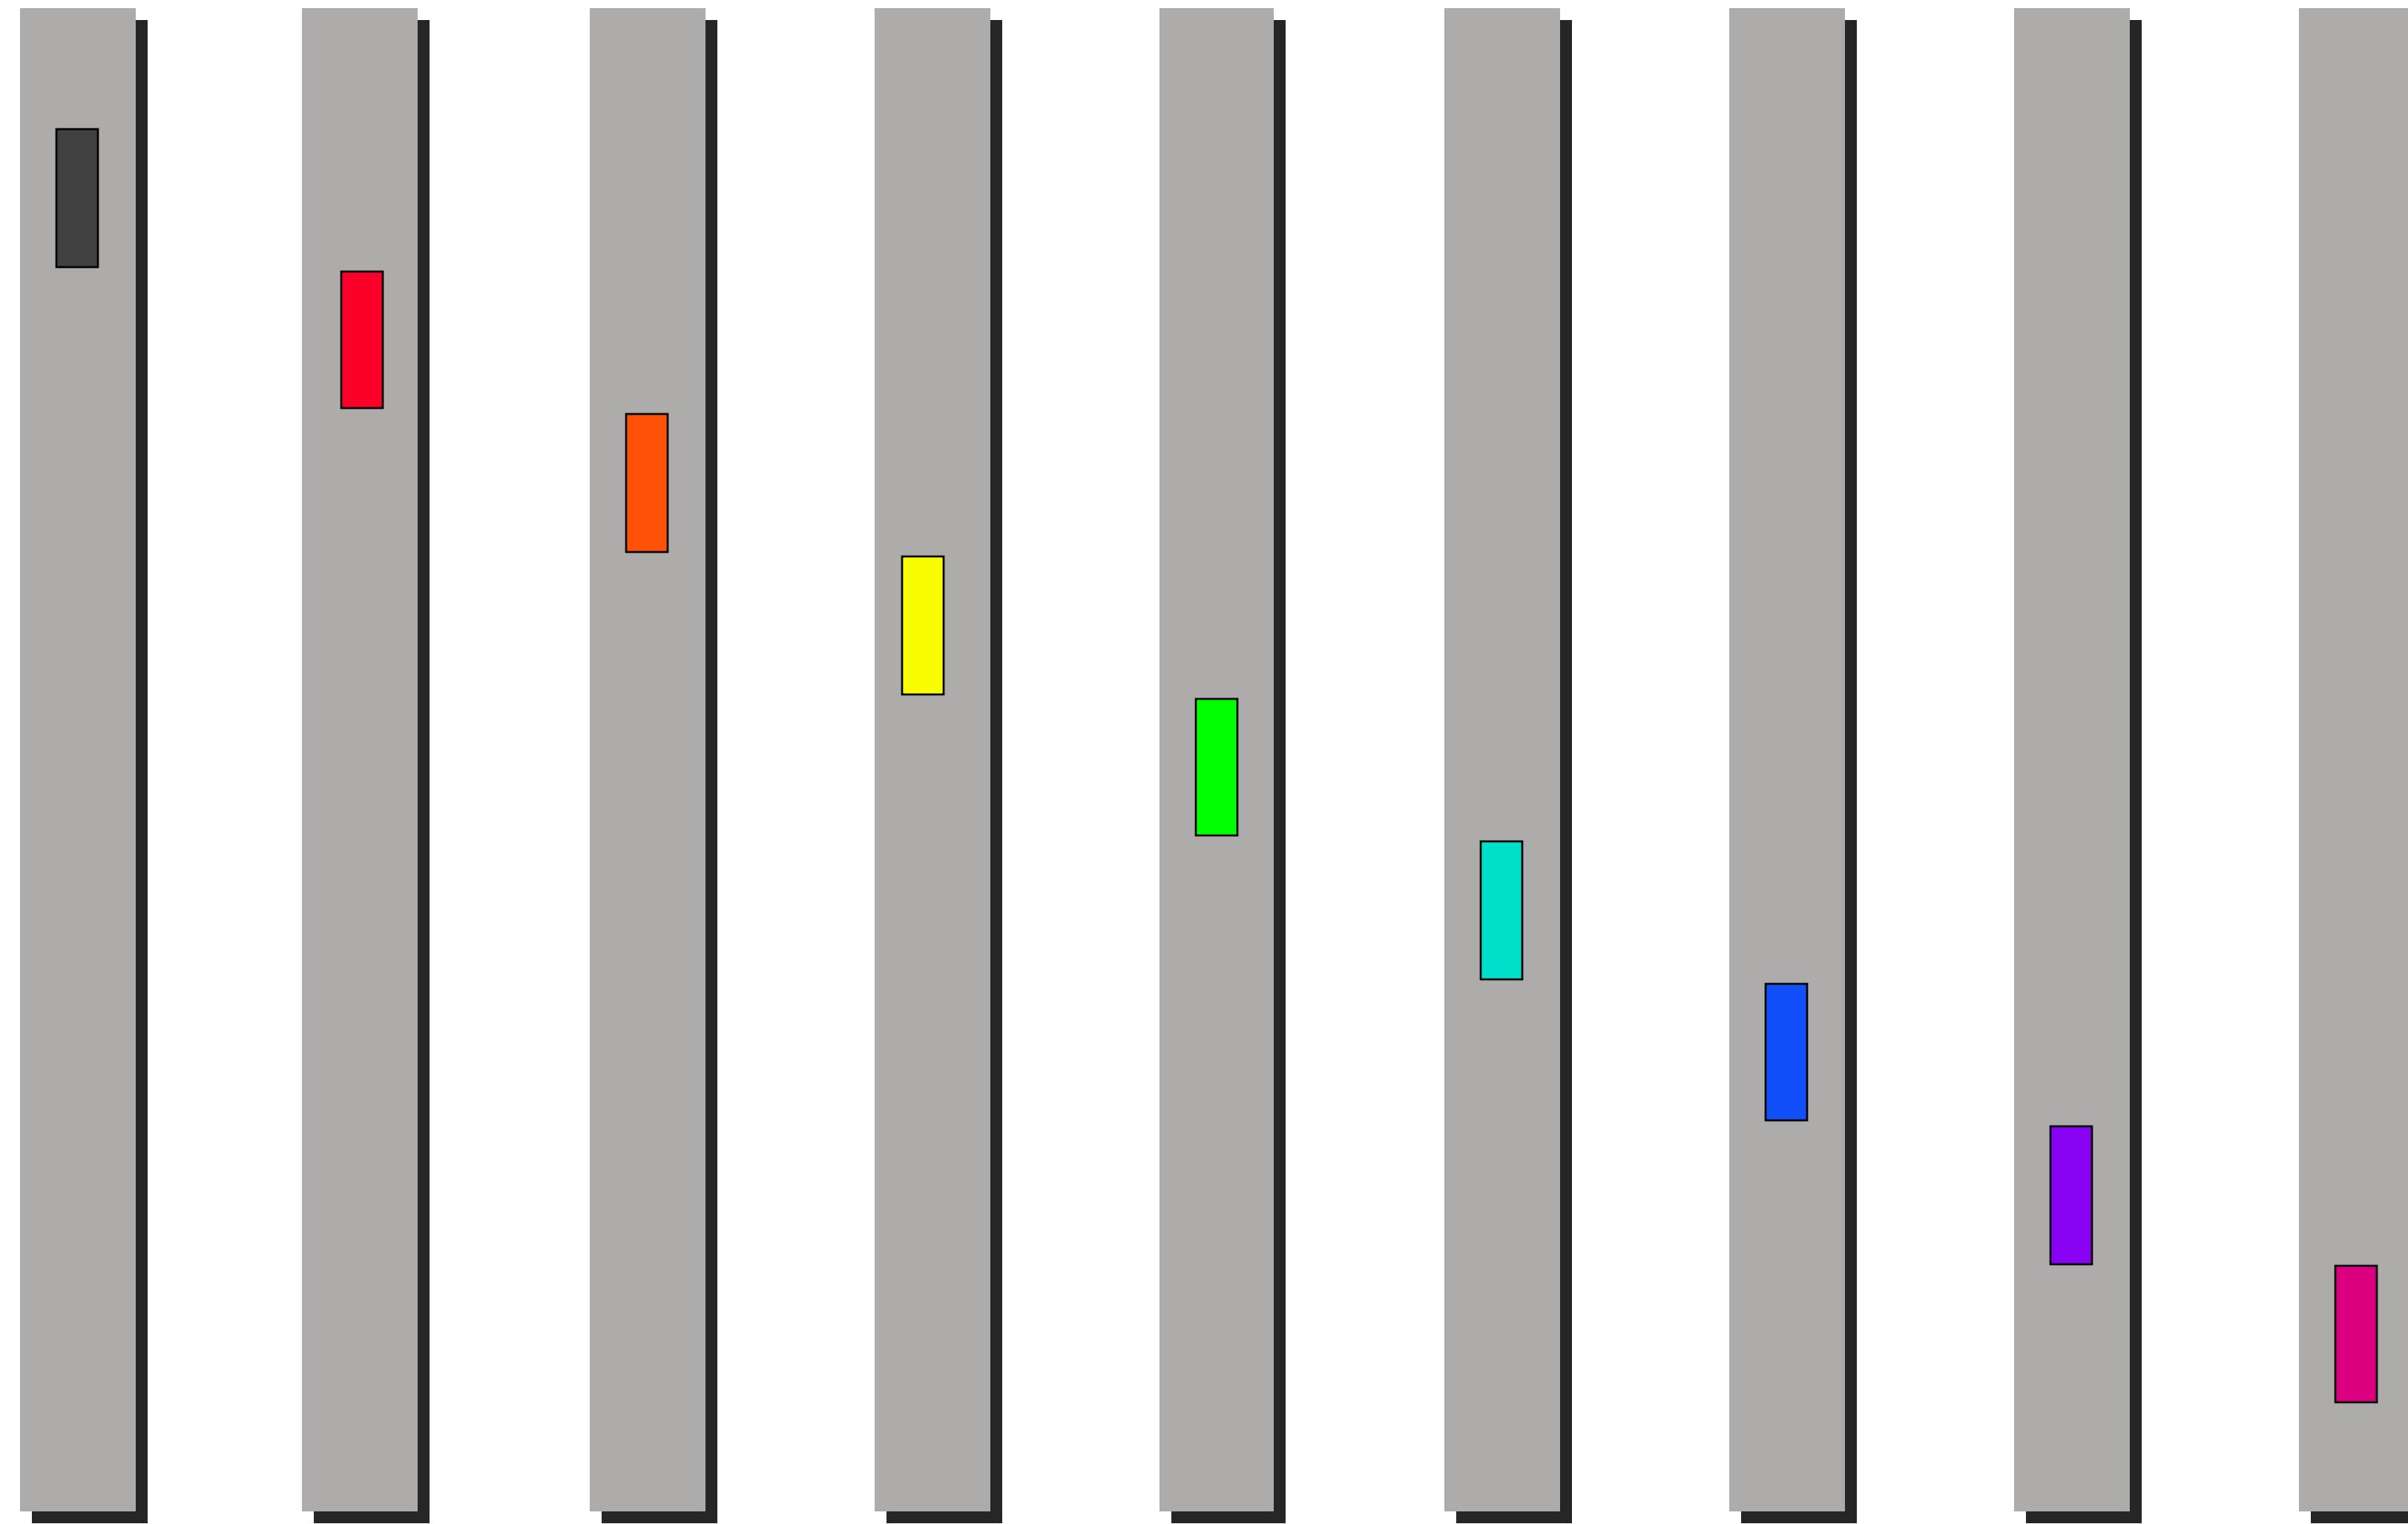
Gather

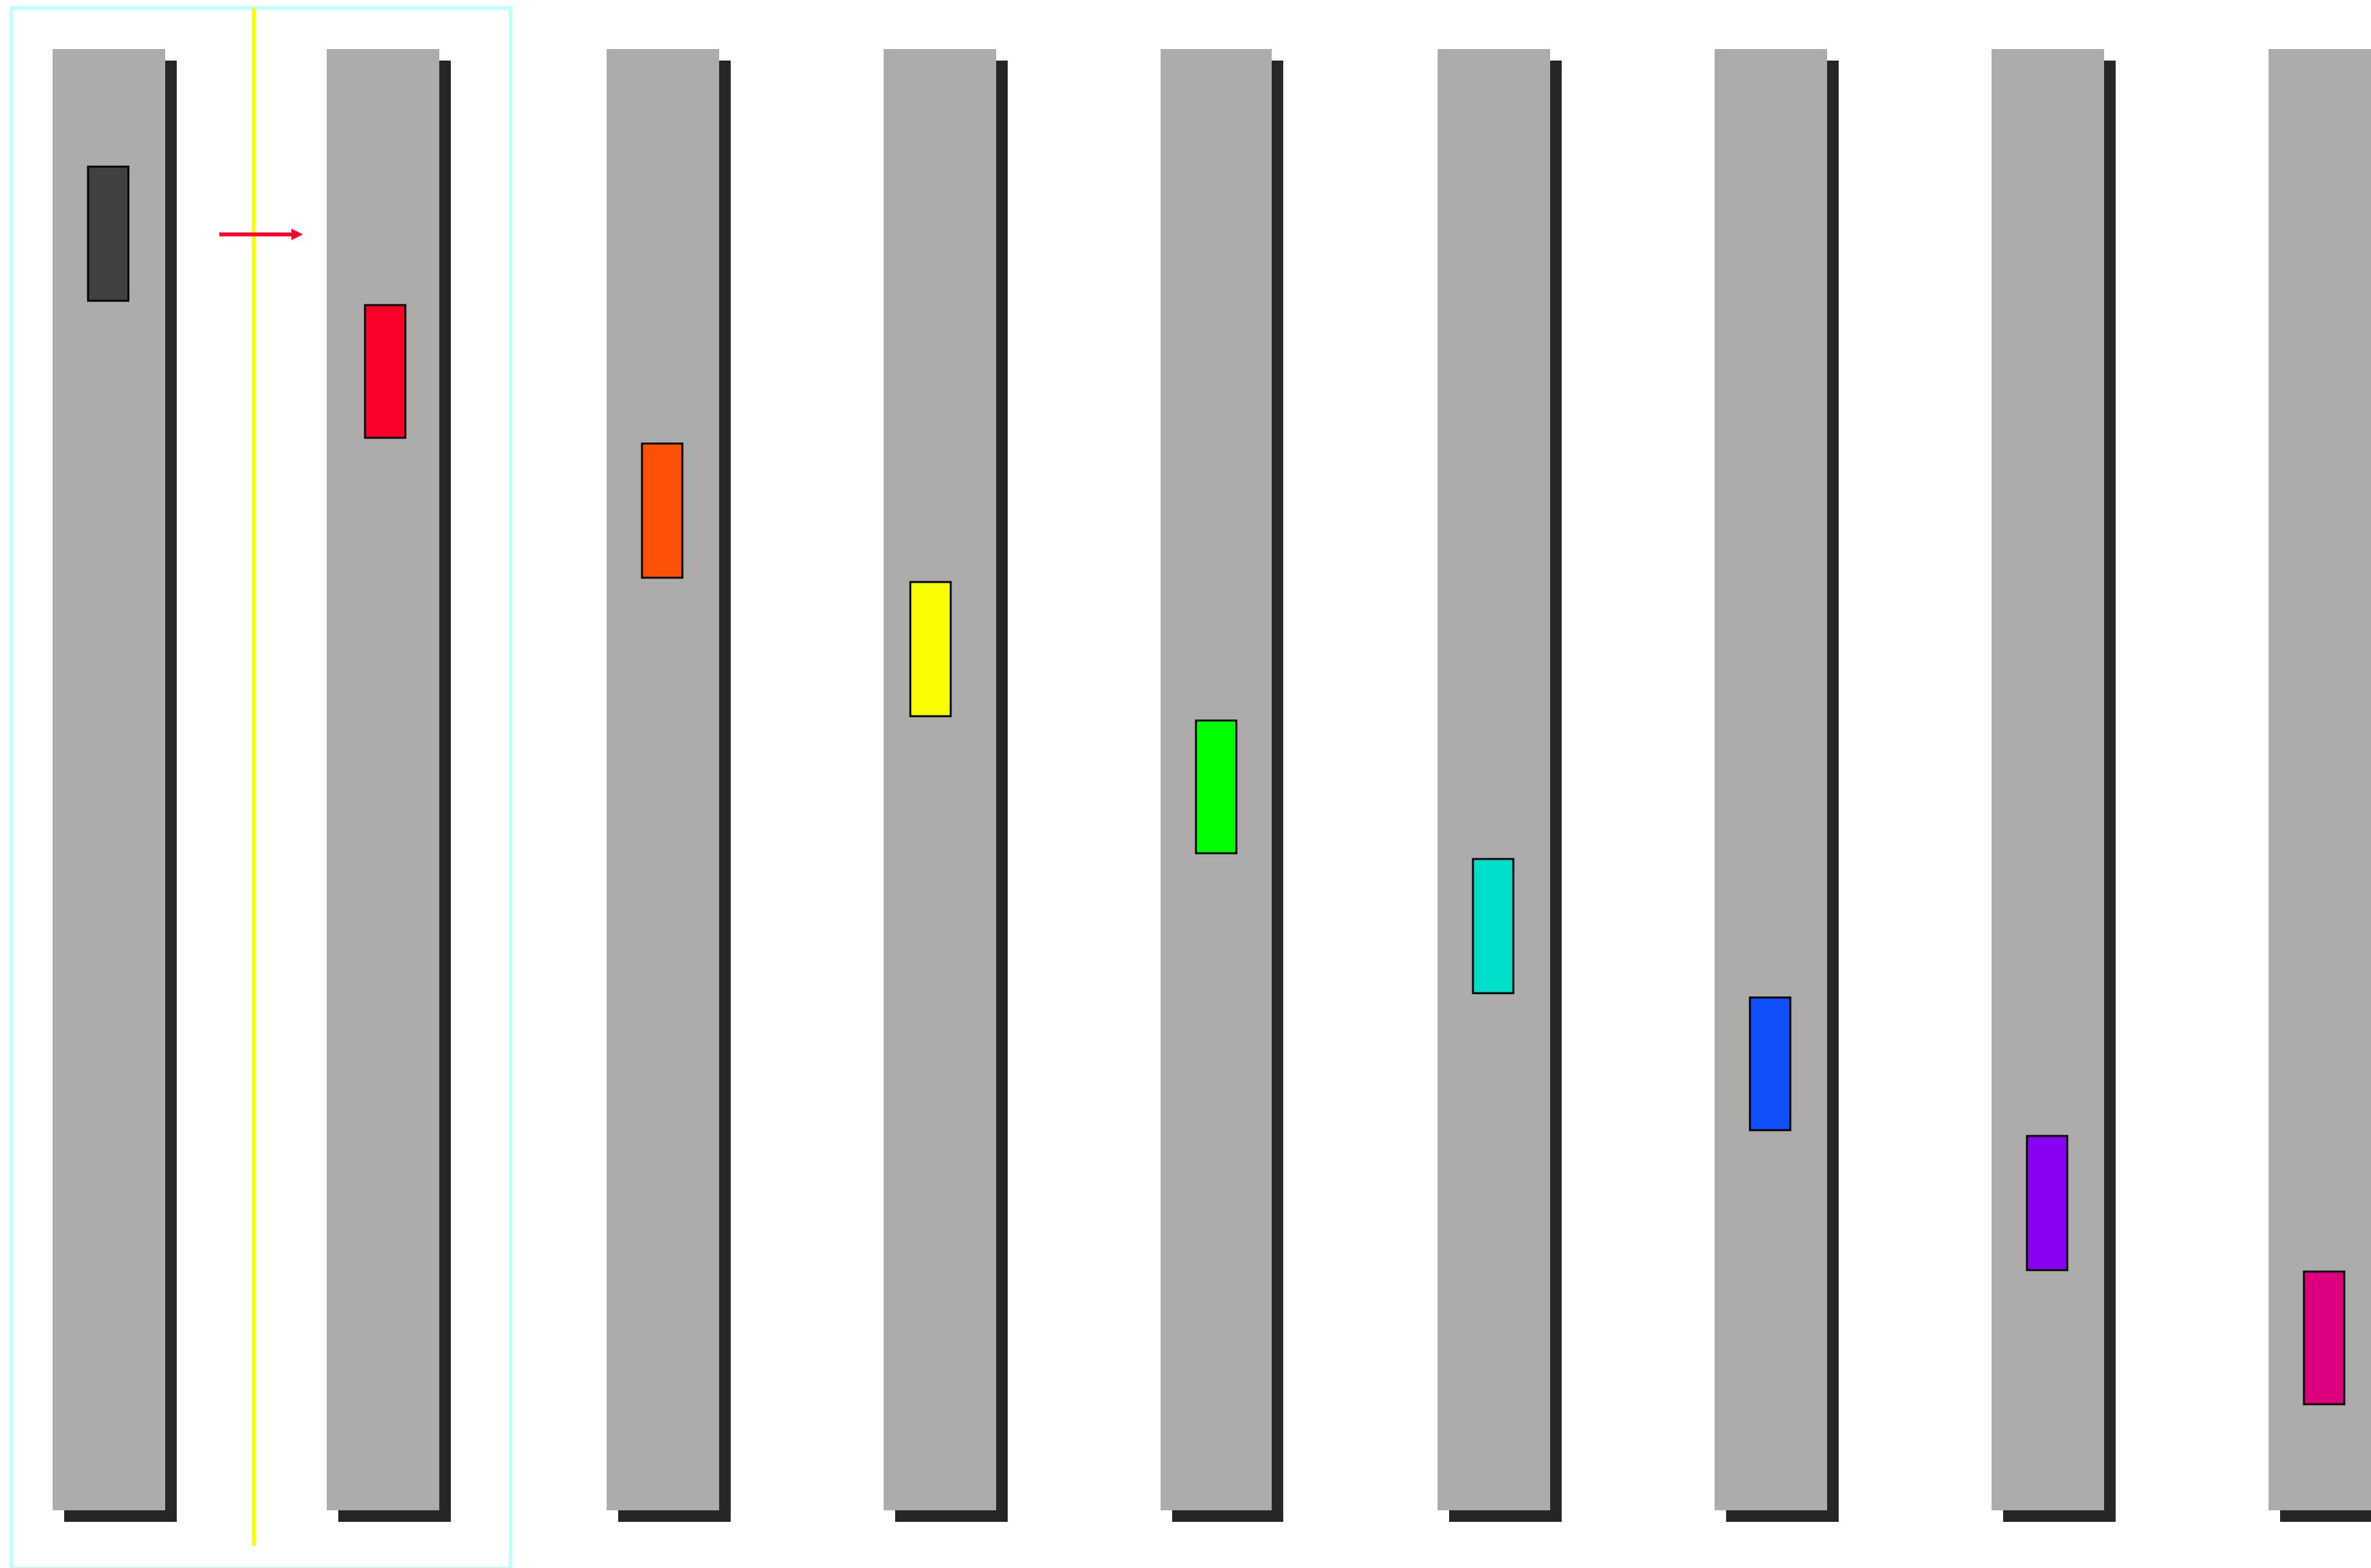
Before

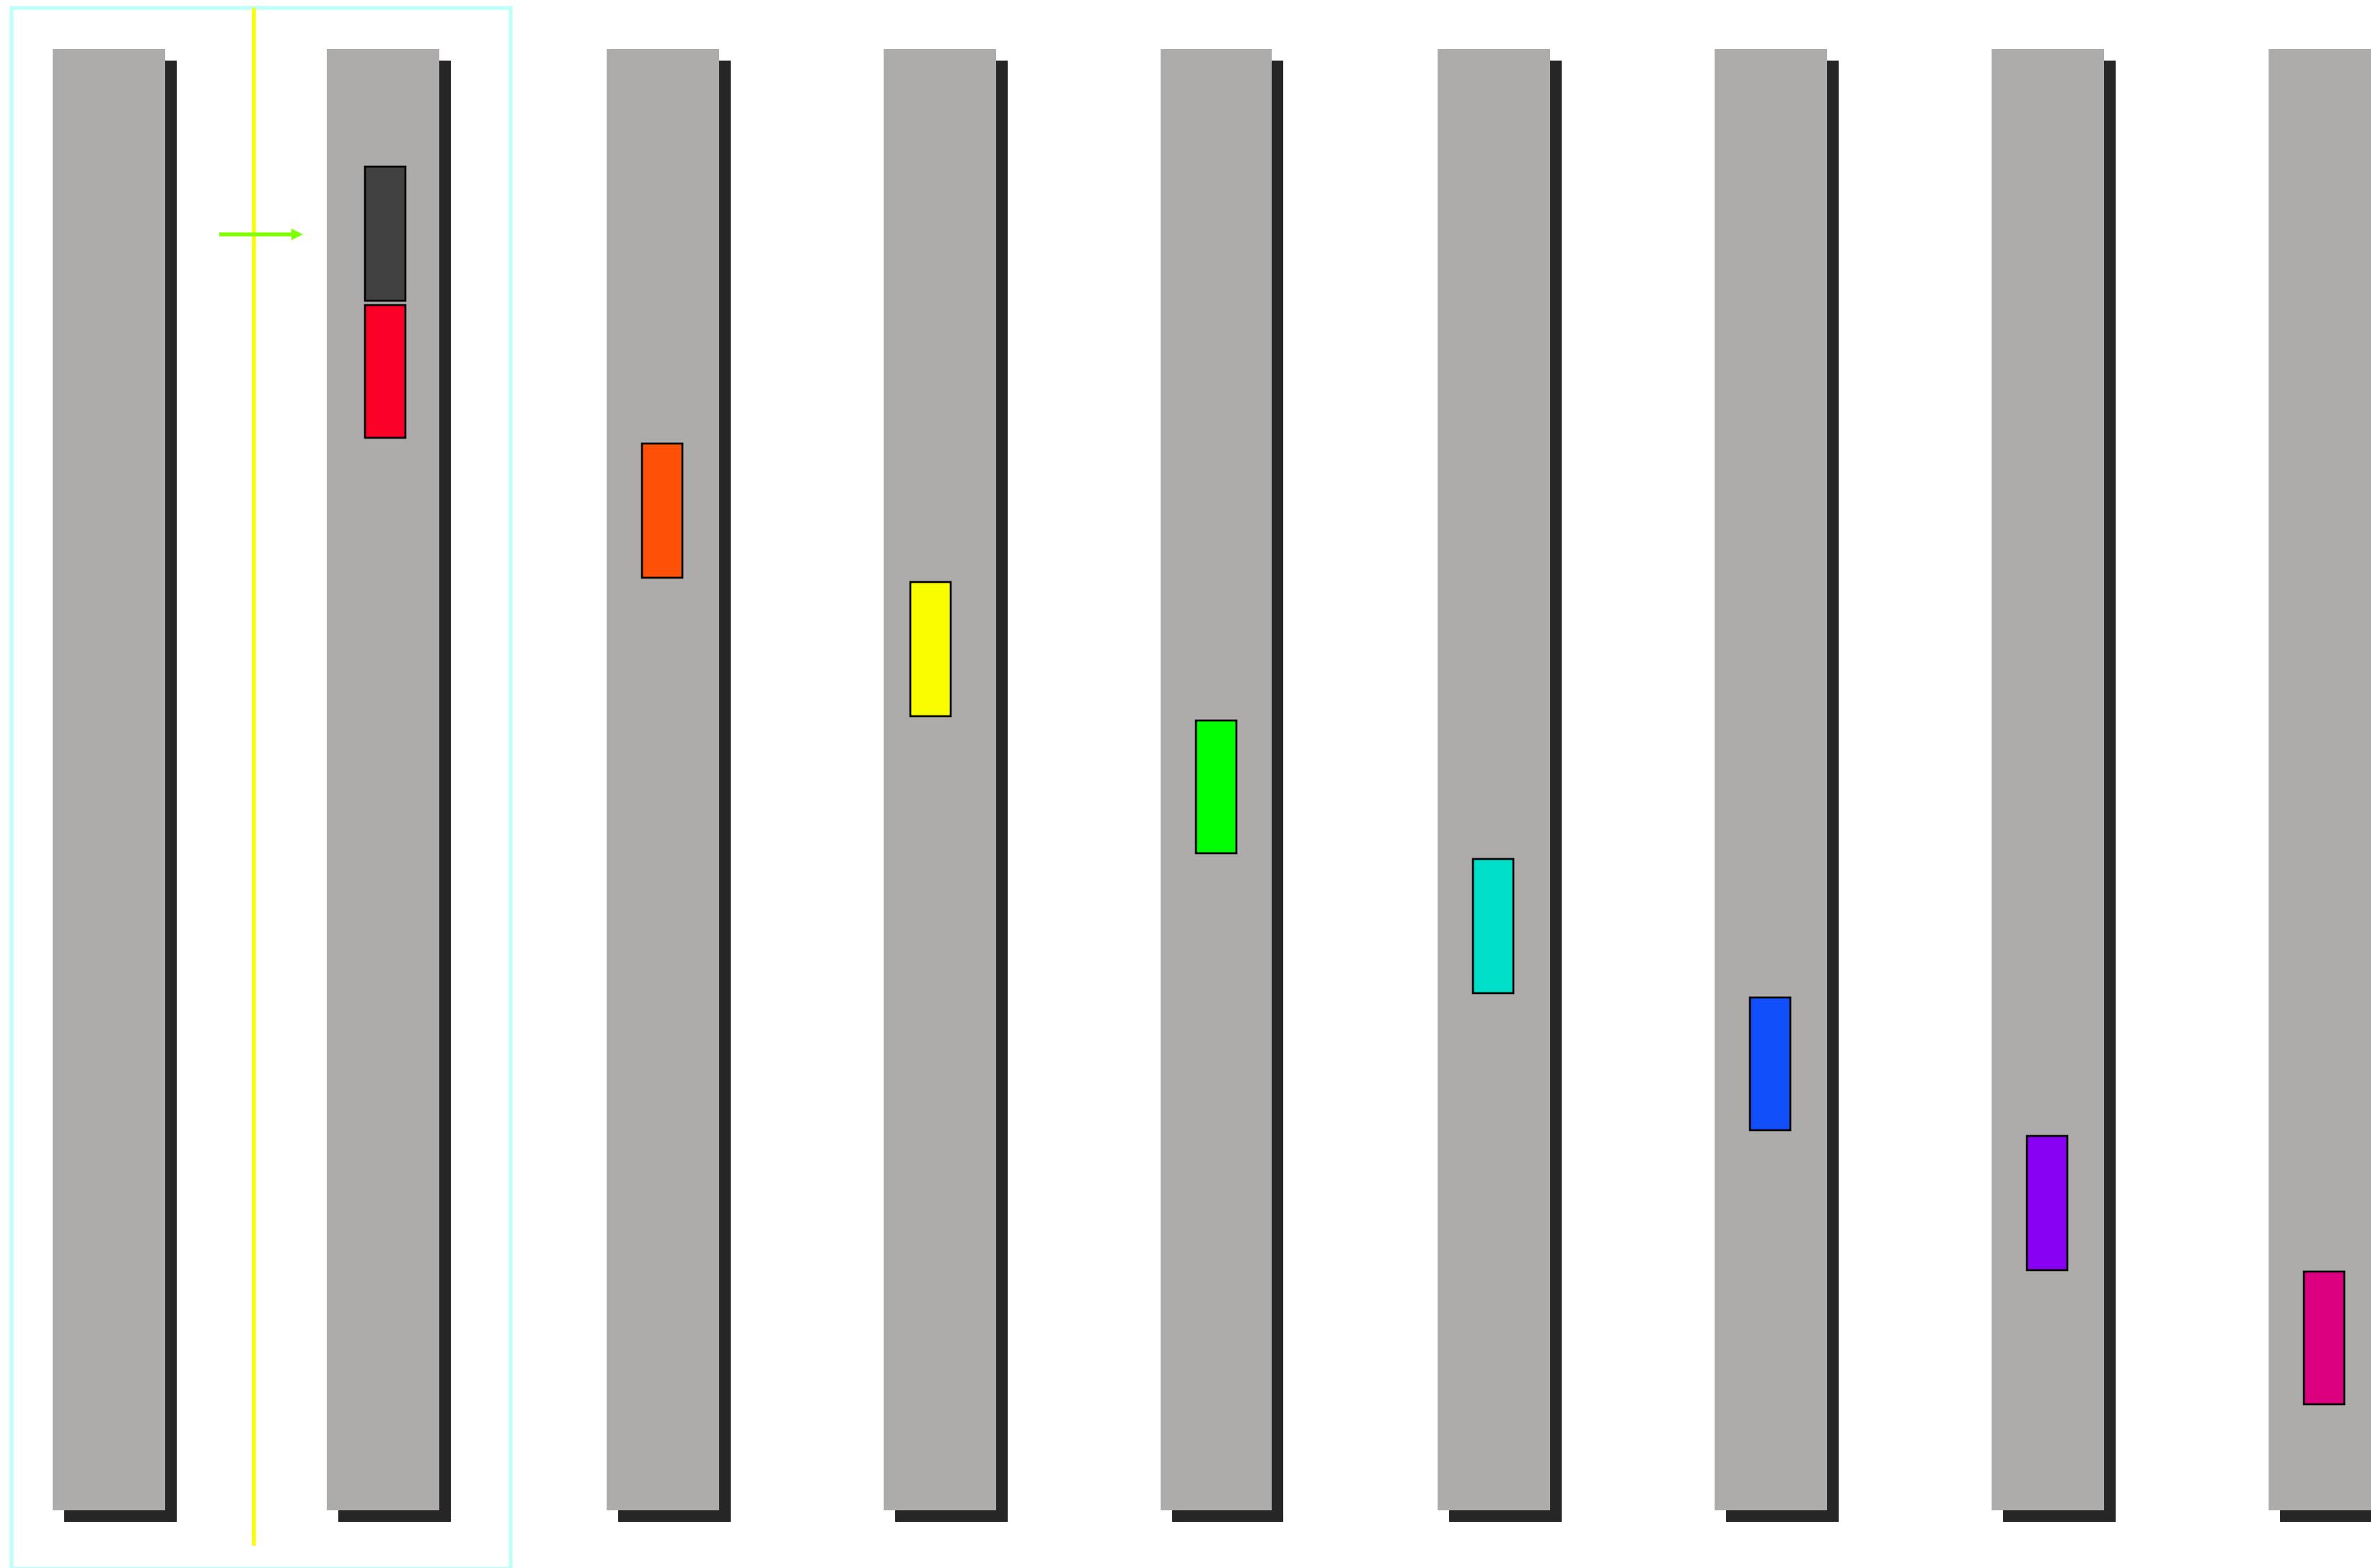


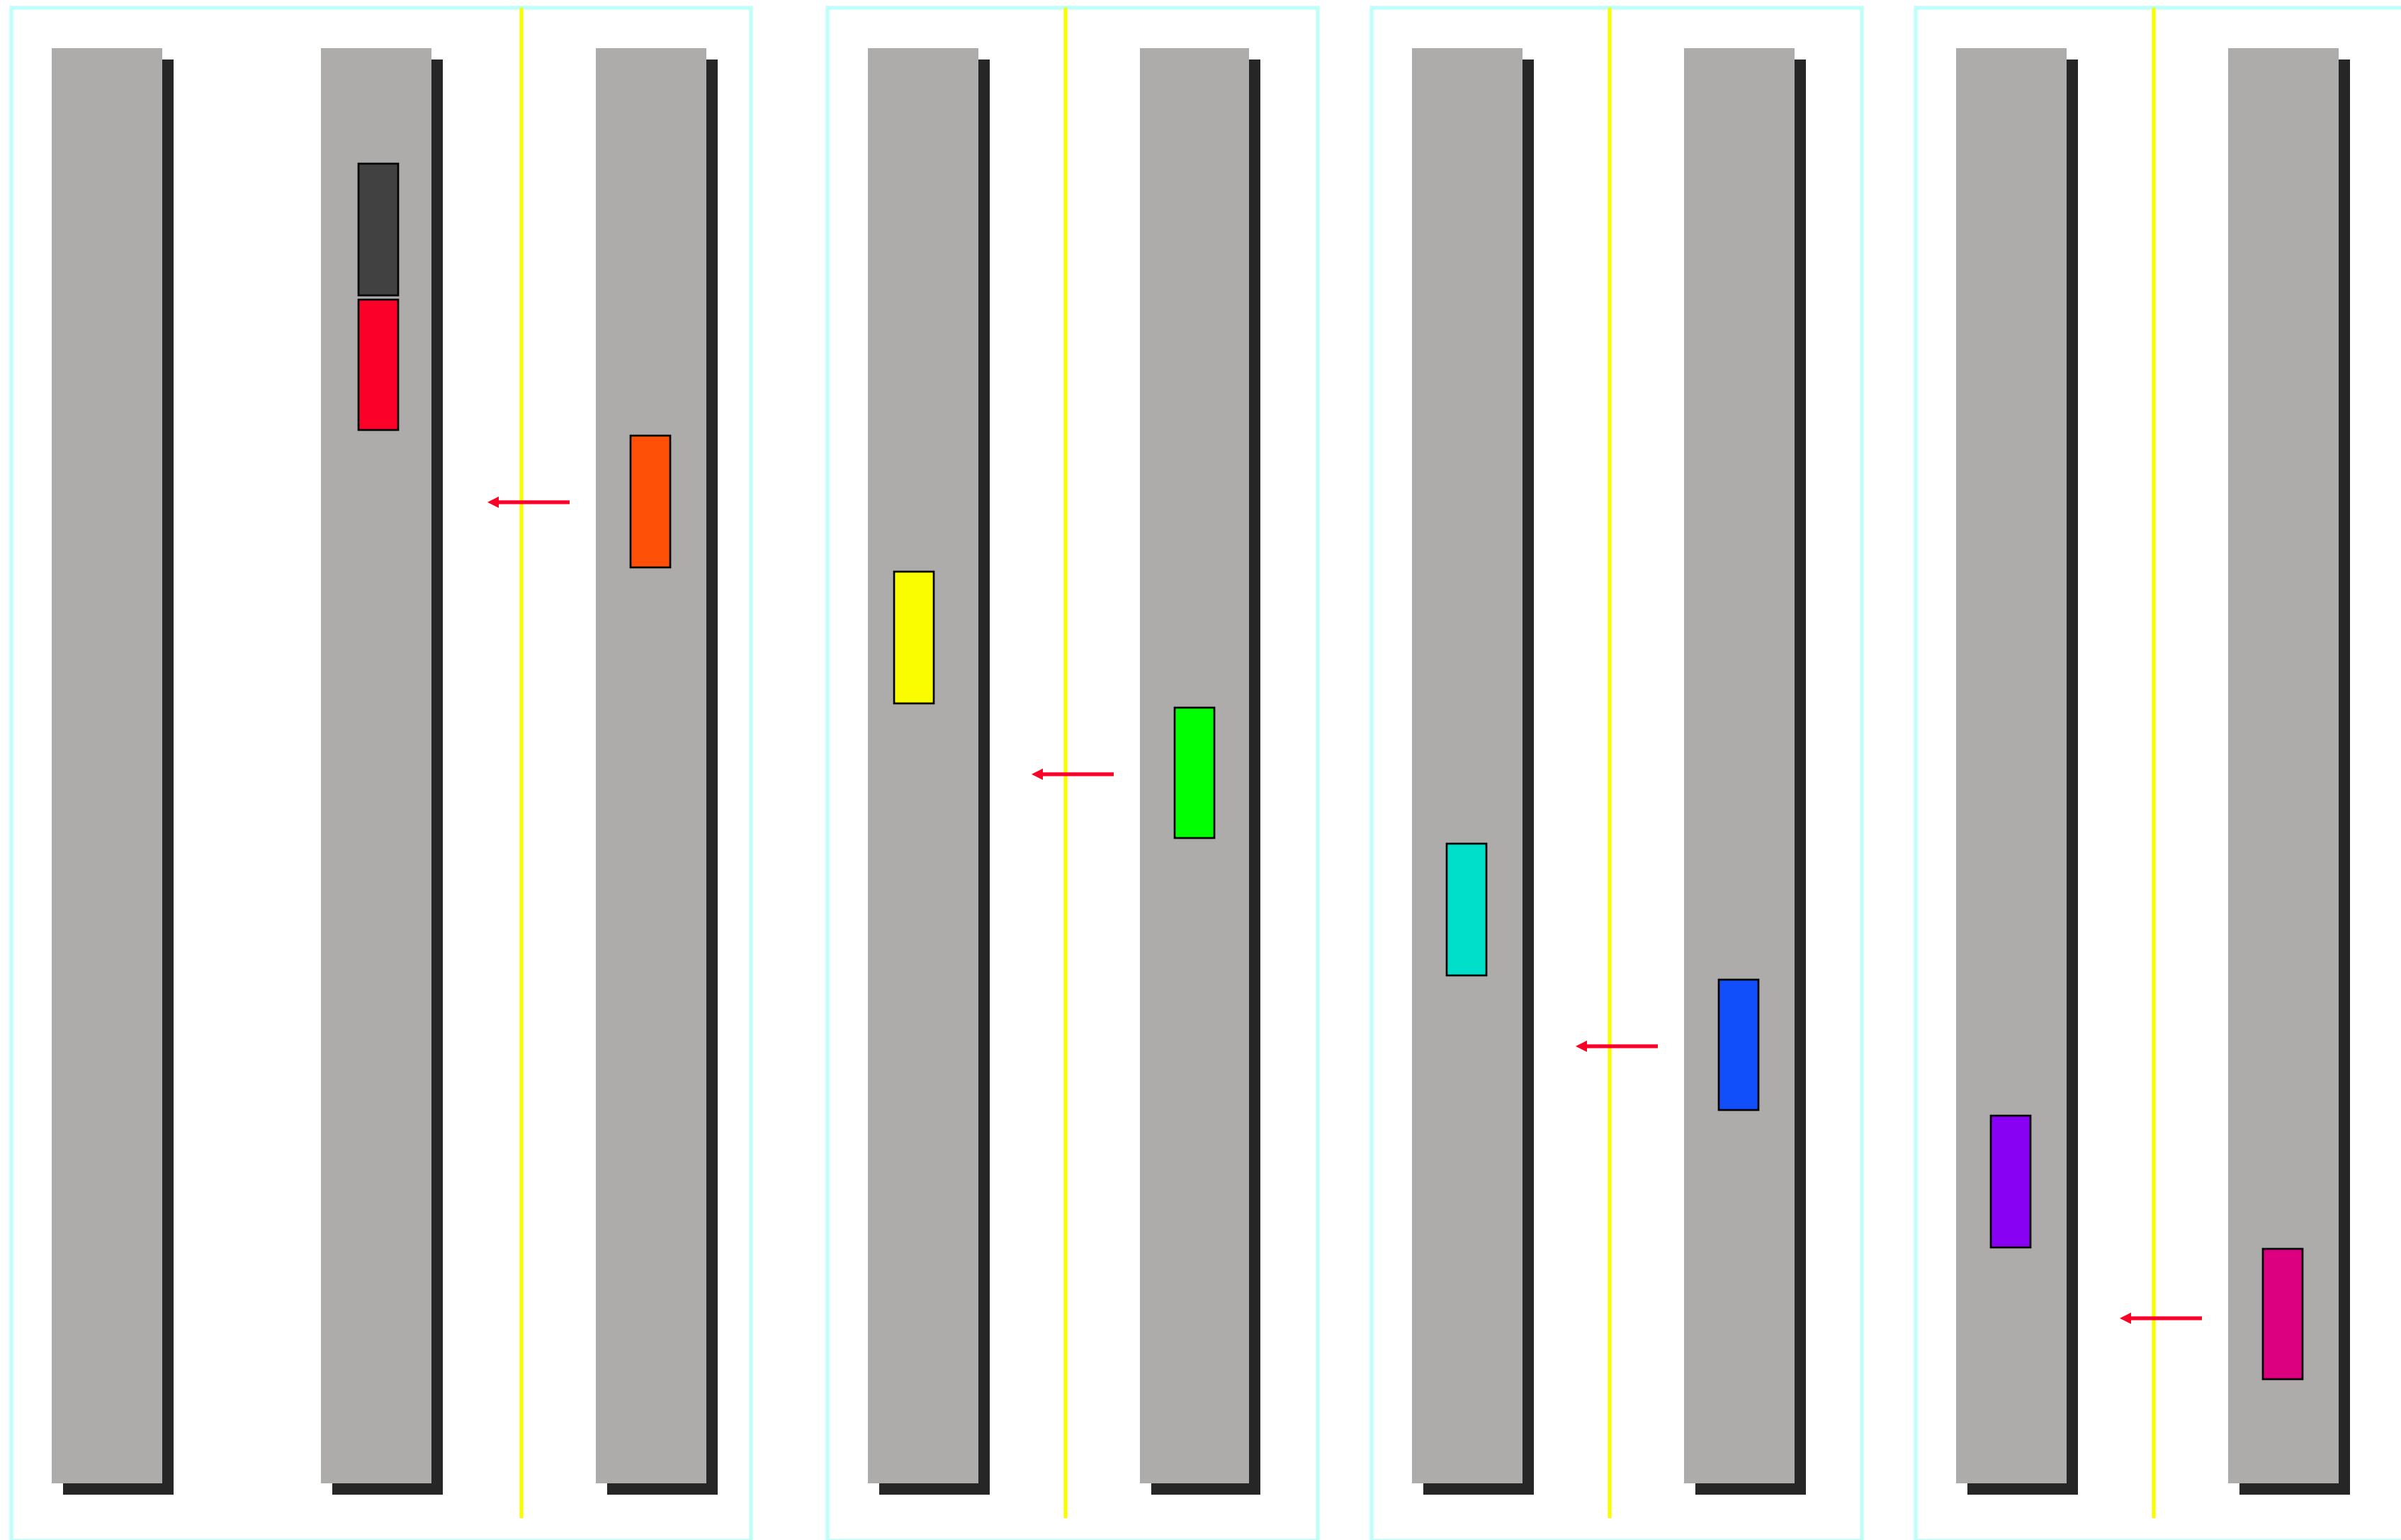
After

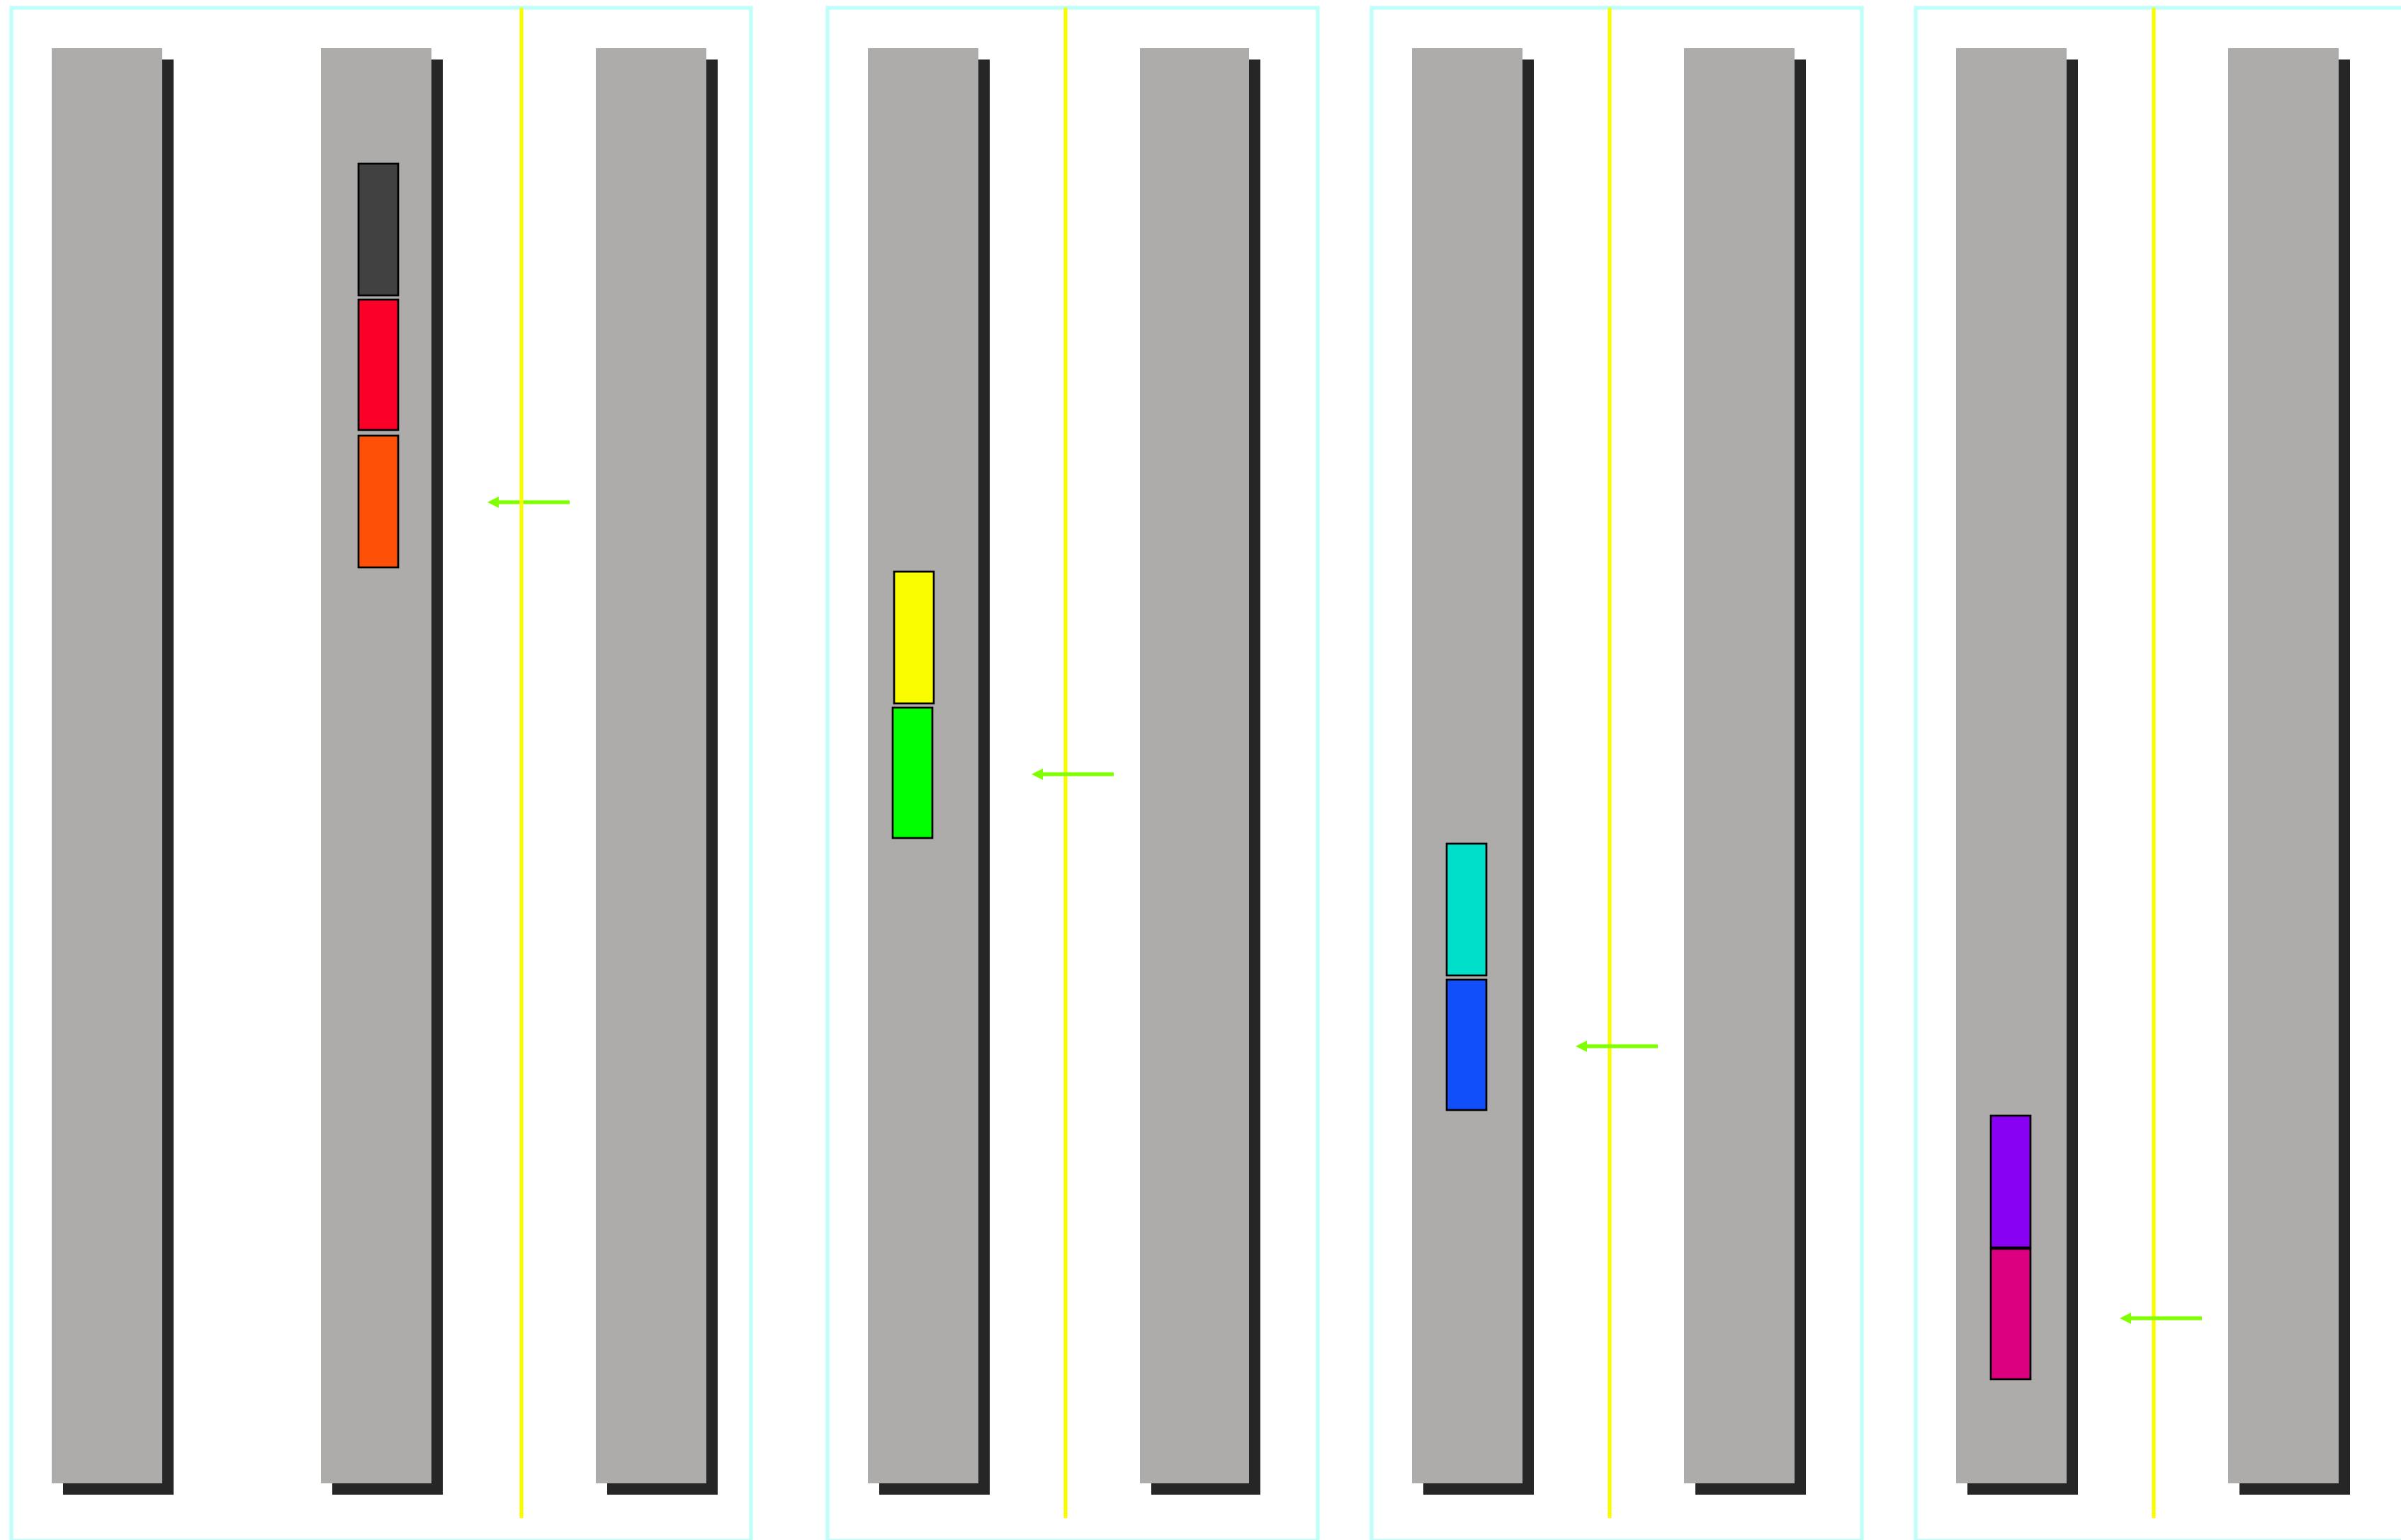


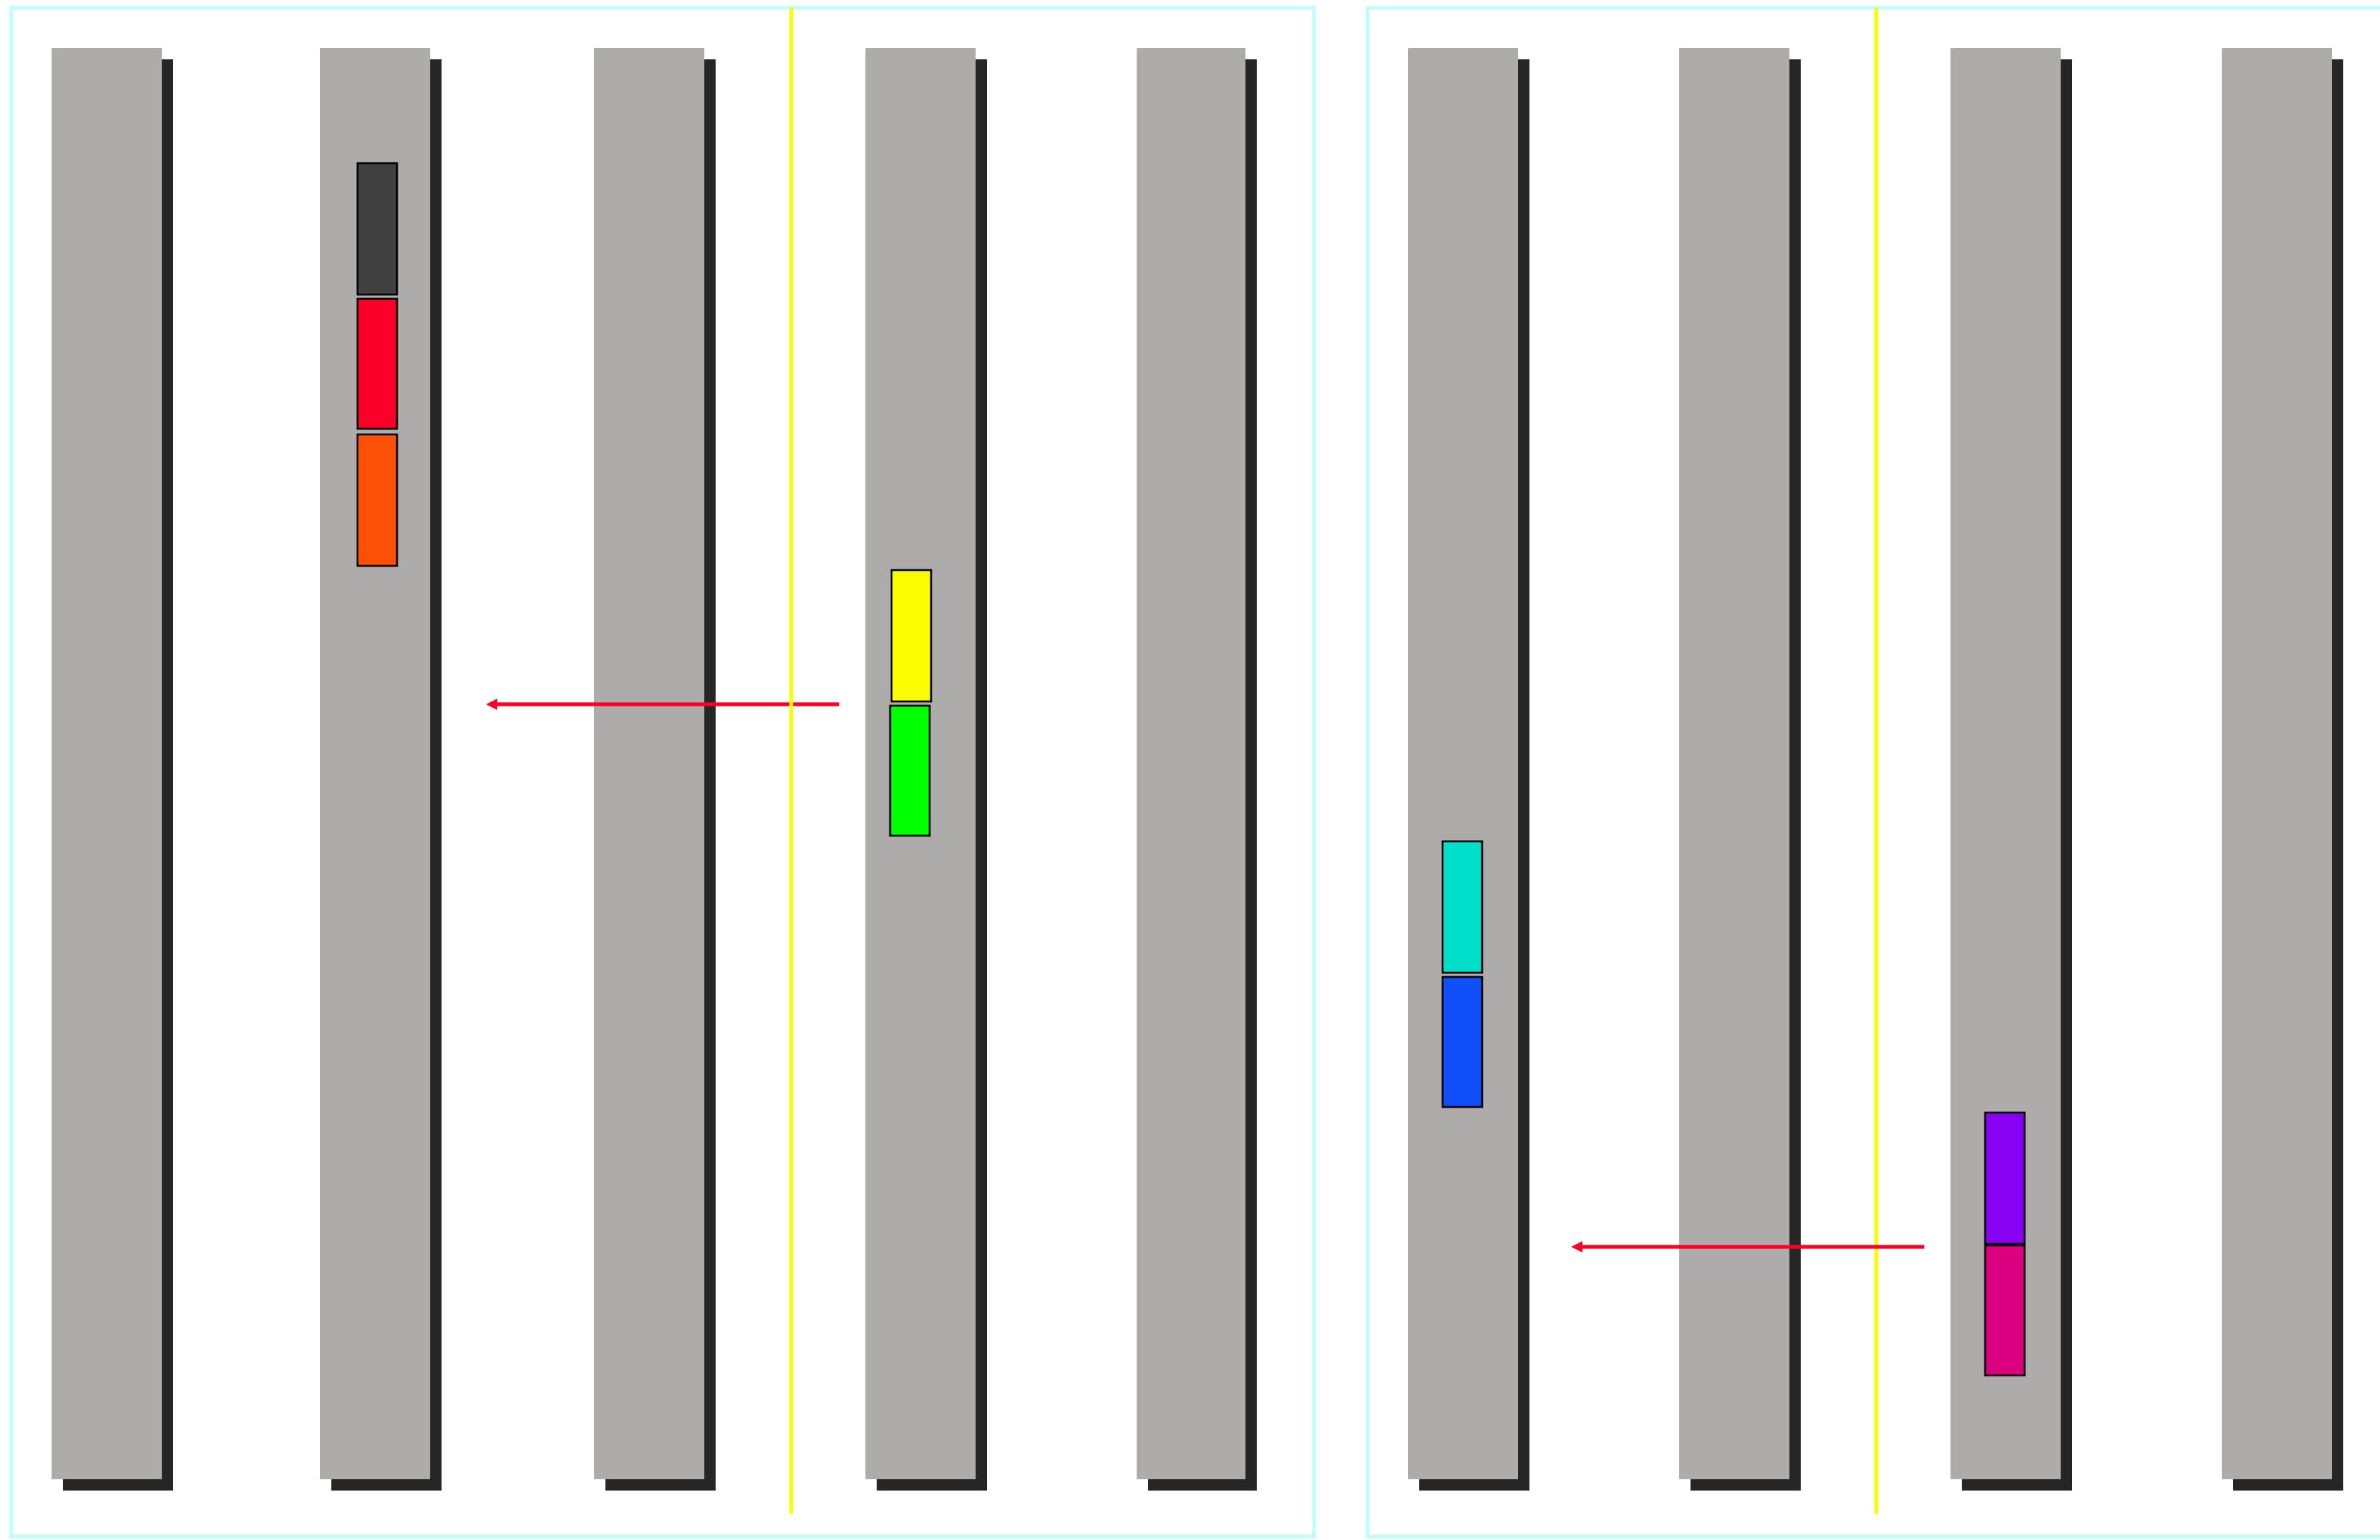


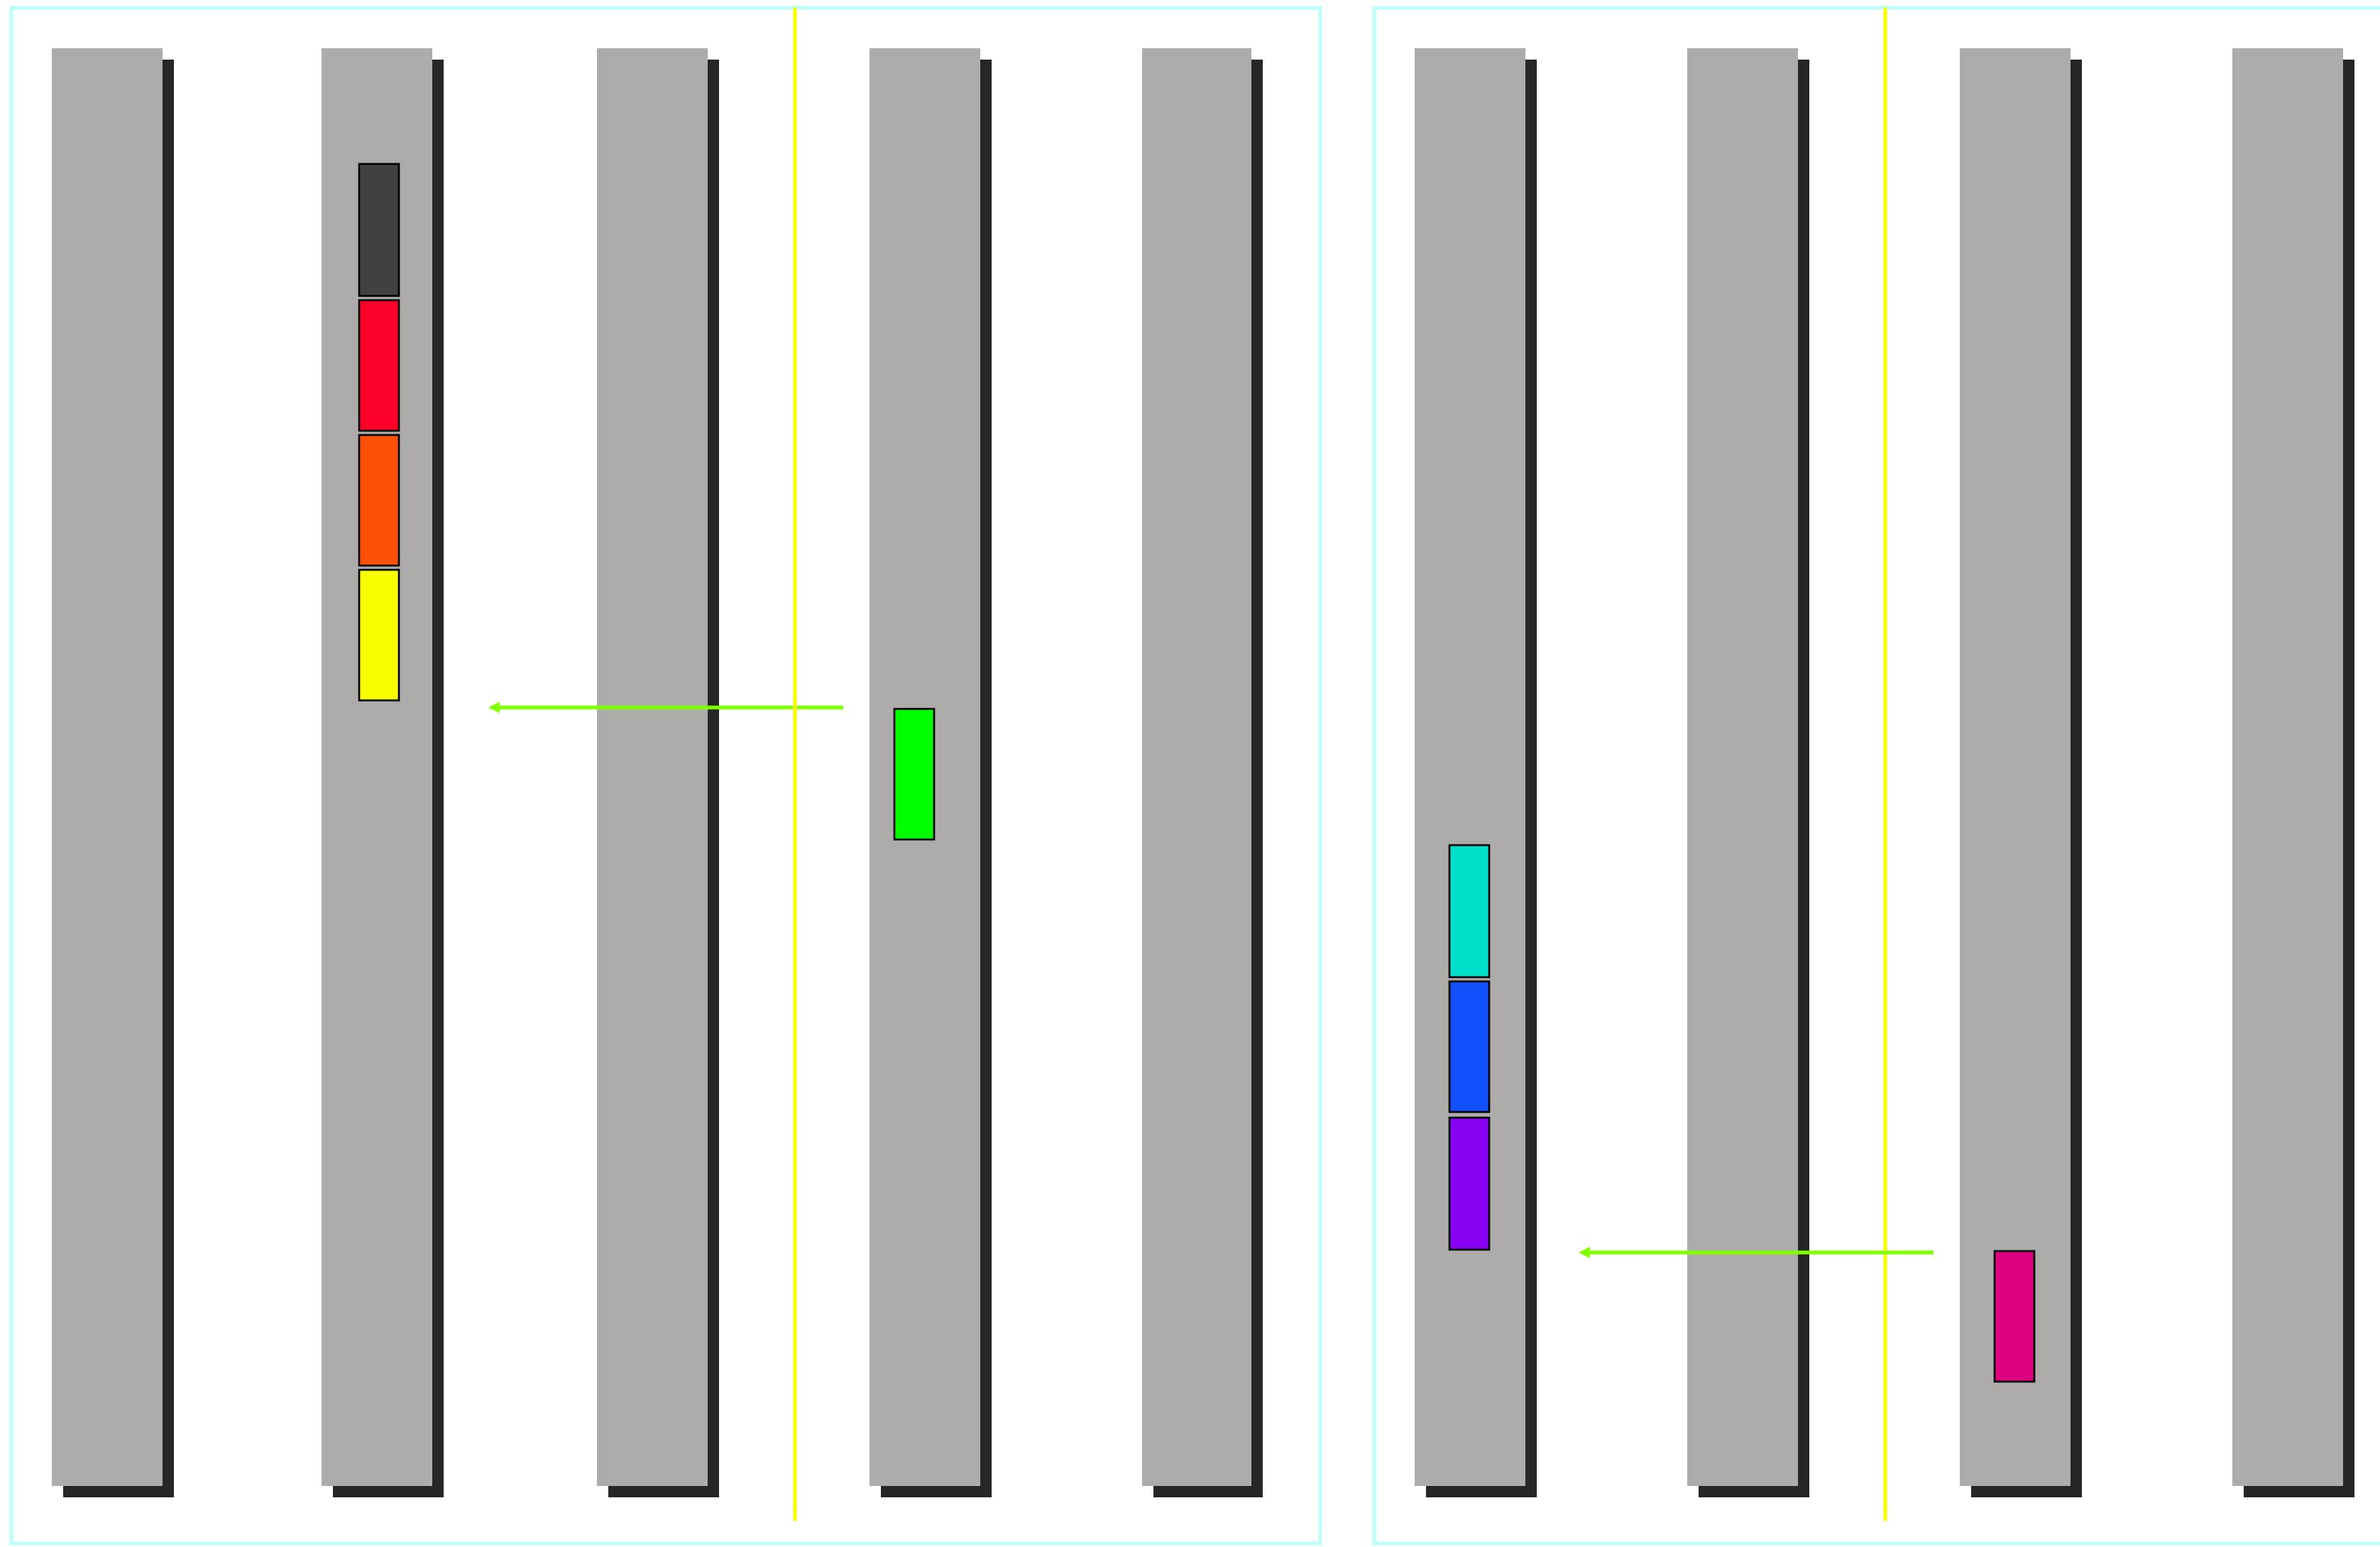


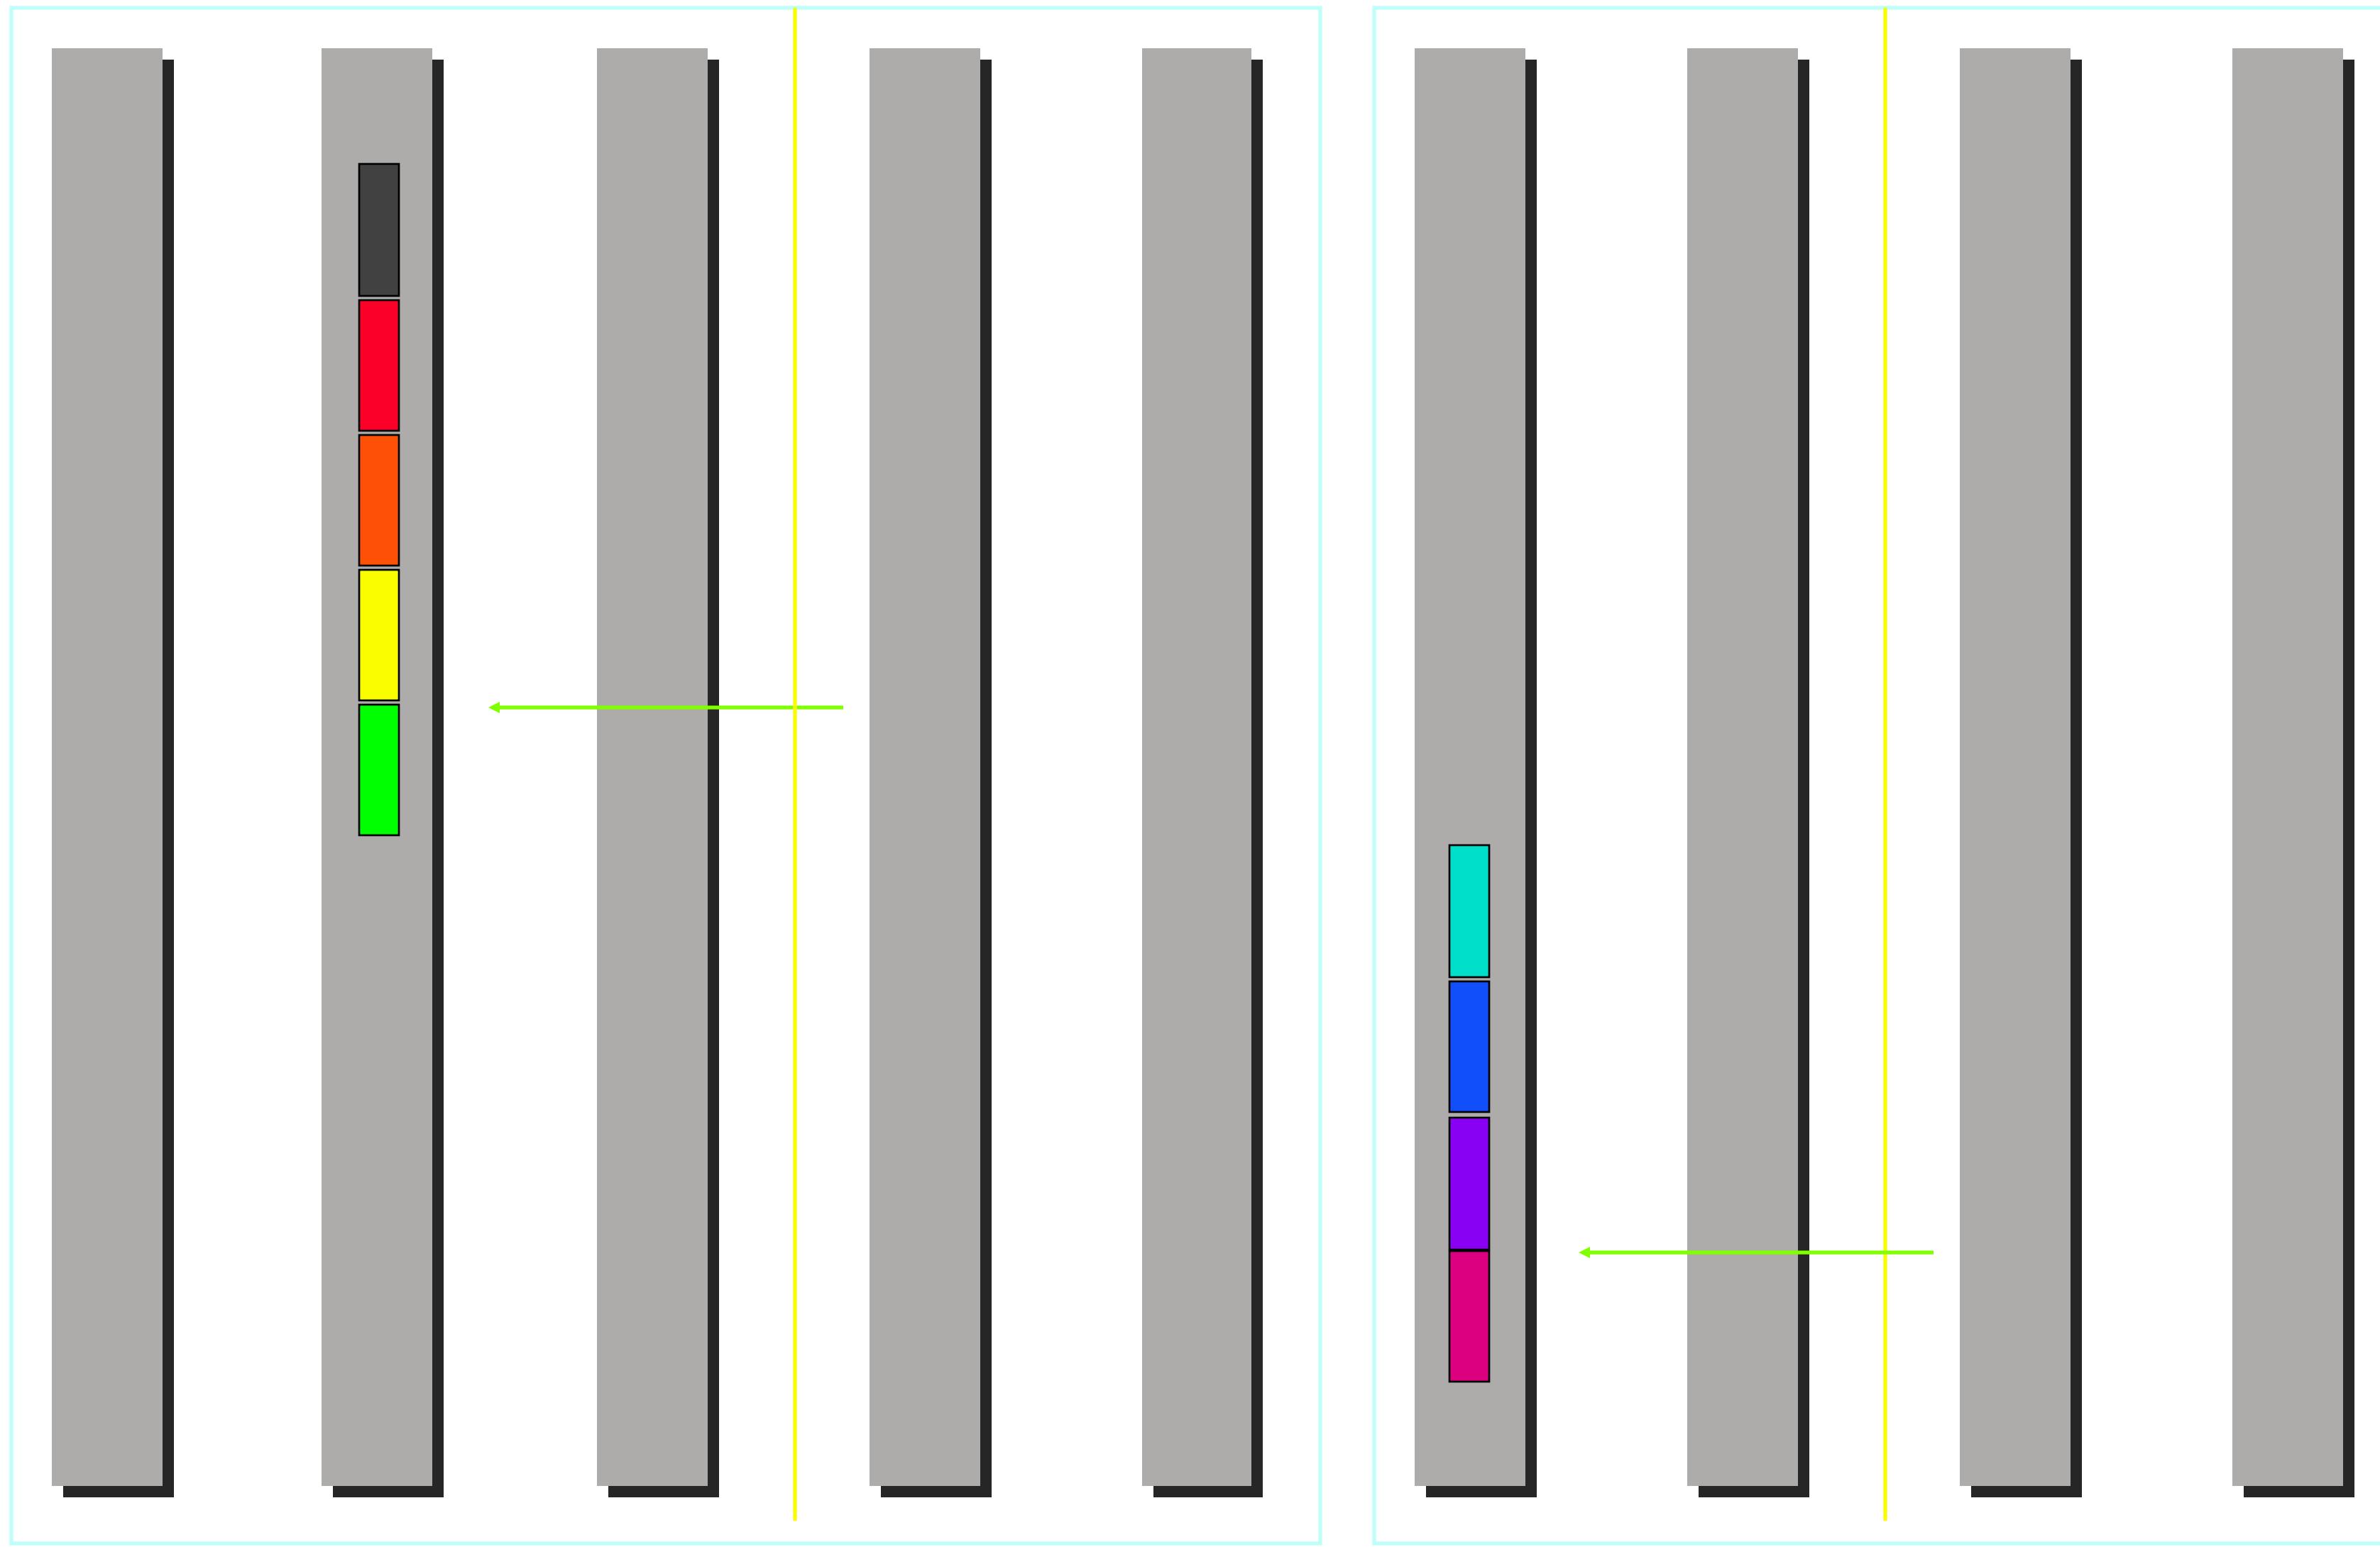


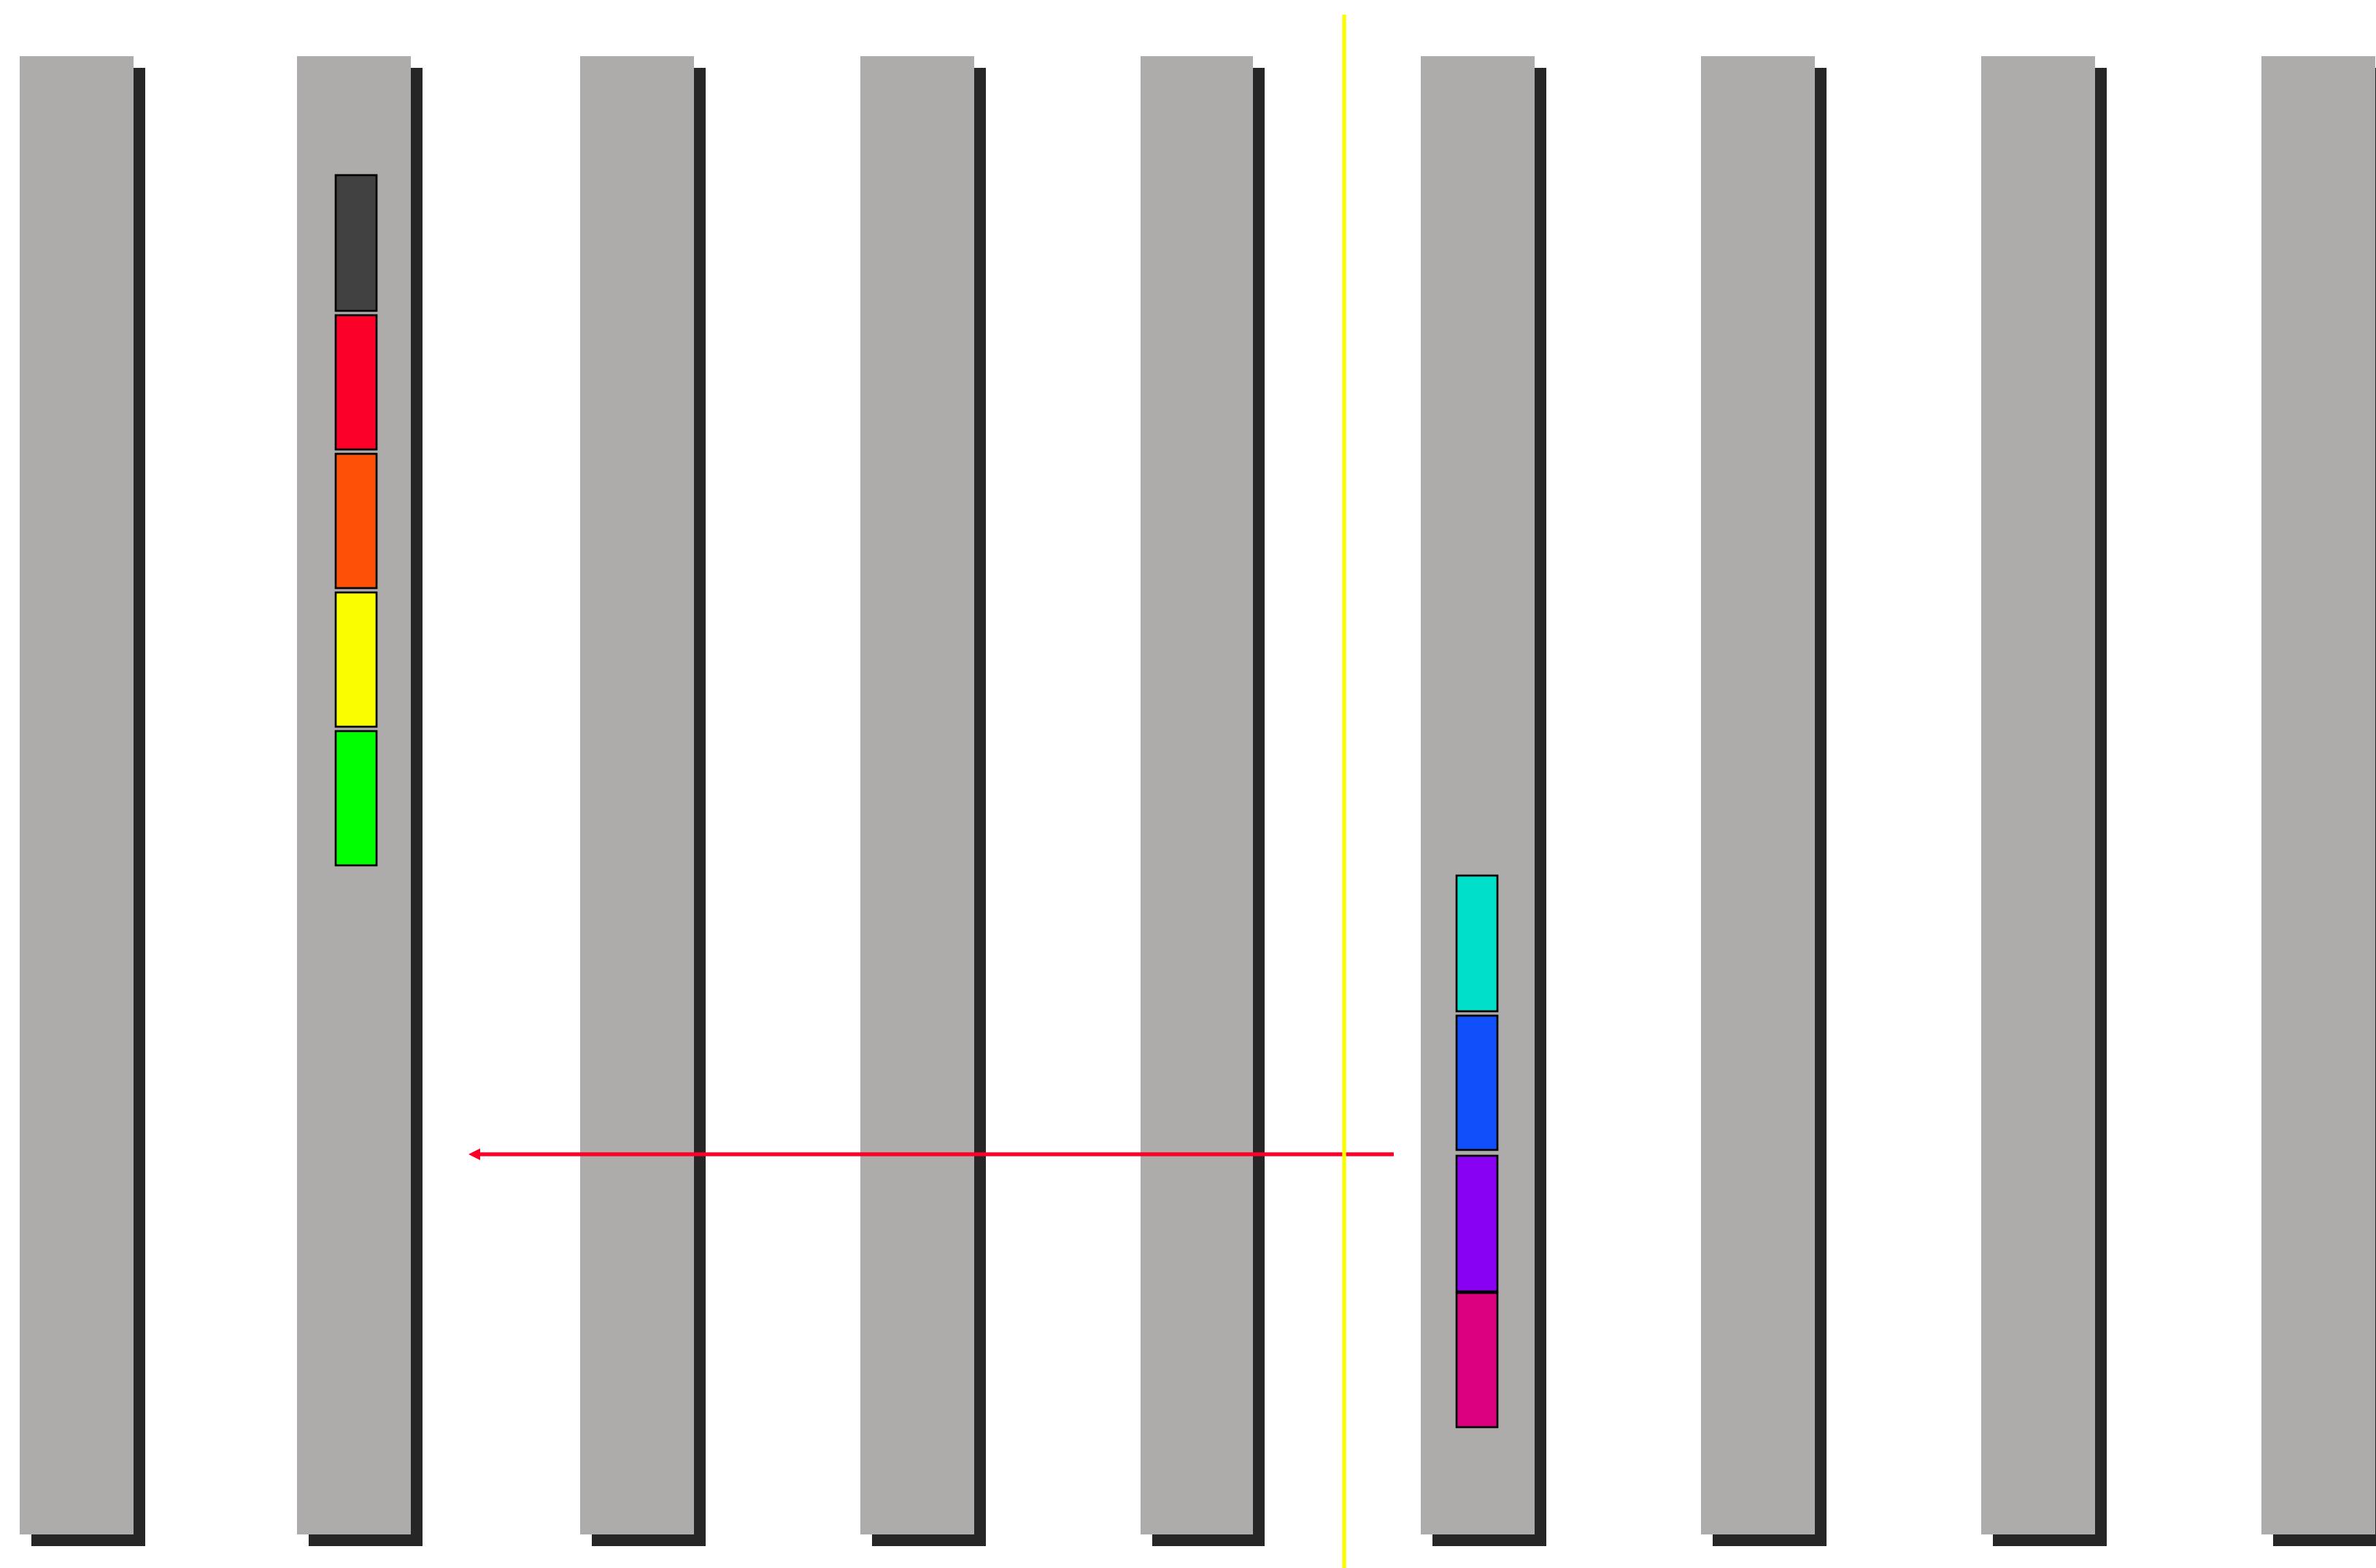


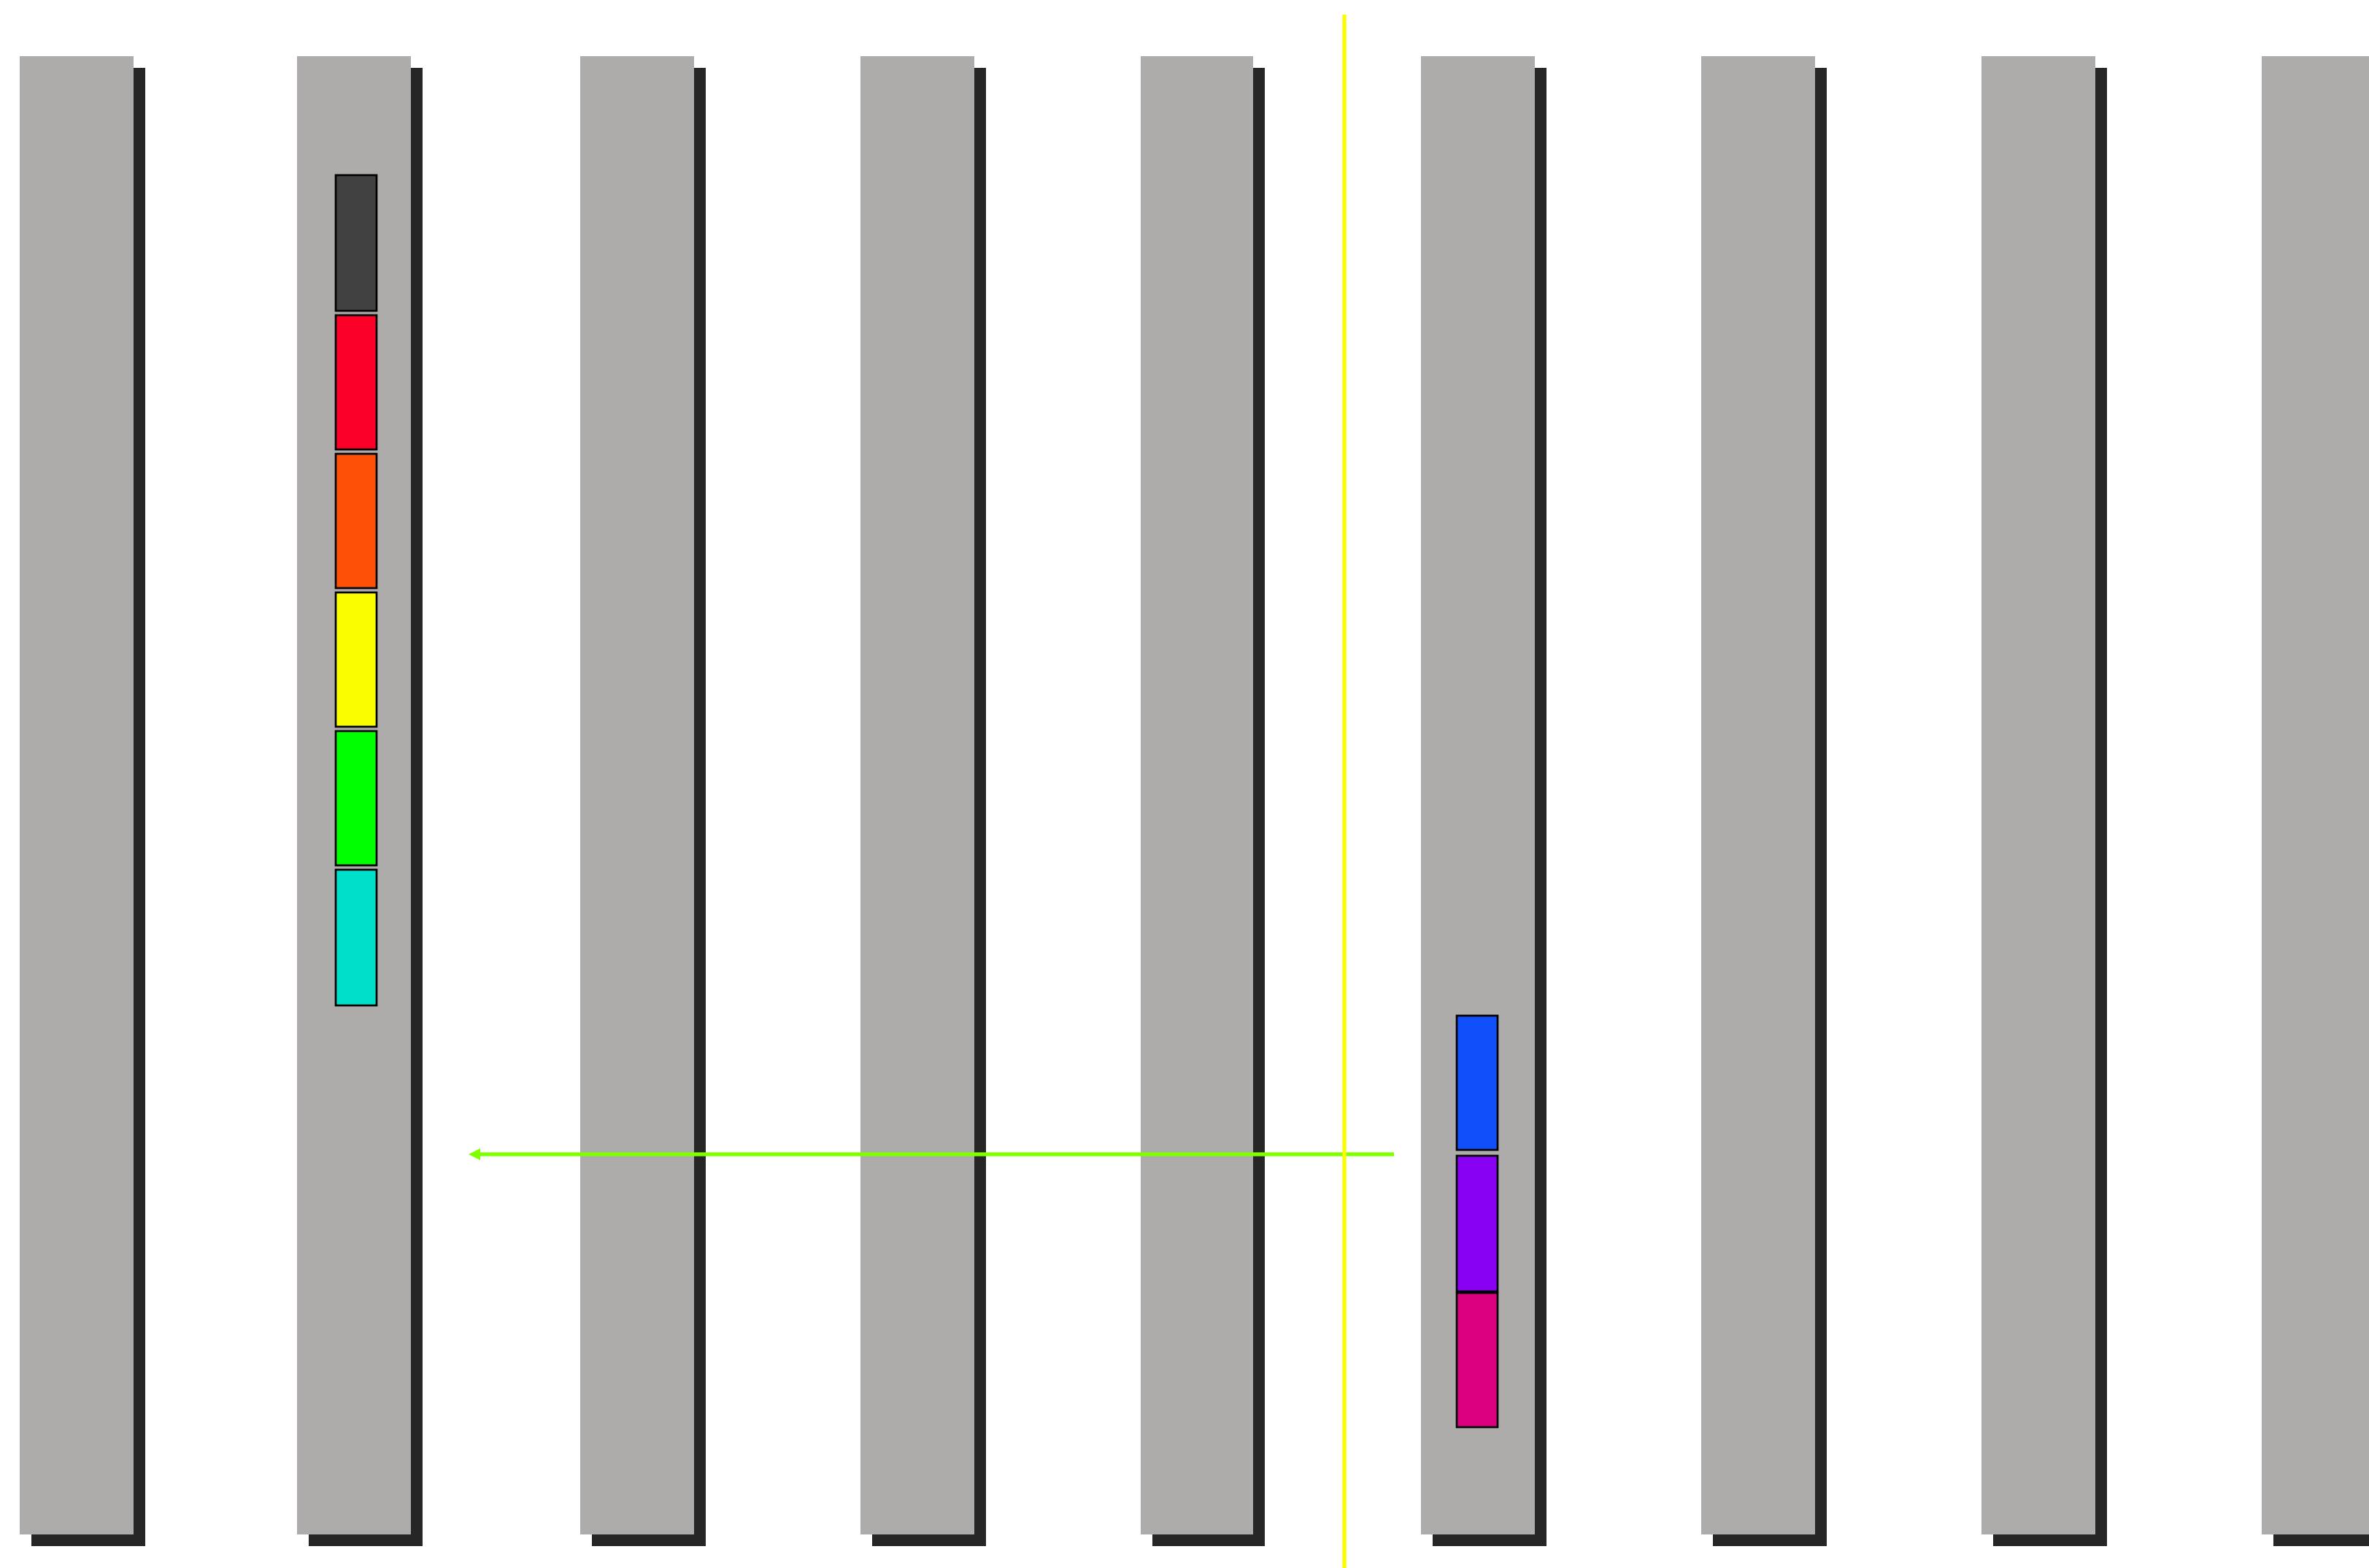


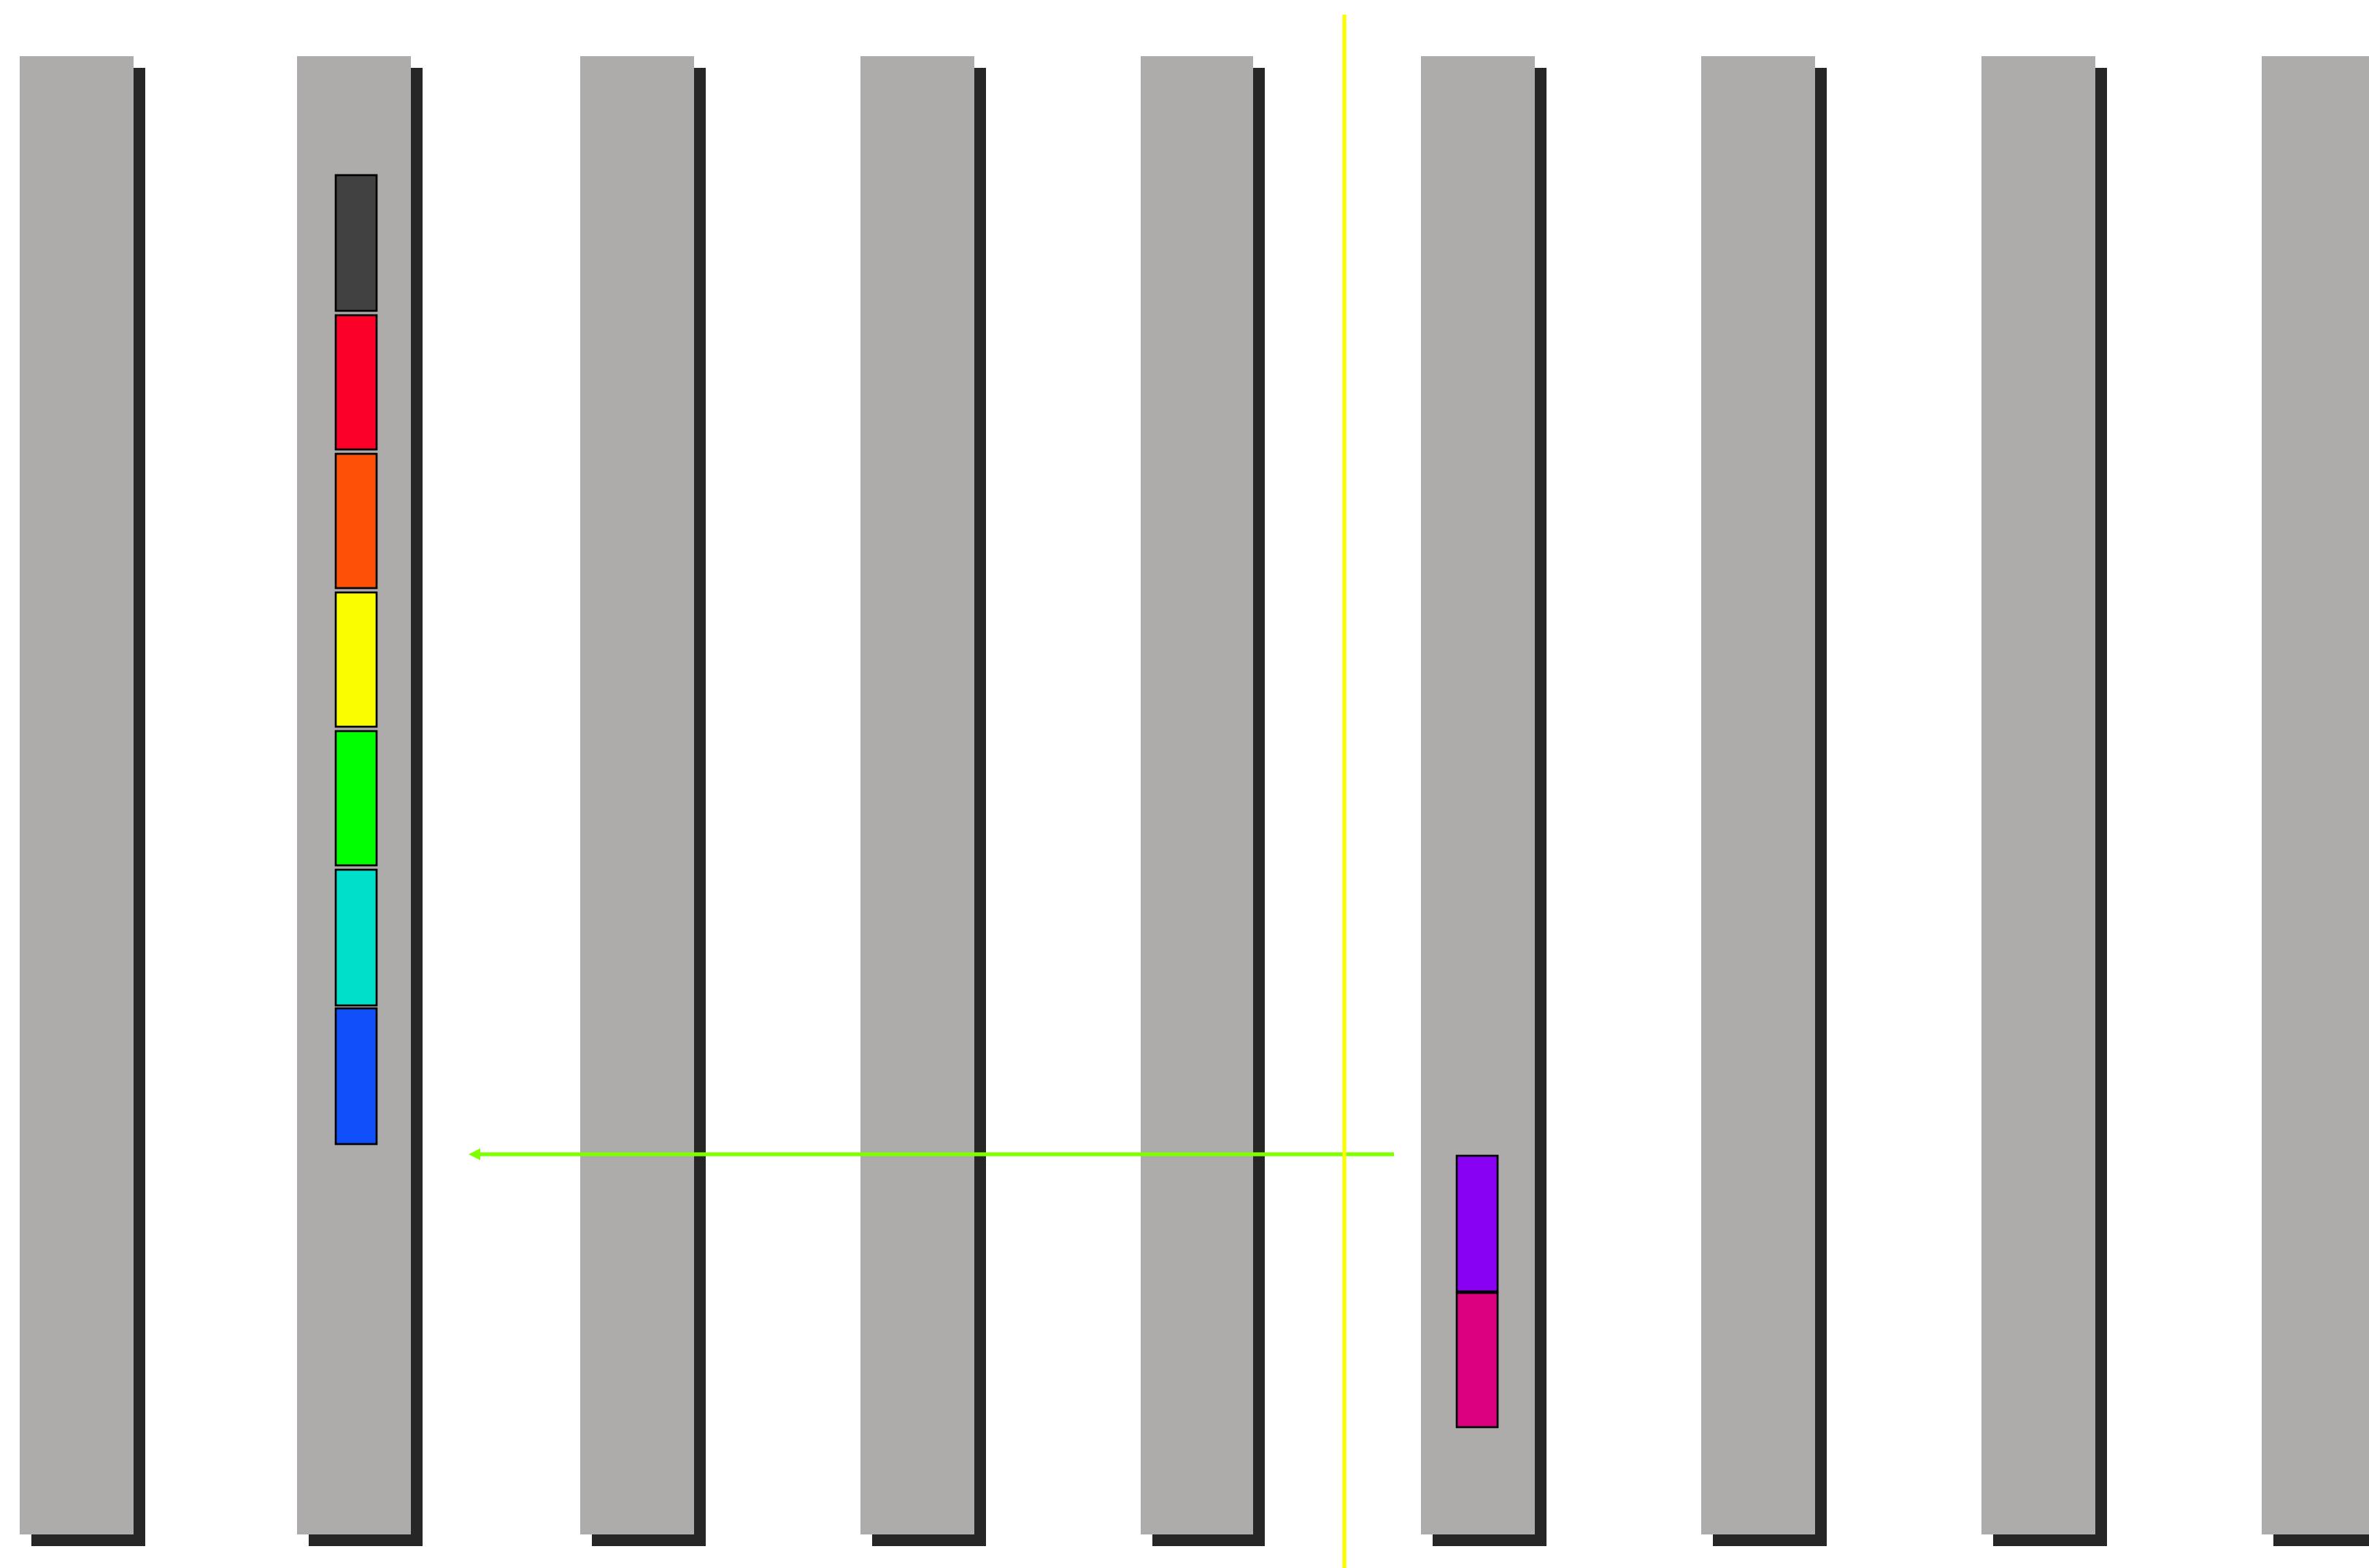


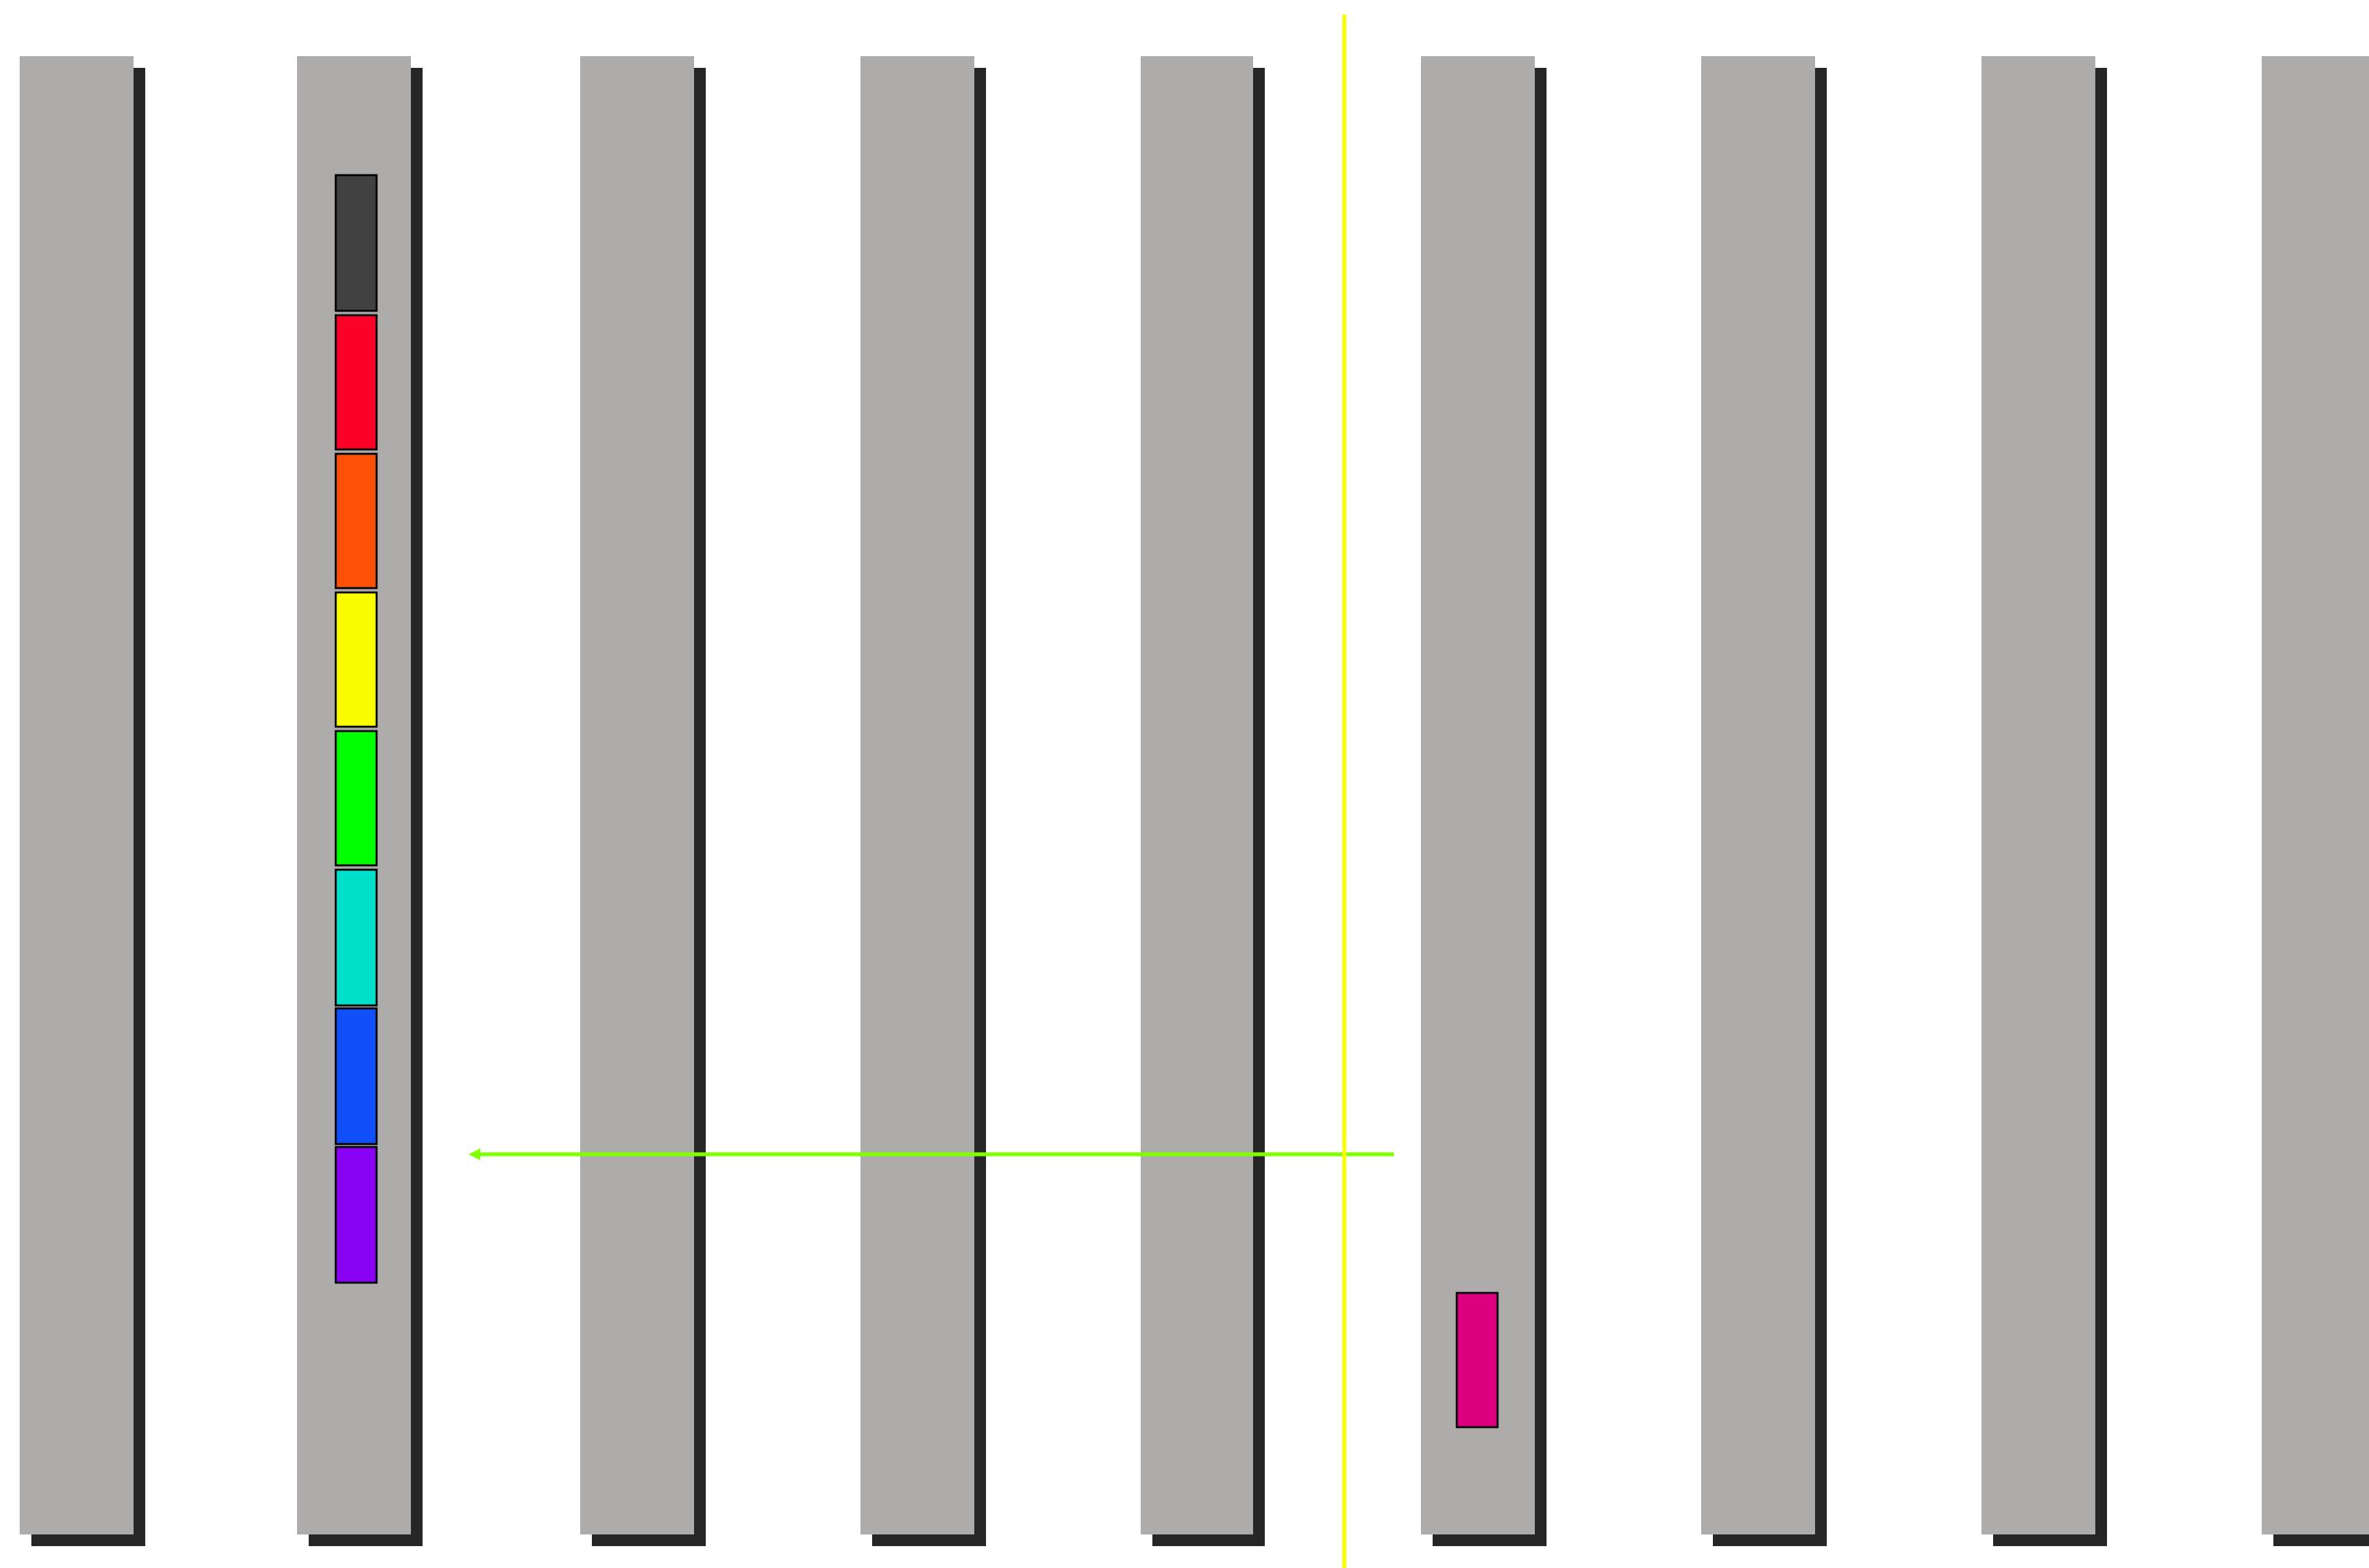


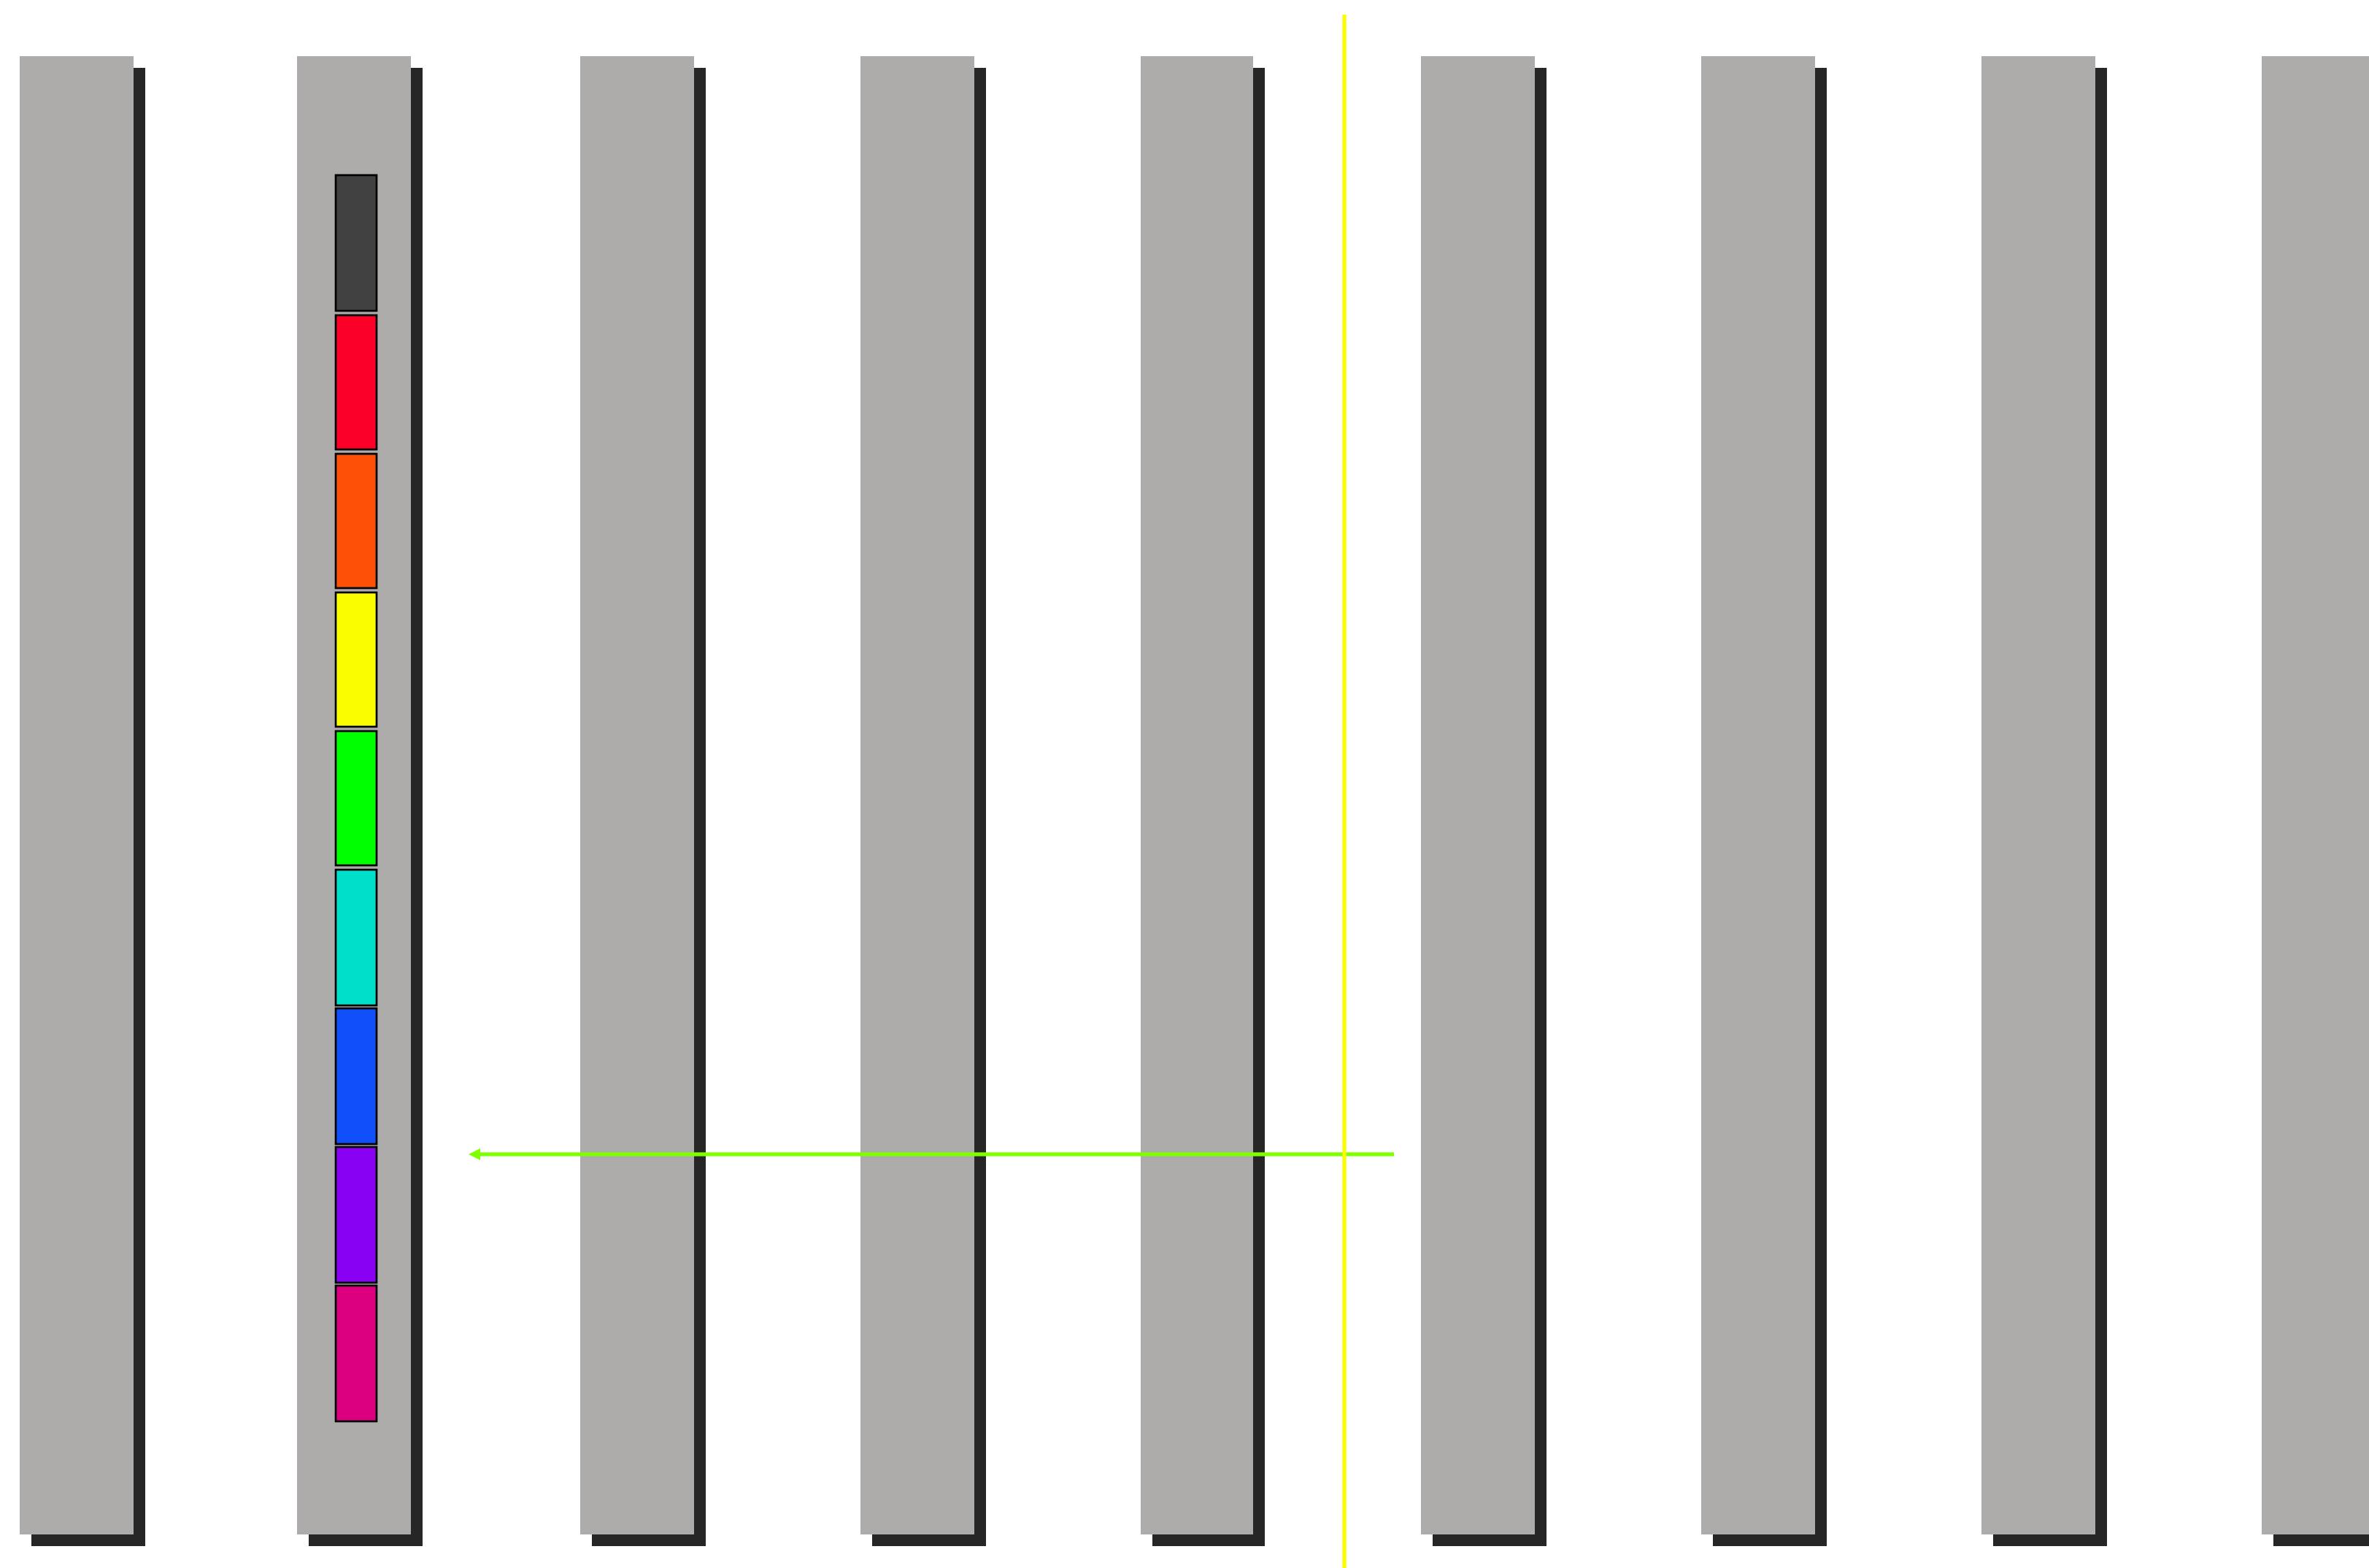


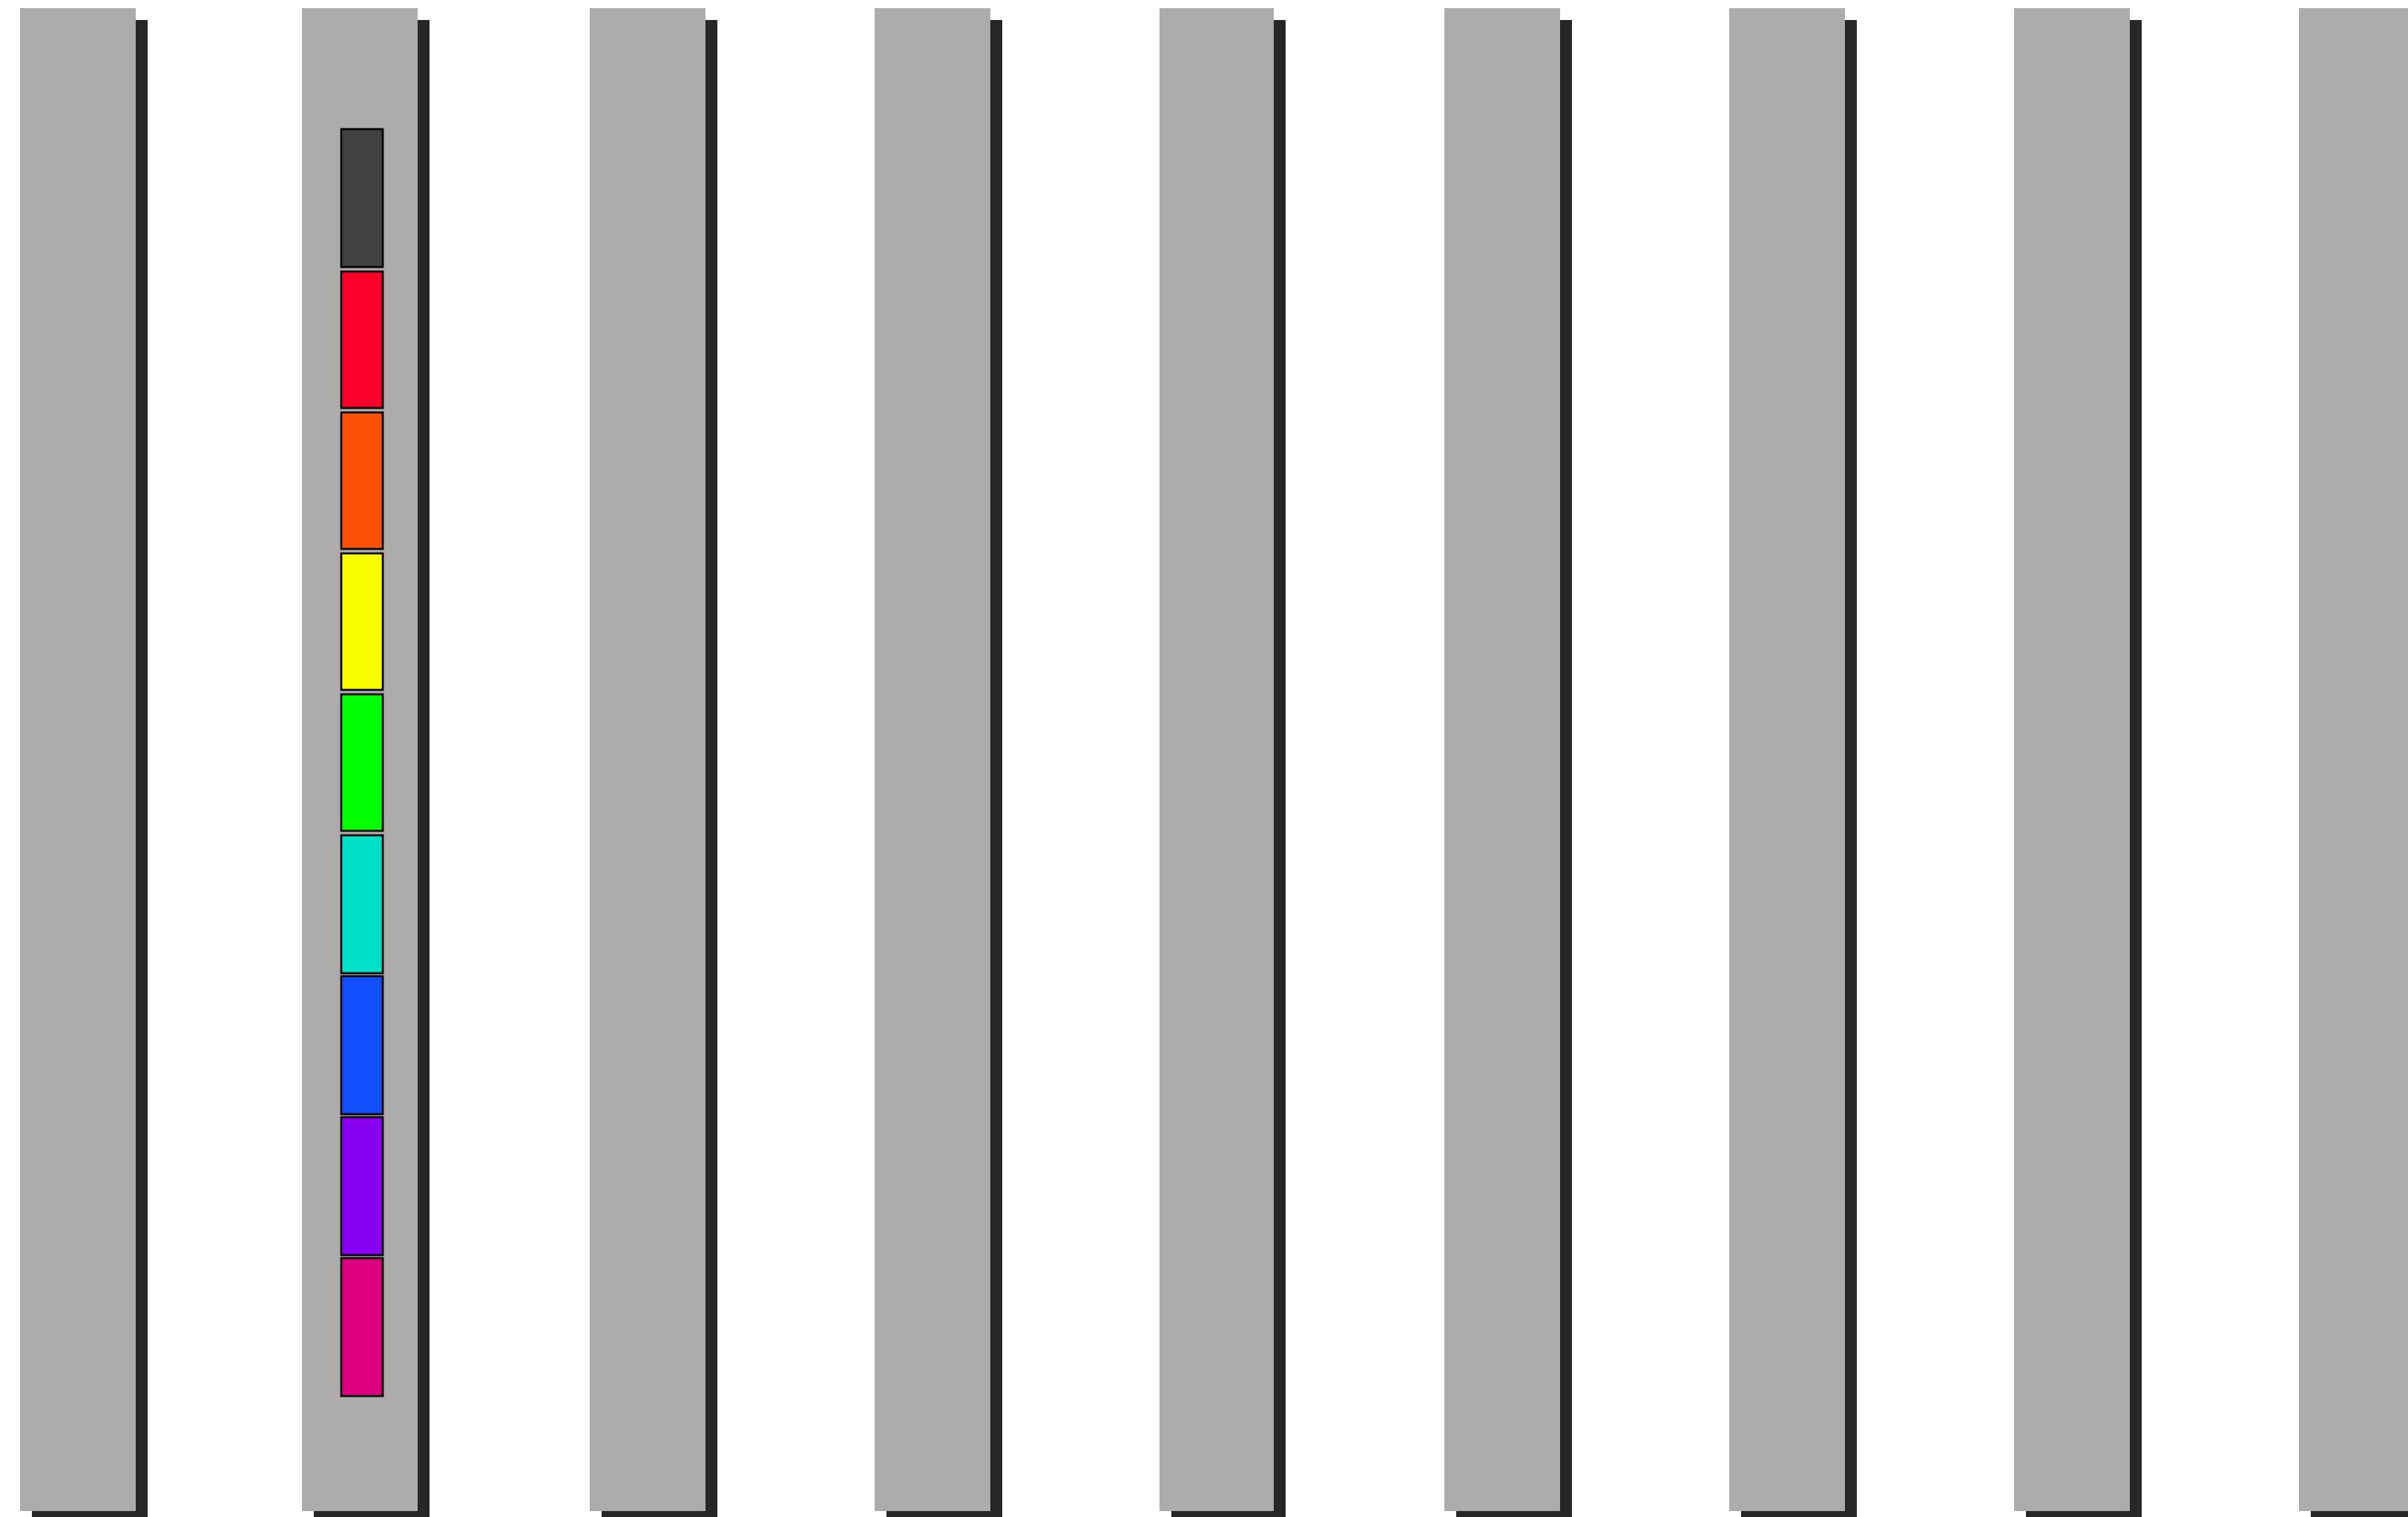












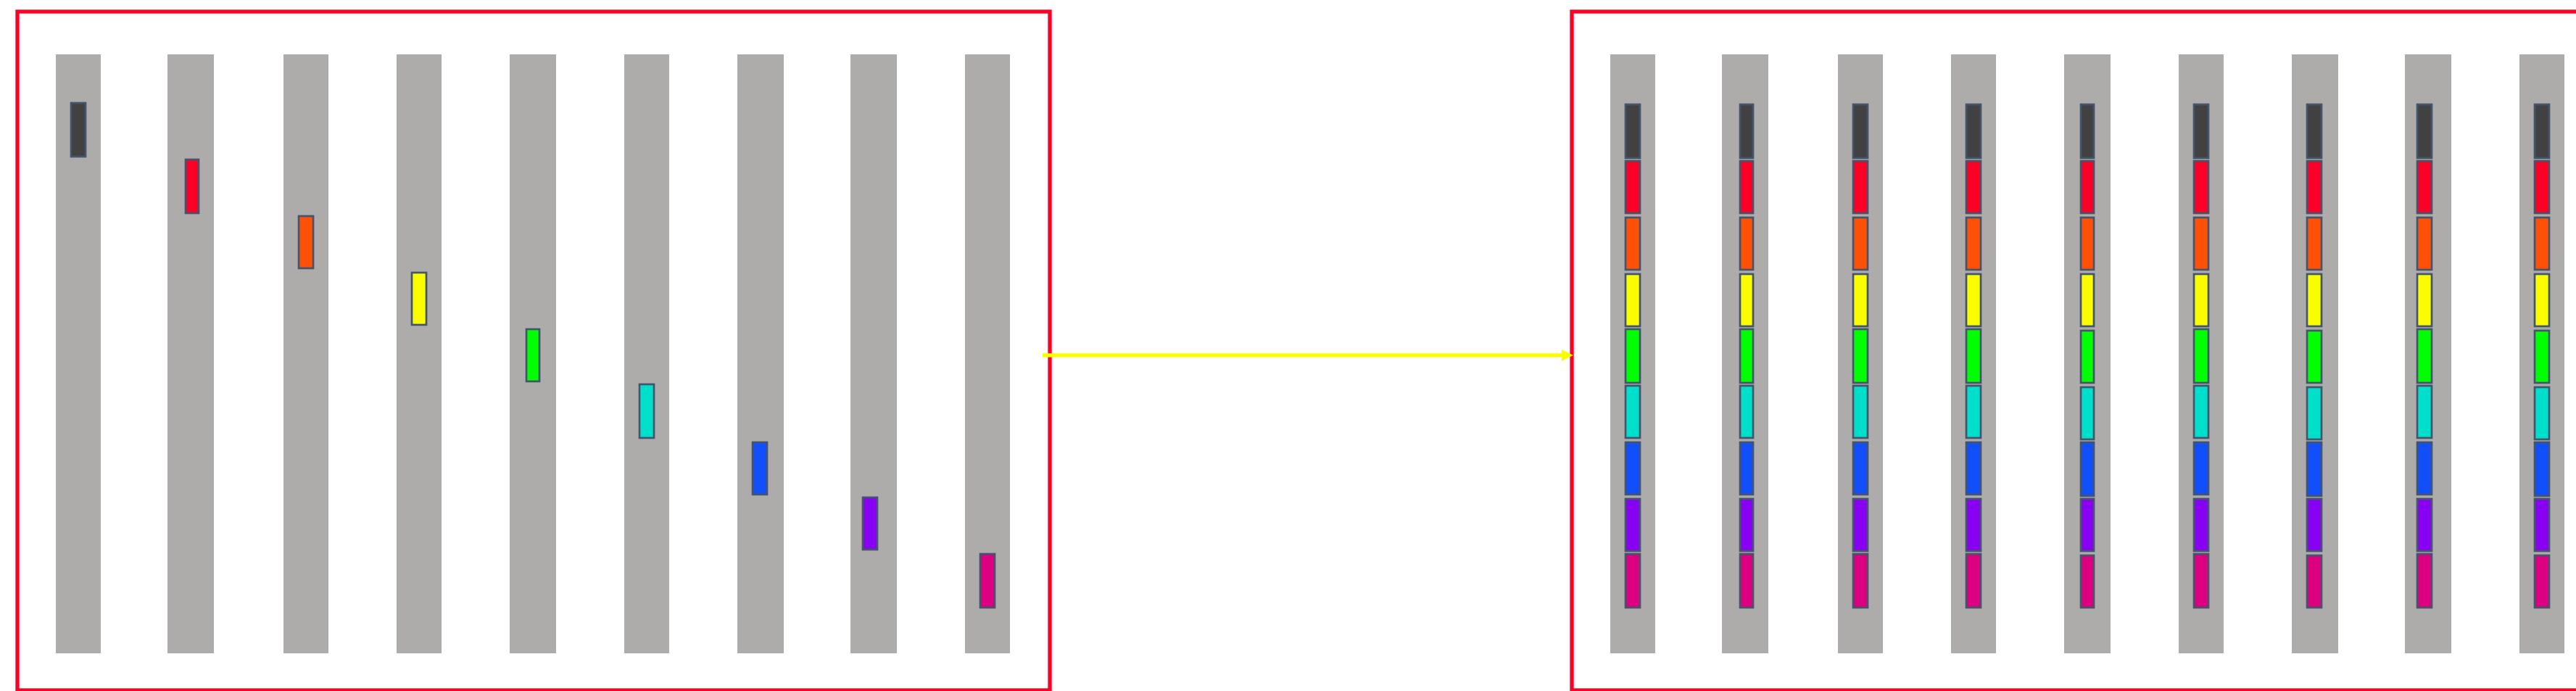
Cost of minimum spanning tree gather

- Assumption: power of two number of nodes

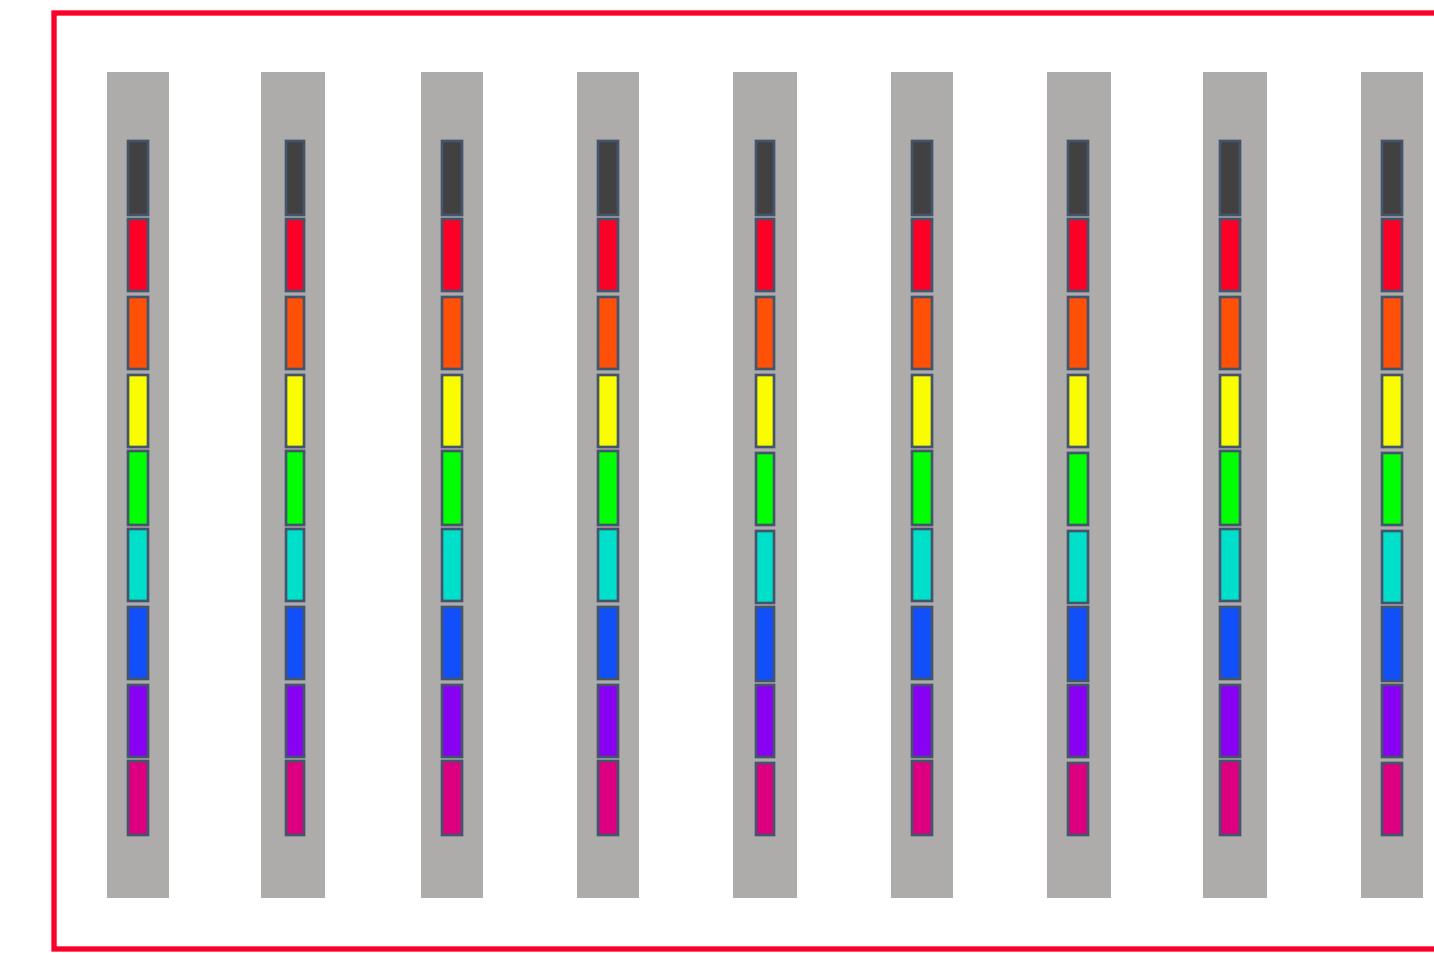
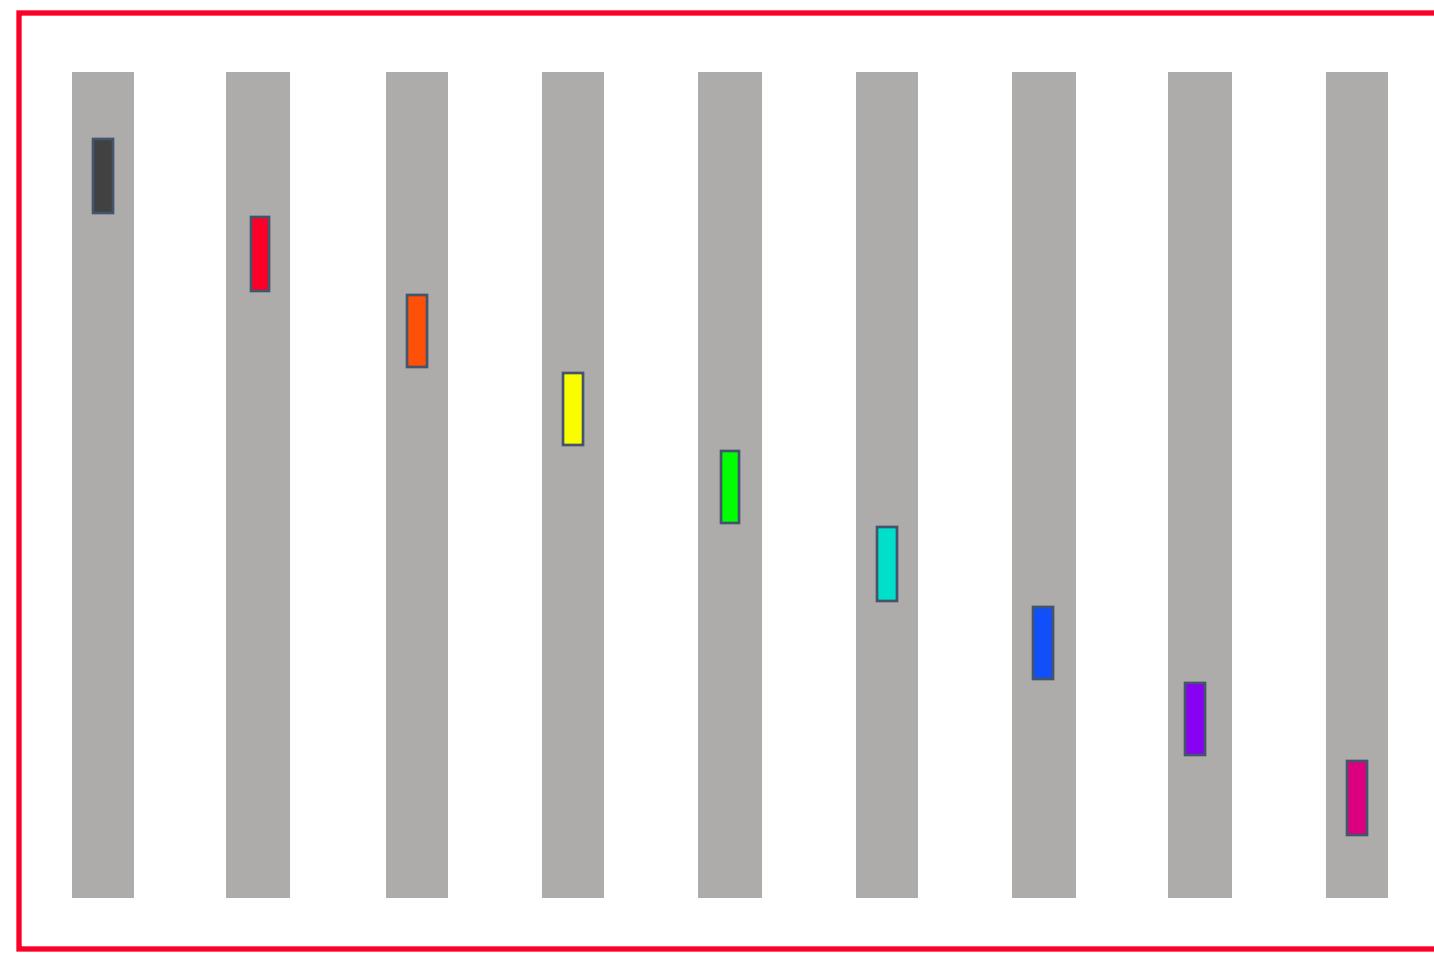
$$\begin{aligned} & \sum_{k=1}^{\log(p)} \left(\alpha + \frac{n}{2^k} \beta \right) \\ &= \\ & \log(p) - \alpha + \frac{p-1}{p} n \beta \end{aligned}$$

Using the building blocks

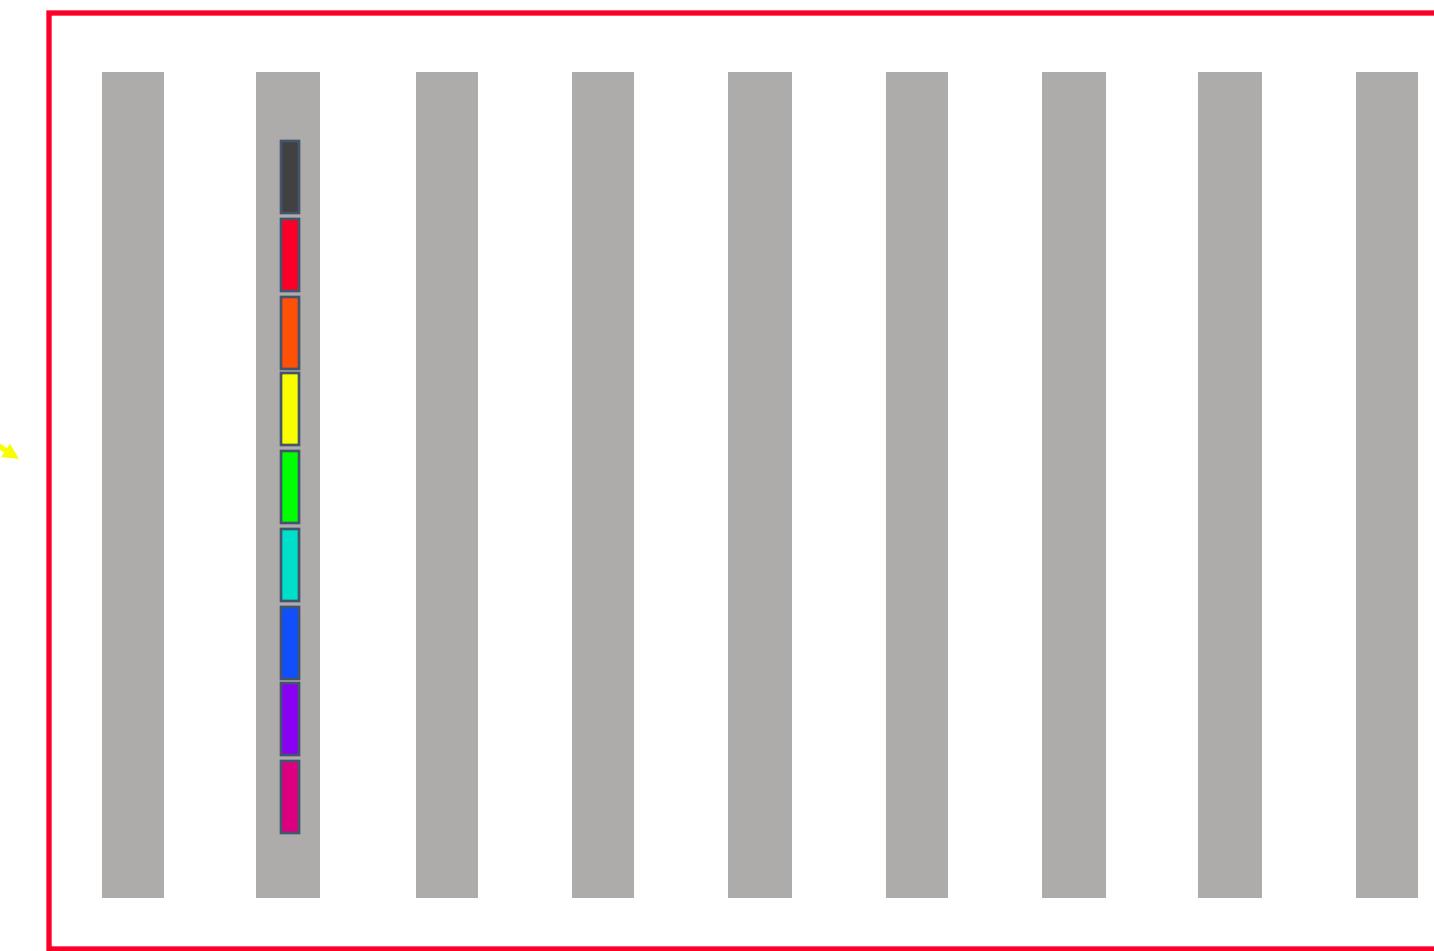
Allgather



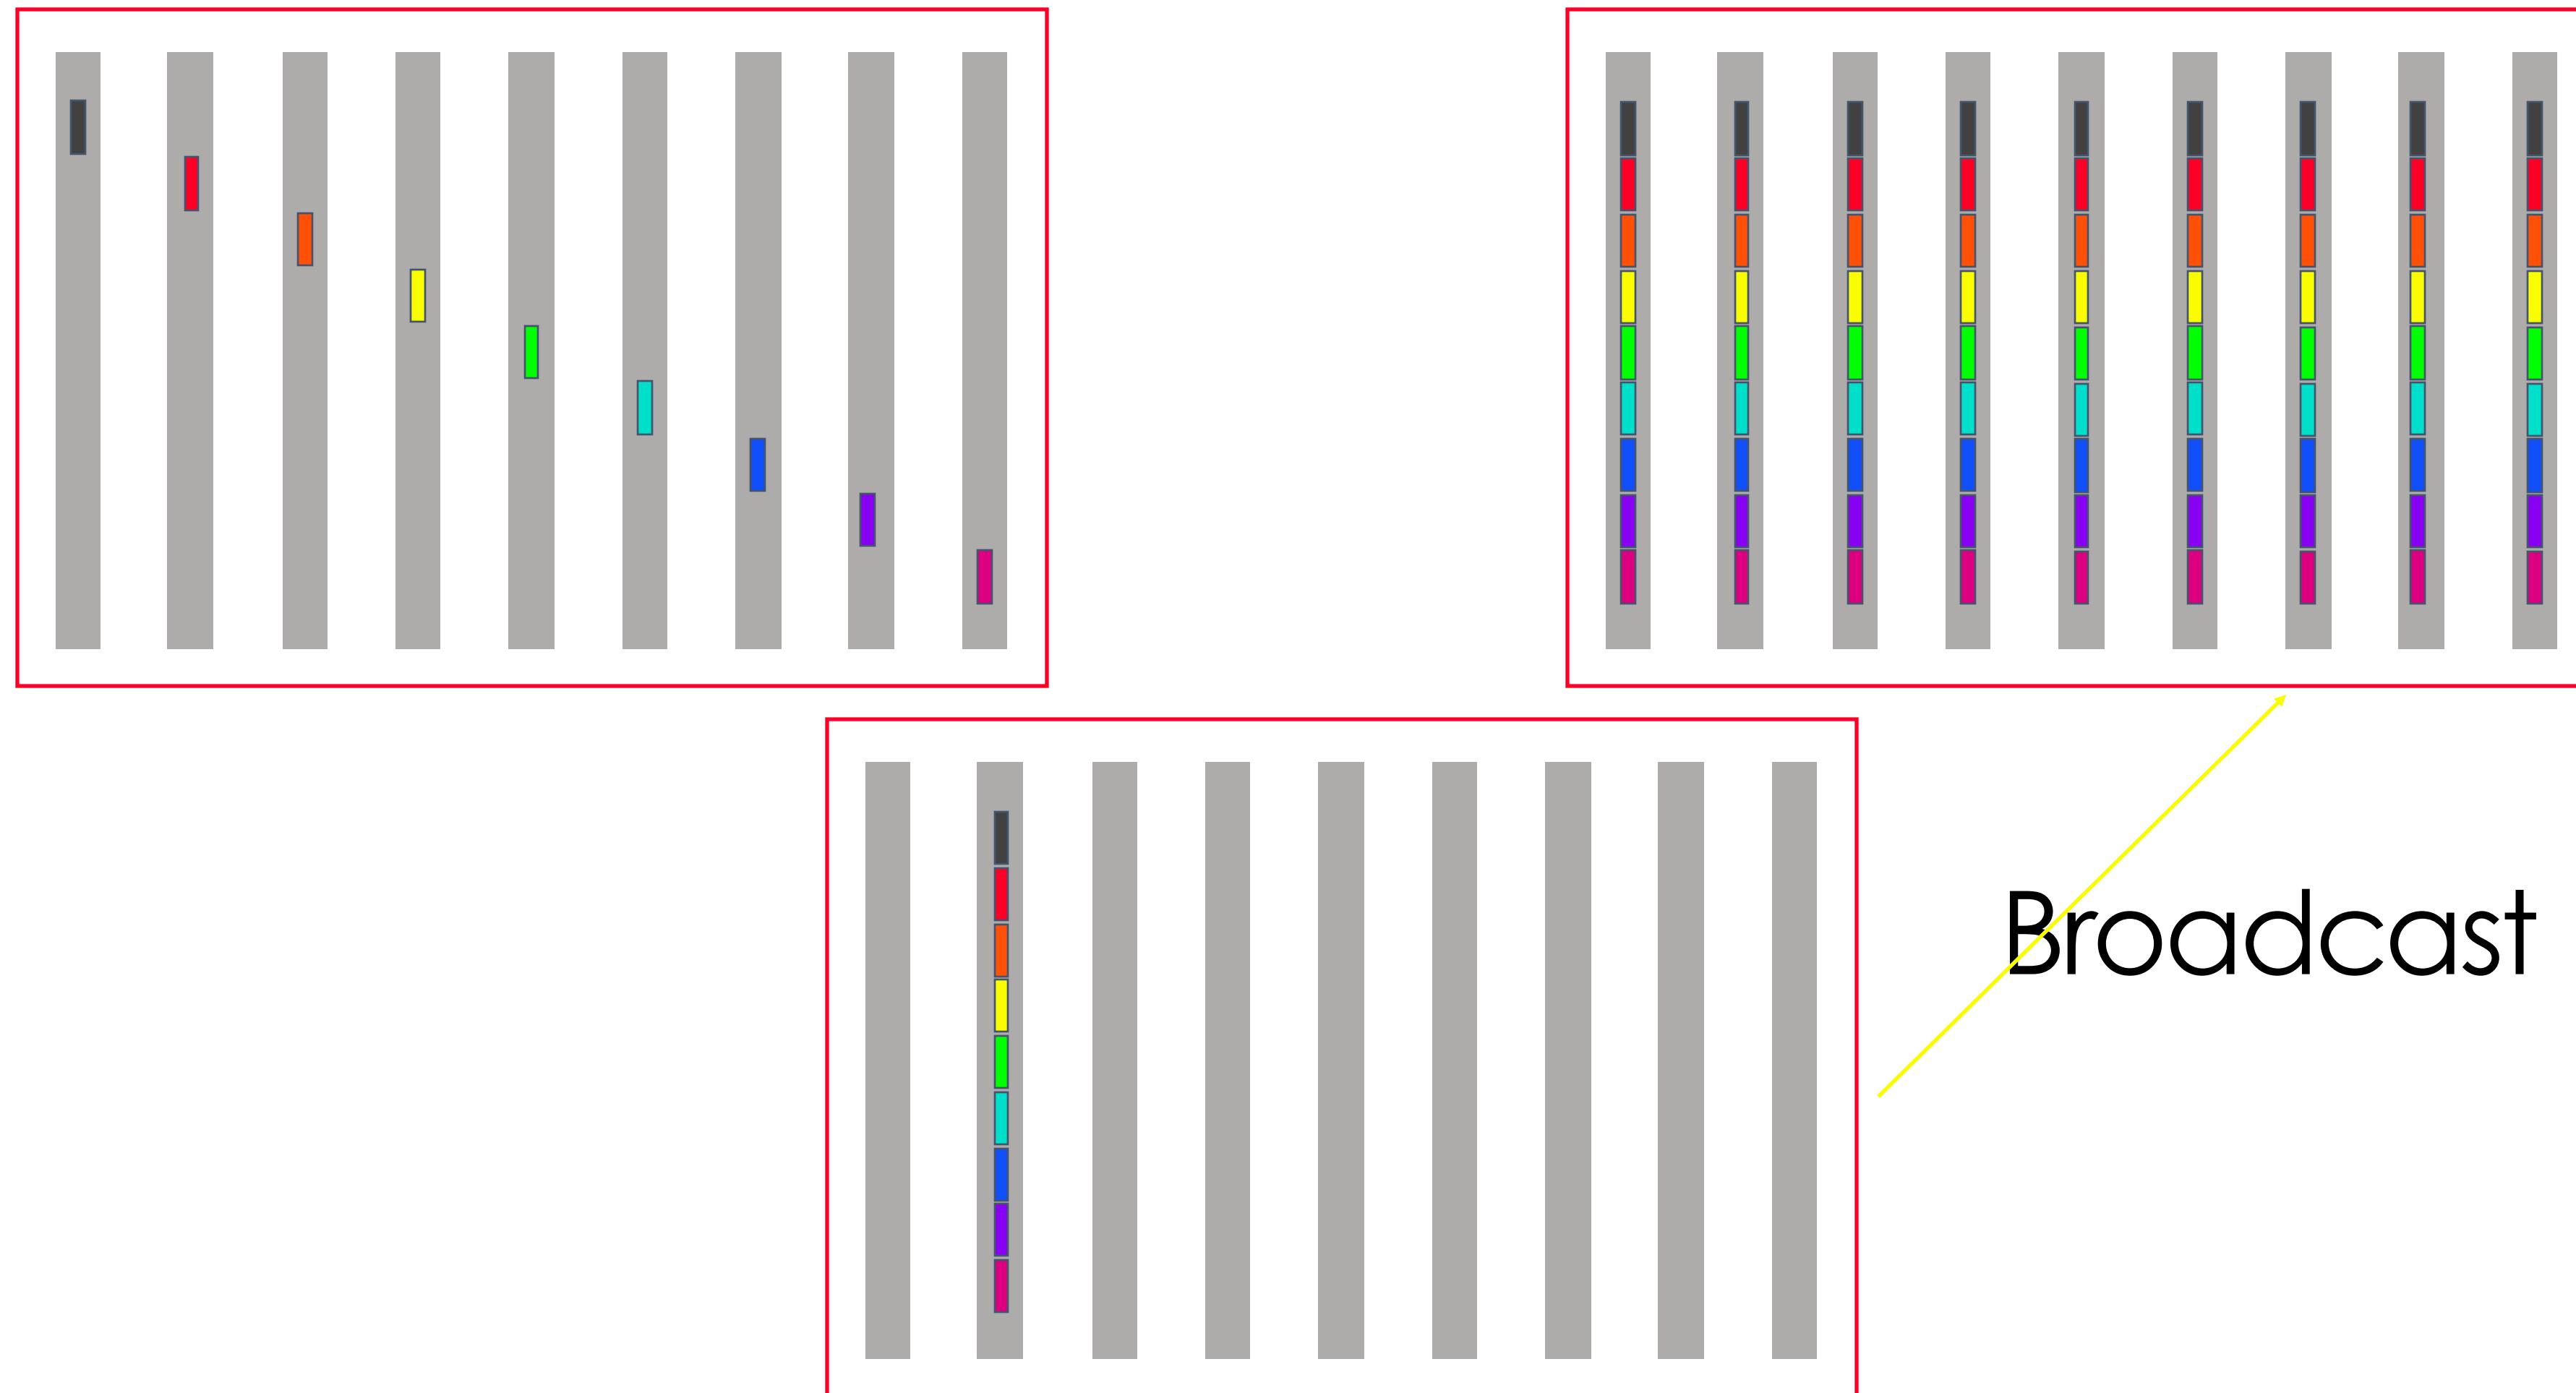
Allgather



Gather



Allgather (short vector)



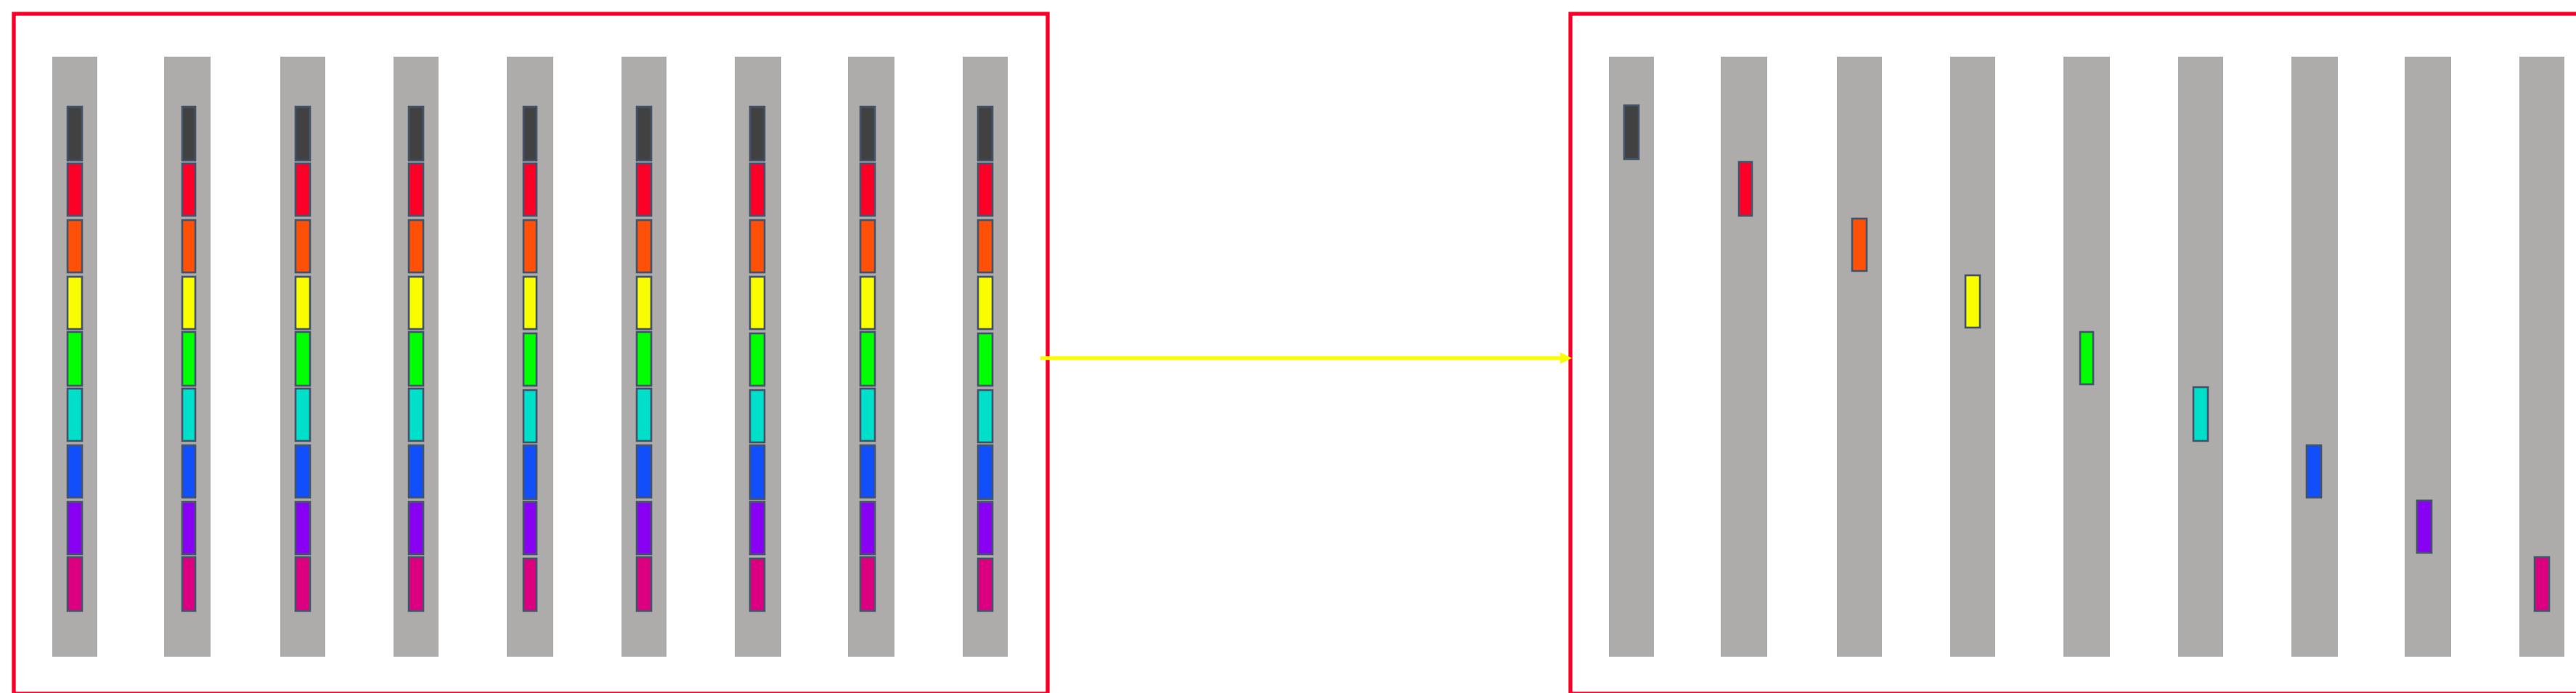
Cost of gather/broadcast allgather

- Assumption: power of two number of nodes

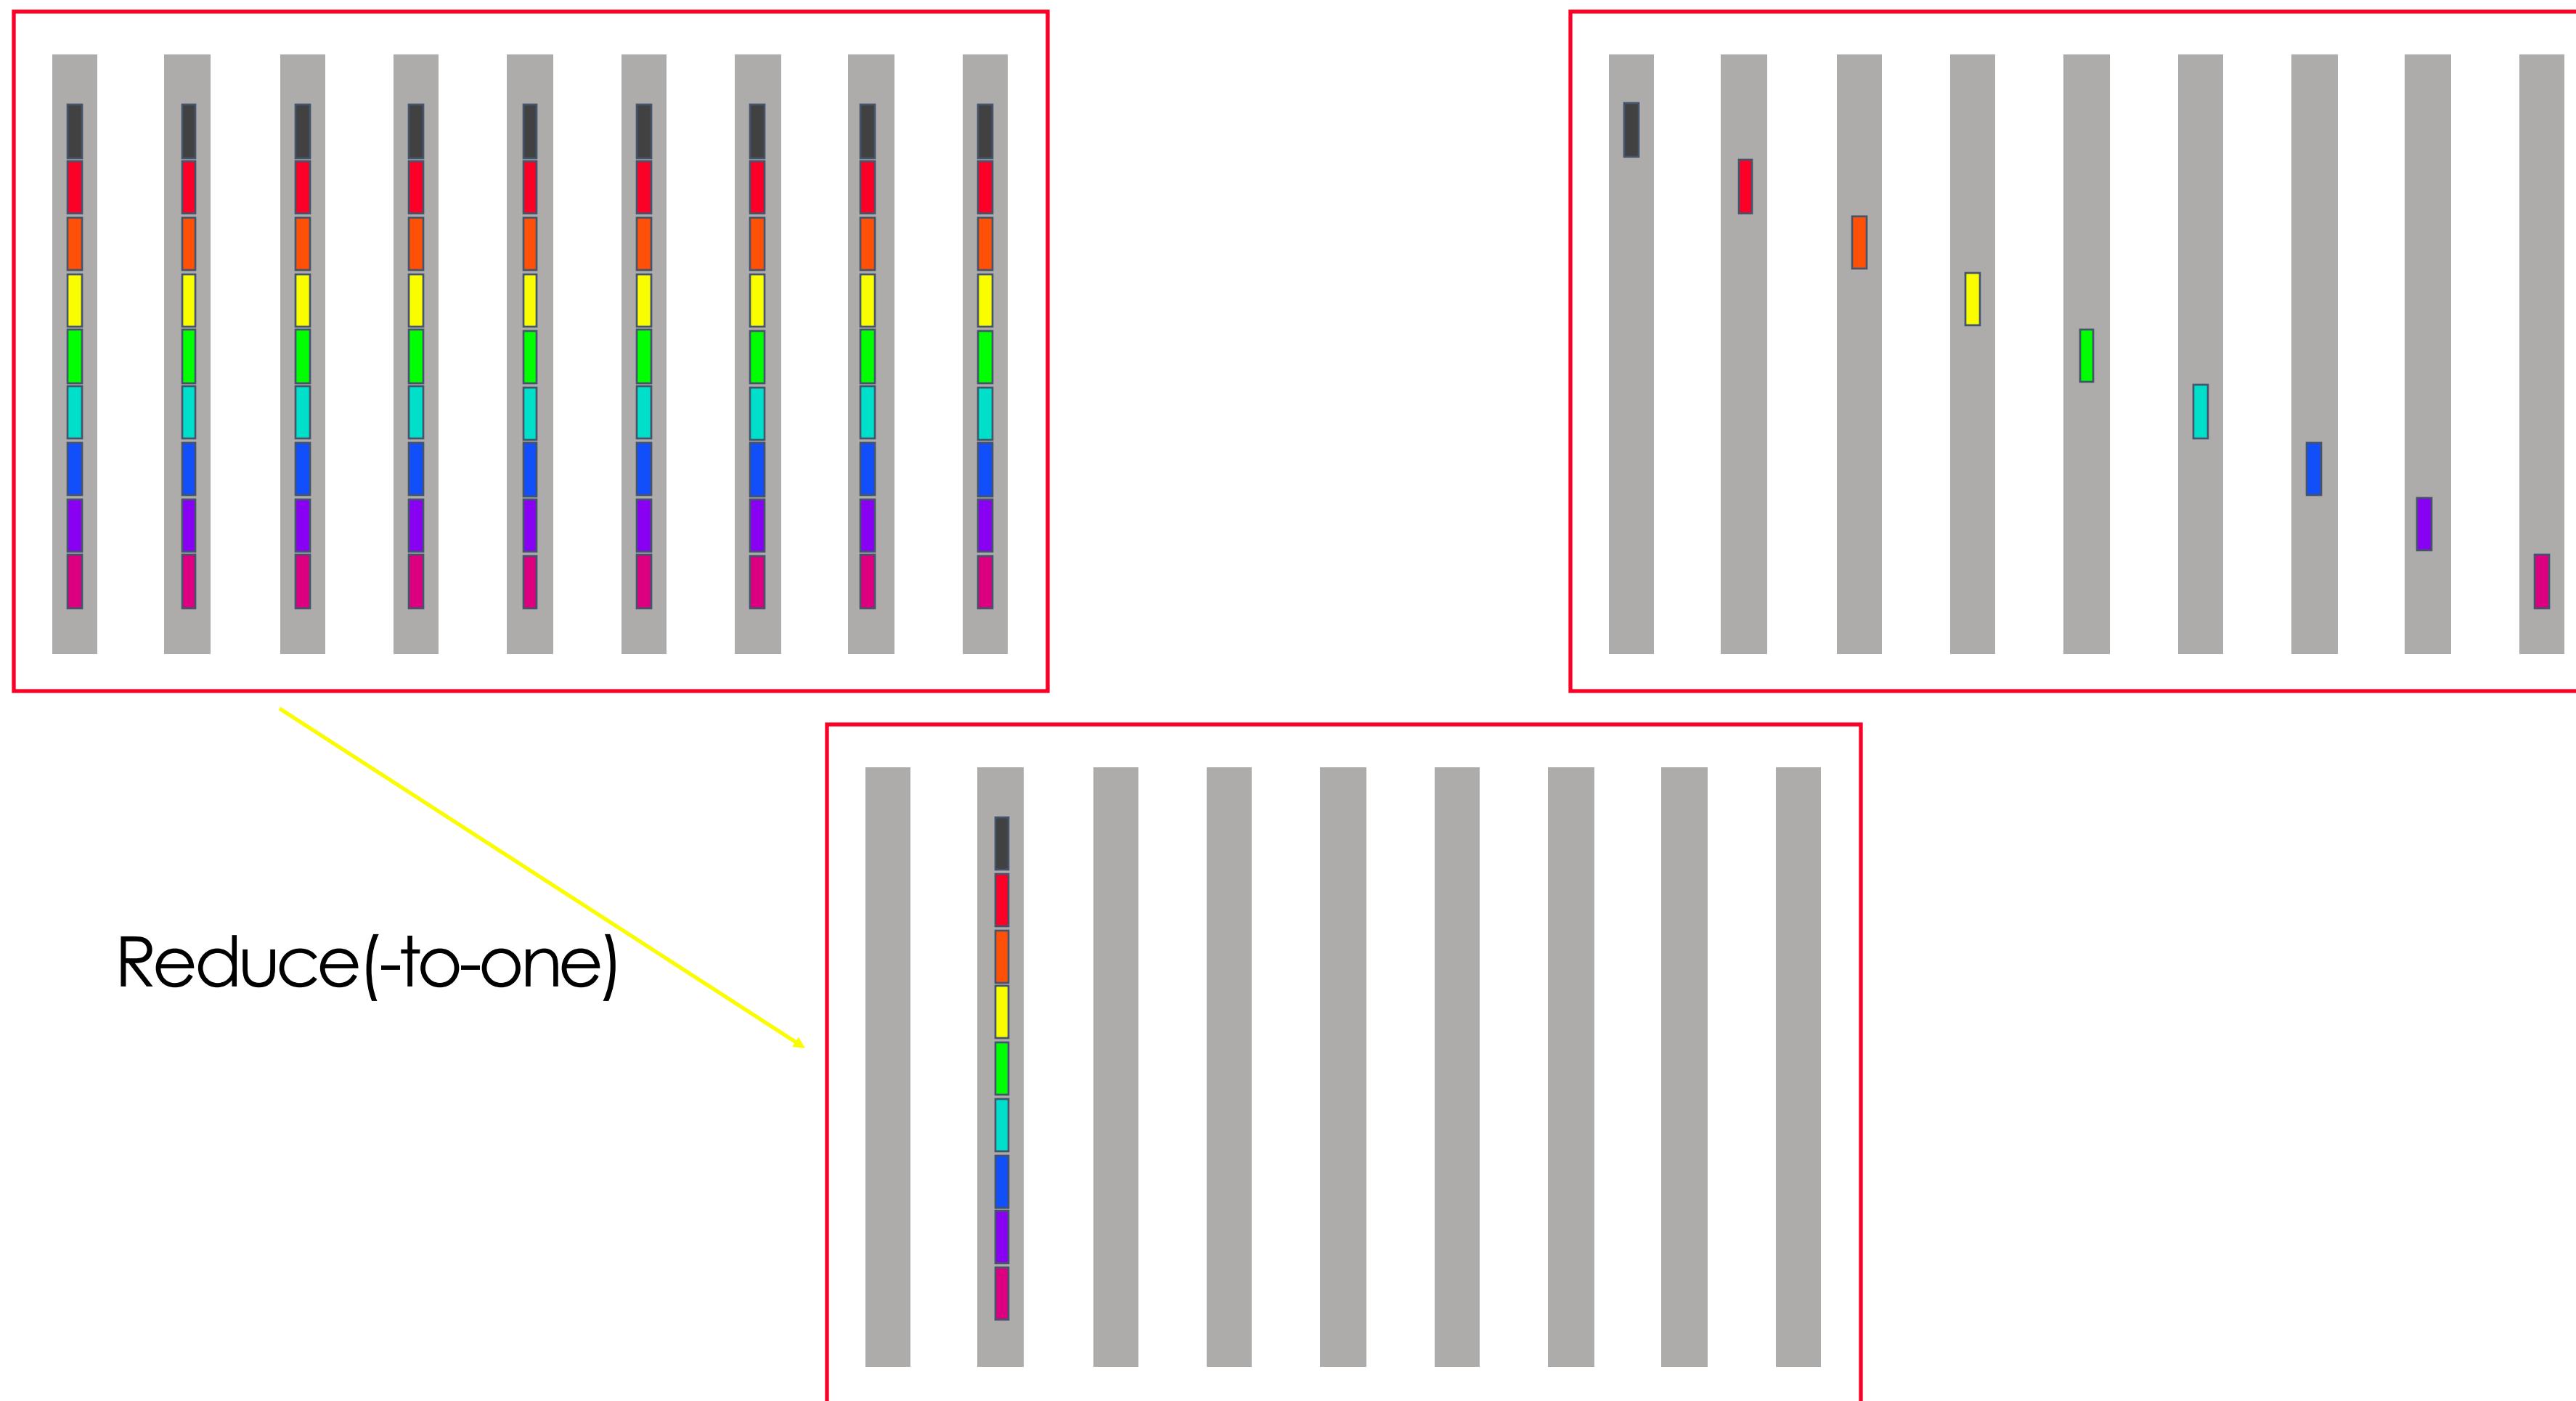
gather $\log(p)\alpha + \frac{p-1}{p}n\beta$

broadcast $\frac{\log(p)(\alpha + n\beta)}{2\log(p)\alpha + \left(\frac{p-1}{p} + \log(p)\right)n\beta}$

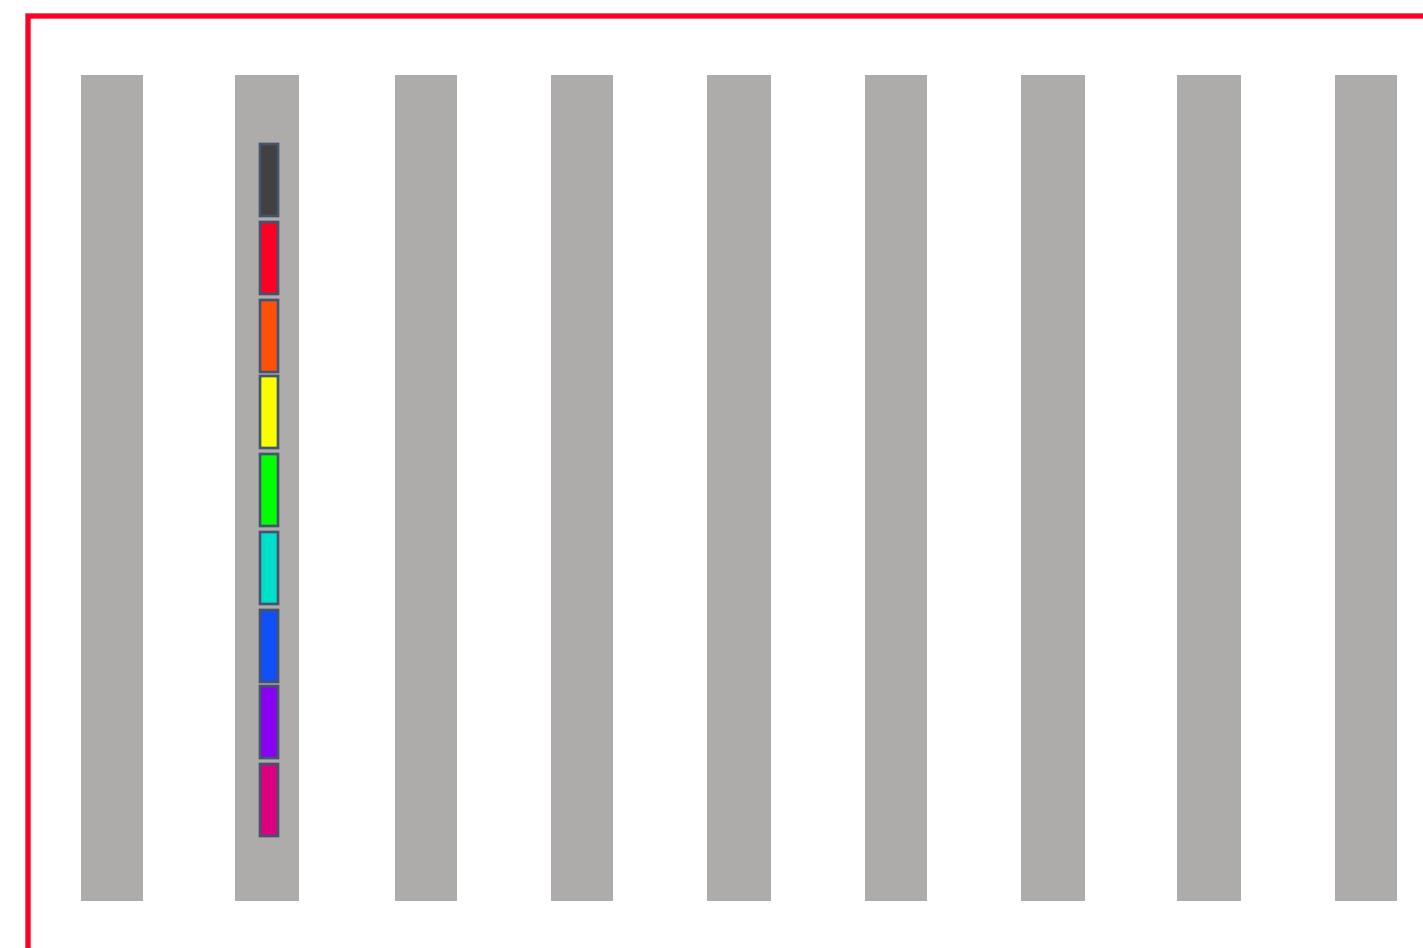
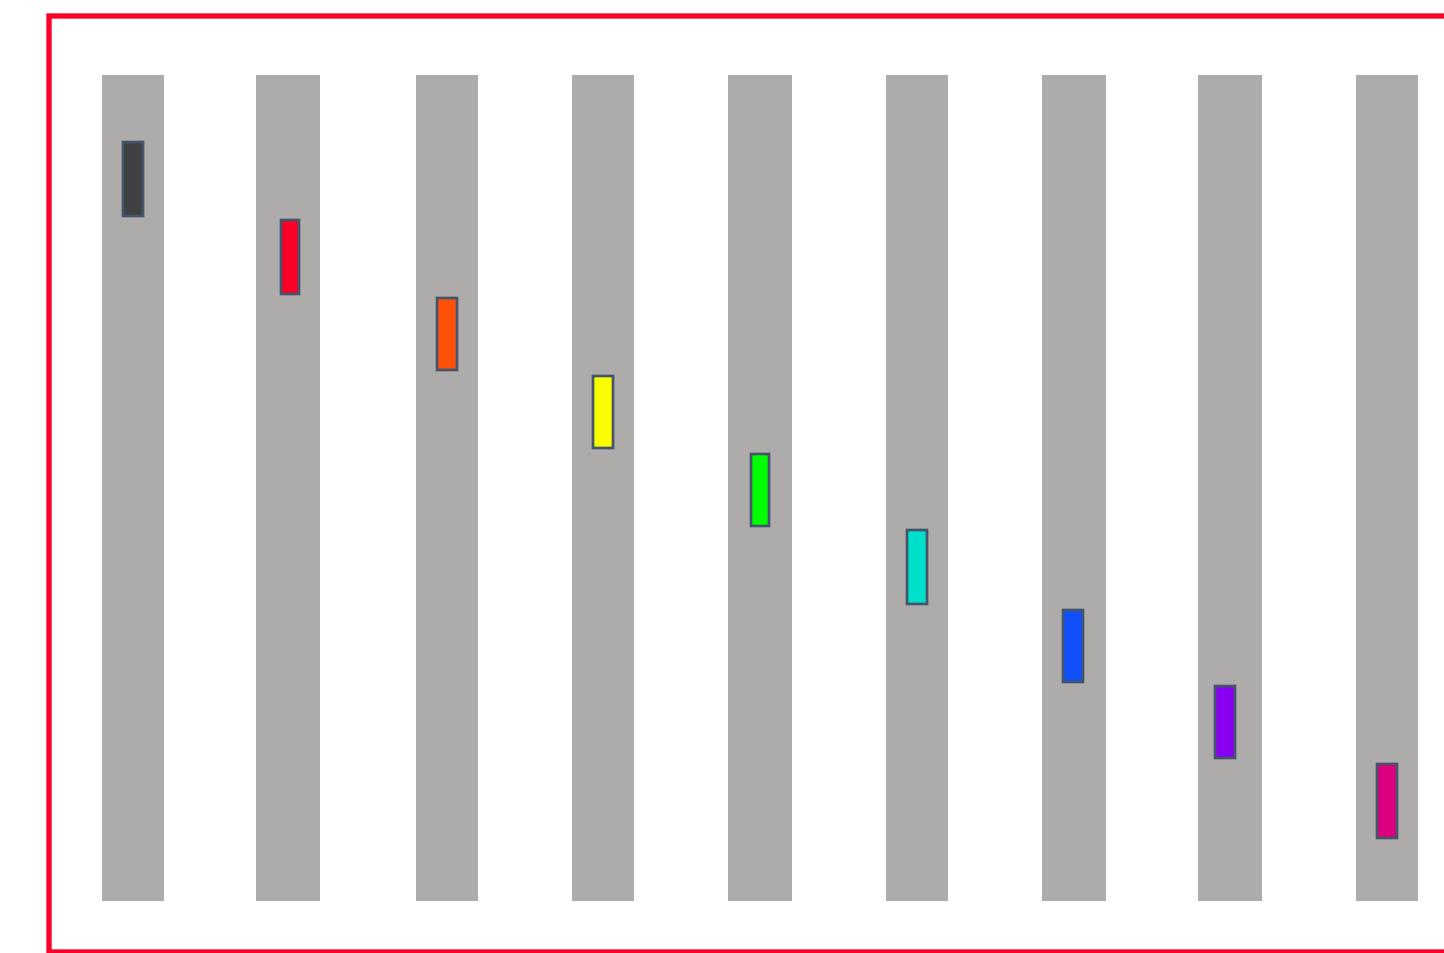
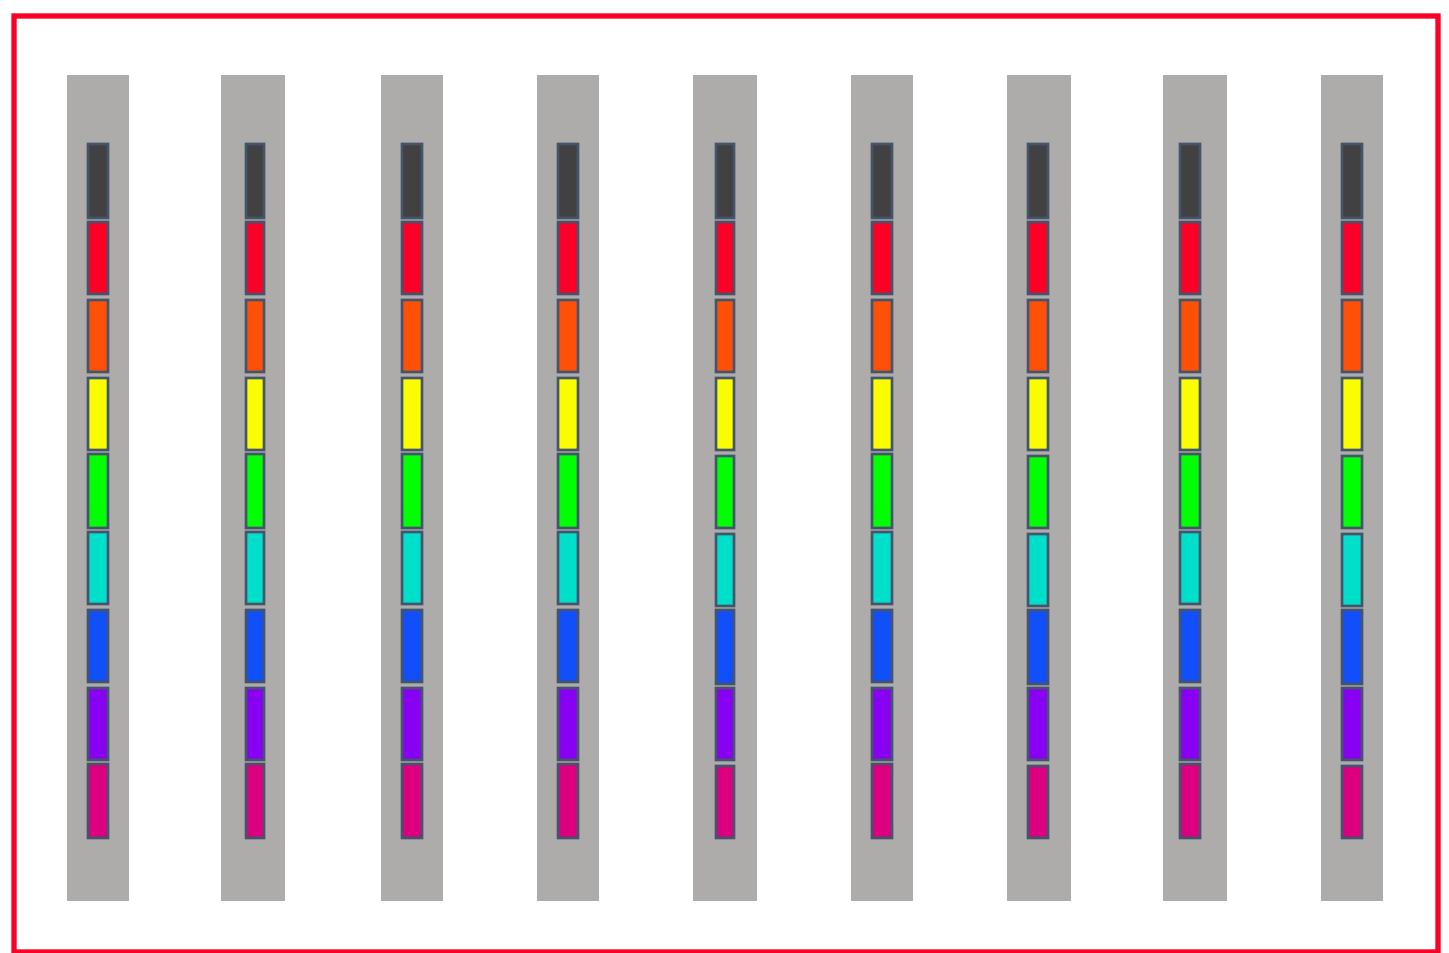
Reduce-scatter (small message)



Reduce-scatter (short vector)



Reduce-scatter (short vector)



Scatter

Cost of Reduce(-to-one)/scatter Reduce-scatter

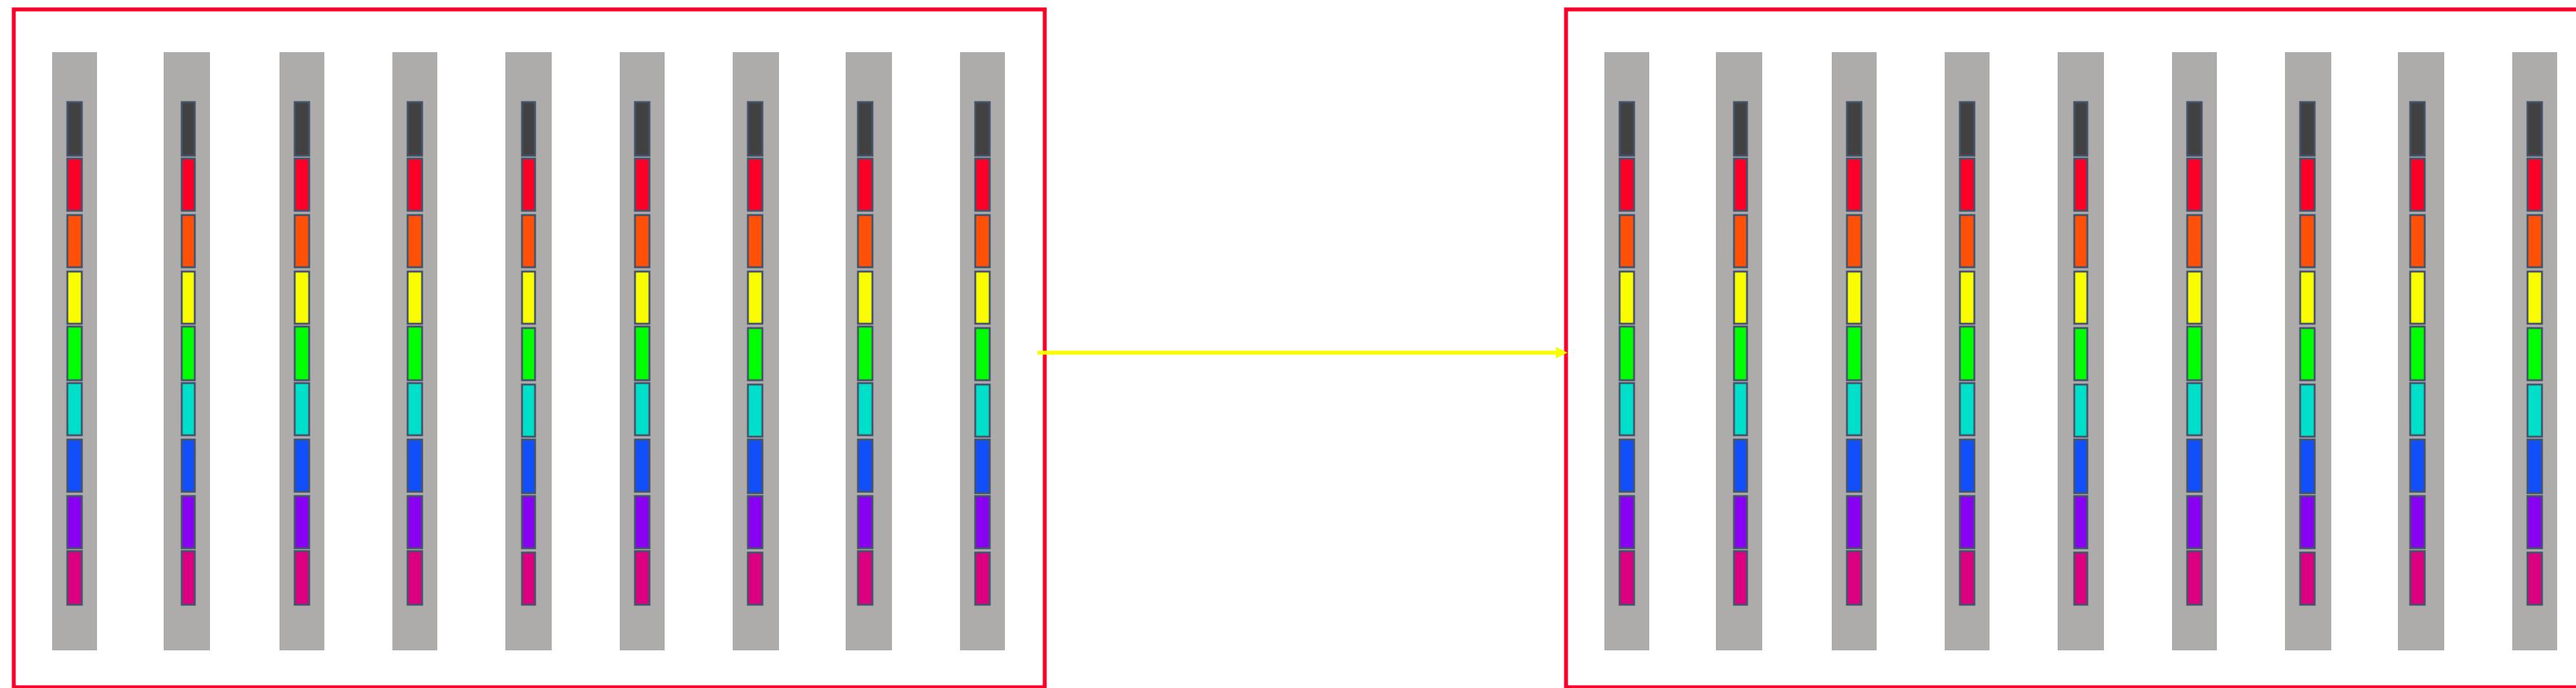
- Assumption: power of two number of nodes

$$\text{Reduce(-to-one)} \quad \log(p)(\alpha + n\beta + n\gamma)$$

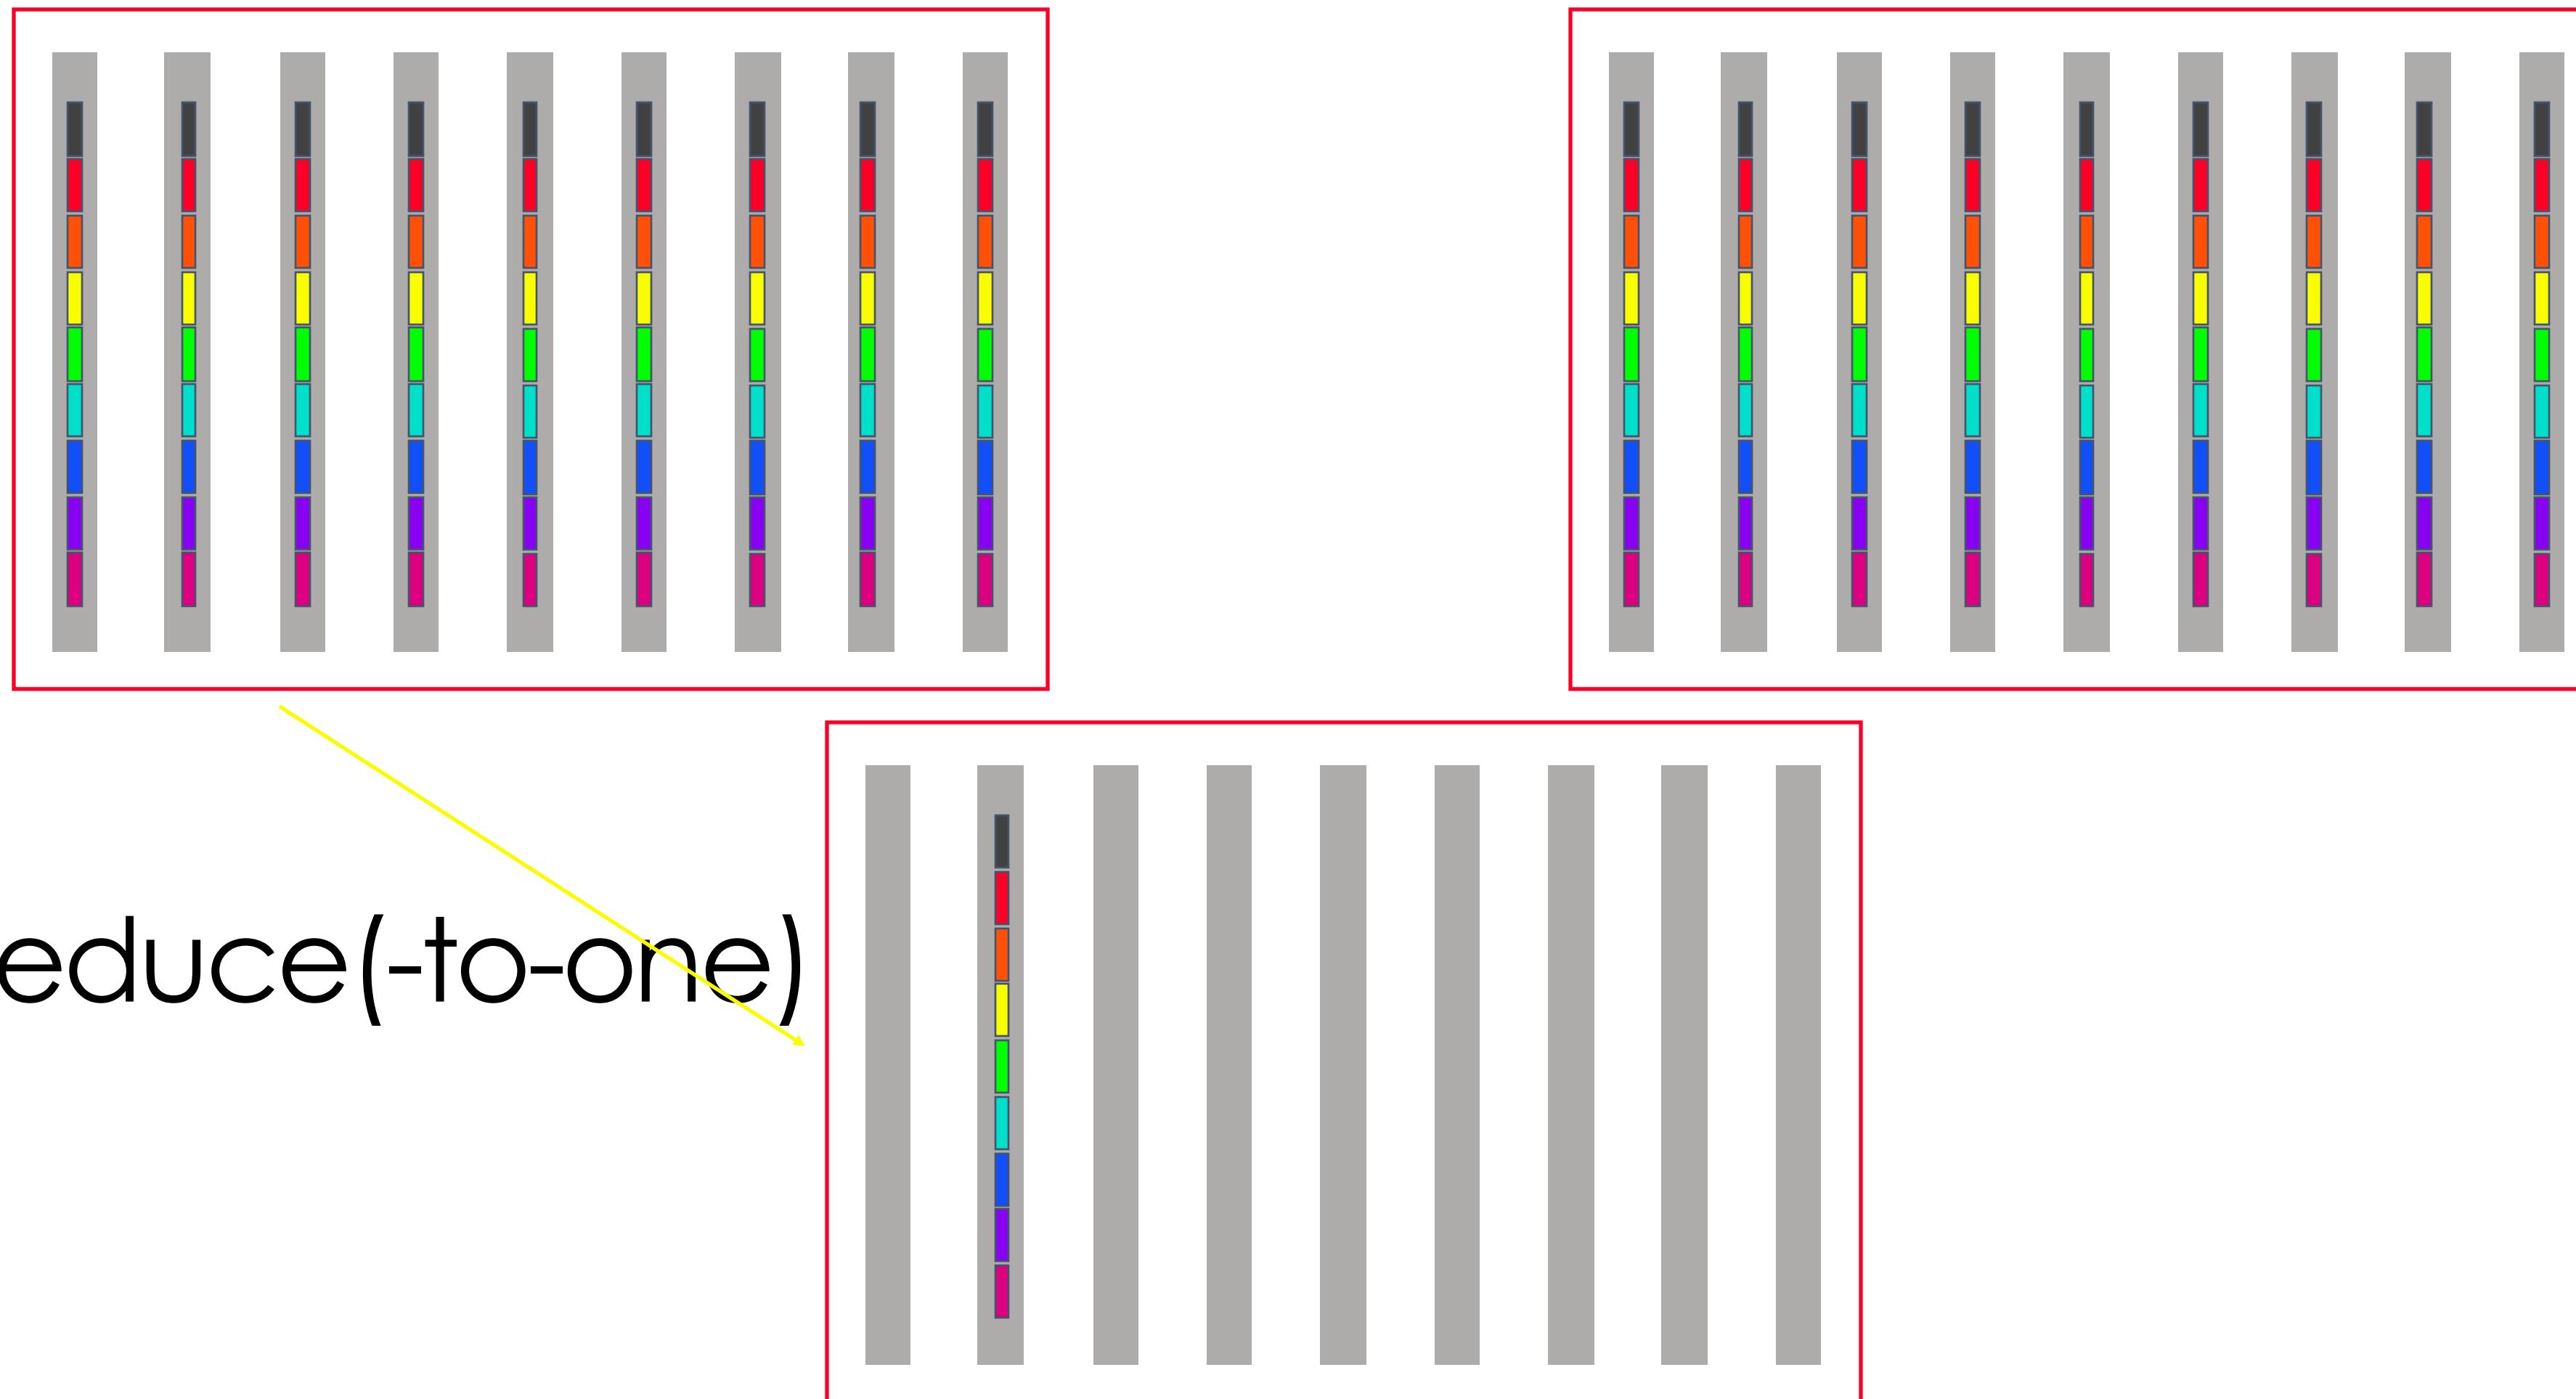
$$\text{scatter} \quad \log(p)\alpha + \frac{p-1}{p}n\beta$$

$$2\log(p)\alpha + \left(\frac{p-1}{p} + \log(p) \right) n\beta + \log(p)n\gamma$$

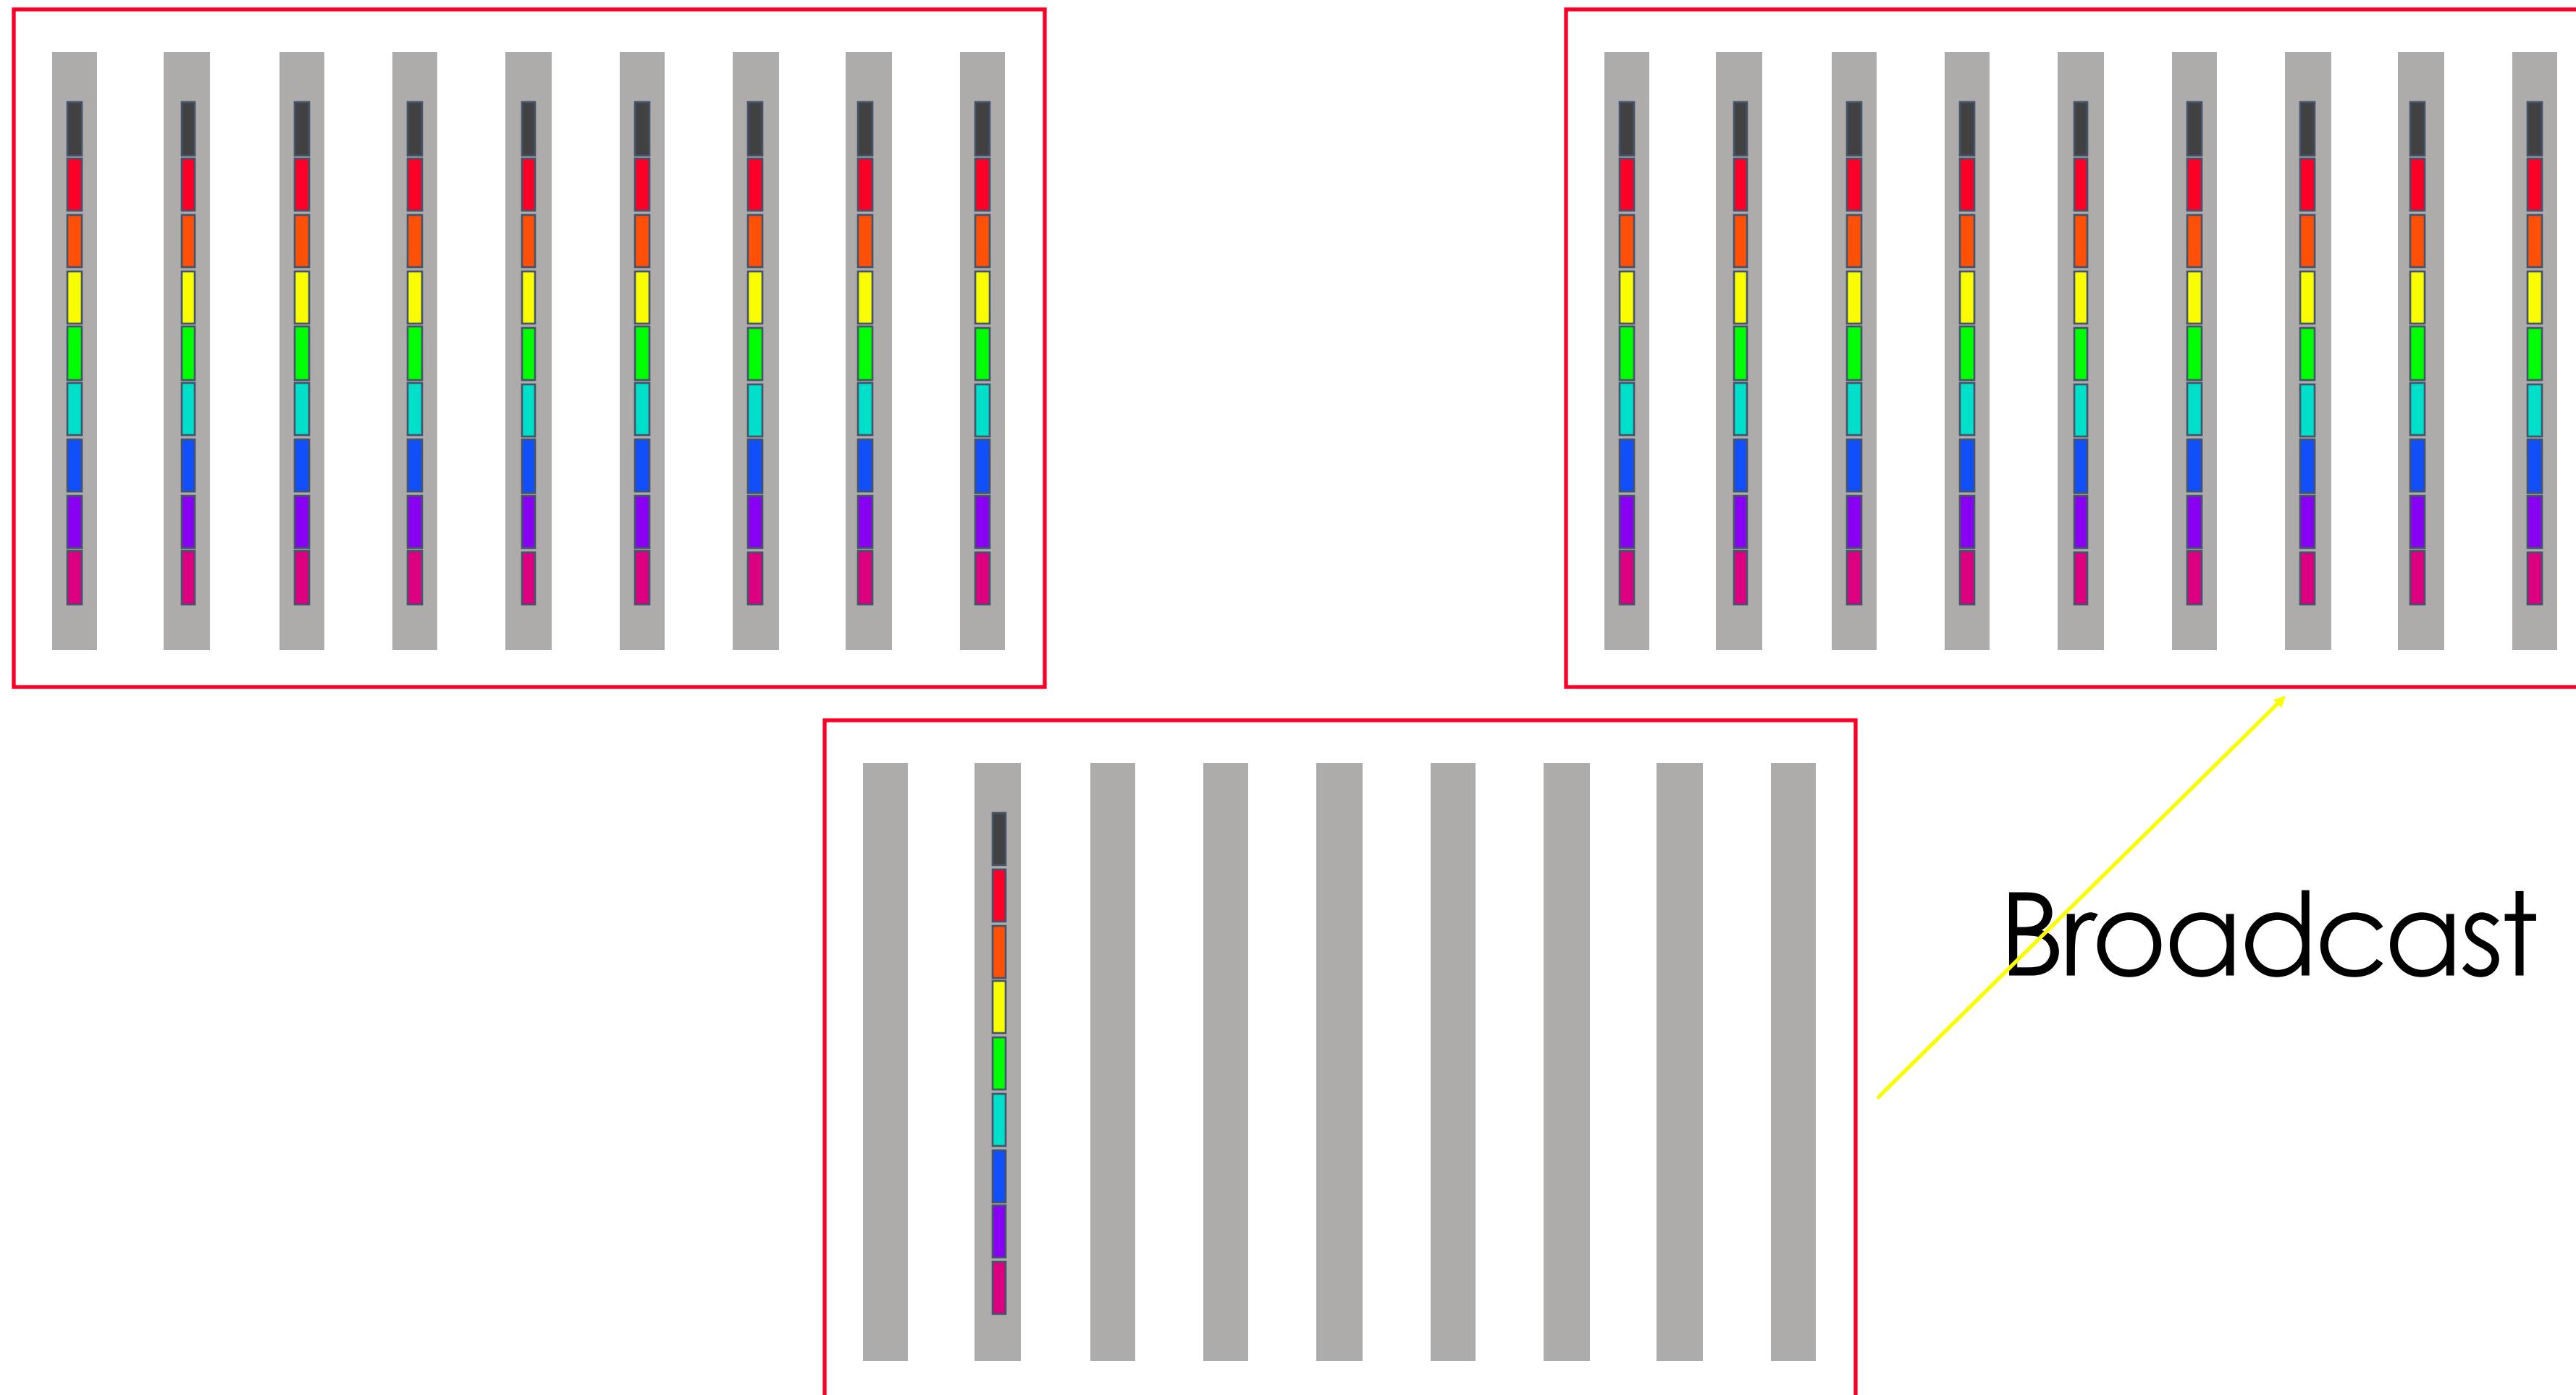
Allreduce (Latency-optimized)



Allreduce (Latency-optimized)



Allreduce (short vector)



Cost of reduce(-to-one)/broadcast Allreduce

- Assumption: power of two number of nodes

Reduce(-to-one) $\log(p)(\alpha + n\beta + n\gamma)$

broadcast $\log(p)(\alpha + n\beta)$

$$2\log(p)\alpha + 2\log(p)n\beta + \log(p)n\gamma$$

Recap

Reduce(-to-one)

$$\log(p)(\alpha + n\beta + n\gamma)$$

Scatter

$$\log(p)\alpha + \frac{p-1}{p}n\beta$$

Gather

$$\log(p)\alpha + \frac{p-1}{p}n\beta$$

Broadcast

$$\log(p)(\alpha + n\beta)$$

Reduce-scatter

Allreduce

Allgather

Recap

Reduce(-to-one)

$$\log(p)(\alpha + n\beta + n\gamma)$$

Scatter

$$\log(p)\alpha + \frac{p-1}{p}n\beta$$

Gather

$$\log(p)\alpha + \frac{p-1}{p}n\beta$$

Broadcast

$$\log(p)(\alpha + n\beta)$$

Reduce-scatter

$$2\log(p)\alpha + \log(p)n(\beta + \gamma) + \frac{p-1}{p}n\beta$$

Allreduce

Allgather

Recap

Reduce(-to-one)

$$\log(p)(\alpha + n\beta + n\gamma)$$

Scatter

$$\log(p)\alpha + \frac{p-1}{p}n\beta$$

Gather

$$\log(p)\alpha + \frac{p-1}{p}n\beta$$

Broadcast

$$\log(p)(\alpha + n\beta)$$

Reduce-scatter

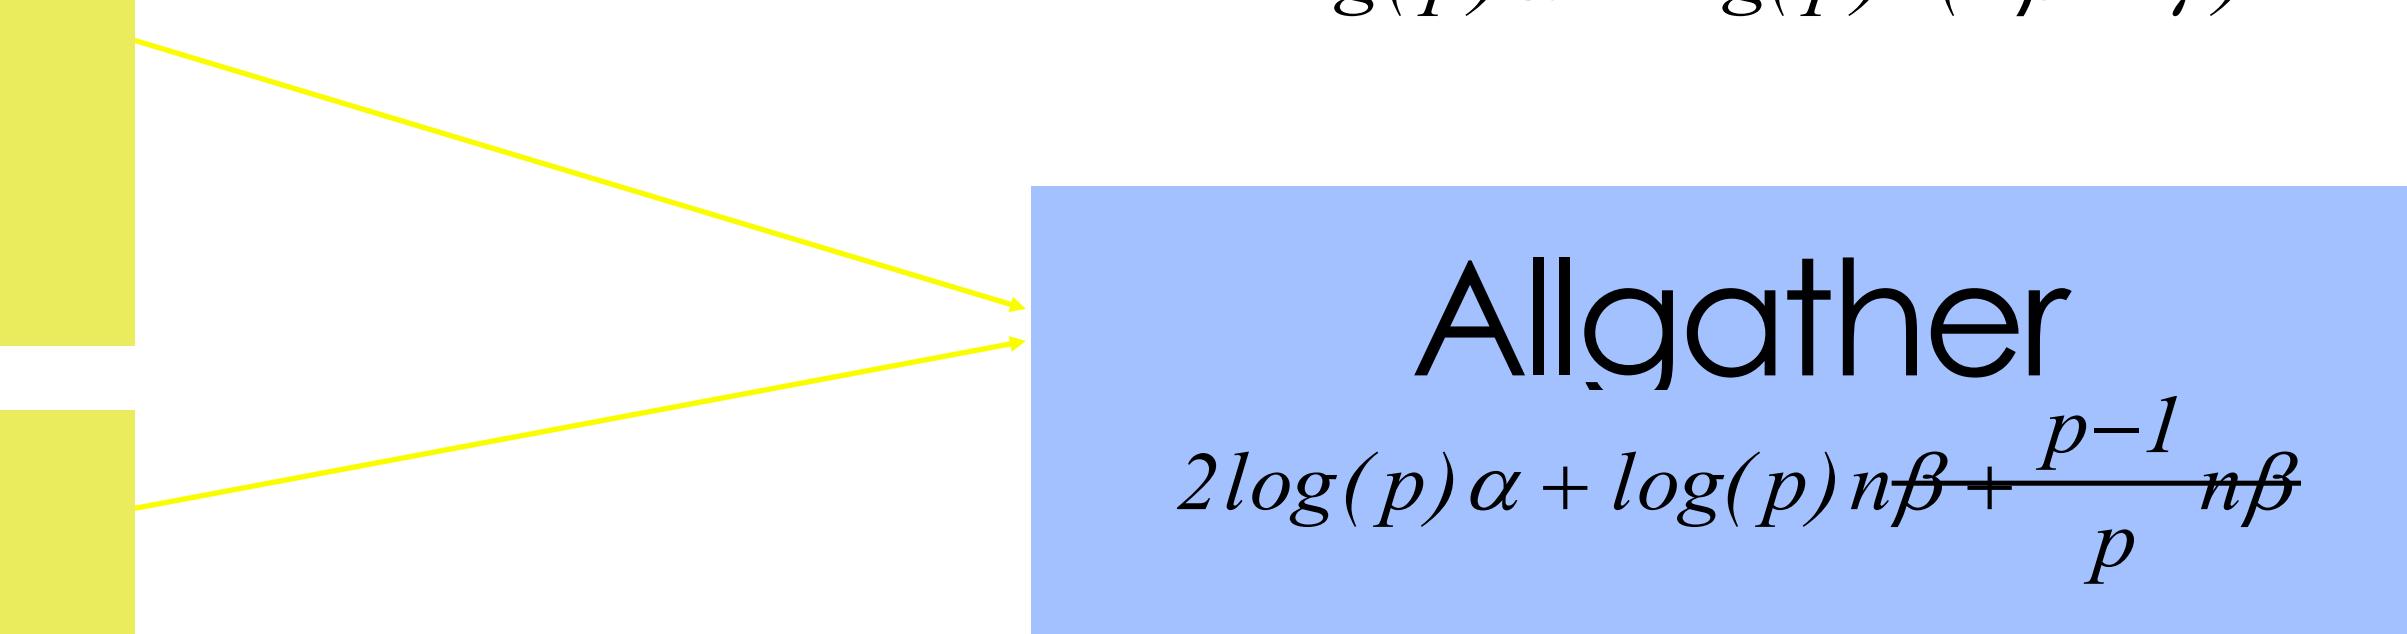
$$2\log(p)\alpha + \log(p)n(\beta + \gamma) + \frac{p-1}{p}n\beta$$

Allreduce

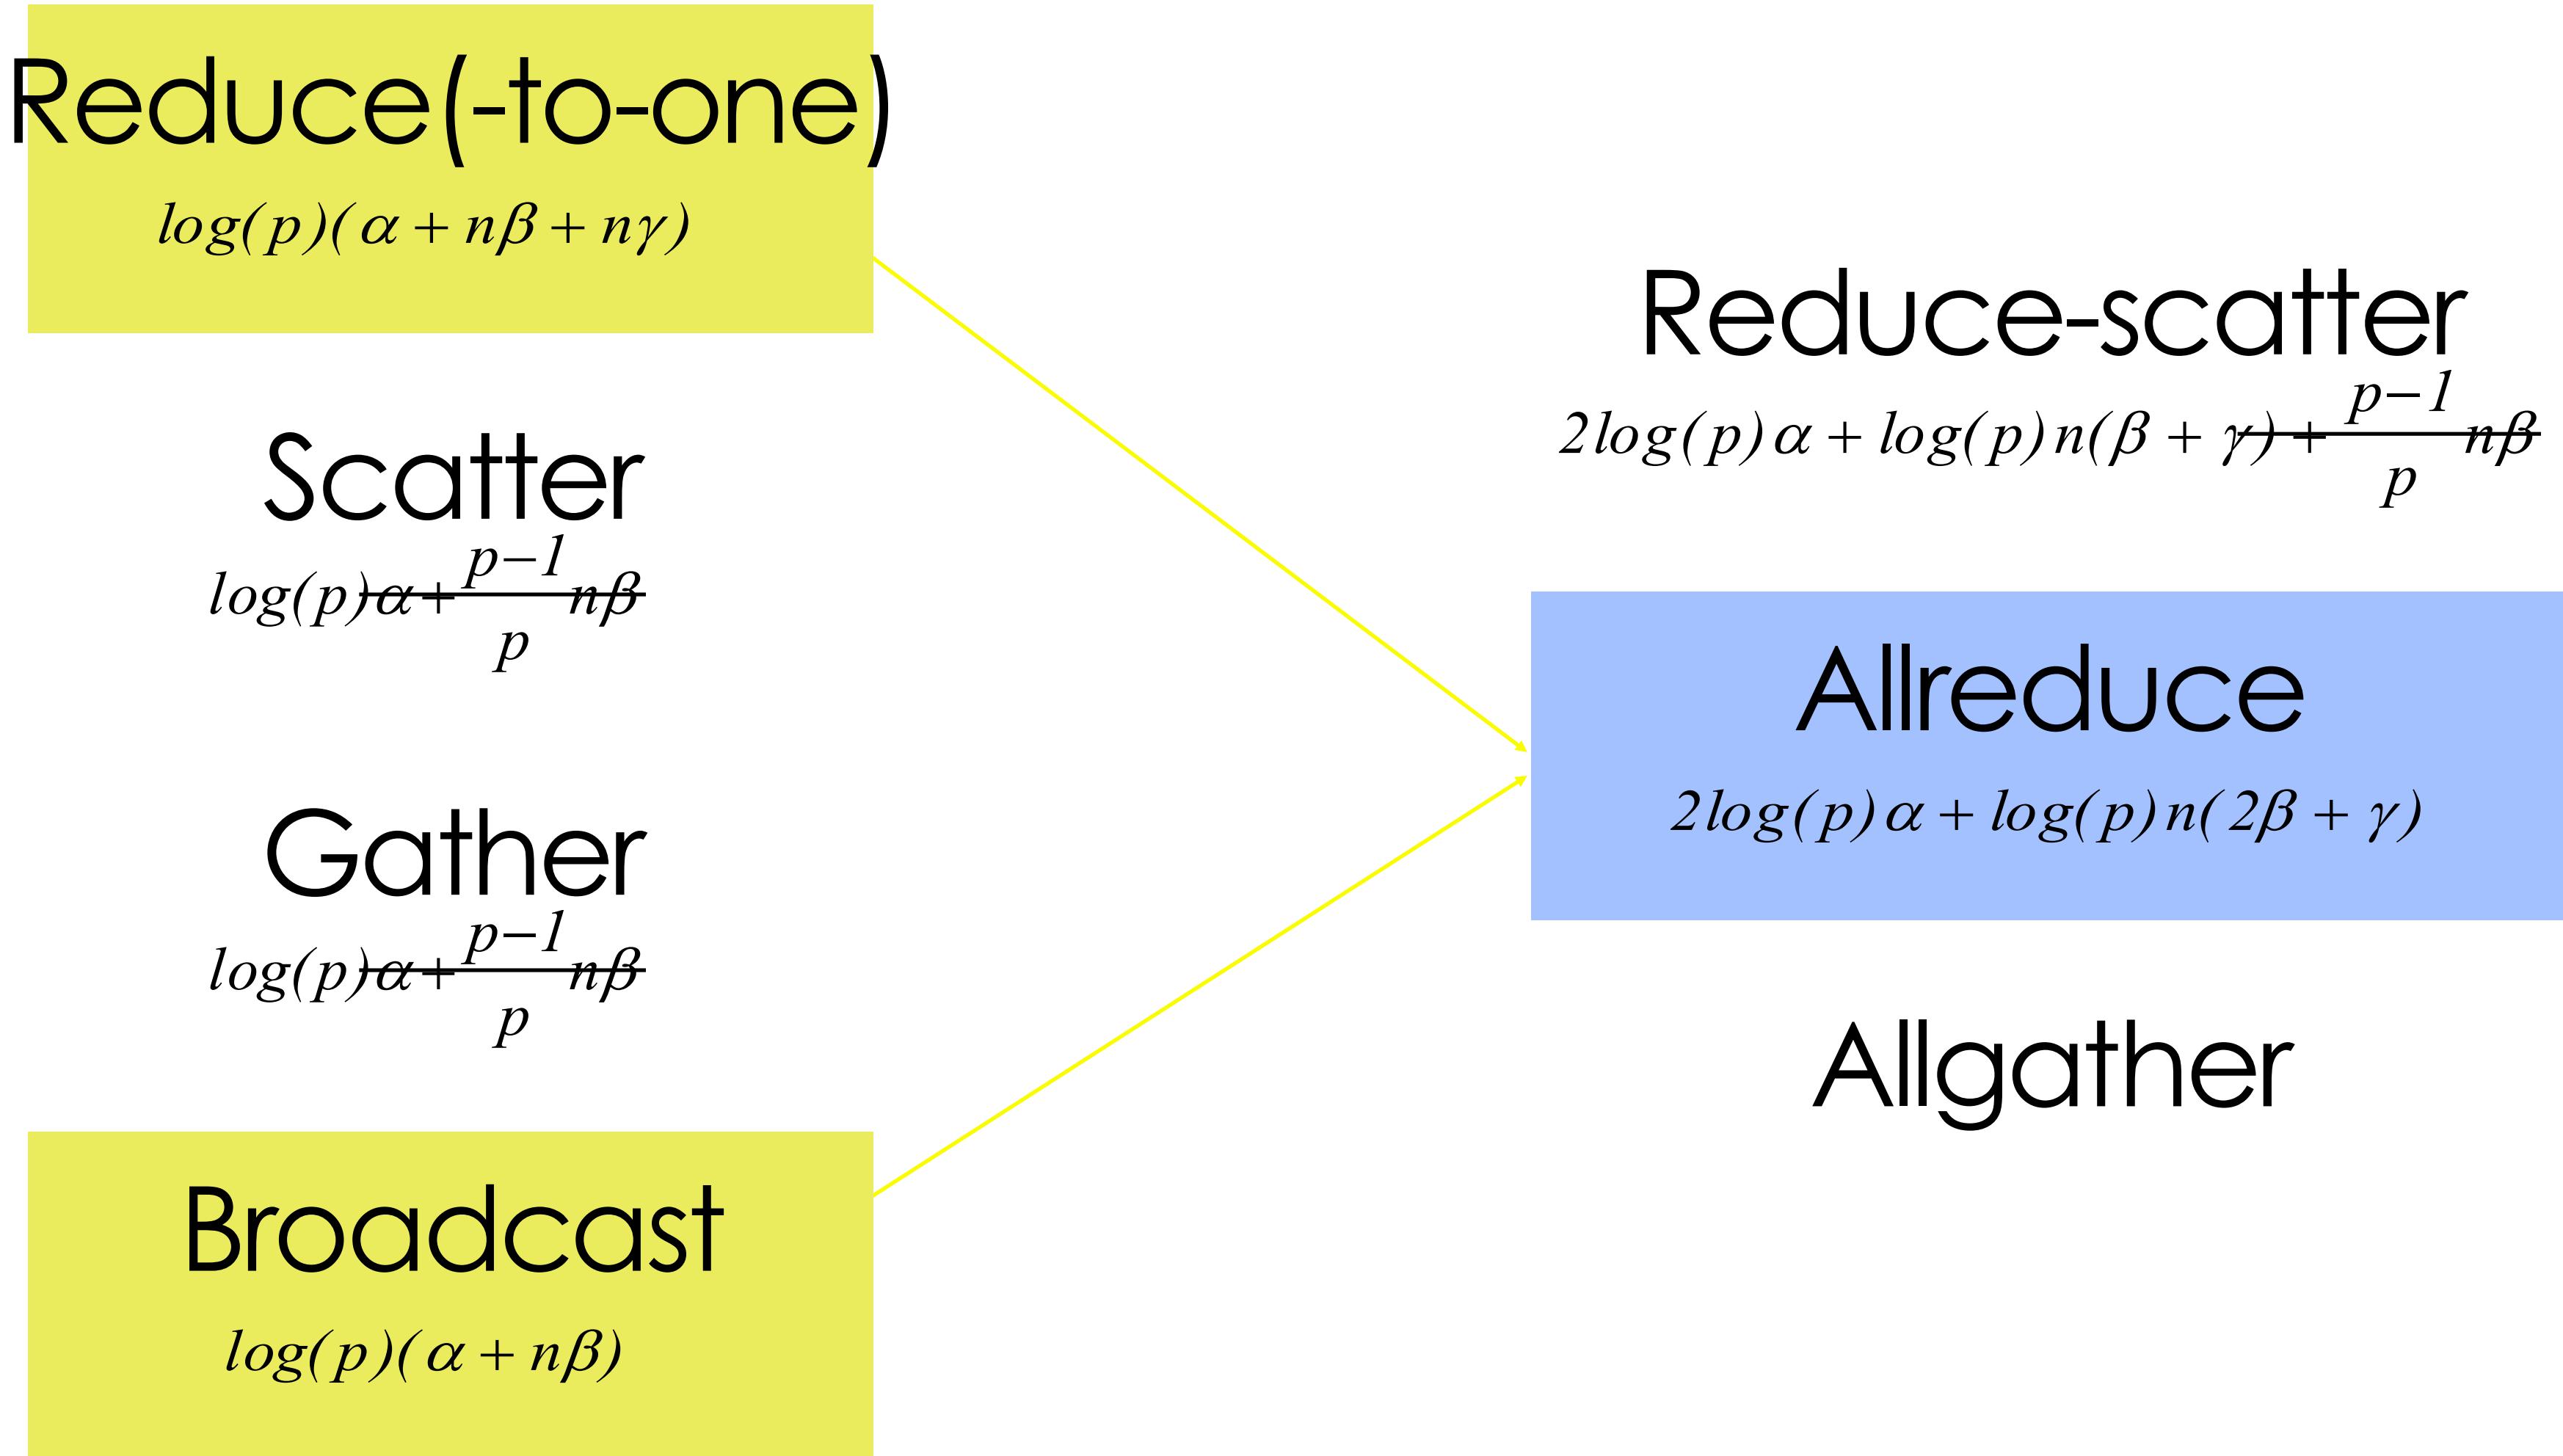
$$2\log(p)\alpha + \log(p)n(2\beta + \gamma)$$

Allgather

$$2\log(p)\alpha + \log(p)n\beta + \frac{p-1}{p}n\beta$$



Recap



Recap

Reduce(-to-one)

$$\log(p)(\alpha + n\beta + n\gamma)$$

Scatter

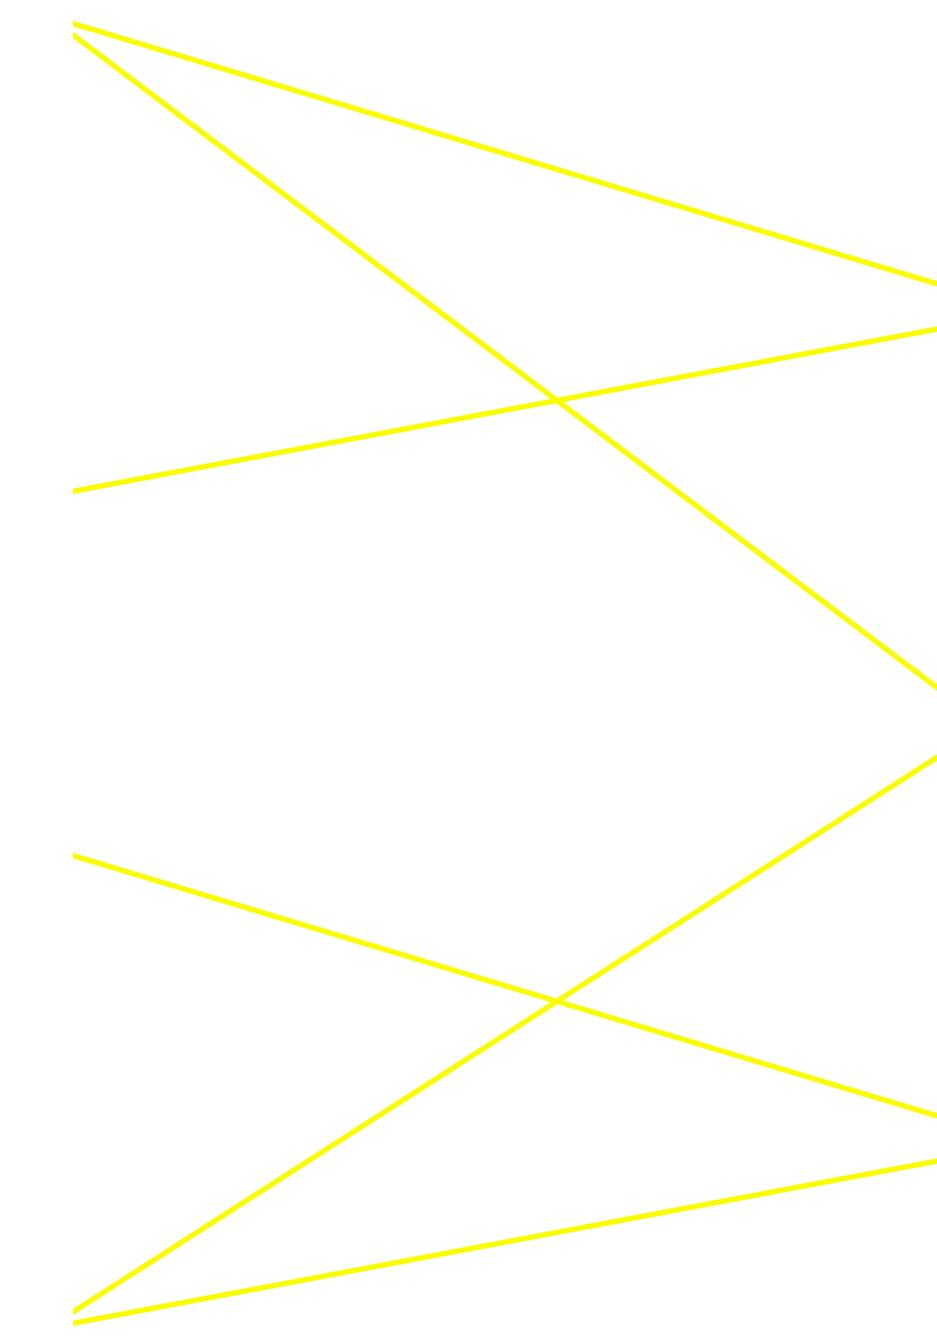
$$\log(p)\alpha + \frac{p-1}{p}n\beta$$

Gather

$$\log(p)\alpha + \frac{p-1}{p}n\beta$$

Broadcast

$$\log(p)(\alpha + n\beta)$$



Reduce-scatter

$$2\log(p)\alpha + \log(p)n(\beta + \gamma) + \frac{p-1}{p}n\beta$$

Allreduce

$$2\log(p)\alpha + \log(p)n(2\beta + \gamma)$$

Allgather

$$2\log(p)\alpha + \log(p)n\beta + \frac{p-1}{p}n\beta$$

Summary of MST algorithms

- Small message: Minimum Spanning Tree algorithm
 - Emphasize **low latency**
- **Can we do better?**
- Problem of Minimum Spanning Tree Algorithm?
 - It prioritize latency rather than bandwidth
 - Hence: Some links are idle
- Next: Large message size algorithm

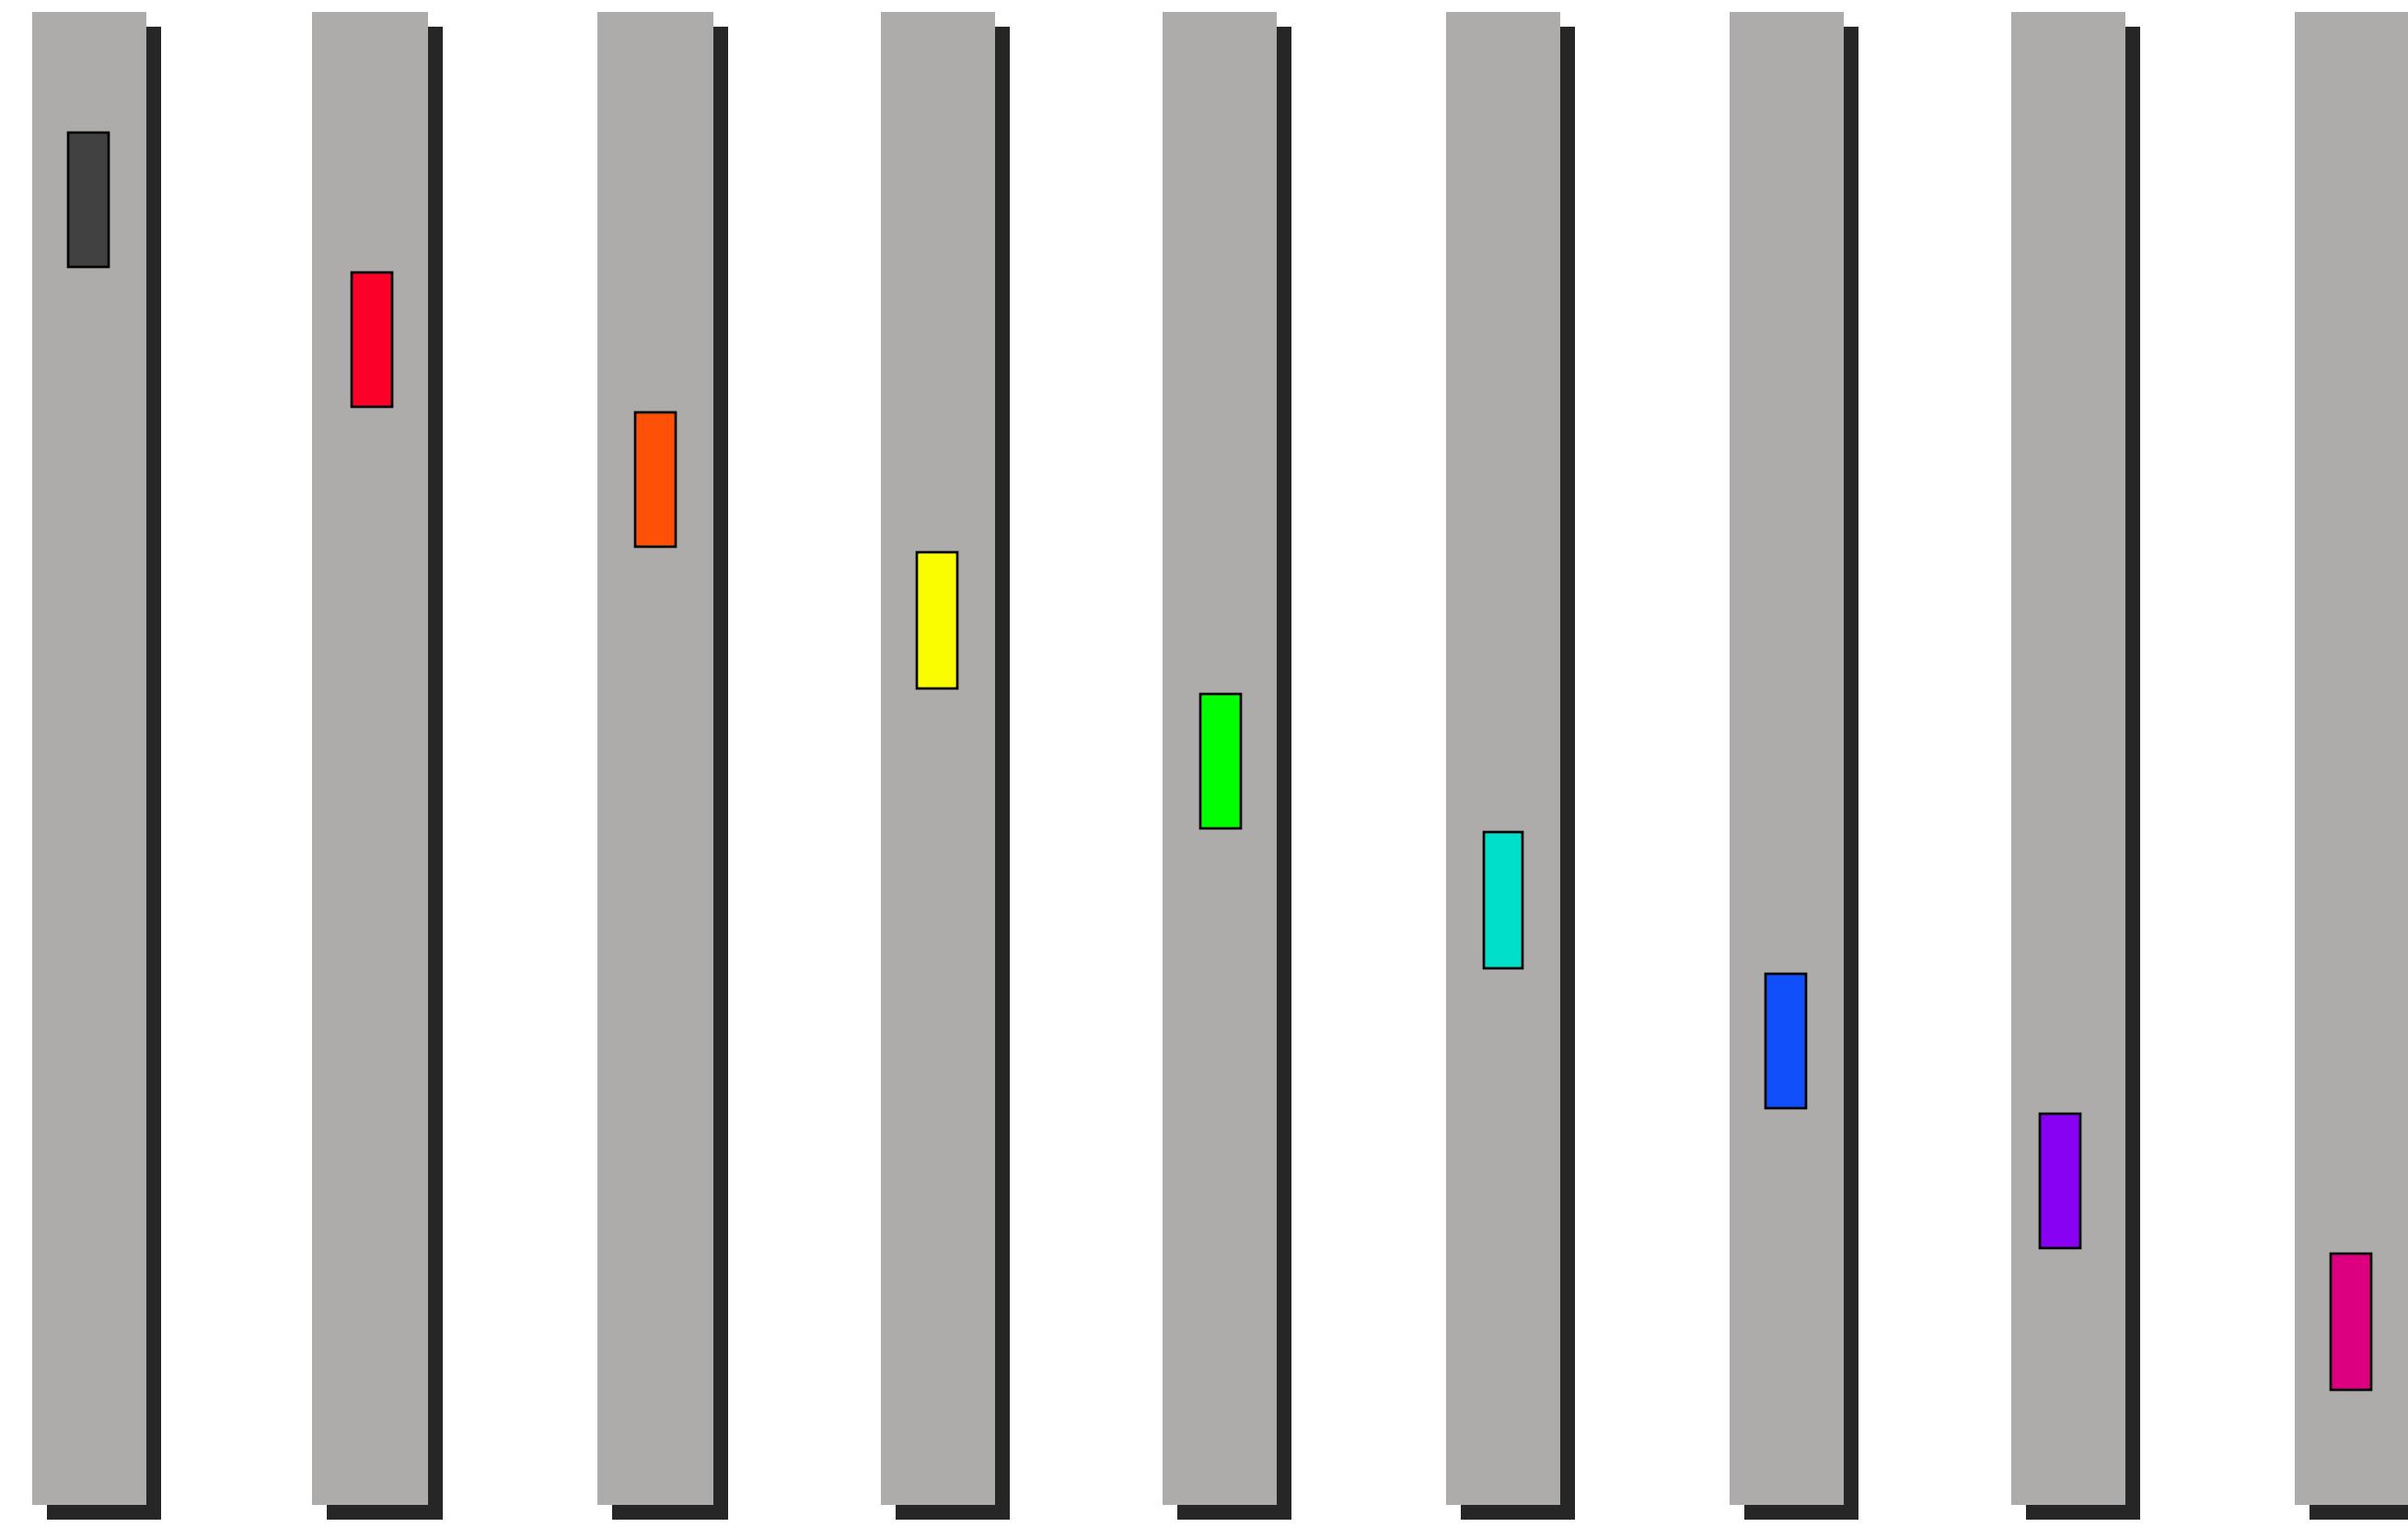
Large Message

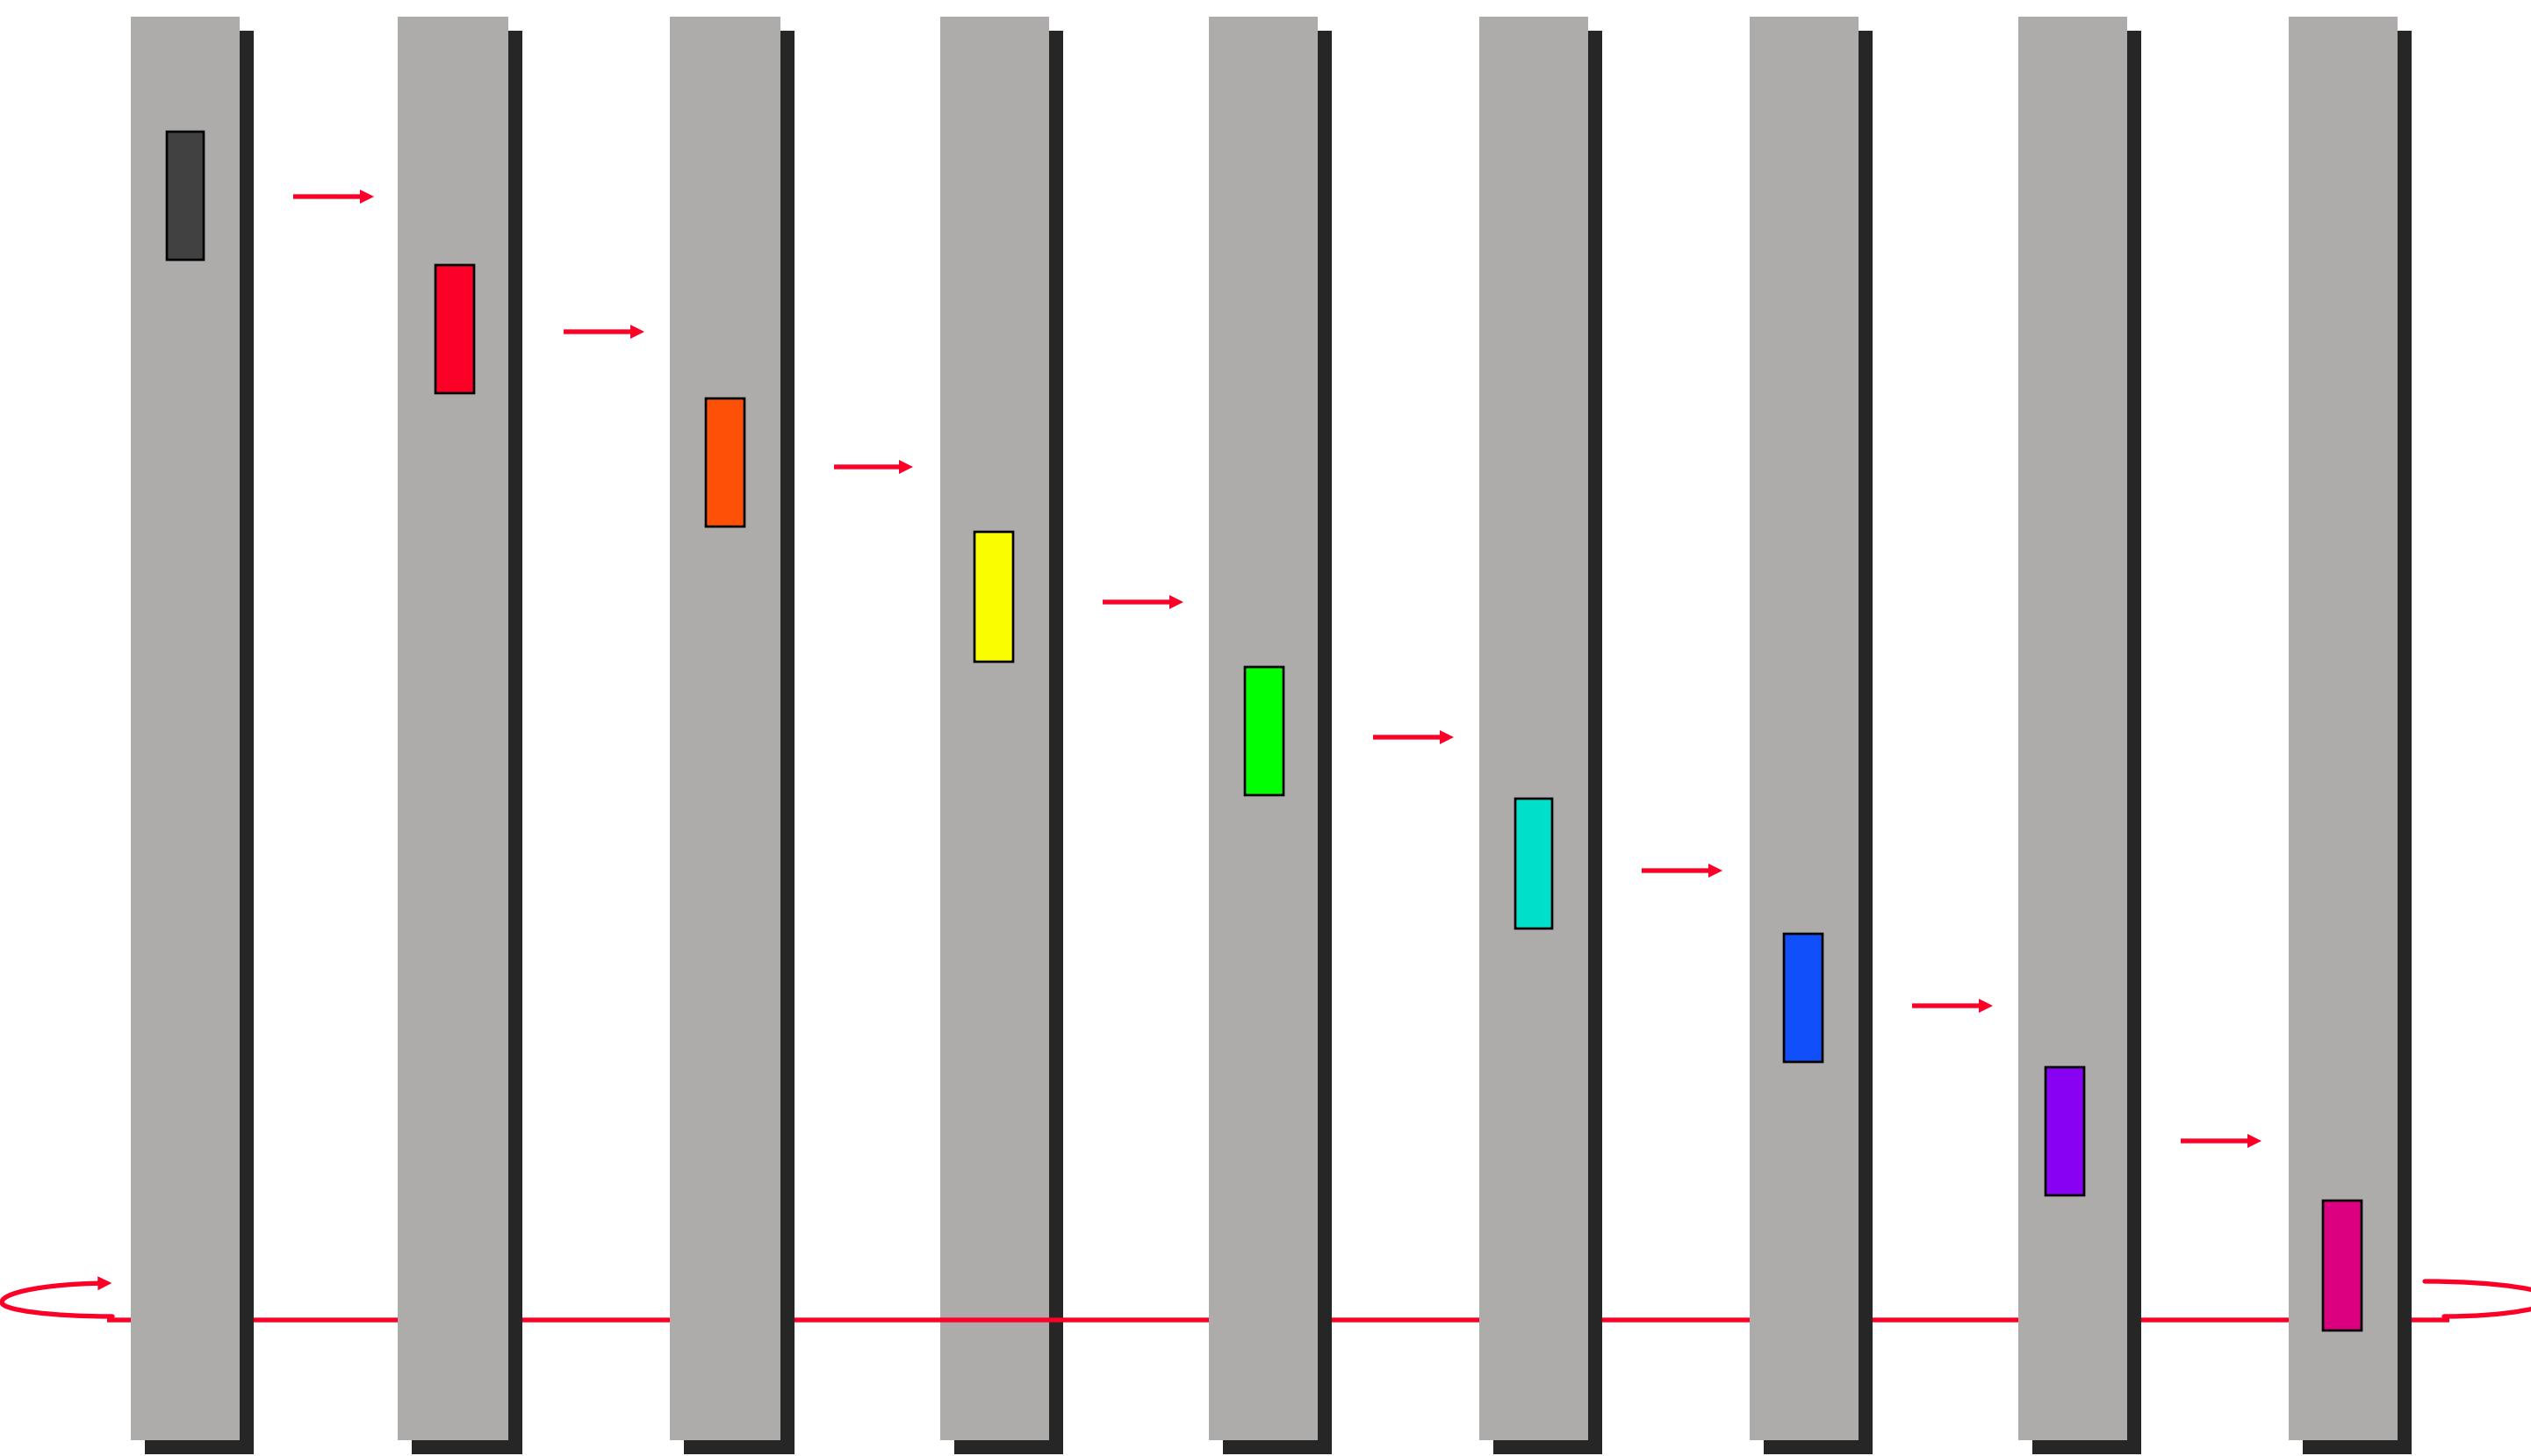
Communication Model: $\alpha + n\beta, \beta = \frac{1}{B}$

- The second term dominates – we want to minimize the second term
- We want to utilize the bandwidth as much as possible

General principles

- Use all the links between every two nodes
- A logical ring can be embedded in a physical linear array with worm-hole routing, since the “wrap-around” message doesn’t conflict





- A logical ring can be embedded in a physical linear array with worm-hole routing, since the “wrap-around” message doesn’t conflict

